

TreeCAM: Decoupling Updates and Lookups in Packet Classification

Balajee Vamanan and T. N. Vijaykumar

School of Electrical & Computer Engineering

CoNEXT 2011



Presenter: Kyle Wang

2012-06-05

Packet Classification

- **Packet Classification:**

- Find highest-priority rule that matches a packet

- Packet classification is key for

- Security, traffic monitoring/analysis, QoS

- **Classifier:** a set of rules

Source IP	Destination IP	Source Port	Dest. Port	Protocol	Action
120...0/24	198...0/2	0:65535	11:17	0xFF/0xFF	Accept
138...1/0	174...0/8	50:10000	0:65535	0x06/0xFF	Deny

Packet classification prevalent in modern routers

Trends in Packet Classification

- Line rates increasing (40 Gbps now, 160 Gbps soon)
- Classifier size (number of rules) increasing
 - Custom rules for VPNs, QoS
- Rules are getting more dynamic too

 Larger classifiers at faster lookup & update rates

- Much work on lookups, but not on updates

Must perform well in lookups and updates at low power

Characteristics of updates

- Two flavors
 - Virtual interfaces: add/remove 10,000s rules per minute
 - QoS: update 10s of rules per flow (milliseconds)
 - For either flavor, update rate (1/ms) \ll packet rate (1/ns)
- Current approaches
 - Either incur high update effort despite low update rates
 - Eat up memory bandwidth
 - Hold-up memory for long
 - \rightarrow packet drops, buffering complexity, missed deadlines
 - Or do not address updates
- Recent OpenFlow, online classification \rightarrow faster updates

Updates remain a key problem

Current Approaches

■ TCAM

- Unoptimized TCAMs search all rules per lookup → high power
- Modern partitioned TCAMs prune search → reduce power
 - Extended TCAMs [ICNP 2003]
- Physically order rules per priority for fast highest-priority match
 - This ordering **fundamentally** affects update effort
- Updates move **many** rules to maintain order
 - E.g., updates 10 per ms; lookups 1 per 10 ns; 100,000 rules
 - If 10 % updates move (read+write) 10 % rules
 - Updates need 20,000 ops/ms = 0.2 op per 10 ns
 - → 20% bandwidth overhead

High-effort updates in TCAM degrade throughput & latency

Current Approaches (Contd.)


- Decision Trees:
 - Build decision trees to prune search per lookup
 - Do not address updates
- No ordering like TCAMs but updates may cause tree imbalance
 - Imbalance increase lookup accesses
 - Re-balancing is costly

Previous schemes are not good in both lookups and updates

Our Contributions


- **TreeCAM**: Three *novel ideas*
 - *Dual tree versions* to decouple lookups and updates
 - *coarse tree* in TCAM → reduce lookup accesses
 - Tree/TCAM hybrid
 - *fine tree* in control memory → reduce update effort
 - *Interleaved layout* of leaves to cut ordering effort
 - *Path-by-path updates* to avoid hold-up of memory
 - Allow non-atomic updates interspersed with lookups
- Performs well in lookups and updates
 - 6-8 TCAM accesses for lookups
 - Close to **ideal** TCAM for updates

Outline


- Introduction
-  Background: Decision Trees
- Dual tree versions
 - Coarse Tree
 - Fine Tree
- Updates using Interleaved Layout
- Path-by-path updates
- Results
- Conclusion

Background: Decision Trees

- Rules are hypercubes in rule space
- Builds decision tree by cutting rule space to separate rules into smaller subspaces (child nodes)
- Stop when a small number of rules at a leaf called **binth** (e.g., 16)
- Packets traverse tree during classification
- Many heuristics
 - Dimension, number of cuts




Outline

- Introduction
- Background: Decision Trees
-  Dual tree versions
 - Coarse Tree
 - Fine Tree
- Updates using Interleaved Layout
- Path-by-path updates
- Results
- Conclusion

TreeCAM Coarse Tree (Version#1 for lookups)

- **Idea**: partition rules among TCAM subarrays using decision trees
- 4k-entry subarrays → coarse tree with each leaf in a subarray
 - 2-deep tree → fast lookup
- Packets traverse subarrays
- Previous heuristics complicated
- We propose simple sort heuristic
 - Sort rules & cut equally
 - *EffiCuts* → min. rule replication


Coarse Trees → Search Pruning (trees) + Parallel Search (TCAM)




TreeCAM Fine Tree (Version#2 for updates)

- **Key observation:** A packet cannot match multiple leaves → **only** rules within the **same leaf** need ordering
- Reduce update effort → Tree with small binth – *fine tree*
 - Not real, just annotations
 - One *coarse-tree* leaf contains some contiguous fine-tree leaves
 - Both trees kept consistent
- Updates slow → store *fine tree* in control memory

Fine Trees: Small binth, reduced ordering effort in TCAM




Outline

- Introduction
- Background: Decision Trees
- Dual tree versions
 - Coarse Tree
 - Fine Tree
-  Updates using Interleaved Layout
- Path-by-path updates
- Results
- Conclusion


Updates

- Add a rule
 - Create empty space at right priority level via repeated swaps
 - With naïve, contiguous layout of leaves, requires **many** repeated swaps
 - As many swaps as #rules
- Observation: **Only** overlapping rules need ordering per priority
 - Too hard in full generality




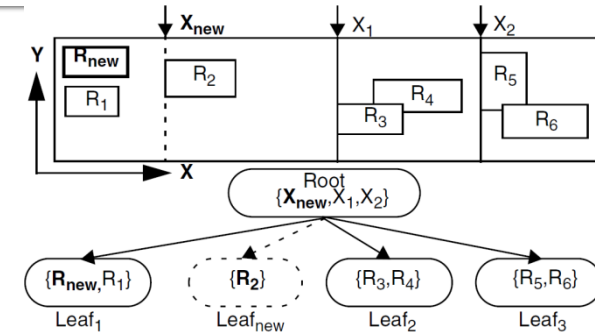
Interleaved Layout

- **Insight:** Leaves are naturally non-overlapping → **only** rules in a leaf need to be ordered
 - Trivial unlike general overlap
- *Interleave* rules across all leaves
 - Contiguously place same priority-level rules from all leaves
 - Order only across priority levels
 - Move **at most one** rule per level
 - binth levels → small for fine tree
 - Update effort \approx binth swaps



Updates with interleaved layout (cont.)


- Updating for three cases
 - No Re-balancing
 - Empty entries exist in target leaf
 - Local Re-balancing
 - No room in target leaf, displace the borderline rules of target leaf to neighboring leaf, or create a new leaf
 - Global Re-balancing
 - No room in target leaf and subarray, displace rules recursively from one to the next, until empty entries exist in target subarray
- Worst case where rules move across TCAM subarrays
 - Interleaved layout effort $\approx 3 * \text{binth} * \text{\#subarrays}$ swaps
 - $\approx 3 * 8 * 25 \approx 600$ swaps (100K rules and 4K rules/subarray)
 - Contiguous layout effort $\approx \text{\#rules}$
 - $\approx 100\text{K}$ swaps (100K rules)



Path-by-path Updates

- Problem: Update moves hold up memory for long
- Make updates non-atomic
 - Packet lookups can be interspersed between updates
- Procedure
 - (1) place the to-be-displaced rule into the back-up TCAM;
 - (2) shrink the boundary of the source subarray;
 - (3) remove (invalidate) the rule from the source;
 - (4) add the rule to the destination subarray;
 - (5) expand the boundary of the destination subarray; and
 - (6) invalidate the back-up TCAM.
- The rule swap in (4) still causes atomic operations, but few
 - 3 (read, write, invalidate) versus 1250 (global re-balancing)


Outline

- Introduction
- Background: Decision Trees
- Dual tree versions
 - Coarse Tree
 - Fine Tree
- Updates using Interleaved Layout
- Path-by-path updates
-  **Results**
- Conclusion

Experimental Methodology

- Key metrics: Lookups & Updates
 - Lookup accesses: #accesses per packet match
 - Update effort: #accesses per one-rule addition
- Software simulator
- EffiCuts, TCAM, and TreeCAM
 - TCAM: Unoptimized & Partitioned TCAM (Extended TCAM)
 - 4K subarrays
 - Decision tree algorithm (EffiCuts)
 - Tree/TCAM hybrid (TreeCAM)
 - Coarse tree binth = 4K, Fine tree binth = 8
- ClassBench generated classifiers – ACL, FW and IPC

Accesses per Lookup



- EffiCuts require many more SRAM accesses than TCAMs
- Extended TCAM and TreeCAM require only at most 8 accesses even for 100,000 rule classifiers
 - Extended TCAM does not handle updates

Update Effort

- Compare TCAM-basic, TCAM-ideal and TreeCAM
 - EffiCuts and Extended TCAM do not discuss updates
- TCAM-basic: periodic empty entries
- TCAM-Ideal: Adapt existing work on longest prefix match for packet classification
 - Identify groups of overlapping rules, and ideally assume
 - Enough empty entries at the end of every group
 - Two groups of overlapping rules DO NOT merge (*ideal*)
- We generate worst case update stream for each scheme
 - Adds rules to the left-most leaf of the largest separable tree
 - Removes rules from the right-most leaf of the tree

Worst-case Update Effort (cont.)

Classifier Type	Classifier Size	TCAM-basic		TCAM-ideal		TreeCAM	
		Empty Slots	Max # TCAM Ops	Max. Overlaps	Max # TCAM Ops	#Sub-arrays	Max # TCAM Ops
ACL	10K	44K	30K	67	134	3	91
	100K	60K	276K	166	332	19	684
FW	10K	68K	64K	90	180	3	112
	100K	82K	567K	295	590	29	1069
IPC	10K	43K	22K	62	124	1	24
	100K	46K	236K	137	274	11	385

TreeCAM is close to Ideal and two orders of magnitude better than TCAM-basic

Conclusion

- Previous schemes do not perform well in both lookups and updates
- **TreeCAM** uses three techniques to address this challenge:
 - **Dual tree versions:** Decouples lookups and updates
 - *Coarse trees* for lookups and *fine trees* for updates
 - **Interleaved layout** bounds the update effort
 - **Path-by-path update** enables non-atomic updates which can be interspersed with packet lookups
- TreeCAM achieves 6 – 8 lookup accesses and close to ideal TCAM for updates, even for large classifiers (100K rules)

TreeCAM scales well with classifier size, line rate, and update rate