

Frenetic: A High-Level Language for OpenFlow Networks

**Nate Foster, Rob Harrison, Matthew L. Meola,
Michael J. Freedman, Jennifer Rexford, David Walker**

Present by Xiang Wang
Venus Team, NSLab
RIIT, Tsinghua Univ.

May. 08, 2012 @ NSLab Seminar

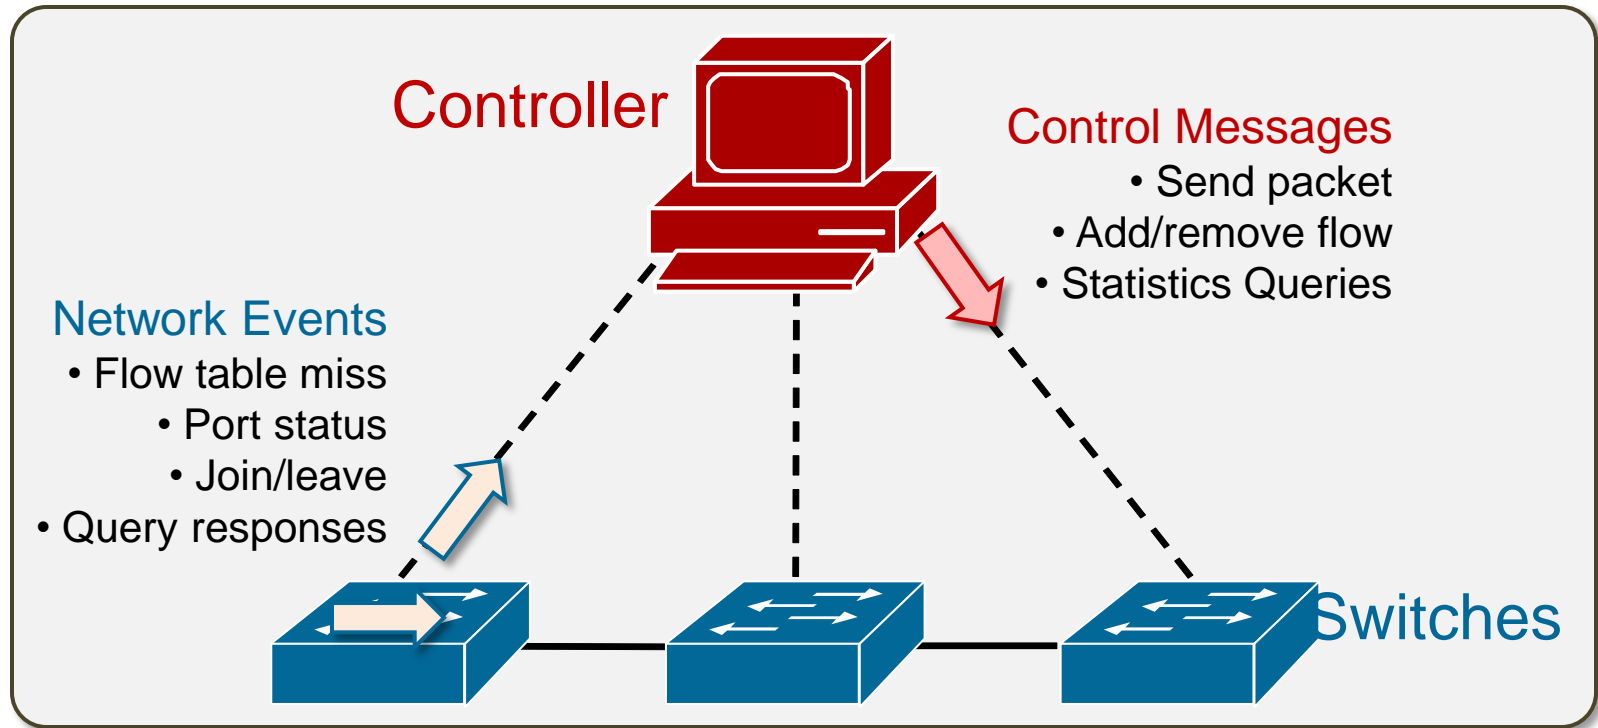


Background

- ❑ Chilling effect on innovation
 - ❑ Hardware and software are closed and proprietary
 - ❑ New requirements brought problems into sharp relief
- ❑ **OpenFlow/NOX** allowed us to take back the network
 - ❑ Direct access to dataplane hardware
 - ❑ Programmable control plane via open API
- ❑ OpenFlow/NOX made innovation possible, not easy
 - ❑ Low level interface mirrors hardware
 - ❑ Thin layer of abstraction
 - ❑ Few built-in features
- ❑ So let's give the network programmer some help...



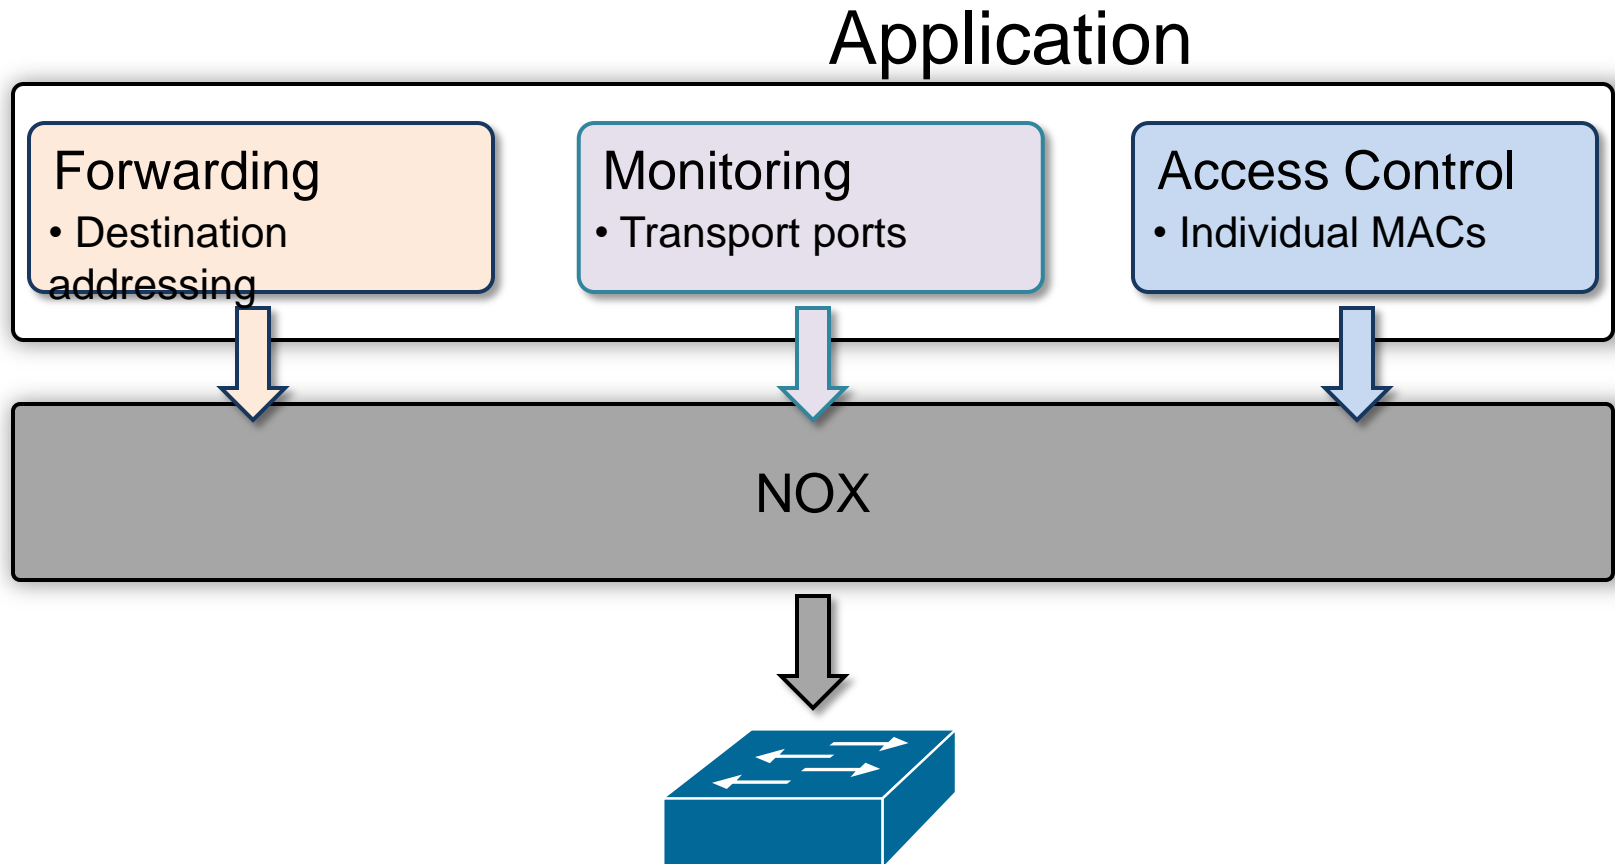
OpenFlow Architecture



OpenFlow Switch Flow Table

Priority	Pattern	Action	Counters
0-65535	Physical Port, Link Source/Destination/Type, VLAN, Network Source/Destination/Type, Transport Source/Destination	Forward Modify Drop	Bytes, Count

Programming Networks with NOX




In general, program modules do not *compose*

OpenFlow/NOX Difficulties

- ❑ Multiple task vs. independent modules
 - ❑ Routing, access control, traffic monitoring
 - ❑ Rules may be overlapping
- ❑ Low-level abstraction
 - ❑ No support for operations like union and intersection
 - ❑ Manual refactoring of rules to compose subprograms
- ❑ Split architecture
 - ❑ Between logic running on the switch and controller
- ❑ Asynchronous interactions
 - ❑ Between switch and controller
 - ❑ Various race conditions



Example




- ❑ Simple Network Repeater
 - ❑ Forward packets received on port 1 out 2; vice versa

Simple Repeater

NOX Program


```
def simple_repeater():  
    # Repeat Port 1 to Port 2  
    p1 = {IN_PORT:1}  
    a1 = [(OFPAT_OUTPUT, PORT_2)]  
    install(switch, p1, HIGH, a1)  
  
    # Repeat Port 2 to Port 1  
    p2 = {IN_PORT:2}  
    a2 = [(OFPAT_OUTPUT, PORT_1)]  
    install(switch, p2, HIGH, a2)
```



Flow Table

Priority	Pattern	Action	Counters
HIGH	IN_PORT:1	OUTPUT:2	(0,0)
HIGH	IN_PORT:2	OUTPUT:1	(0,0)

Interactions between modules



- ❑ Simple Network Repeater **with Host Monitoring**
 - ❑ Forward packets received on port 1 out 2; vice versa
 - ❑ Monitor incoming HTTP traffic totals per host

Interactions between modules

Repeat port 1 to 2

```
def port1_to_2():  
    p1 = {IN_PORT:1}  
    a1 = [(OFPAT_OUTPUT, PORT_2)]  
    install(switch, p1, HIGH, a1)
```

Callback to generate rules per host

```
def packet_in(switch, inport, pkt):  
    p    = {DL_DST:dstmac(pkt)}  
    pweb = {DL_DST:dstmac(pkt),  
           DL_TYPE:IP,NW_PROTO:TCP,  
           TP_SRC:80}  
    a = [(OFPAT_OUTPUT, PORT_1)]  
    install(switch, pweb, HIGH, a)  
    install(switch, p, MEDIUM, a)
```

```
def main():  
    register_callback(packet_in)  
    port1_to_2()
```

```
def simple_repeater():
```

Port 1 to port 2

```
p1 = {IN_PORT:1}  
a1 = [(OFPAT_OUTPUT, PORT_2)]  
install(switch, p1, HIGH, a1)
```

Port 2 to Port 1

```
p2 = {IN_PORT:2}  
a2 = [(OFPAT_OUTPUT, PORT_1)]  
install(switch, p2, HIGH, a2)
```

Priority	Pattern	Action	Counter s
HIGH	{IN_PORT:1}	OUTPUT:2	(0,0)
HIGH	{DL_DST:mac,DL_TYPE:IP_TYPE,NW_PROTO:TCP, TP_SRC:80}	OUTPUT:1	(0,0)
MEDIUM	{DL_DST:mac}	OUTPUT:1	(0,0)

Low-level programming interface

Switch-level API:

- Derived from switch feature, rather than being designed for ease of use
- Multiple rules are manually adjusted with restrict priority

```
def repeater_monitor(switch):  
    pat1 = {in_port:1}  
    pat2 = {in_port:2}  
    pat2web = {in_port:2,tp_src:80}  
    install(switch,pat1,[output(2)],DEFAULT)  
    install(switch,pat2web,[output(1)],HIGH)  
    install(switch,pat2,[output(1)],DEFAULT)  
    query_stats(switch,pat2web)
```


```
def repeater_monitor_noserver(switch):  
    pat1 = {in_port:1}  
    pat2 = {in_port:2}  
    pat2web = {in_port:2,tp_src:80}  
    pat2srv = {in_port:2,nw_dst:10.0.0.9,tp_src:80}  
    install(switch,pat1,DEFAULT,None,[output(2)])  
    install(switch,pat2srv,HIGH,None,[output(1)])  
    install(switch,pat2web,MEDIUM,None,[output(1)])  
    install(switch,pat2,DEFAULT,None,[output(1)])  
    query_stats(switch,pat2web)
```

Two-tiered Programming Model

Tricky problem:

- Controller activity is driven by packets sent from switches
- Efficient applications install rules on switches to forward packets in hardware

Controller




Constant questions:

- “Is that packet going to come to the controller to trigger my computation?”
- “Or is it already being handled invisibly on the switch?”

Network Race Conditions

A challenging chain of events:

- Switch
 - sends packet to controller
- Controller
 - analyzes packet
 - updates its state
 - initiates installation of new packet-processing rules
- Switch
 - hasn't received new rules
 - sends new packets to controller
- Controller
 - confused
 - packets in the same flow handled inconsistently



Problems – One Common Cause

- Problems:
 - **Non-modular programming:** Programs can't be divided into modules for monitoring and forwarding
 - **Network race conditions:** The controller sees more events (packets) than it anticipates
 - **Two-tiered programming:** Will the controller be able to see the appropriate events given the forward rules installed?
- One common cause:
 - No effective *abstractions* for *reading* network state



The Solution

- Separate network programming into two parts:
 - Abstractions for reading network state
 - Reads should have *no effect* on forwarding policy
 - Reads should be able to *see every packet*
 - Abstractions for specification of forwarding policy
 - Forwarding *policy* must be separated from *implementation mechanism*
- A natural decomposition that mirrors the two fundamental tasks of network management
 - Monitoring and forwarding



Key principles

- ❑ Declarative Design
 - ❑ Consider what the programmer might want
 - ❑ Rather than how the hardware implements it
- ❑ Modular Design
 - ❑ Limited network-wide effects primitives and semantics
 - ❑ Independently used in different contexts
- ❑ Single-tier Programming
 - ❑ Support see-every-packet abstraction
 - ❑ Side-step many complexities of two-tiered model
- ❑ Race-free Semantics
 - ❑ Auto race detection
 - ❑ Packet suppression
- ❑ Cost Control
 - ❑ Gives programmers guidance concerning the costs
 - ❑ Query language is carefully defined



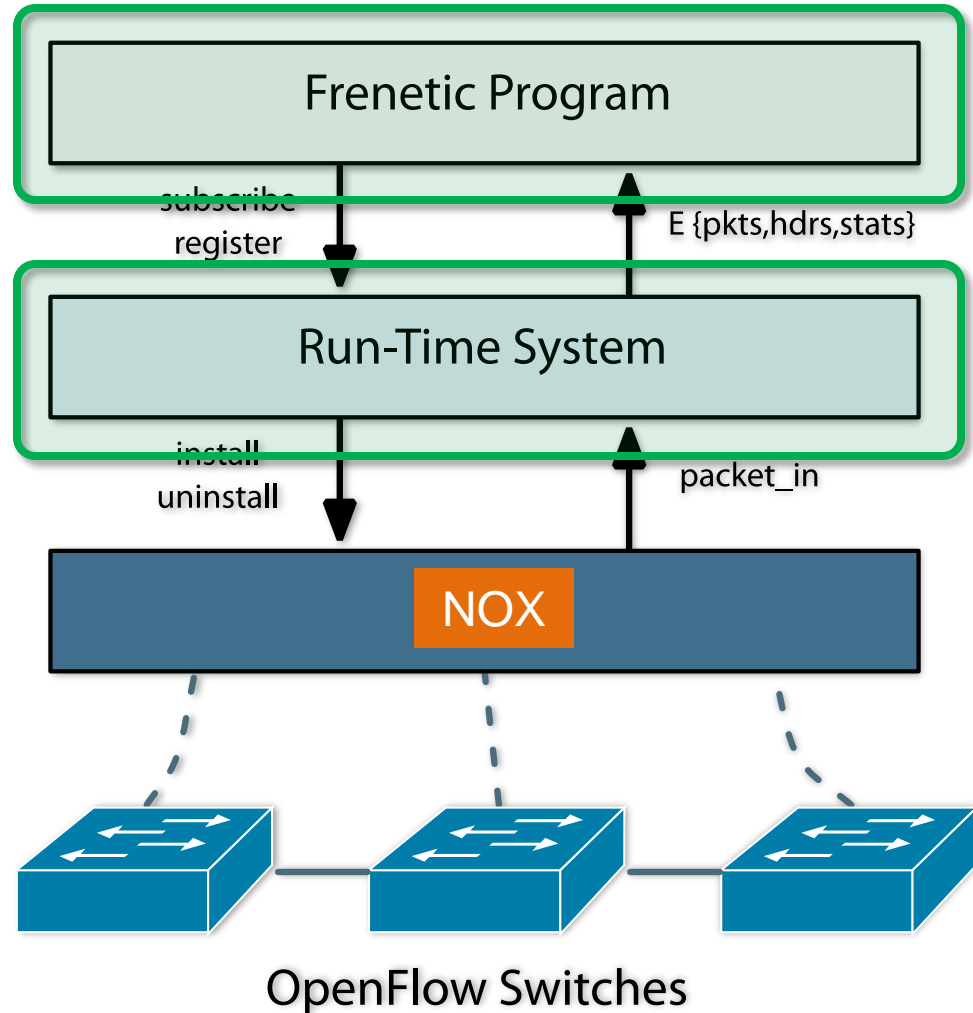
Frenetic

A High-level Language

- High-level patterns to describe flows
- Unified abstraction
- Composition

A Run-time System

- Handles module interactions
- Deals with asynchronous behavior



Frenetic Design


The Network Query Language

- Low-level rules on switches
- High-level abstraction for programmers
- Result is an event stream

```
Queries       $q ::= \text{Select}(a) * \text{Where}(fp) * \text{GroupBy}([qh_1, \dots, qh_n]) * \text{SplitWhen}([qh_1, \dots, qh_n]) * \text{Every}(n) * \text{Limit}(n)$   
  
Aggregates   $a ::= \text{packets} \mid \text{sizes} \mid \text{counts}$   
Headers      $qh ::= \text{inport} \mid \text{srcmac} \mid \text{dstmac} \mid \text{ethtype} \mid \text{vlan} \mid \text{srcip} \mid \text{dstip} \mid \text{protocol} \mid \text{srcport} \mid \text{dstport} \mid \text{switch}$   
  
Patterns     $fp ::= \text{true\_fp}() \mid qh\_fp(n) \mid \text{and\_fp}([fp_1, \dots, fp_n]) \mid \text{or\_fp}([fp_1, \dots, fp_n]) \mid \text{diff\_fp}(fp_1, fp_2) \mid \text{not\_fp}(fp)$ 
```



Frenetic Queries



Goal: measure the total bytes of web traffic arriving on port 2,
every 30 seconds

data to be returned from query
(options: sizes, counts, packets)


filter based on packet headers
(web traffic in on port 2)

```
def web_query():  
    return (Select (sizes) *  
            Where (inport_fp (2) & srcport_fp (80)) *  
            Every (30))
```

period: 30 seconds

Key Property: Query semantics independent of other program parts

Frenetic Queries




Goal: sum the number of packets, per host (ie: mac address), traveling through port 2, every minute

```
def host_query():  
    return (Select (counts) *  
            Where (inport_fp(2)) *  
            GroupBy ([srcmac]) *  
            Every (60))
```

categorize results by
srcmac address



Frenetic Queries



Goal: report the hosts connected to each switch port;
report a host each time it moves from one port to the next

get packets for analysis

```
def learning query():  
    return (Select (packets) *  
            GroupBy ([srcmac]) *  
            SplitWhen ([inport]) *  
            Limit (1))
```

at most one packet per flow

categorize by srcmac

sub-categorize when the inport
changes (the host moves)

Key Property: Query implementation handles network race conditions

Using Queries

- ▣ Query results, or other streams, are piped in to **listeners**

```
def web_query(): ...  
def host_query(): ...  
def learning_query(): ...
```


```
def web_stats():  
    web_query() >> Print()
```

```
def all_stats():  
    Merge(web_query(), host_query()) >> Print()
```

Key Property: Queries compose



Frenetic Forwarding Policies



Goal: implement a repeater switch

rule
actions

```
rules = [Rule(inport_fp(1), [forward(2)]),
         Rule(inport_fp(2), [forward(1)])]
```

```
def repeater():
    return (SwitchJoin() >>
            Lift(lambda switch: {sw(defined over headers))
```

```
def main():
    repeater() >> register()
```

construct repeater policy for that switch

Key Property: Policy semantics independent of other queries/policies
register policy with run time

Program Composition

Goal: implement both the stats monitor and the repeater

```
def main():  
    repeater() >> register()  
    all_stats()
```

Key Property: Queries and policies compose



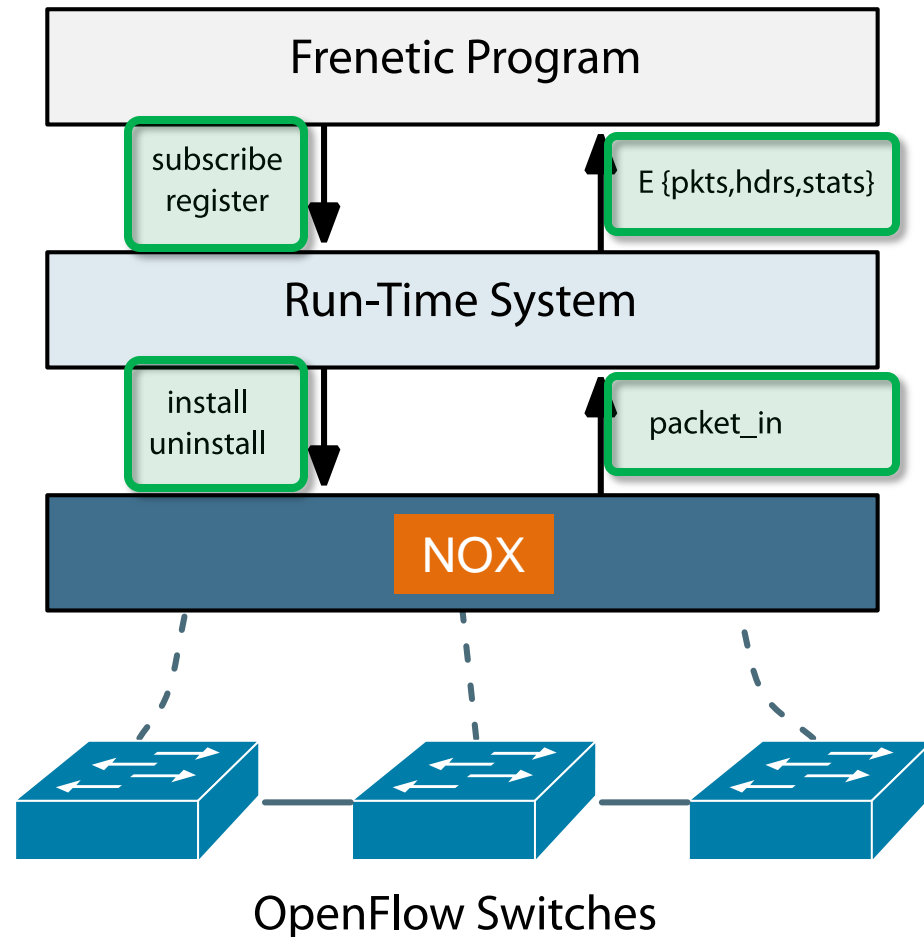
Frenetic Run-time System

Frenetic programs interact only with the run-time

- Programs create *subscribers*
- Programs *register* rules

Run-time handles the details

- Manages switch-level rules
- Handles NOX events
- Pushes values onto the appropriate event streams



Implementation Options

□ Rule Granularity

- **microflow** (exact header match)
 - simpler; more rules generated
- **wildcard** (multiple header match in single rule)
 - more complex; fewer rules (may be) generated

Frenetic 1.0

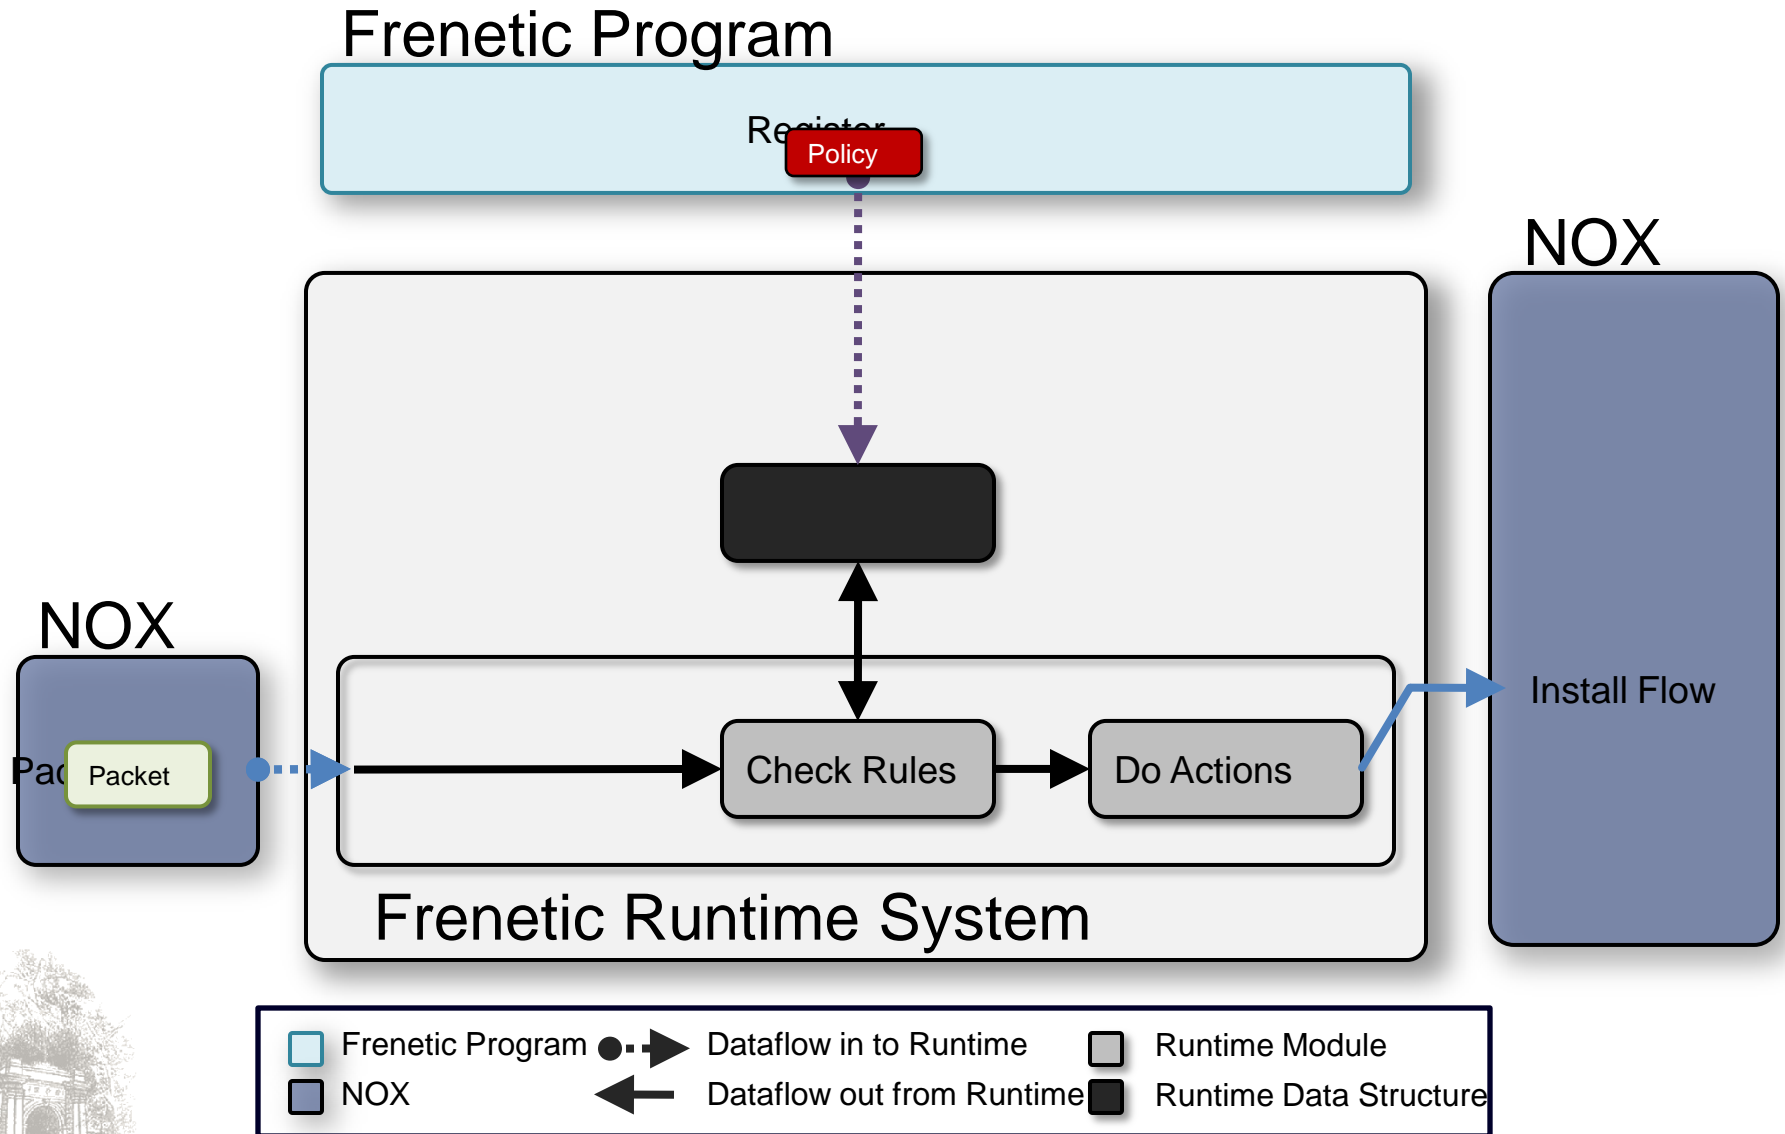
□ Rule Installation

- **reactive** (lazy)
 - first packet of each new flow goes to controller
- **proactive** (eager)
 - new rules pushed to switches

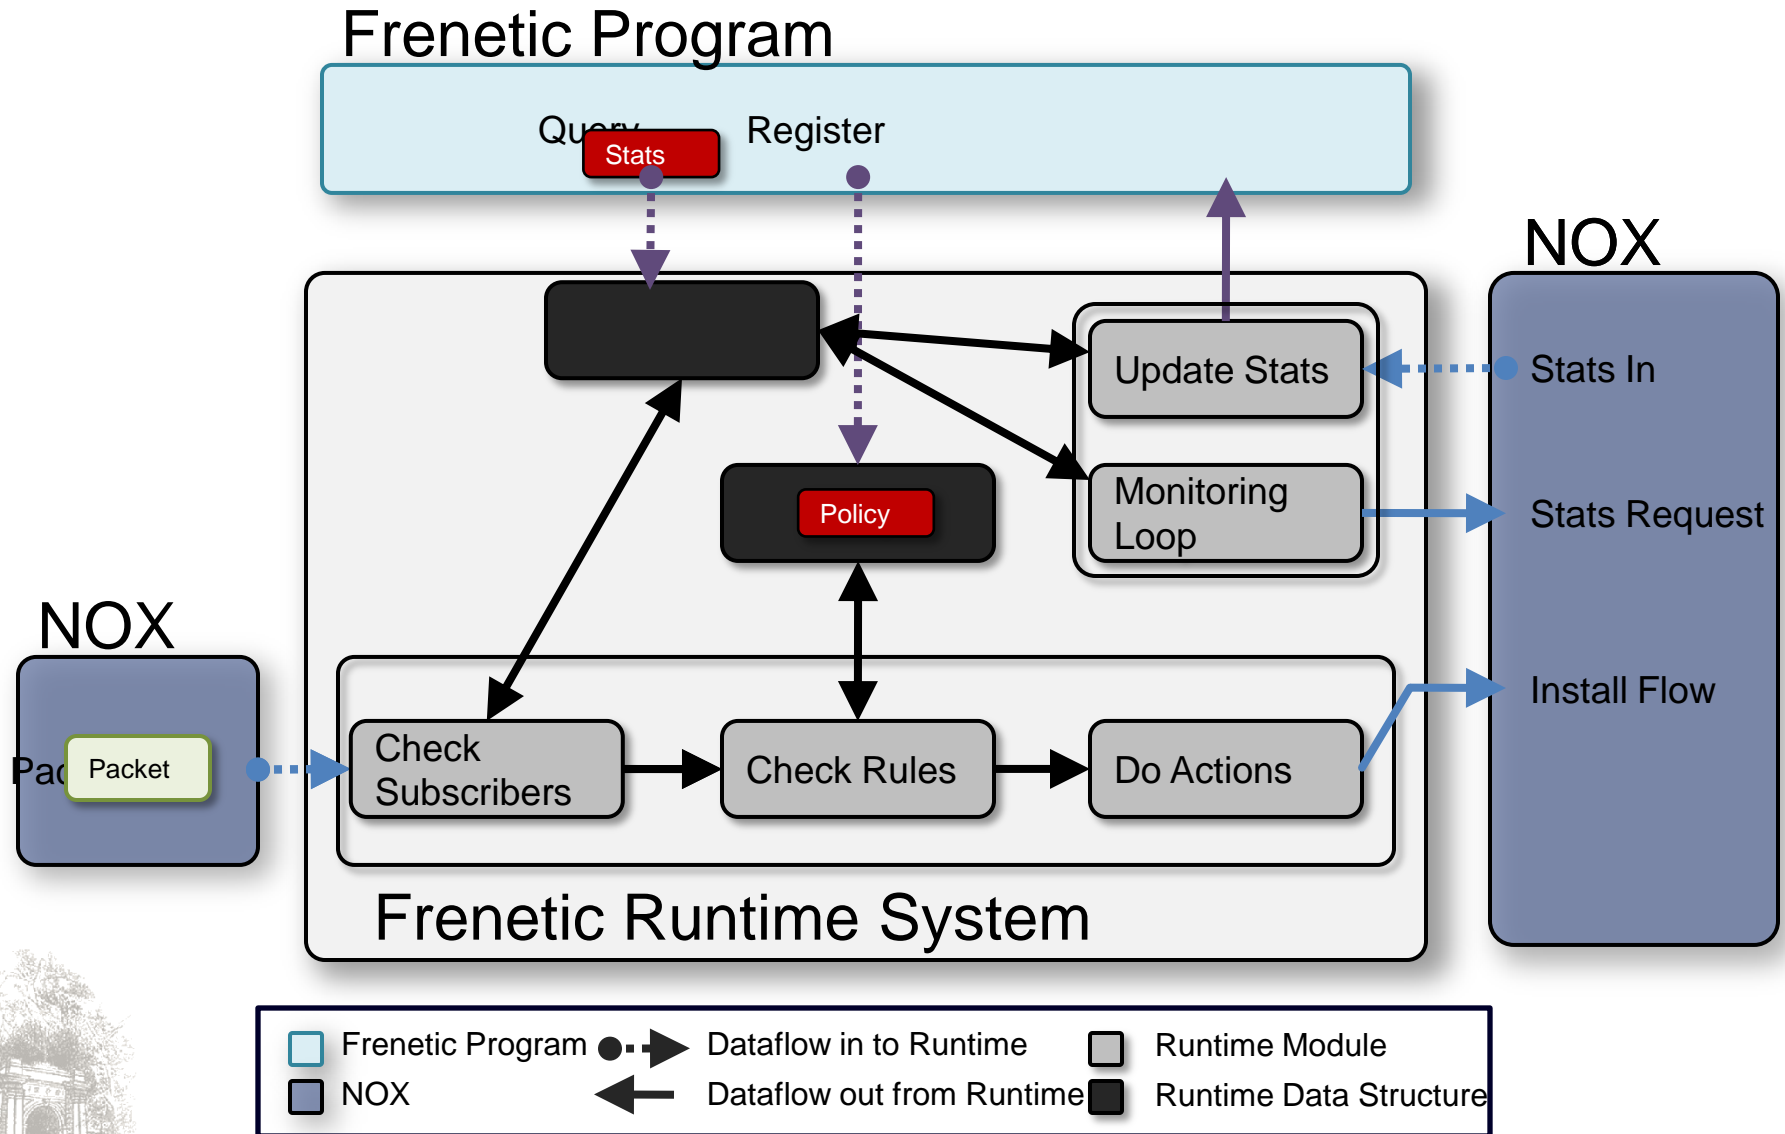
Frenetic 2.0



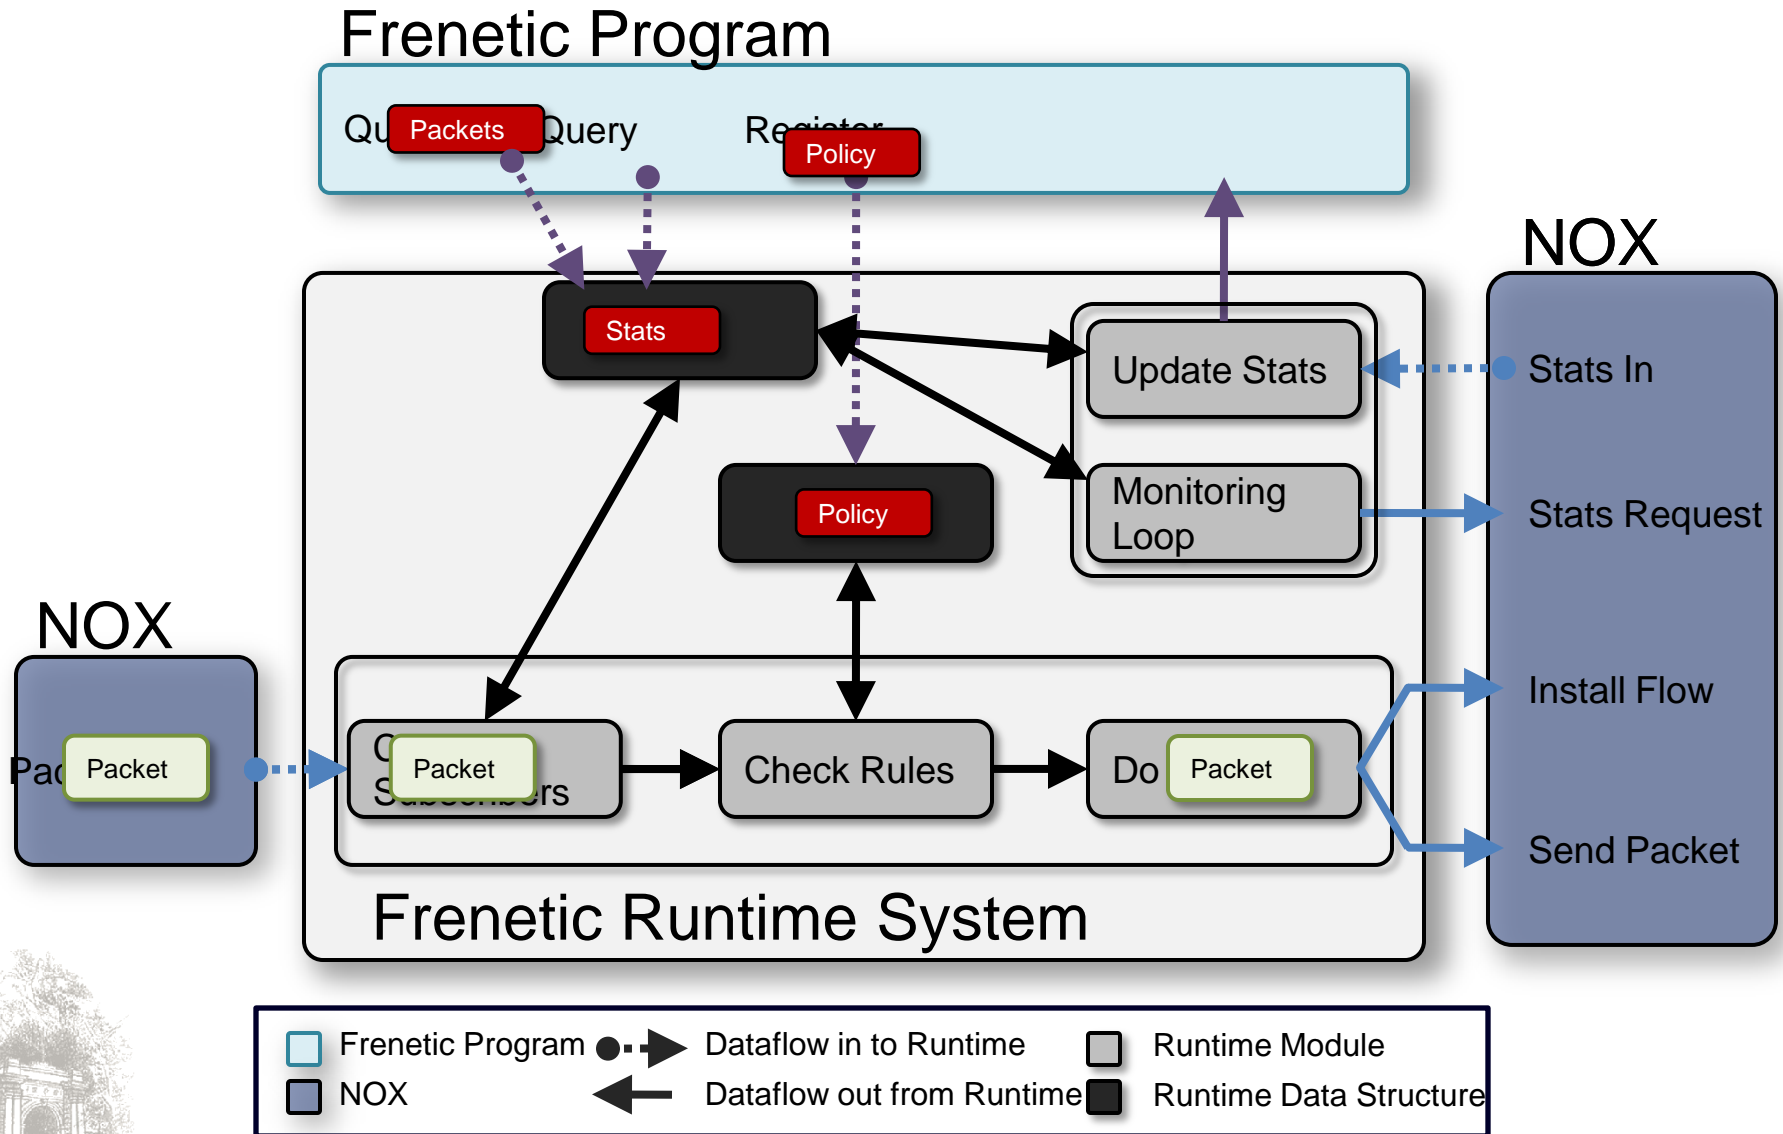
Run-time Activities



Run-time Activities



Run-time Activities



Evaluation

Metrics:

- Lines of code
- Traffic to controller
- Total traffic

Microbenchmarks:

- All-Pairs Connectivity
- Web Statistics
- Heavy Hitters

Forwarding Policy:


- Hub
- Learning Switch
- Loop-Free Learning Switch




Evaluation

		Connectivity			Heavy Hitters			Web Stats		
		<i>HUB</i>	<i>LSW</i>	<i>LFLSW</i>	<i>HUB</i>	<i>LSW</i>	<i>LFLSW</i>	<i>HUB</i>	<i>LSW</i>	<i>LFLSW</i>
NOX	Lines of Code	20	55	75	110	198		104	135	
	Controller Traffic (kB)	12.8	13.5	31.3	9.3	10.3	*	4.5	4811	*
	Aggregate Traffic (kB)	69.2	42.3	64.1	57.2	36.1		14.1	9.0	
Frenetic	Lines of Code	6	30	58	29	53	81	13	37	65
	Controller Traffic (kB)	9.1	12.0	12.4	11.1	10.6	10.9	4.5	5.1	5.8
	Aggregate Traffic (kB)	65.6	41.0	41.5	55.0	36.4	36.9	13.6	9.20	9.9


Table 1. Experimental results.



(a) All-Pairs Connectivity




(c) Heavy Hitters



(b) Web Statistics

Figure 9. Scalability experimental results.


MicroBench: Lines of Code



Forwarding Policy:
HUB: Floods out other ports
LSW: Learning Switch

Monitoring Policy

MicroBench: Controller Traffic

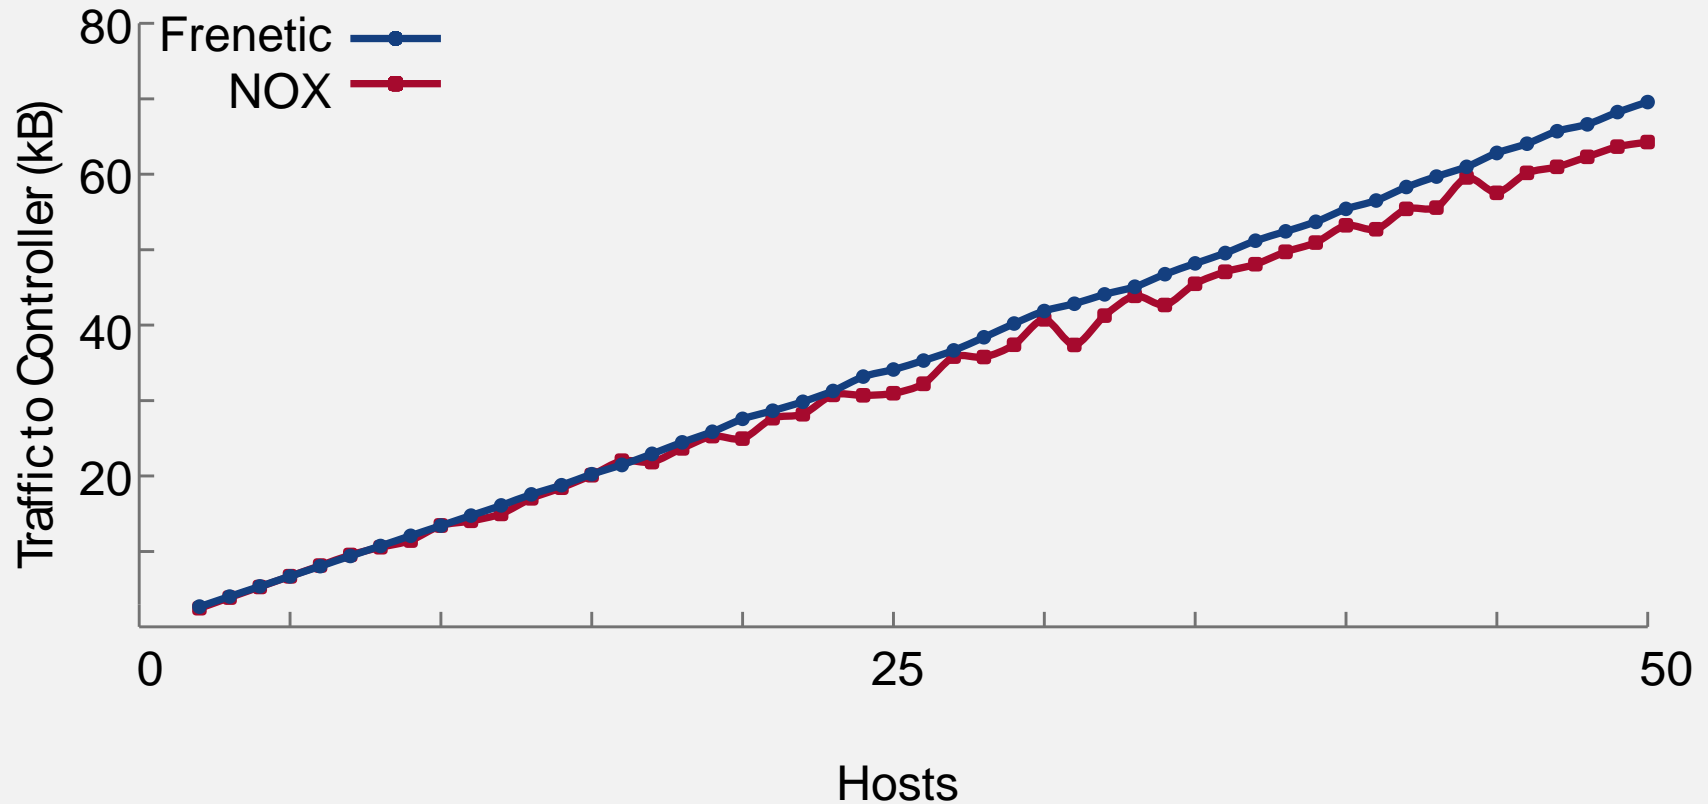


Forwarding Policy:
HUB: Floods out other ports
LSW: Learning Switch

Monitoring Policy

Frenetic Scalability

Frenetic scales to larger networks comparably with NOX



Ongoing and Future Work

❑ Surface Language

- ❑ Current prototype is in Python – to ease transition
- ❑ Would like a **standalone language**

❑ Optimizations

- ❑ More programs can also be implemented **efficiently**
- ❑ Would like a **compiler** to identify and rewrite **optimizations**

❑ Proactive Strategy

- ❑ Current prototype is **reactive**, based on microflow rules
- ❑ Would like to enable **proactive**, wildcard rule installation

❑ Network Wide Abstractions

- ❑ Current prototype focuses only on a **single** switch
- ❑ Need to expand to **multiple** switches



Acknowledgements

- ❑ Almost the whole content comes from authors' slides presented at PRESTO 2010, Microsoft Research 2011 and also their paper at ICFP 2011
- ❑ This slides is only for seminar use in NSLab
- ❑ For more information, please refer to the following links:
 - ❑ <http://www.cs.princeton.edu/~jrex/papers/icfp11.pdf>
 - ❑ <http://www.cs.princeton.edu/~jrex/talks/presto11.pptx>
 - ❑ <http://www.cs.princeton.edu/~dpw/talks/frenetic-05-11.pptx>





Discussion