

Split / Merge

System Support for Elastic Execution in Virtual Middleboxes



Shriram **RAJAGOPALAN** IBM Research & UBC

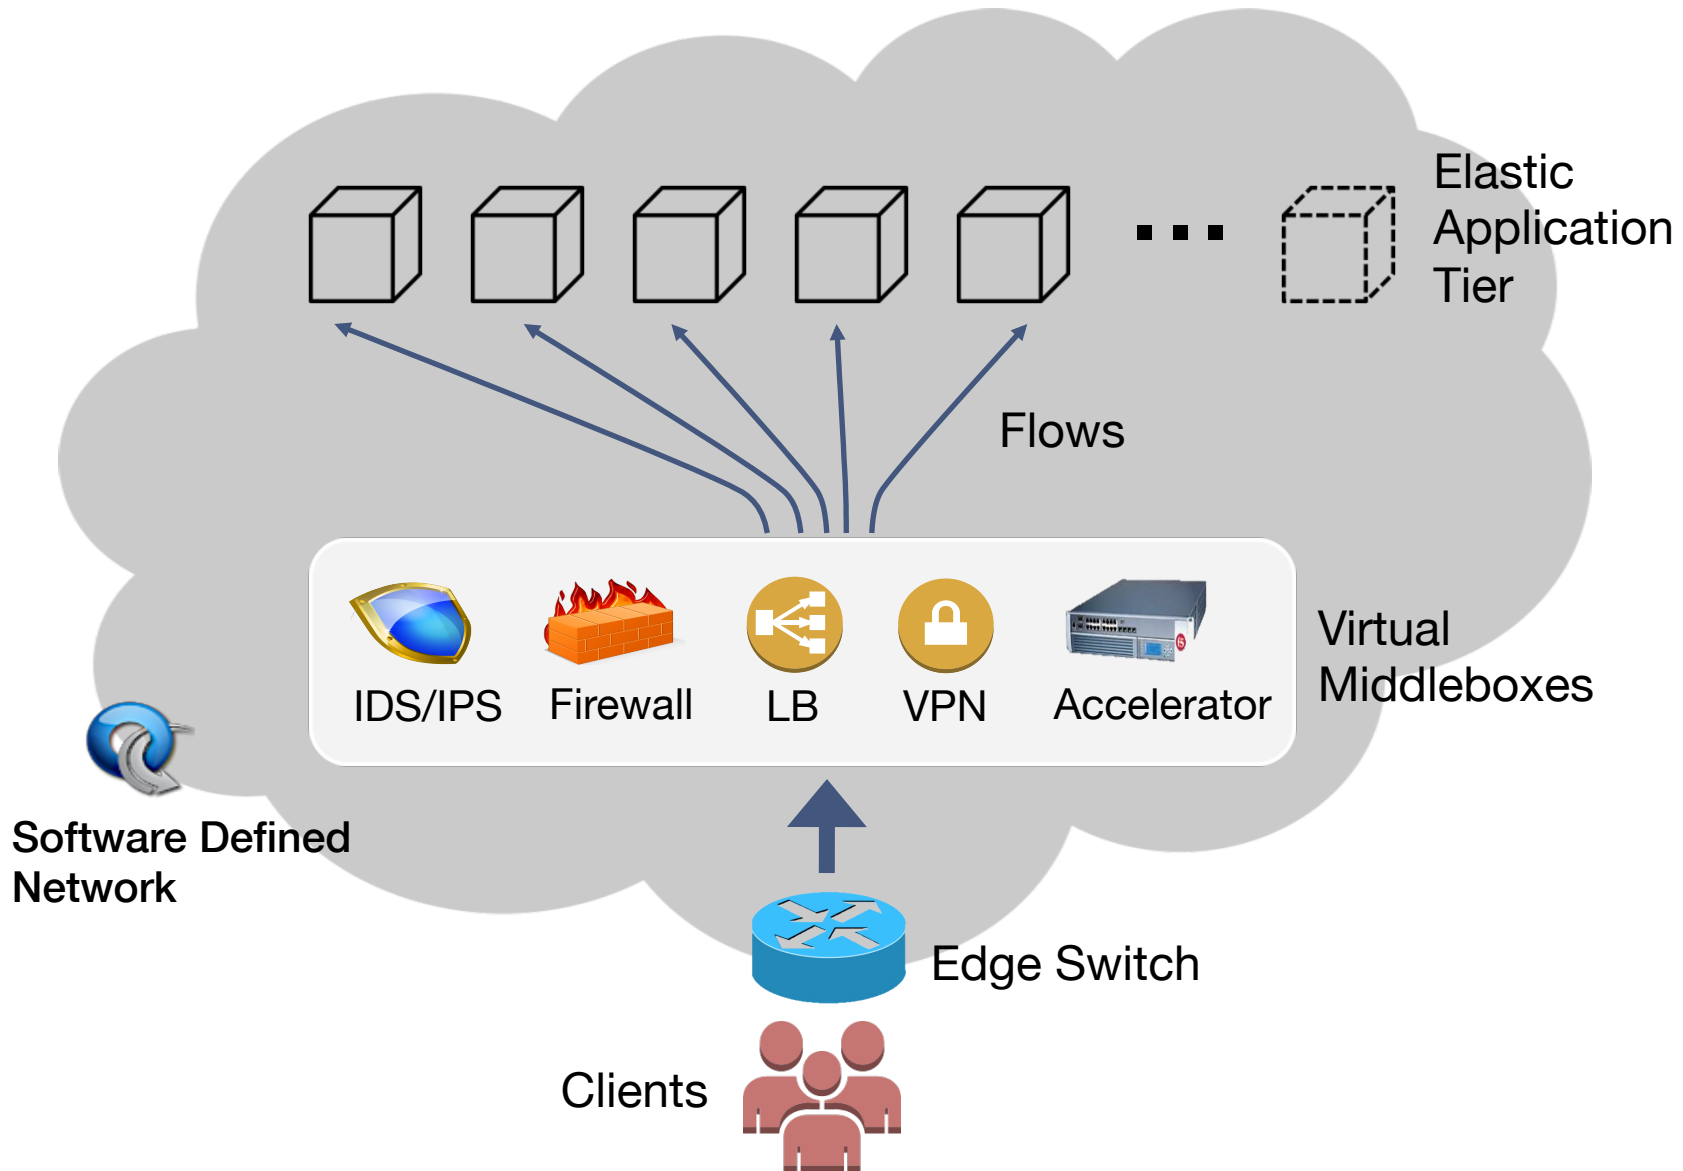
Dan **WILLIAMS** IBM Research

Hani **JAMJOOM** IBM Research

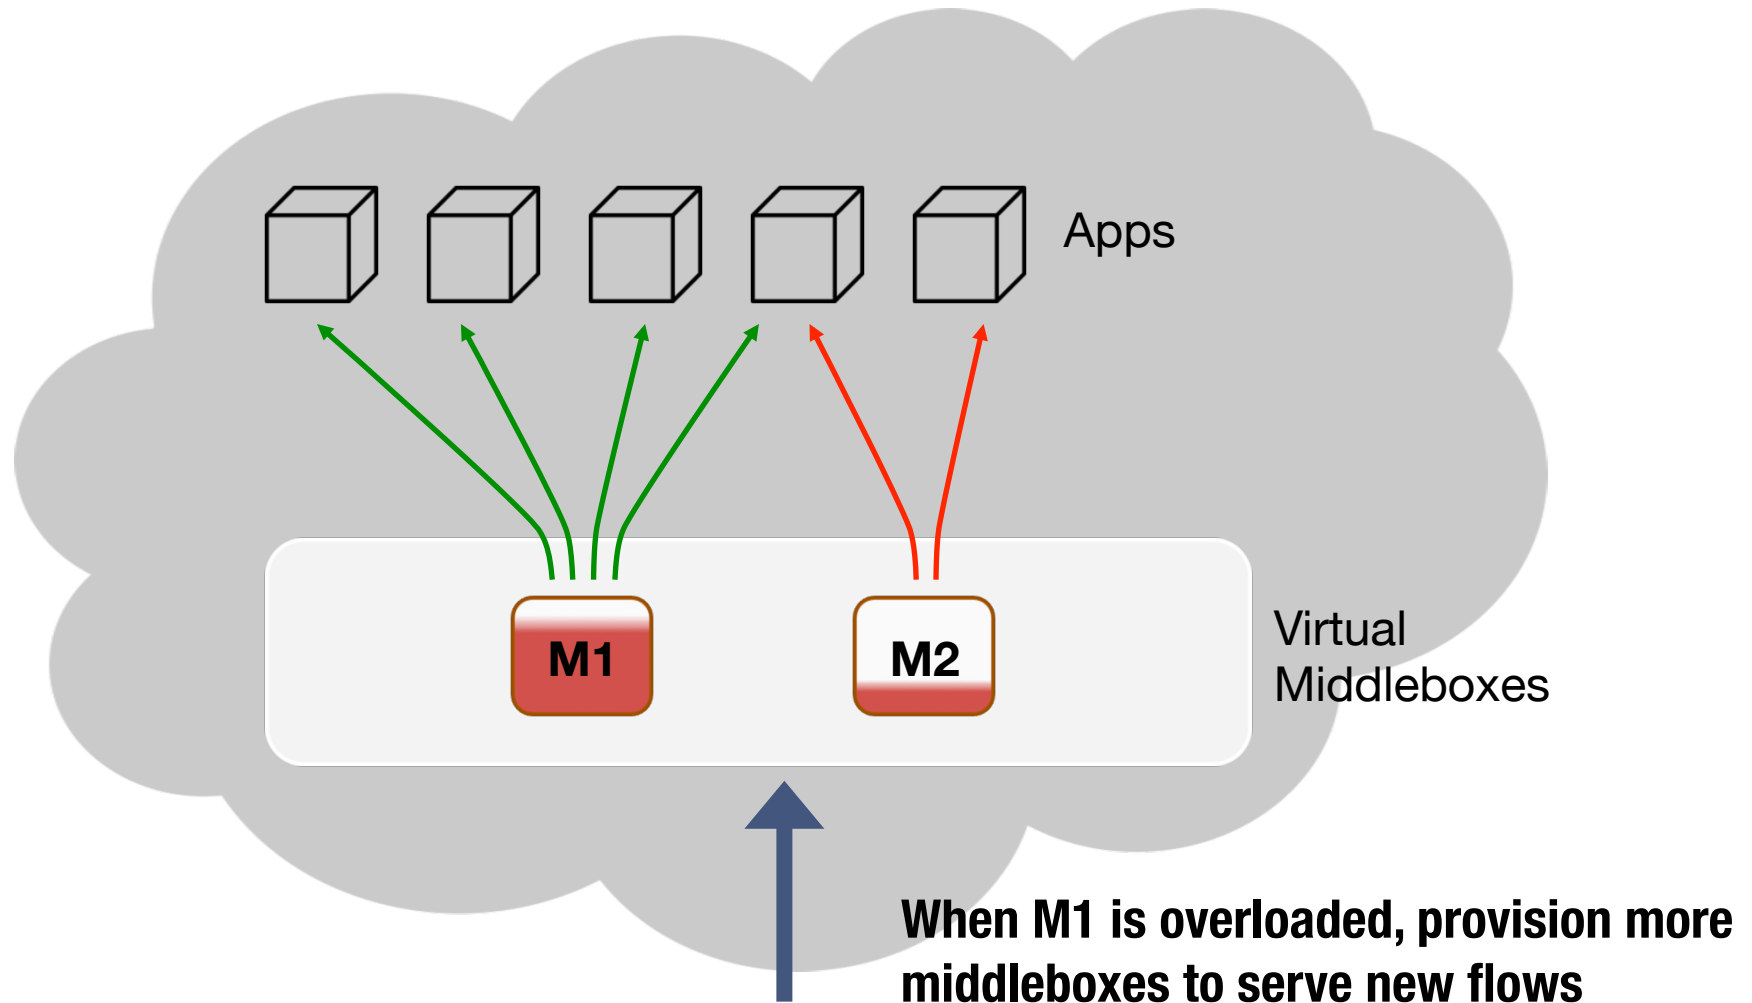
Andrew **WARFIELD** UBC

The Problem

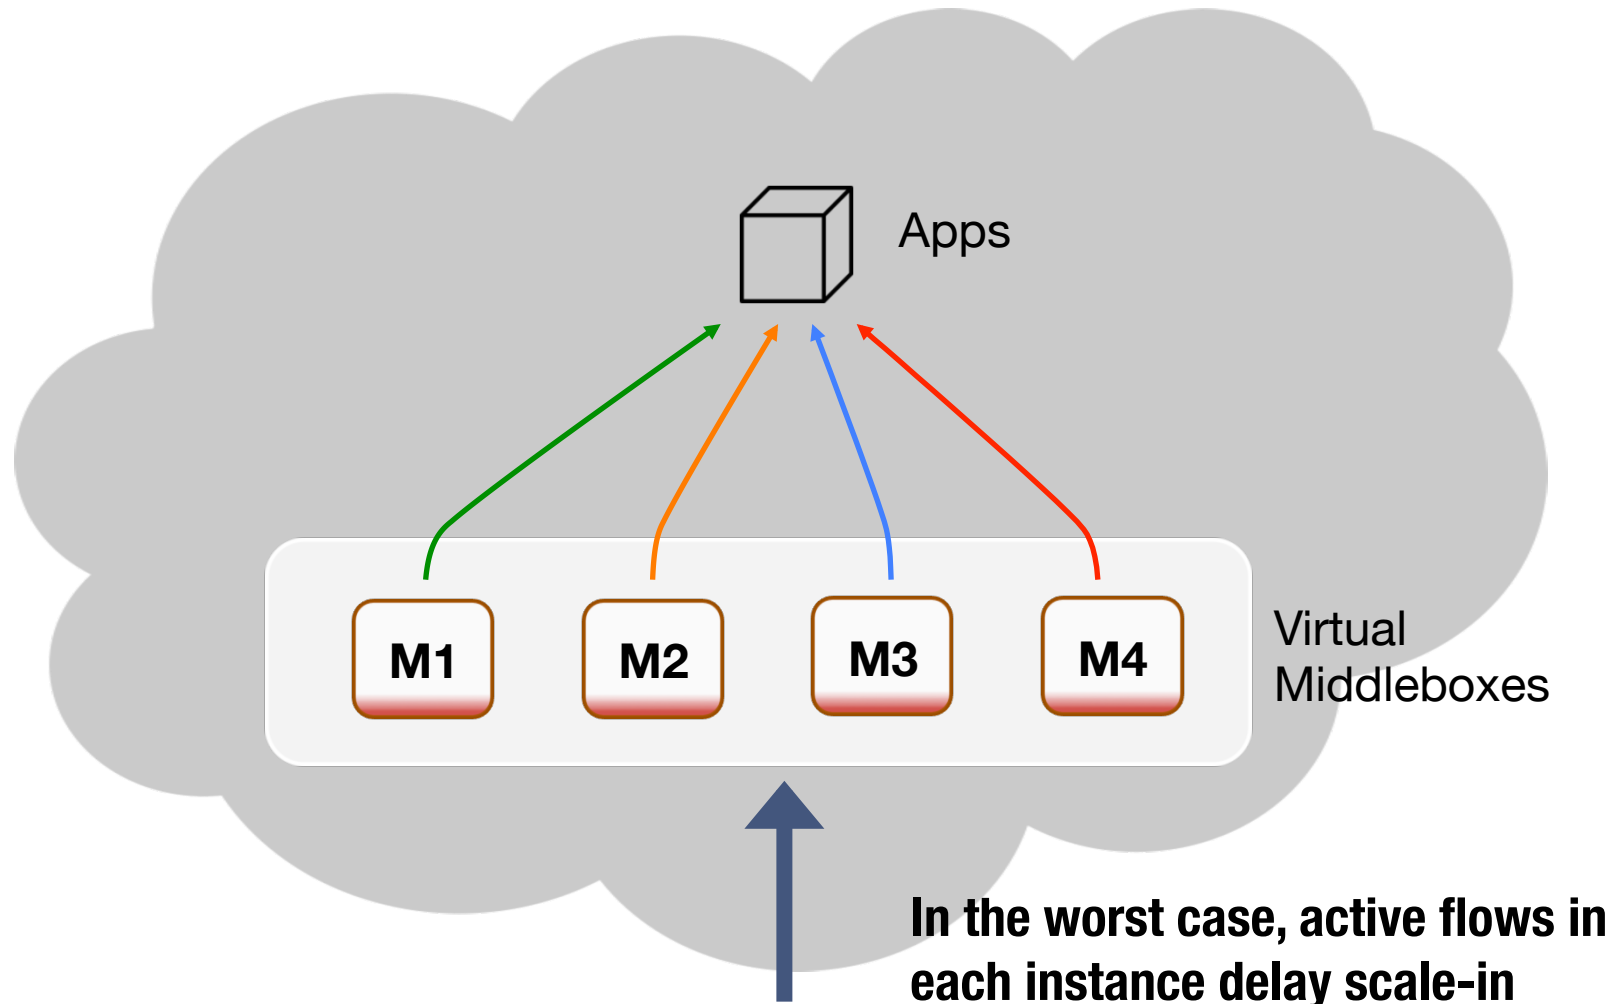
Elastic Applications Need Elastic Middleboxes



Hotspots Cannot be Alleviated Quickly




Scaling Inefficiencies Lead to Poor Utilization

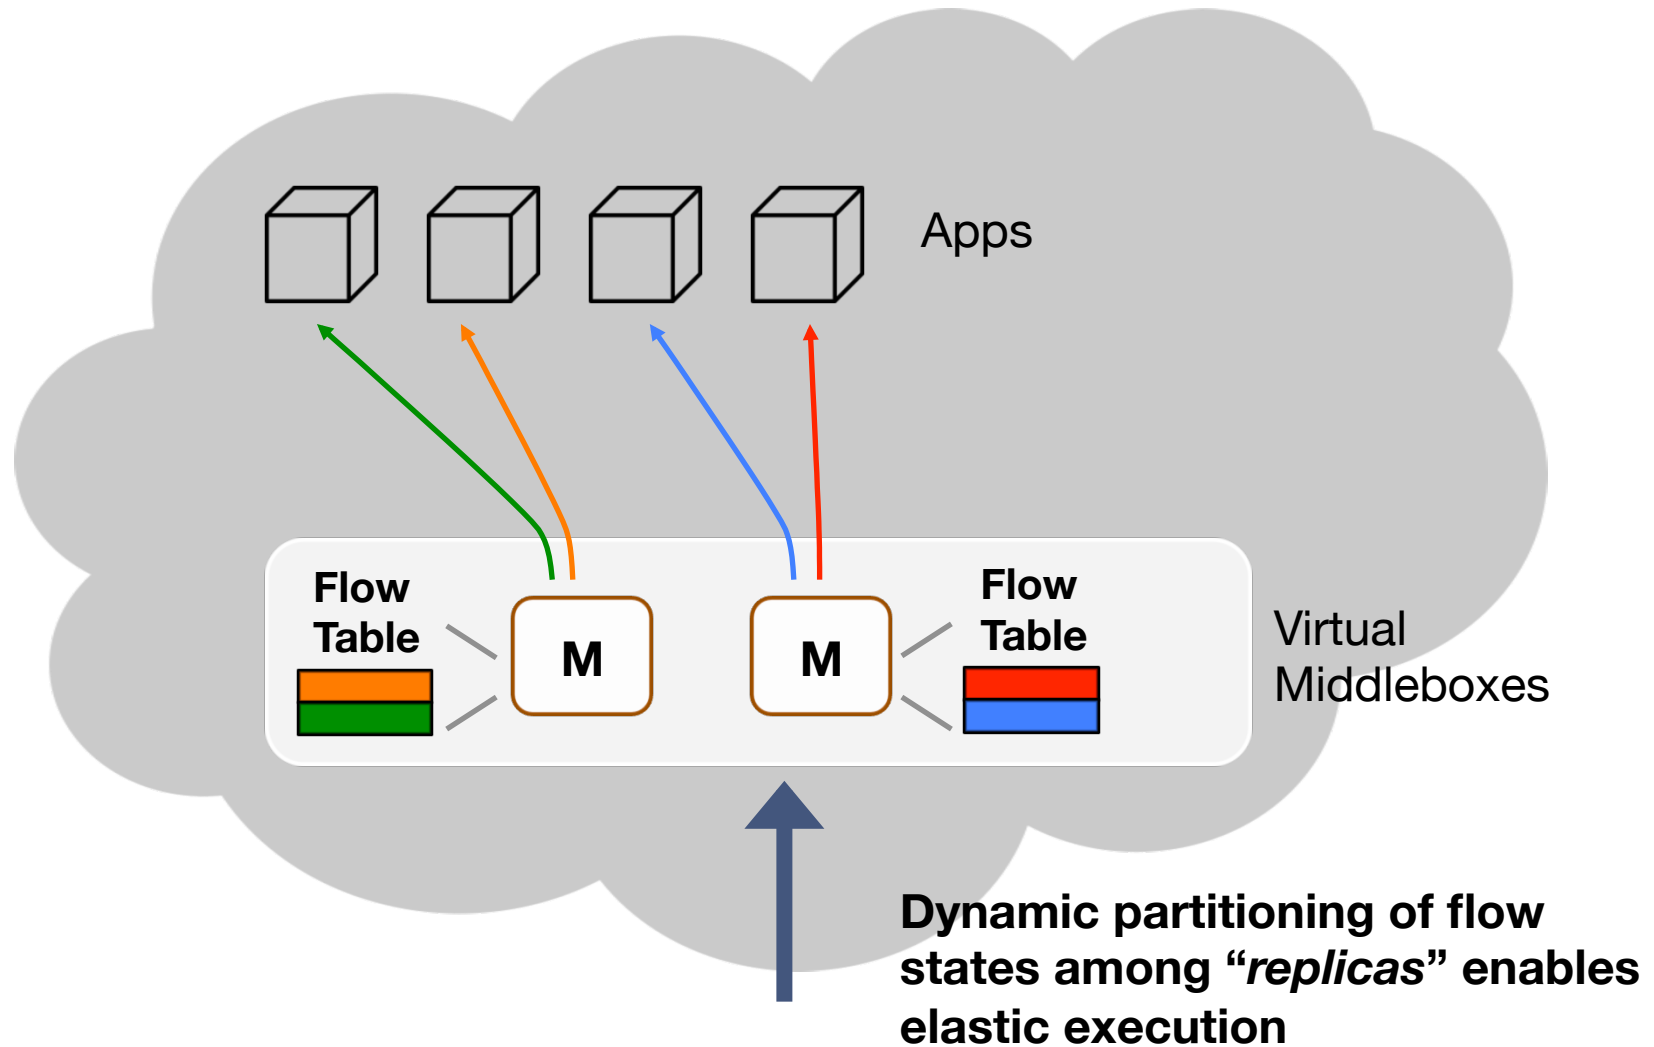


The Insight


Flow State is Naturally Partitioned



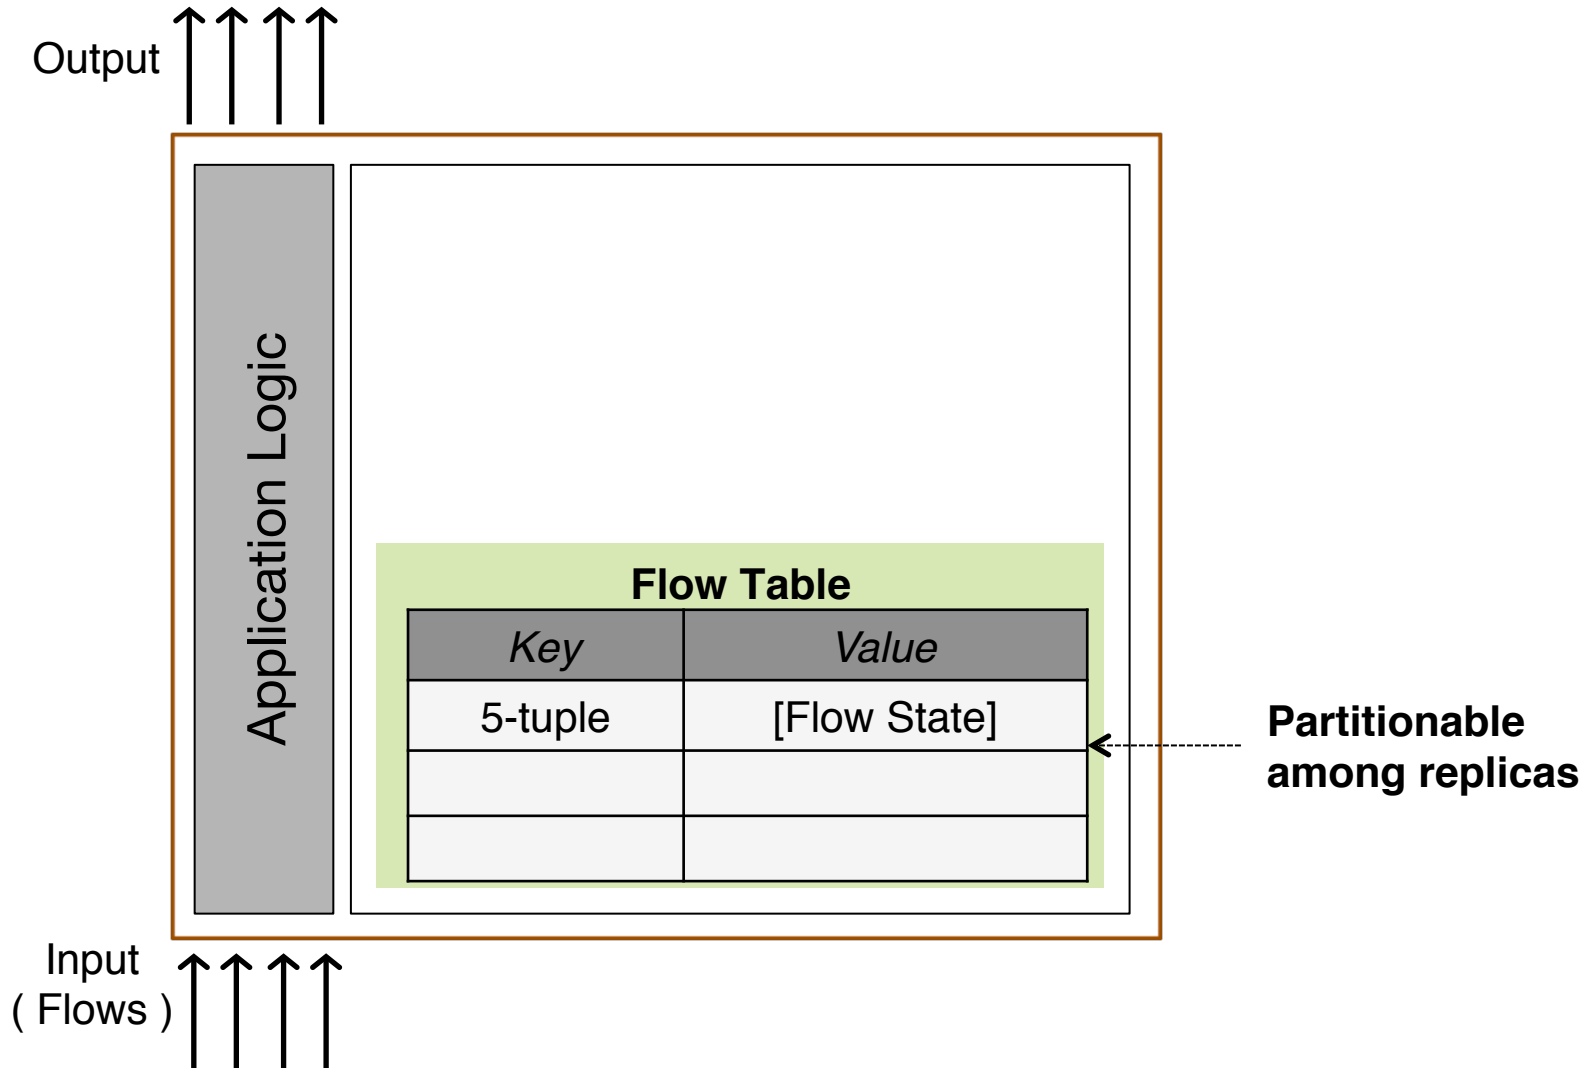
Enabling Elasticity in Virtual Middleboxes



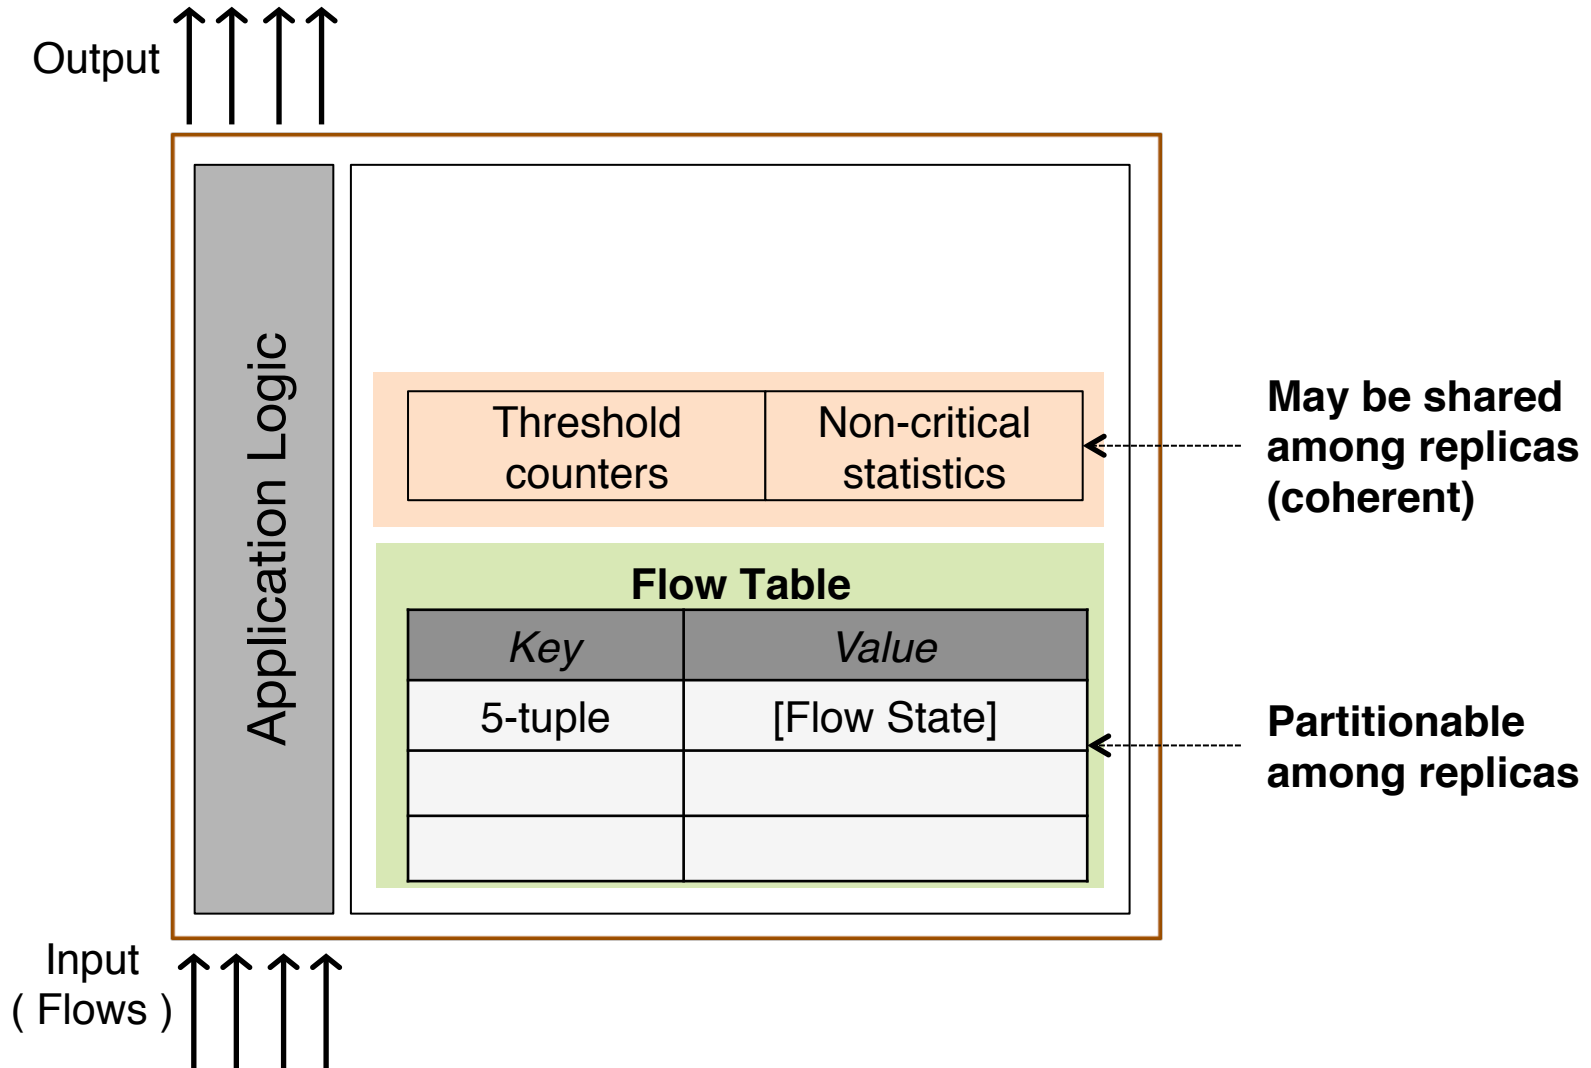
Understanding the State Inside a Middlebox



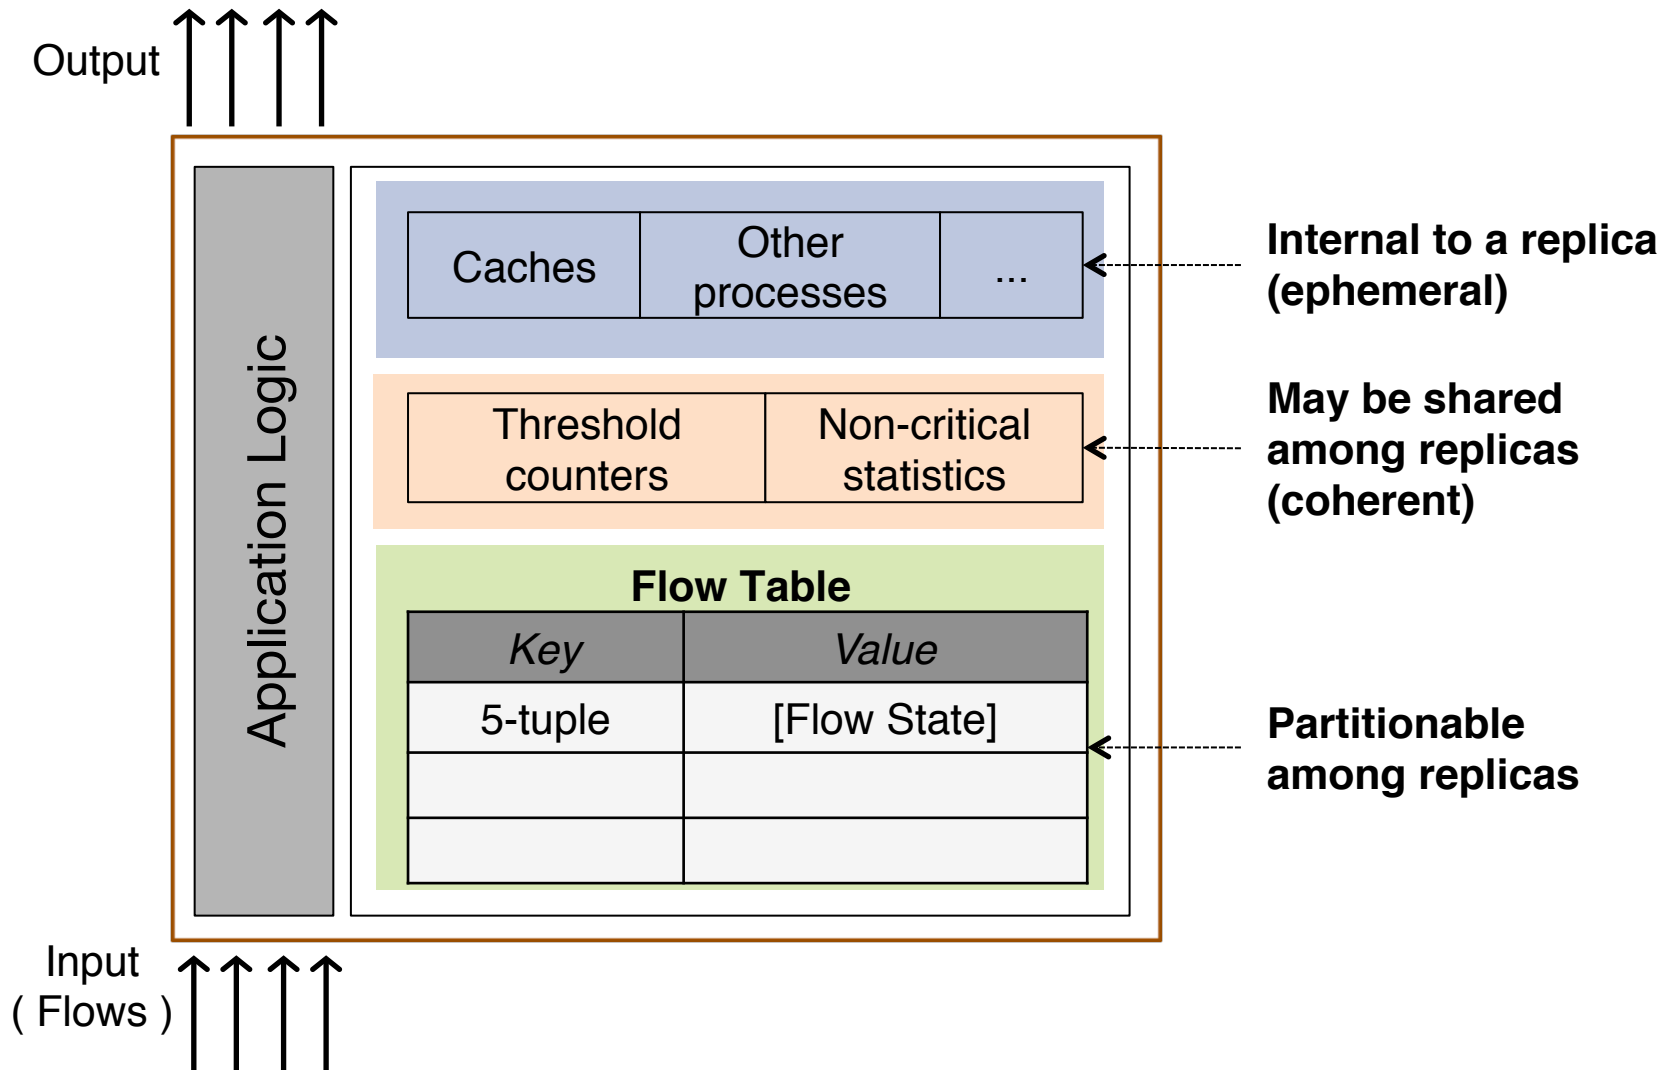
Understanding the State Inside a Middlebox



Understanding the State Inside a Middlebox




Understanding the State Inside a Middlebox




Our Contribution

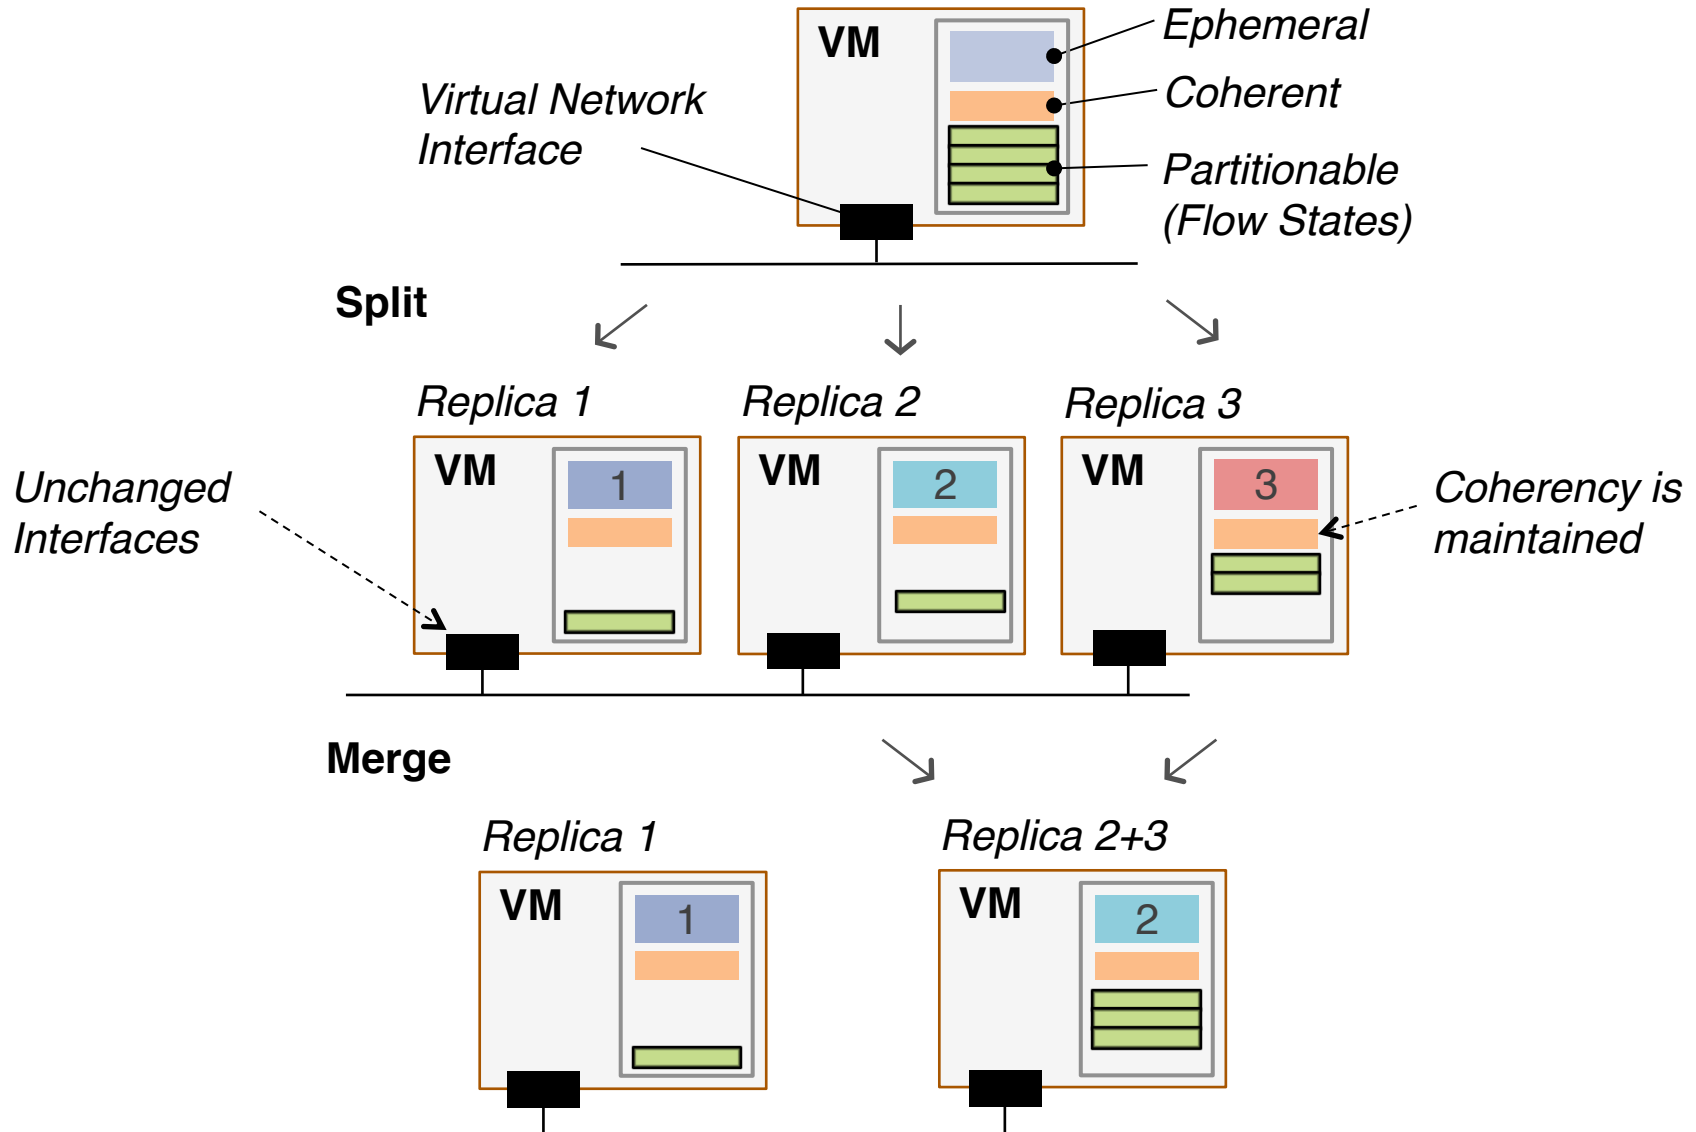
Split/Merge: A State-Centric Approach to Elasticity



Split/Merge: A State-Centric Approach to Elasticity



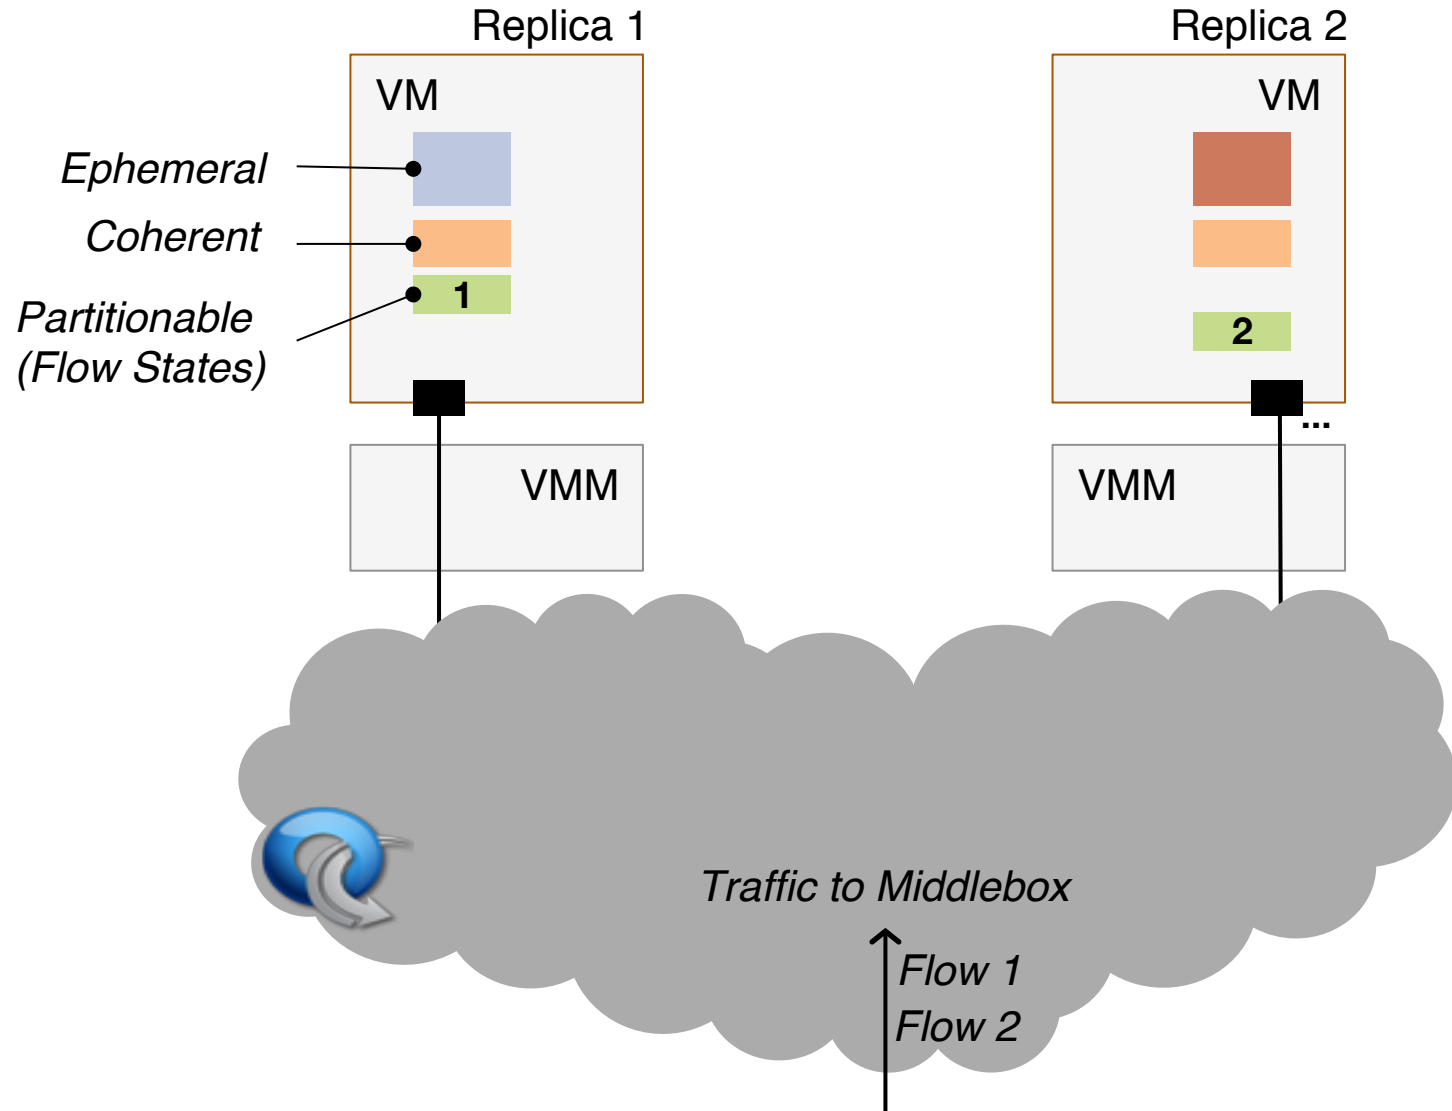
Split/Merge: A State-Centric Approach to Elasticity



Implementation


- A VMM based runtime that provides Split/Merge abstraction to applications
- Developers modify application code to annotate flow state
- FreeFlow takes care of the rest!

FreeFlow: A Split/Merge Implementation




FreeFlow: A Split/Merge Implementation

- Need to manage application state



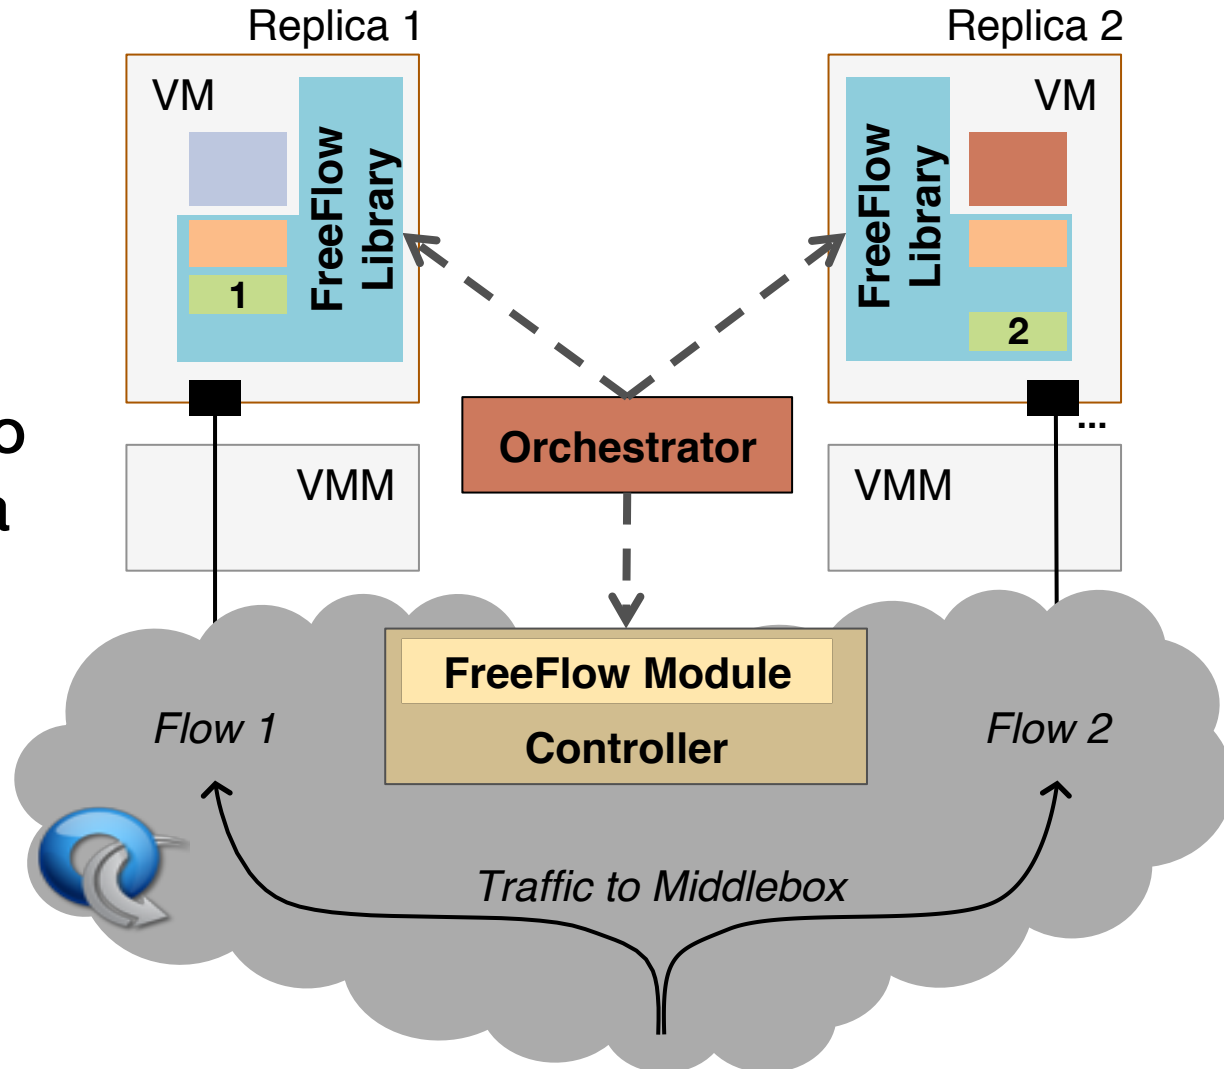
FreeFlow: A Split/Merge Implementation

- Need to manage application state
- Need to ensure flows are routed to the correct replica



FreeFlow: A Split/Merge Implementation

- Need to manage application state
- Need to ensure flows are routed to the correct replica
- Need to decide when to split or merge a replica



Annotating State using FreeFlow API

```
create_flow(flow_key, size)
```

Partitioned State API

```
delete_flow(flow_key)
```

```
flow_state get_flow(flow_key)
```

```
put_flow(flow_key)
```

```
flow_timer(flow_key, timeout, callback)
```

```
create_shared(key, size, cb)
```

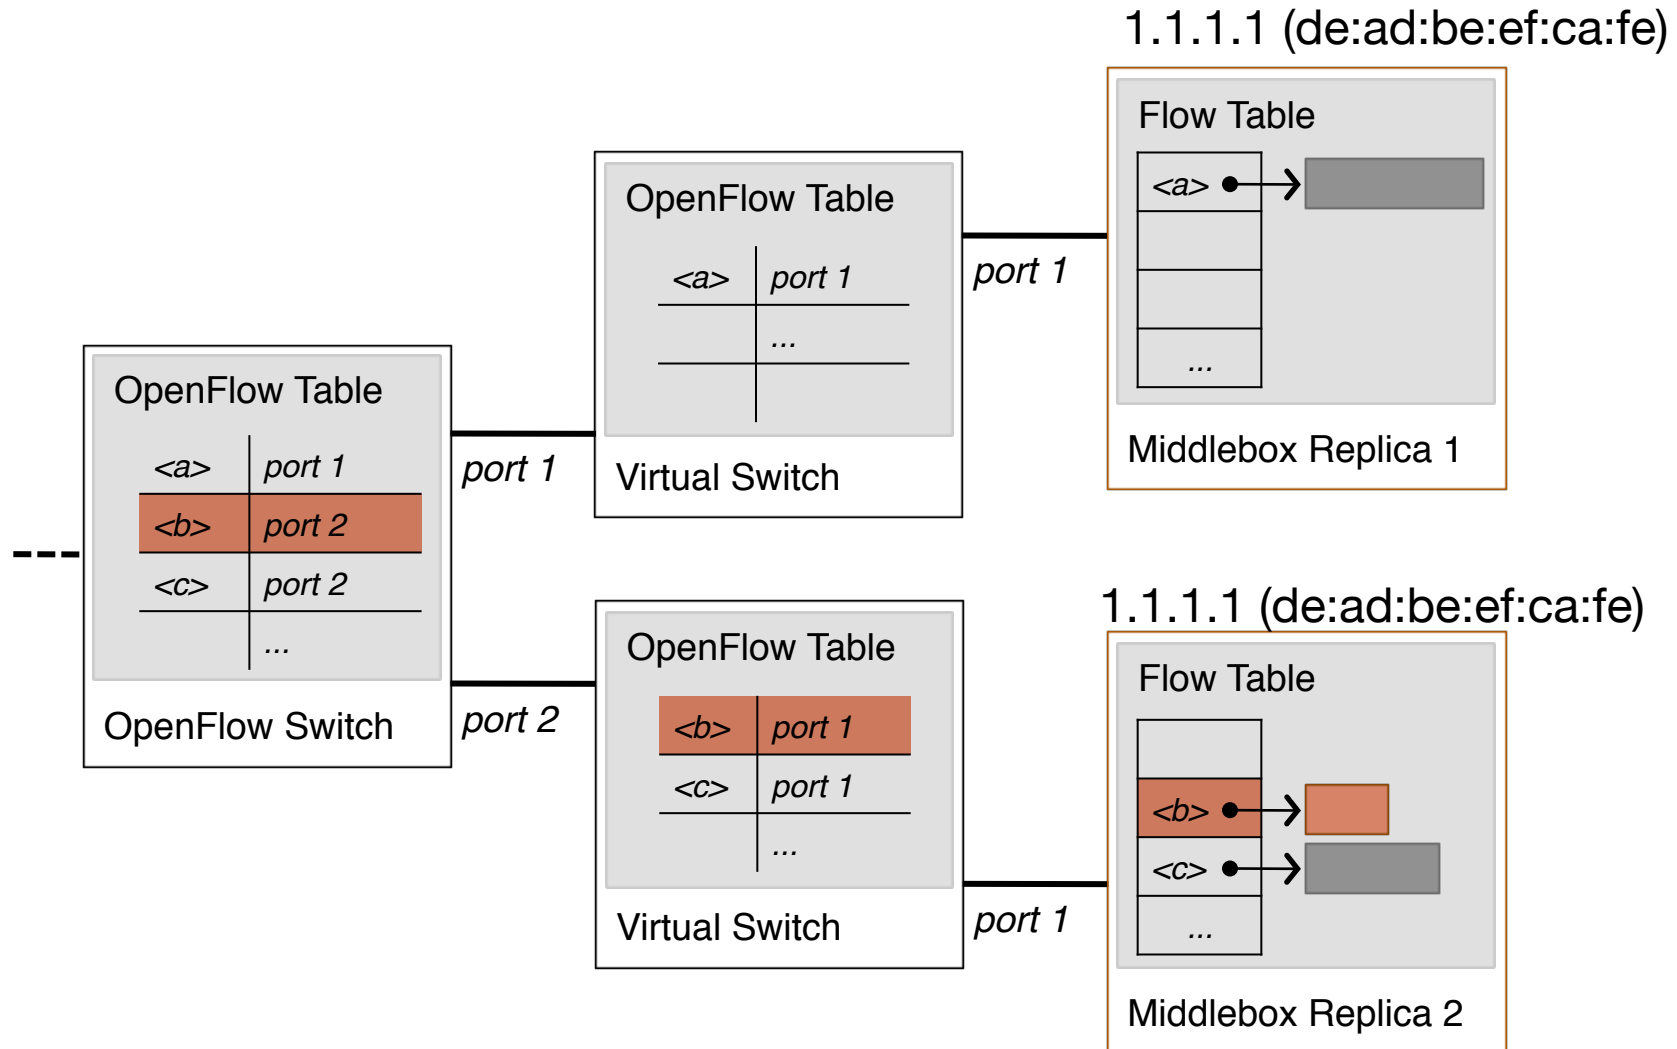
Coherent State API

```
delete_shared(key)
```

```
state get_shared(key, flags) // synch | pull | local
```

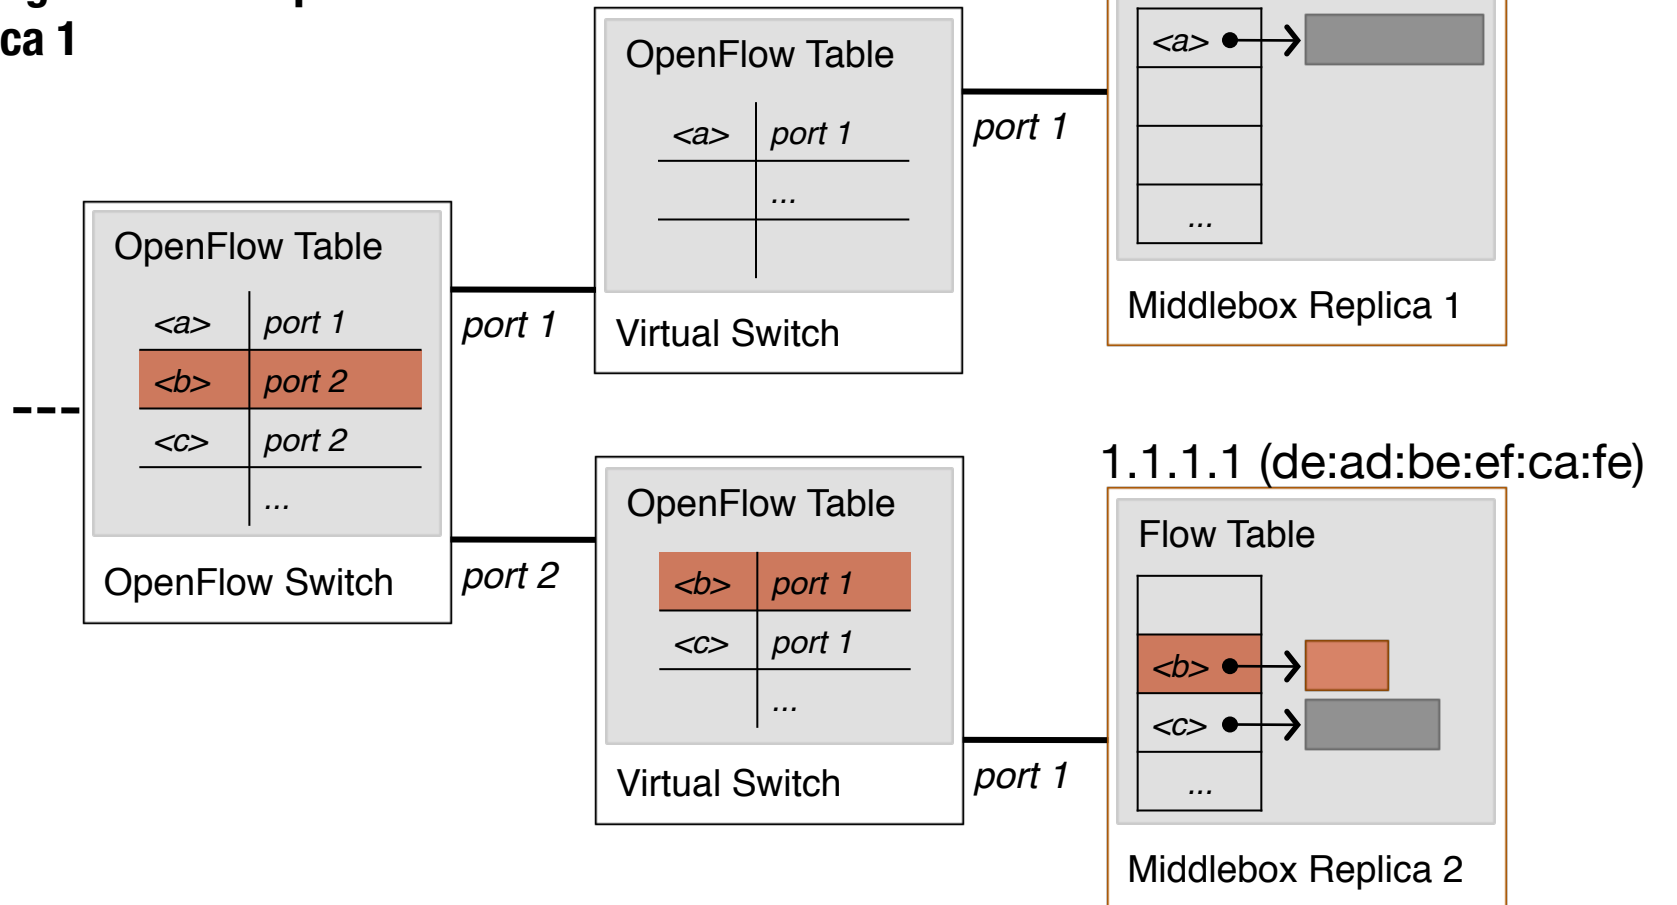
```
put_shared(key, flags) // synch | push | local
```

Forwarding Flows Correctly using OpenFlow



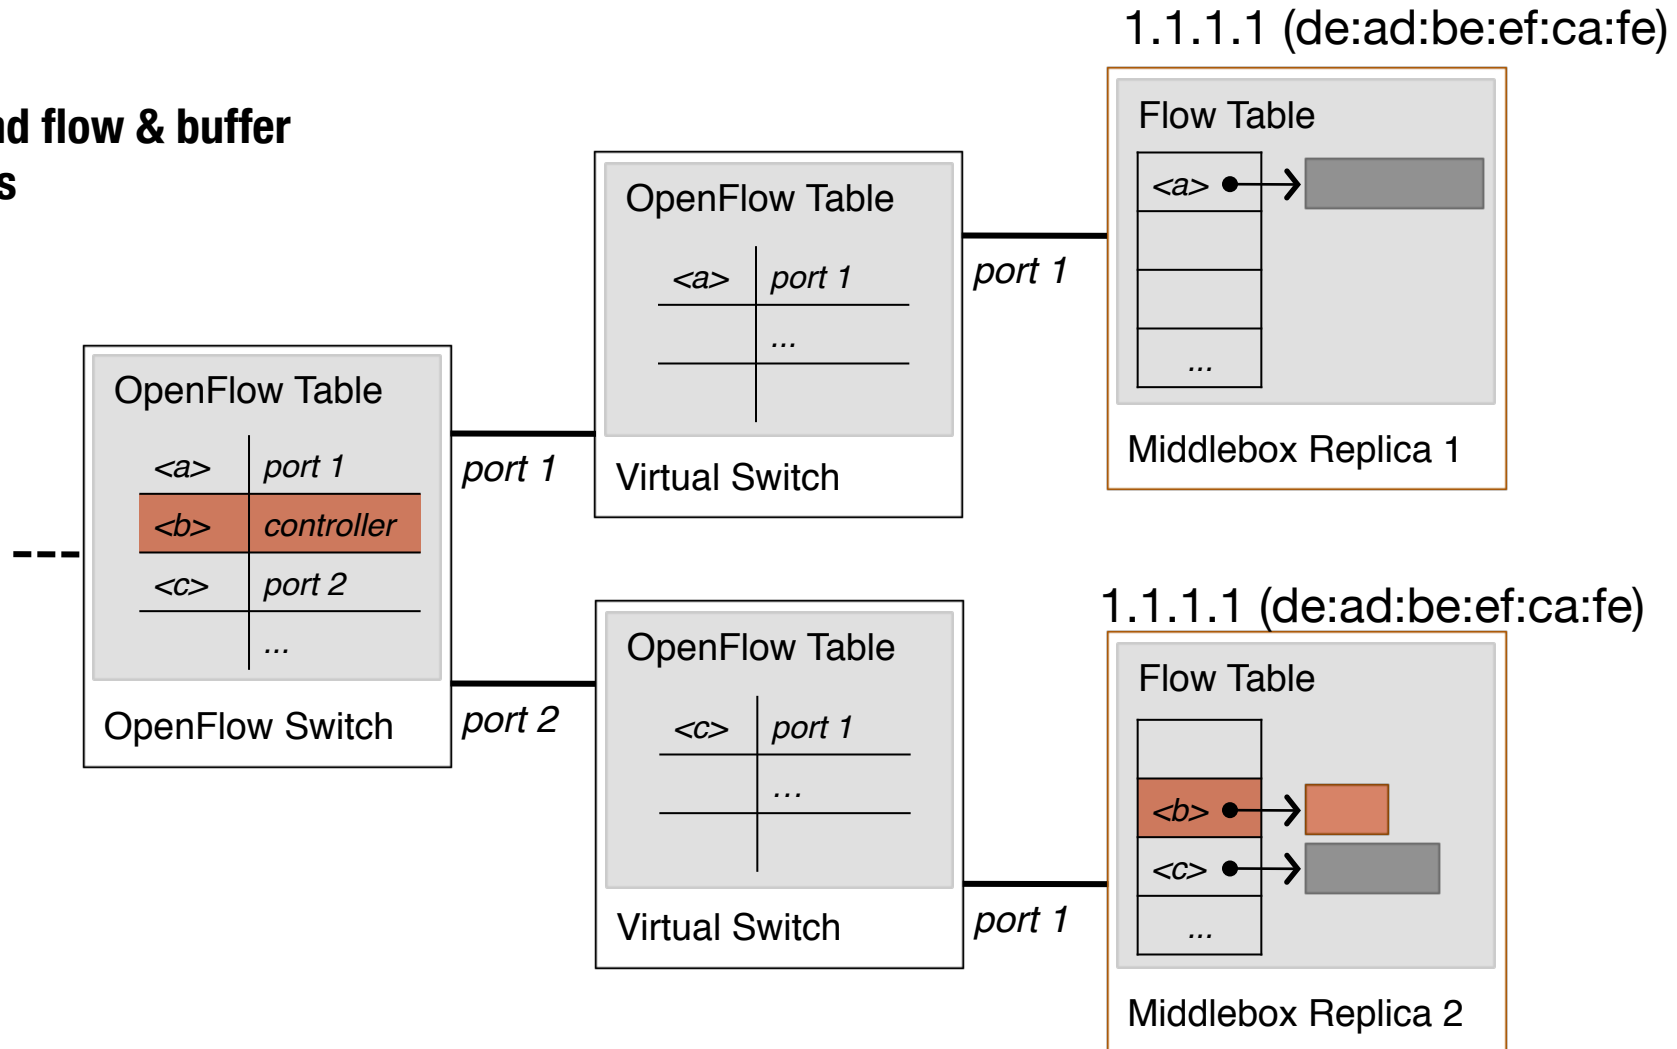
Flow Migration

Migrating from replica 2 to replica 1



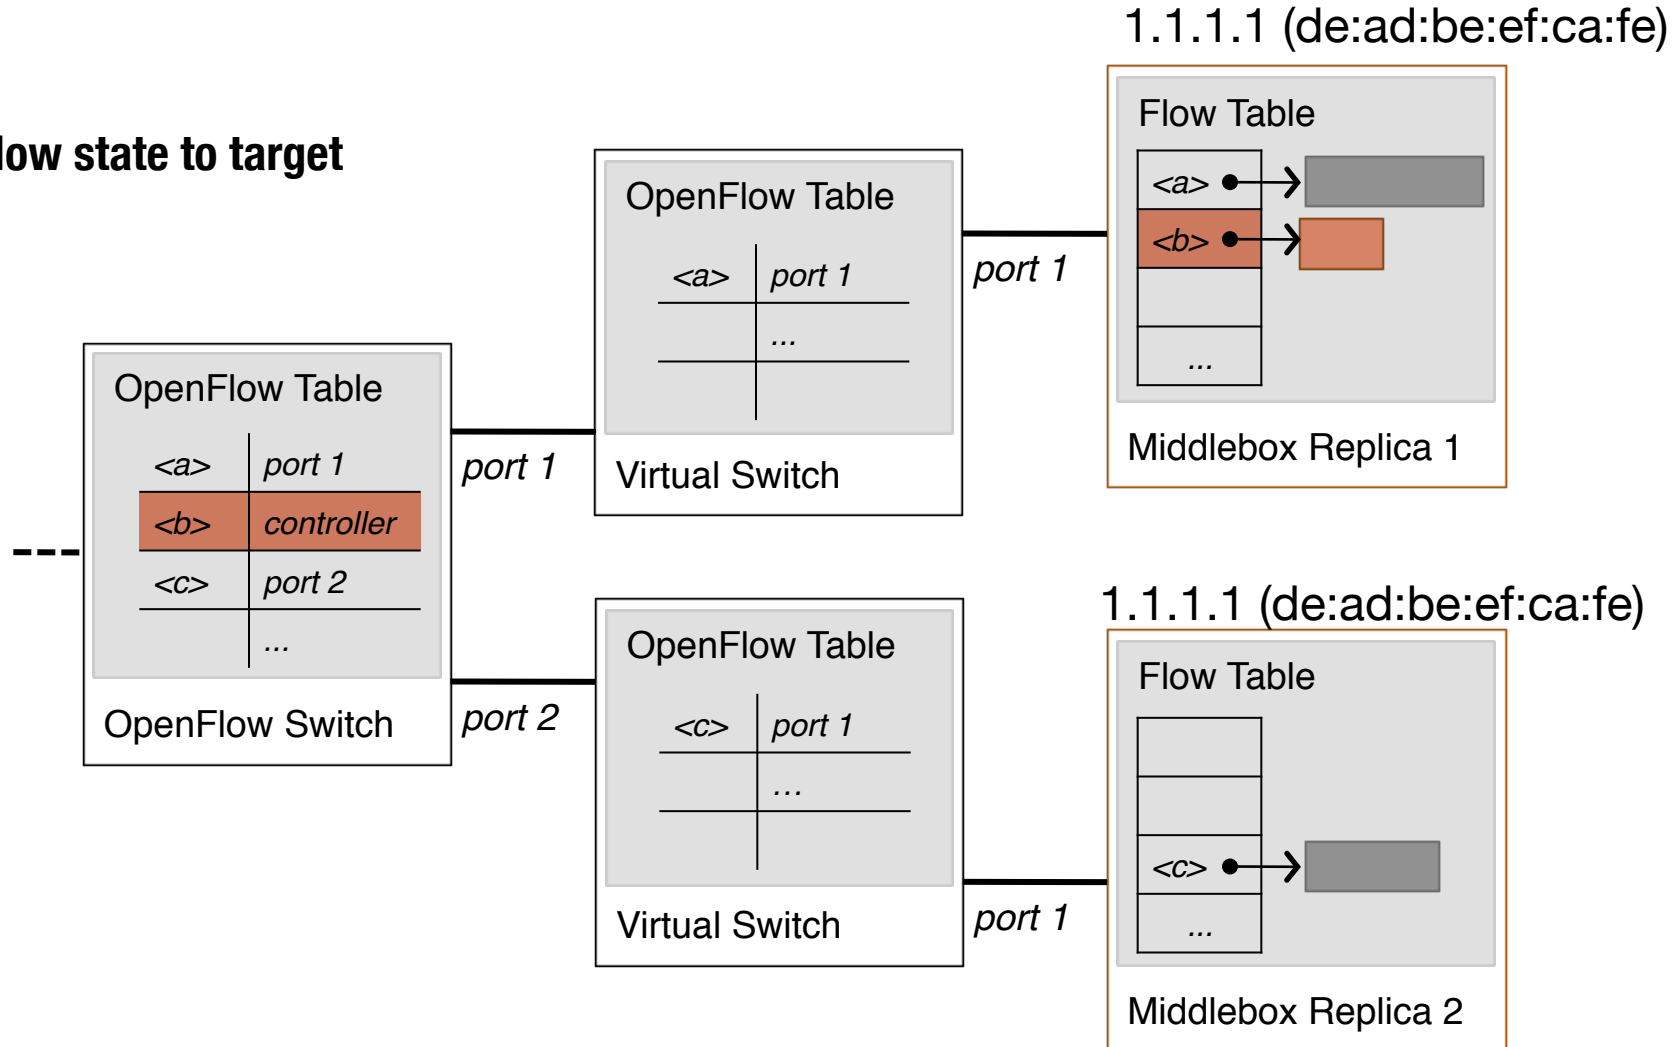
Flow Migration

Suspend flow & buffer packets



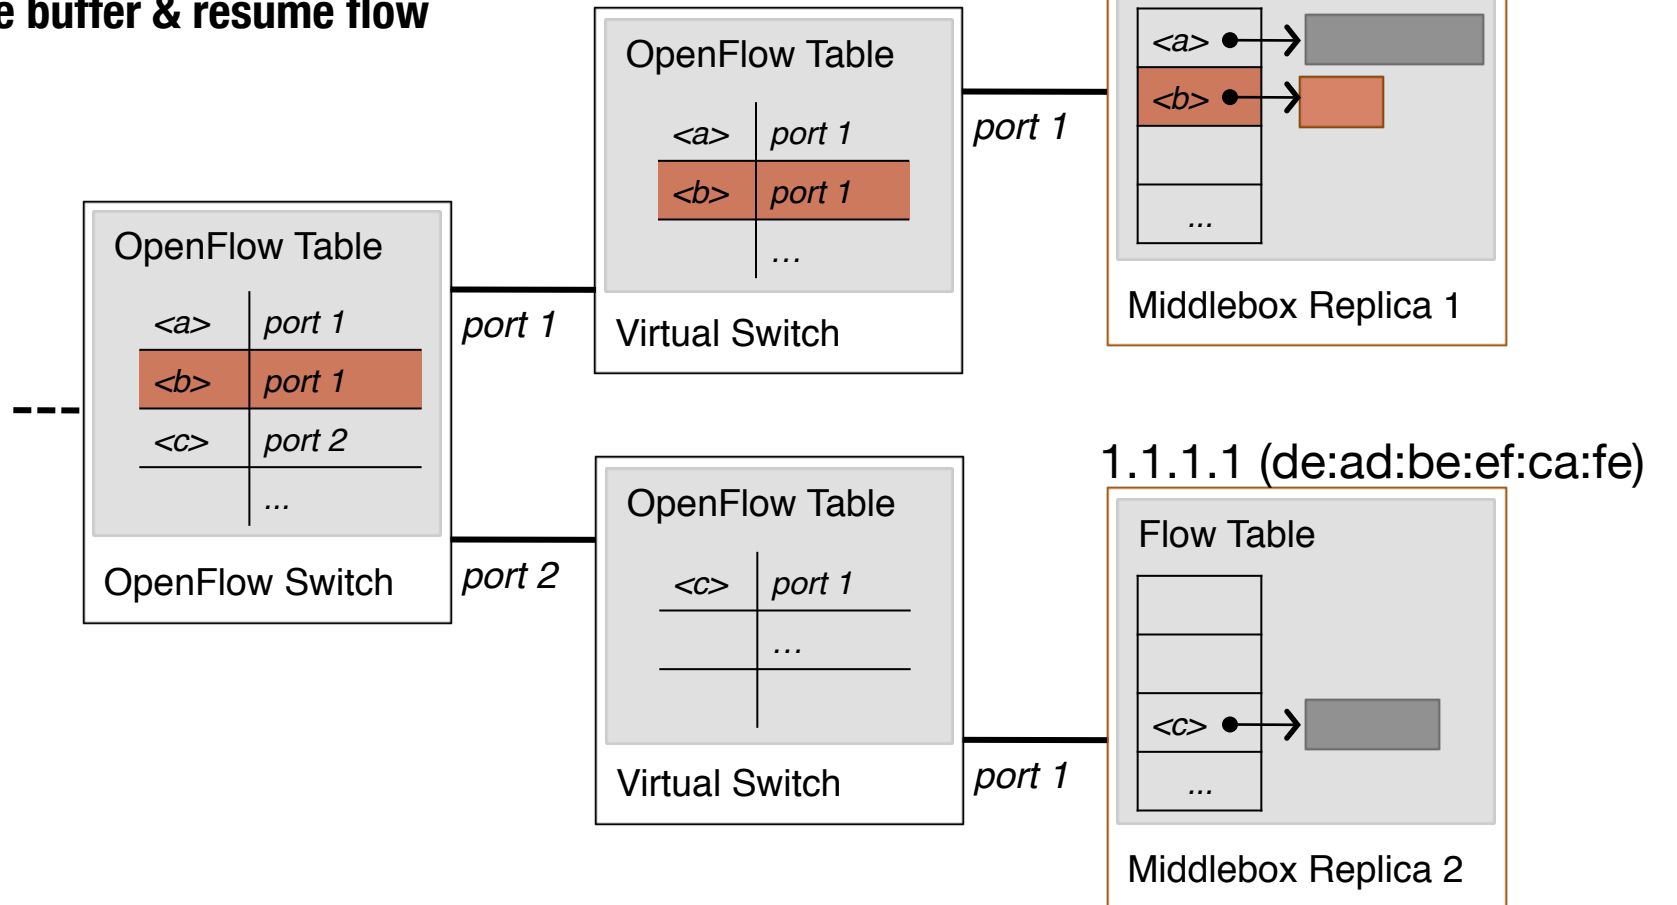
Flow Migration

Move flow state to target



Flow Migration

Release buffer & resume flow



Managing Coherent State

```
create_shared(key, size, cb)
```

```
delete_shared(key)
```

```
state get_shared(key, flags) // synch | pull | local
```

```
put_shared(key, flags) // synch | push | local
```

Managing Coherent State

Strong Consistency

Distributed lock
for every update

```
create_shared("foo", 4, NULL)
while (1)
    process_packet()
    p_foo = get_shared("foo", synch)
    val = (*p_foo)++
    put_shared("foo", synch)

    if (val > threshold)
        bar()
```

Middlebox applications rarely need strong consistency!

Managing Coherent State

Eventual Consistency

```

create_shared("foo", 4, merge_fn)
while (1)
    process_packet()
    p_foo = get_shared("foo", local)
    val = (*p_foo)++
    put_shared("foo", local)

    if (val > threshold)
        bar()
        put_shared("foo", push)
    
```

Hi frequency
local updates

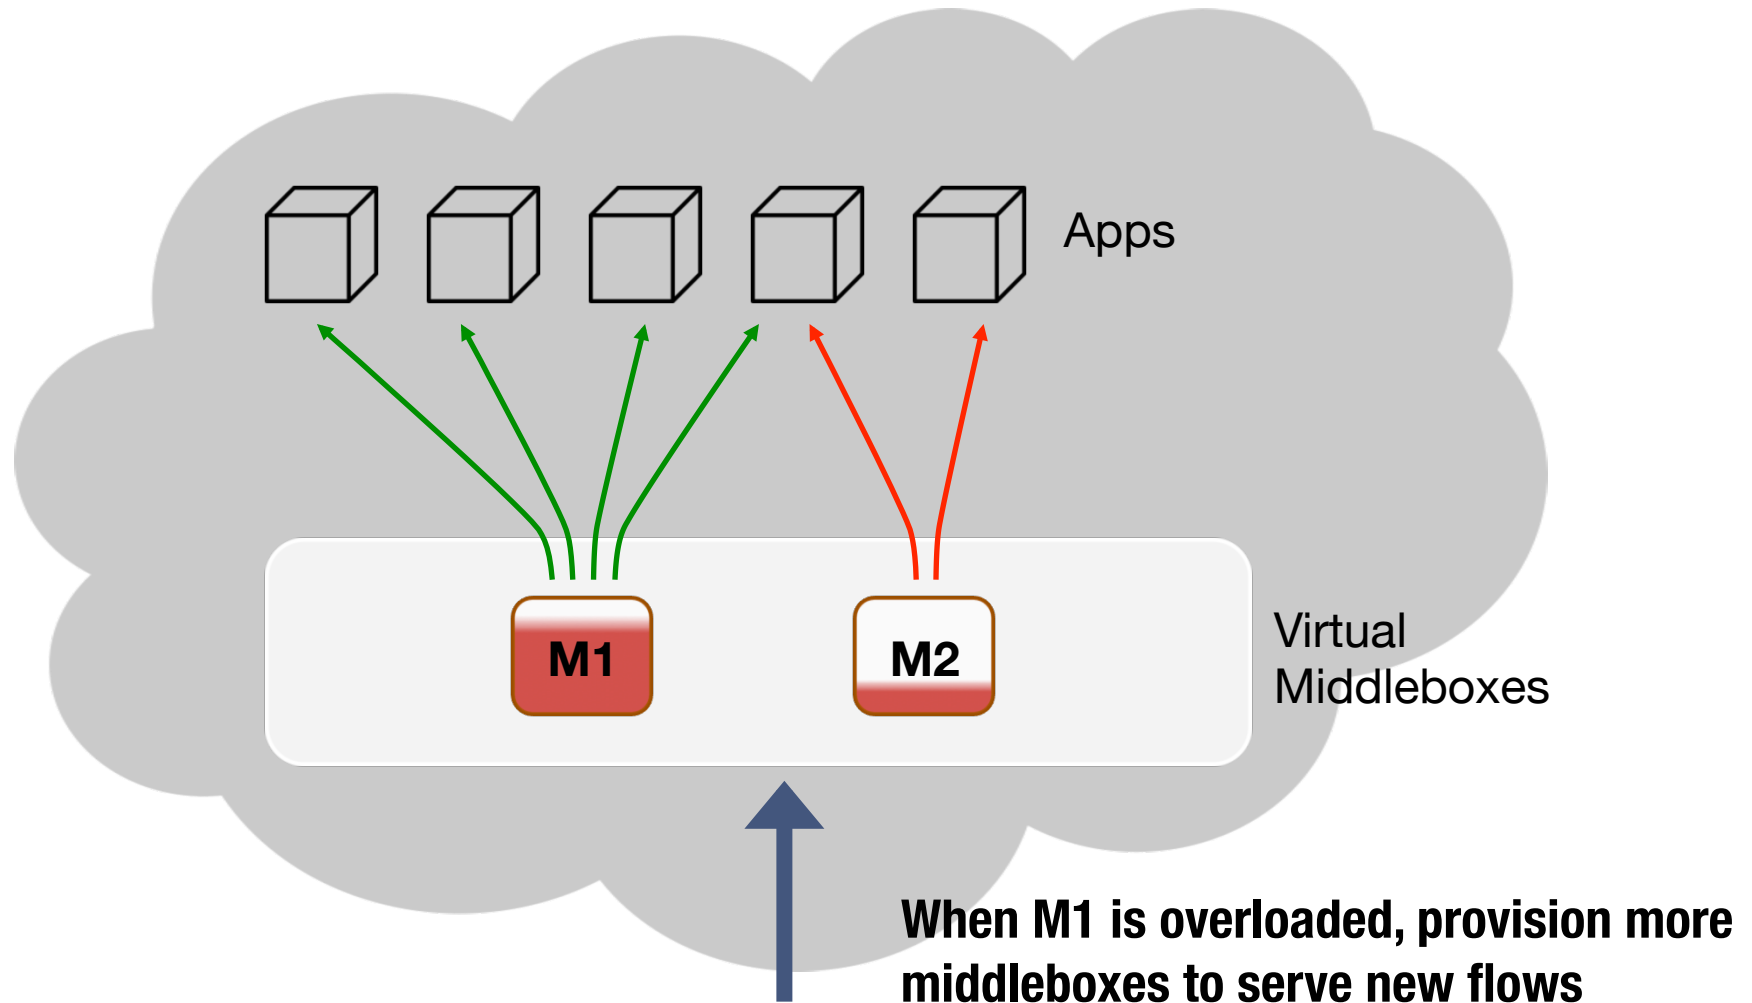
Periodic global
updates

Evaluation


Evaluation Overview

- Eliminating hotspots during scale-out
- Fast and efficient scale-in
- Split/Merge Bro during a load burst


Hotspots Cannot be Alleviated Quickly




Eliminating Hotspots by Shedding Load




Eliminating Hotspots by Shedding Load



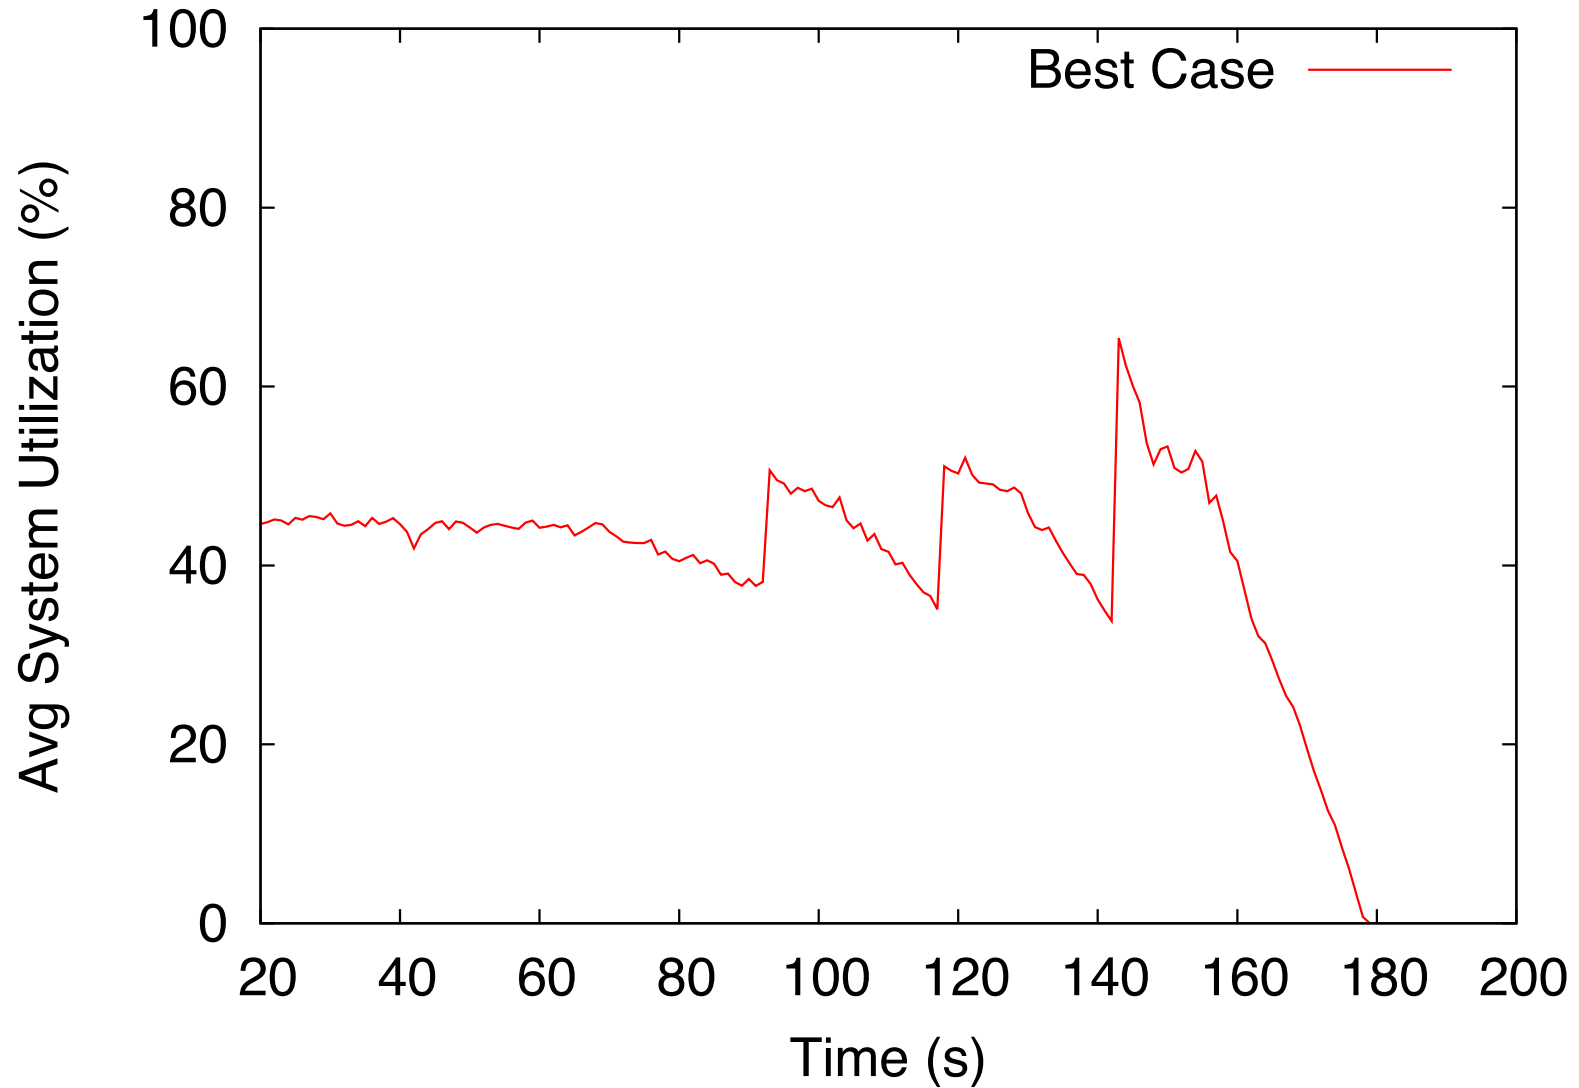
Scaling-In a Deployment : Best Case Scenario




Scaling-In a Deployment : Best Case Scenario




Scaling-In: Best Case Scenario



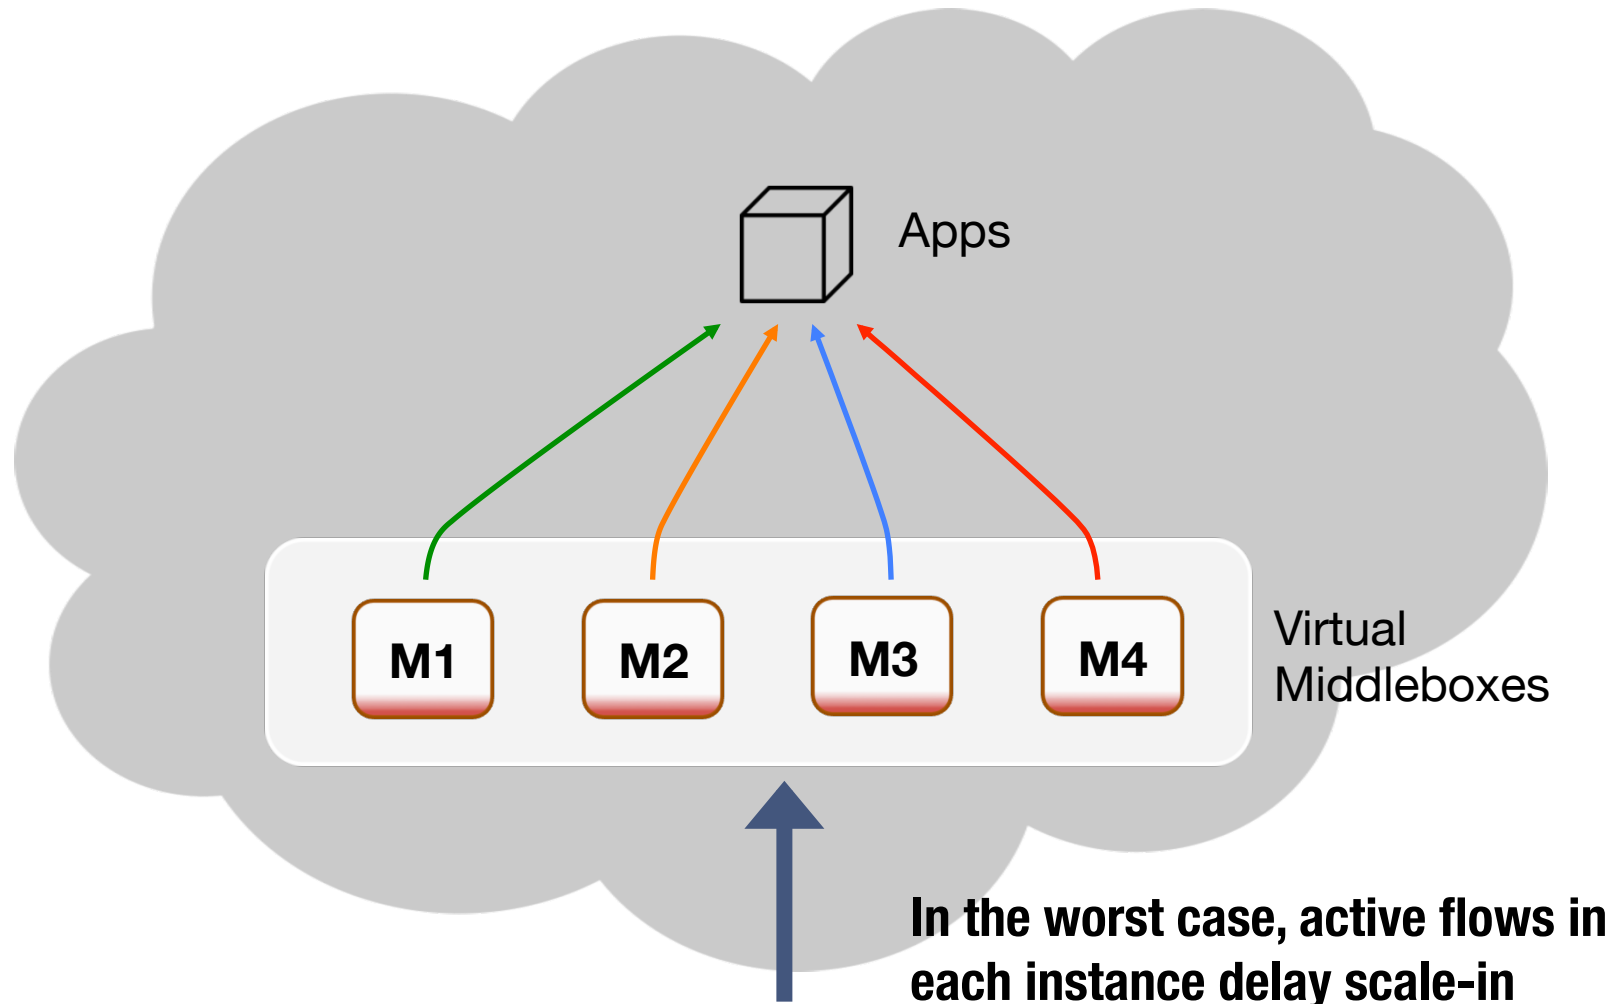
Scaling-In using *kill* : Worst Case Scenario




Scaling-In using *kill* : Worst Case Scenario




Scaling-In using *kill* : Worst Case Scenario




Scaling-In using *kill* : Slow & Inefficient




Scaling-In using *merge* : Worst Case Scenario



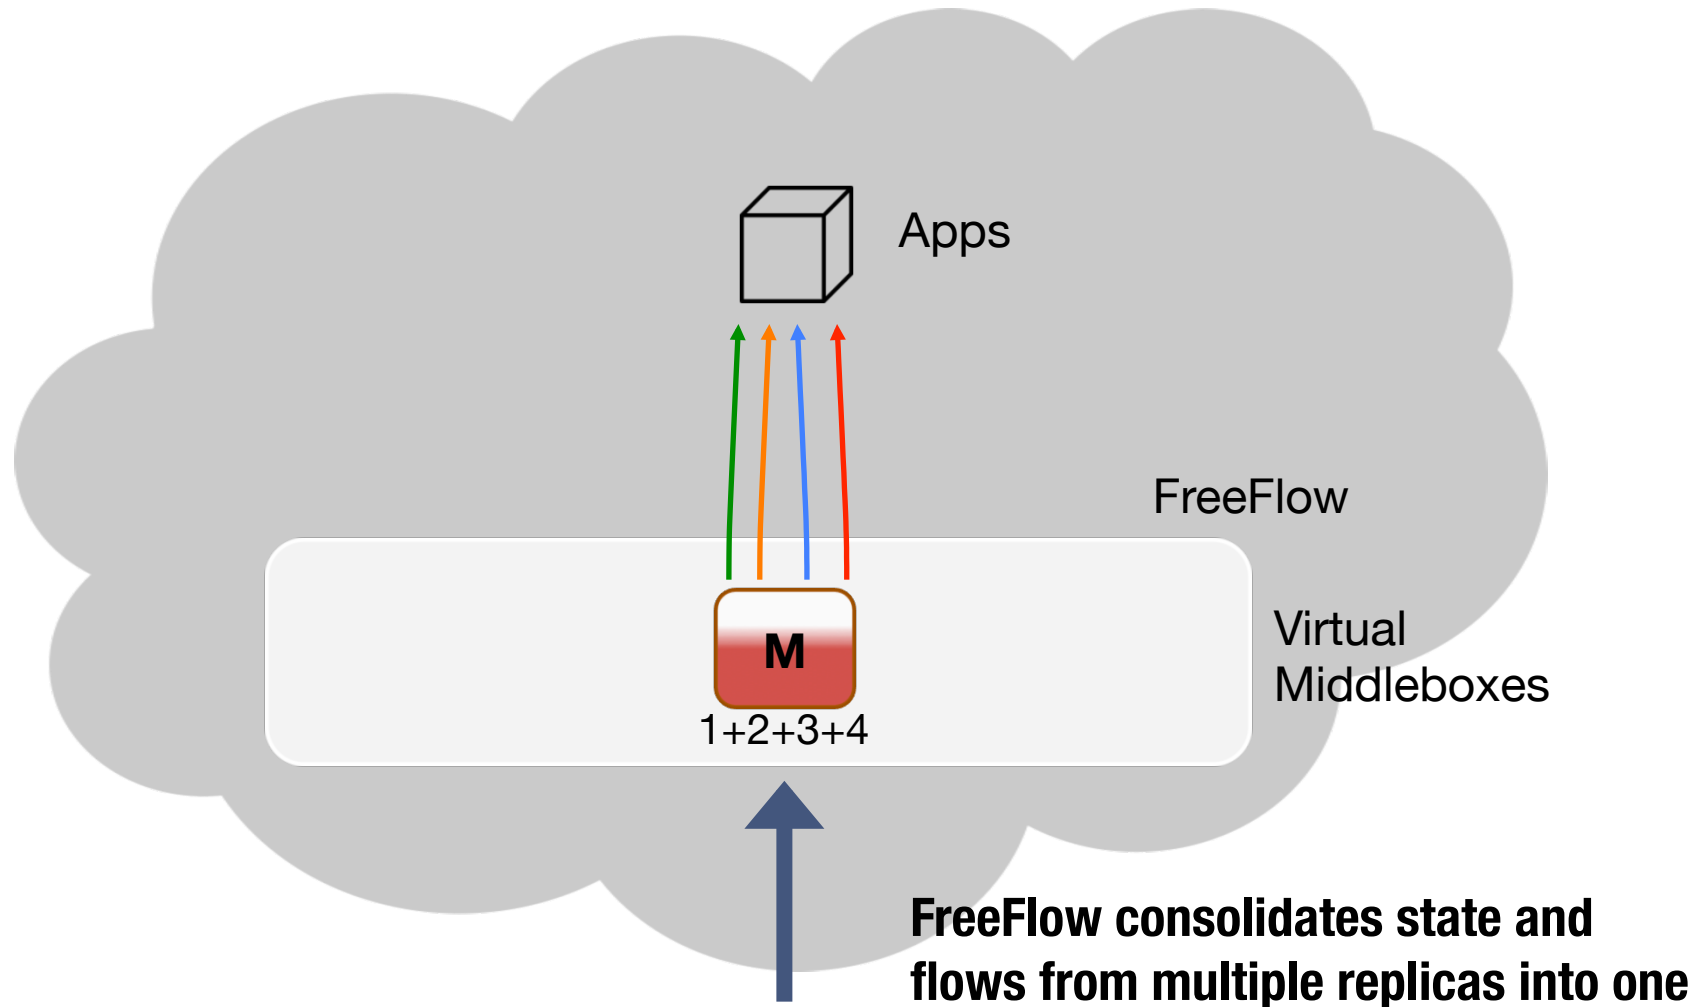
Scaling-In using *merge* : Worst Case Scenario




Scaling-In using *merge* : Worst Case Scenario



Scaling-In using *merge* : Worst Case Scenario




Scaling-In using *merge* : Fast & Efficient





Splitting & Merging Bro IDS



- Ported the Event Engine to FreeFlow
- Support for UDP, TCP/ HTTP protocols
- SQL Injection Detection plugin



Handling a Load Burst




Handling a Load Burst : No Scaling




Without enough capacity to handle the load burst, the system performance degrades severely


Handling a Load Burst : Pre-Scaled




Two instances are provisioned apriori, enough to handle a load burst, if any




Handling a Load Burst : Pre-Scaled




Load burst has no impact on system performance, as there is enough capacity to handle the load





Handling a Load Burst : Split/Merge



One replica handles the load well,
before the load burst




Handling a Load Burst : Split/Merge




When load burst starts, the Orchestrator *splits* the replica and rebalances the load



Handling a Load Burst : Split/Merge



With the load rebalanced,
performance returns to normal





Handling a Load Burst : Split/Merge

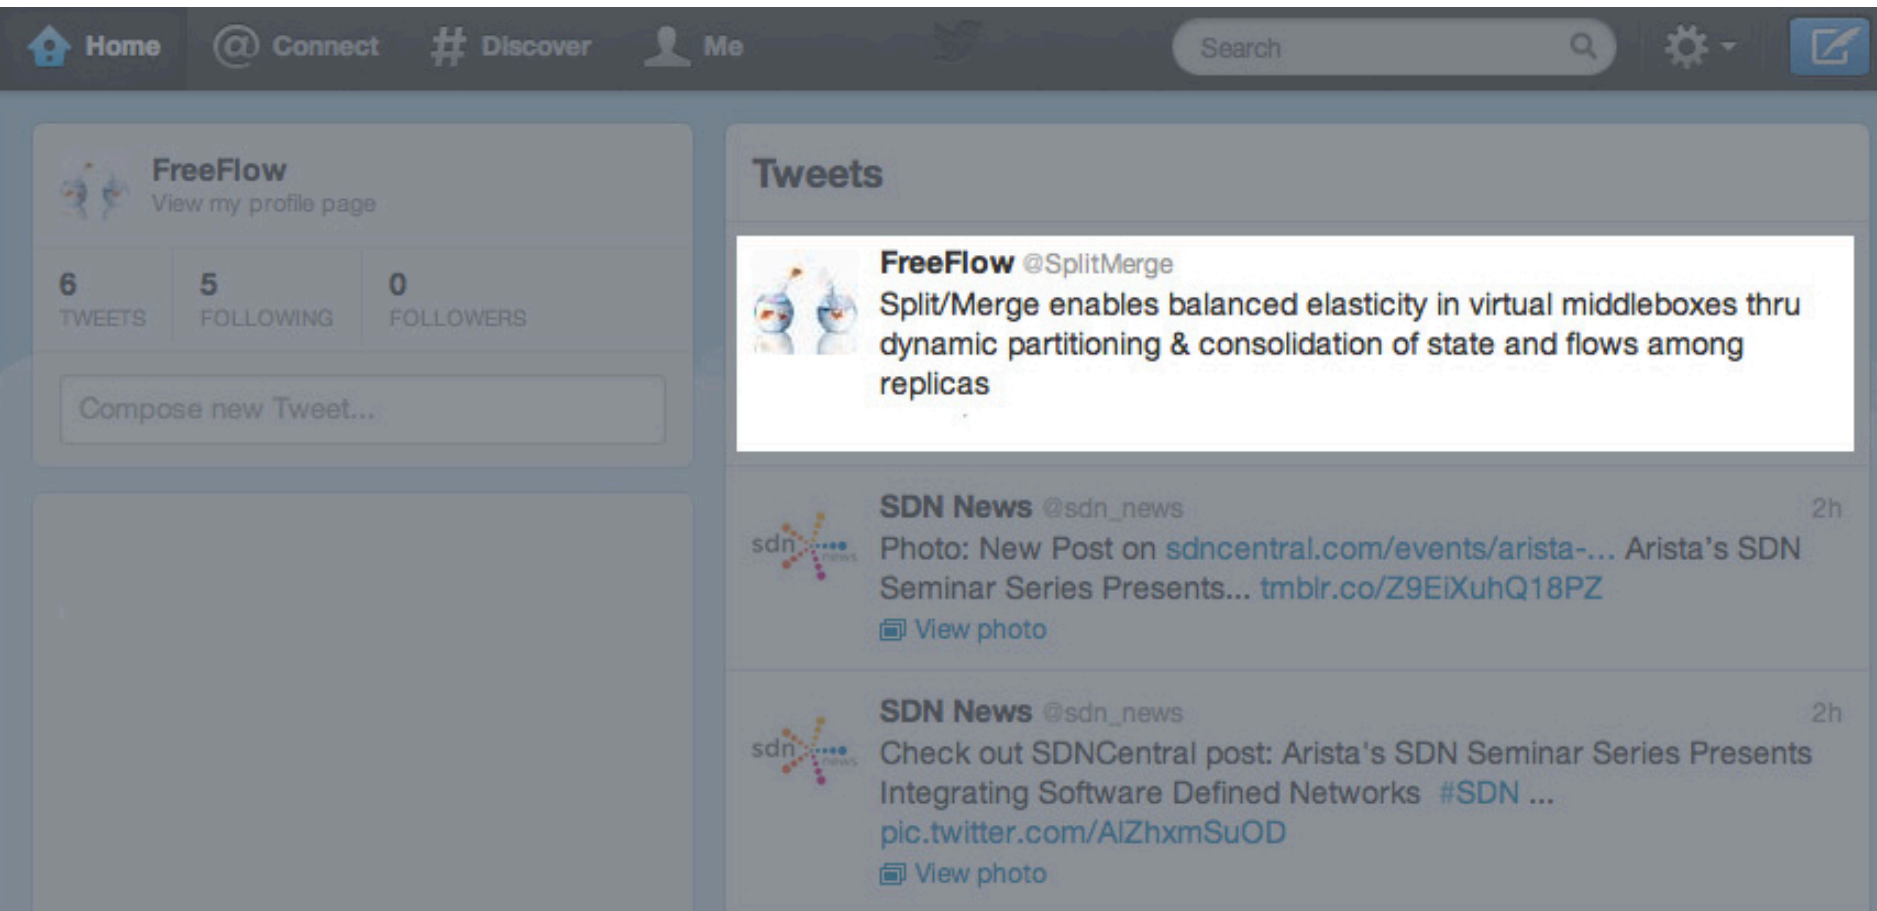


When system utilization drops after the load burst, the Orchestrator *merges* the two replicas

Handling a Load Burst : Split/Merge



Summary



FreeFlow
View my profile page

6 TWEETS 5 FOLLOWING 0 FOLLOWERS

Compose new Tweet...


Tweets

FreeFlow @SplitMerge
Split/Merge enables balanced elasticity in virtual middleboxes thru dynamic partitioning & consolidation of state and flows among replicas

SDN News @sdn_news 2h
Photo: New Post on sdncentral.com/events/arista-... Arista's SDN Seminar Series Presents... tumblr.co/Z9EiXuhQ18PZ
[View photo](#)

SDN News @sdn_news 2h
Check out SDNCentral post: Arista's SDN Seminar Series Presents Integrating Software Defined Networks #SDN ...
pic.twitter.com/AiZhxmSuOD
[View photo](#)

Flow Migration Overhead - TCP



Flow Migration Overhead - UDP

