

# Leveraging Parallelism for Multi-dimensional Packet Classification on Software Routers



THE UNIVERSITY  
*of*  
**WISCONSIN**  
MADISON

Yadi Ma  
Shan Lu

Suman Banerjee  
Cristian Estan\*

Presenter: Chang Chen  
Mar 19, 2014

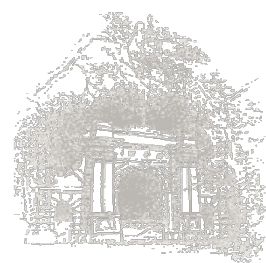




# Question



- *If we were to implement classification using state-of-art desktops, given that it may find applications in router design, then what classification speeds can be achieved on it, by designing a fully software-based classification system that can exploit the degree of parallelism provided by this particular platform?*

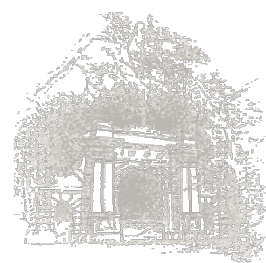




# Outline



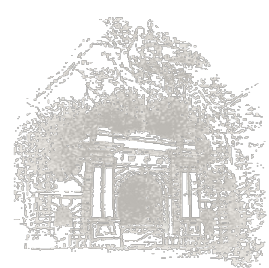
- Introduction
- Overview of Storm
- Storm Design
  - Thread Assignment
  - Rule Cache Updater
  - Micro Benchmarks
- Experimental Results





# Introduction

- **Storm:** a software-based solution to the multi-dimensional packet classification problem.
- A new software system which
  - takes existing classification algorithms;
  - utilizes a common idea of caching;
  - partitions critical tasks into multiple threads that effectively leverage desktop platforms to meet various computation and memory access needs.





# Prior Work

Q. Dong, et al. Wire Speed Packet Classification Without TCAMs: A Few More Registers (And A Bit of Logic) Are Enough, in ACM SIGMETRICS, 2007.

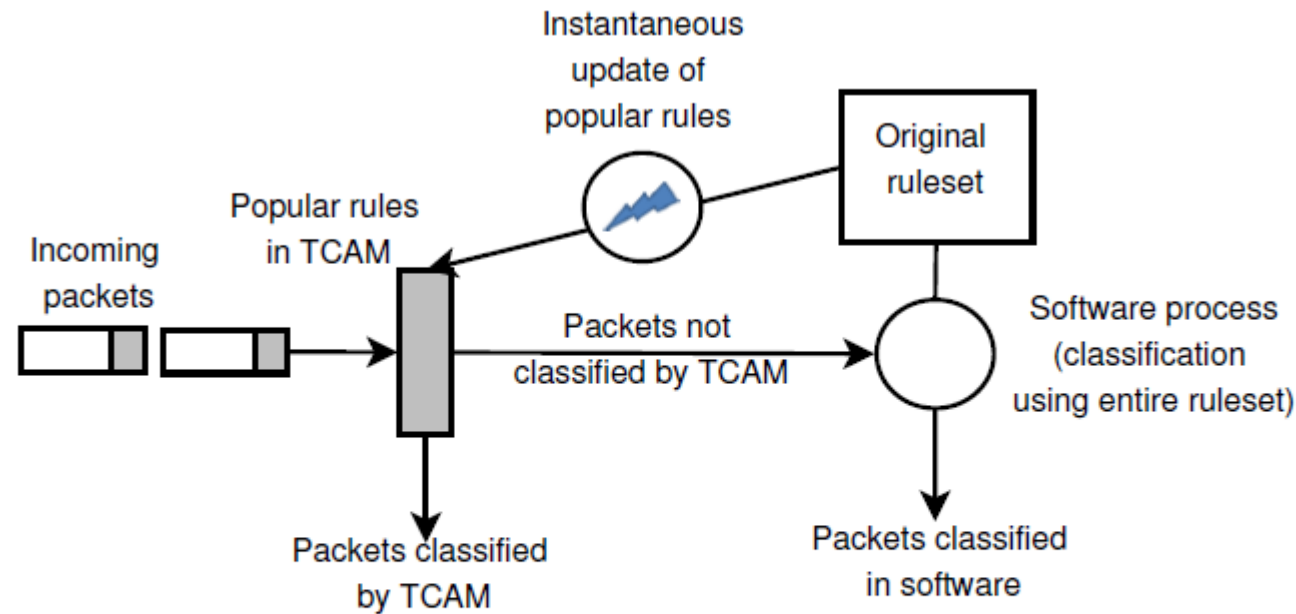
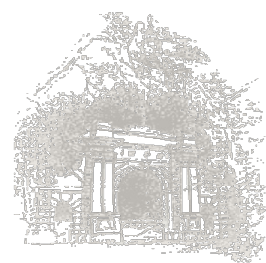


Figure 1: Framework of SRC, its use of TCAMs, and simplifying assumptions.



# Overview of Storm

- A parallelized software system using multi-core desktop platforms.
- Uses a combination of task, data and pipeline parallelism.

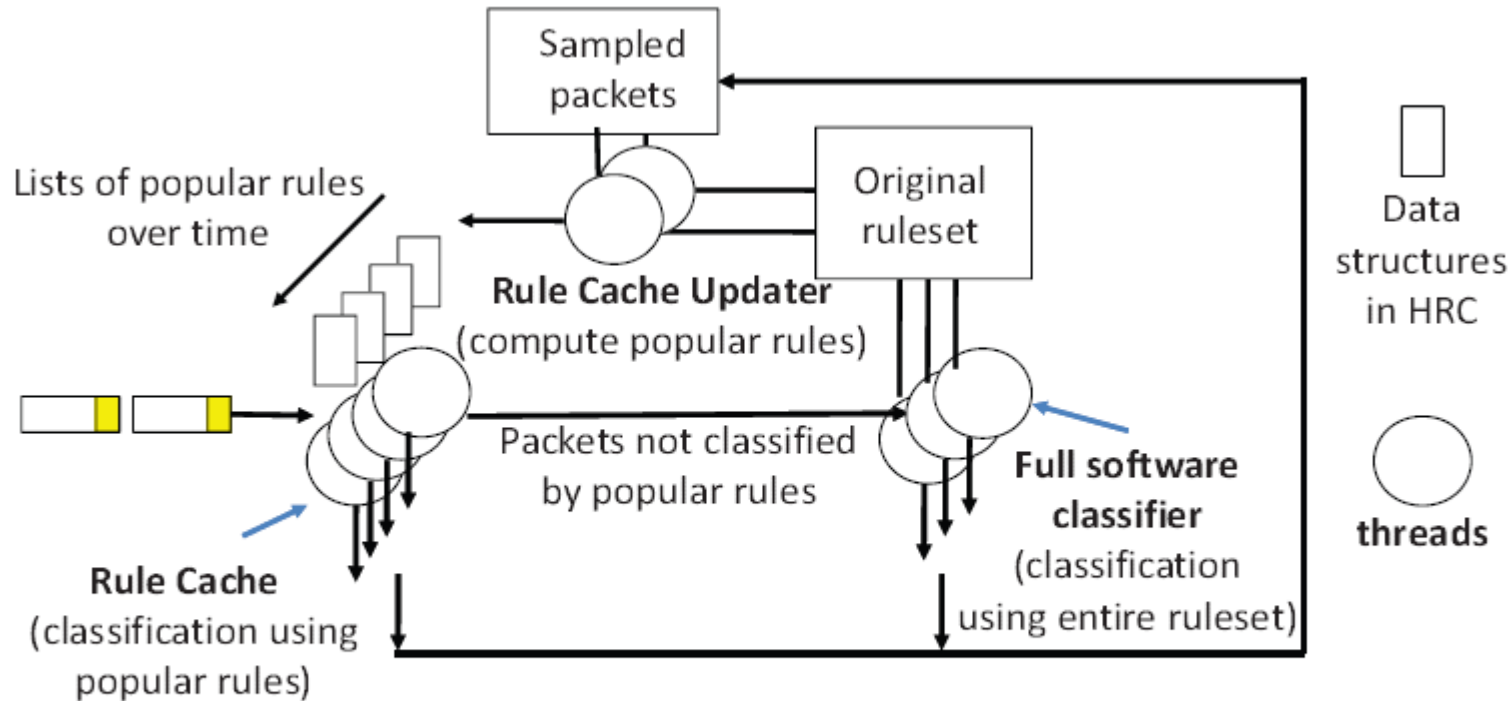
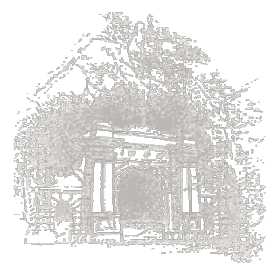


Figure 2: Architecture of Storm.

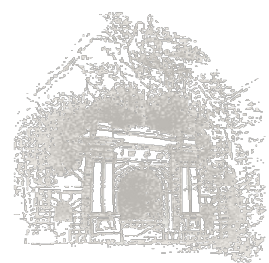




# Overview of Storm



- Task Parallelism
  - Rule cache threads
    - ✓ Match incoming packets against cached rules.
  - Full software classifier threads
    - ✓ Use the entire rule set to carry out classification (with HyperCuts algorithm).
  - Rule cache updater threads
    - ✓ Continuously sample incoming traffic, identify the evolving set of popular rules and update the rule caches.







# Overview of Storm

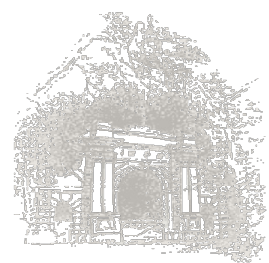


- Data Parallelism

- By creating multiple instances of each thread and allowing them to operate on different packets.

- Pipeline Parallelism

- Multiple tasks executed in a specific pre-defined order for each incoming packet.
- Following a producer-consumer pattern between the two classification stages.



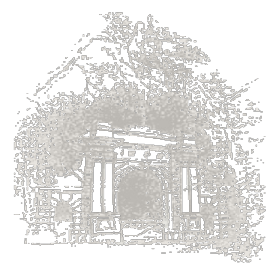




# Overview of Storm



- Main contributions of Storm:
  - ❑ A practical, multi-threaded, software-only packet classification system.
  - ❑ Design of dynamic balancing of computation resources to tasks.

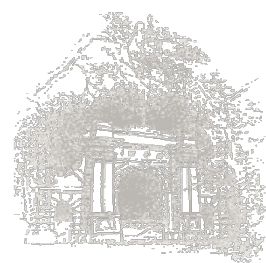




# Outline



- Introduction
- Overview of Storm
- Storm Design
  - Thread Assignment
  - Rule Cache Updater
  - Micro Benchmarks
- Experimental Results





# Thread Assignment

- Three different tasks threads assignment
  - Rule cache lookup
  - Full software classification
  - Rule cache updating
- Theoretical model

$$\frac{1}{T} = \frac{d_1}{t_1} \times r + \frac{d_2}{t_2} \times (1 - r) + C$$

T: throughput

$d_1$ : average delay of rule cache lookup

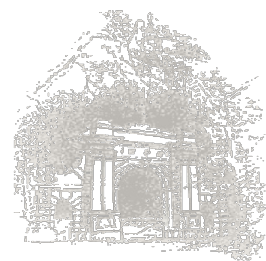
$d_2$ : average delay of full software classification

$t_1$ : number of rule cache threads

$t_2$ : number of full software classifier threads

$r$ : average rule cache hit ratio

$C$ : const (queuing delay, synchronization overhead, etc.)





# Thread Assignment

- Theoretical model

$$\frac{\partial \frac{1}{T}}{\partial t_1} = -\frac{d_1 r}{t_1^2} + \frac{d_2(1-r)}{(N-t_1)^2} = 0$$

$d_1$ : average delay of rule cache lookup

$d_2$ : average delay of full software classification

$t_1$ : number of rule cache threads

$t_2$ : number of full software classifier threads

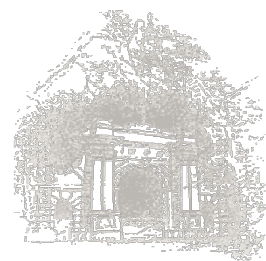
$N$ :  $t_1 + t_2$

$r$ : average rule cache hit ratio



$$t_1 = N \frac{\sqrt{d_1 r}}{\sqrt{d_2(1-r)} + \sqrt{d_1 r}}$$

- ▣ Use only one thread to carry out the sampling and rule cache updating task.





# Thread Assignment

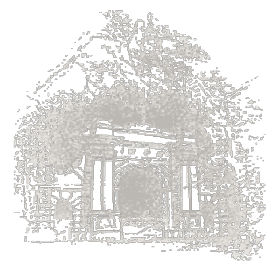
- Static thread assignment (on 8-core machines)

$$t_1 = N \frac{\sqrt{d_1 r}}{\sqrt{d_2(1-r)} + \sqrt{d_1 r}}$$

- Control total number of thread around 8
- $d_1$ : about 200 ns (rule cache lookup)
- $d_2$ : about 2000 ns (full classification using HyperCuts)
- $r$ : around 0.95
- minimize synchronization overhead



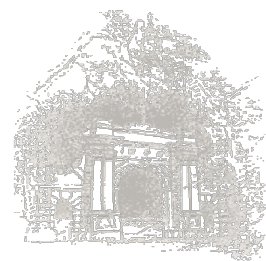
- Potential thread partitions: 3-3-1, 4-2-1, 2-4-1





# Thread Assignment

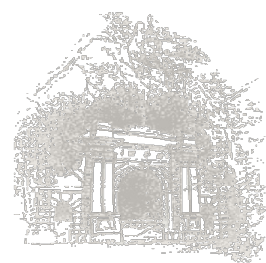
- Dynamic thread assignment (on 8-core machines)
  - ▣ Initially create 4 rule cache threads and 4 full software classification threads.
  - ▣ Define the total buffer size in between to be  $B$ ; define two thresholds,  $r1$  and  $r2$ , where  $0 < r1 < r2 < 1$ .
  - ▣ Depending on the available buffer size, switch between 4-2-1 (if  $< r1 \times B$ ) and 2-4-1 ( $> r2 \times B$ ).





# Rule Cache Updater

- A rule cache updater thread is responsible for periodically sampling the incoming packets and updating the rules in the rule cache.
- Each entry in the rule cache stores an **evolving rule**.







# Rule Cache Updater



Table 3: A simple ruleset on 2 fields.

Rule	Field <sub>1</sub>	Field <sub>2</sub>	Action
R0	1-9	4-10	action <sub>0</sub>
R1	7-14	3-8	action <sub>1</sub>
R2	3-11	1-6	action <sub>2</sub>
R3	0-15	0-11	action <sub>3</sub>

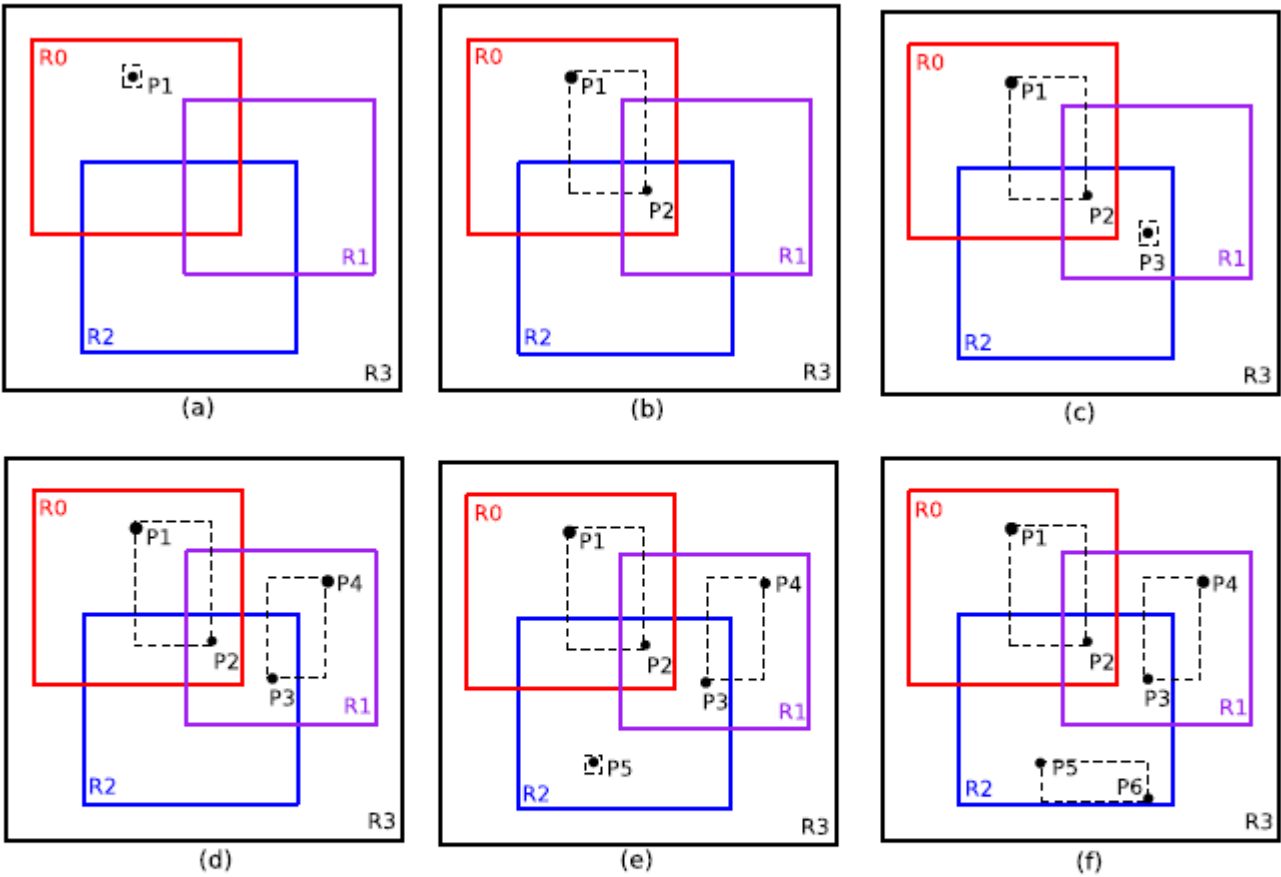
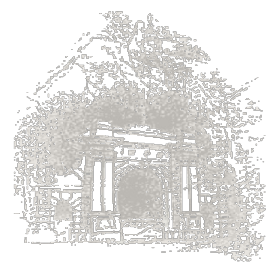


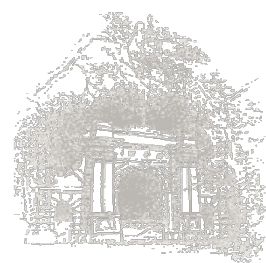
Figure 4: Constructing evolving rules, rule priority:  $R0 > R1 > R2 > R3$ ,  $action_0 \neq action_1 \neq action_2 \neq action_3$ .





# Rule Cache Updater

- Sliding window
  - ▣ A FIFO queue that stores recently sampled packets.
- Evolving rule list
  - ▣ A list of proposed evolving rules waiting to be checked for conflicts and to be transferred into rule cache;
  - ▣ Each evolving rule includes a weight field;

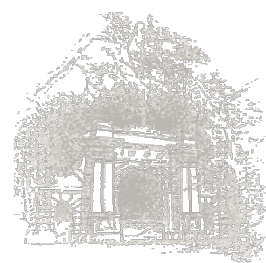




# Rule Cache Updater



- Five properties of evolving rules:
  - ❑ Represents a d-dimensional hypercube.
  - ❑ Be associated with a single action consistent with the original ruleset.
  - ❑ Each sample packet in the sliding window is assigned to one evolving rule that matches it.
  - ❑ Evolving rules either have the same action or are non-overlapping.
  - ❑ Lies entirely inside one of the rules in the original ruleset.





# Rule Cache Updater



- Check conflicts
  - Start with root
  - Runs recursively on overlapping child node until a leaf node is identified
  - Check each rule in the leaf node is a match or a conflict

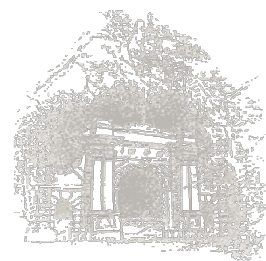
---

Algorithm 1 CheckConflict(Cnode, ExpandedRule, MatchID, ConflictID)

---

```
1: if Cnode is a leaf node then
2:   for each rule  $r$  in the rule list of Cnode do
3:     if  $r.ID > \min(MatchID, ConflictID)$  then
4:       return
5:     end if
6:     if  $r$  overlaps with ExpandedRule then
7:       if  $r.action \neq ExpandedRule.action$  then
8:          $ConflictID = r.ID$ 
9:         return
10:      end if
11:      if ExpandedRule lies entirely inside  $r$  then
12:         $MatchID = r.ID$ 
13:        return
14:      end if
15:    end if
16:  end for
17: end if
18: if Cnode is not a leaf node then
19:   for each child  $c$  of Cnode that overlaps with ExpandedRule do
20:     CheckConflict( $c$ , ExpandedRule, MatchID, ConflictID);
21:   end for
22: end if
```

---





- Check conflicts

Table 1: A simple example with 8 rules on 5 fields

Rule	$F_1$	$F_2$	$F_3$	$F_4$	$F_5$	Action
$R_0$	000*	111*	10	*	UDP	$action_0$
$R_1$	000*	10*	01	10	TCP	$action_1$
$R_2$	000*	01*	*	11	TCP	$action_0$
$R_3$	0*	1*	*	01	UDP	$action_2$
$R_4$	0*	0*	10	*	UDP	$action_1$
$R_5$	000*	0*	*	01	UDP	$action_1$
$R_6$	*	*	*	*	UDP	$action_3$
$R_7$	*	*	*	*	TCP	$action_4$

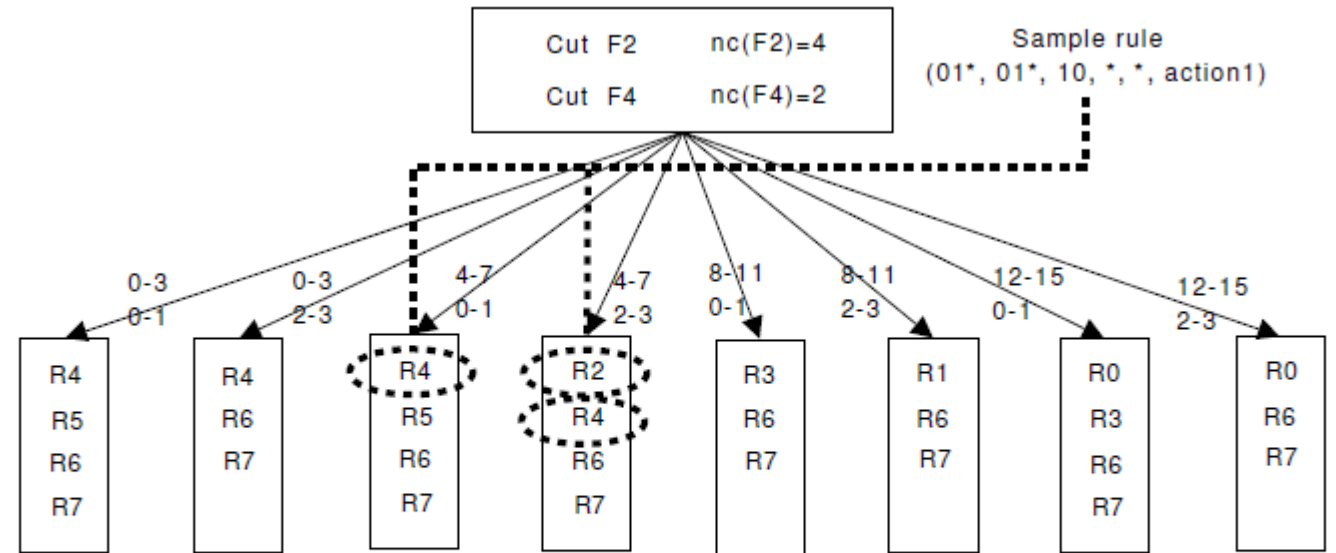
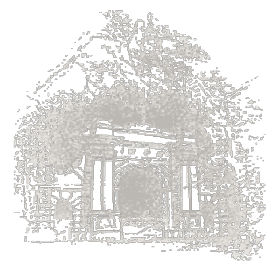


Figure 5: Check an expanded rule for conflicts using HyperCuts decision tree.



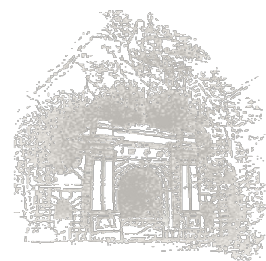


# Rule Cache Lookup

- How to lookup rules in rule cache?
  - ▣ Use linear search to search rules in cache to find a match or cache miss.
- What the rule cache size should be?

Table 5: Performance with different cache sizes on ruleset R3

Cache size	Delay(ns/p)	Hit ratio(%)
10	95.21	93.92
15	93.27	95.31
20	82.01	96.55
25	83.28	95.92
30	96.52	96.52



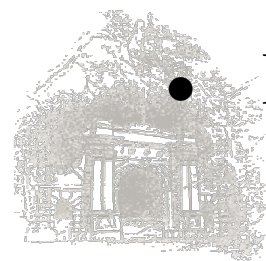


# Experimental Results

- Rulesets

Rule set	Type	Size
R1	real	460
R2	real	711
R3	real	852
R4	real	1036
R5	real	1802
R6	synthetic	4415
R7	synthetic	9603

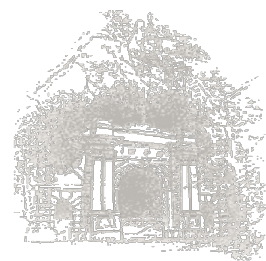
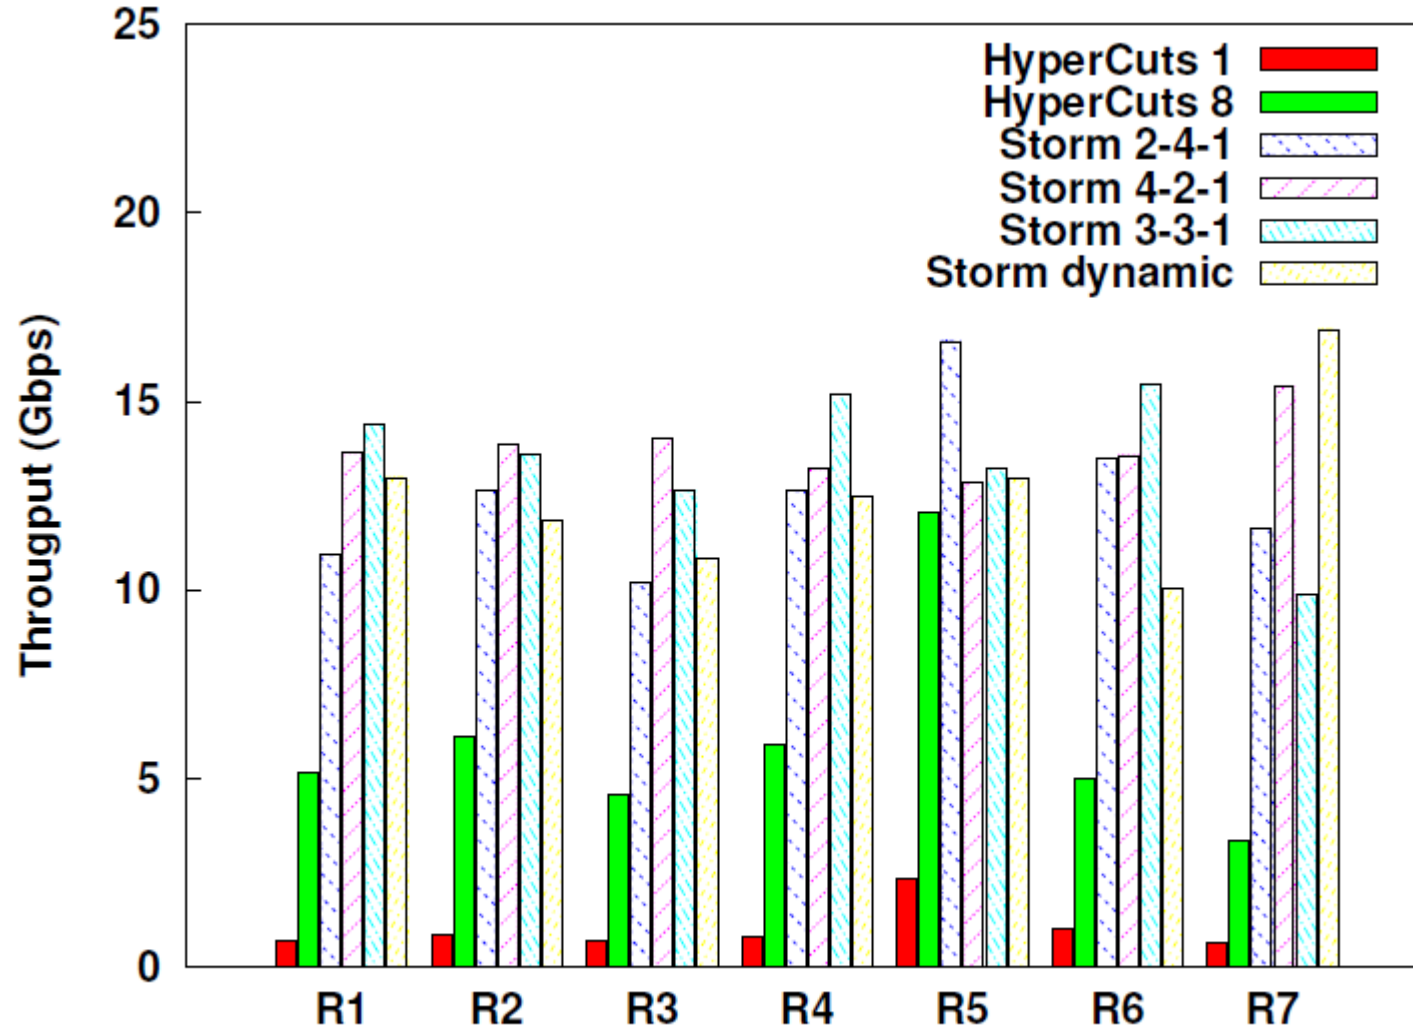
- 5 tuples; action is either permit or deny.
- Traces
  - Five traces each contain about 7M packets (packet size: 128 bytes).
  - Different incoming traffic rates.





# Experimental Results

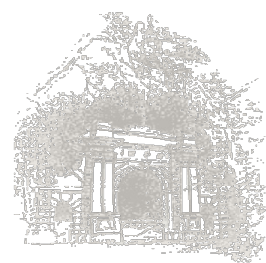
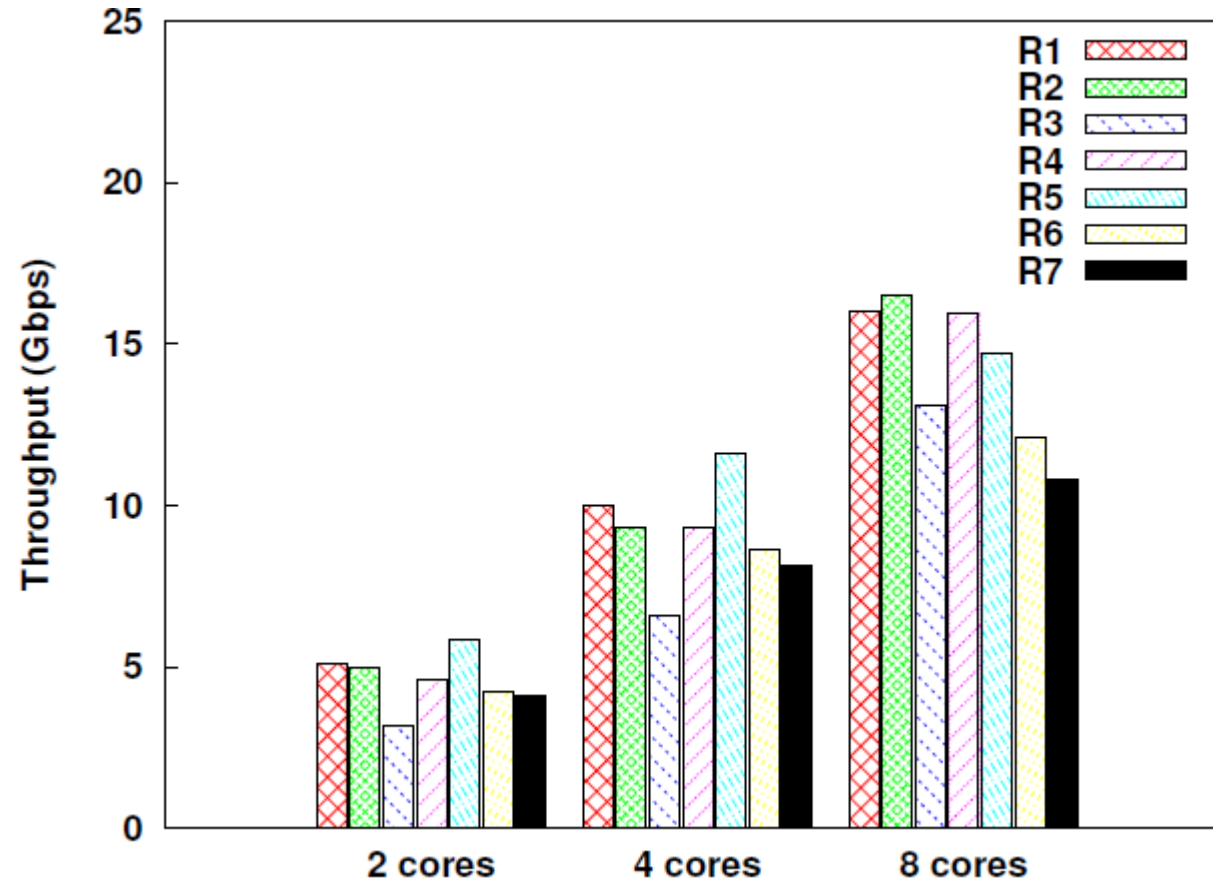
- Maximum achievable throughput





# Experimental Results

- Scalability with number of cores





- Cache hit ratio

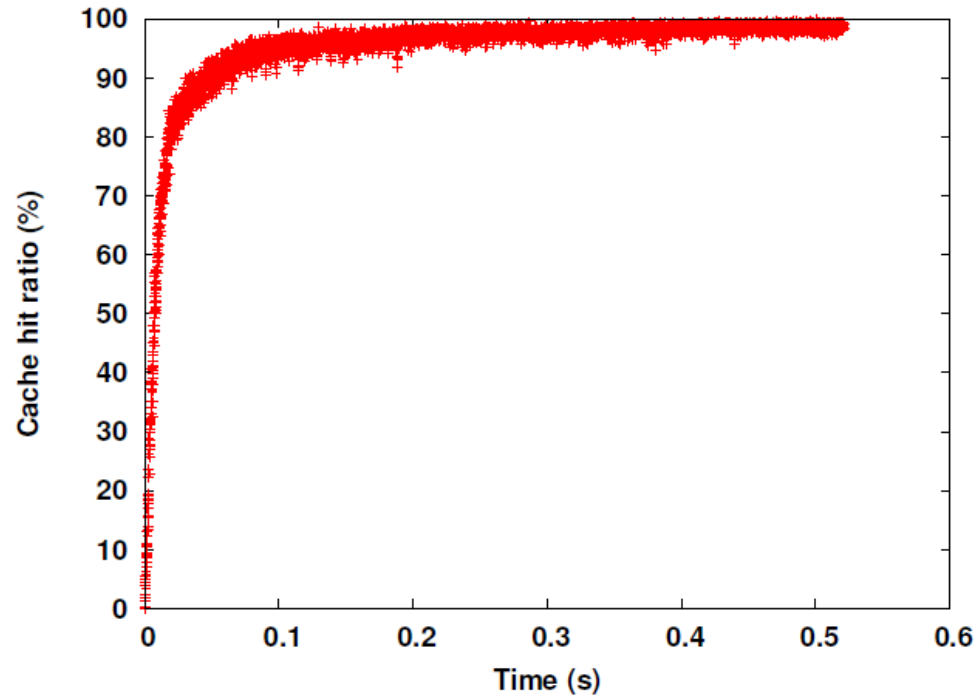


Figure 8: Storm's rule cache hit ratio ramps up quickly (example uses ruleset R3).

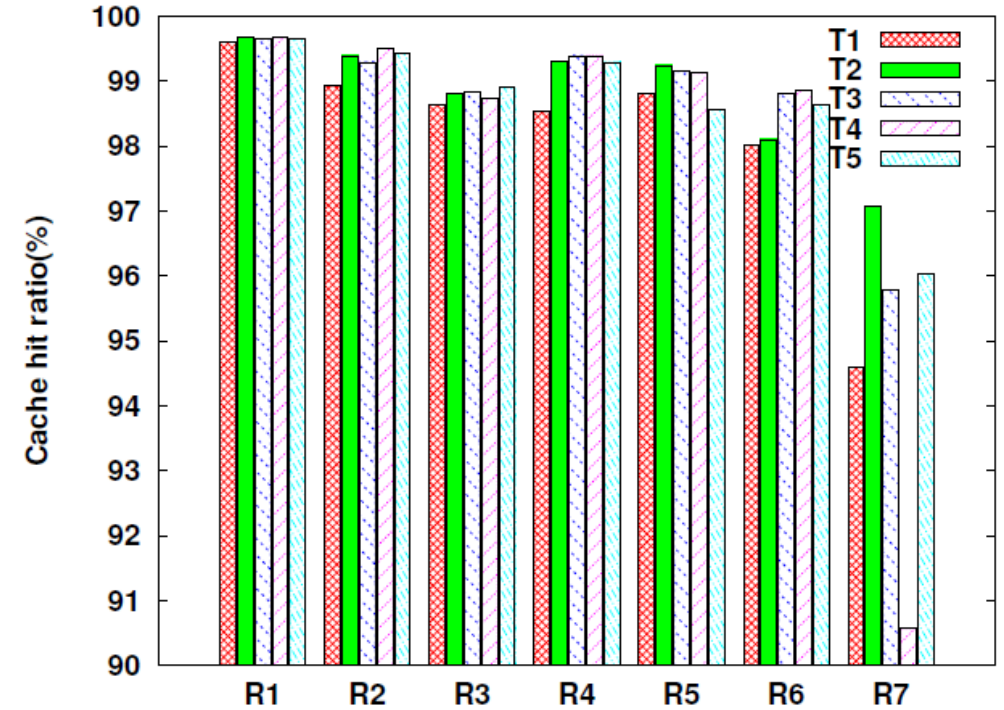


Figure 9: Cumulative rule cache hit ratios of rulesets R1 through R7 with traces T1 through T5 (Y-axis starts at 90%).



# Experimental Results



- Micro benchmarks

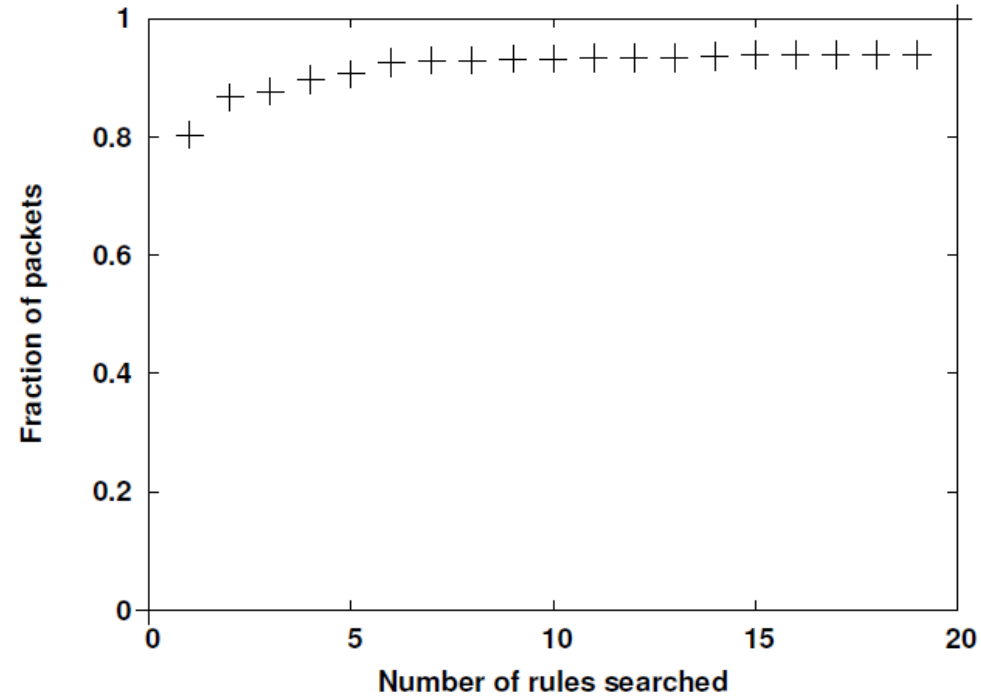
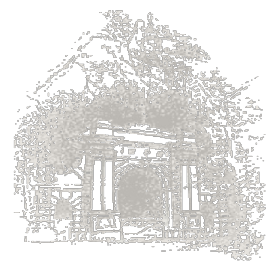


Figure 10: Cumulative distribution of number of entries searched in Storm's rule cache for 1000 randomly sampled packets (using ruleset R2).





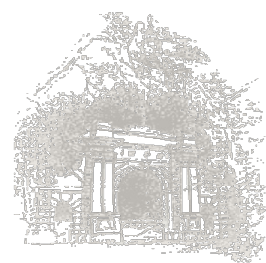
# Experimental Results



- Micro benchmarks

Table 6: Warmup time (in ms) of the seven rulesets using thread partition 2-4-1 and 4-2-1 for Storm.

Rule set	Storm 2-4-1	Storm 4-2-1
R1	3 ms	5 ms
R2	2 ms	4 ms
R3	10 ms	13 ms
R4	6 ms	12 ms
R5	1 ms	1 ms
R6	8 ms	26 ms
R7	36 ms	79 ms

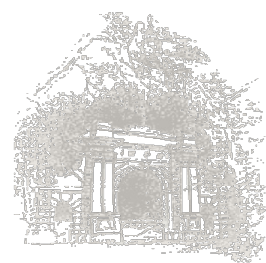




# Conclusions



- **Storm:** a software-based solution to the multi-dimensional packet classification problem.
- A new software system which
  - ▣ takes existing classification algorithms;
  - ▣ utilizes a common idea of caching;
  - ▣ partitions critical tasks into multiple threads that effectively leverage desktop platforms to meet various computation and memory access needs.





**Thank You!**

