

Scaling Up Clustered Network Appliances with ScaleBricks

DongZhou

Bin Fan

Hyeontaek Lim

David G. Andersen

Michael Kaminsky



Focus

- Throughput Scaling

 - scale with number of servers

- FIB Scaling


 - total size of forwarding table(the number of supported keys)

- Update Scaling


 - update rate of the FIB

FIB (comprised Entries for Forwarding or Processing)


Cluster Architecture




(a) RouteBricks



(b) Full Duplication



(c) Hash Partitioning



(d) ScaleBricks


- FIB size (full or slice)
- Intermidate (Server or Switch)
- Internal Bindwidth

ScaleBricks' Design

Maintain entire routing information in **RIB**

Distribute RIB across the cluster

Generate FIB & GPT



▣ Partial FIB

Each handling node stores FIB entries that point to it
Based on prior work (leveraging CuckHashing)

▣ Global Partition Table(GPT)

Used for forward packet to handling node
Replicated to every ingress node

Must be compact

▣ RIB partition and updates



Global Partition Table

Attributes of switch-based “middle box” cluster

- ❑ Total number of nodes is typically modest
- ❑ They can handle one-side error

Key Question: How to map millions of input key to nodes?
correct

Basic idea in high level:


- ❑ Maintain a Hash function **families**
- ❑ Use **brute force** to find the suitable hash function
- ❑ Store the **indices** rather than key/value

Group & Set Separation(**SetSep**)

Global Partition Table

How to divide a group of n keys into two disjoint subsets when n is small?


Binary Separation



0

1

Hash function families



⋮


Shape



$H_i(x)$

Color

Store the index i
if all match



If no function succeeds for $i < I$, a fallback mechanism is triggered.



Global Partition Table

Why SetSep Save Space?

Optimistically assume hash function produce fully random hash values:

Probability all n keys are properly mapped is $p = (1/2)^n$

The number of tested functions is a random variable with Geometric distribution with

Entropy: $-(1-p)\log_2(1-p) - p\log_2 p \approx -\log_2 p = n$

Storing a function for binary set separation requires 1 bit per key

16bits for a group of 16 keys



Global Partition Table

Generate the hash function family


- $H_i(x) = G_1(x) + i \cdot G_2(x)$
- In Practice, only the most significant bit are used
- Construct fast but theoretically weak(lack sufficient independence) Empirically feasible

Horrible iterations for finding a hash function(2^n), how to **speed up** construction ?

Global Partition Table


Trading Space for faster Construction

- Adds an array of m bits ($m > 2$)
- intuitive thinking: more buckets, fewer collisions and increase odds of success.



Global Partition Table

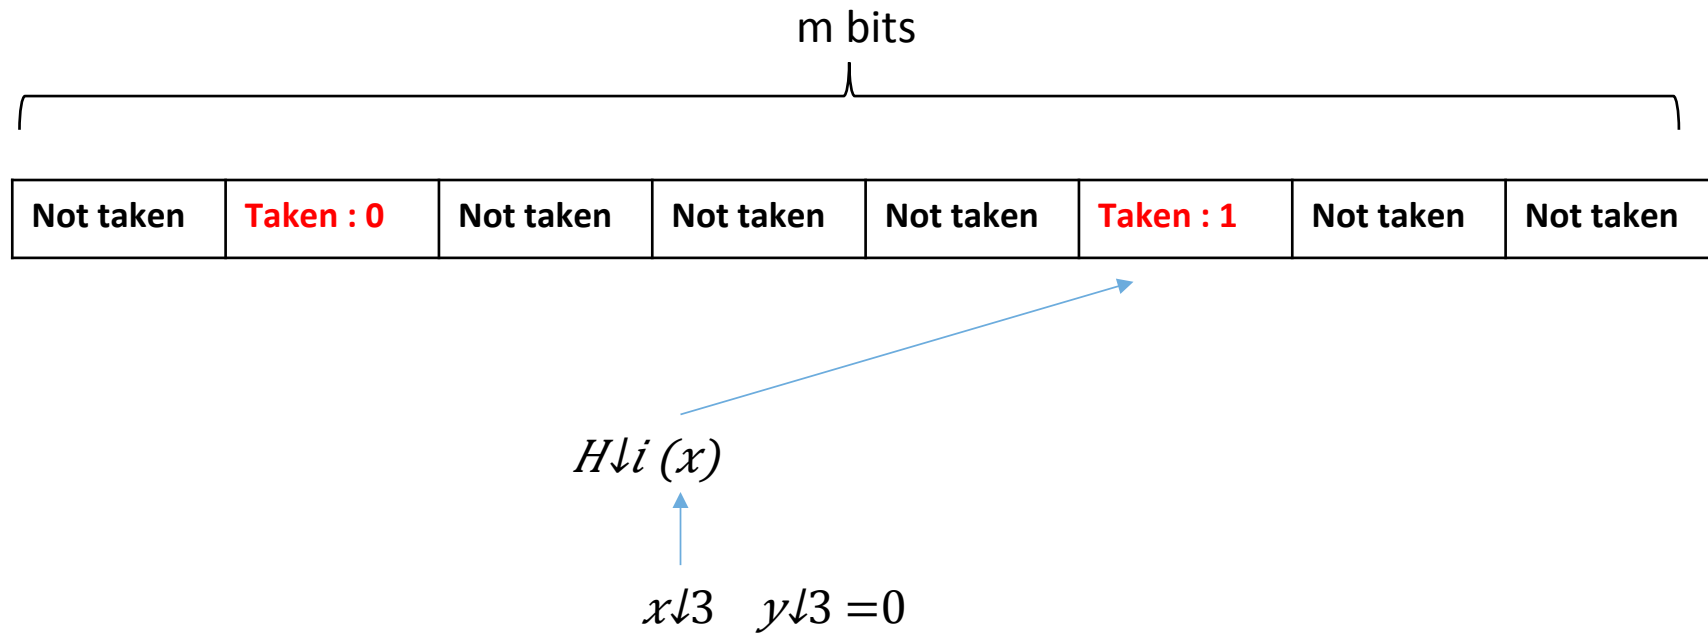
Trading Space for faster Construction



Global Partition Table

Trading Space for faster Construction


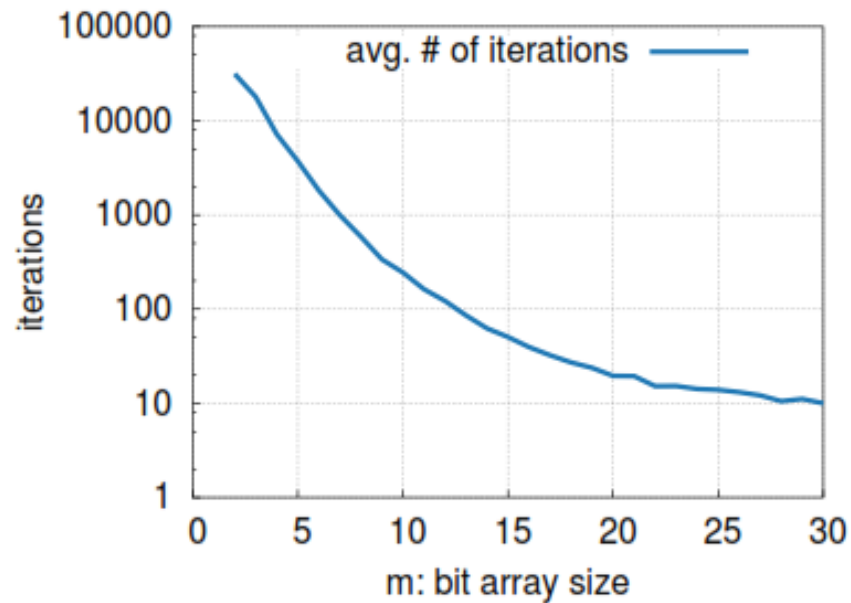
Oops! Conflicts.



Global Partition Table

Representing the **SetSep**

- Fixed 24-bit representation per group
- 16 bits represent hash index and $m=8$



Global Partition Table

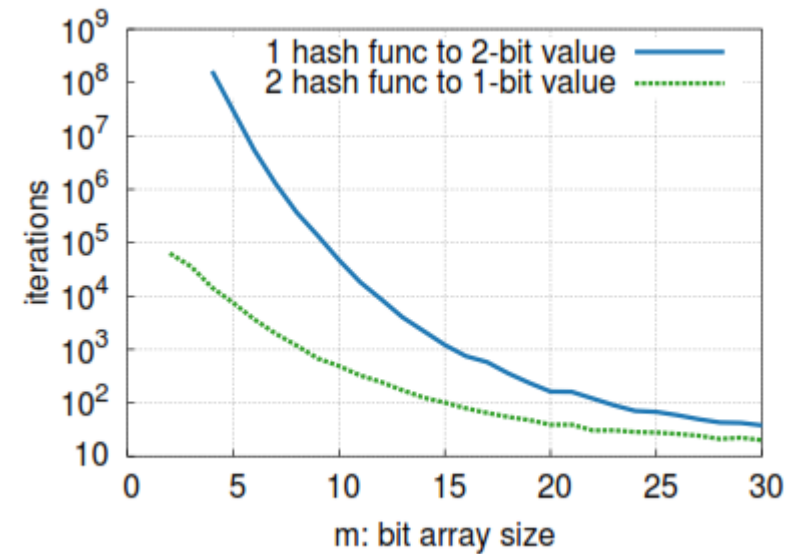
Representing the Non-Boolean values

□ j-th hash function is responsible for generating j-th bit of final mapping value.

□ Mapping value: {0,1,2,3}

“foo” maps to 1, “bar” maps to 2

Then hash function 1 maps “foo”, “bar” to 0,1 respectively; hash function 2 maps “foo”, “bar” to 1,0 respectively.




Global Partition Table

Group(how to map a key to a group)

- ❑ Low variance in group size(strong hash function/ sort and assign both failed)
- ❑ Mapping should add little space


Two level hashing



Global Partition Table

Two level hashing

- ❑ Long range of small buckets shows less variance
- ❑ Map 1024-key block to 64 groups of average size 16
- ❑ Pre-assigned “Candidate groups” & additional storage for choice
- ❑ Keys assignment is an NP-hard variant of knapsack problem



Implementation & Optimization

Global Partition Table using Setsep

- ❑ Intel DPDK
- ❑ Batched look-ups and prefetch
- ❑ Hardware Accelerate Construction
 - SIMD or GPU may help

Partial FIB using Cuckoo Hashing

Algorithm 1: Batched SetSep lookup with prefetching

```
BatchedLookup(keys[1..n])  
begin  
  for i ← 1 to n do  
    bucketID[i] ← keys[i]'s bucket ID  
    prefetch(bucketIDToGroupID[bucketID[i]])  
  for i ← 1 to n do  
    groupID[i] ← bucketIDToGroupID[bucketID[i]]  
    prefetch(groupInfoArray[groupID[i]])  
  for i ← 1 to n do  
    groupInfo ← groupInfoArray[groupID[i]]  
    values[i] ← LookupSingleKey(groupInfo, keys[i])  
  return values[1..n]
```


Evaluation

Micro-Benchmark


Construction Throughput

Construction setting			Construction throughput	Fallback ratio	Total size	Bits/key
<i>x + y bits to store a hash function, x-bit hash function index and y-bit array</i>						
16+8	1-bit value	1 thread	0.54 Mkeys/sec	0.00%	16.00 MB	2.00
8+16	1-bit value	1 thread	2.42 Mkeys/sec	1.15%	16.64 MB	2.08
16+16	1-bit value	1 thread	2.47 Mkeys/sec	0.00%	20.00 MB	2.50
<i>increasing the value size</i>						
16+8	2-bit value	1 thread	0.24 Mkeys/sec	0.00%	28.00 MB	3.50
16+8	3-bit value	1 thread	0.18 Mkeys/sec	0.00%	40.00 MB	5.00
16+8	4-bit value	1 thread	0.14 Mkeys/sec	0.00%	52.00 MB	6.50
<i>using multiple threads to generate</i>						
16+8	1-bit value	2 threads	0.93 Mkeys/sec	0.00%	16.00 MB	2.00
16+8	1-bit value	4 threads	1.56 Mkeys/sec	0.00%	16.00 MB	2.00
16+8	1-bit value	8 threads	2.28 Mkeys/sec	0.00%	16.00 MB	2.00
16+8	1-bit value	16 threads	2.97 Mkeys/sec	0.00%	16.00 MB	2.00

Evaluation

Micro-Benchmark


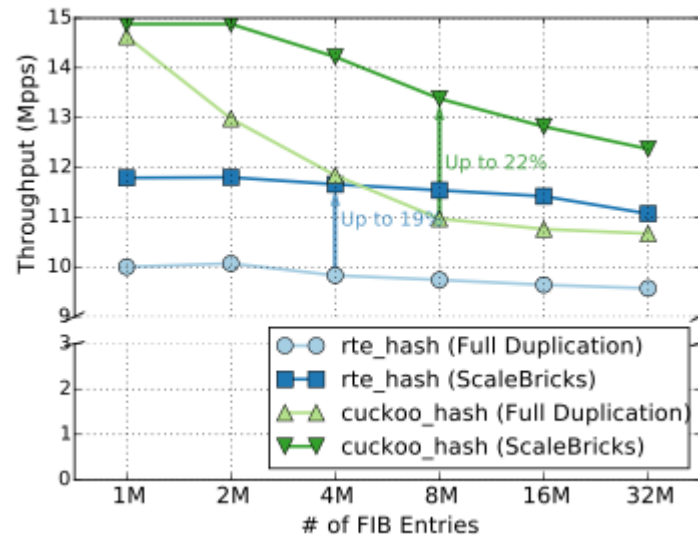
Local Lookup Throughput



Evaluation

Macro-Benchmark

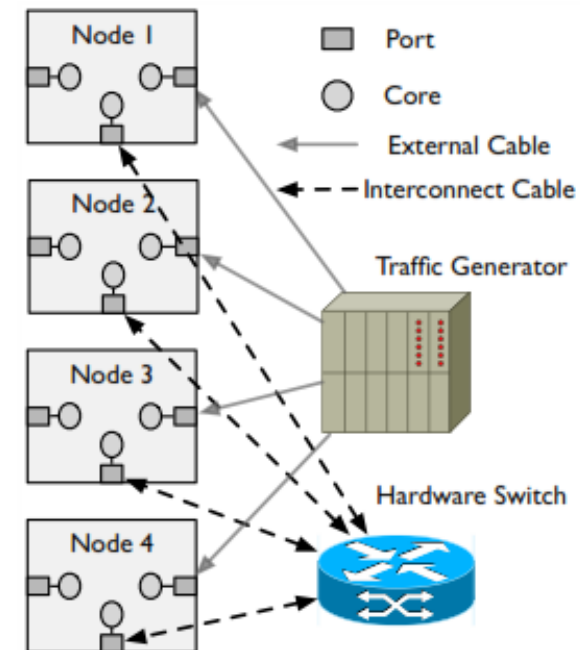

- ❑ rte_hash & extended cuckoo hash table
- ❑ Single node throuput /Full duplication & ScaleBricks
- ❑ Improve the through put and core utilization.



Evaluation

Scalability `rte_hash` & extended cuckoo hash table

- Single node throughput / Full duplication & ScaleBricks
- Improve the throughput and core utilization.



Evaluation

Macro-Benchmark

- ❑ Share the CPU cache with other application/
launch thread to consume cache
- ❑ Latency
- ❑ Update rate
 - Single core handle 60K updates/sec, 4-node
cluster for aggregated rate of 240K updates/sec.

