

RecMax: Exploiting Recommender Systems for Fun and Profit

Amit Goyal

Laks V. S. Lakshmanan




University of British Columbia

<http://cs.ubc.ca/~goyal>

Recommender Systems

2



Movies



Products



Music



Videos



News



Websites



RecMax – Recommendation Maximization

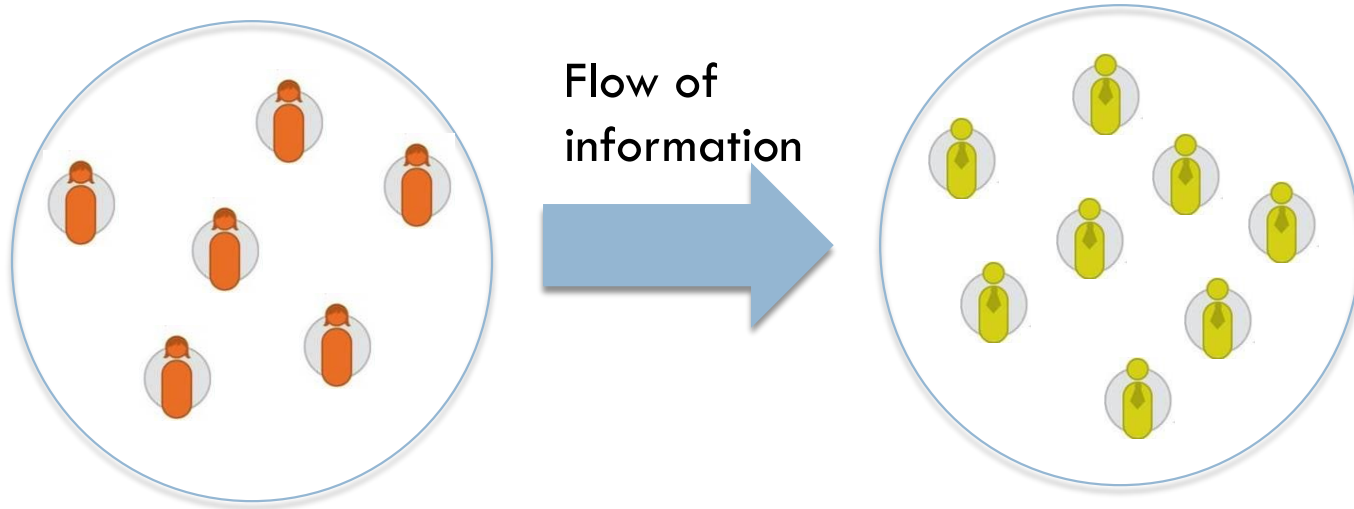
3

- Previous research mostly focused on improving accuracy of recommendations.
- In this paper, we propose a novel problem RecMax (short for Recommendation Maximization).

Can we launch a targeted marketing campaign over an existing operational Recommender System?

Consider an item in a Recommender System

4



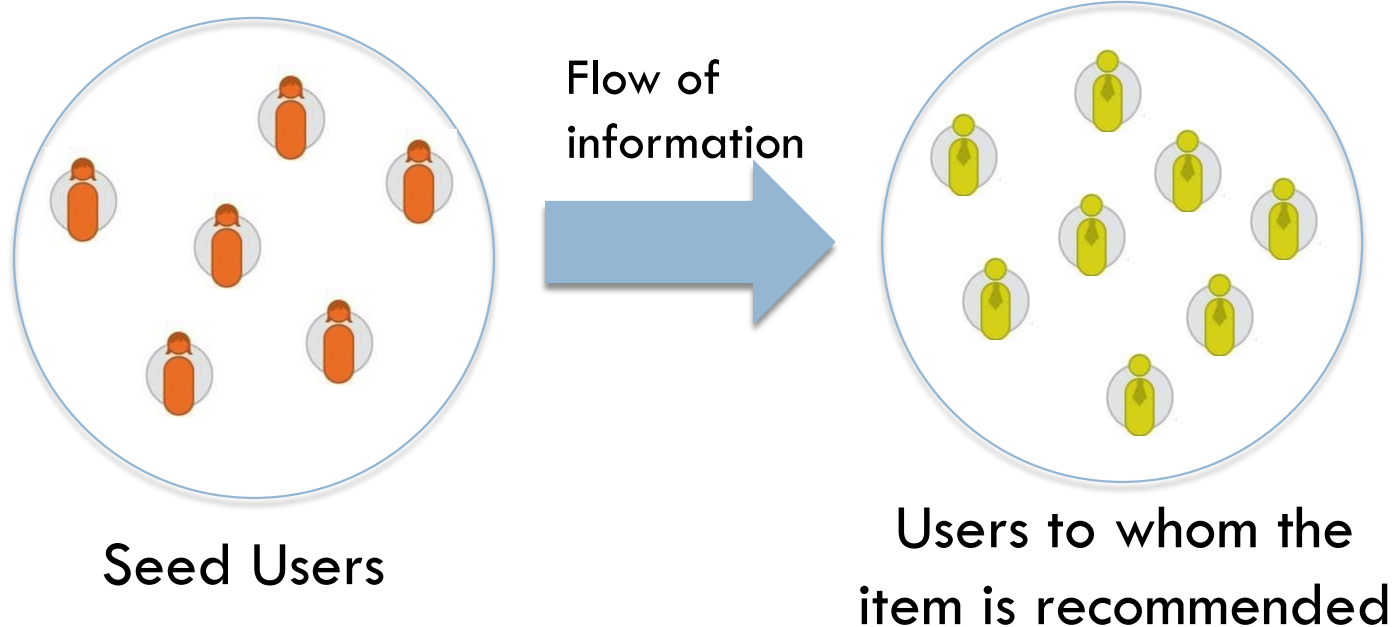
Some users rate the item
(seed users)

Because of these ratings,
the item may be
recommended to some
other users.

RecMax: Can we strategically select the seed users?

RecMax

5



Select k seed users such that
if they provide high ratings to a new product,
then the number of other users to whom the product is
recommended (hit score) by the underlying
recommender system algorithm is maximum.

RecMax – Problem Formulation

6

Recommendation List for user v

Recommendations	Expected Rating
Harry Potter	4.8
American Pie	4.3
...	
...	
The Dark Knight	3.2

rating threshold of user v
(denoted by θ_v)

For a new item i , if expected rating $R(v,i) > \theta_v$, then the new item is recommended to v

$$f(S) = \sum_{v \in V-S} I(R(v,i) > q_v)$$

The goal of RecMax is to find a seed set S such that hit score $f(S)$ is maximized.

Number of recommendations are 1

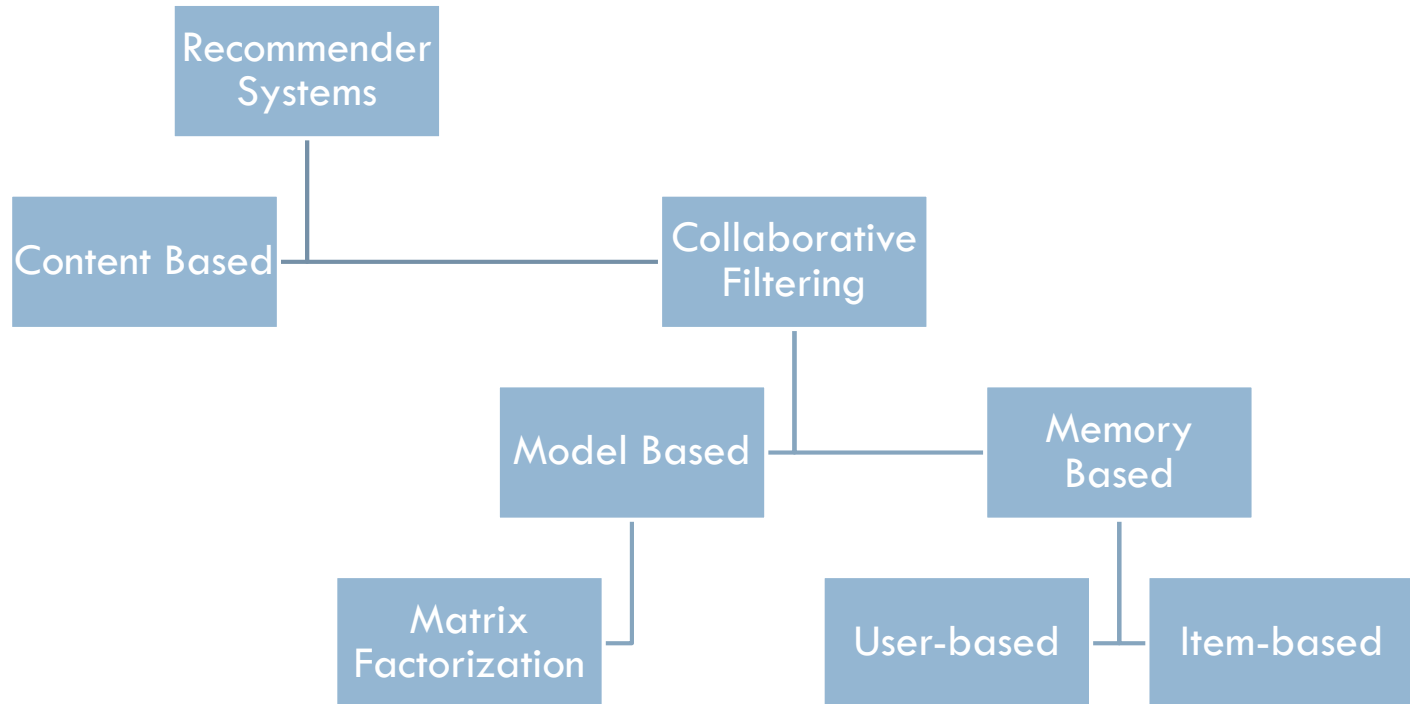
Benefits of RecMax

7

- Targeted marketing in Recommender Systems
 - ▣ Marketers can effectively advertise new products on a Recommender System platform.
 - ▣ Business opportunity to Recommender System platform.
 - ▣ Similar to Influence Maximization problem in spirit.
- Beneficial to seed users
 - ▣ They get free/discounted samples of a new product.
- Helpful to other users
 - ▣ They receive recommendations of new products – solution to cold start problem.

A key Challenge – Wide diversity of Recommender Systems

8




Similarity functions: Cosine, Pearson, Adjusted Cosine etc

Due to this wide diversity, it is very difficult to study RecMax

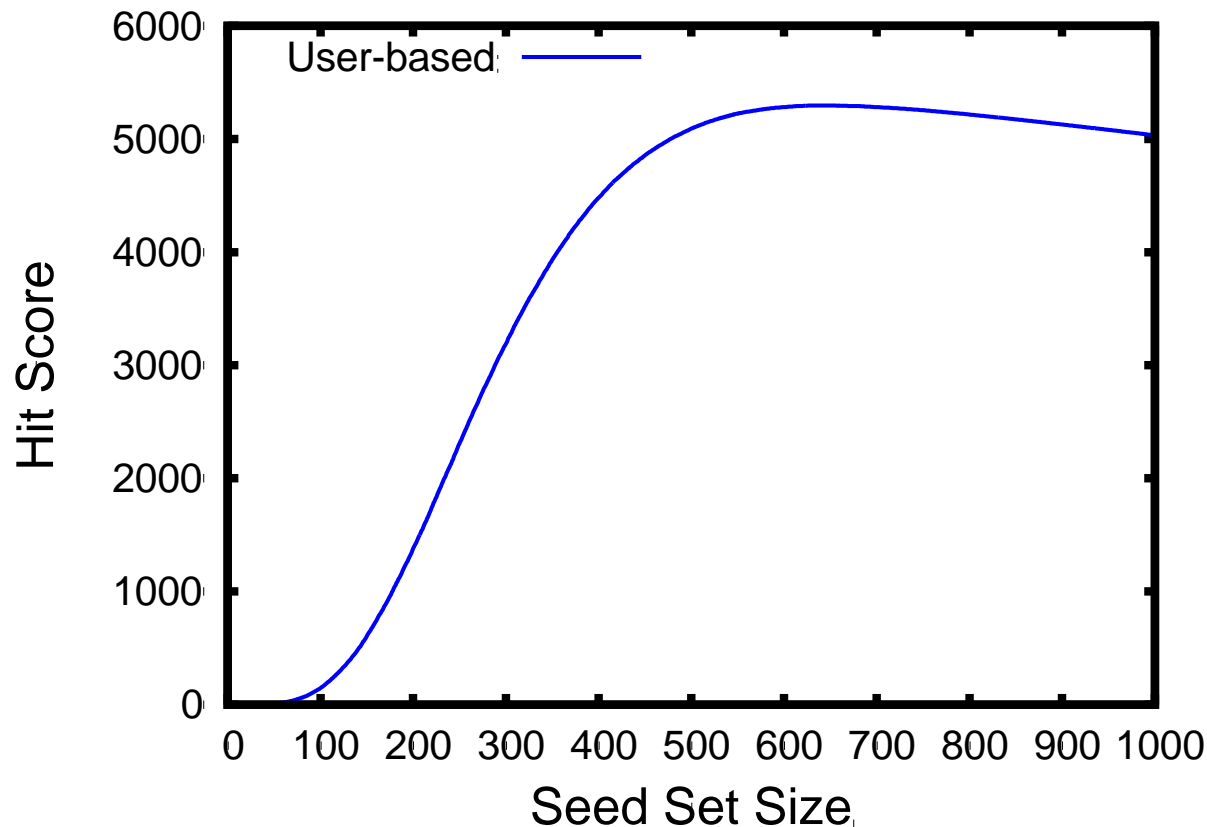
Outline

9

- What is RecMax? 
- Does Seeding Help? – Preliminary Experiments.
- Theoretical Analysis of RecMax.
- Experiments.
- Conclusions and Future Work.

Does Seeding Help?

10

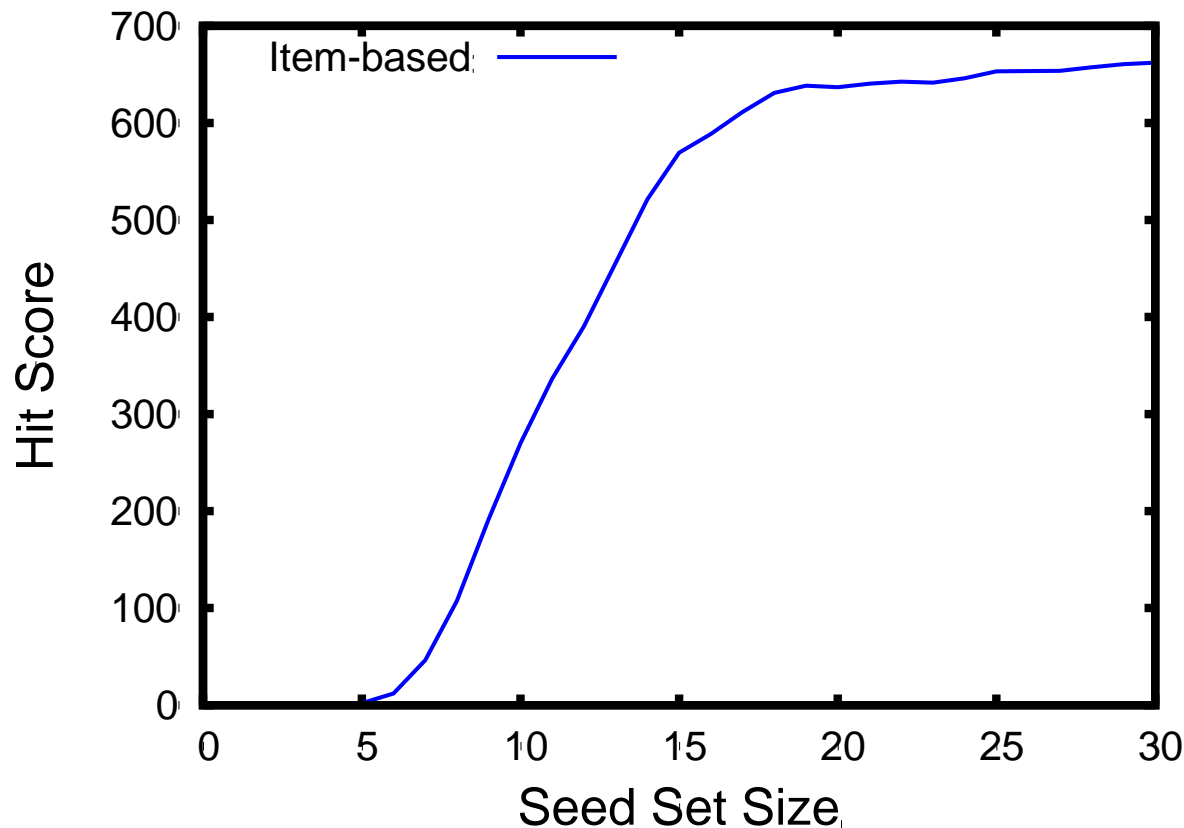


- Dataset: MovieLens
- Recommender System: User-based
- Seeds are picked randomly.
- Recall that Hit Score is the number of users to whom the product is recommended.

**A budget of 500 can get a hit score of 5091 (10x)
(User-based)**

Does Seeding Help?

11





- Dataset: Movielens
- Recommender System: Item-based
- Seeds are picked randomly.
- Recall that Hit Score is the number of users to whom the product is recommended.

**A budget of 20 can get a hit score of 636 (30x)
(Item-based)**

Outline

12

- What is RecMax? 
- Does Seeding Help? – Preliminary Experiments. 
- Theoretical Analysis of RecMax.
- Experiments.
- Conclusions and Future Work.

Key Theoretical Results


13

- RecMax is NP-hard to solve exactly.
- RecMax is NP-hard to approximate within a factor to $1/|V|^{(1-\epsilon)}$ for any $\epsilon > 0$.
 - ▣ No reasonable approximation algorithm can be developed.
- RecMax is as hard as Maximum Independent Set Problem.
- Under both User-based and under Item-based.

Why is RecMax so hard? (1 / 2)

14

- We introduce a helper problem – **Maximum Encirclement Problem**
 - ▣ find a set S of size k such that it encircles maximum number of nodes in the graph.




- Nodes $\{B, C\}$ encircle node A.
- Nodes $\{B, C, E\}$ encircle node D.
- Thus, $\{B, C, E\}$ encircle A and D.

Why is RecMax so hard? (2/2)

15

- Set $\{B, C, E\}$ is a solution to Maximum Encirclement Problem (for $k=3$).
- Nodes $\{A, D\}$ form Maximum Independent Set.
- **Reduction:** Nodes $\{B, C, E\}$ must rate the new item highly for the item to be recommended to A and D.
- RecMax is as hard as Maximum Independent Set, and hence NP-hard to approximate within a factor to $1/|V|^{(1-\epsilon)}$



Discussion (1 / 2)

16

- We show hardness for User-based and Item-based methods.
 - ▣ What about Matrix Factorization?
 - ▣ Most likely hardness would remain (future work).




Discussion (2/2)

17

- Since the problem is hard to approximate, does it make sense to study?
 - ▣ YES, as we saw earlier, even a random heuristic fetches impressive gains.
- We explore several natural heuristics and compare them.
 - ▣ What about sophisticated heuristics (future work).

Outline

18

- What is RecMax? 
- Does Seeding Help? – Preliminary Experiments. 
- Theoretical Analysis of RecMax. 
- Experiments.
- Conclusions and Future Work.

Datasets

19

	Movielens	Yahoo	Jester
#Users	6040	10K	25K
#Items	3706	5069	100
#Ratings	1M	1M	1.8M
Mean	3.58	51.9	10.88
Std. Dev.	1.12	39.9	5.23

Heuristics

20

- **Random**: Seed set is selected randomly. The process is repeated several times and average is taken.
- **Most-Active**: Top-k users with most number of ratings.
- **Most-Positive**: Top-k users with most positive average ratings.
- **Most-Critical**: Top-k users with most critical average ratings.

Heuristics

21

- **Most-Central**: Top-k central users.

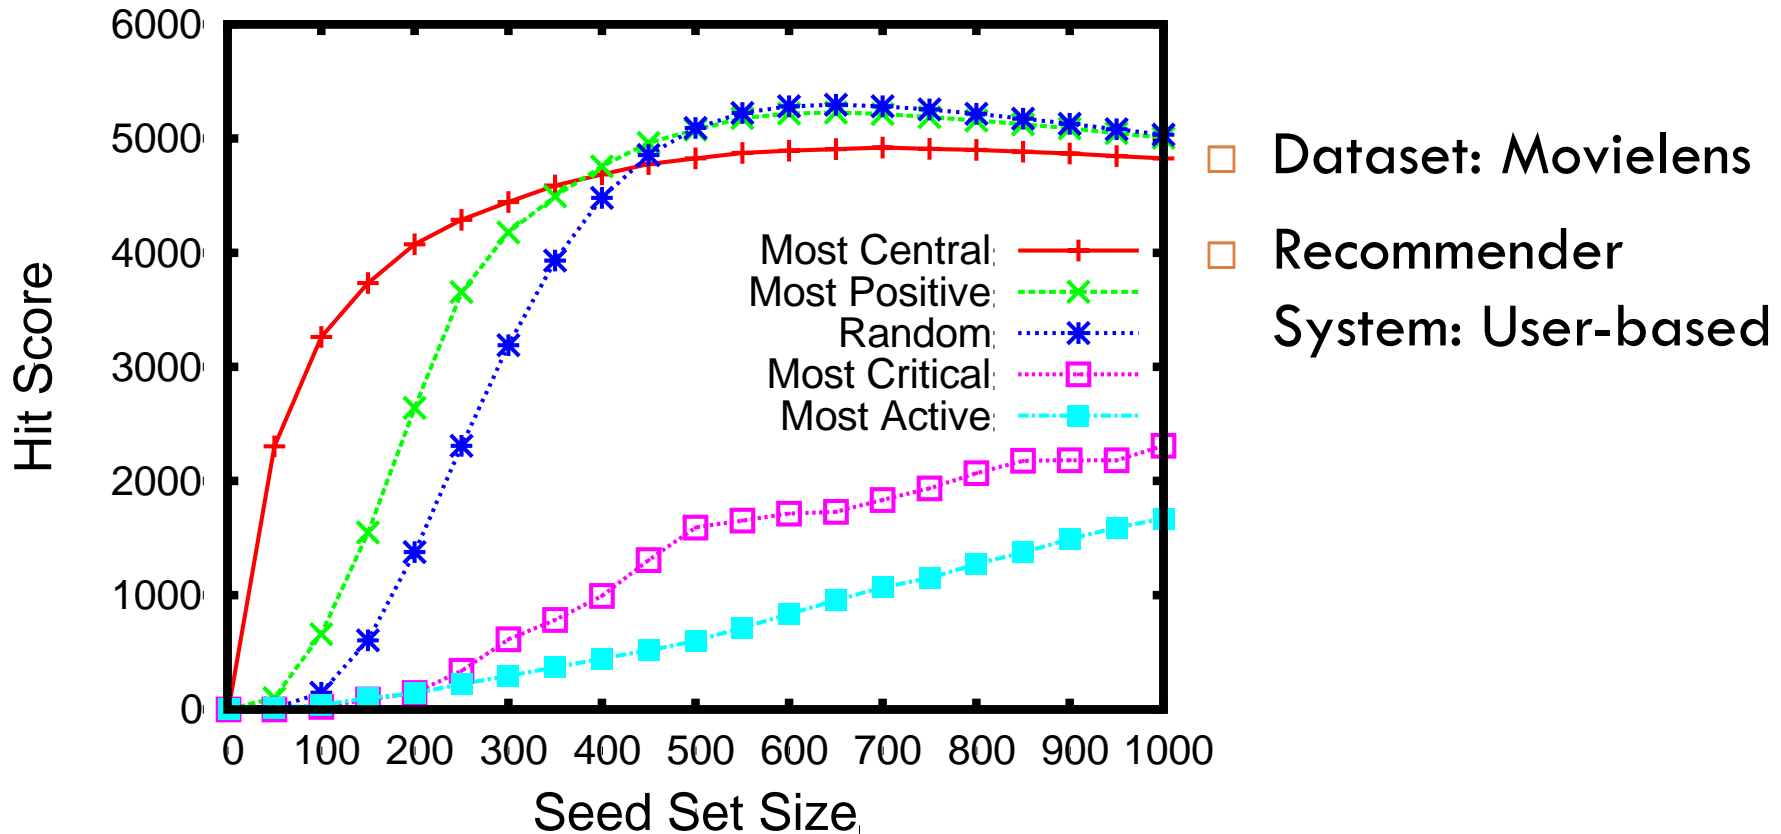
$$agg(u) = \sum_{v \in V - u} sim(u, v)$$

22

User-Based Recommender Systems

Comparison – Hit Score achieved

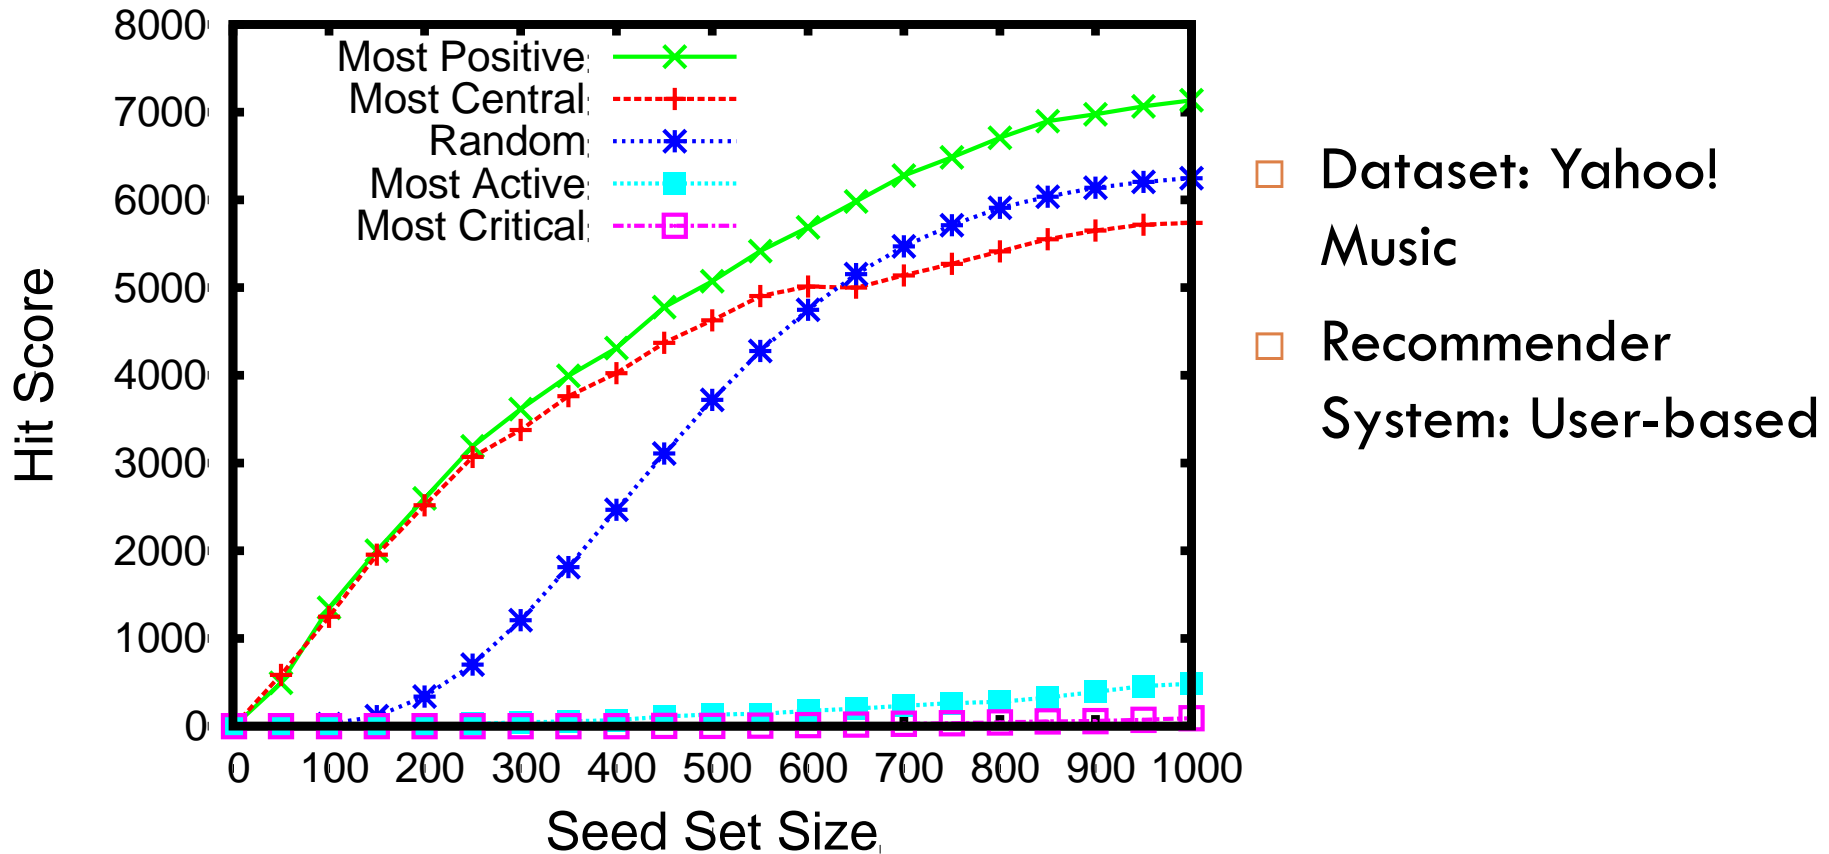
23



Most Central, Most Positive and Random perform good here.

Comparison – Hit Score achieved

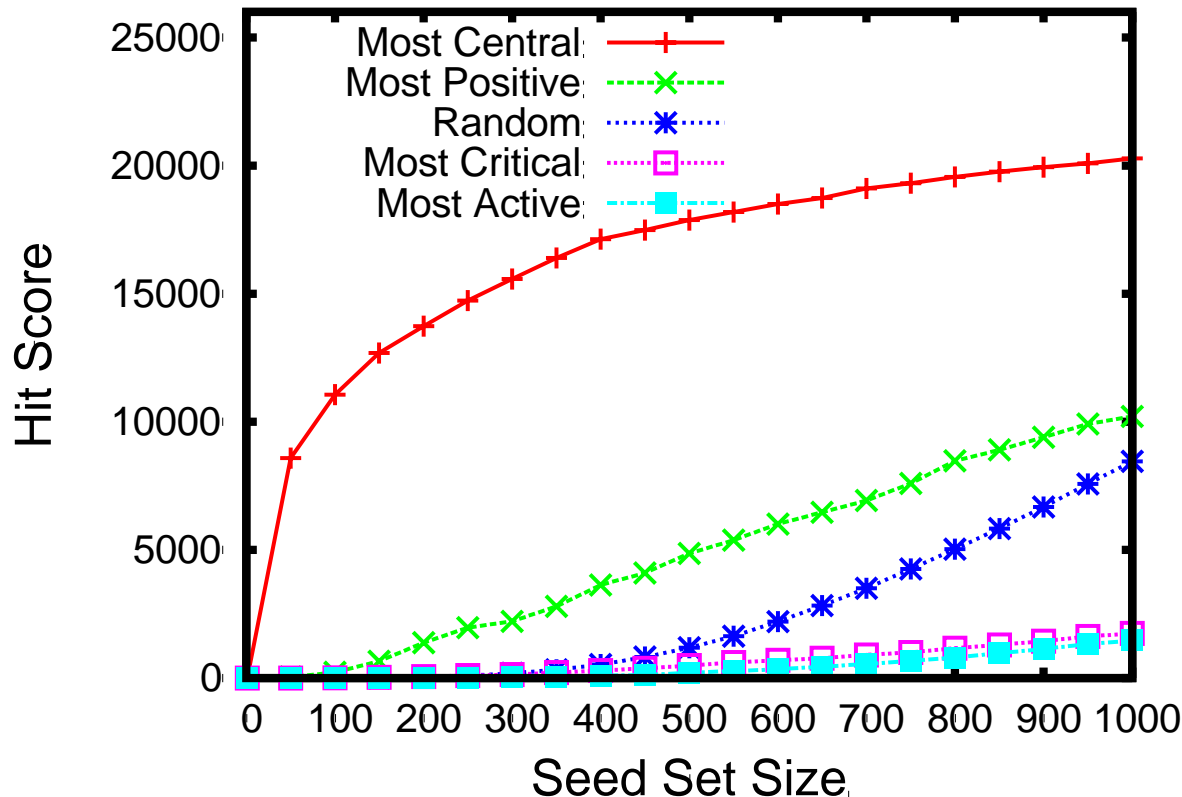
24



Most Positive, Most Central perform good here.

Comparison – Hit Score achieved

25



Dataset: Jester Joke
Recommender System: User-based

Most Central out-performs all other heuristics.

Key Takeaways

26

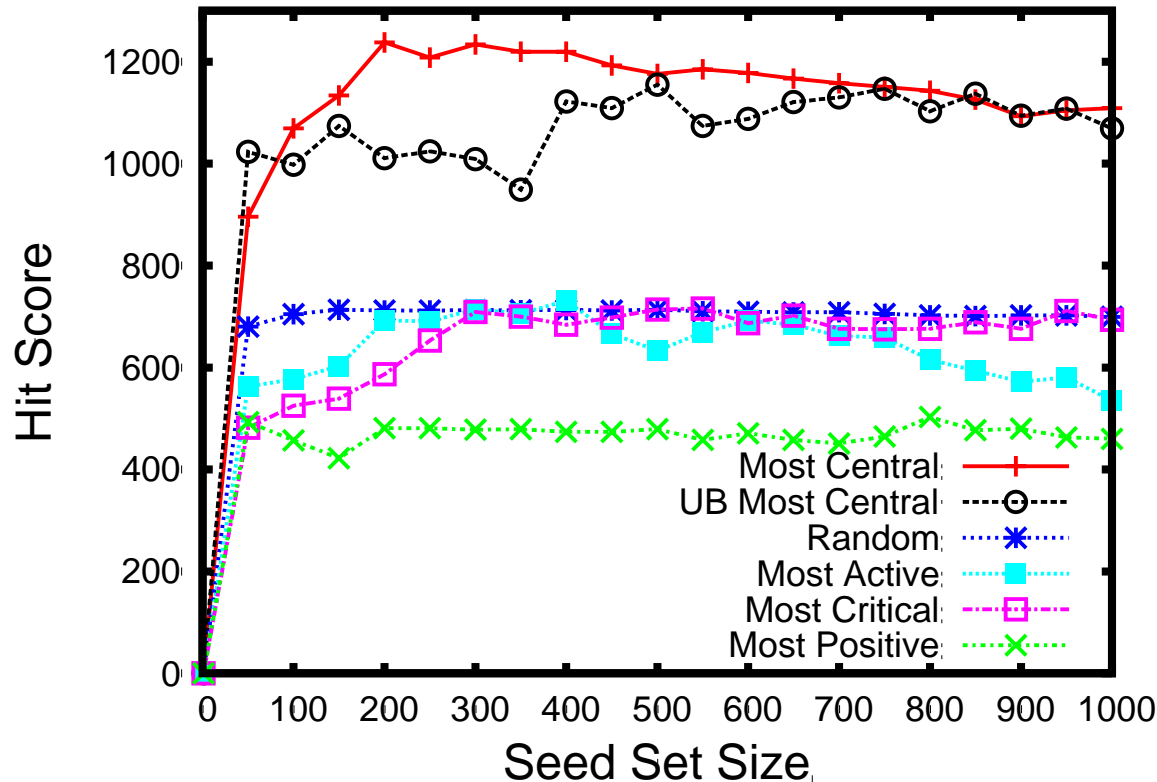
- Even the simple heuristics perform well.
- With a budget of 300, Most-Central heuristic achieves hit score of 4.4K, 3.4K and 15.6K on Movielens, Yahoo! and Jester respectively.
- Depending on the data set, we may encounter a “tipping point” – a minimum seeding is needed for the results to be impressive.

27

Item-Based Recommender Systems

Comparison – Hit Score achieved

28

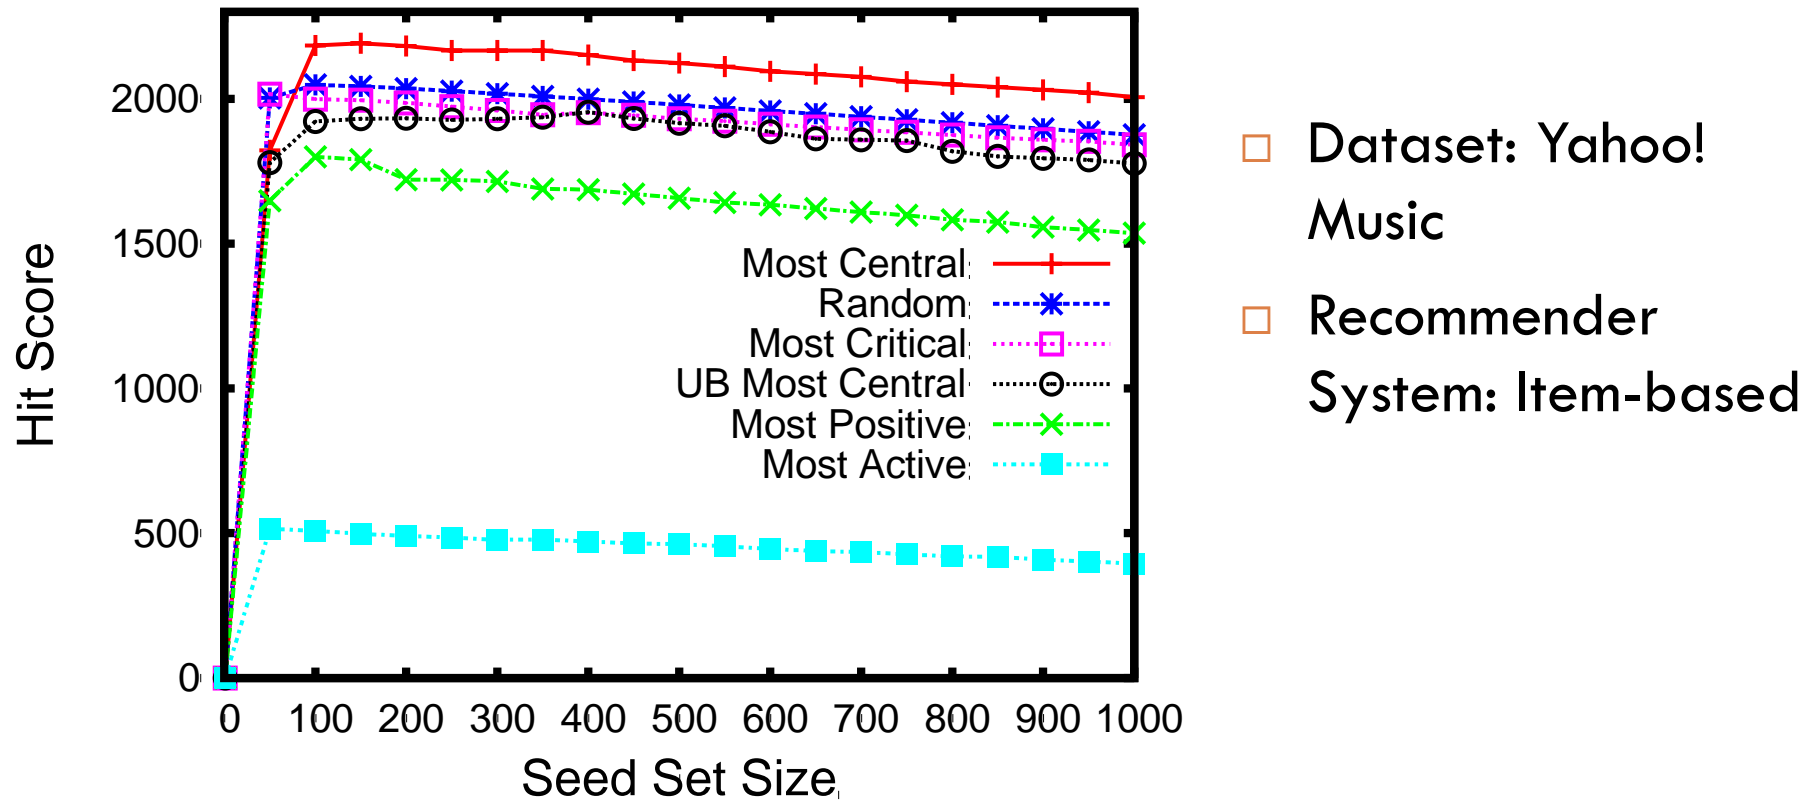


Dataset: Movielens
Recommender System: Item-based

Most Central performs good here.

Comparison – Hit Score achieved

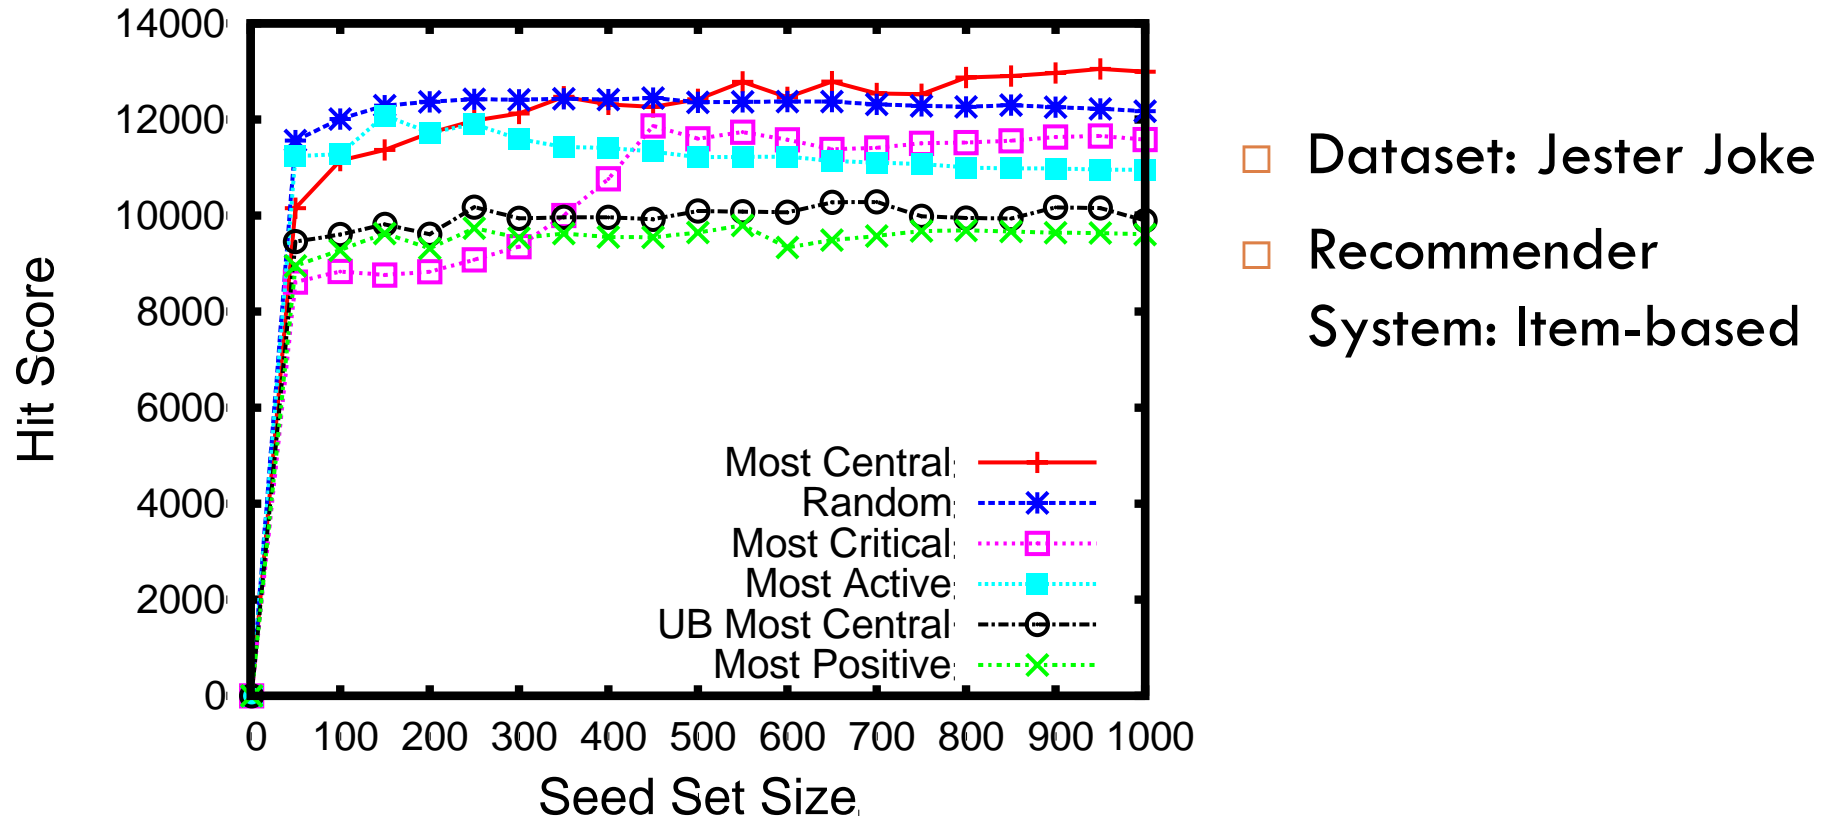
29



Most Central performs good here.

Comparison – Hit Score achieved

30



Most Central, Random and Most-Active performs good here.

Key Takeaways

31

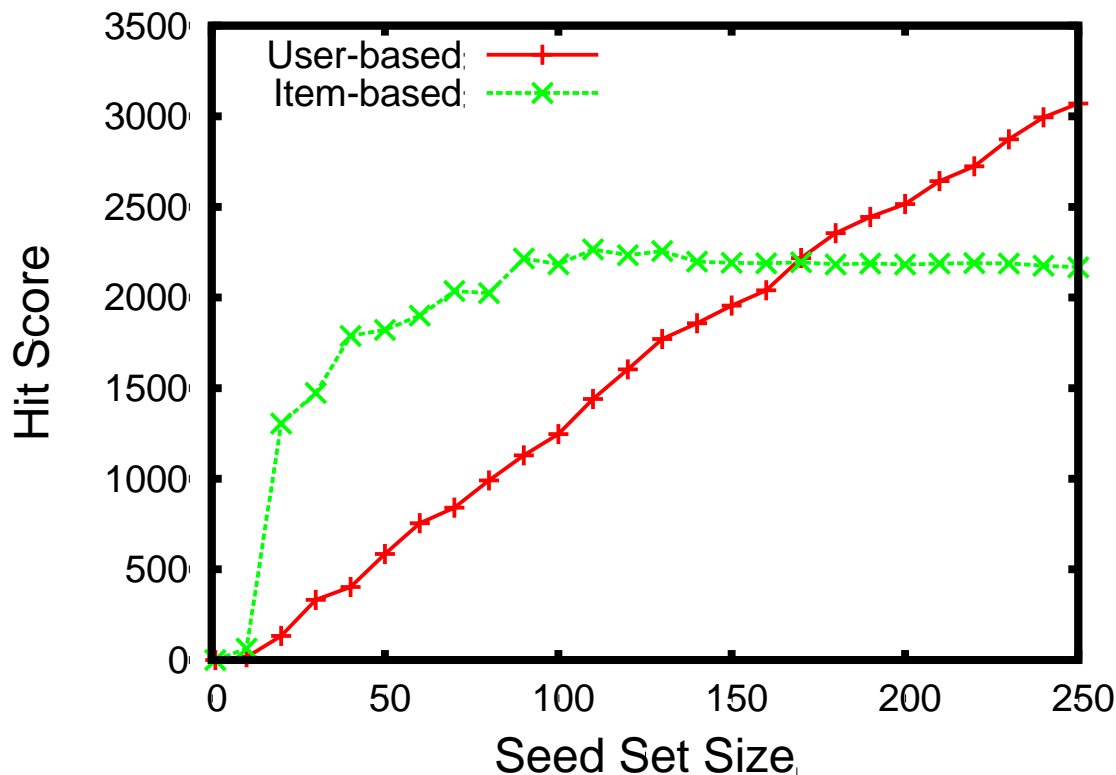
- Again, the simple heuristics perform well.
- Hit score achieved in Item-based is much lower than in User-based.
- Thus, much less seeding is required to achieve maximum possible hit score.
- Overall, Most-Central performs well.
- The difference of Most-Central with baseline Random is not much.
 - ▣ We need better heuristics (future work).

32

User-Based vs Item-Based

User-based vs Item-based

33



- Dataset: Yahoo! Music
- Initial rise of hit score is steeper in Item-based.
- Hit score saturates much earlier in Item-based.

Eventual hit score that can be achieved is much more in User-based.

User-based vs Item-based





34

	Common Seeds (out of 1000 seeds)
Movielens	103 (10.3%)
Yahoo! Music	219 (21.9 %)
Jester Joke	62 (0.62 %)

Seed Sets are different in both methods.

Outline

35

- What is RecMax? 
- Does Seeding Help? – Preliminary Experiments. 
- Theoretical Analysis of RecMax. 
- Experiments. 
- Conclusions and Future Work.

Our Contributions

36

- The main goal of the paper is to propose and study a novel problem that we call RecMax
 - ▣ Select k seed users such that if they endorse a new product by providing relatively high ratings, the number of users to whom the product is recommended (hit score) is maximum.
 - ▣ We focus on User-based and Item-based recommender systems.
- We offer empirical evidence that seeding does help in boosting the number of recommendations

Our Contributions

37

- We perform a thorough theoretical analysis of RecMax.
 - ▣ RecMax is NP-hard to solve exactly.
 - ▣ RecMax is NP-hard to approximate within any reasonable factor.
- Given this hardness, we explore several natural heuristics on 3 real world datasets and report our findings.
- Even simple heuristics like Most-Central provide impressive gains
 - ▣ This makes RecMax an interesting problem for targeted marketing in recommender systems.

Future Work

38

- RecMax is a new problem and has real applications
 - our work is just the first work.
- Developing better heuristics.
- Studying RecMax on more sophisticated recommender systems algorithms
 - ▣ Matrix Factorization.

Thanks and Questions

RecMax: Exploiting Recommender Systems for Fun and Profit

Amit Goyal

Laks V. S. Lakshmanan




University of British Columbia

Heuristics

40

- Most-Central: Top-k central users.



$$sim(u, v) = \frac{\sum_{i \in I(u), j \in I(v)} w(i, j) \cdot sim(R(u, i), R(v, j))}{\sum_{i \in I(u), j \in I(v)} w(i, j)}$$