

# Heading Off Correlated Failures through Independence-as-a-Service

Ennan Zhai<sup>1</sup>

Ruichuan Chen<sup>2</sup>, David Isaac Wolinsky<sup>1</sup>, Bryan Ford<sup>1</sup>

<sup>1</sup> Yale University

<sup>2</sup> Bell Labs/Alcatel-Lucent



Bell Labs 

# Background

- Cloud services ensure reliability by redundancy:

# Background

- Cloud services ensure reliability by redundancy:
  - Amazon S3 replicates data on multiple racks

# Background

- Cloud services ensure reliability by redundancy:
  - Amazon S3 replicates data on multiple racks
  - iCloud rents EC2 and Azure redundantly

# Background

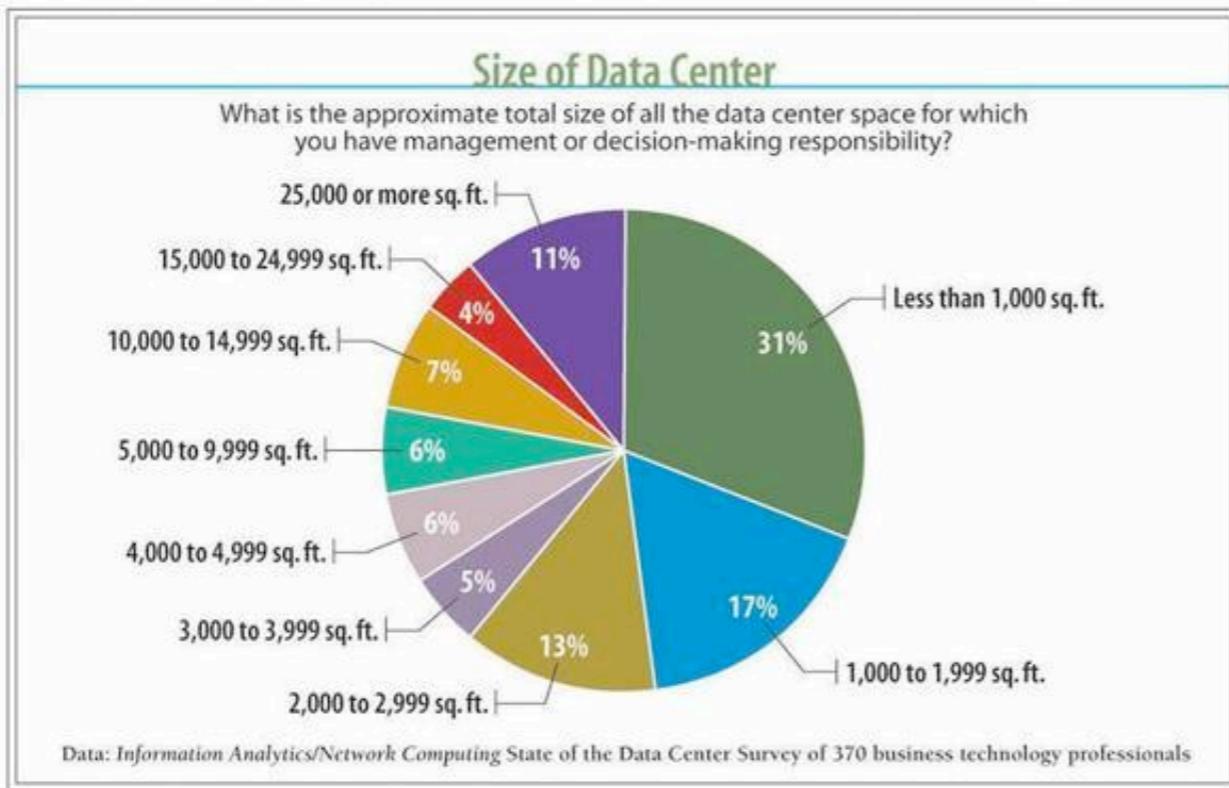
- Cloud services ensure reliability by redundancy:
  - Amazon S3 replicates data on multiple racks
  - iCloud rents EC2 and Azure redundantly

**Unexpected common dependencies!**

# Service Outage Losses

## Data Center Outages Generate Big Losses

Downtime in a data center can cost an average of \$505,500 per incident, according to a Ponemon Institute study.

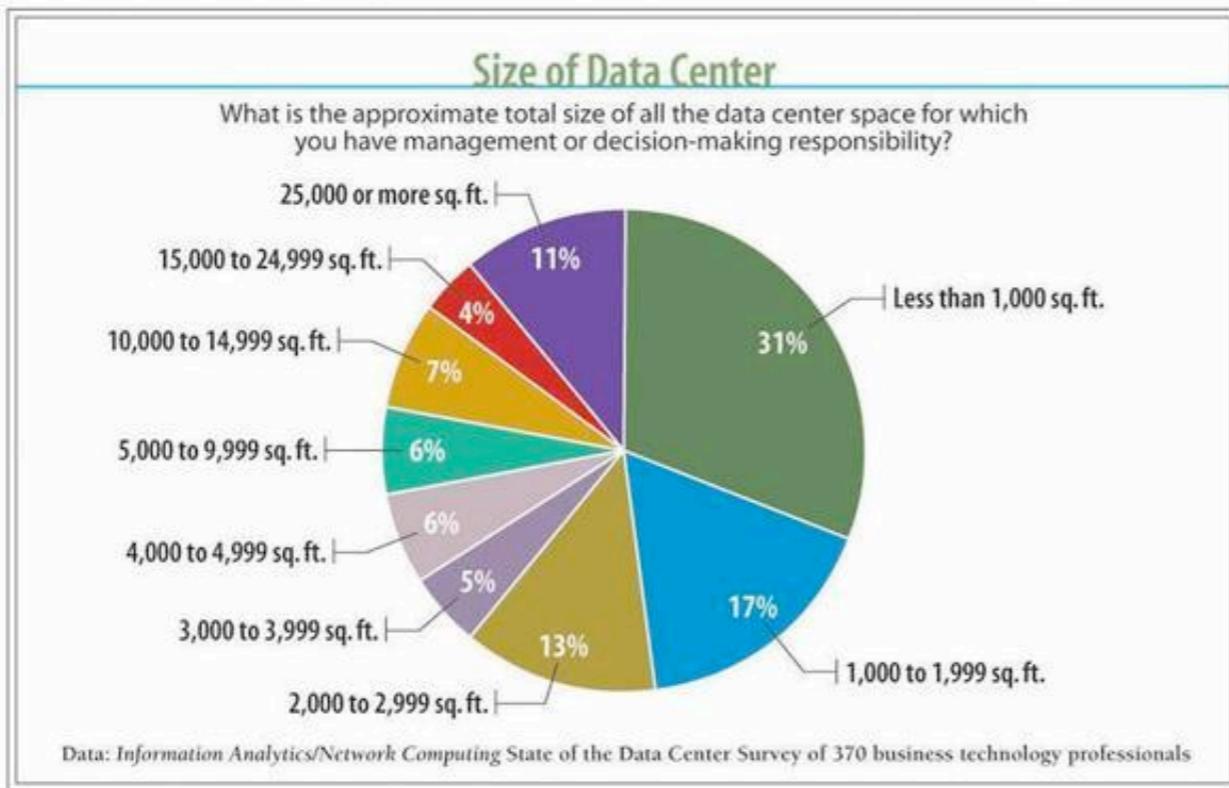


**Analytics Slideshow:  
2010 Data Center  
Operational Trends  
Report**

# Service Outage Losses

## Data Center Outages Generate Big Losses

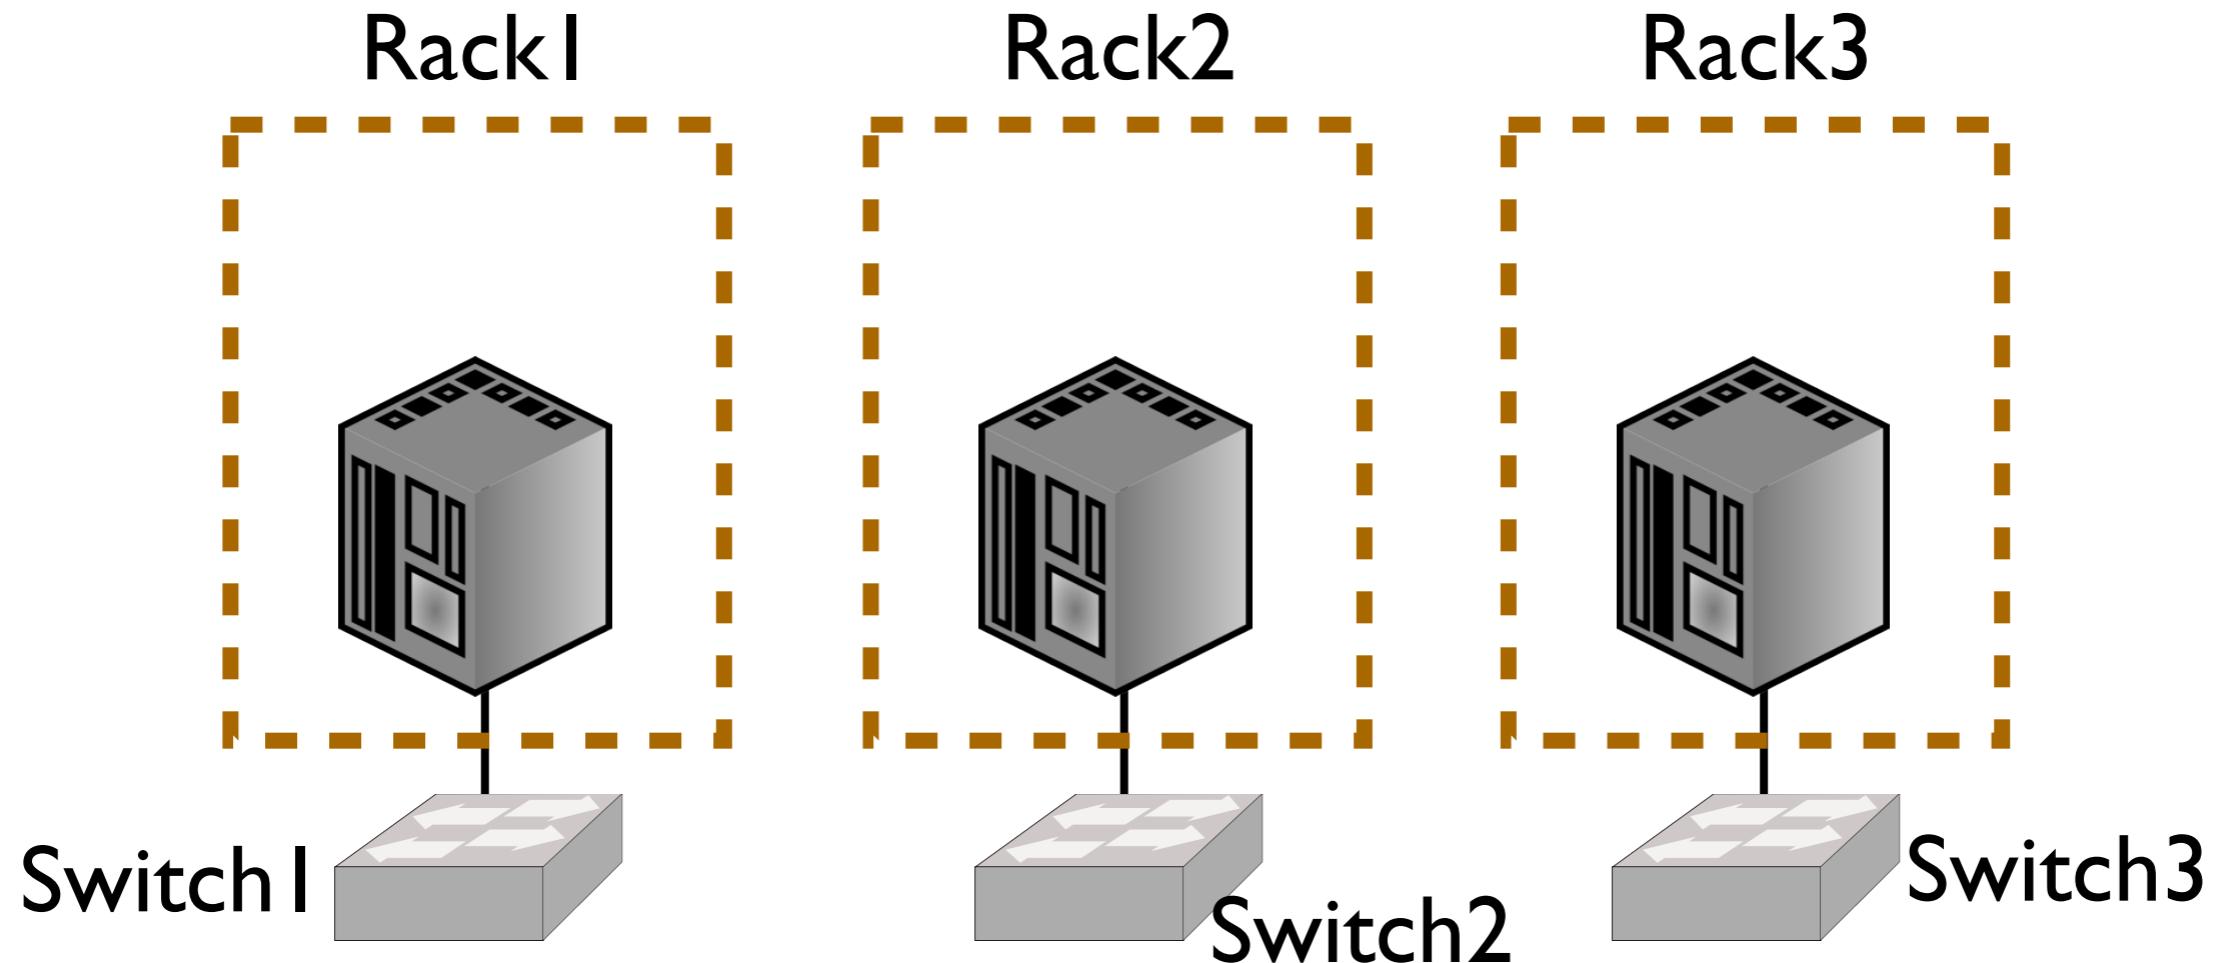
Downtime in a data center can cost an average of \$505,500 per incident, according to a Ponemon Institute study.



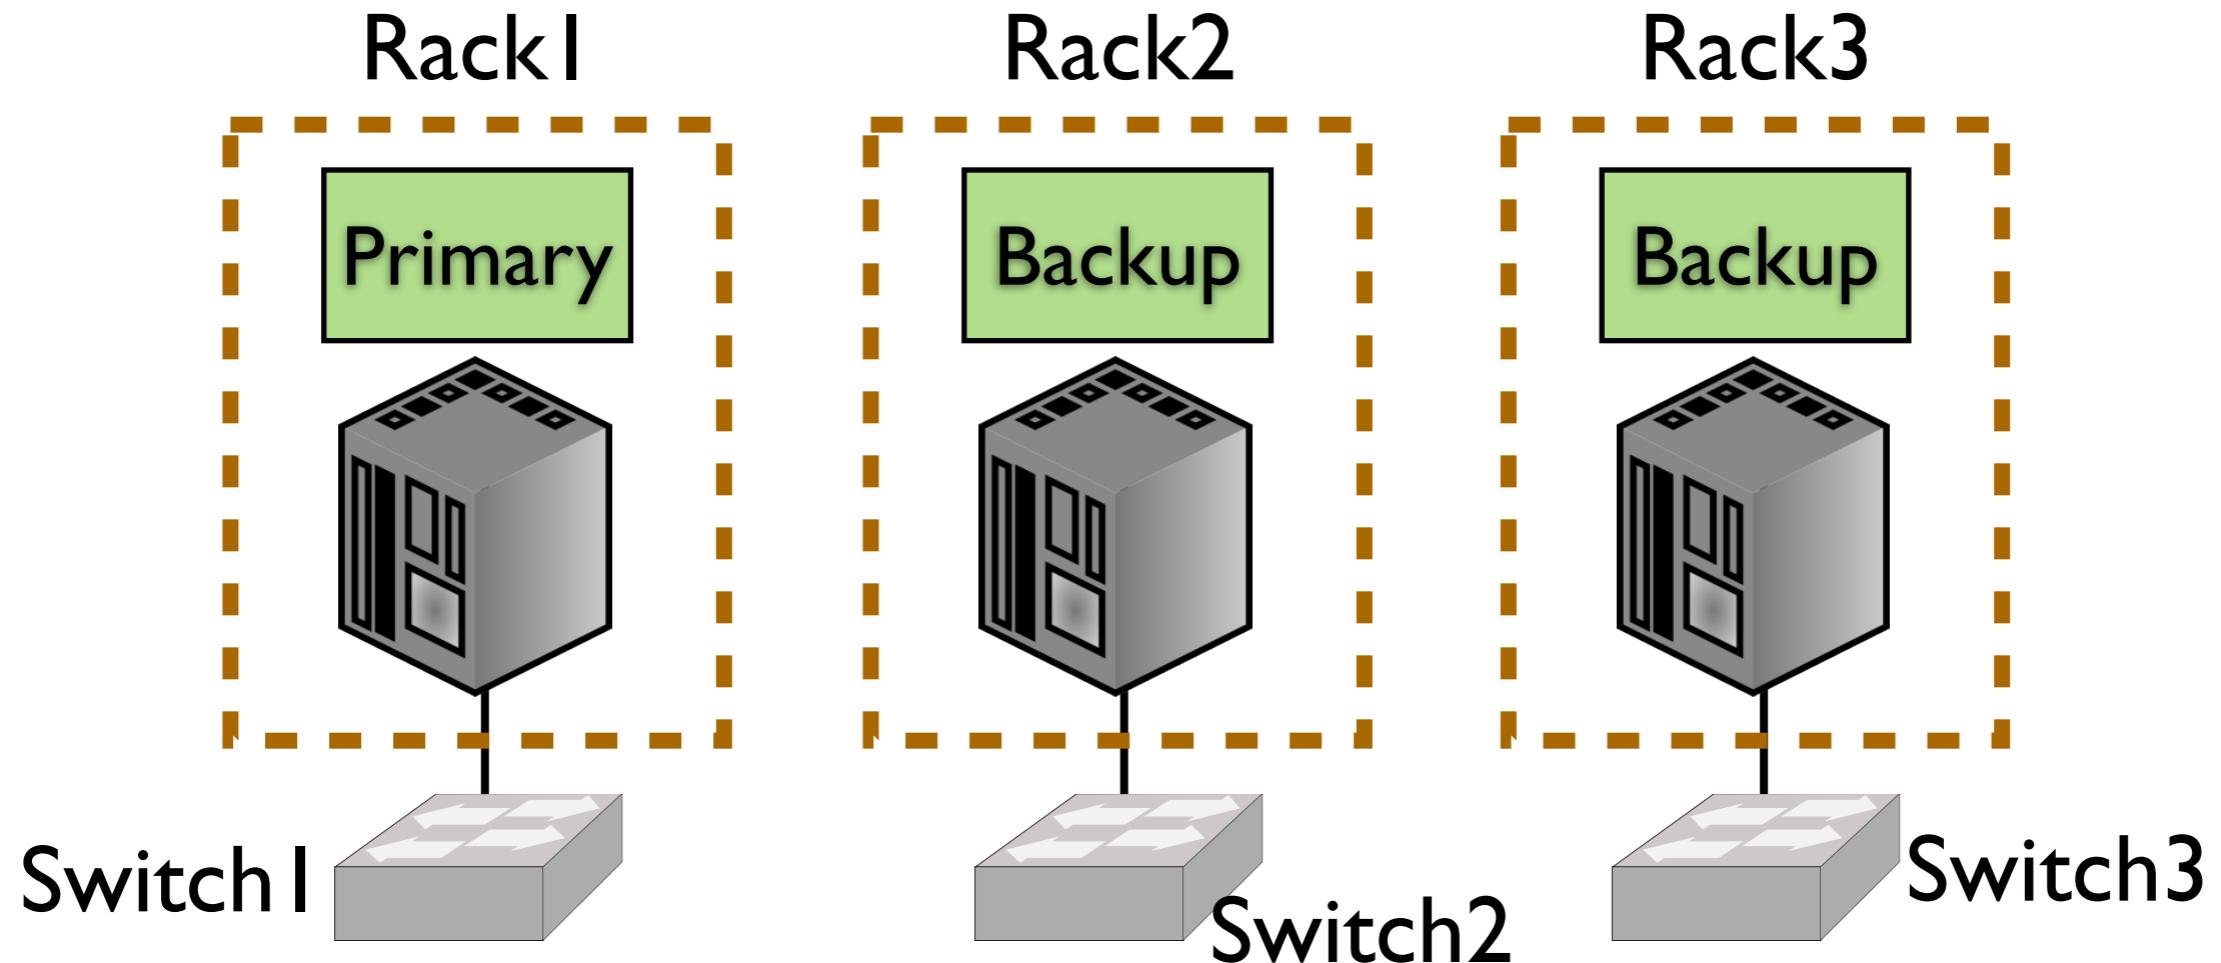
**Analytics Slideshow:  
2010 Data Center  
Operational Trends  
Report**

# What is correlated failure?

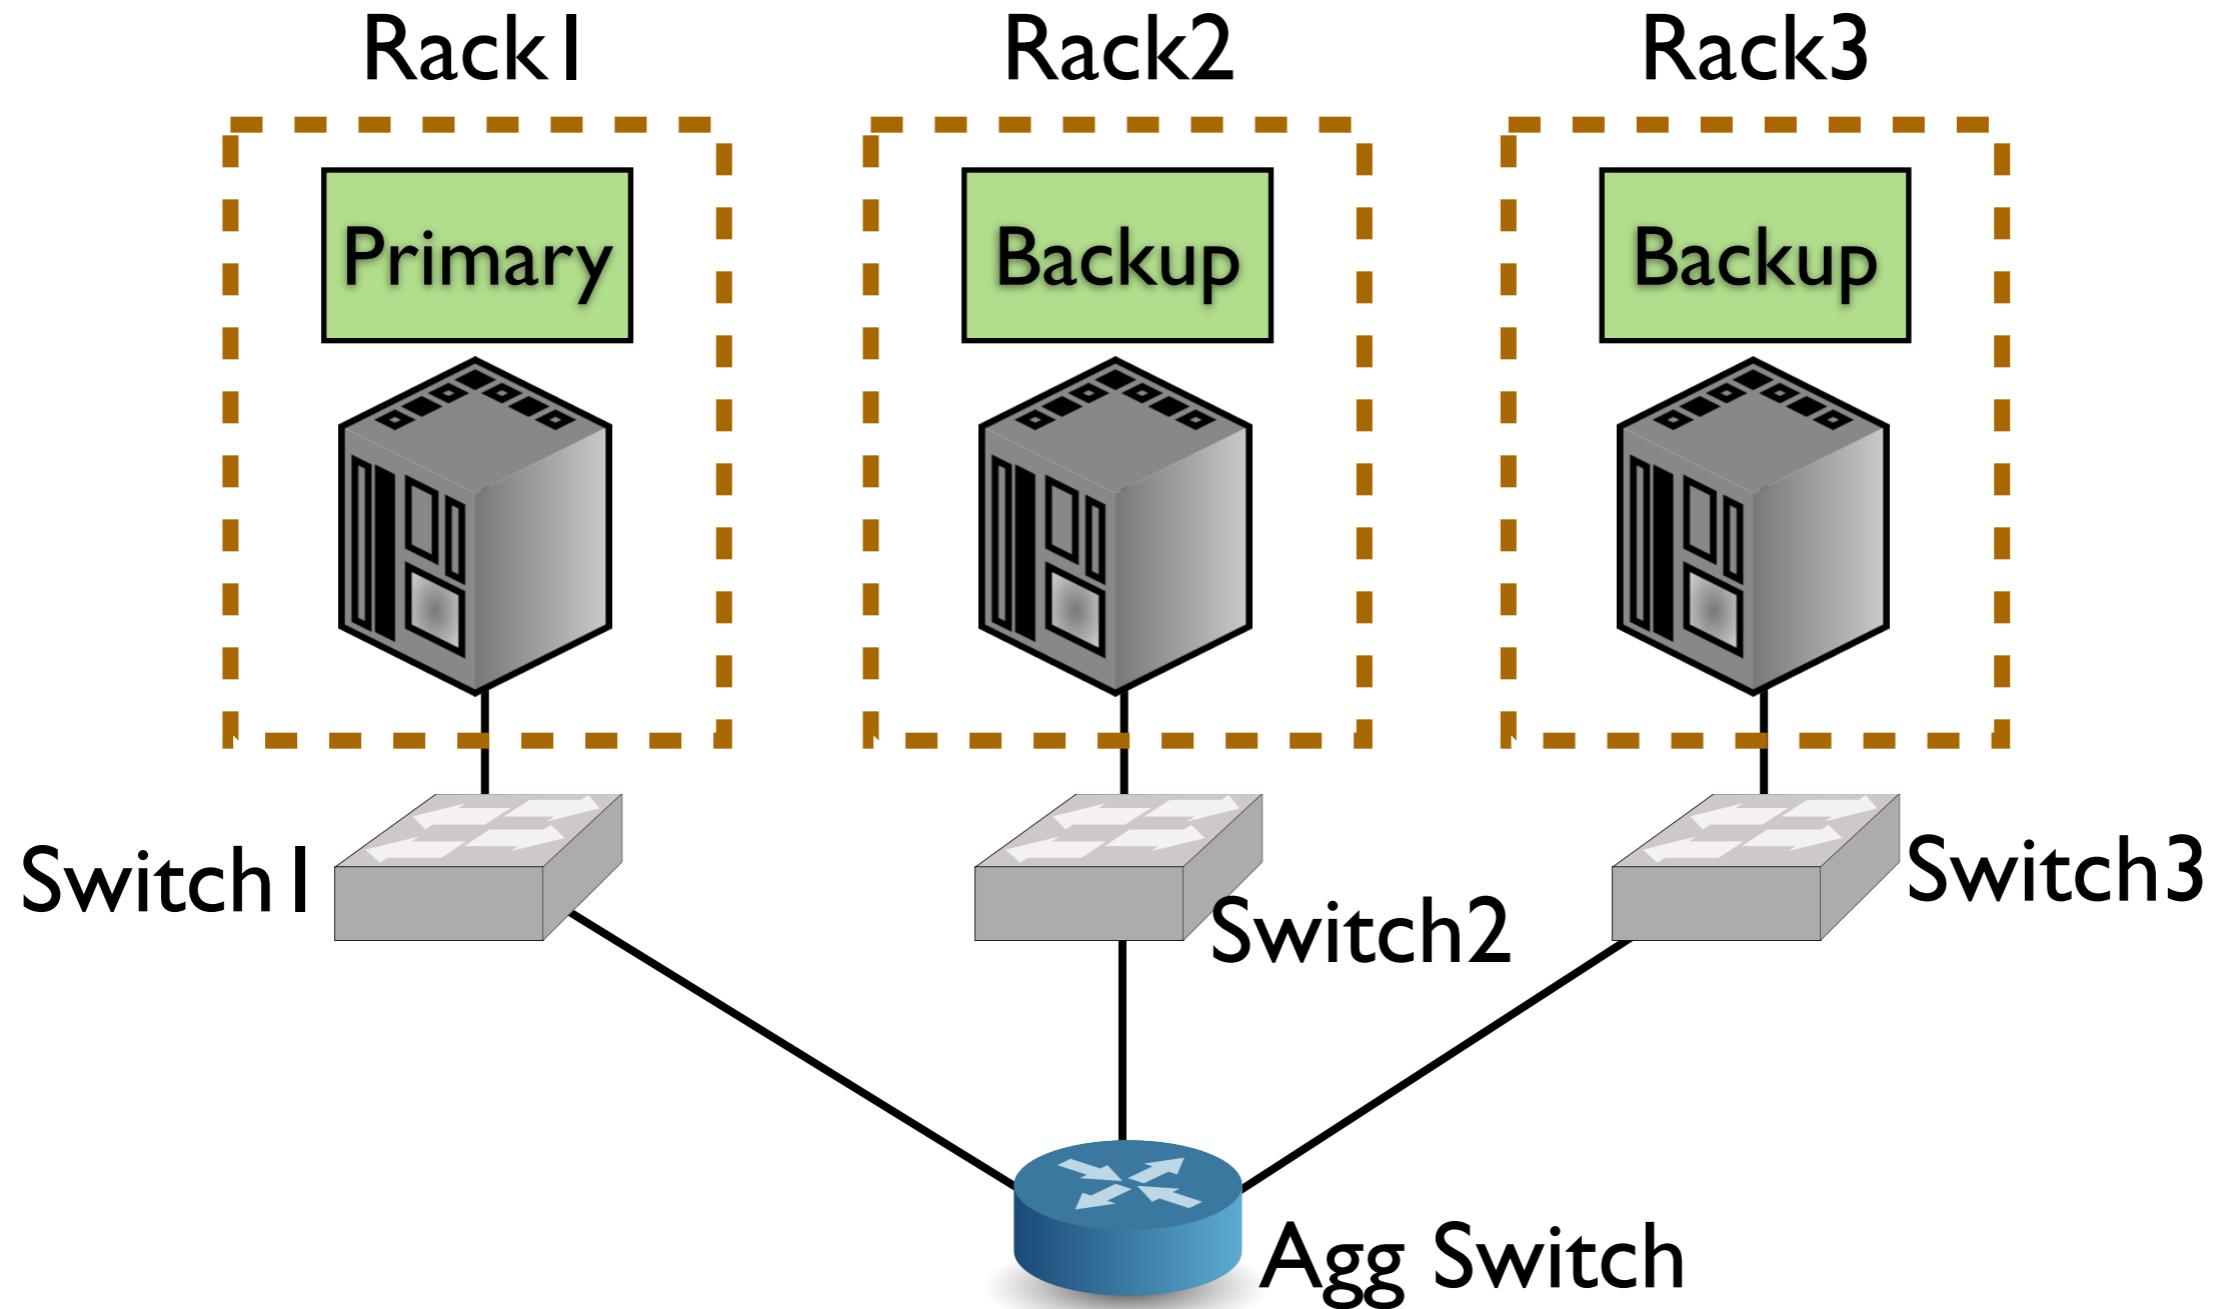
# What is correlated failure?



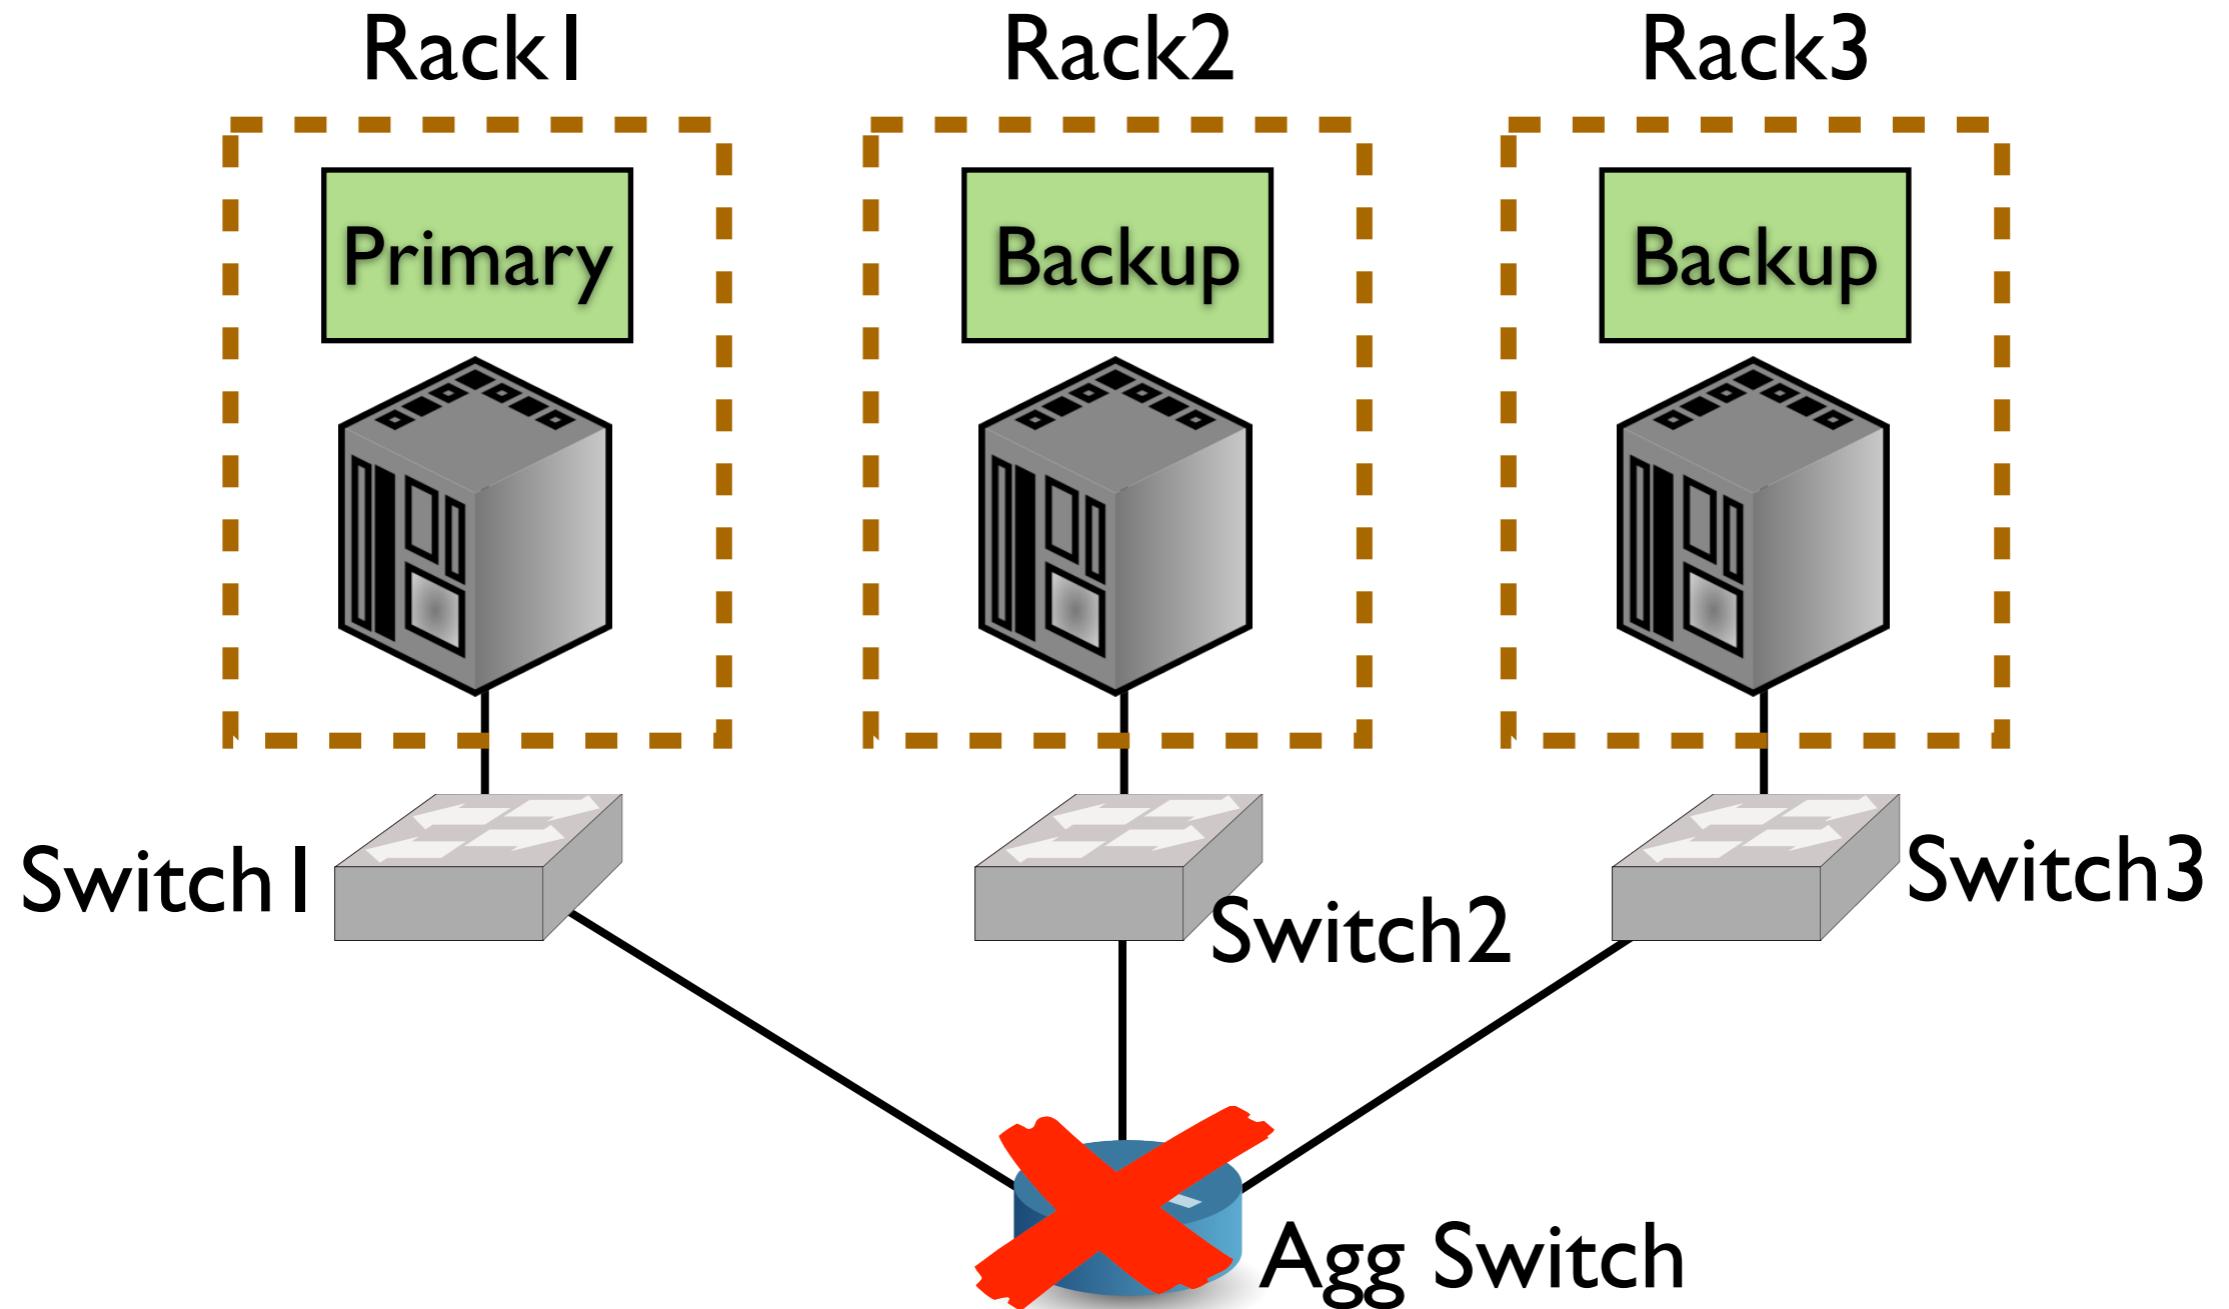
# What is correlated failure?



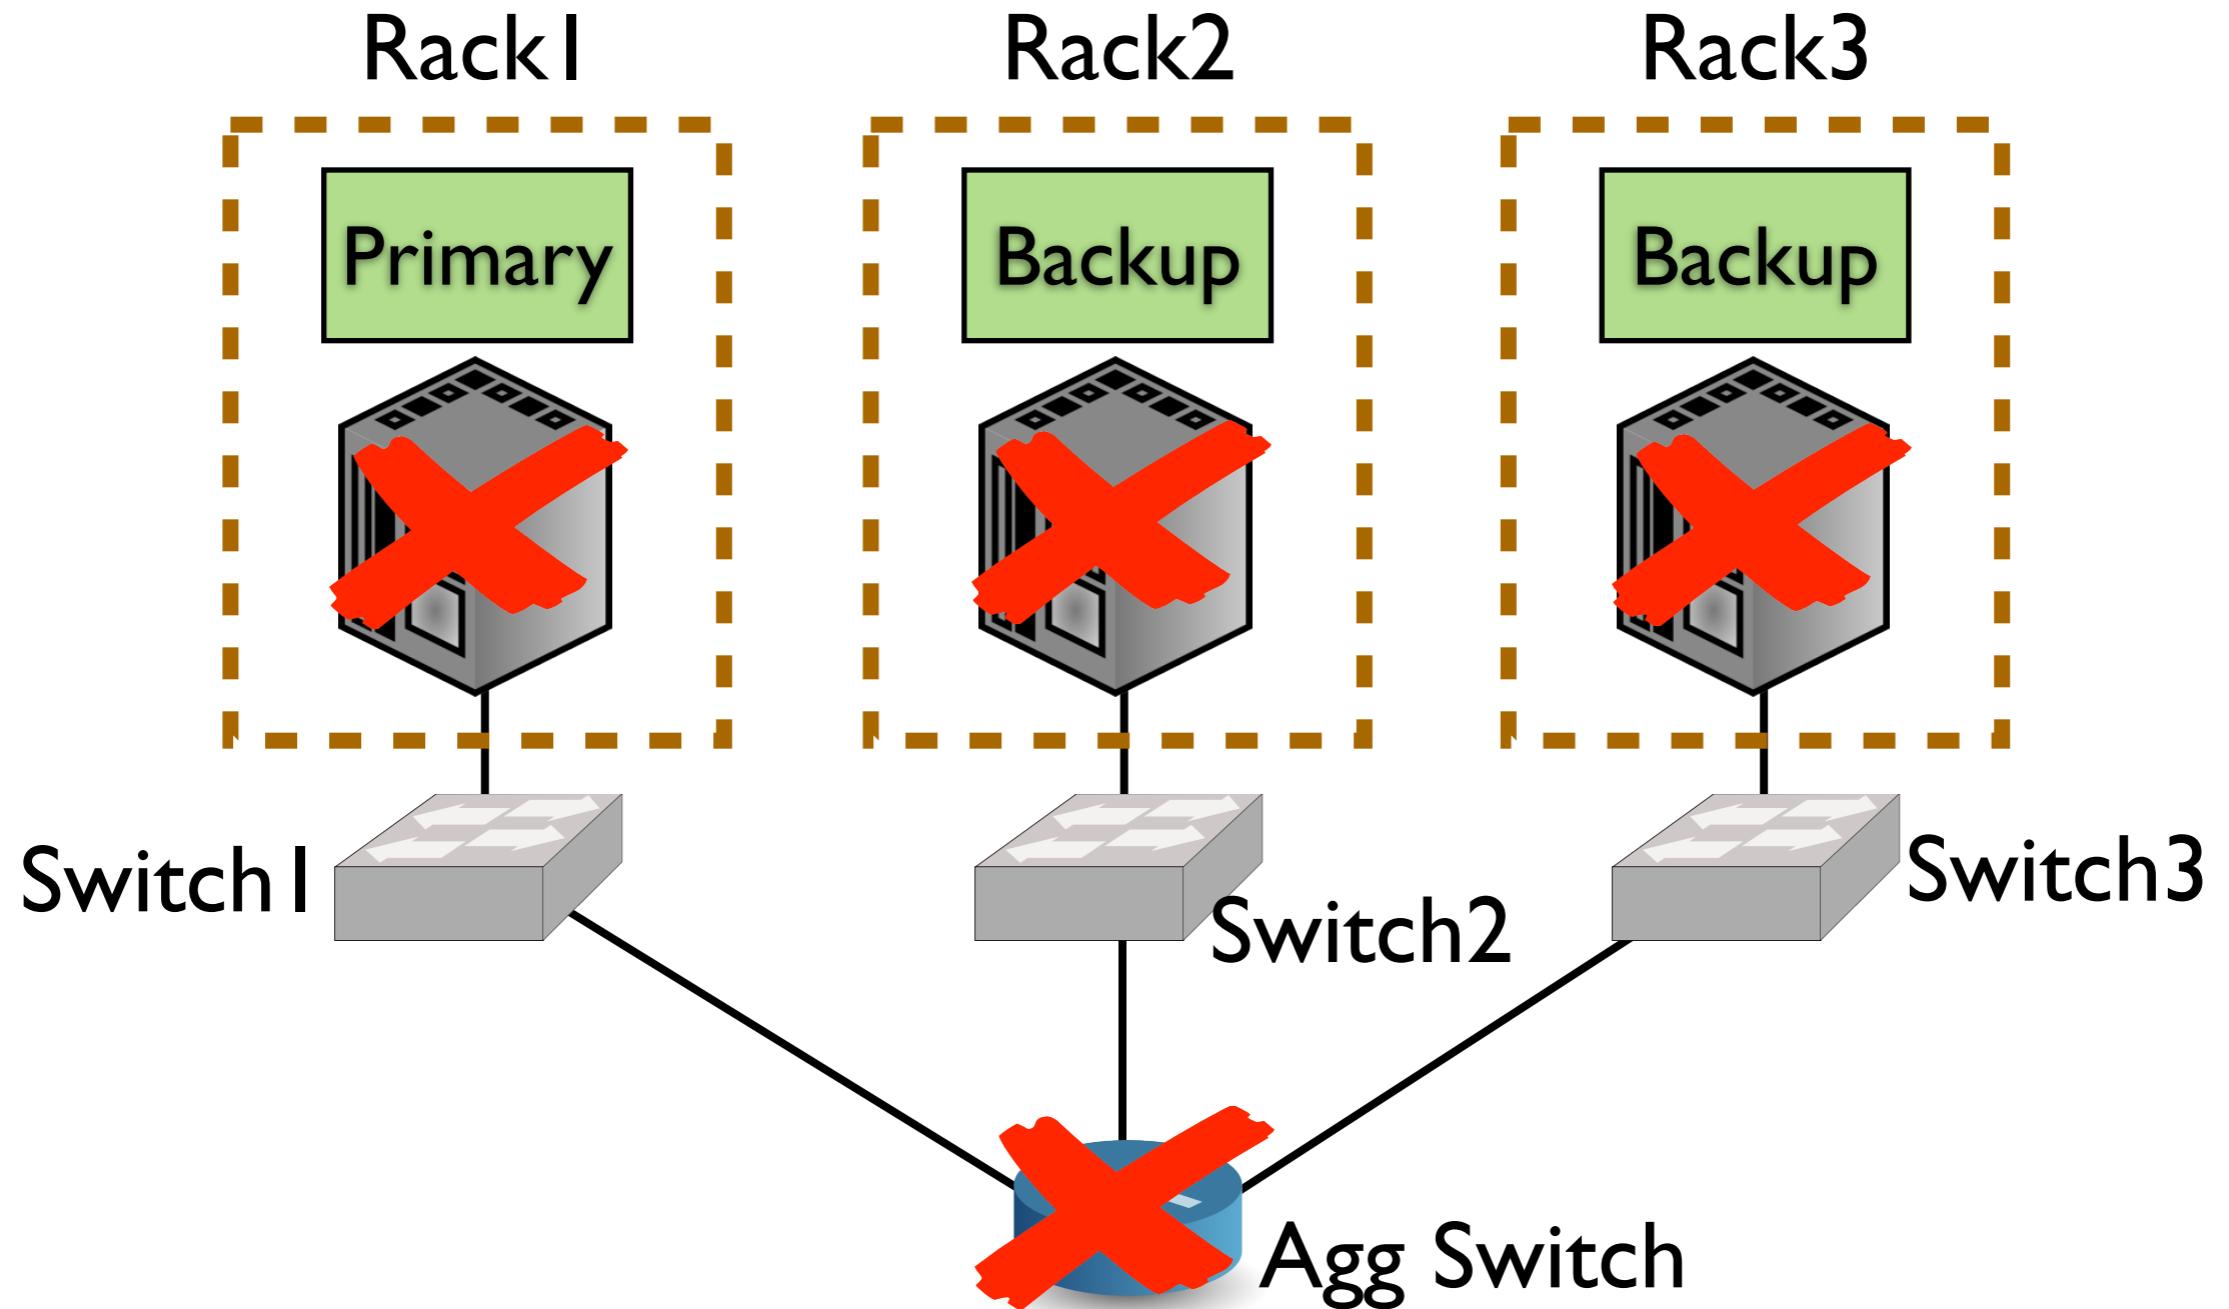
# What is correlated failure?



# What is correlated failure?



# What is correlated failure?



# Realistic Example

# Realistic Example



## Correlated failures resulting from EBS due to bugs in one EBS server

Summary of the October 22, 2012 AWS Service Event in the US-East Region

We'd like to share more about the service event that occurred on Monday, October 22nd in the US-East Region. We have now completed the analysis of the events that affected AWS customers, and we want to describe what happened, our understanding of how customers were affected, and what we are doing to prevent a similar issue from occurring in the future.

**The Primary Event and the Impact to Amazon Elastic Block Store (EBS) and Amazon Elastic Compute Cloud (EC2)**

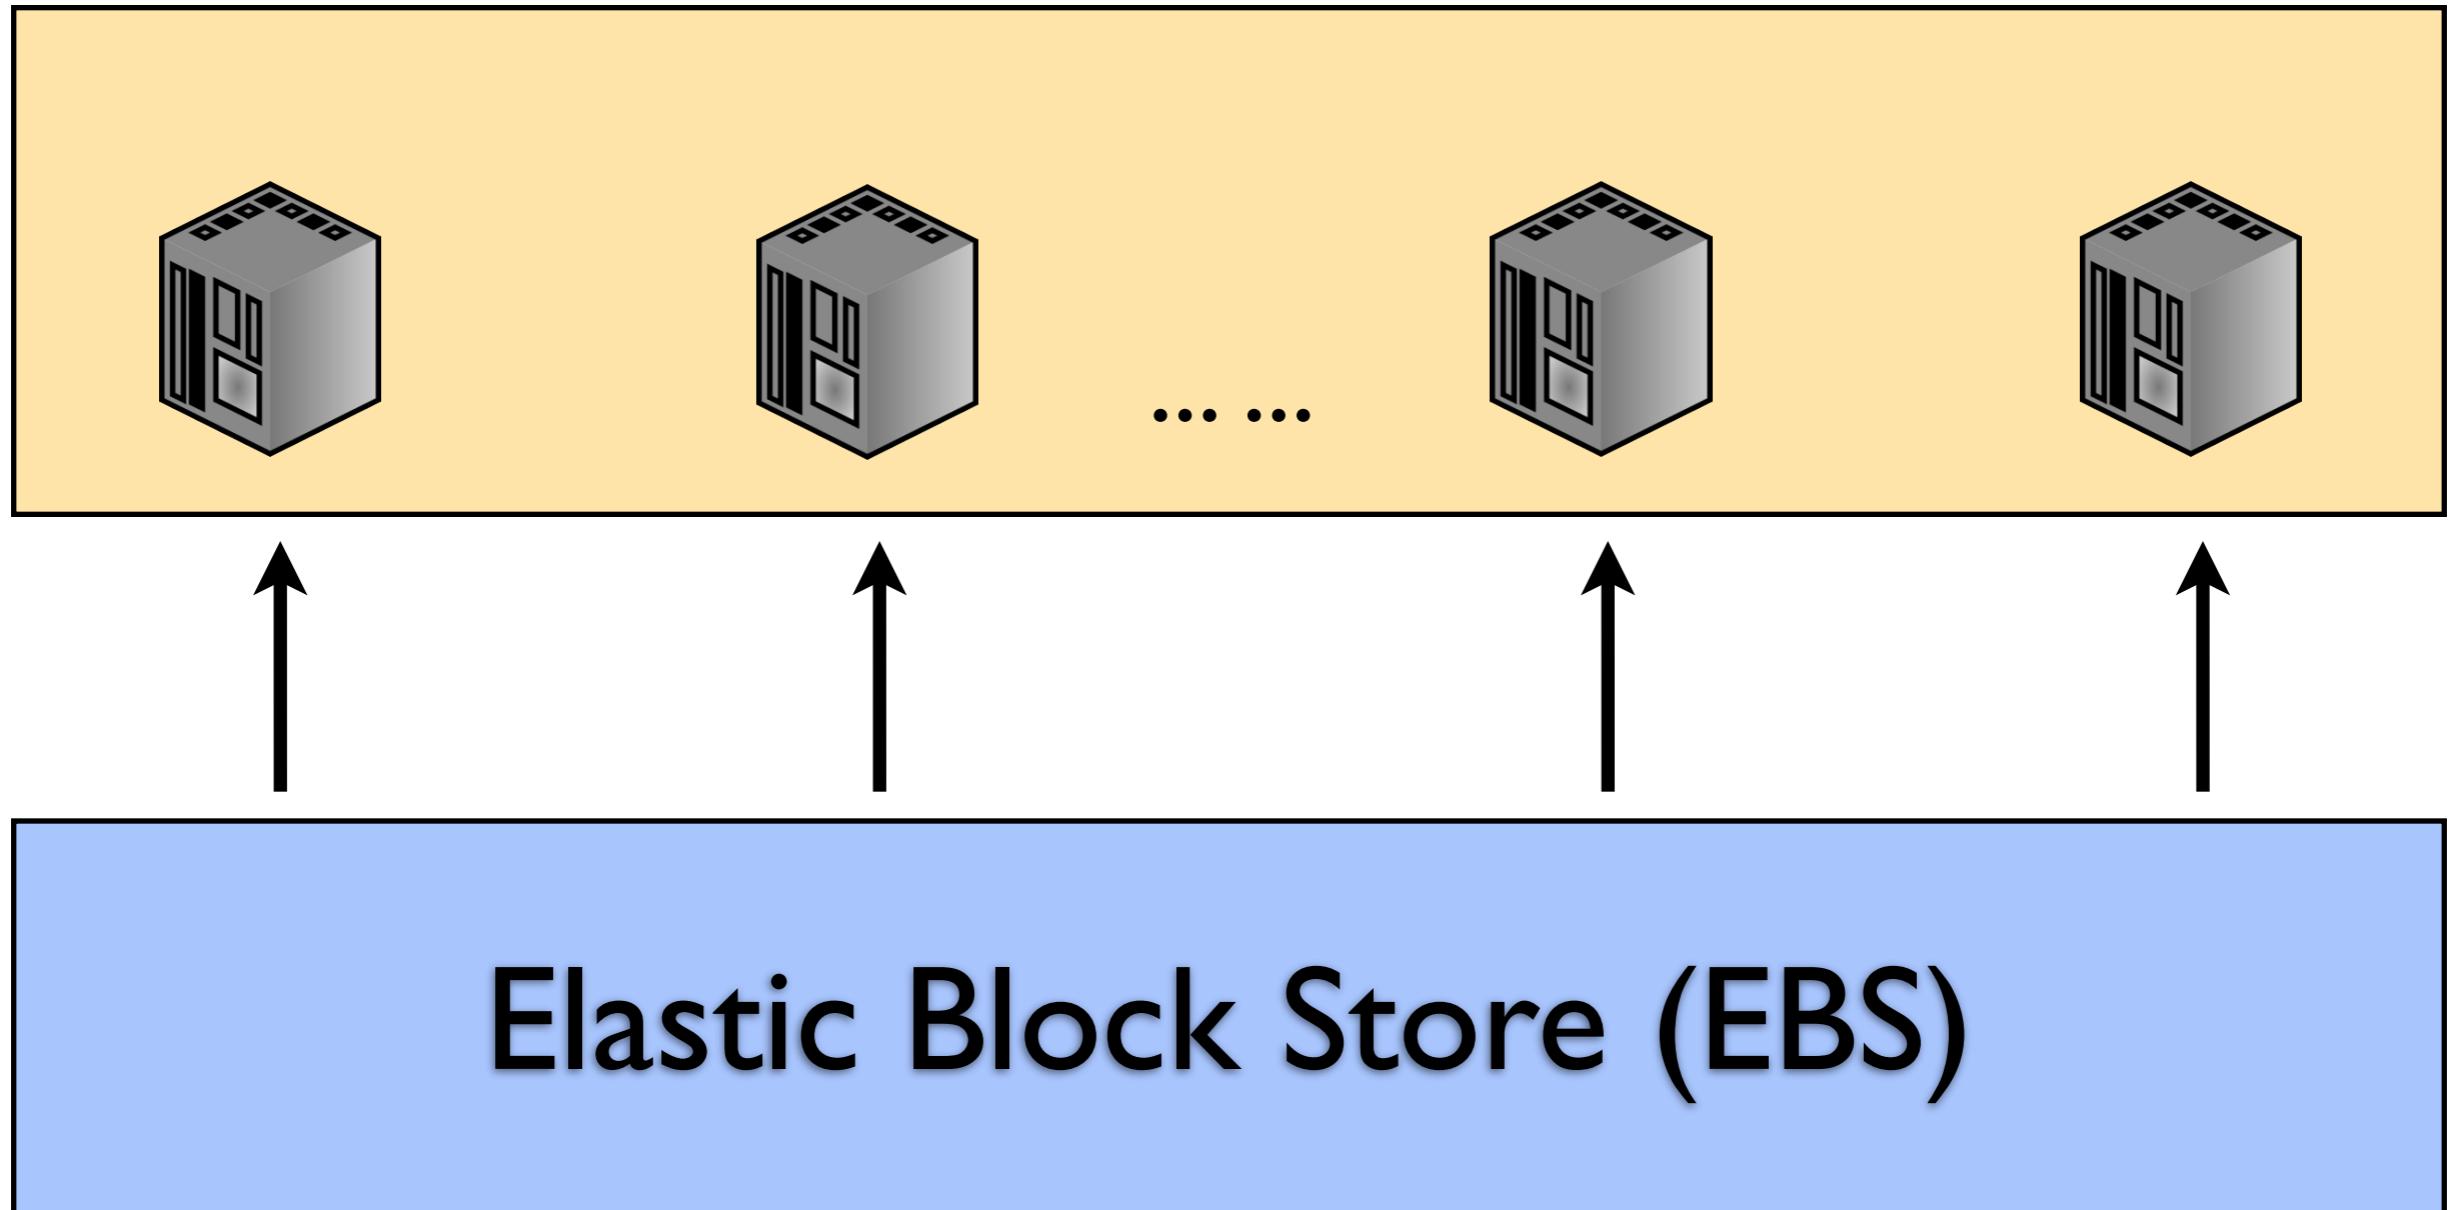
# Realistic Example

Elastic Compute Cloud (EC2)

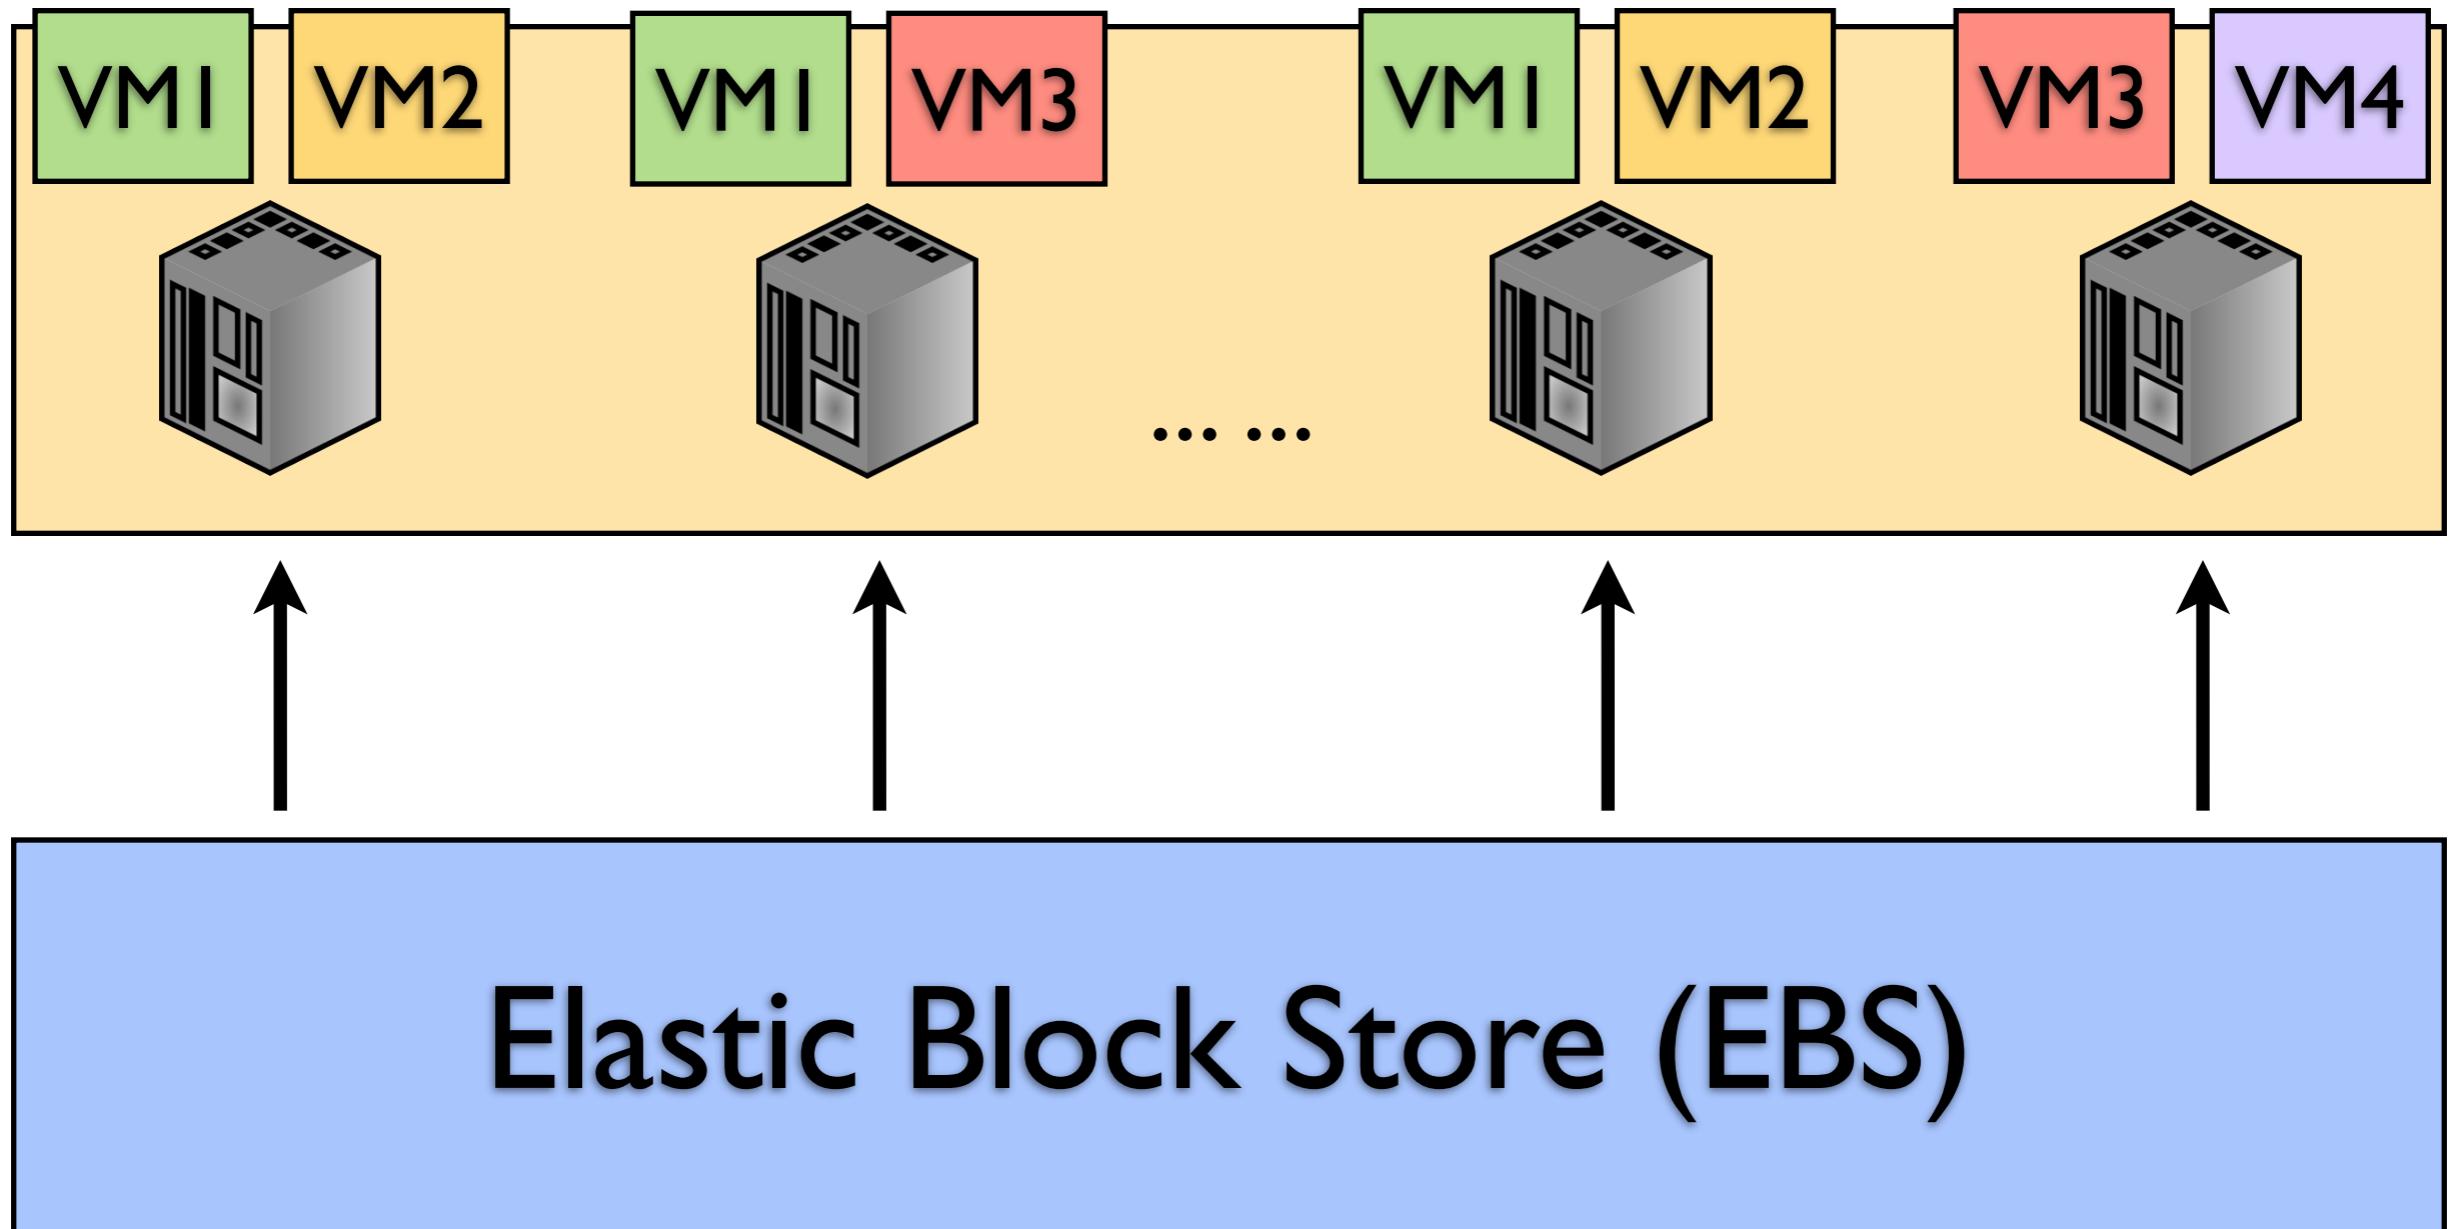


Elastic Block Store (EBS)

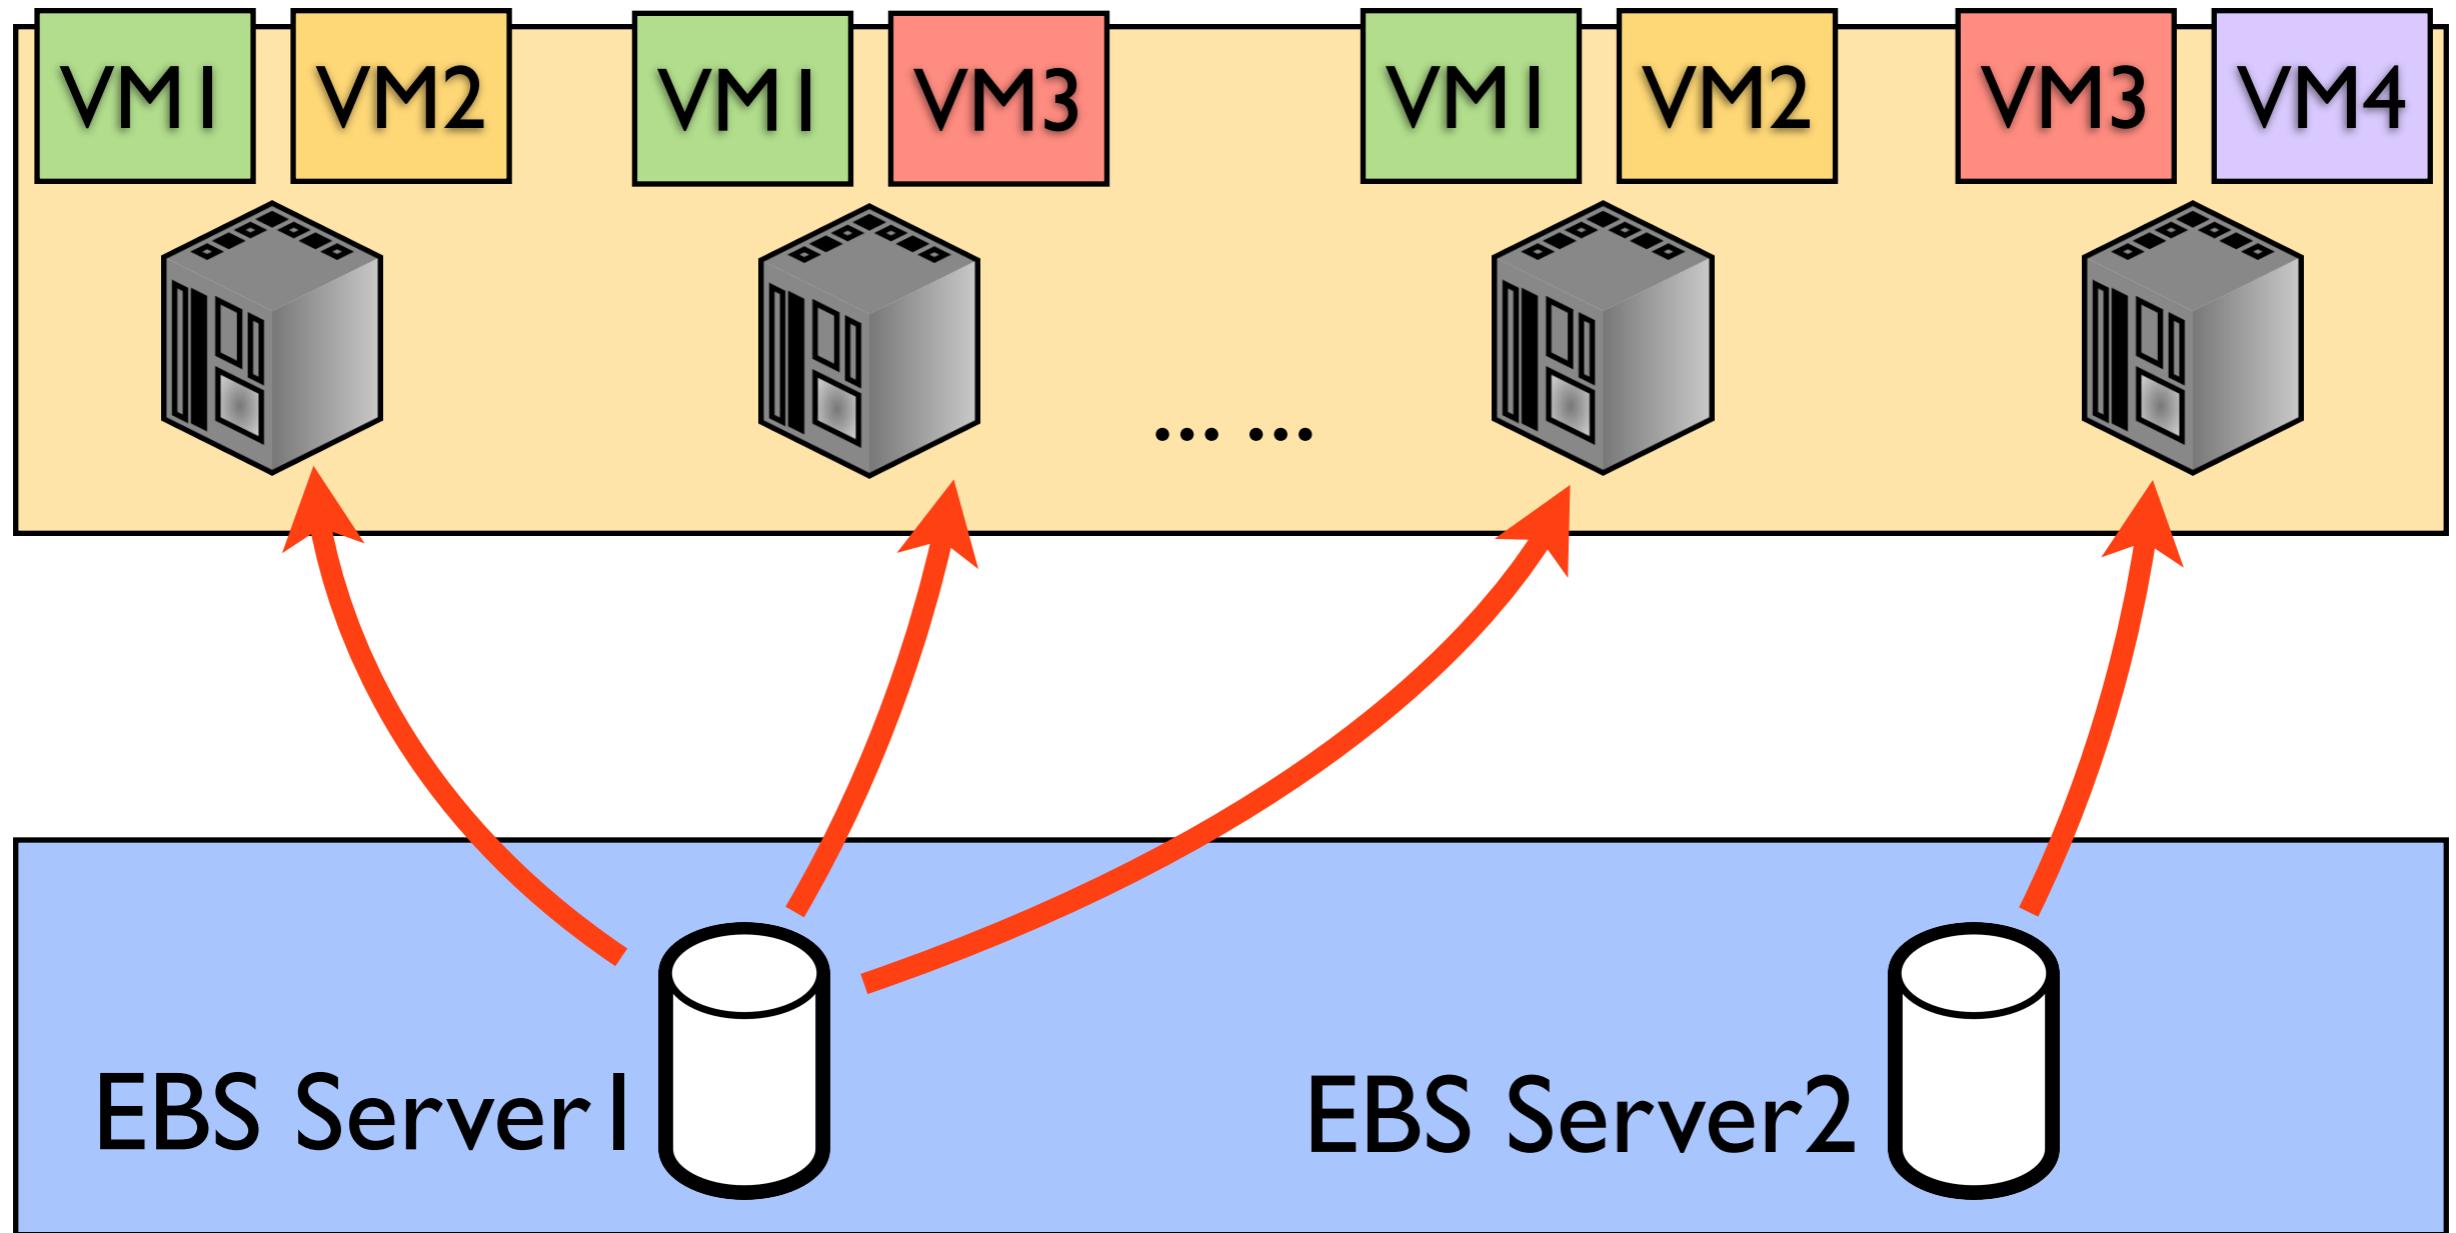
# Realistic Example



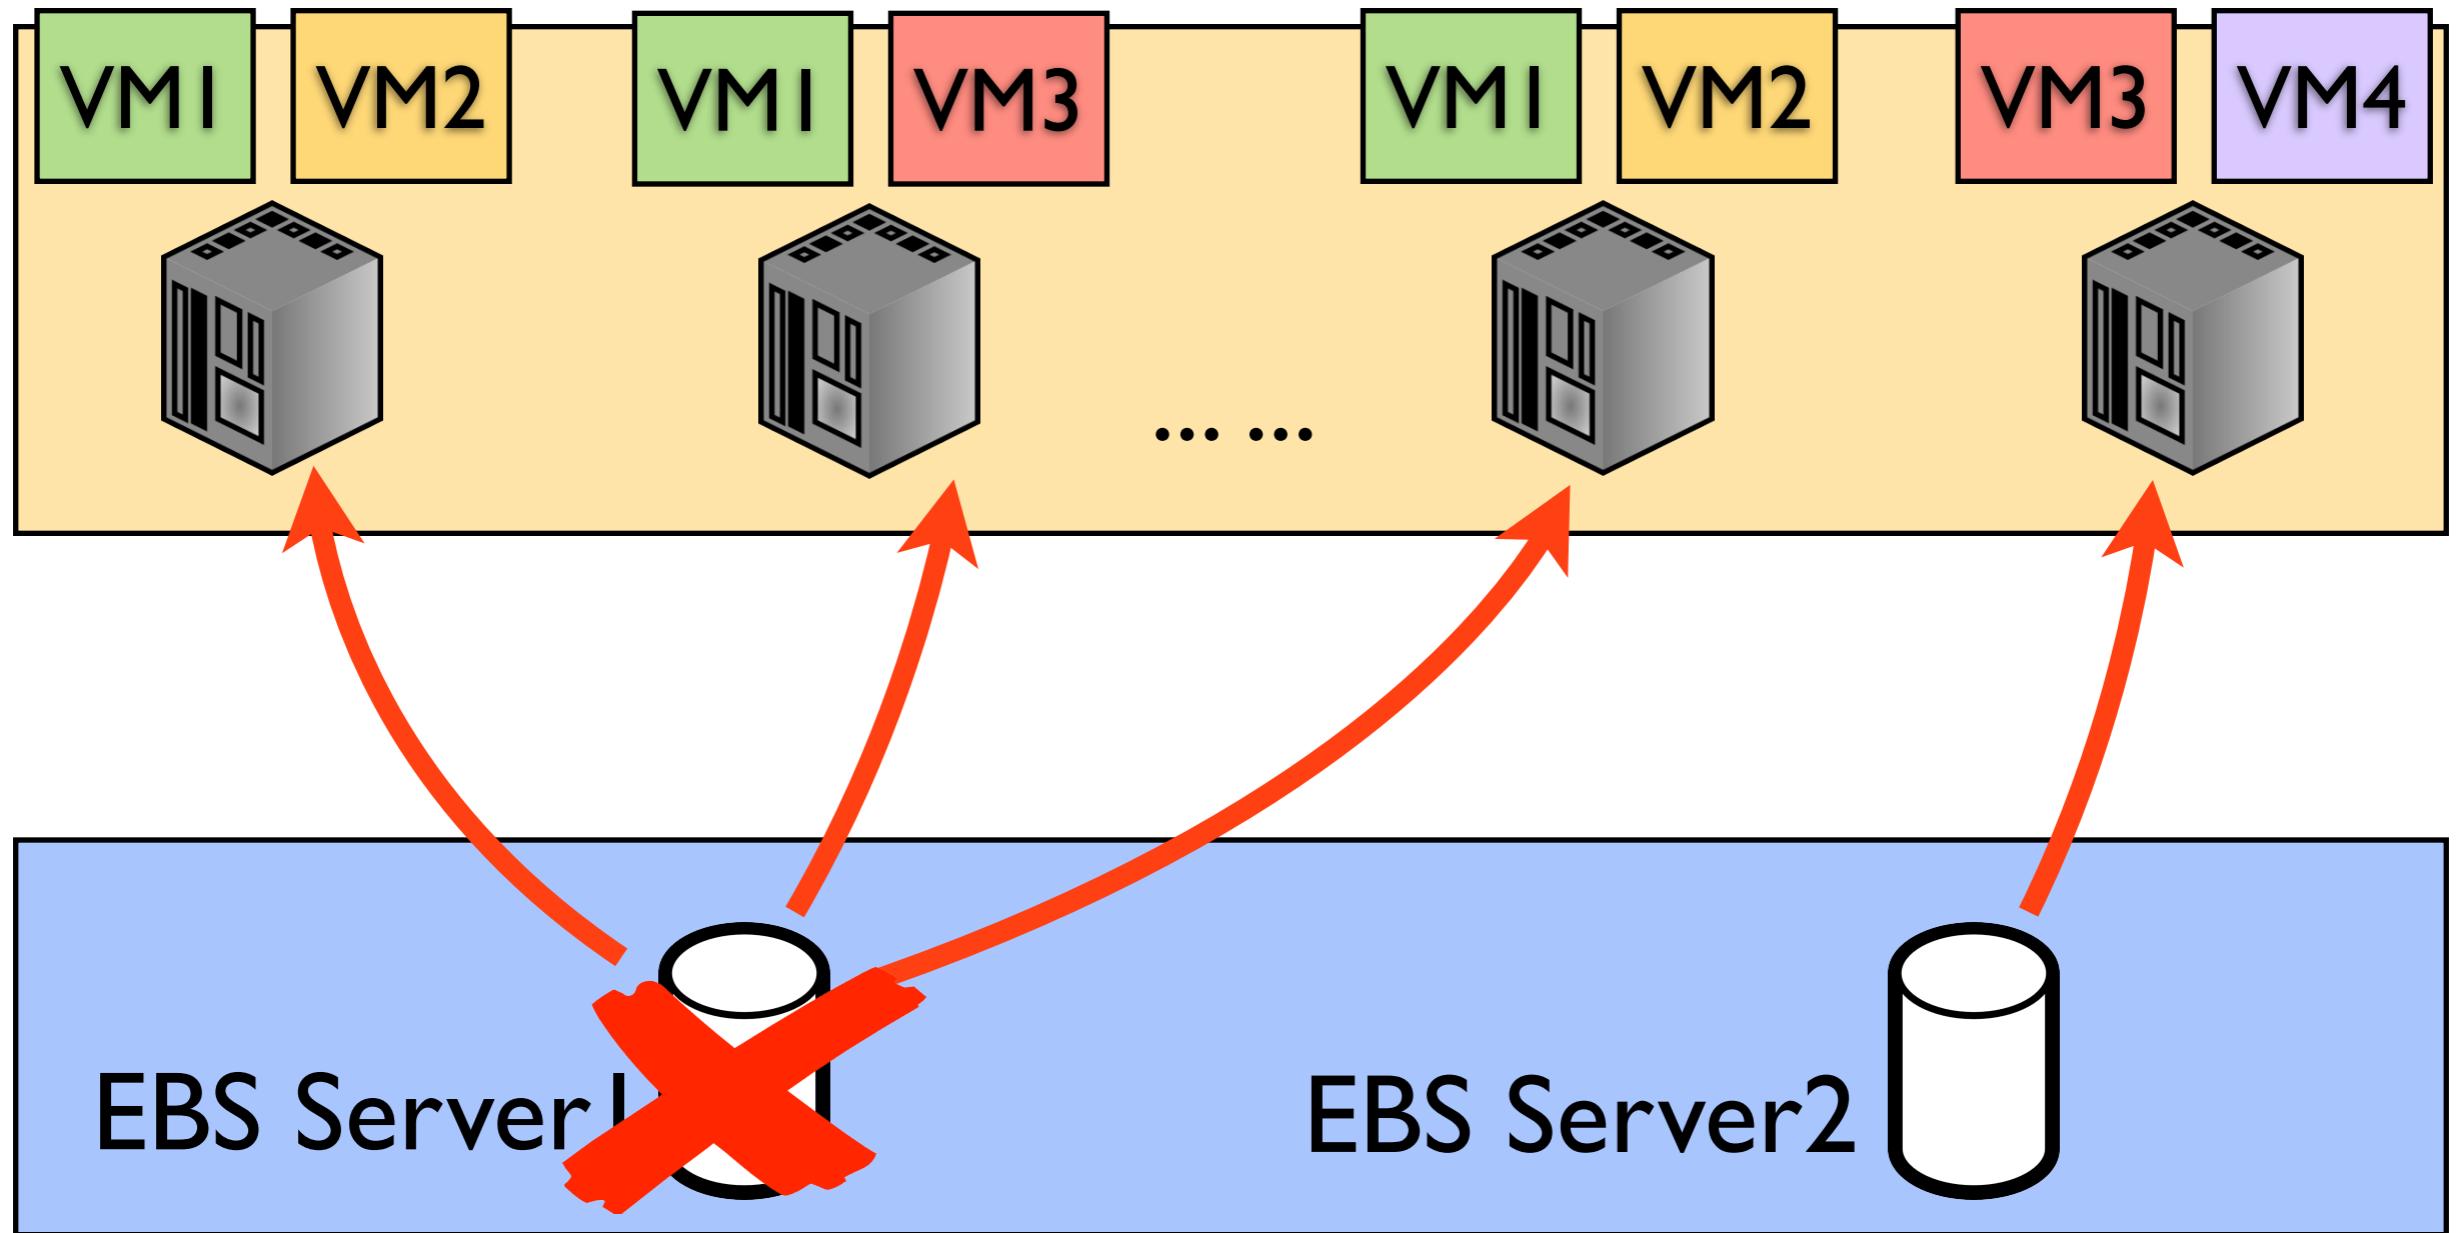
# Realistic Example



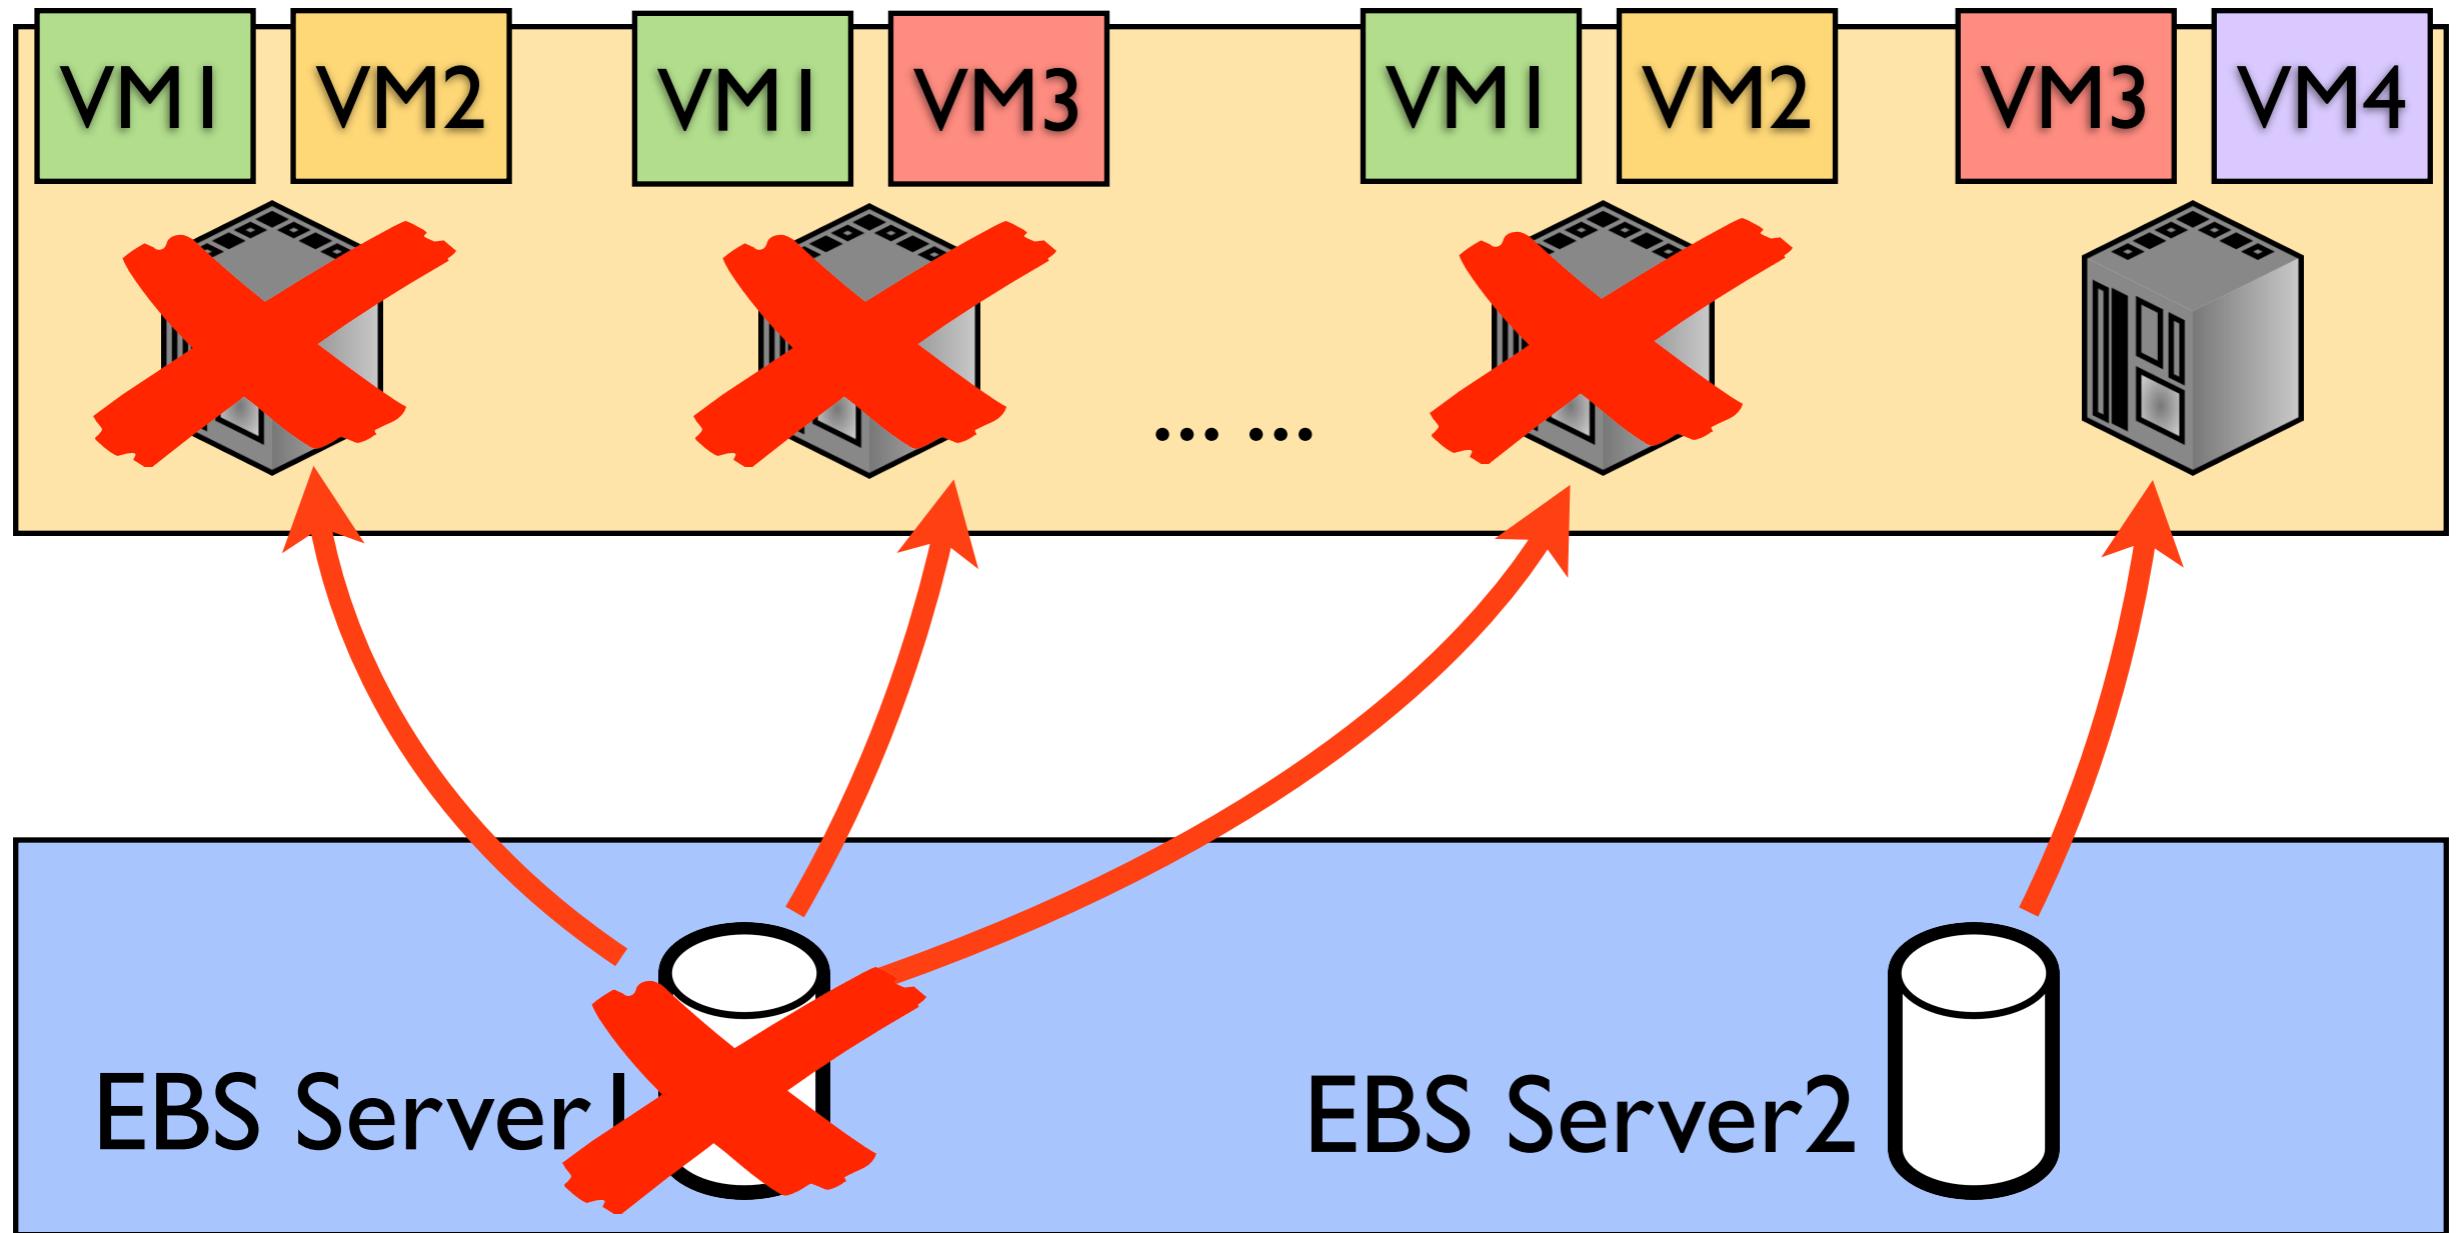
# Realistic Example



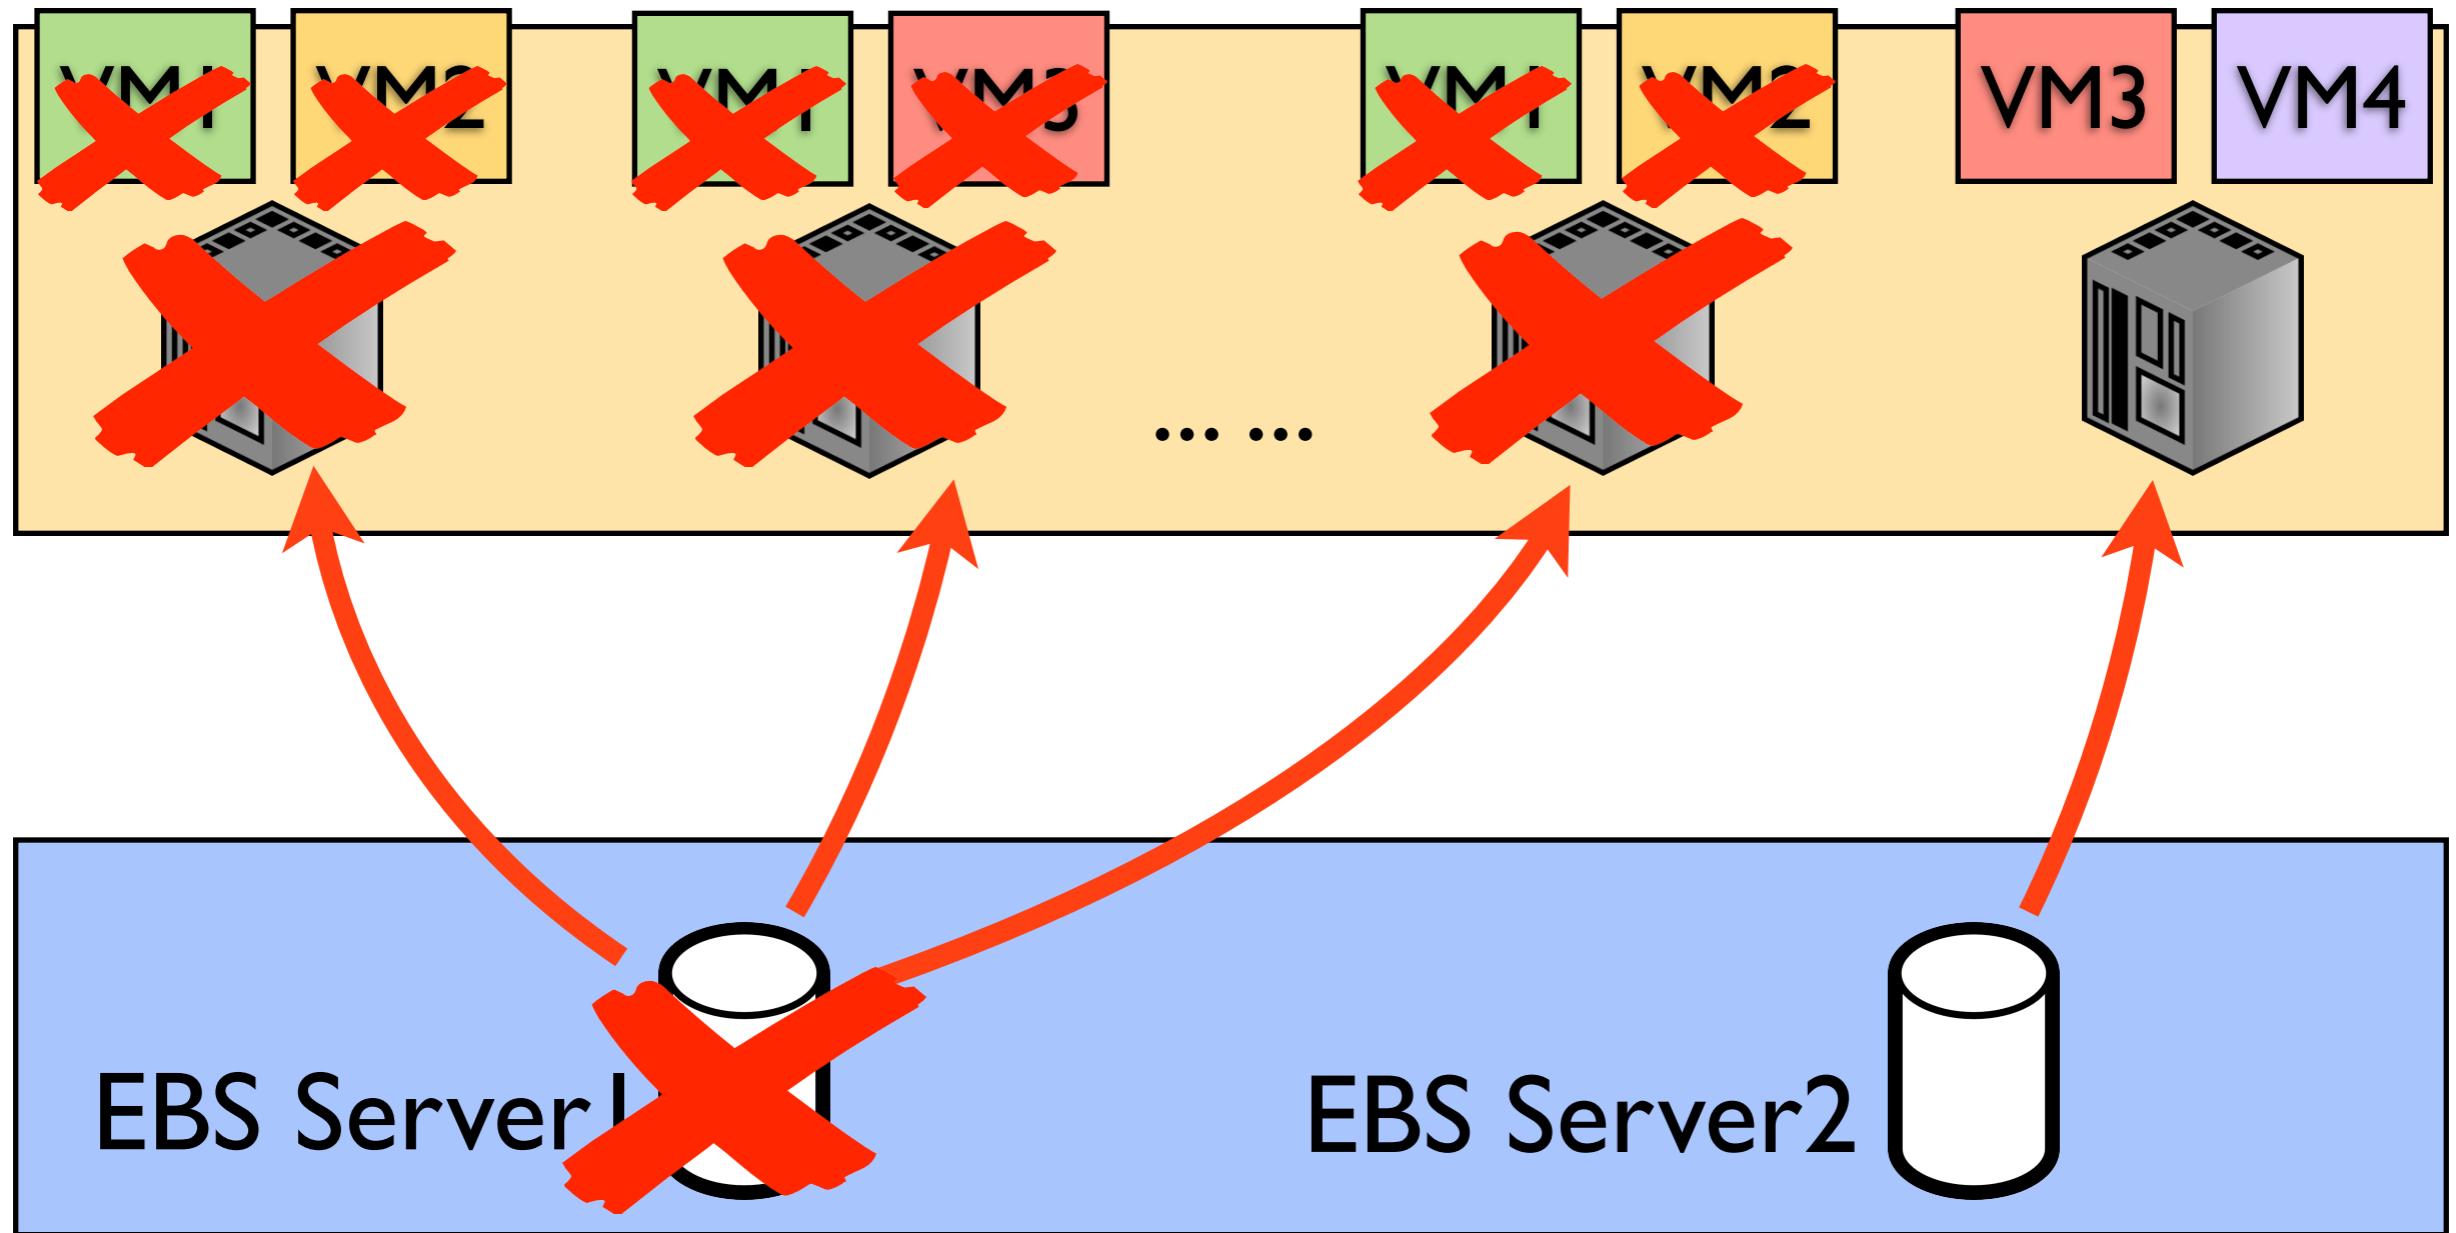
# Realistic Example



# Realistic Example



# Realistic Example

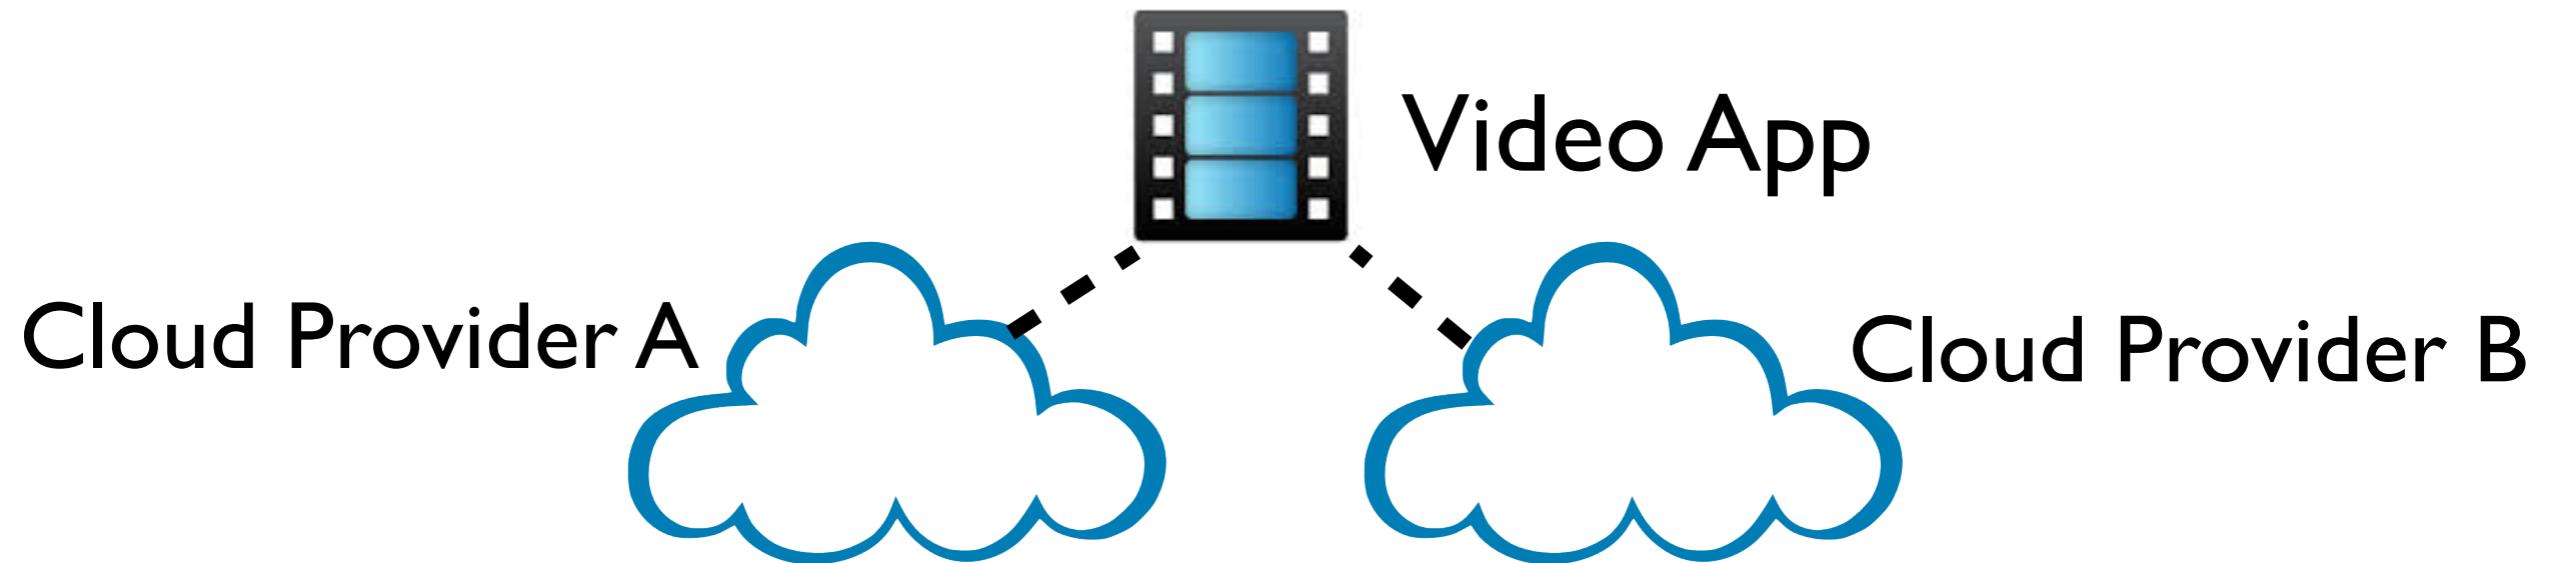


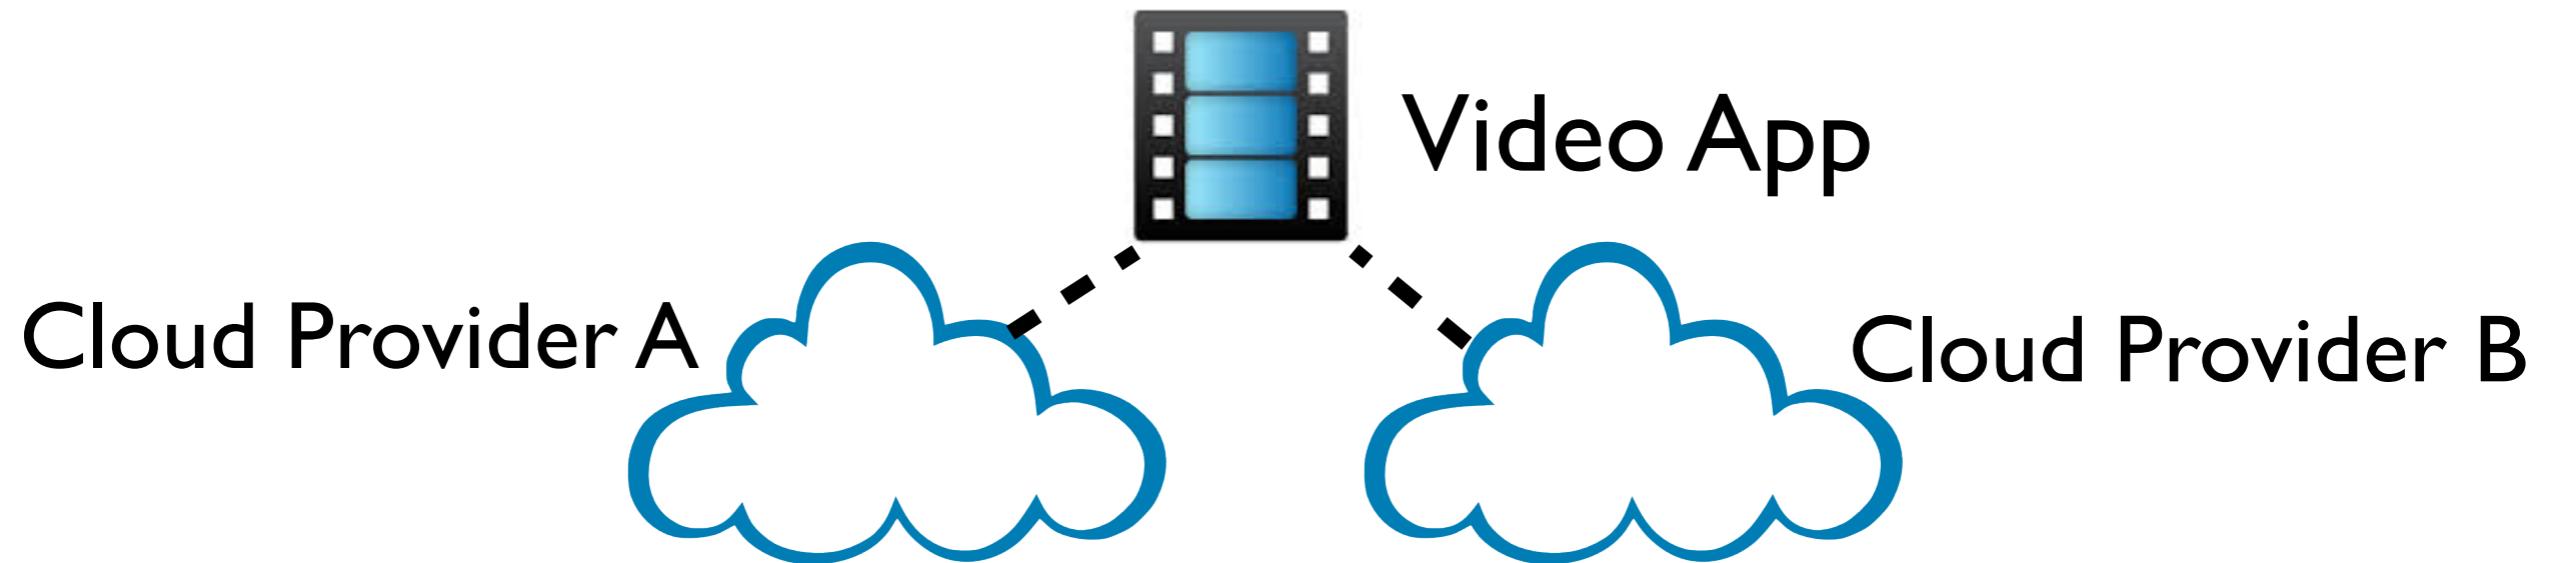
# Even Worse

# Lightning strikes Amazon's European cloud

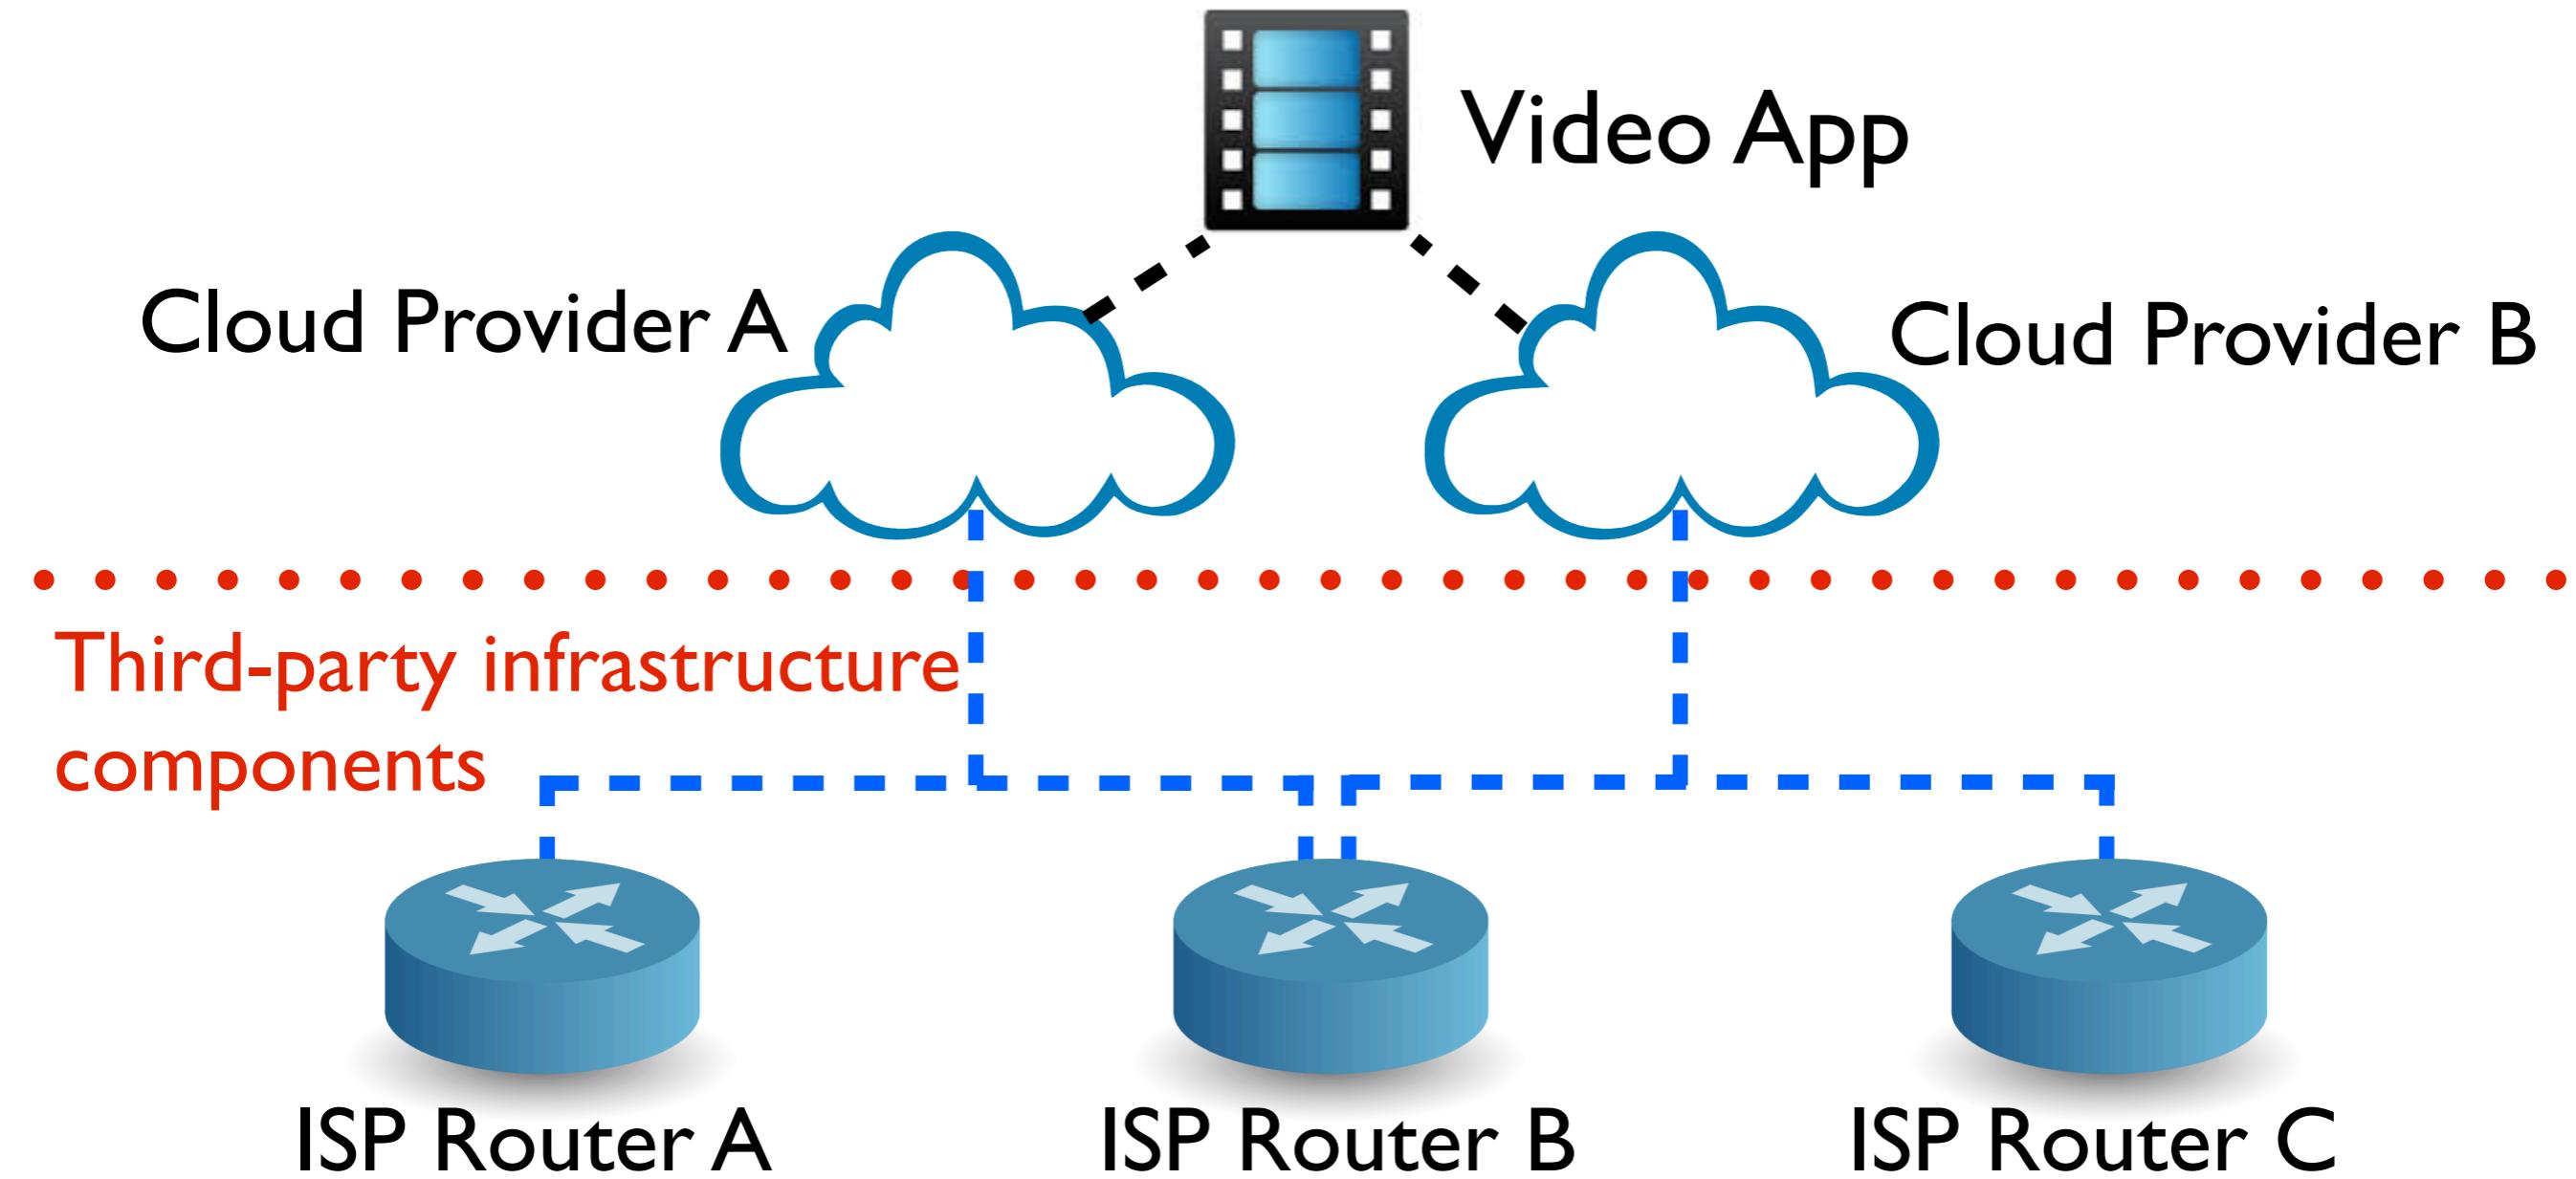
**Summary:** *The lightning strike damaged a power company's transformer, causing disruption to Amazon Web Services's European cloud, and may have affected Microsoft's BPOS as well*

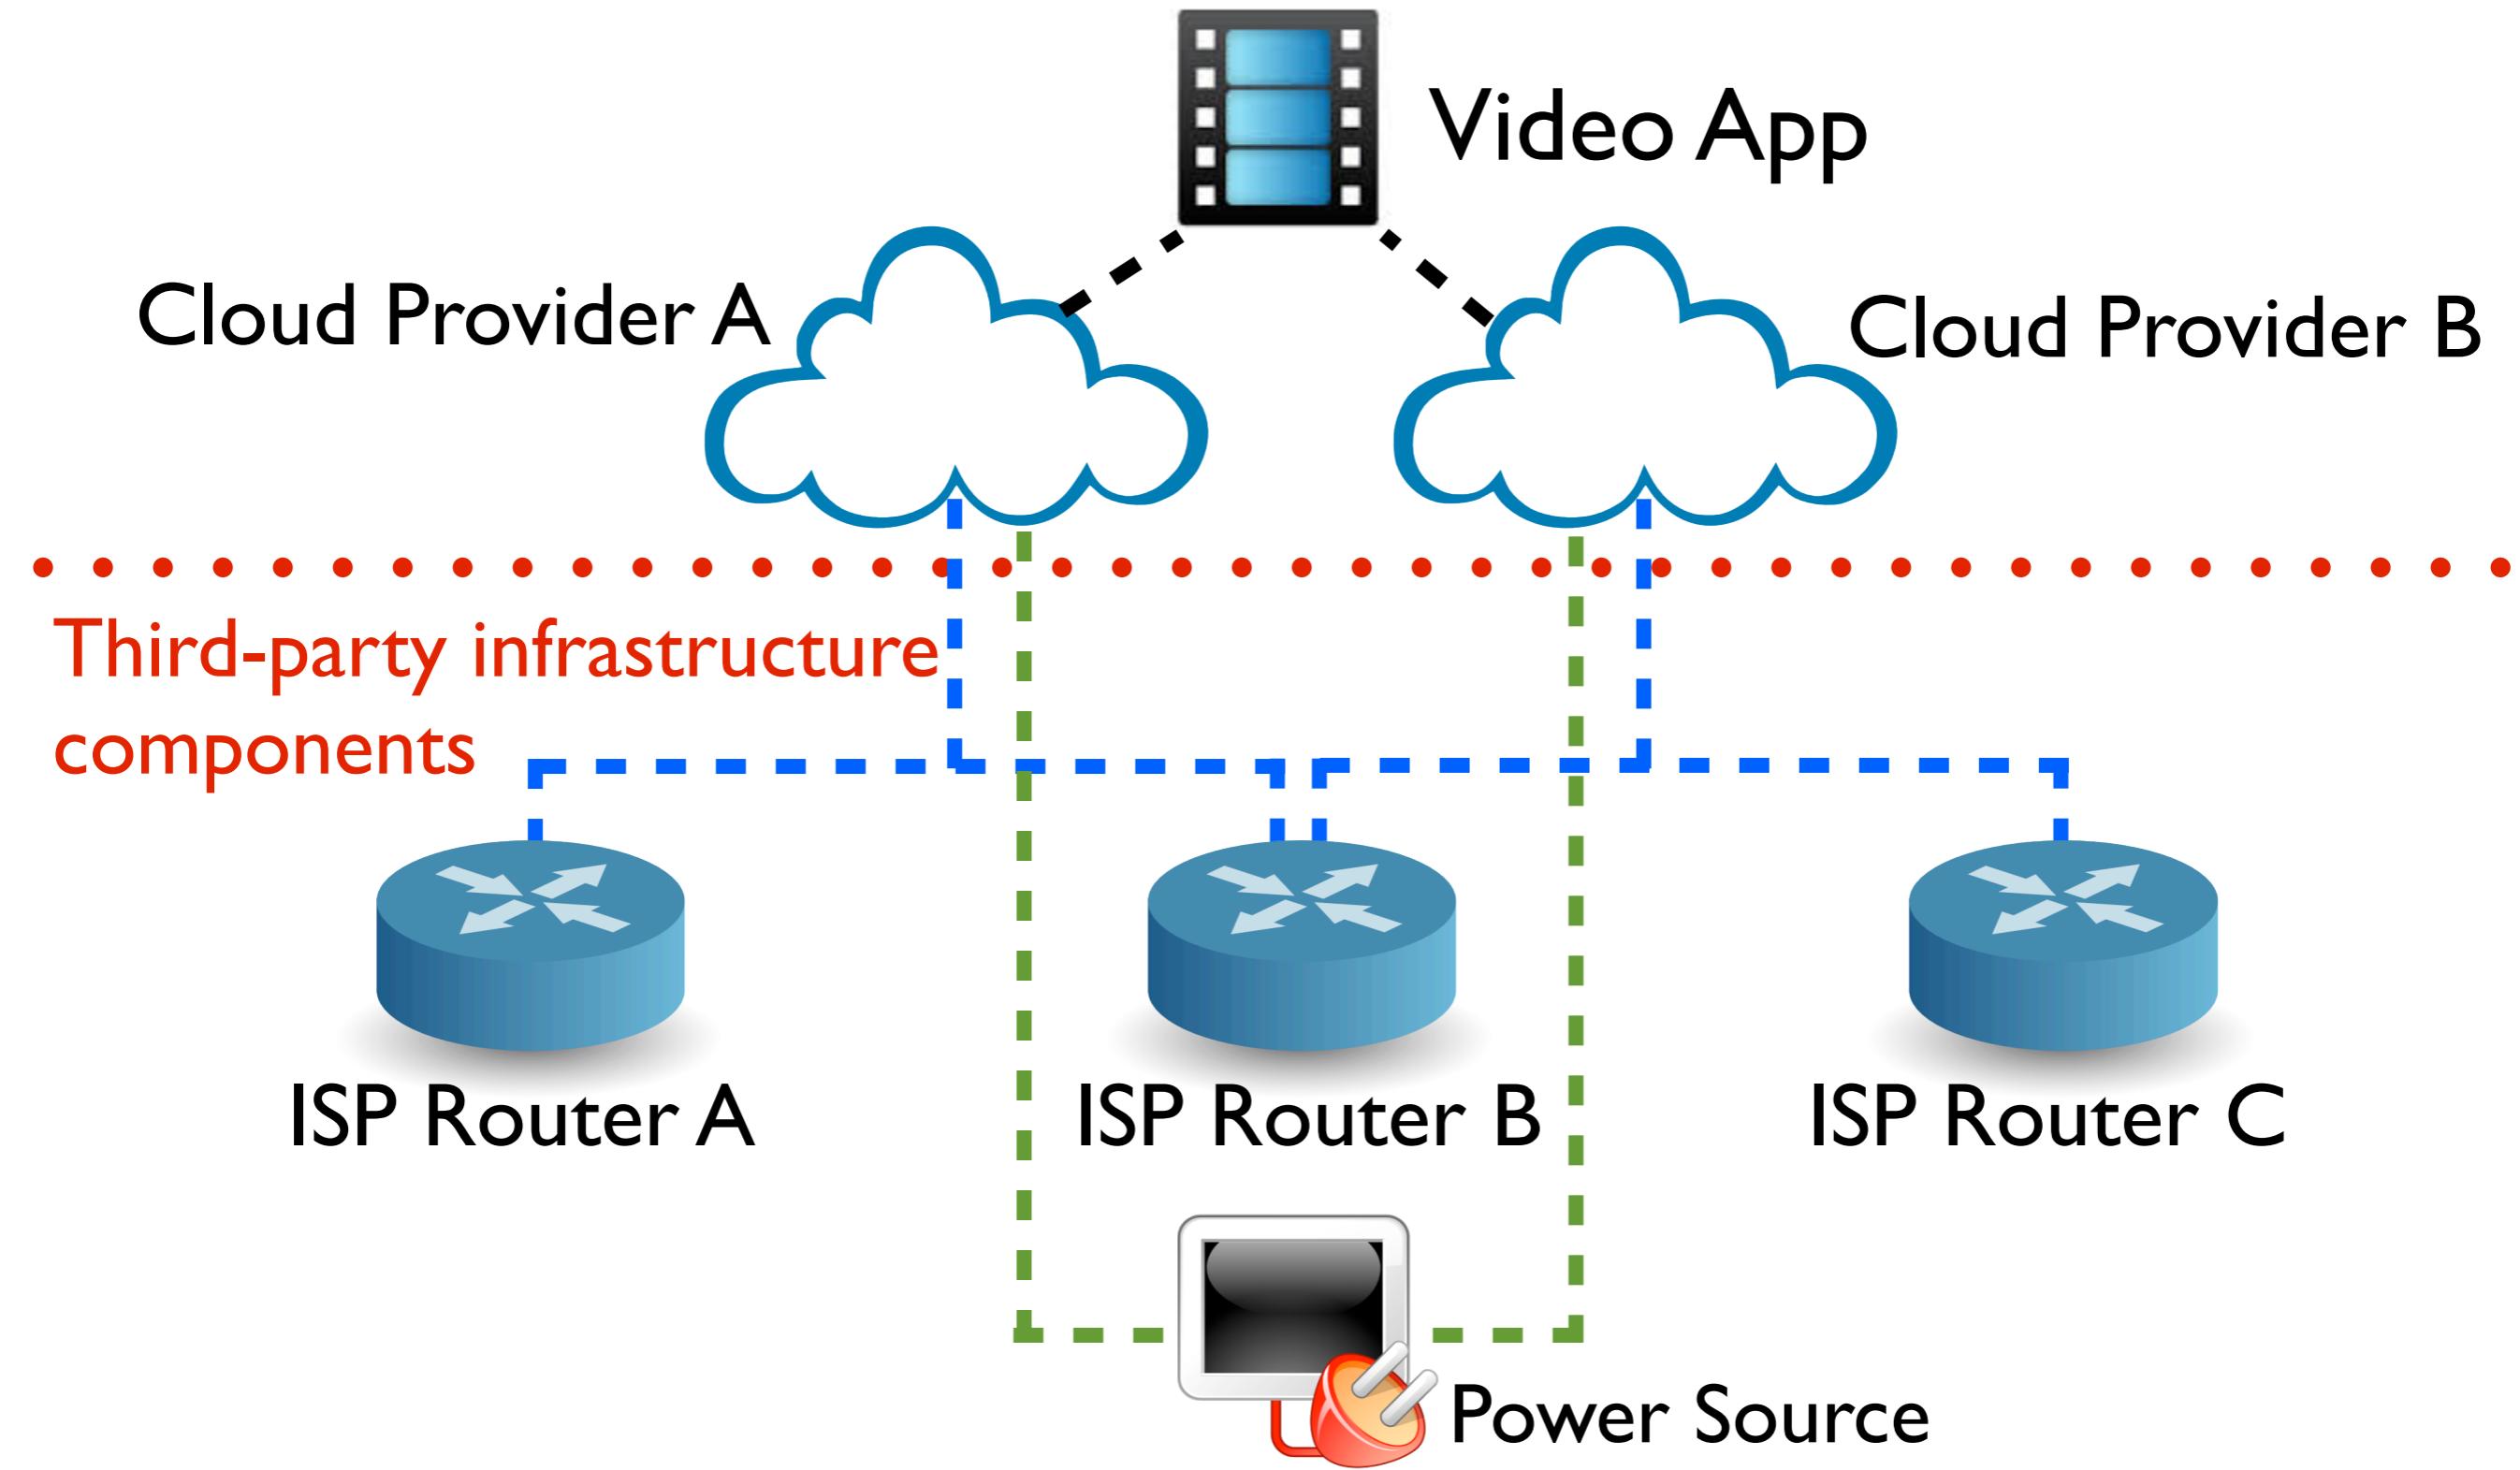
The outage, which [Amazon Web Services \(AWS\)](#) acknowledged on Sunday evening, affected its Dublin-based Elastic Compute Cloud (EC2) and Relational Database Service (RDS) cloud services, among others. The damage to the electricity infrastructure may have affected Microsoft's [Business Productivity Online Services \(BPOS\)](#) cloud as well, Microsoft said in a separate statement.

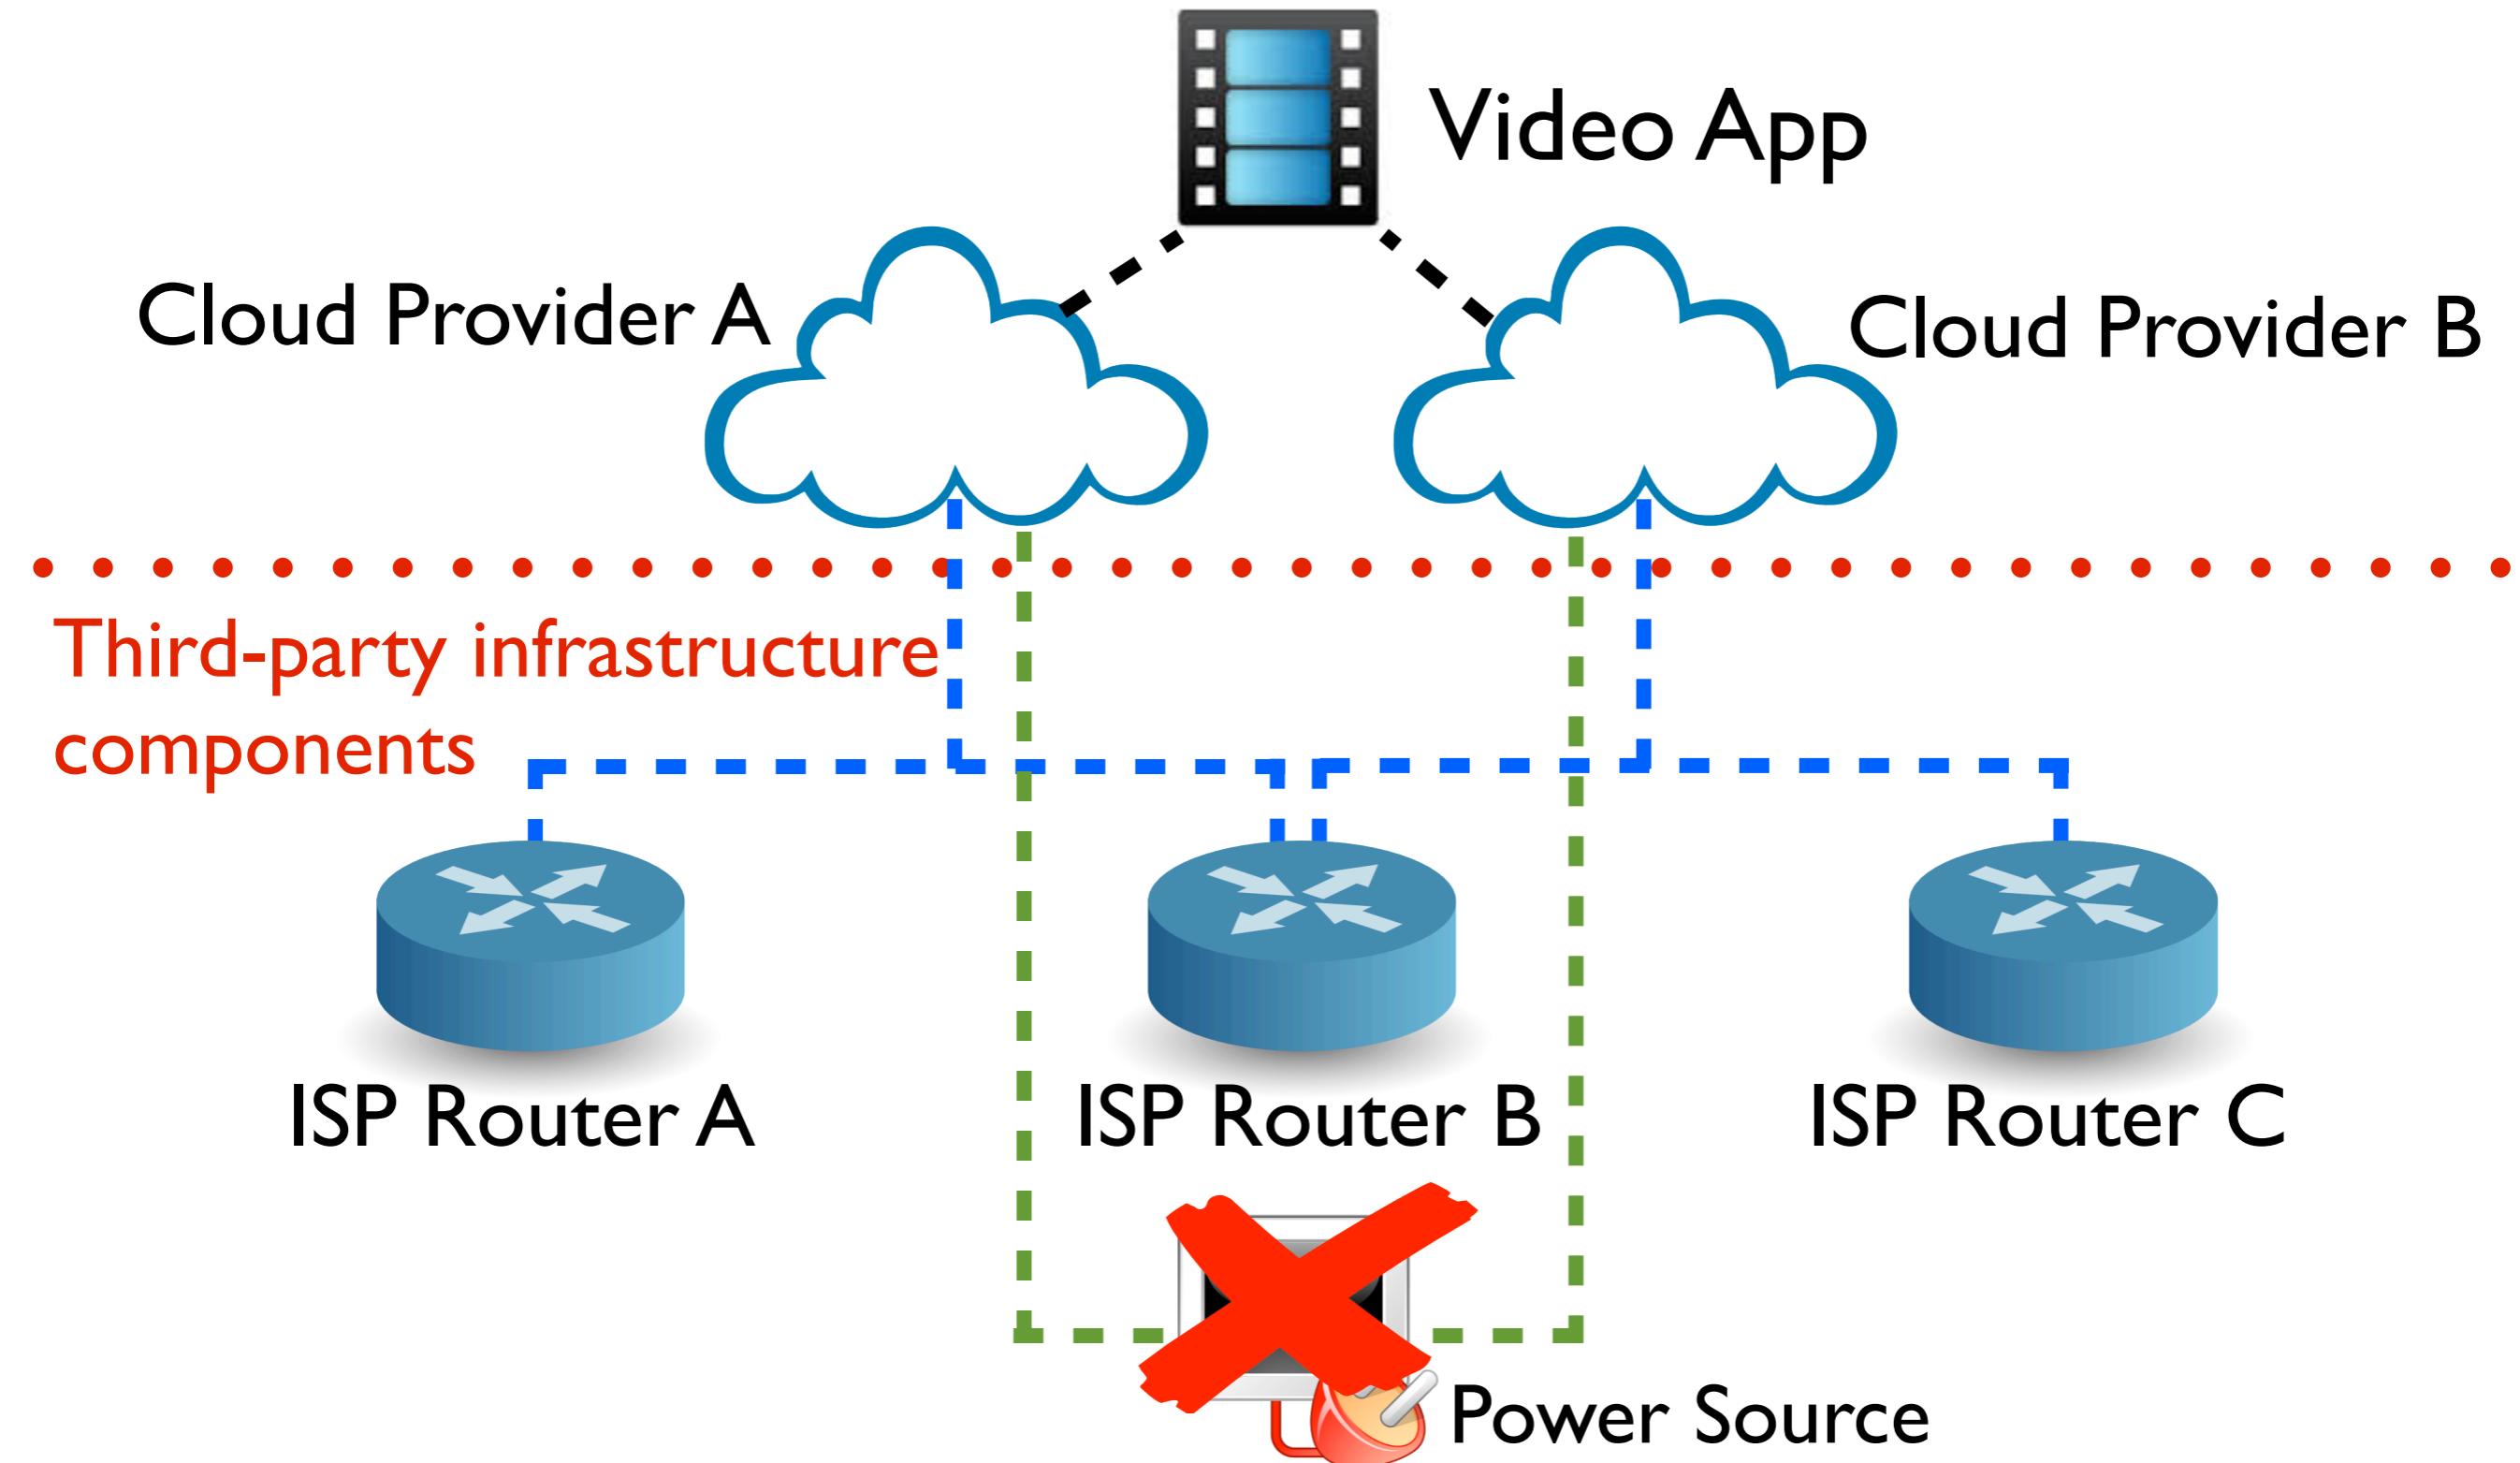


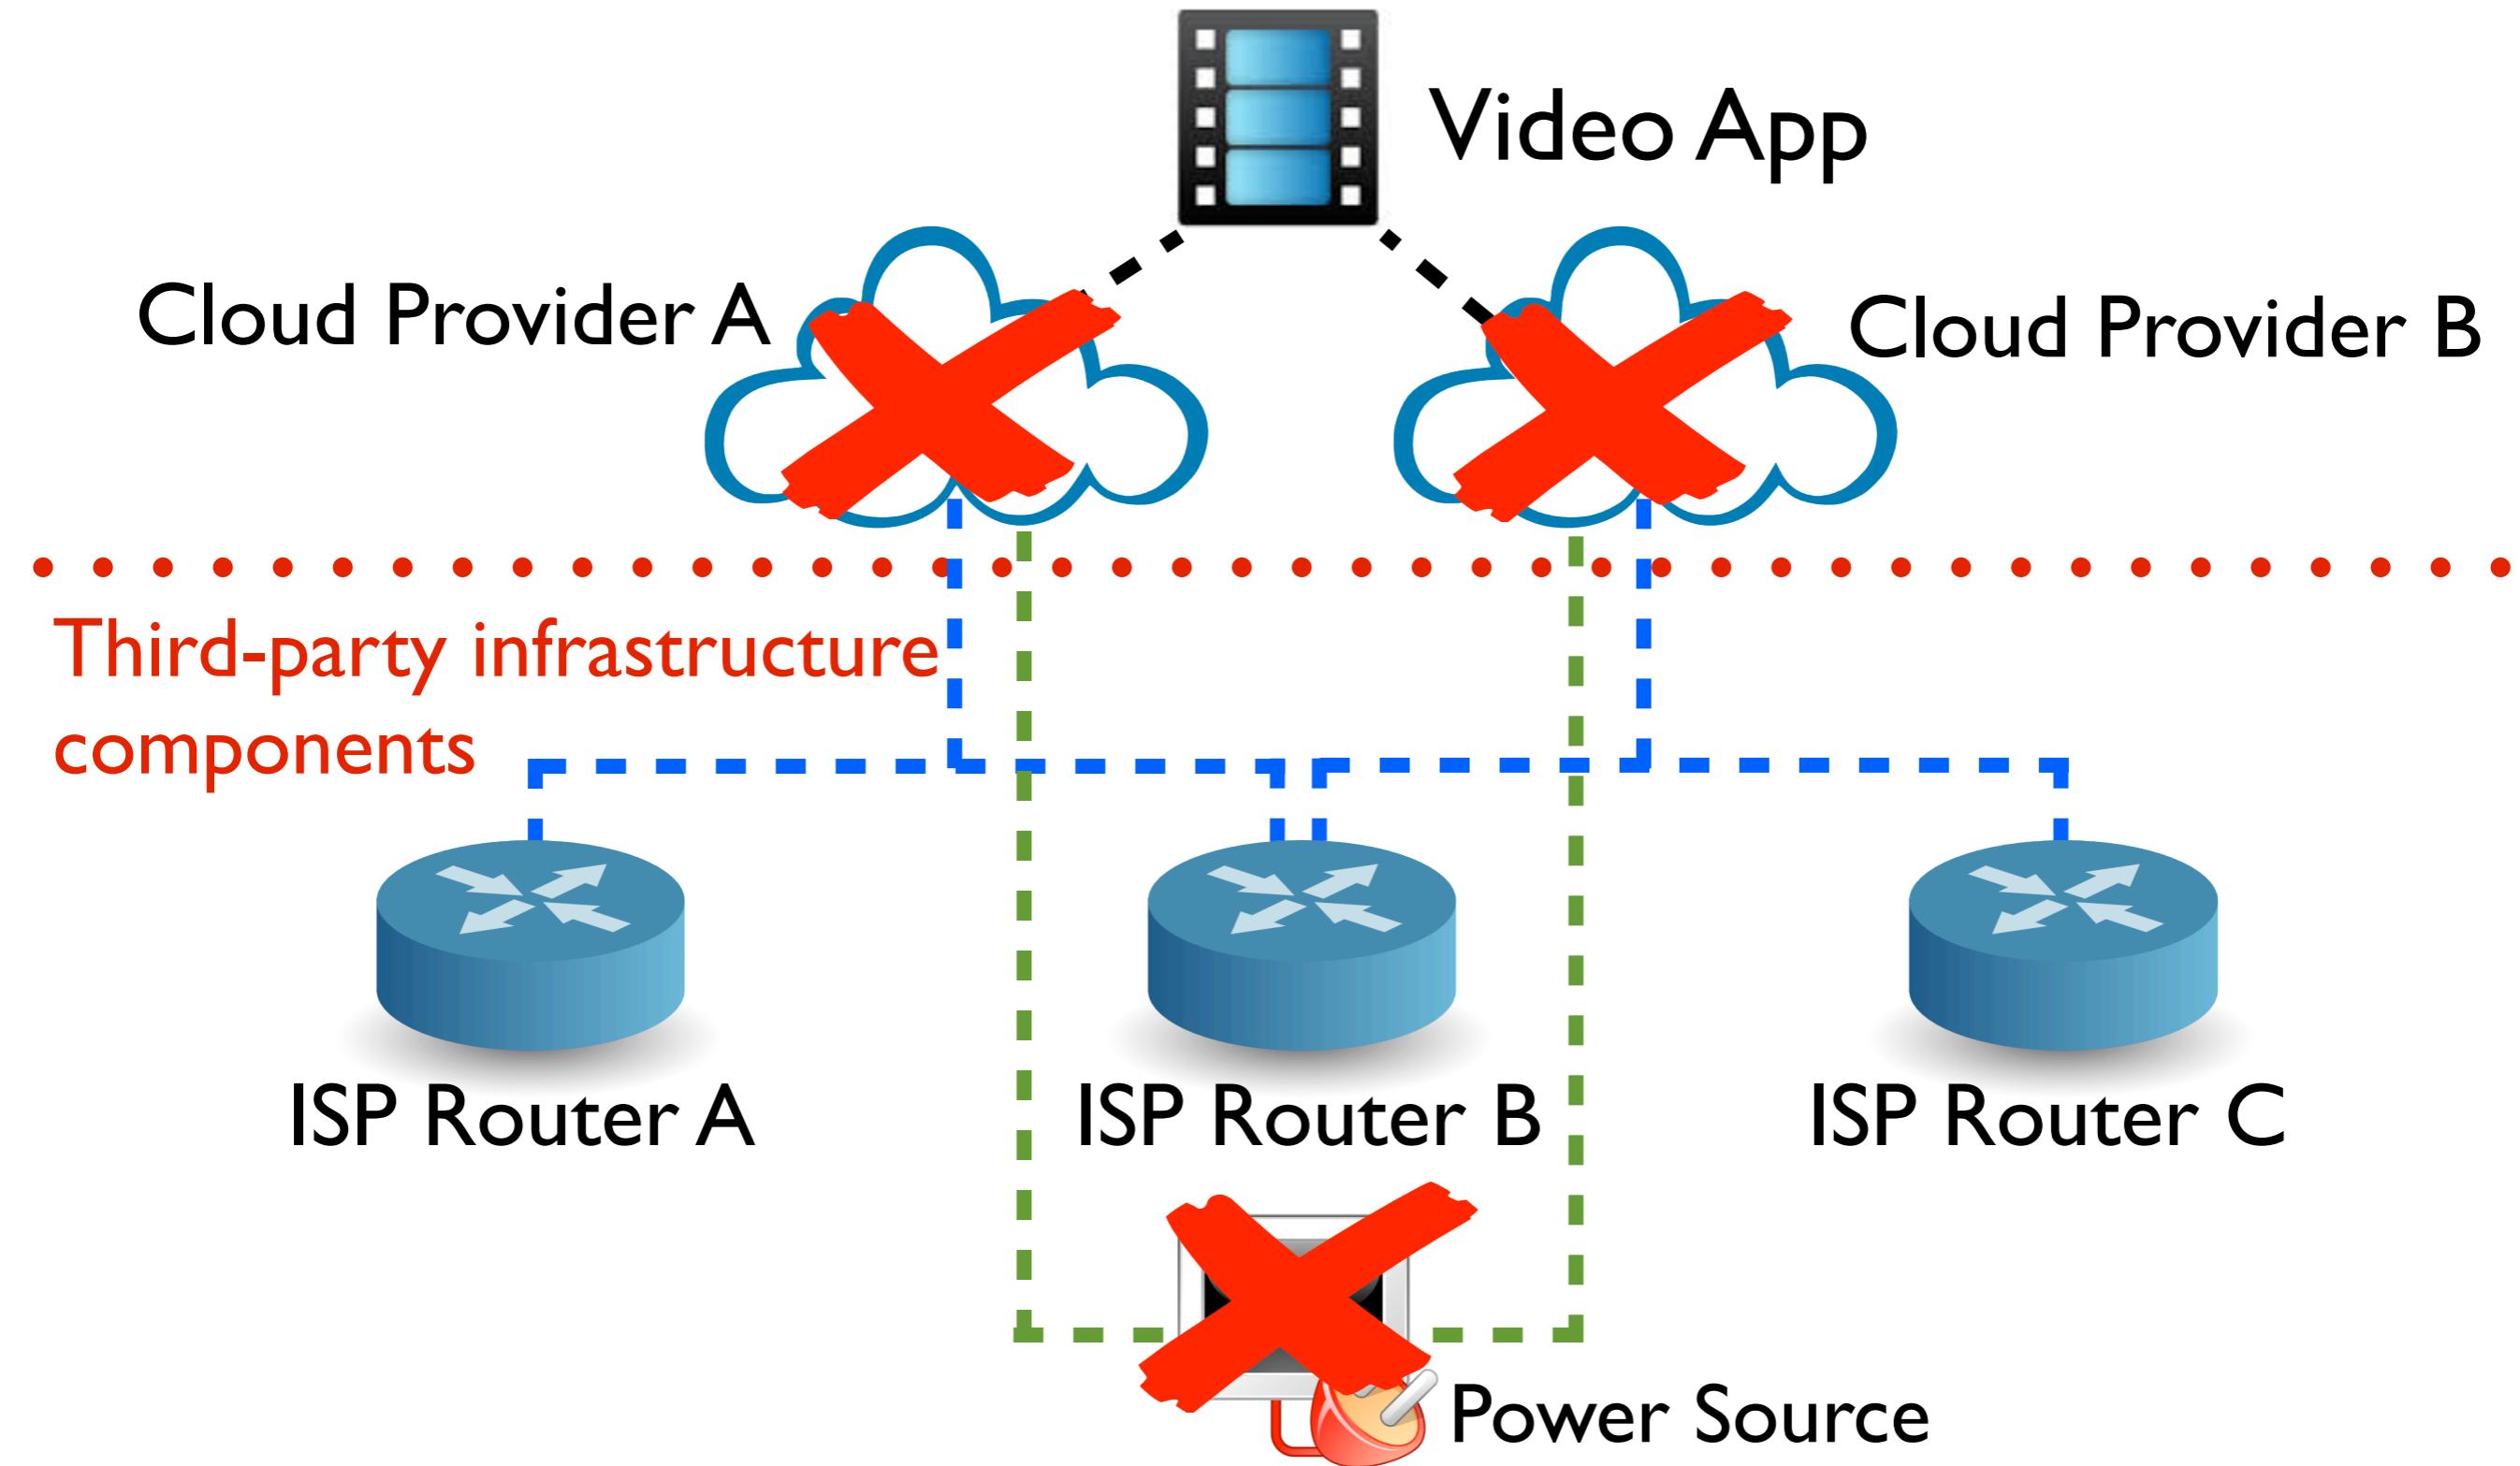


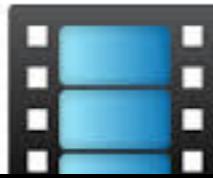
Third-party infrastructure  
components











Video App

Cloud  
Provider A

Cloud  
Provider B

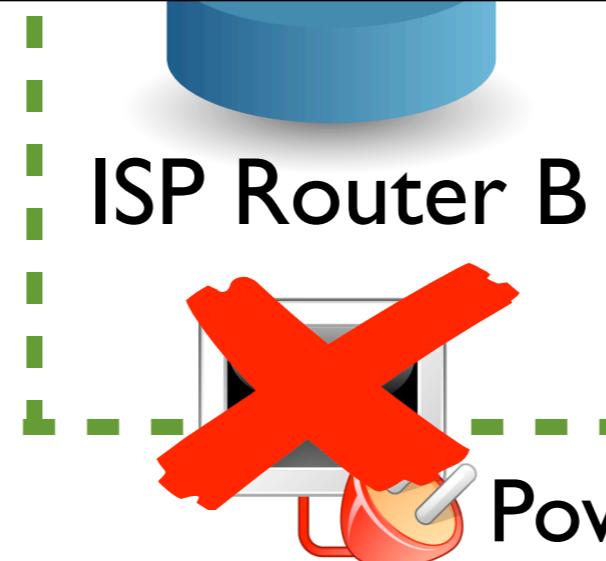
Cloud providers do not usually share  
information about their dependencies

Third-party  
components

ISP Router A

ISP Router B

ISP Router C



# Existing Efforts

- Cloud providers allocate or tolerate failures via:
  - diagnosis systems;
  - fault-tolerant systems.

# Existing Efforts

- Cloud providers allocate or tolerate failures via:
  - diagnosis systems;
  - fault-tolerant systems.
- Solving the problem after outage occurs.

# Existing Efforts

- Cloud providers allocate or tolerate failures via:
  - diagnosis systems;
  - fault-tolerant systems.
- Solving the problem after outage occurs.
- Prevent correlated failures before outage occurs.

# Existing Efforts

- Cloud providers allocate or tolerate failures via:
  - diagnosis systems;
  - fault-tolerant systems.
- Solving the problem after outage occurs.
- Prevent correlated failures before outage occurs.

Independence-as-a-Service  
(INDaaS)

# INDaaS Workflow

Service Provider, Alice



INDaaS

A Given Redundancy Configuration

# INDaaS Workflow

Service Provider, Alice



INDaaS

Two-Way Redundancy Configuration

Dependency  
Data Source1

Dependency  
Data Source2

# INDaaS Workflow

Service Provider, Alice



Independence of this  
two-way redundancy ?

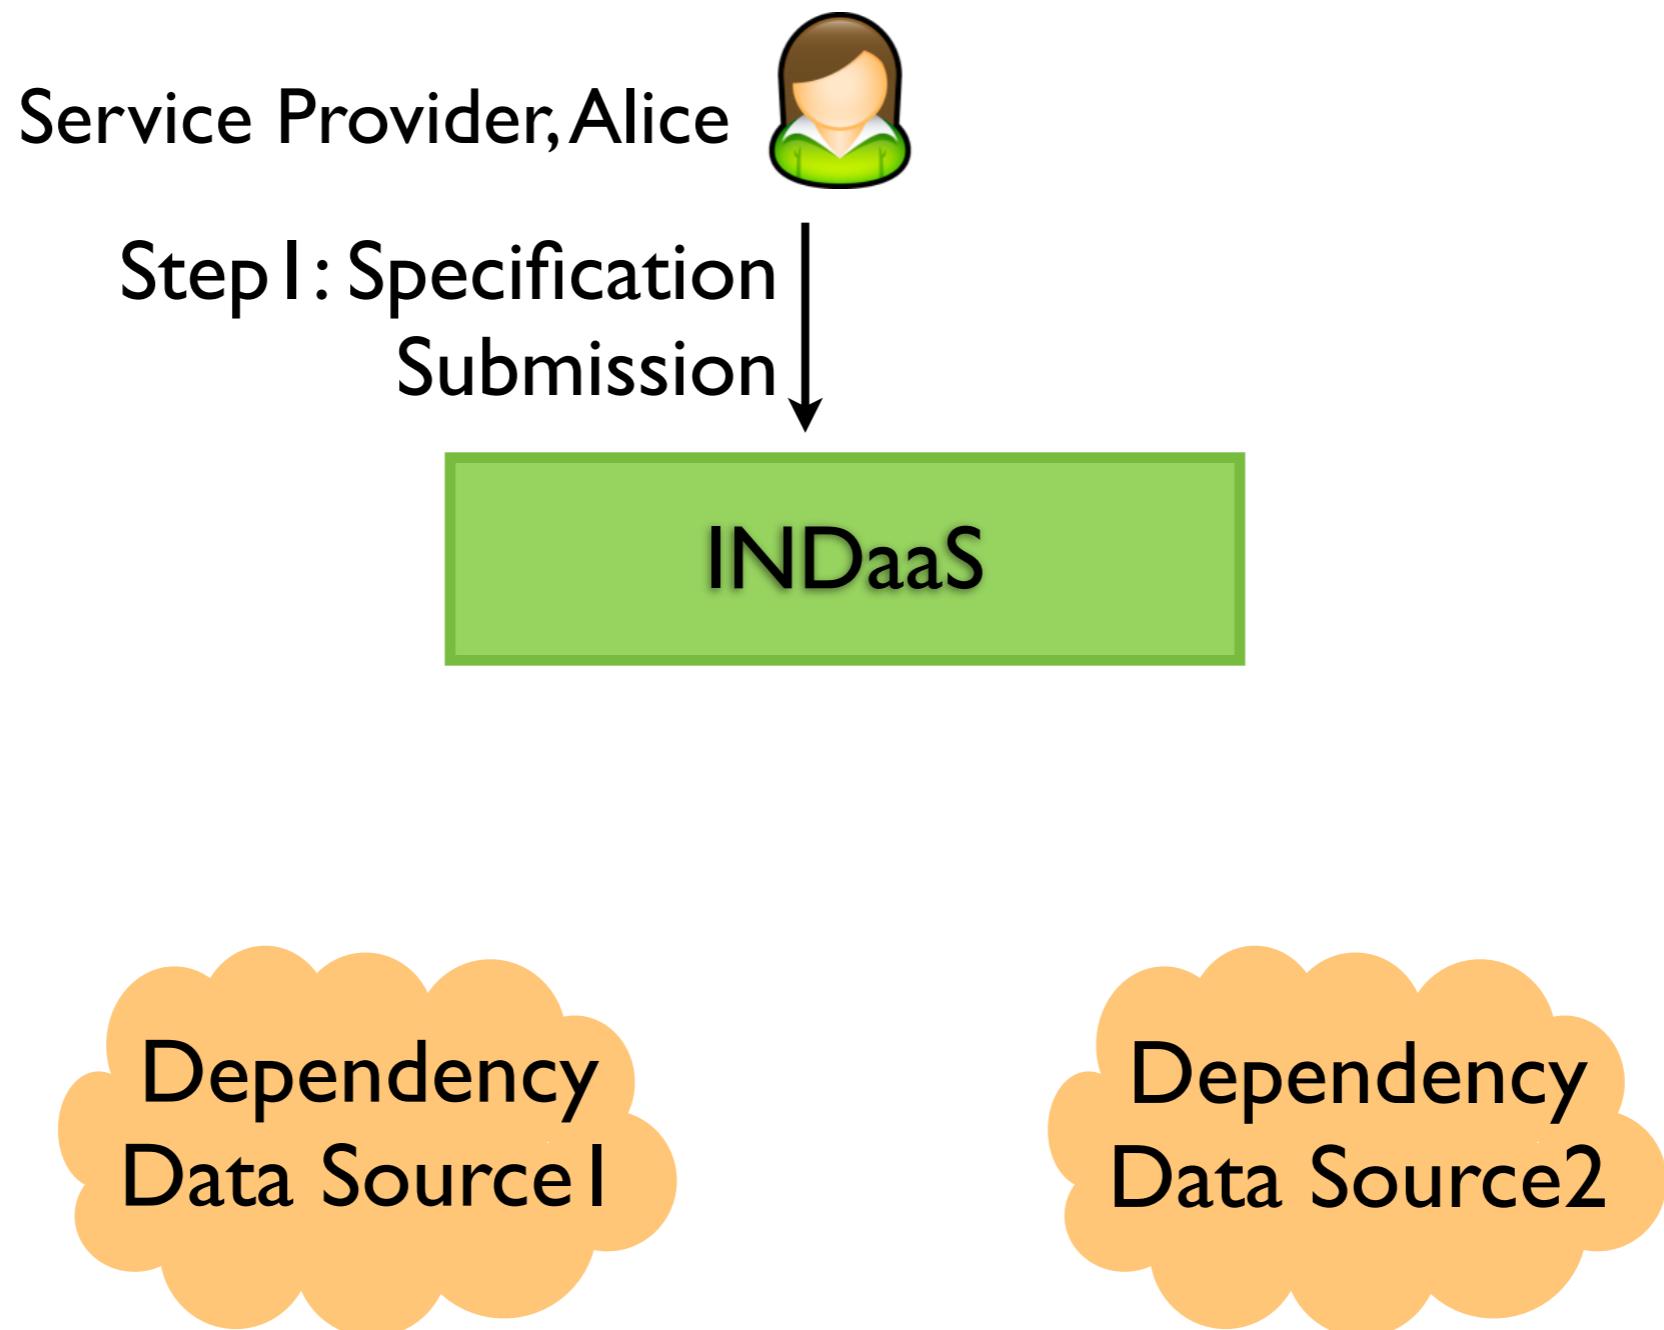
INDaaS

Two-Way Redundancy Configuration

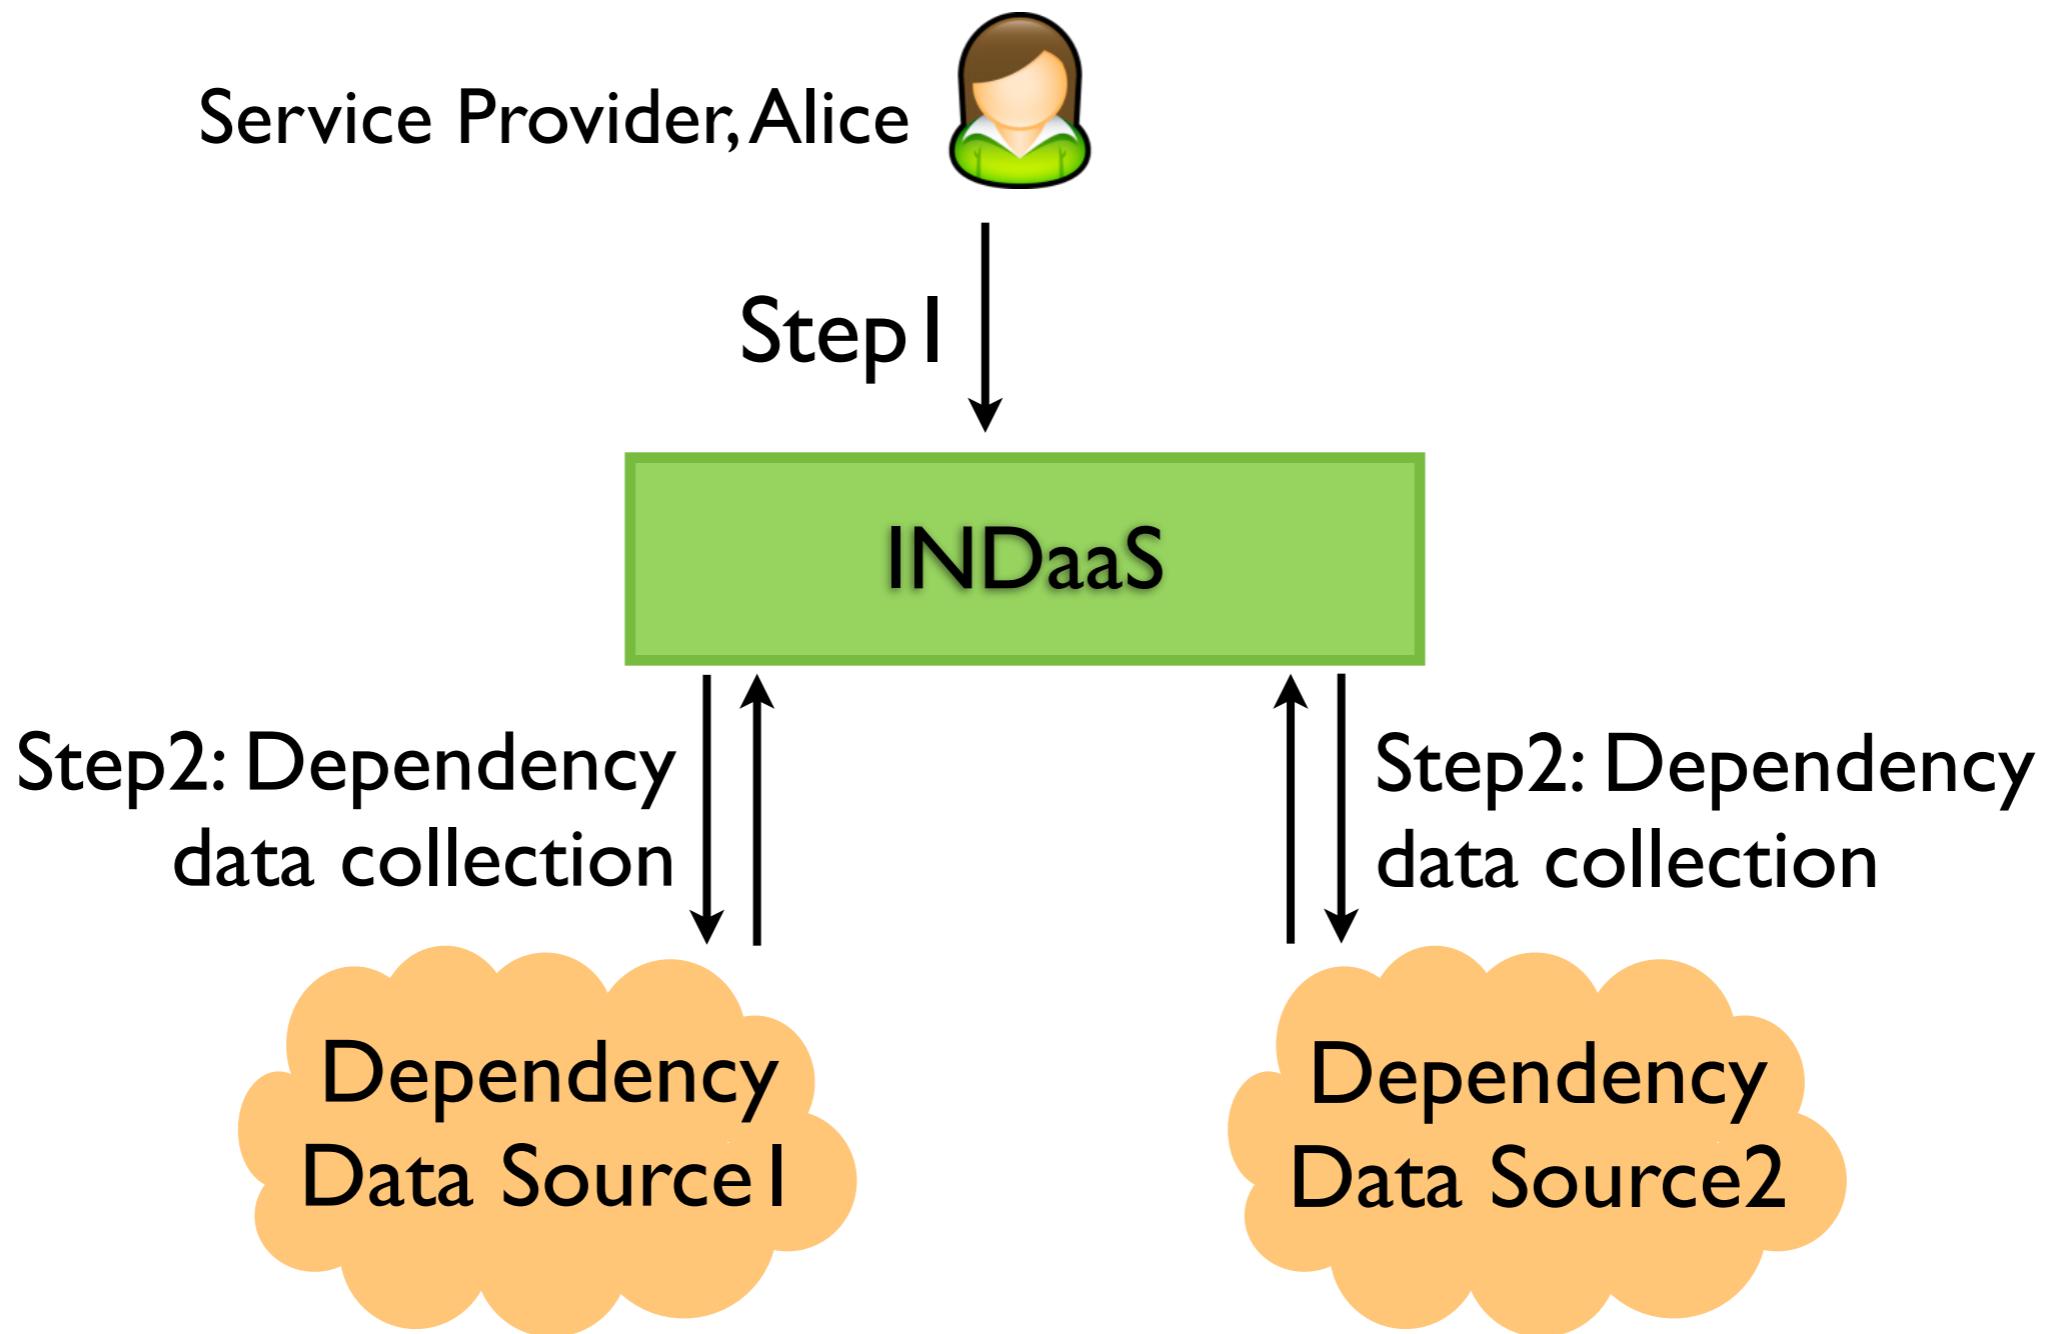
Dependency  
Data Source1

Dependency  
Data Source2

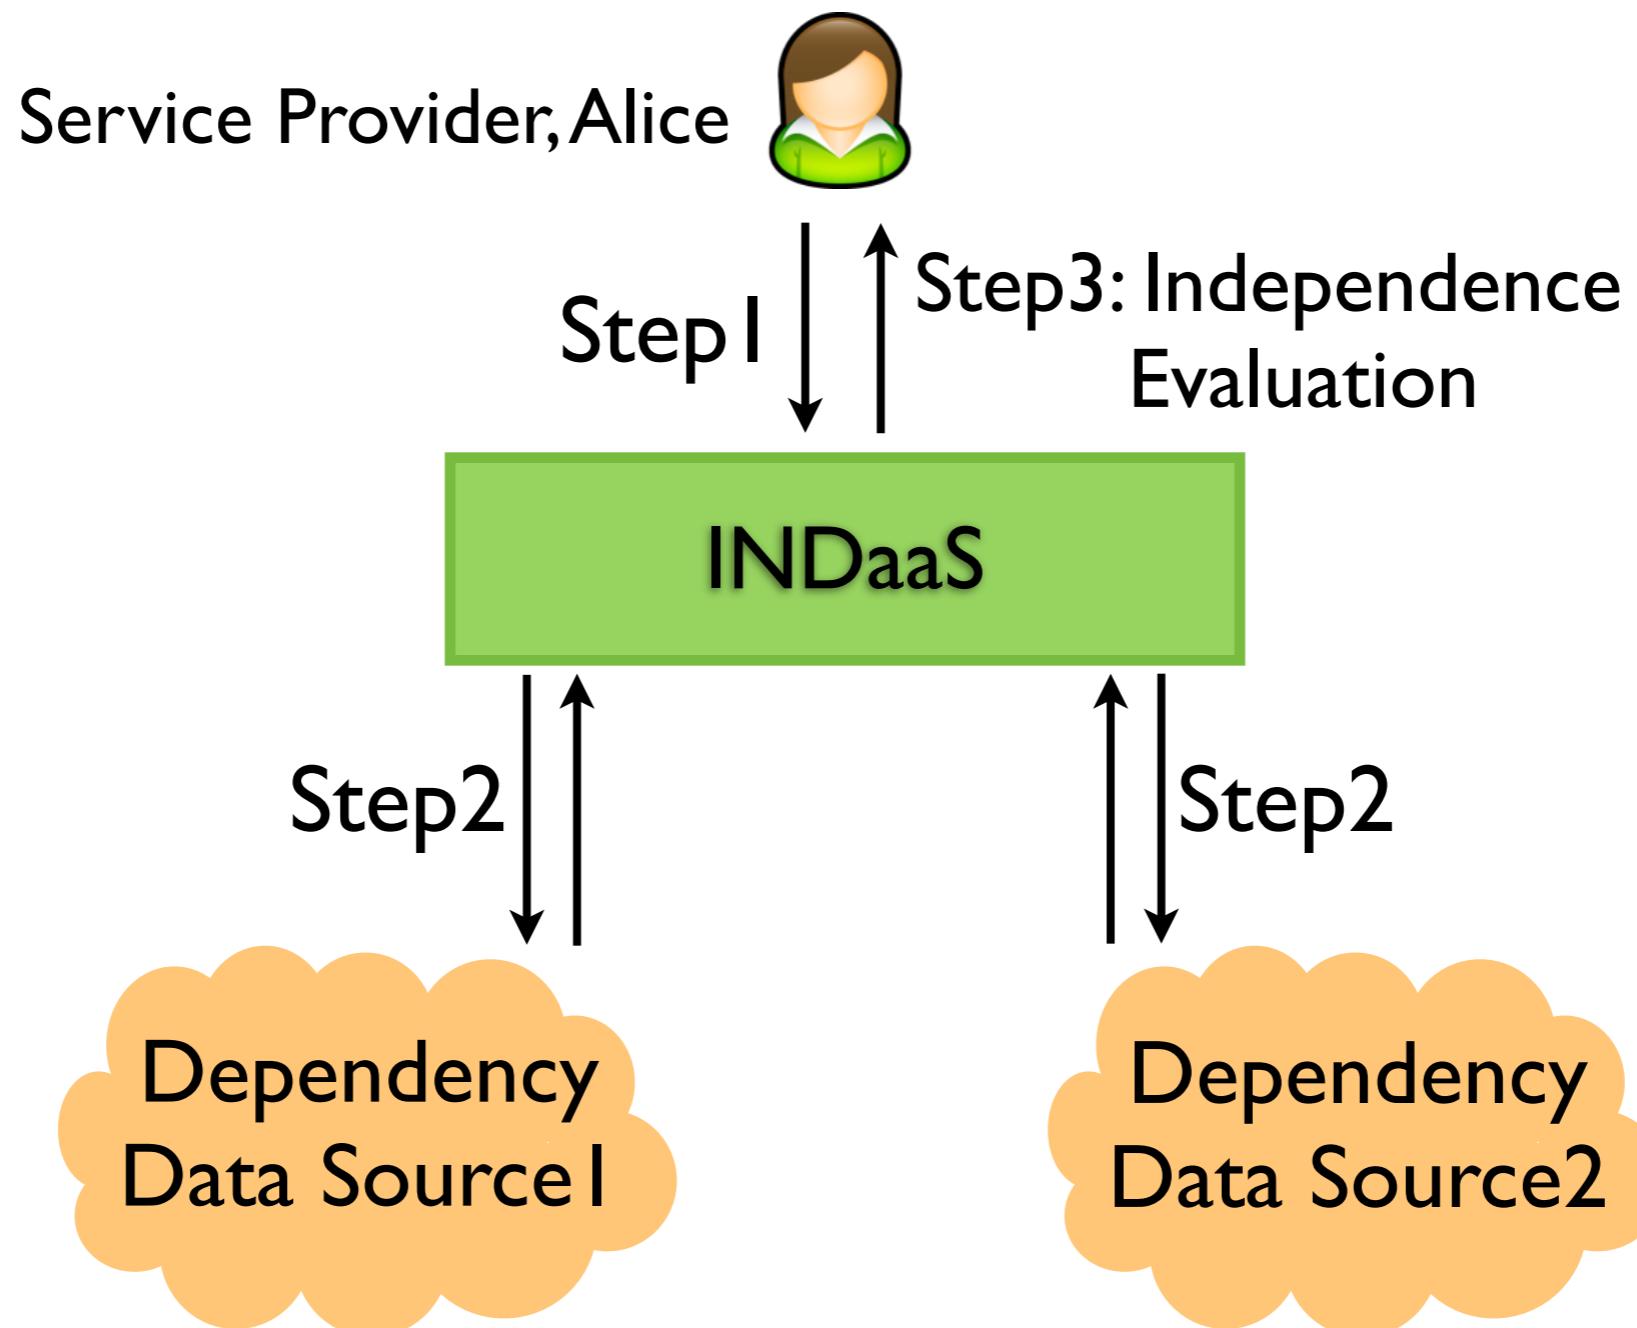
# INDaaS Workflow



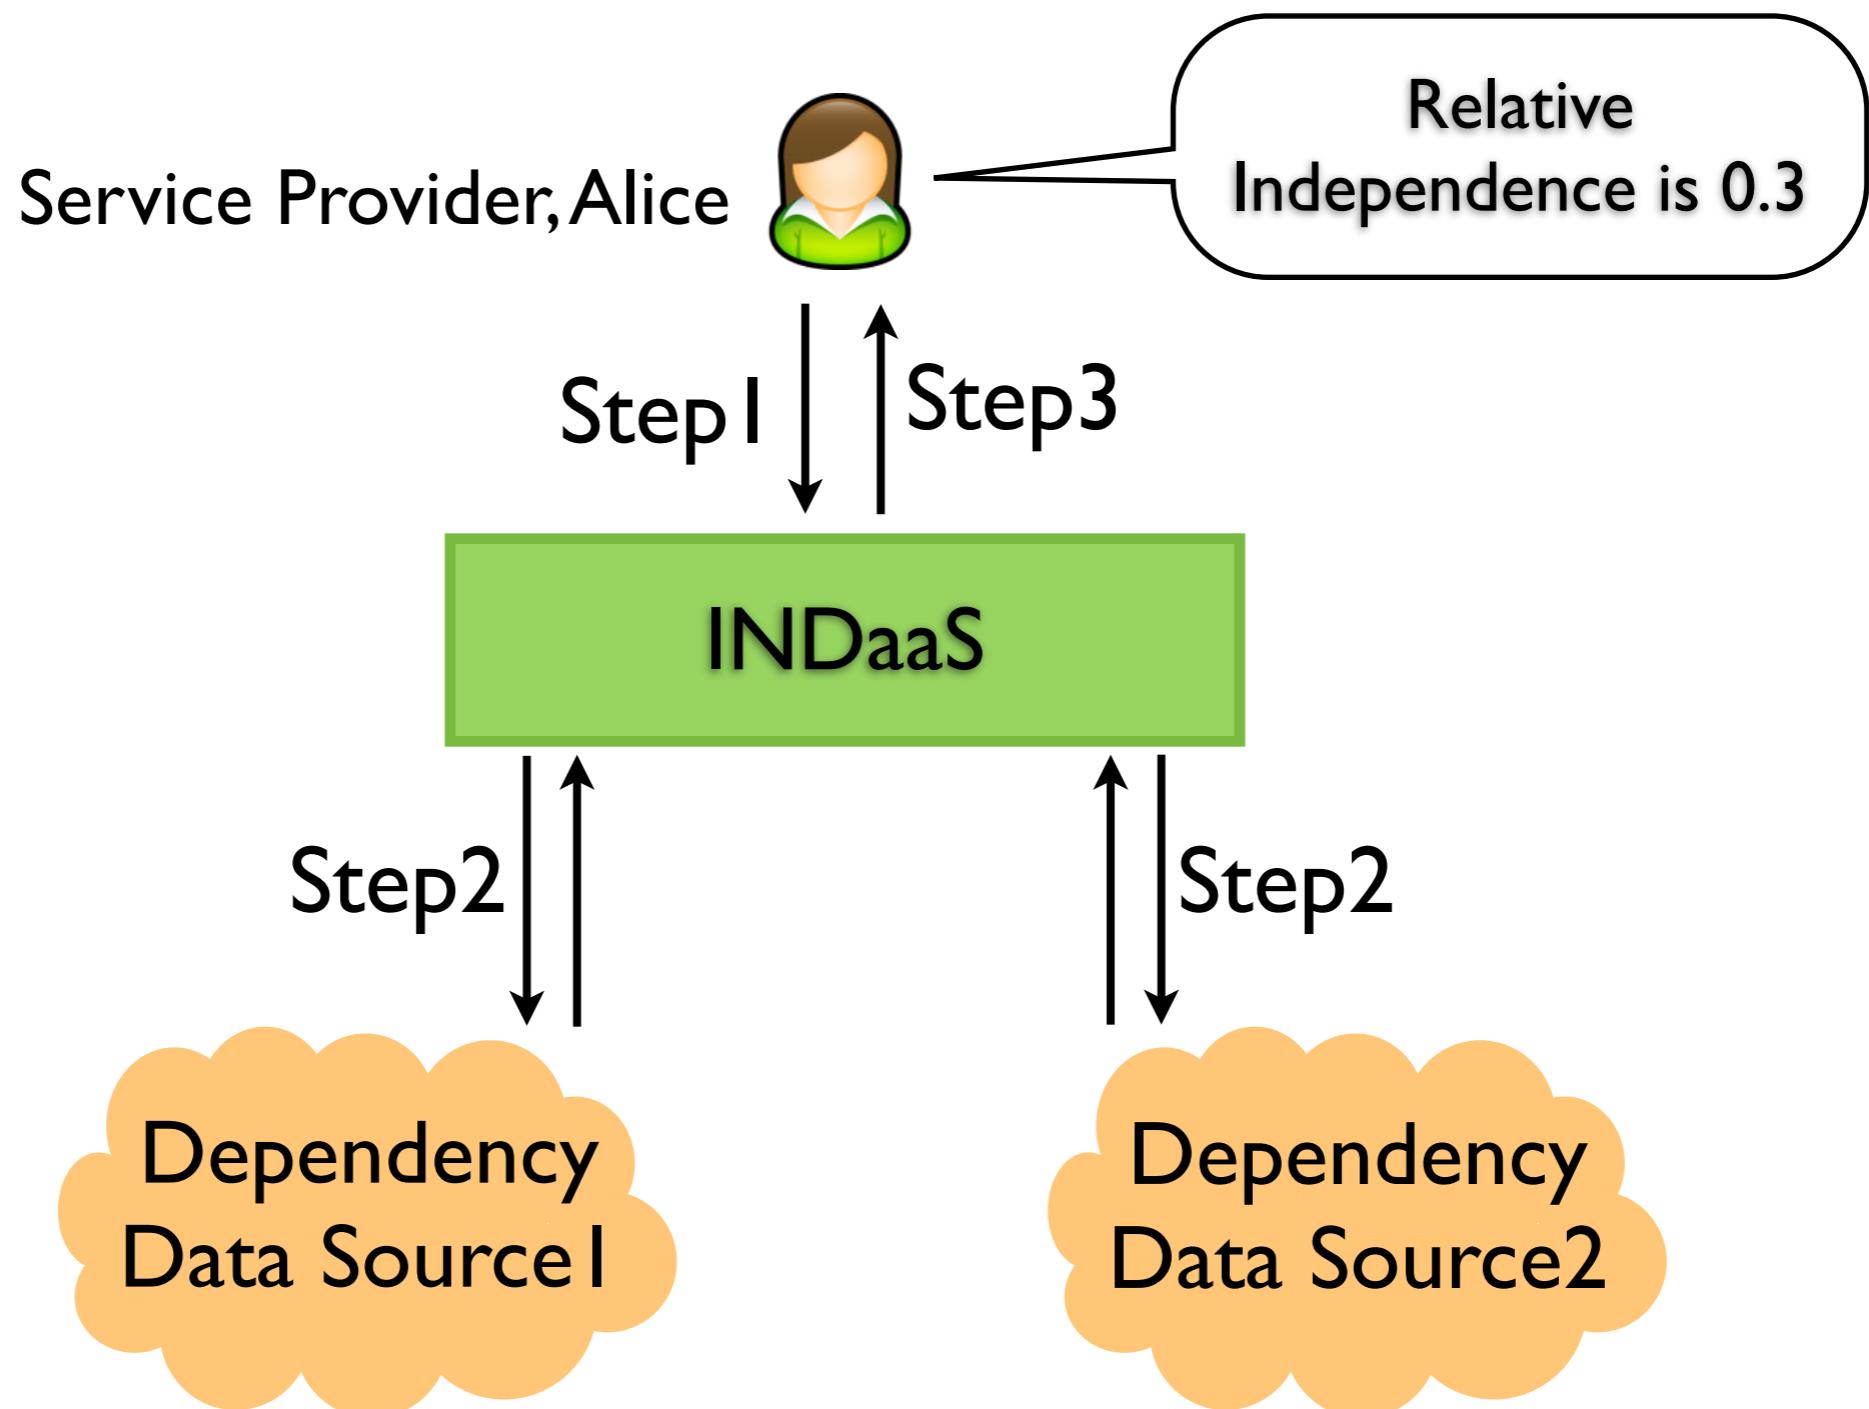
# INDaaS Workflow



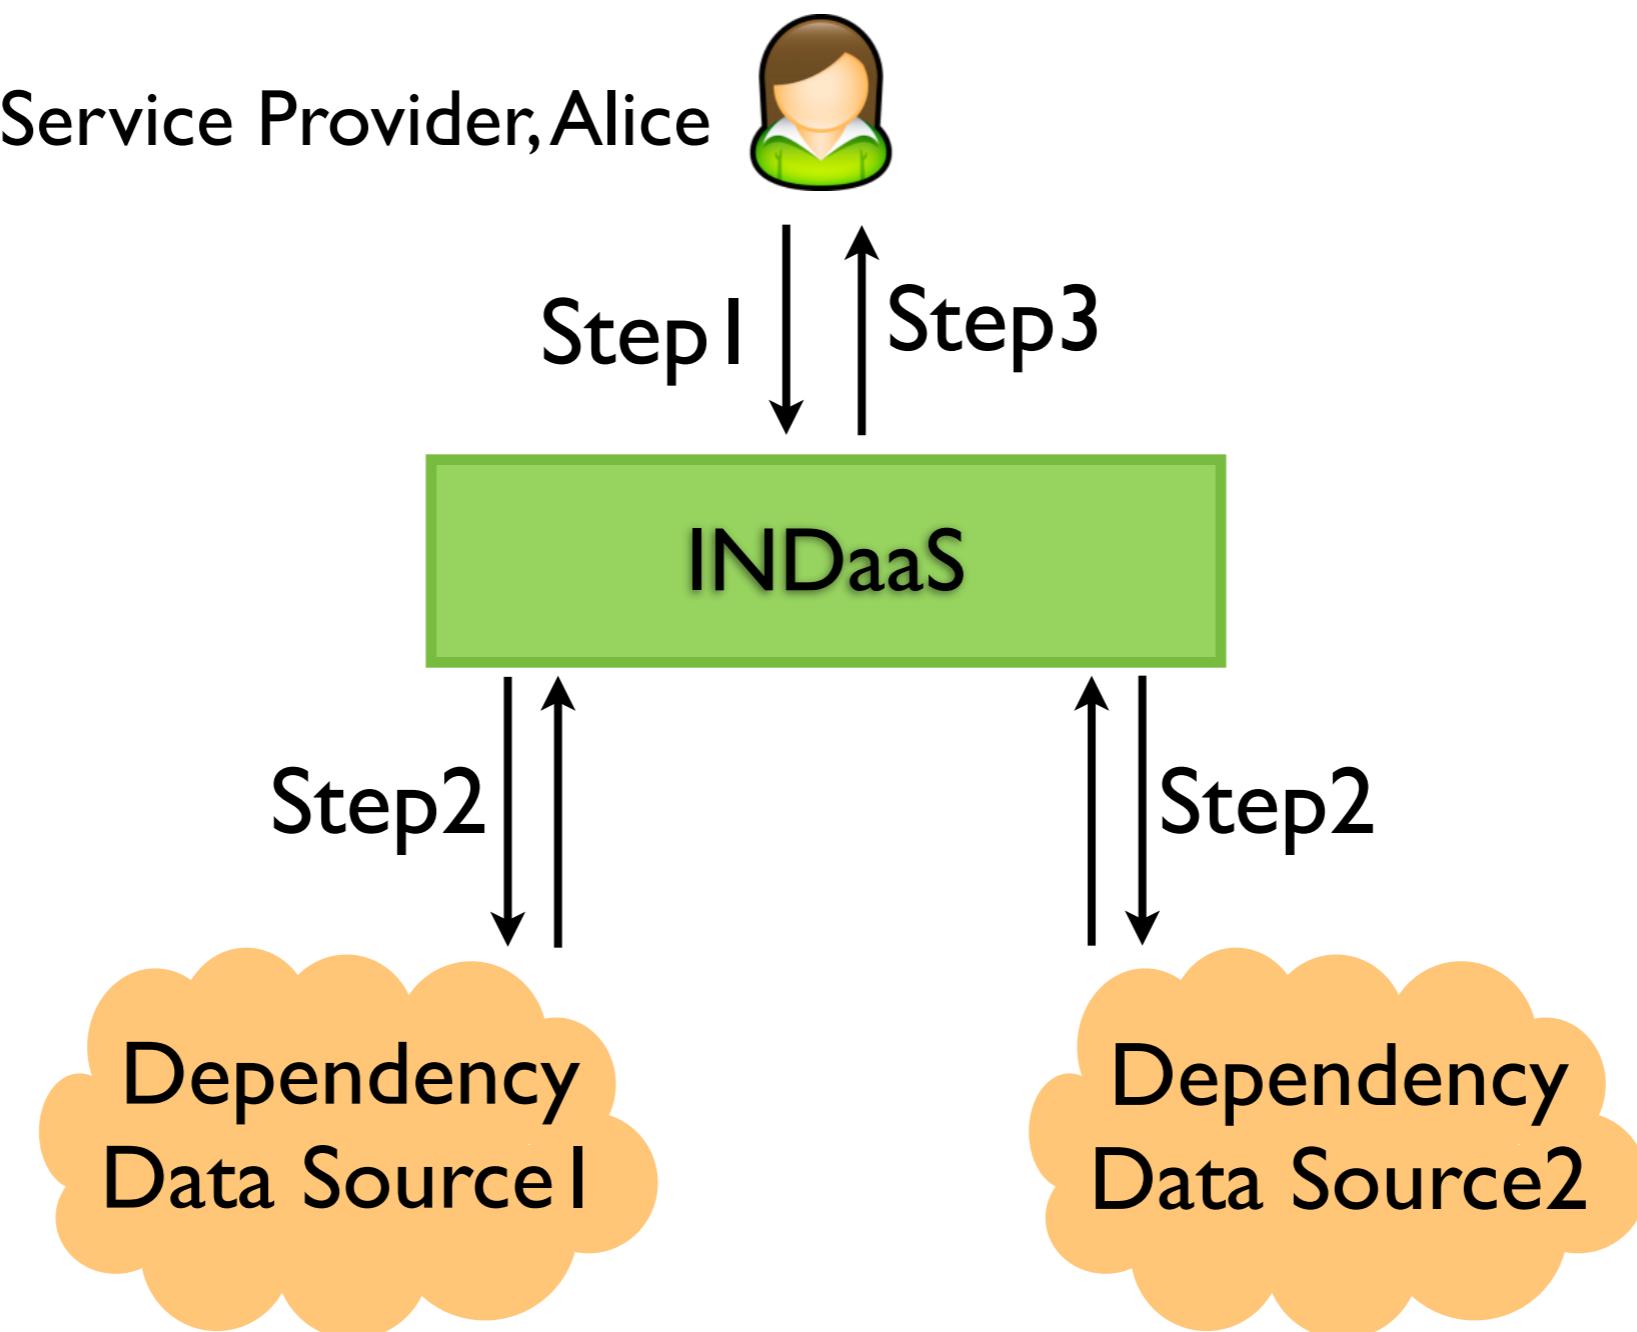
# INDaaS Workflow



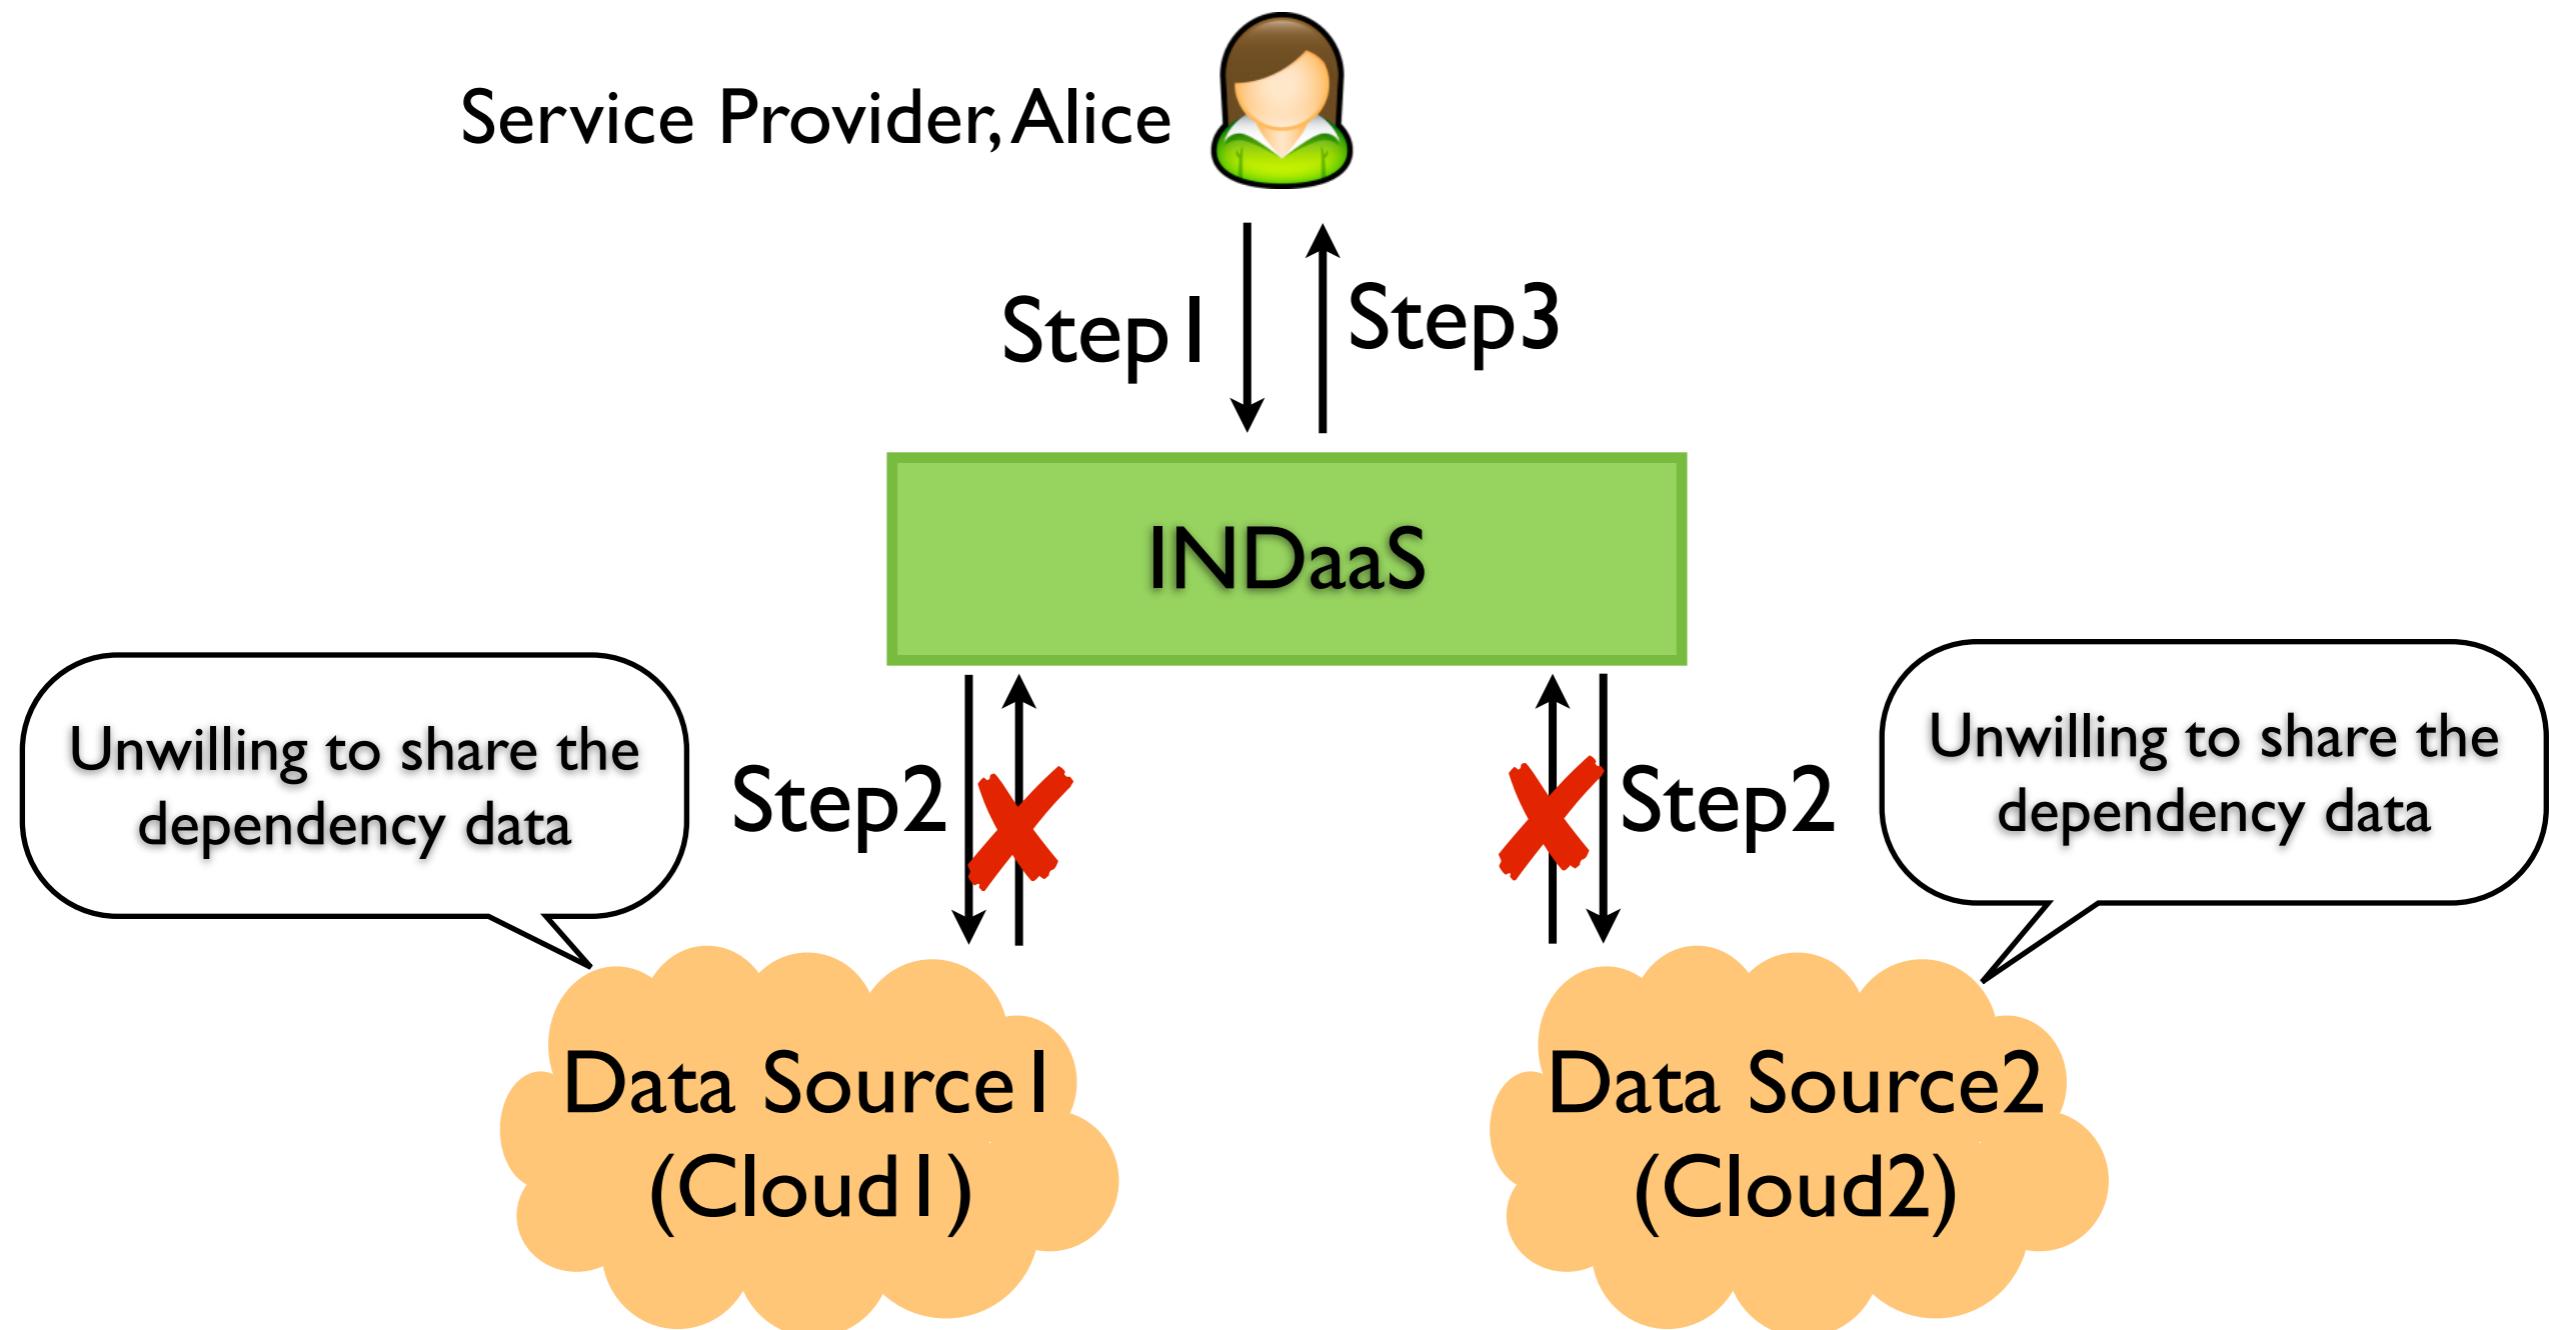
# INDaaS Workflow



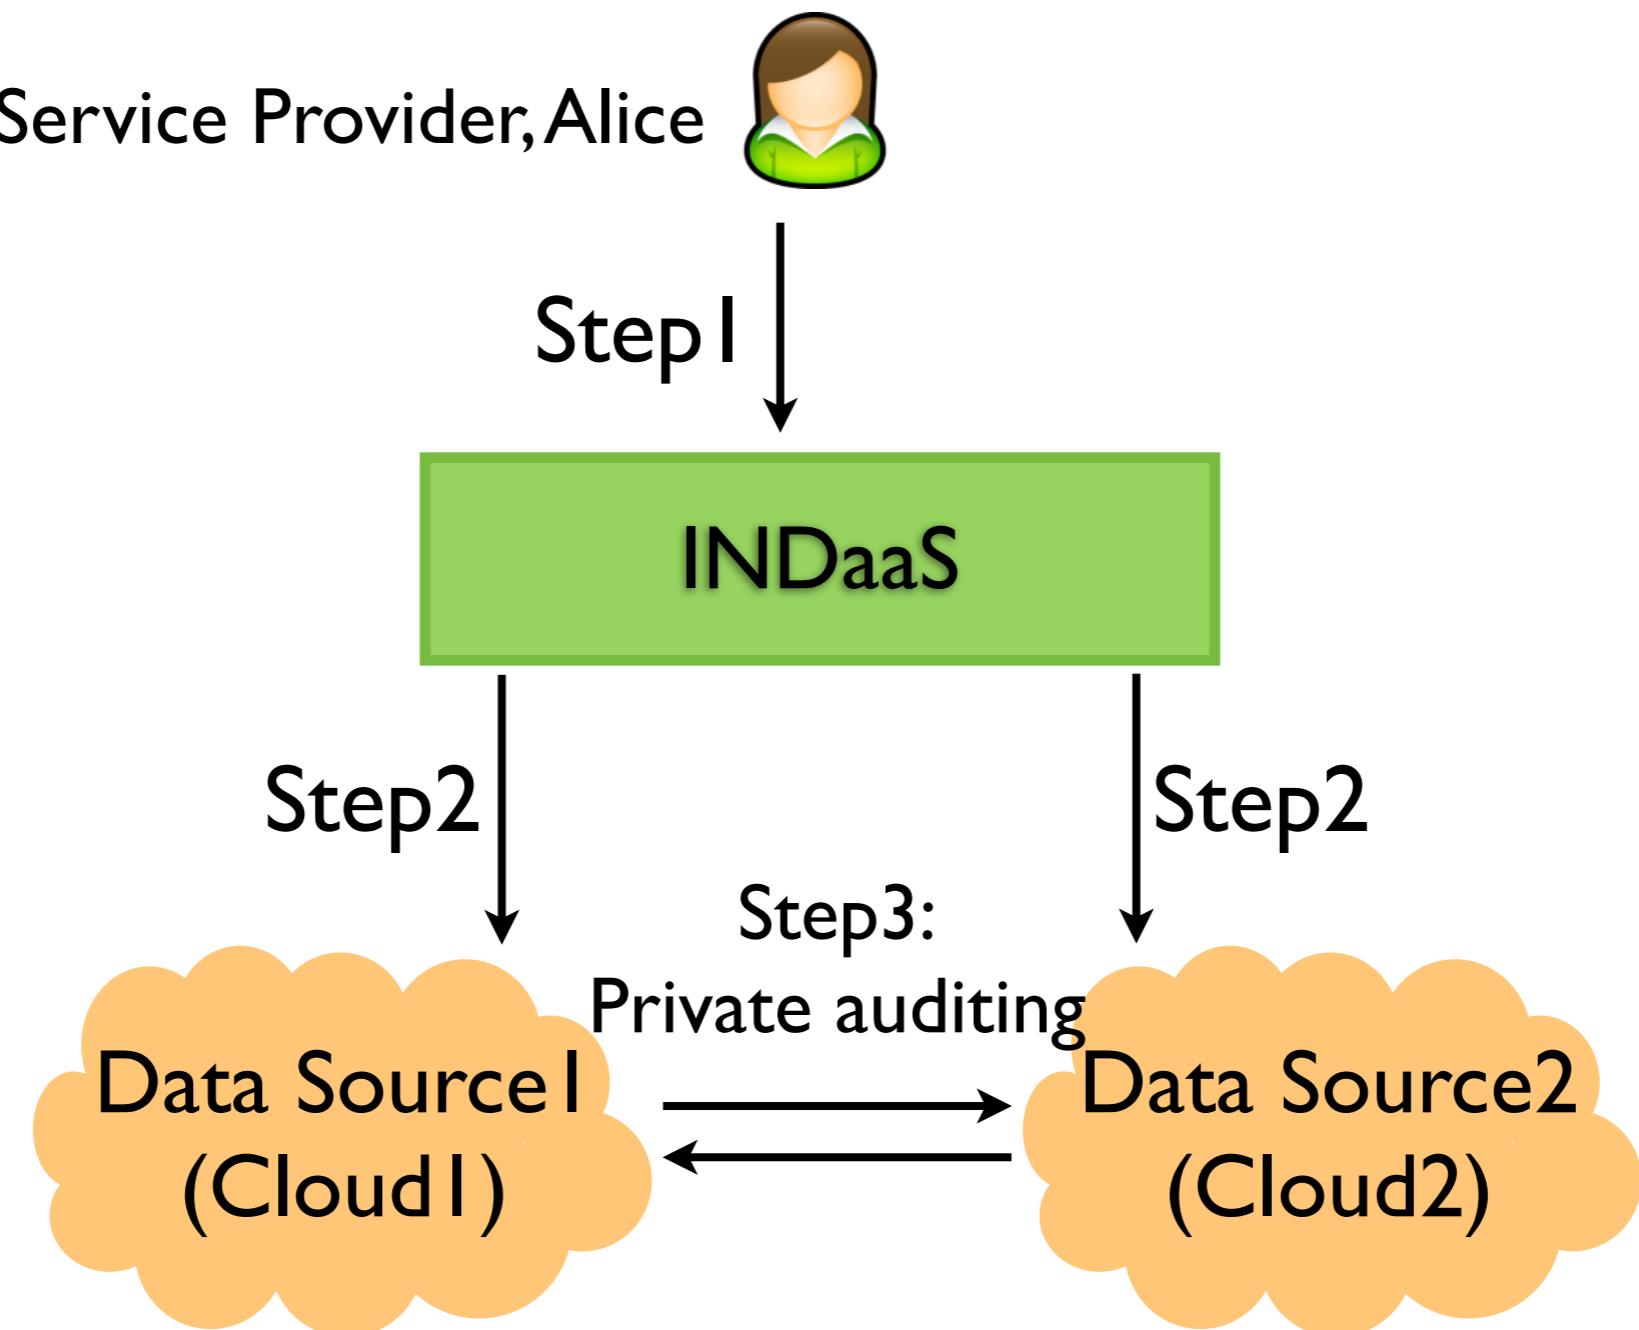
# Multiple Cloud Providers



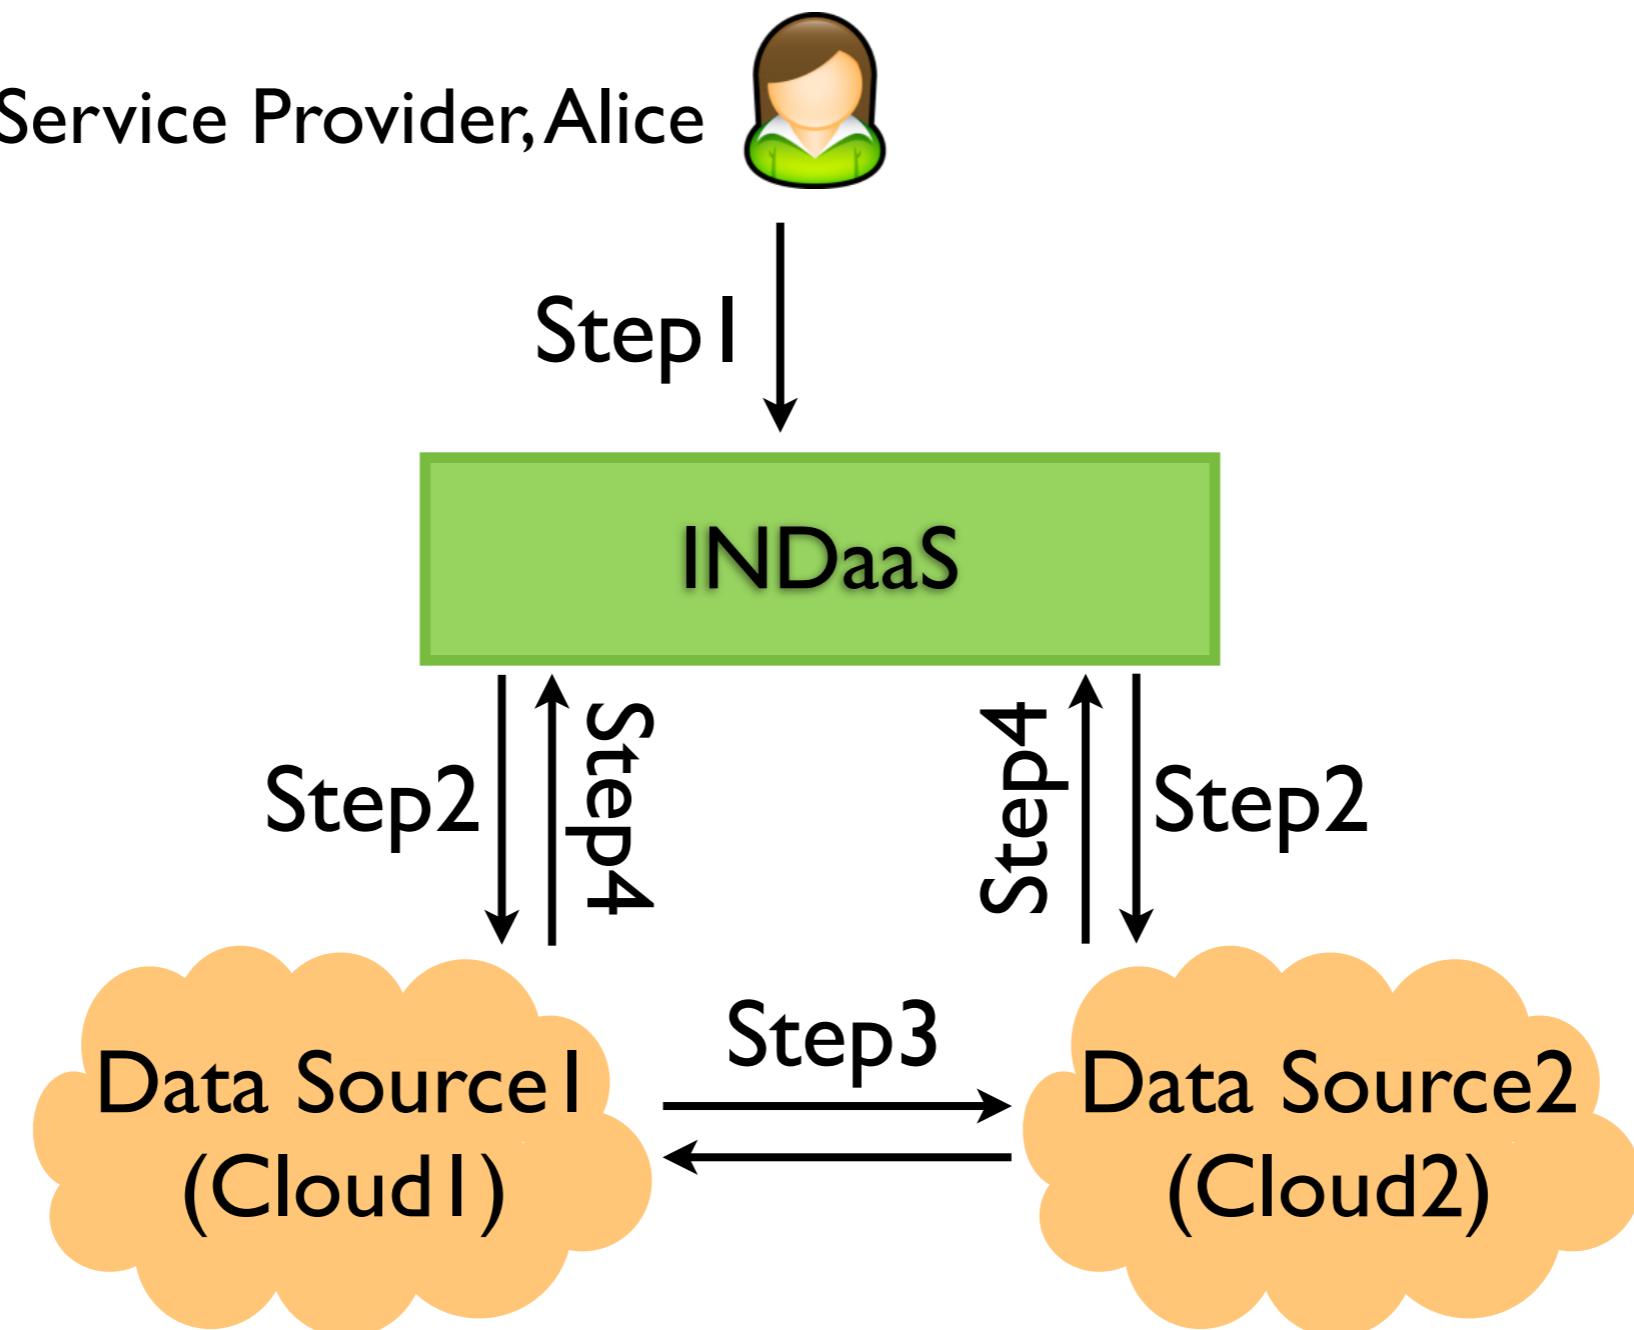
# Multiple Cloud Providers



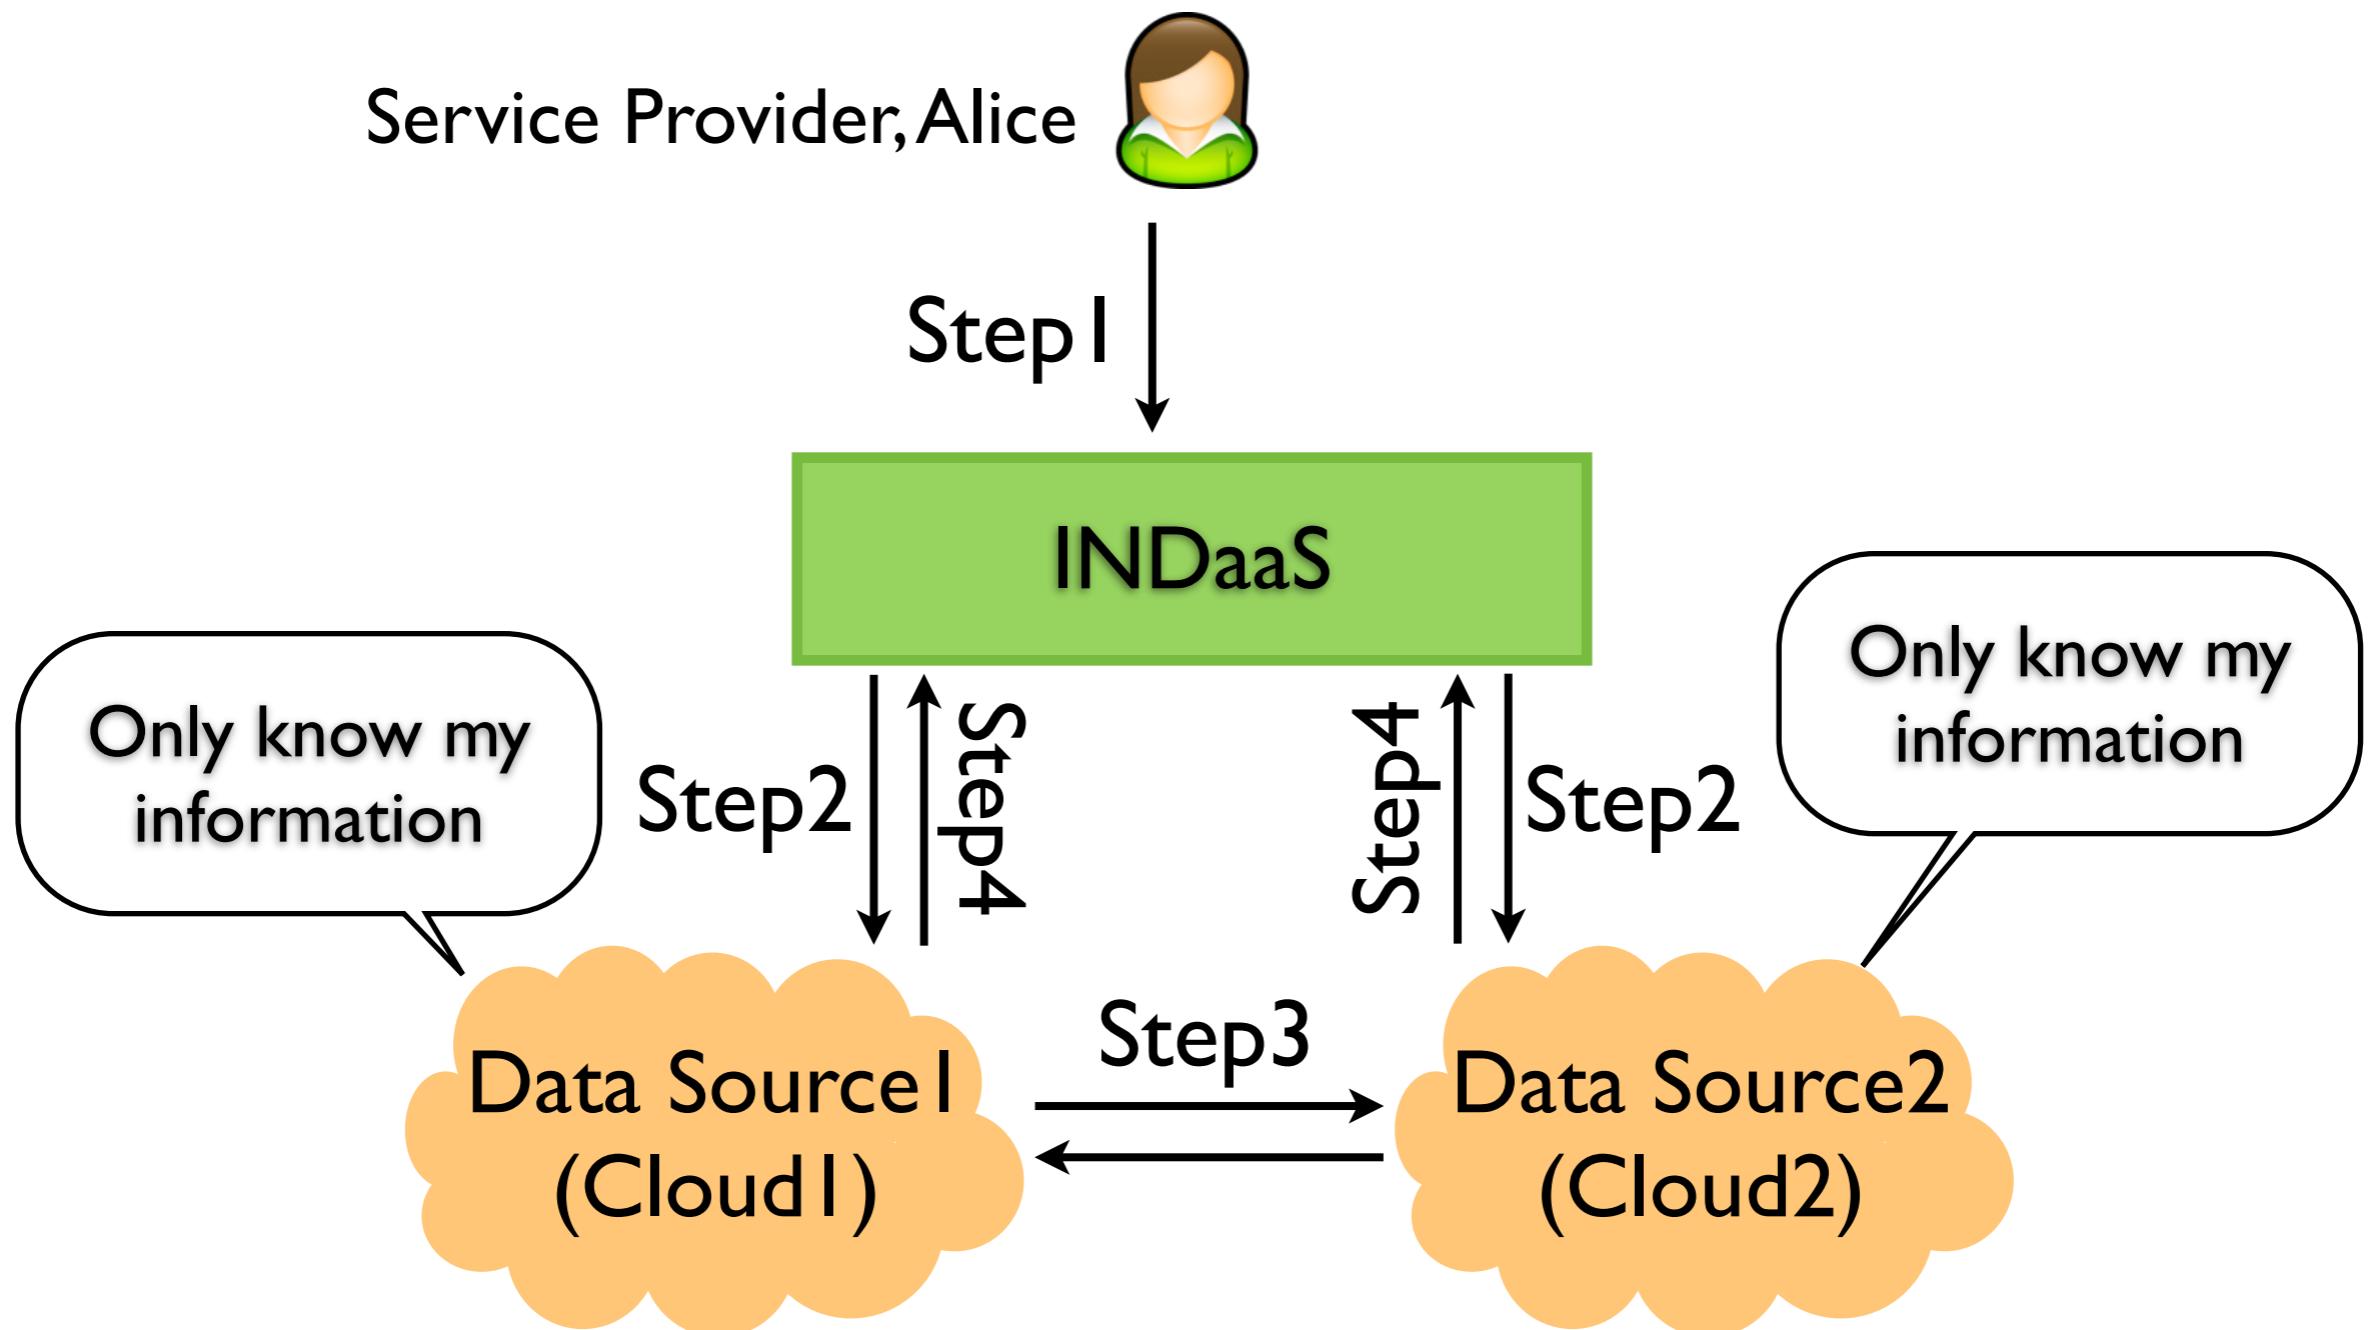
# Private Independence Evaluation



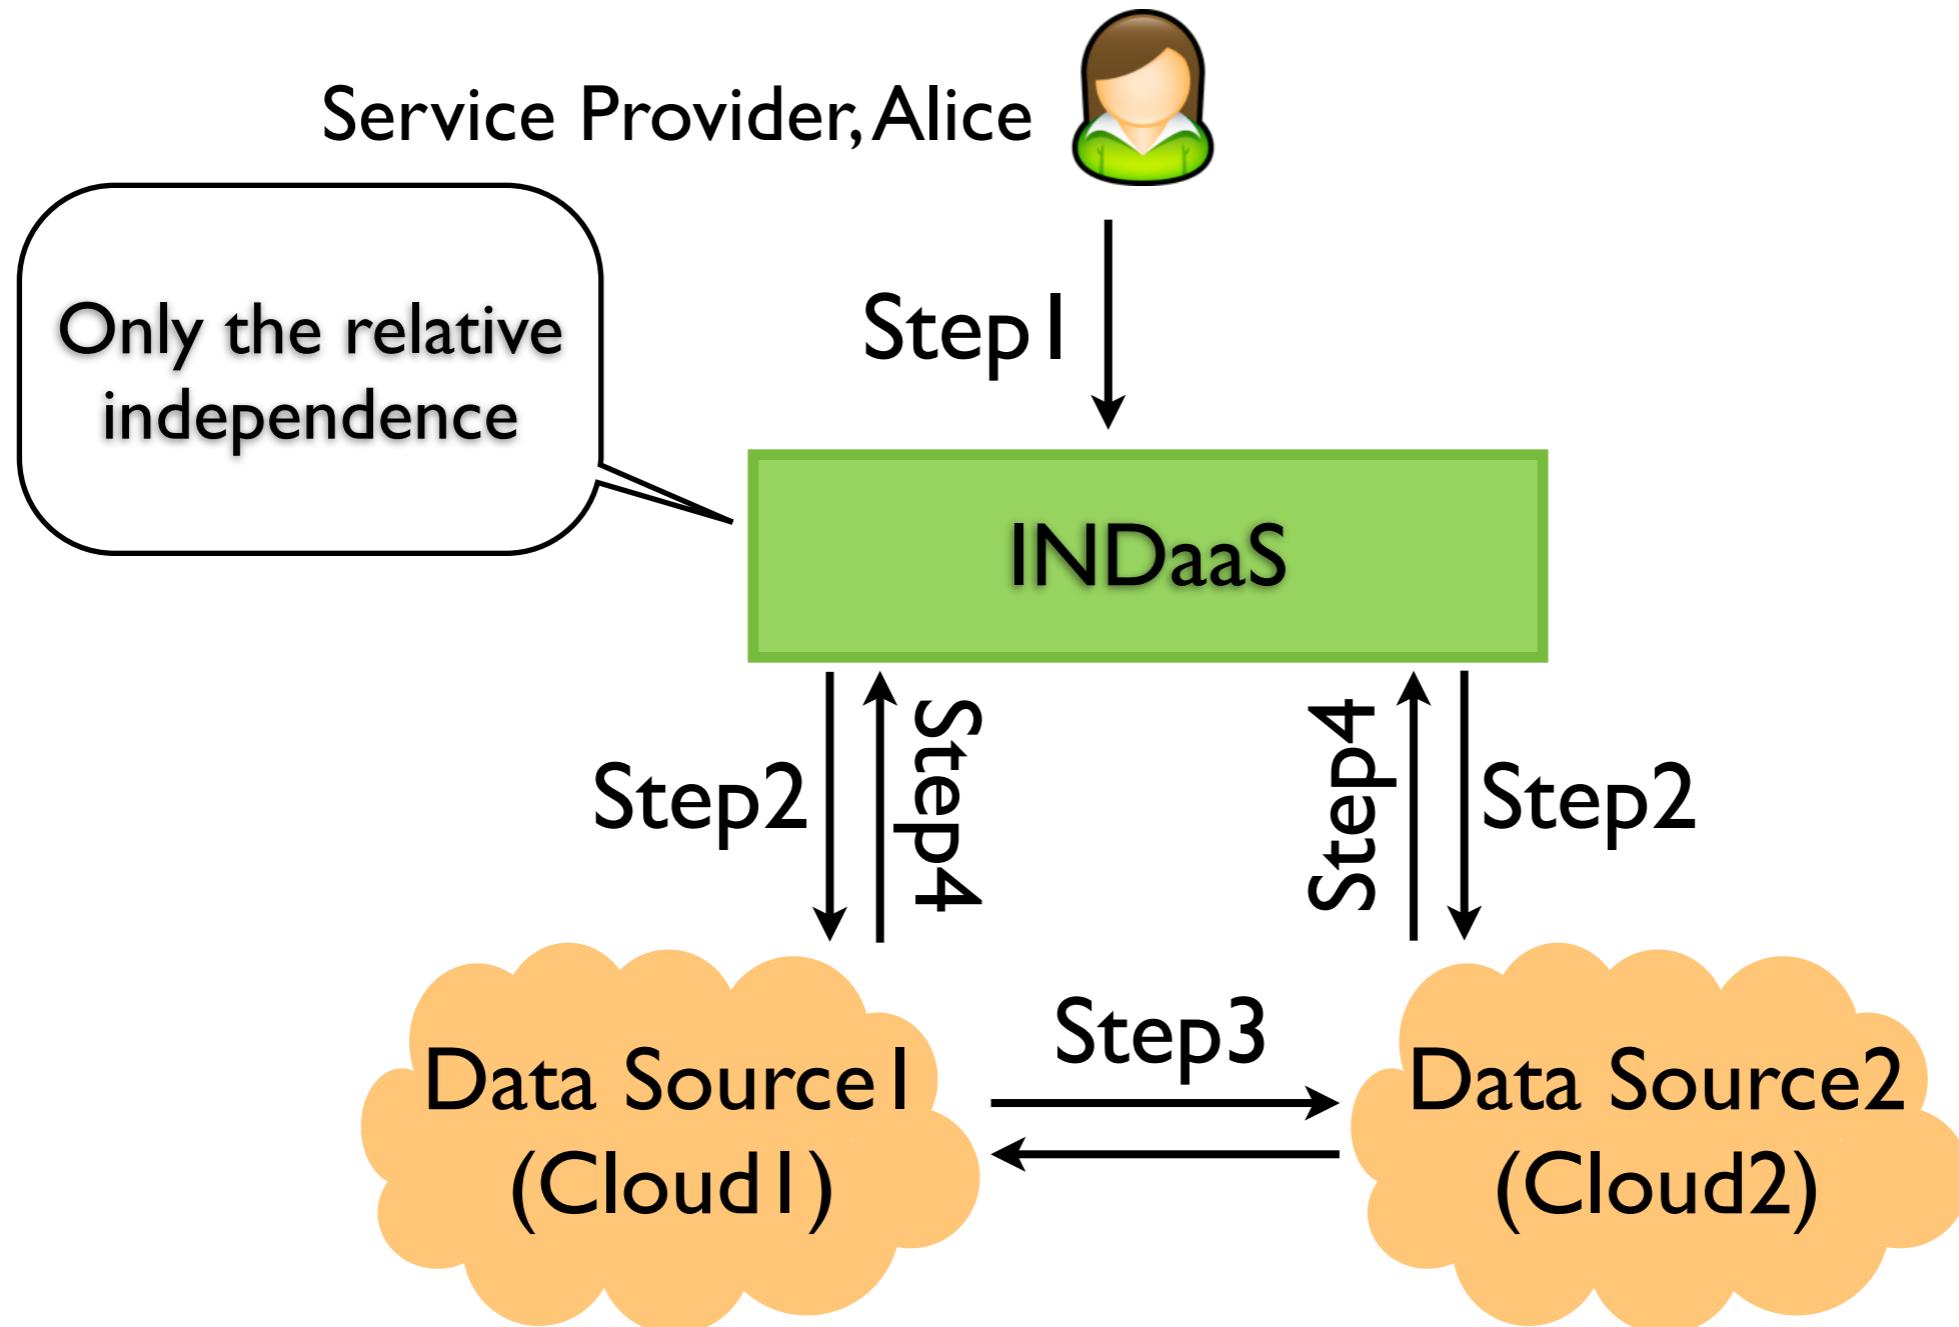
# Private Independence Evaluation



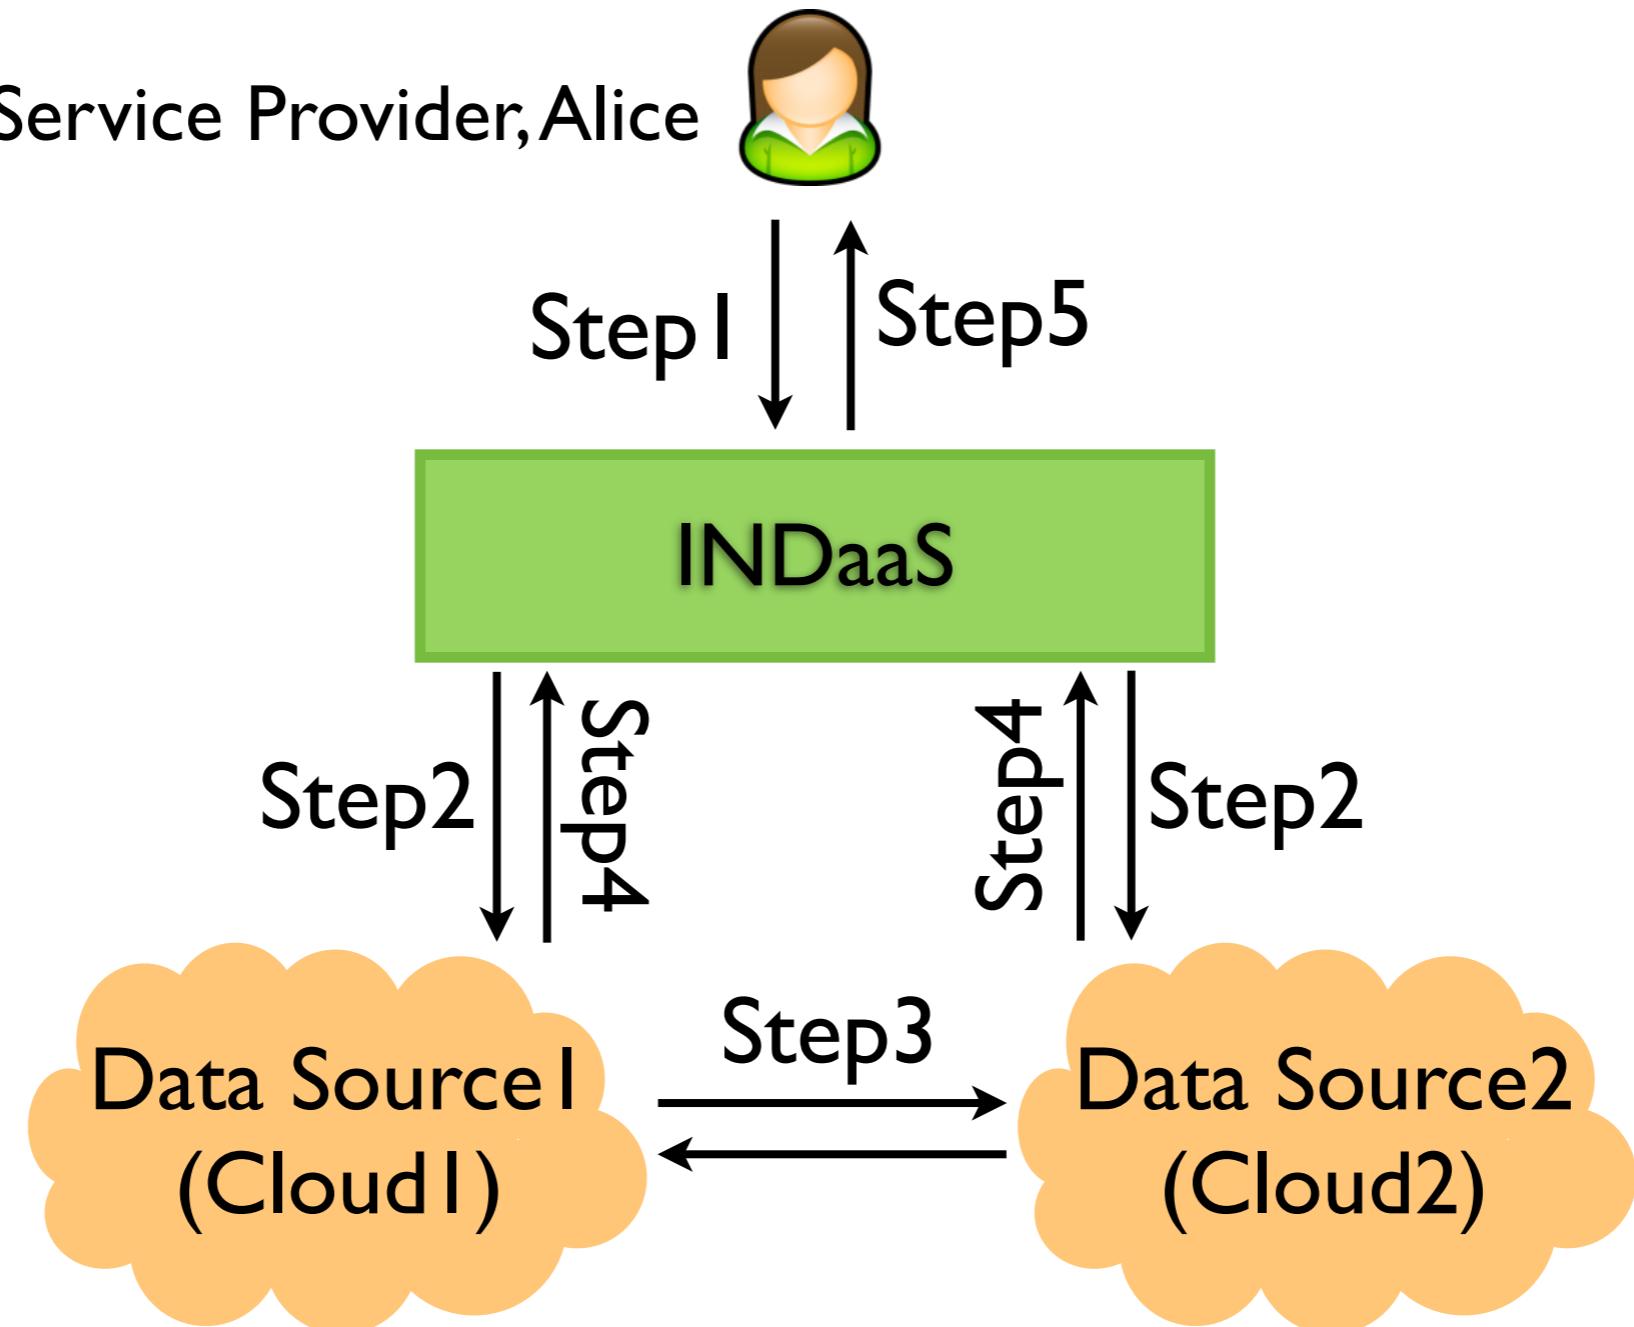
# Private Independence Evaluation



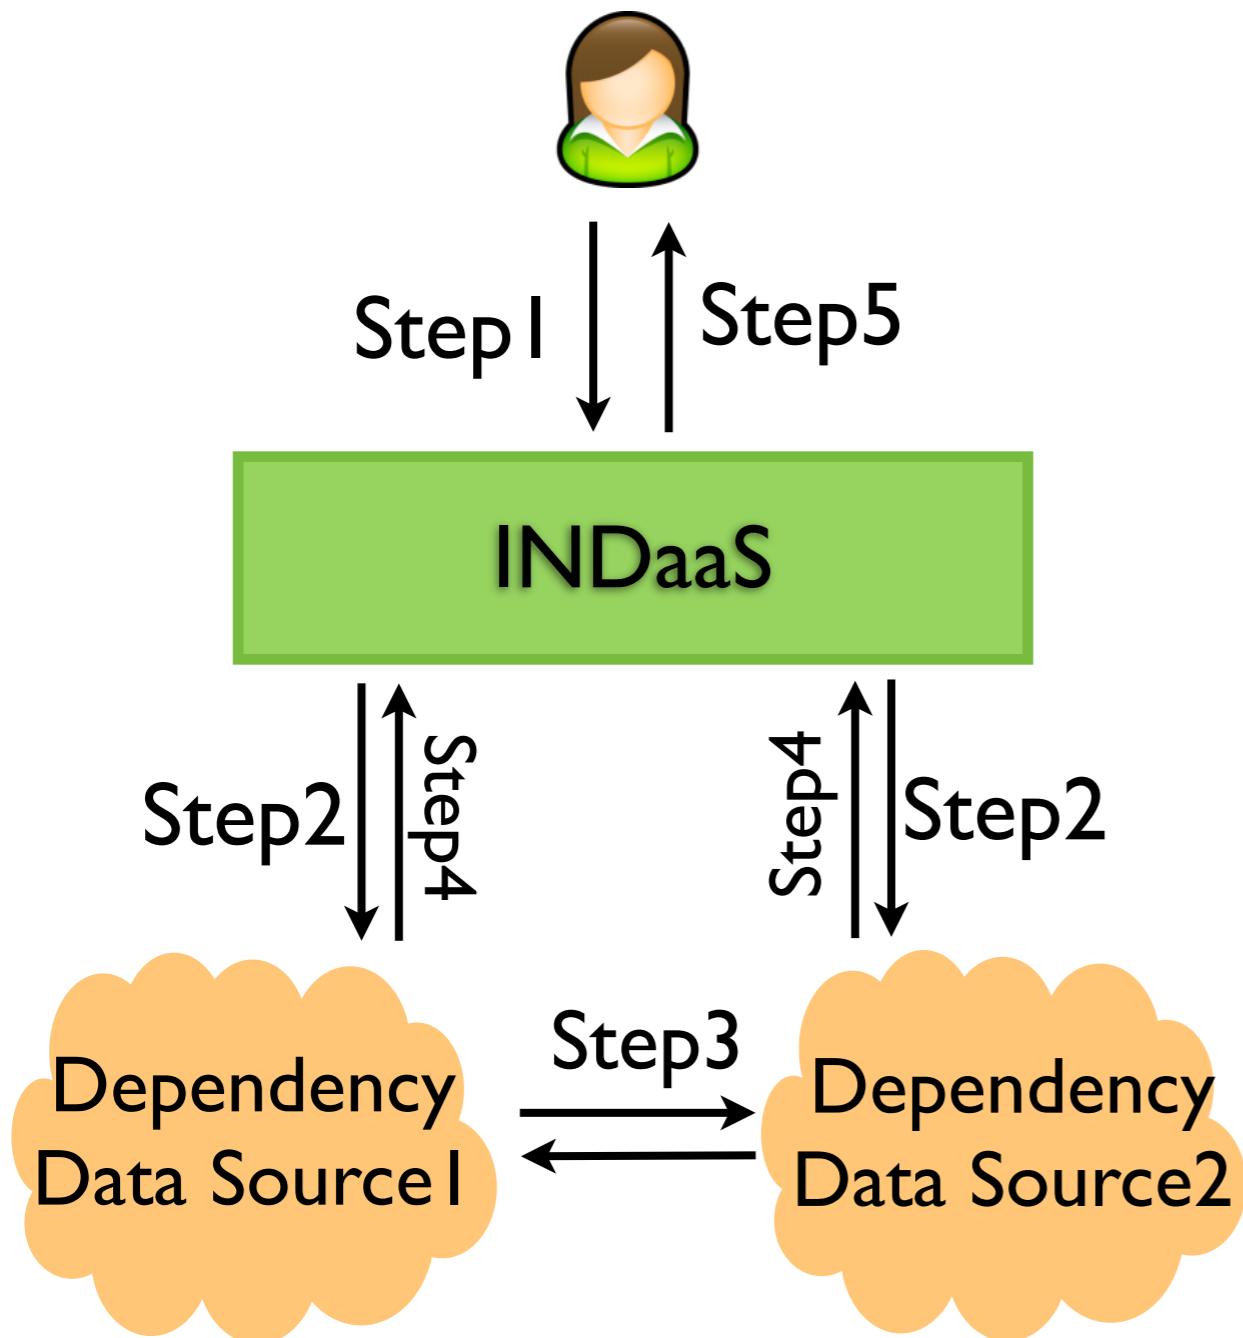
# Private Independence Evaluation



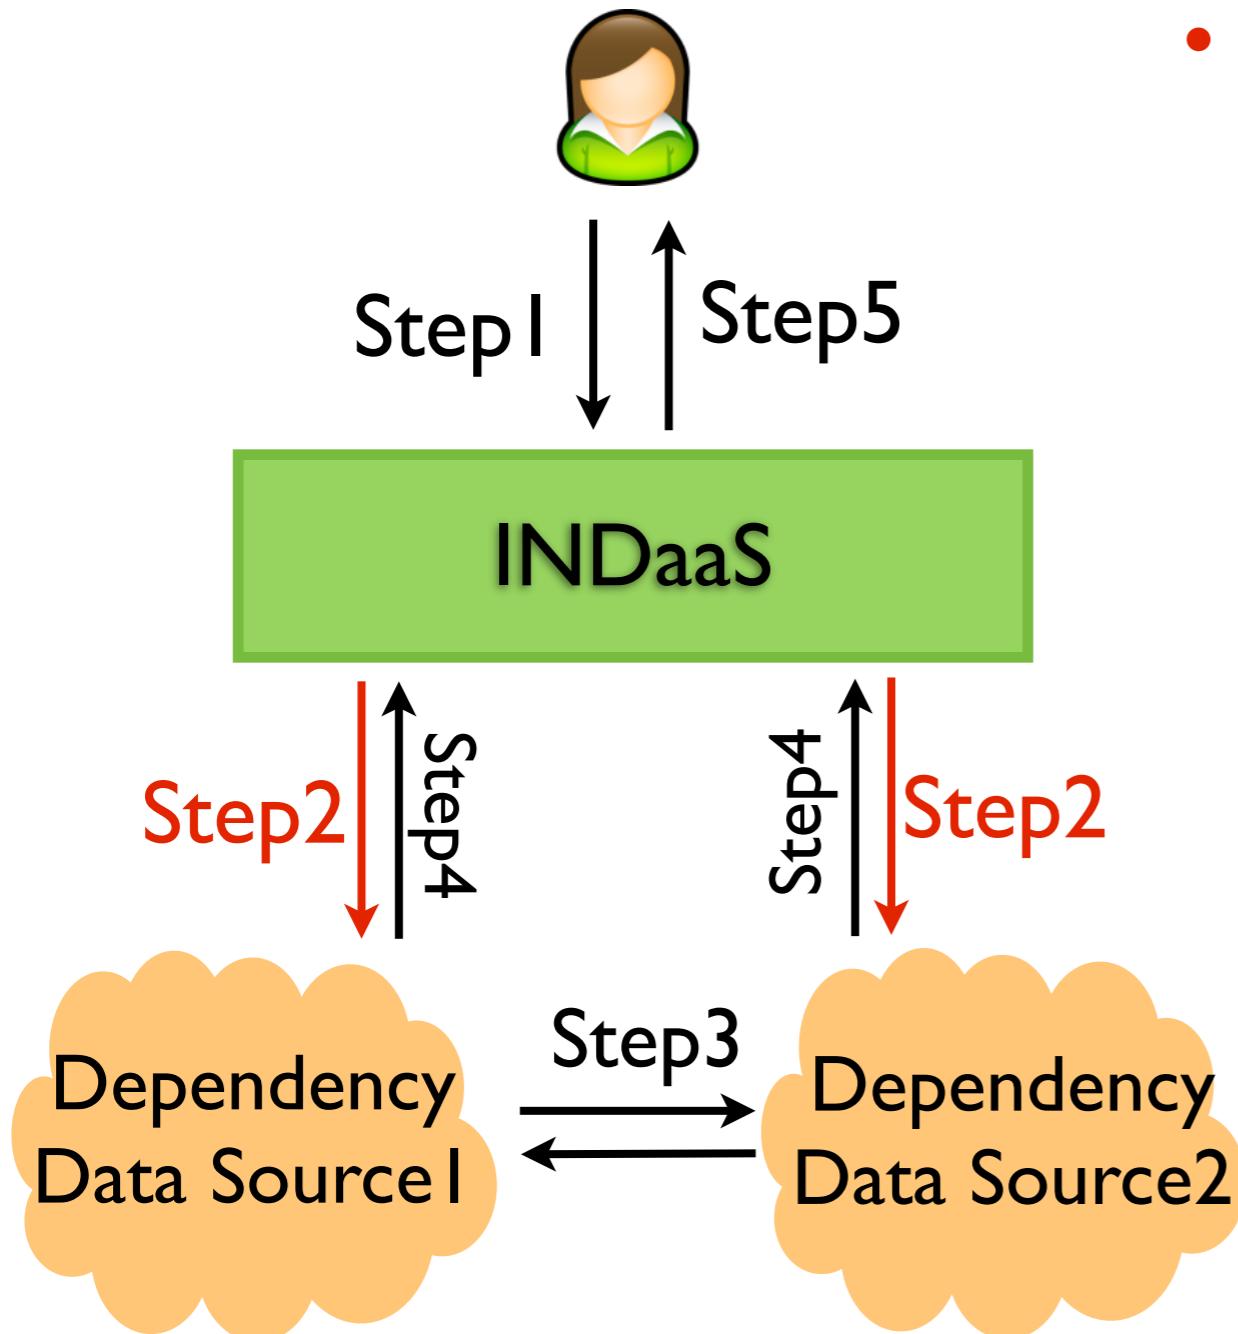
# Private Independence Evaluation



# Technical Challenges

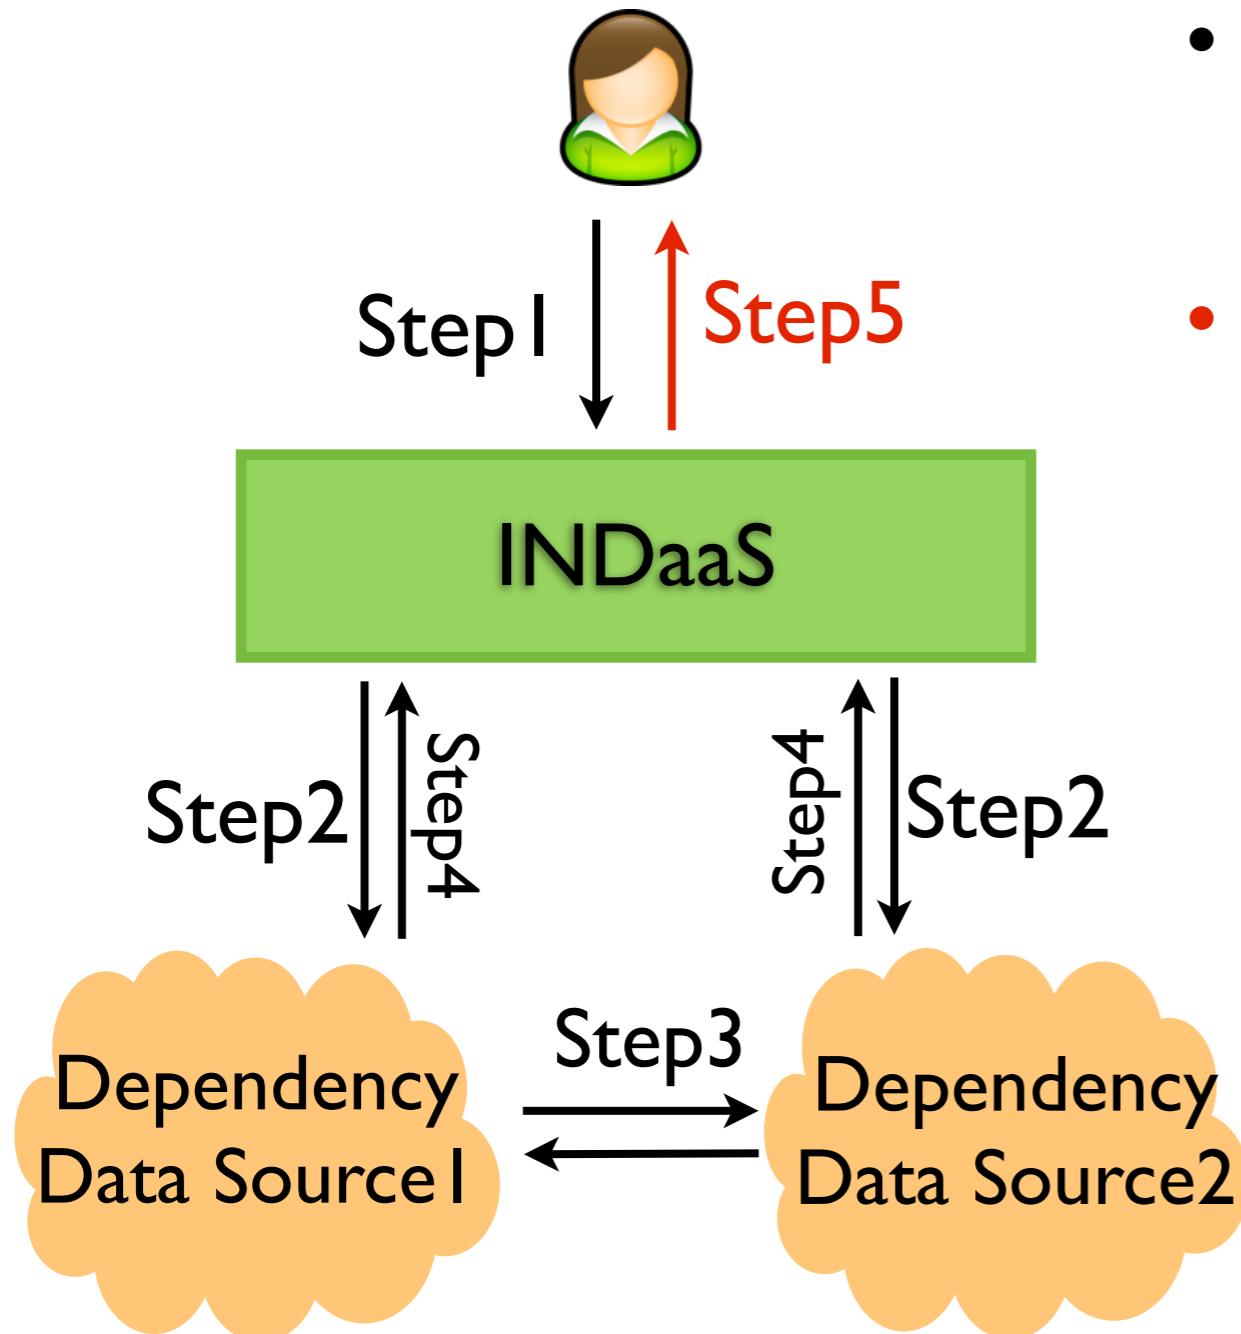


# Technical Challenges



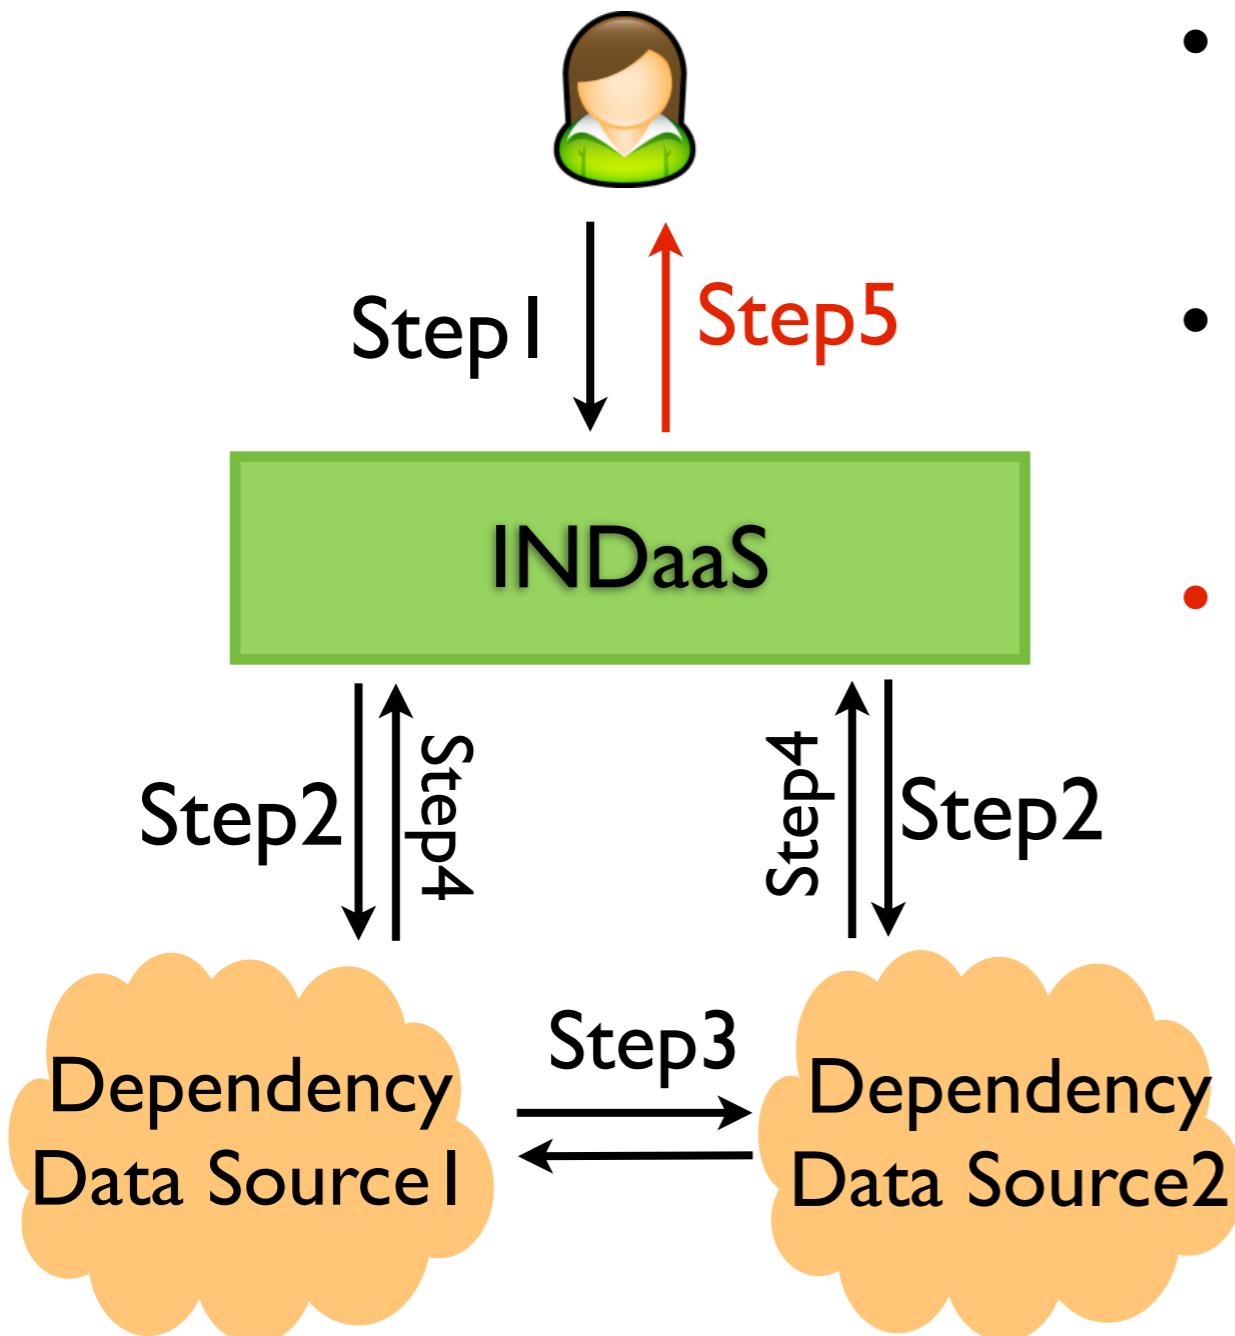
- #1: Dependency collections
  - Solution: Reusing existing tools

# Technical Challenges



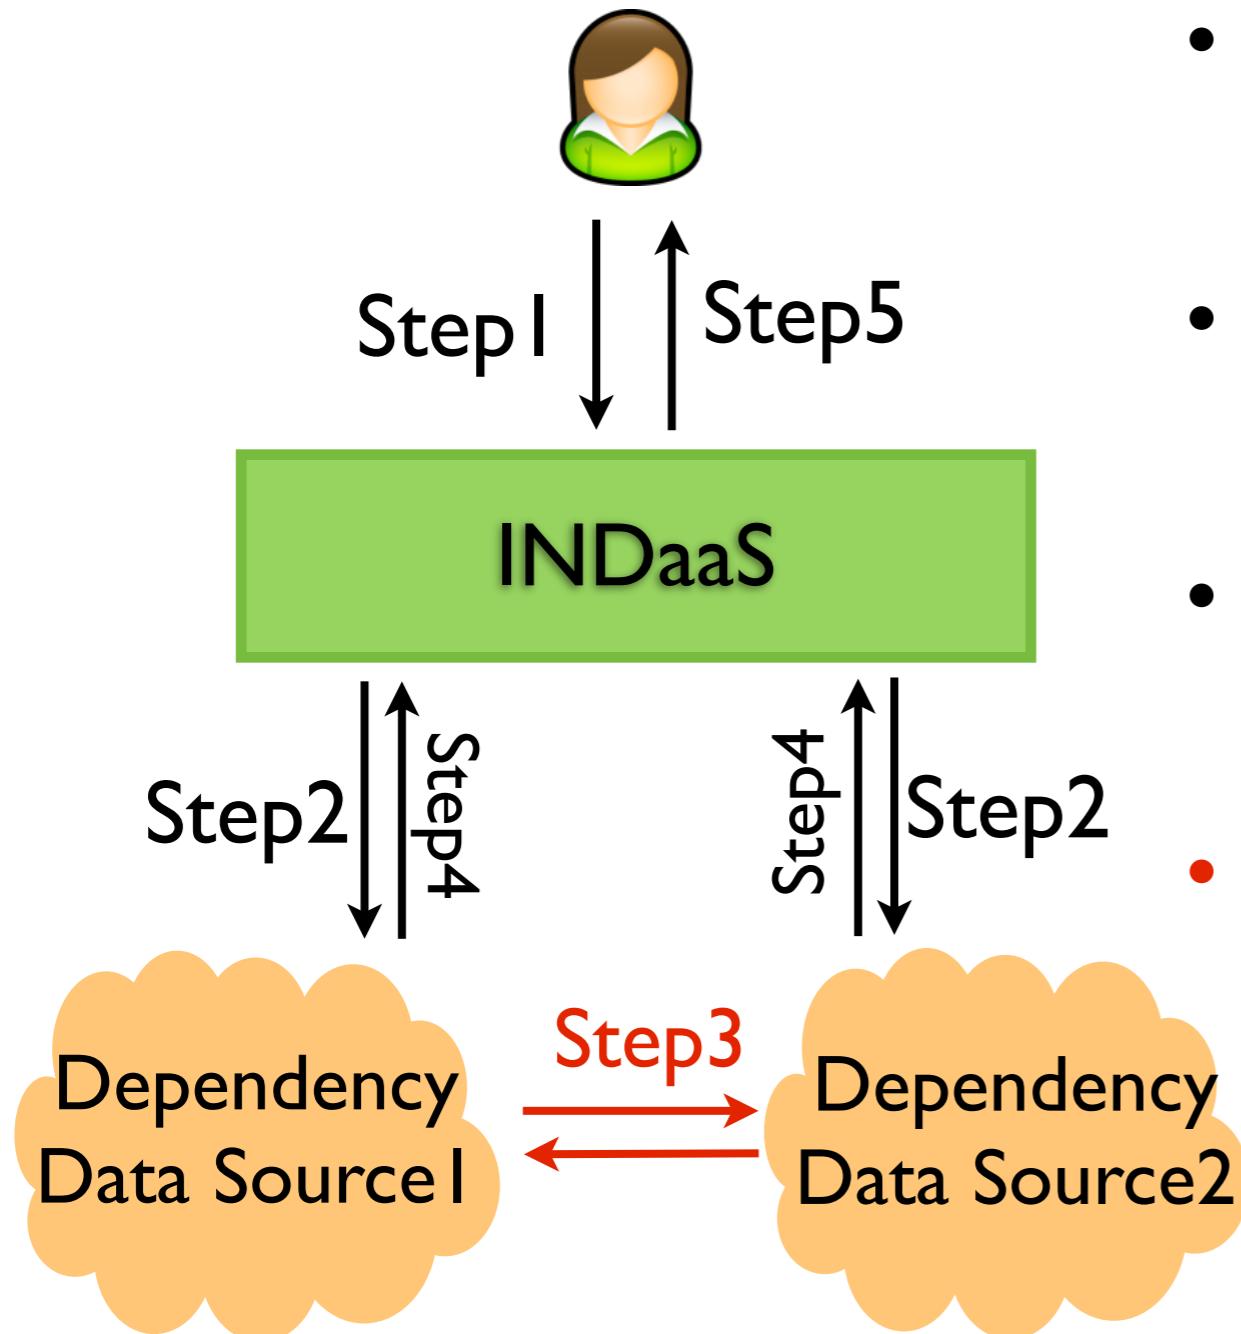
- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs

# Technical Challenges



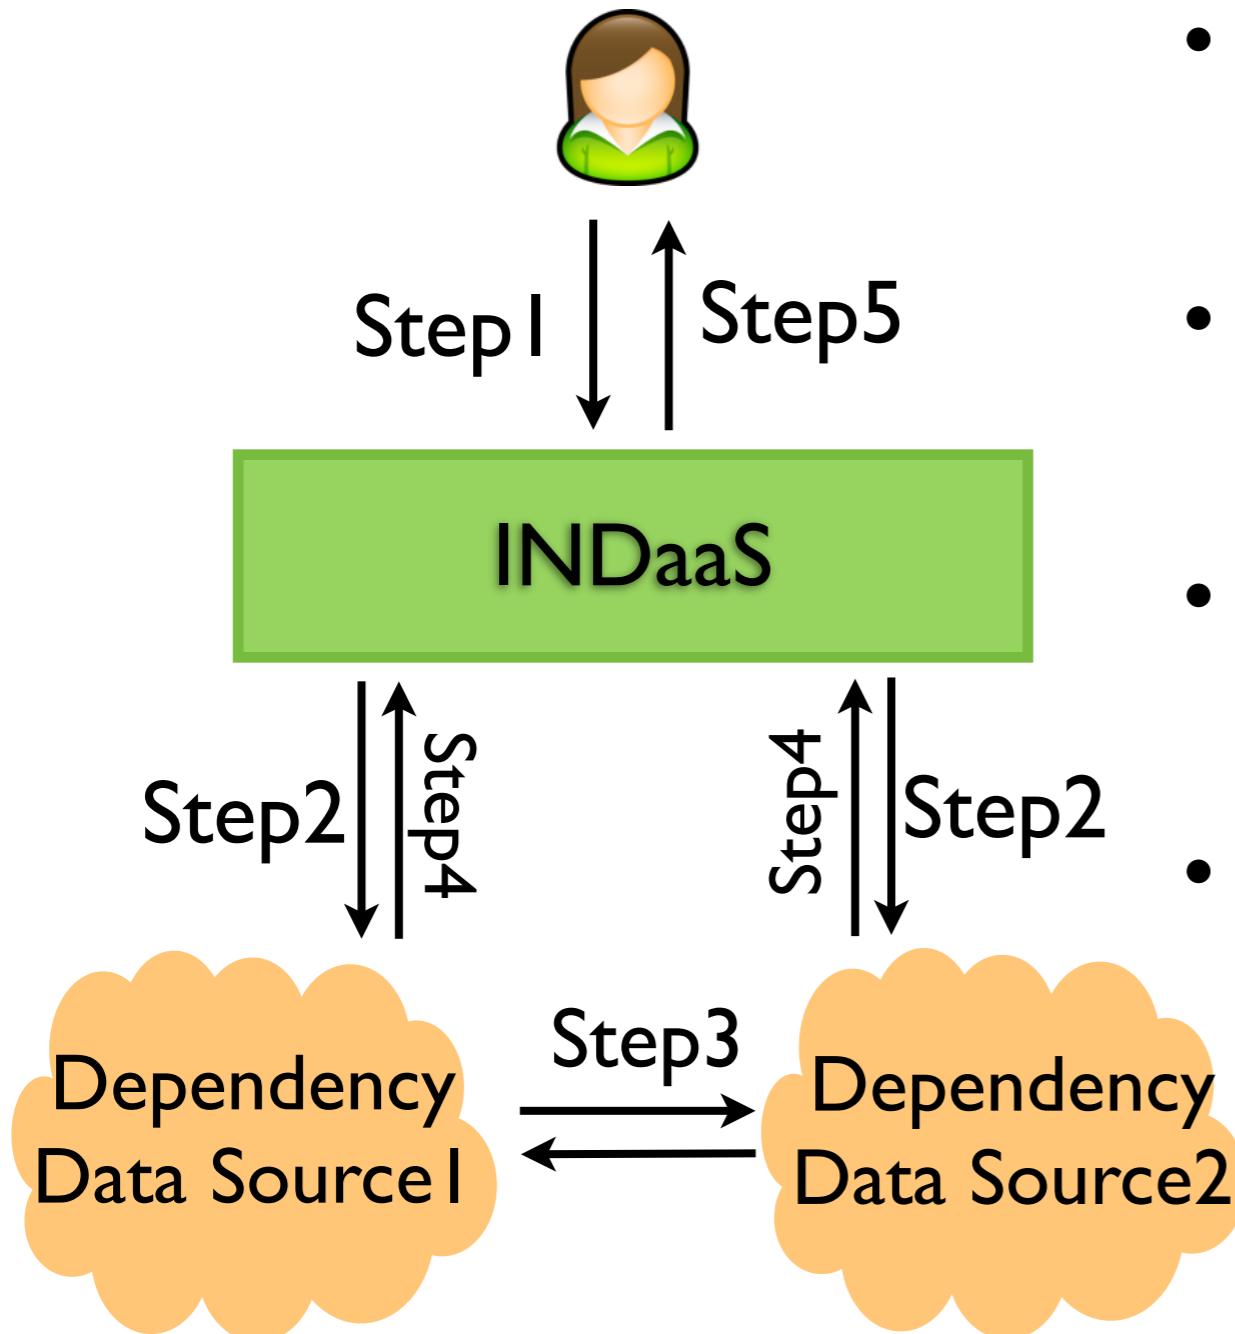
- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm

# Technical Challenges



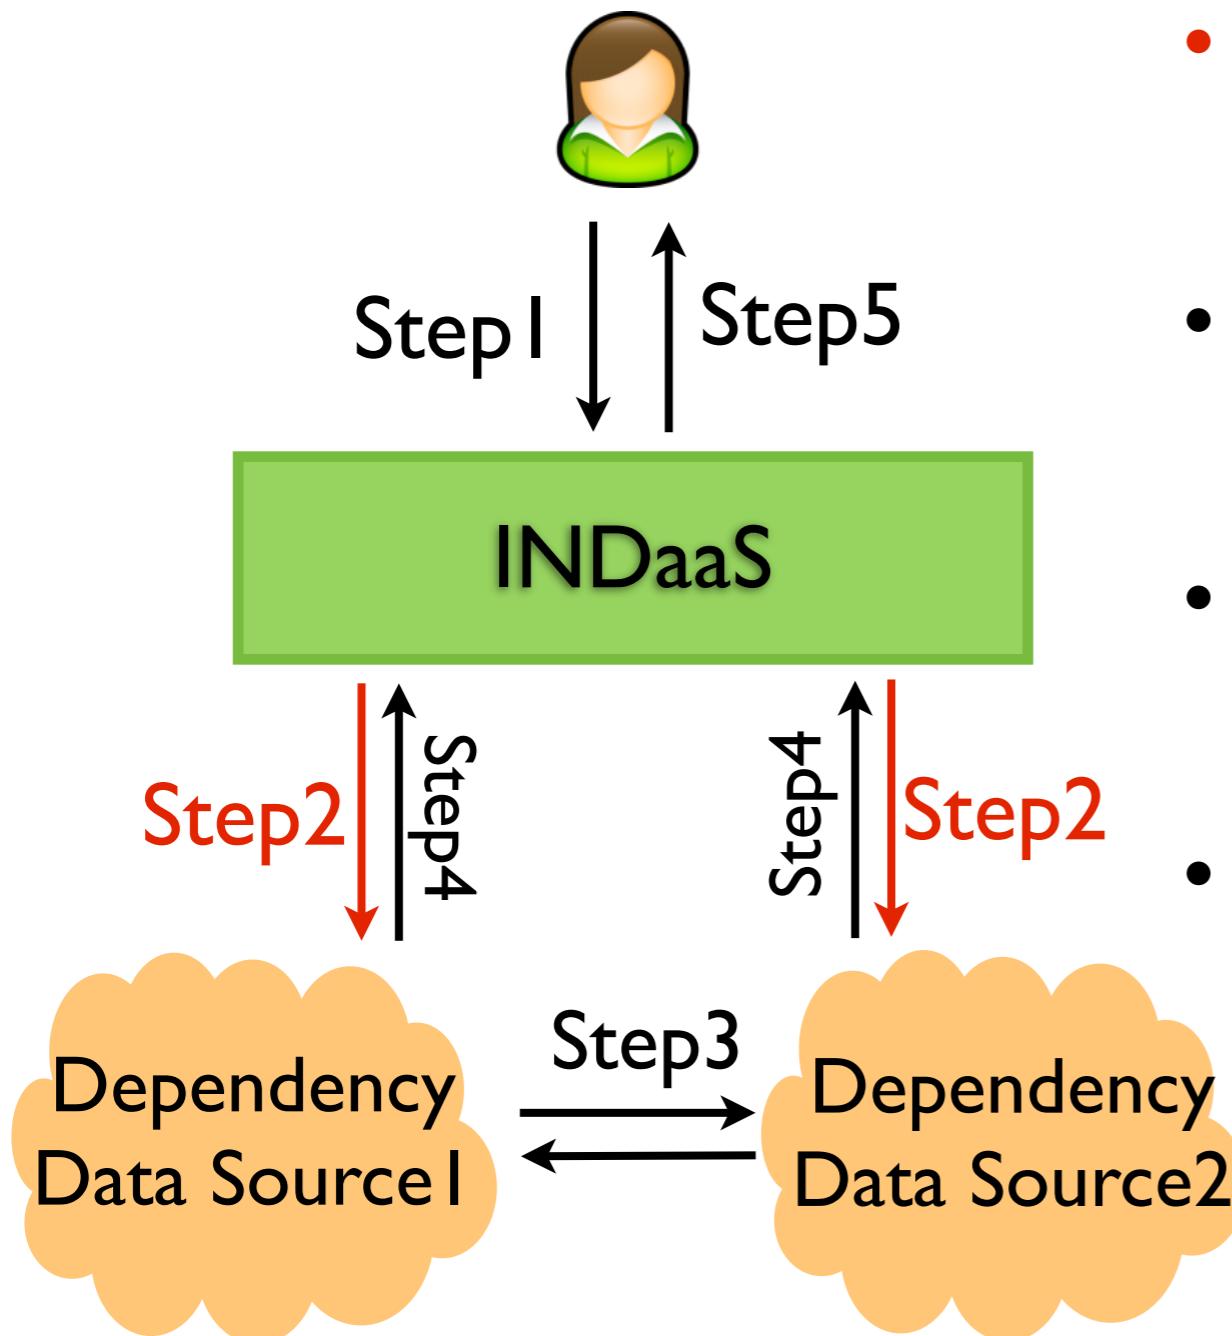
- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

# RoadMap



- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

# RoadMap



- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

# Dependency Data Collections

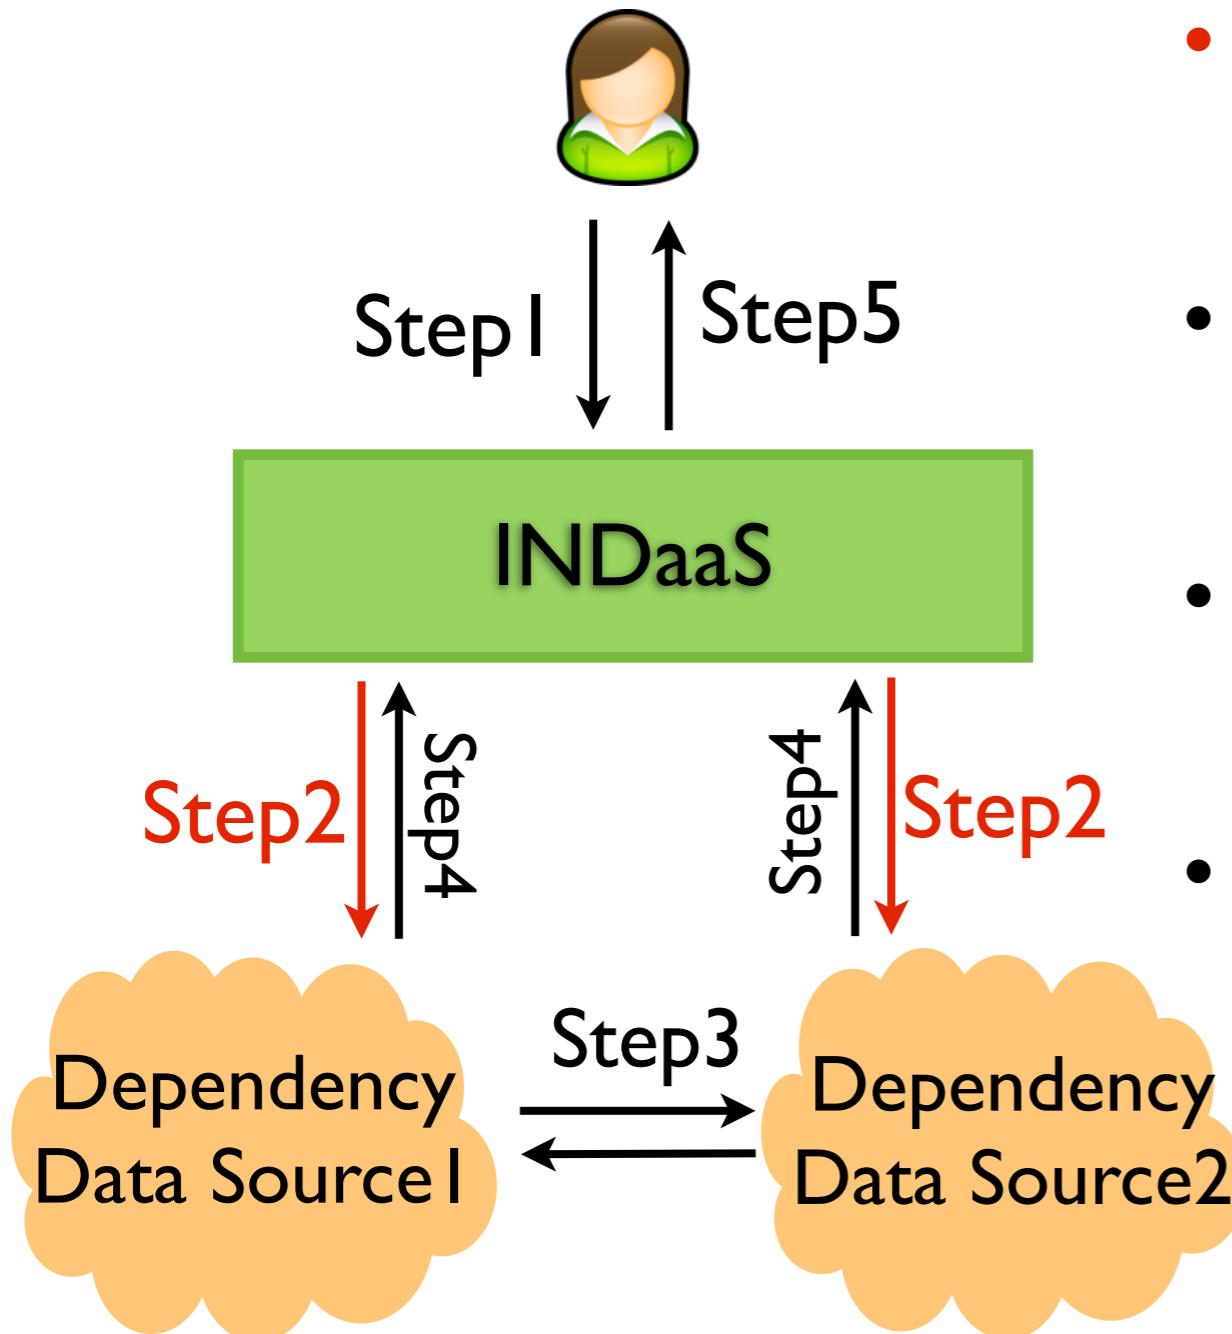
- Reuse existing data collection tools:
  - Convert the outputs to uniform format.
  - Three types of format: NET, HW and SW.

## Our defined format

| Type     | Dependency Expression            |
|----------|----------------------------------|
| Network  | <src="S" dst="D" route="x,y,z"/> |
| Hardware | <hw="H" type="T" dep="x"/>       |
| Software | <pgm="S" hw="H" dep="x,y,z"/>    |

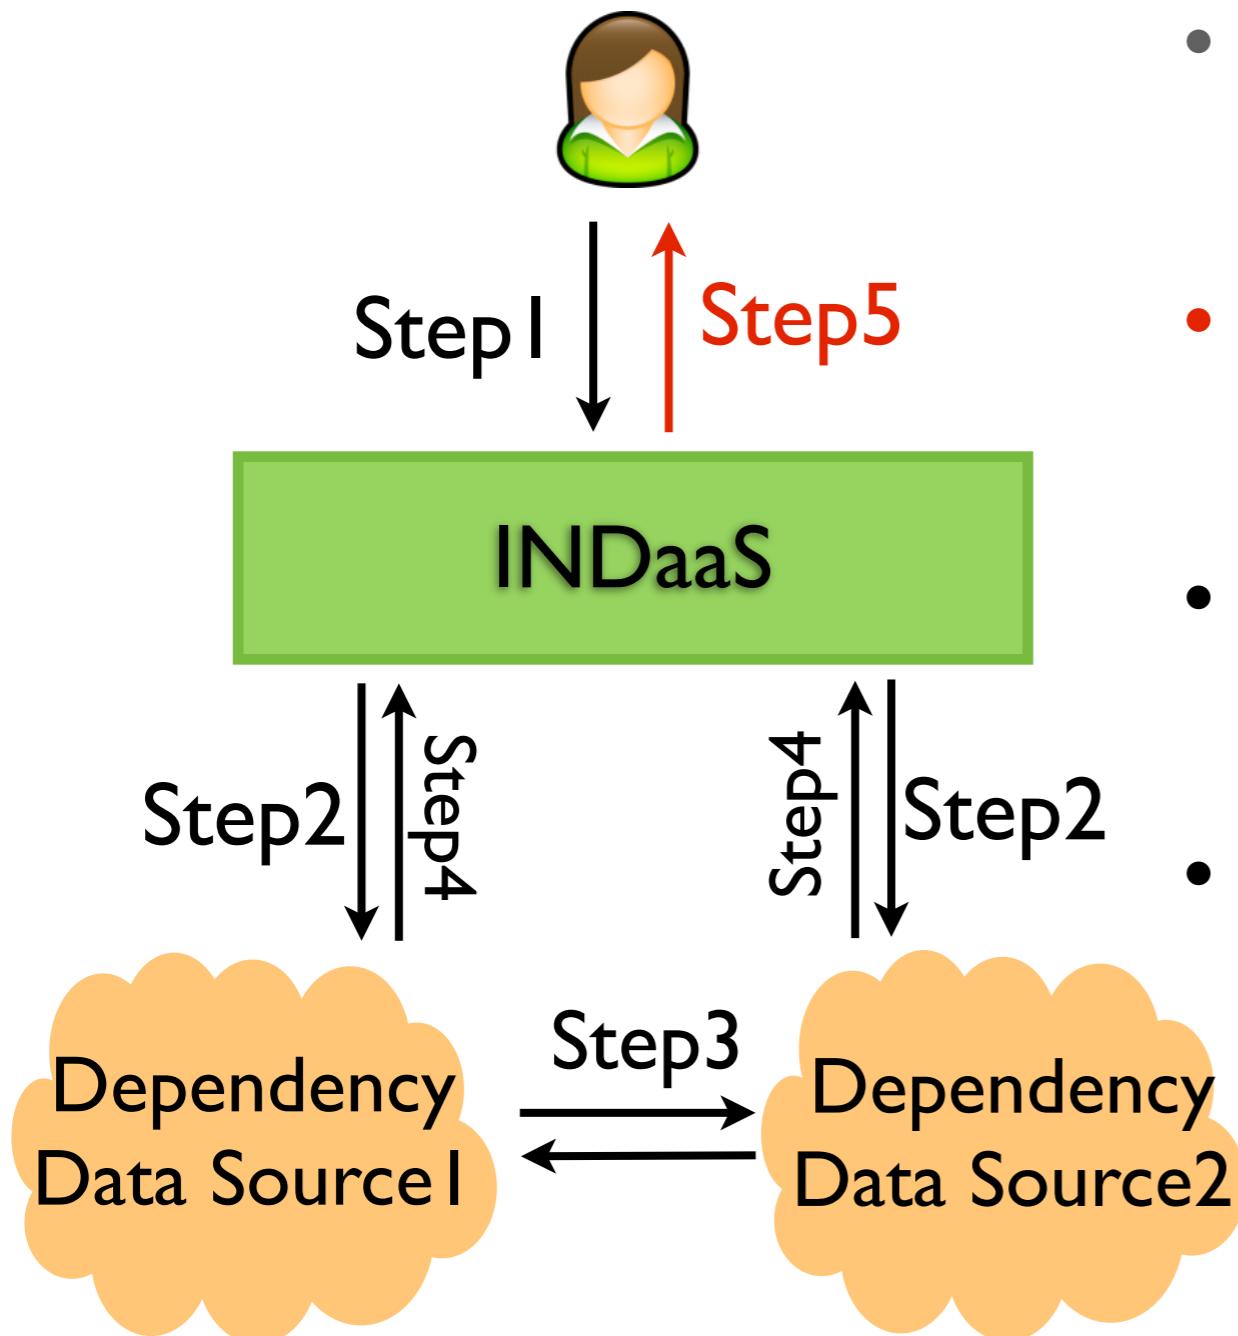
Please see our paper for more details

# RoadMap



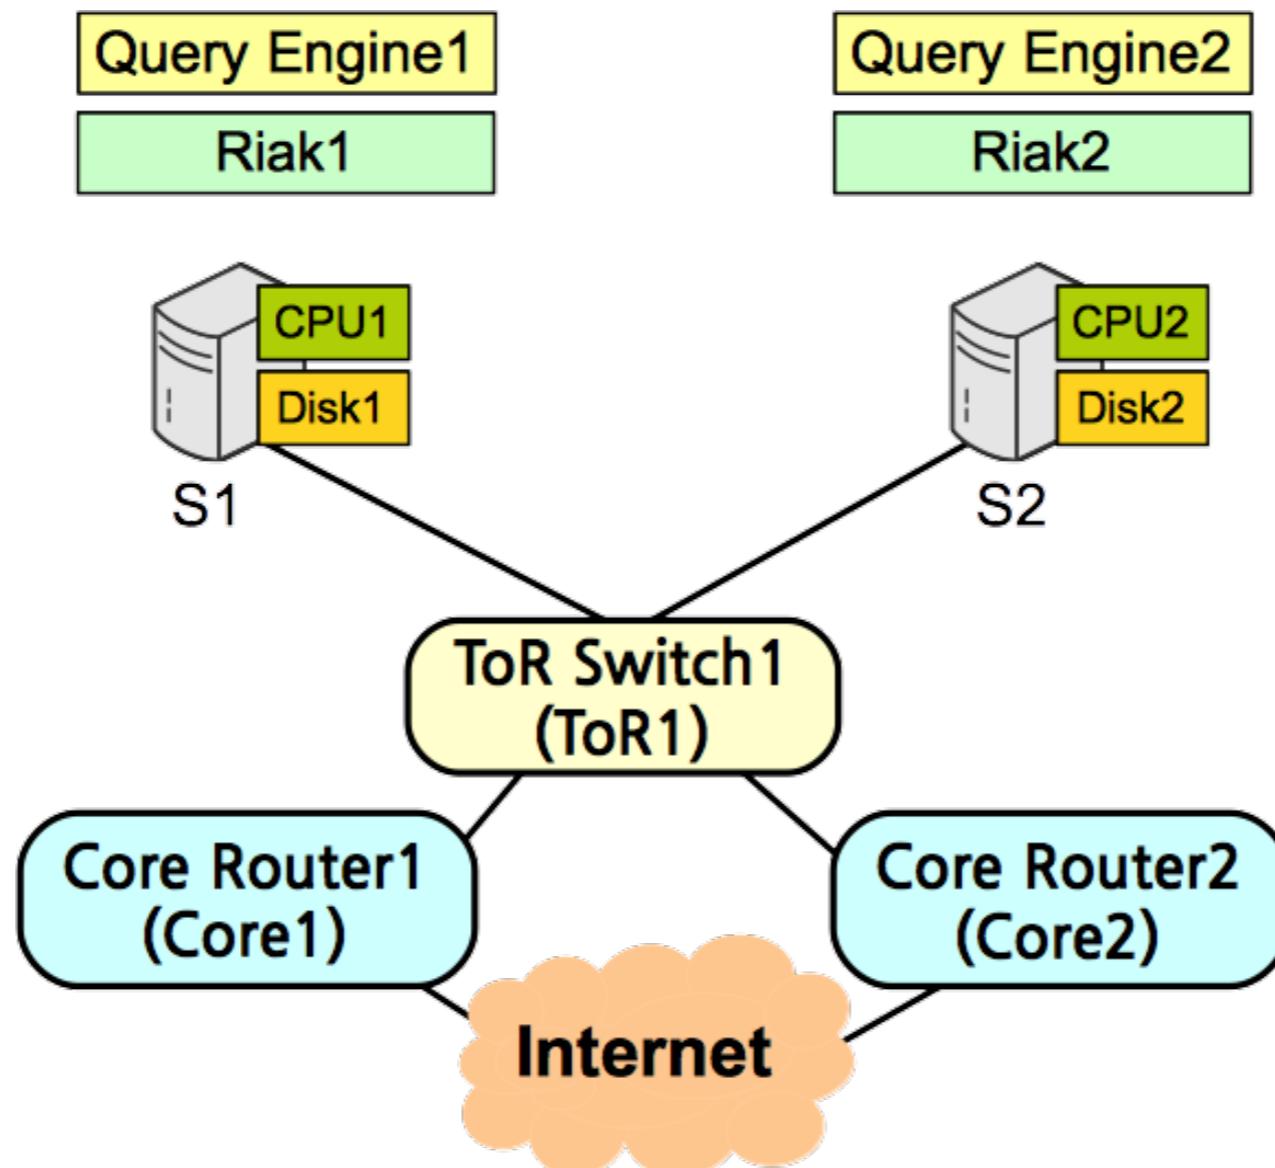
- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

# RoadMap

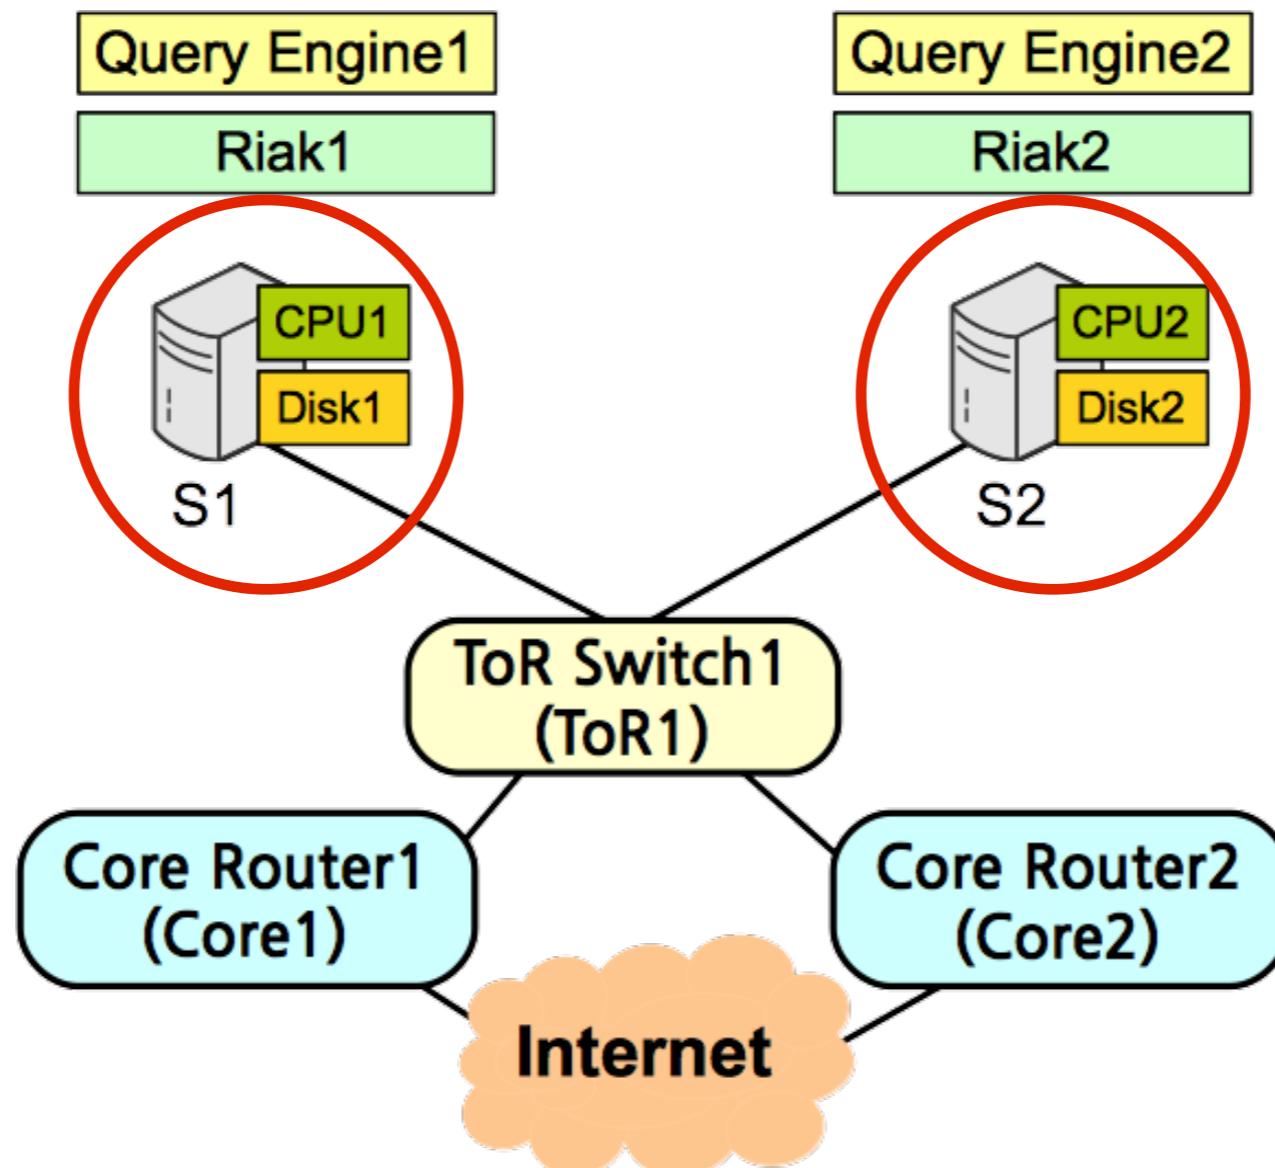


- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

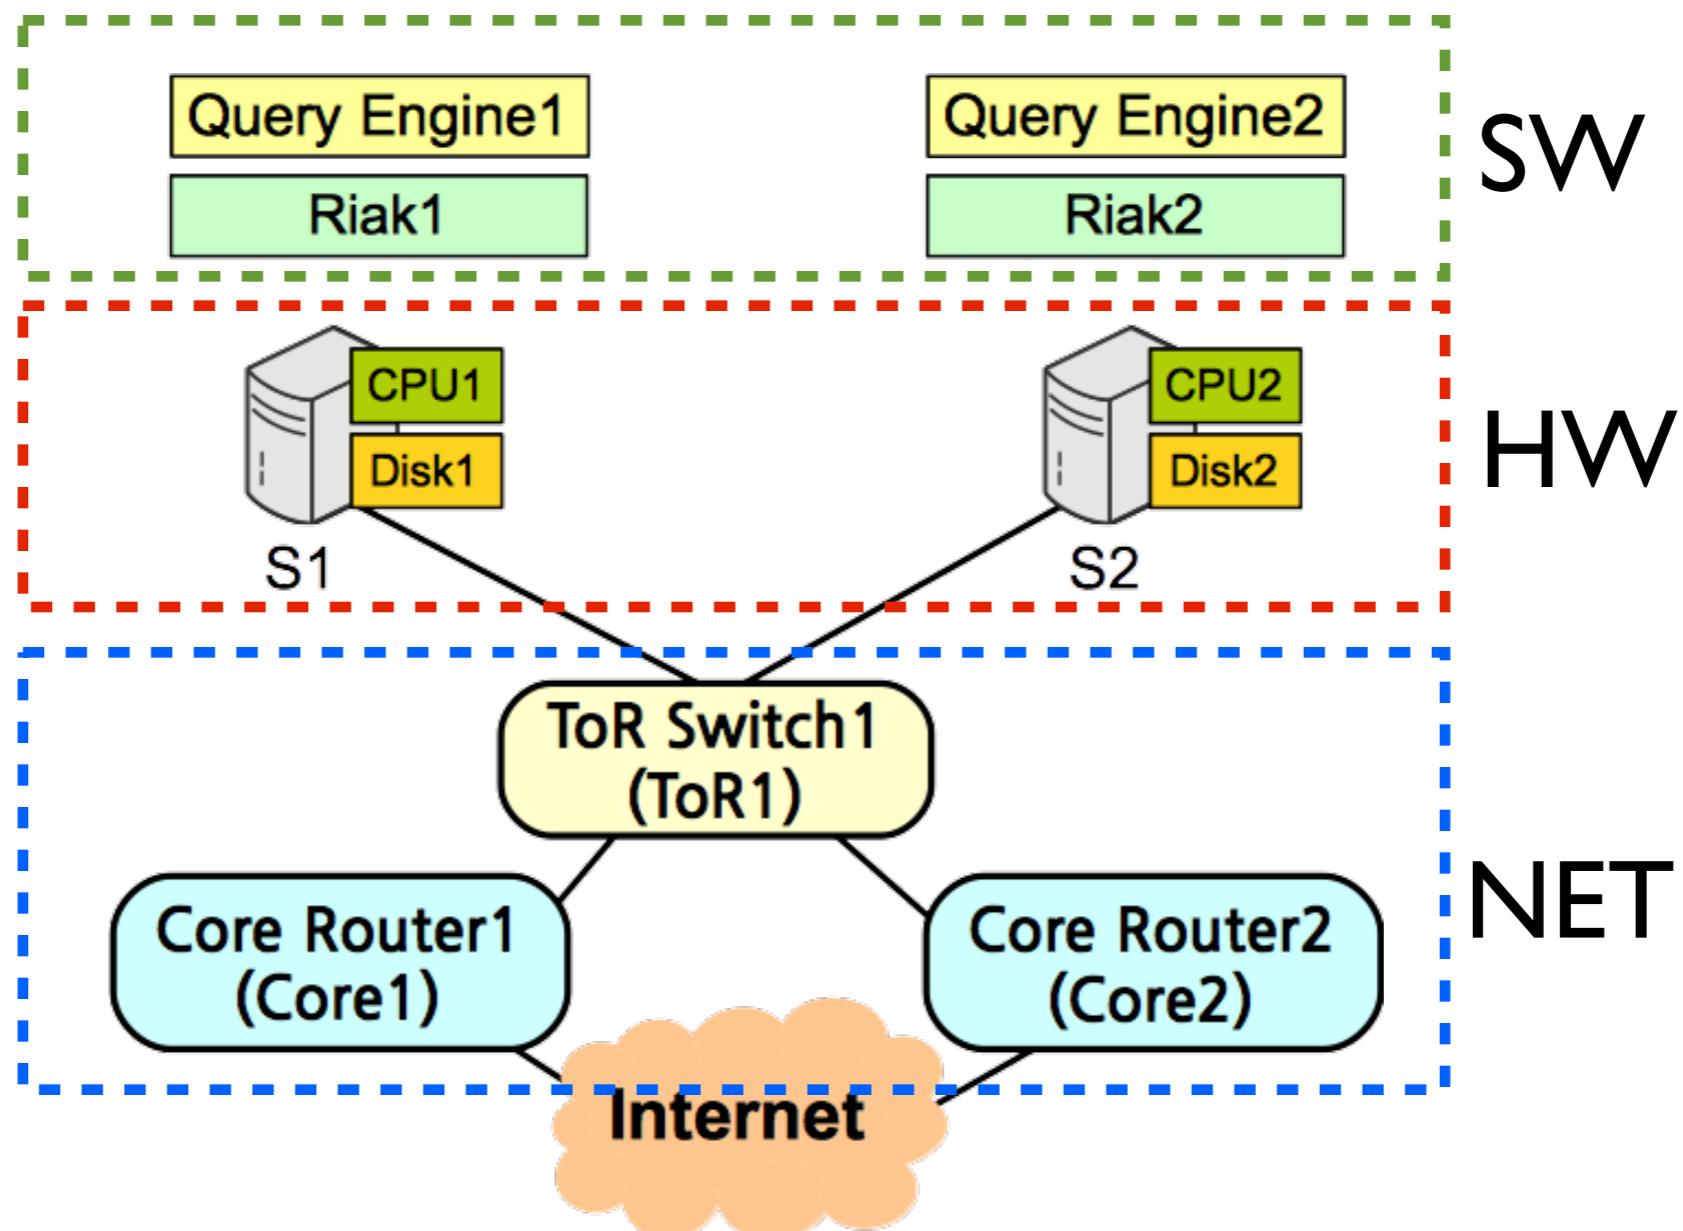
# Example Redundancy



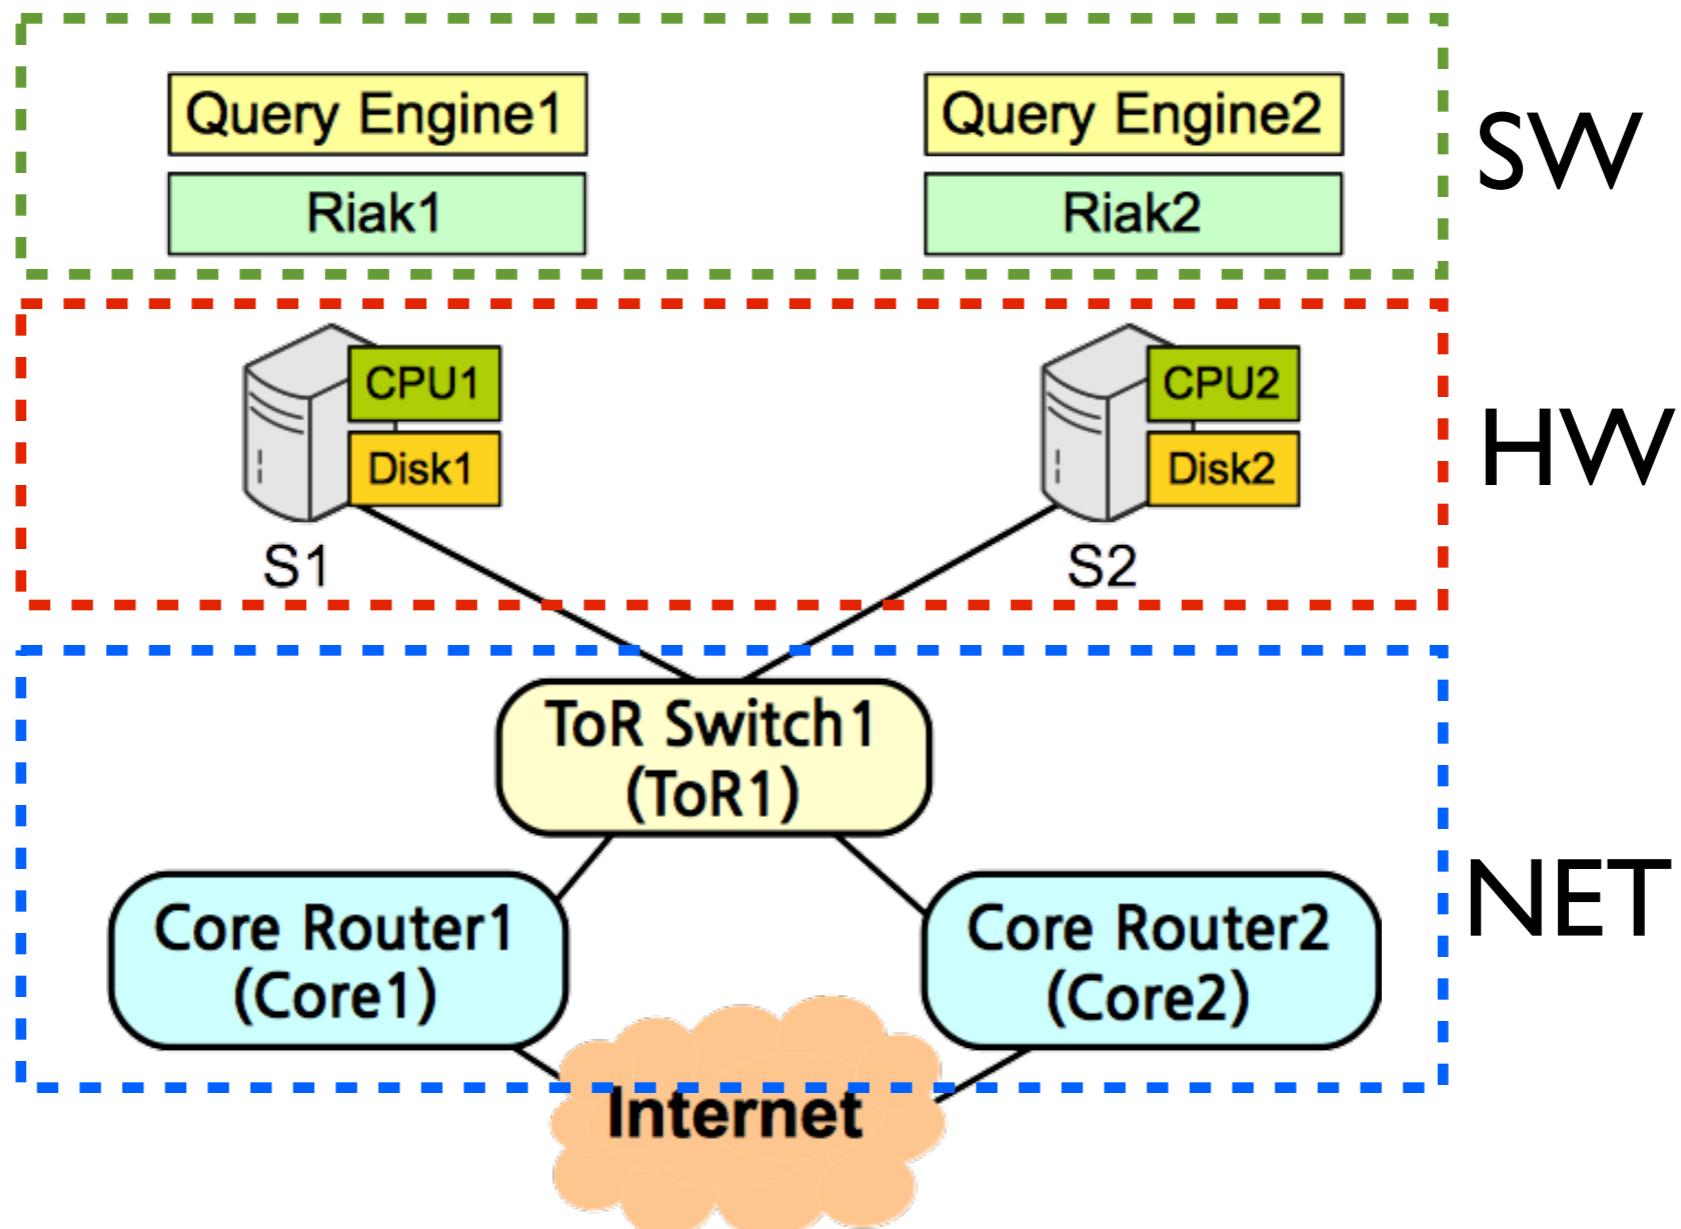
# Example Redundancy



# Example Redundancy



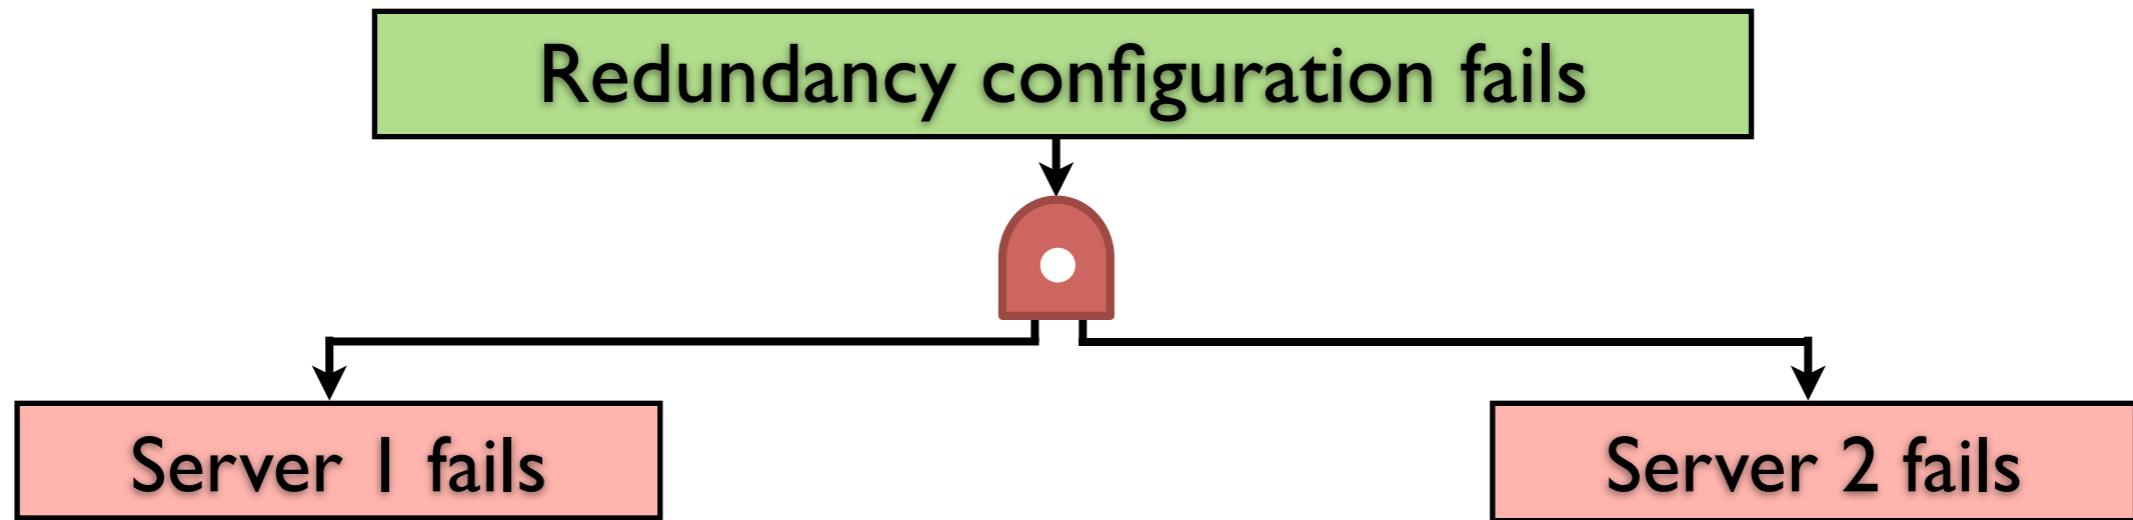
# Building Fault Graph Top-to-Bottom



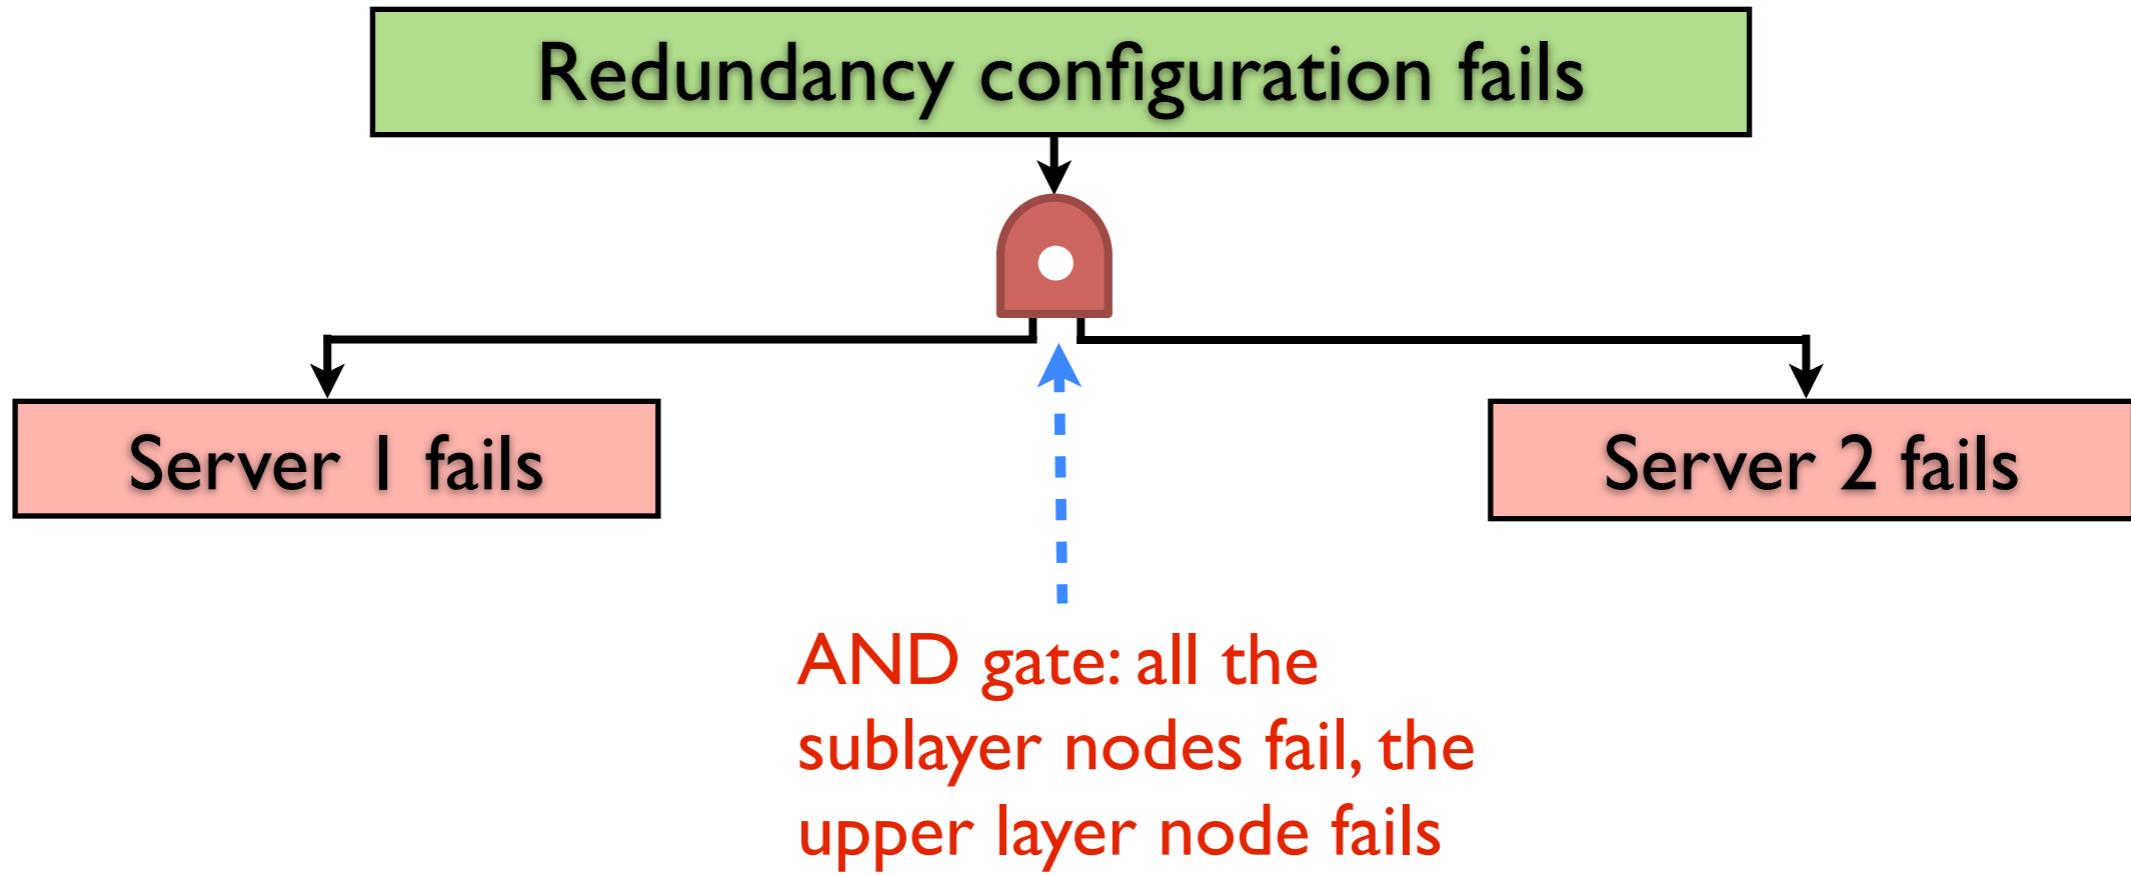
# Step1: Root Node

Redundancy configuration fails

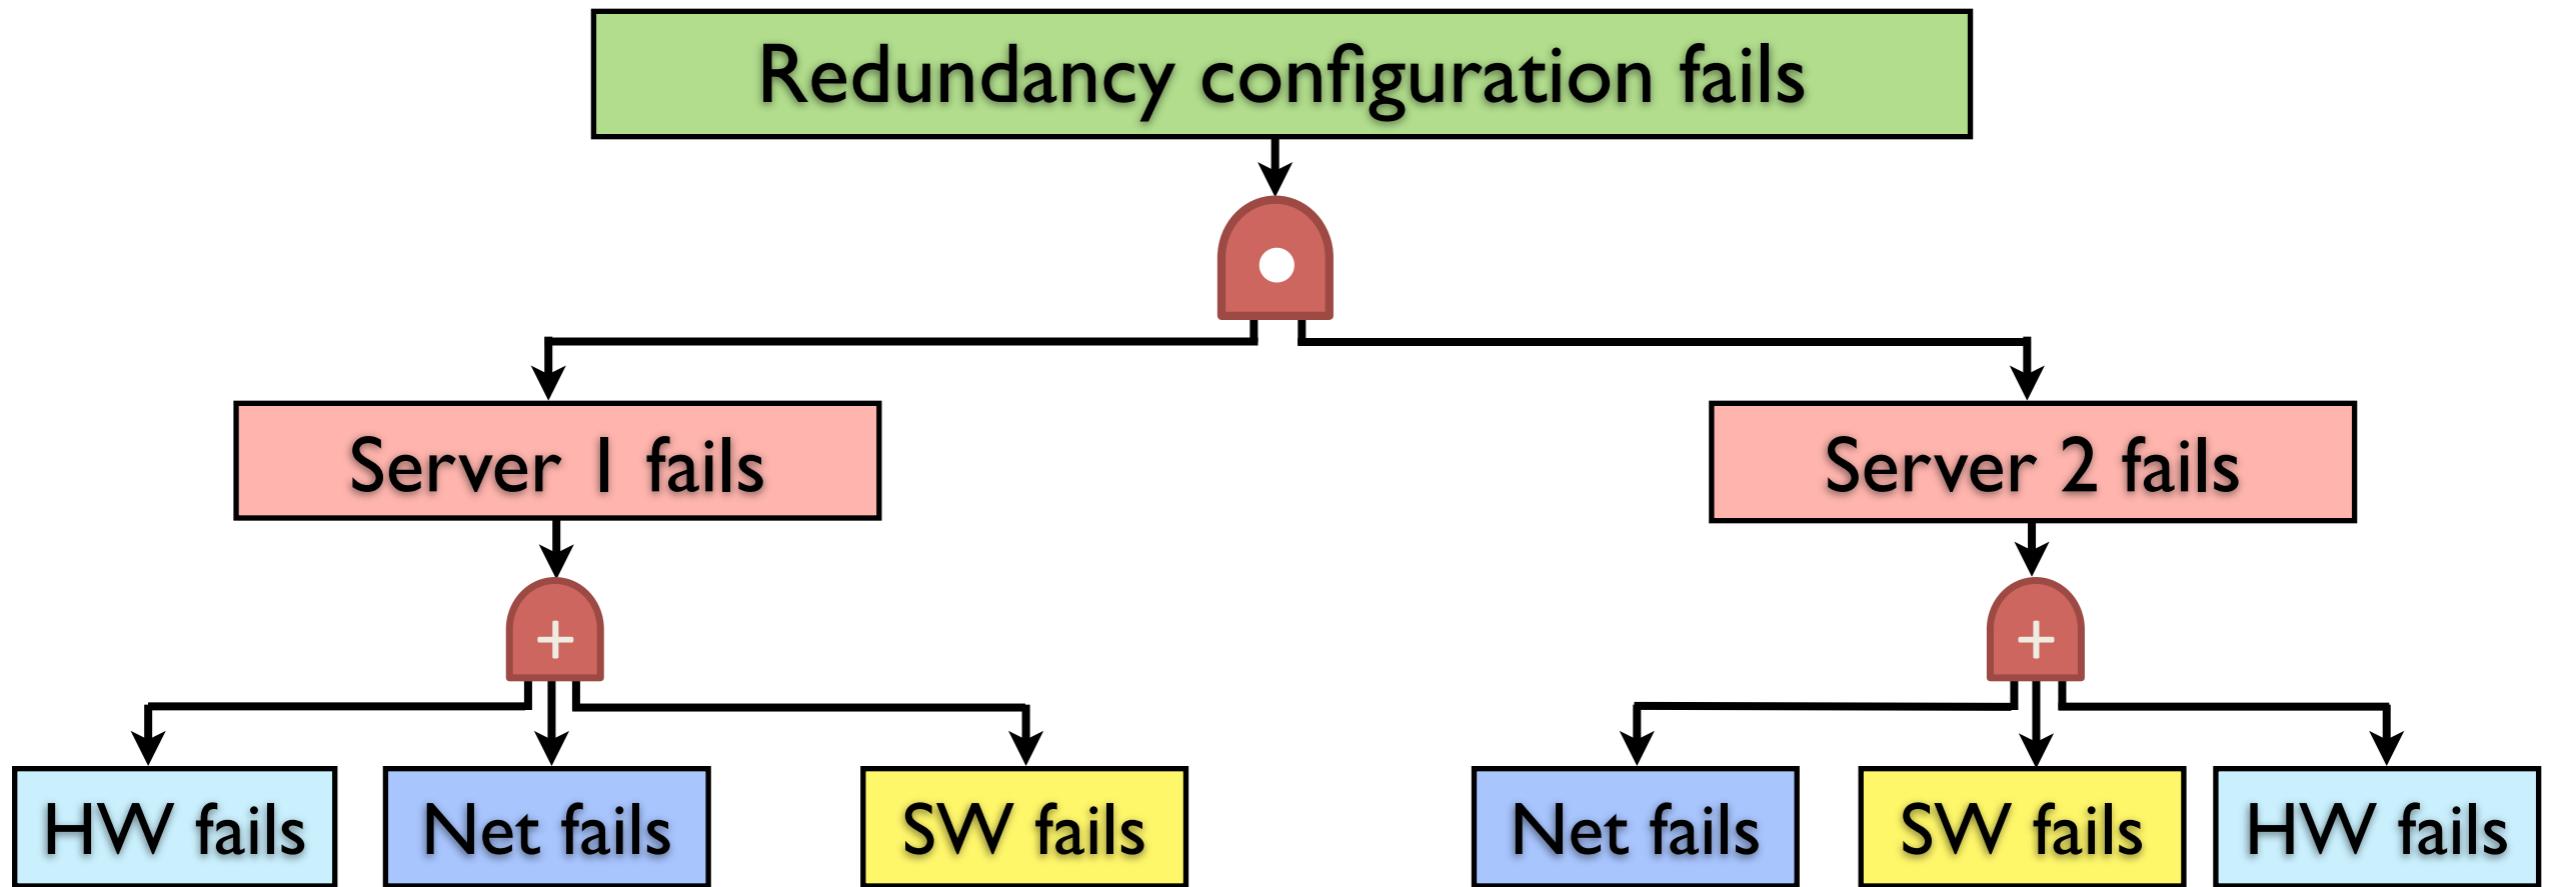
# Step2: Server Nodes



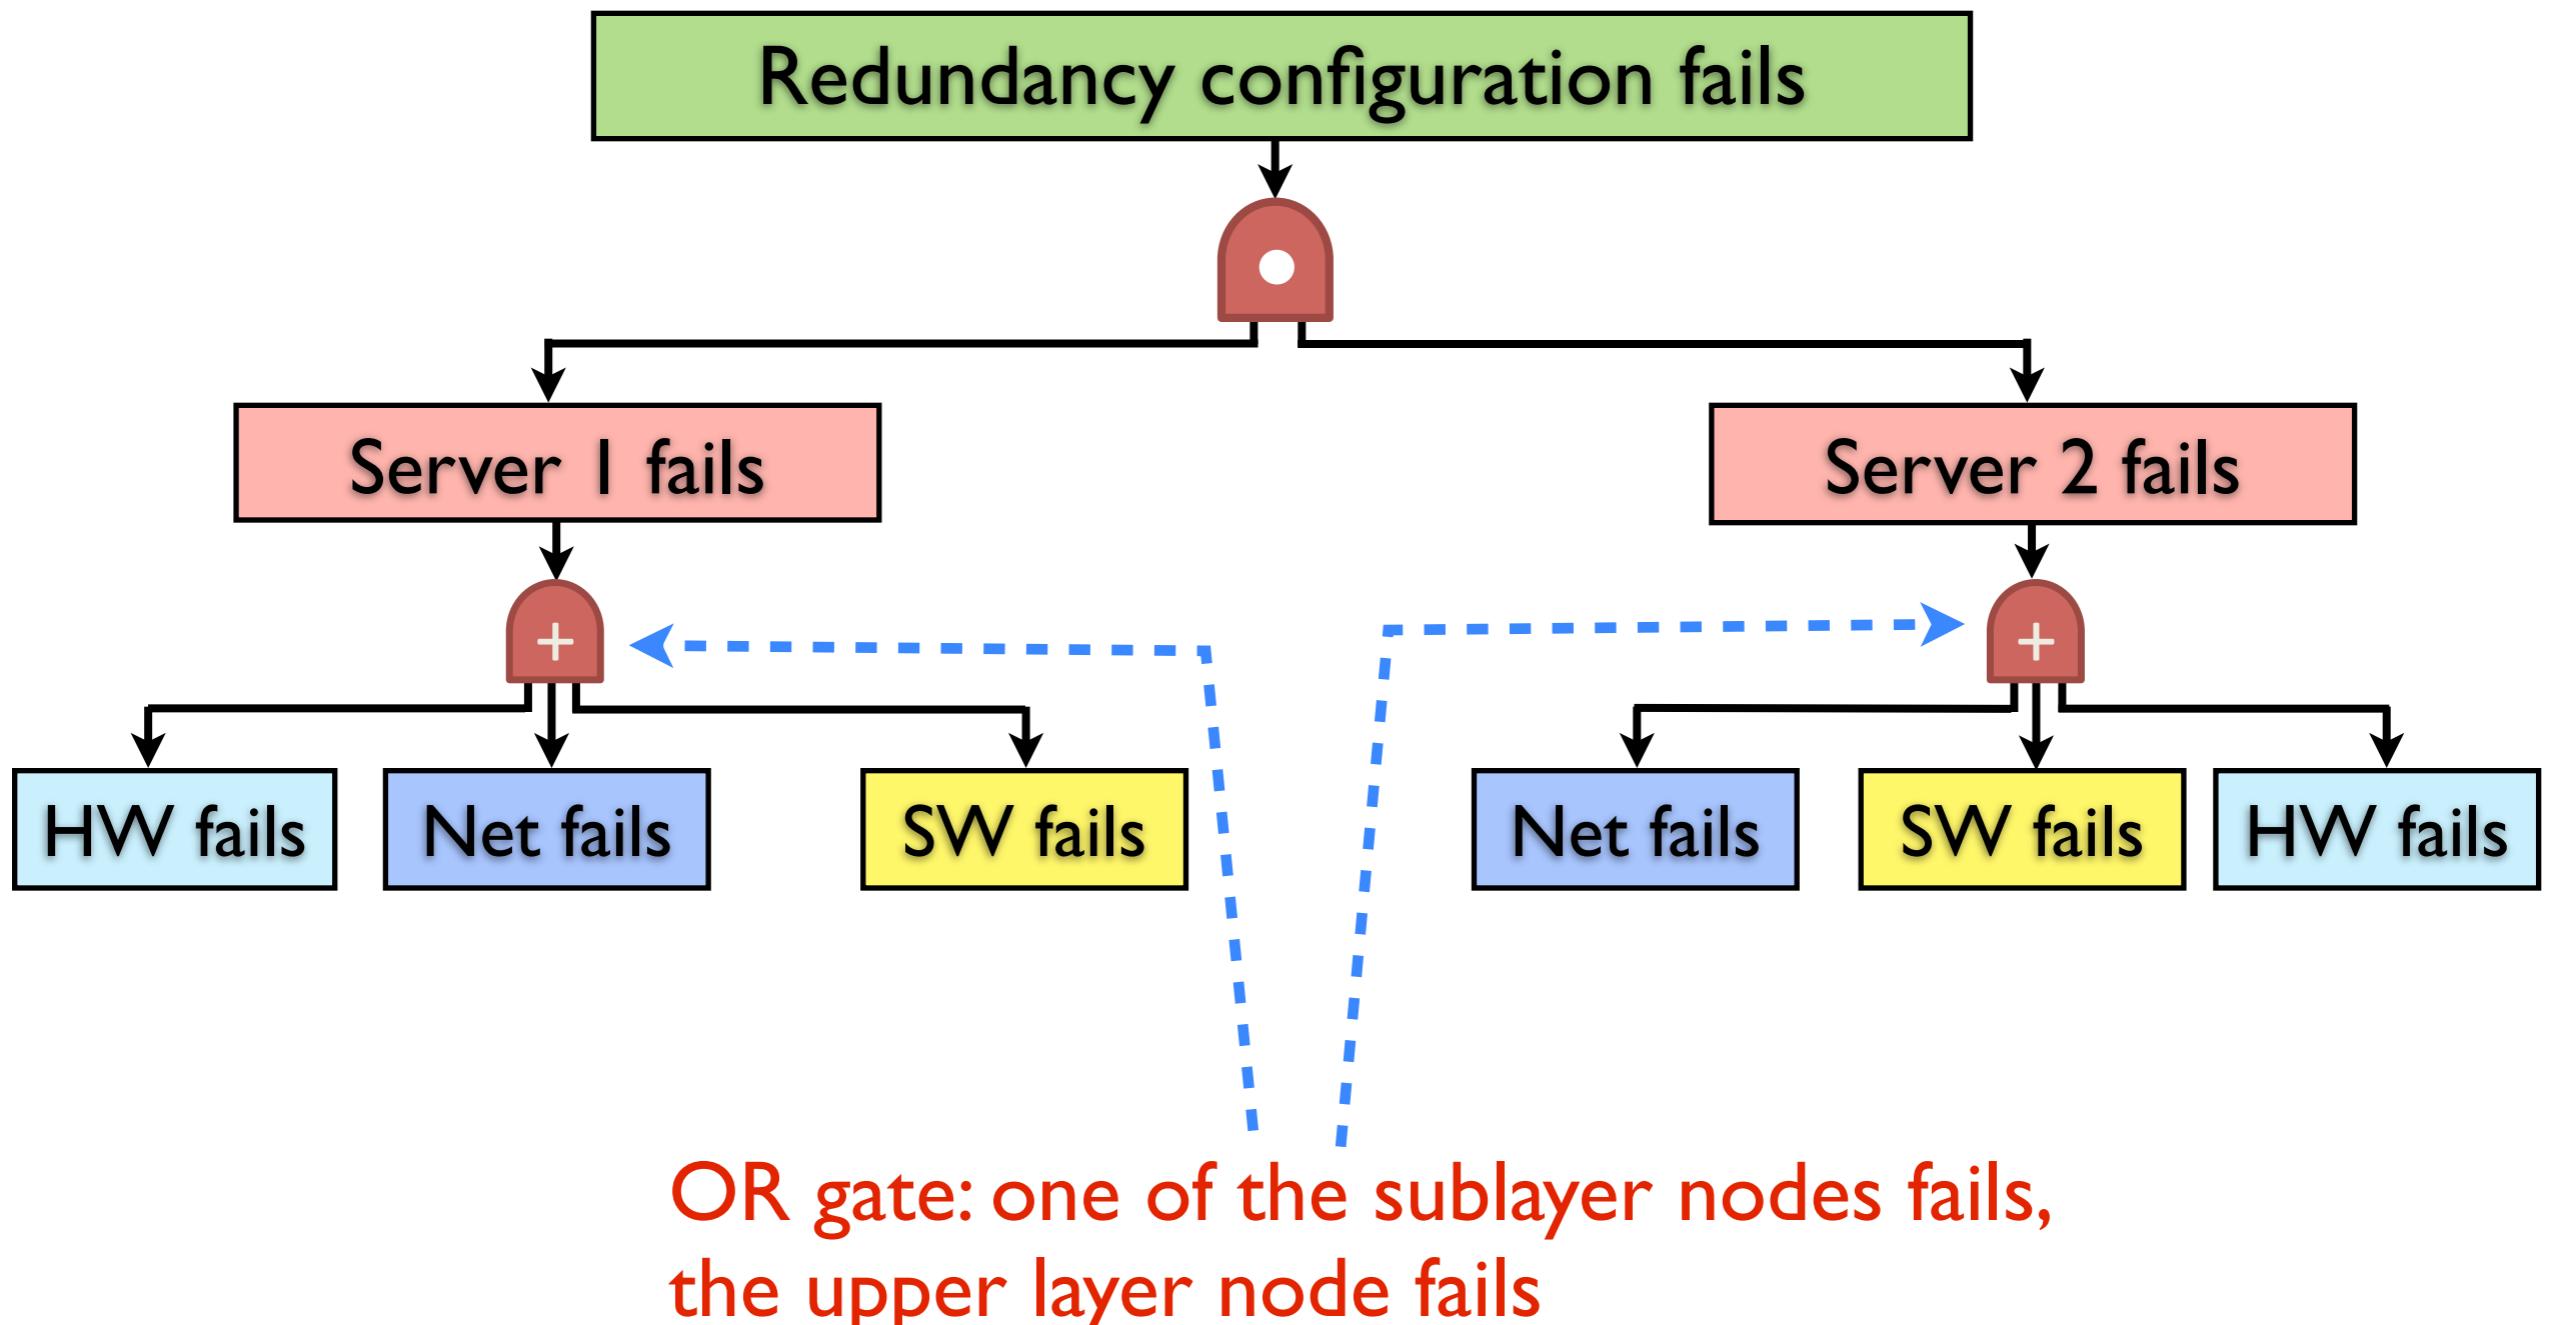
# Step2: Server Nodes



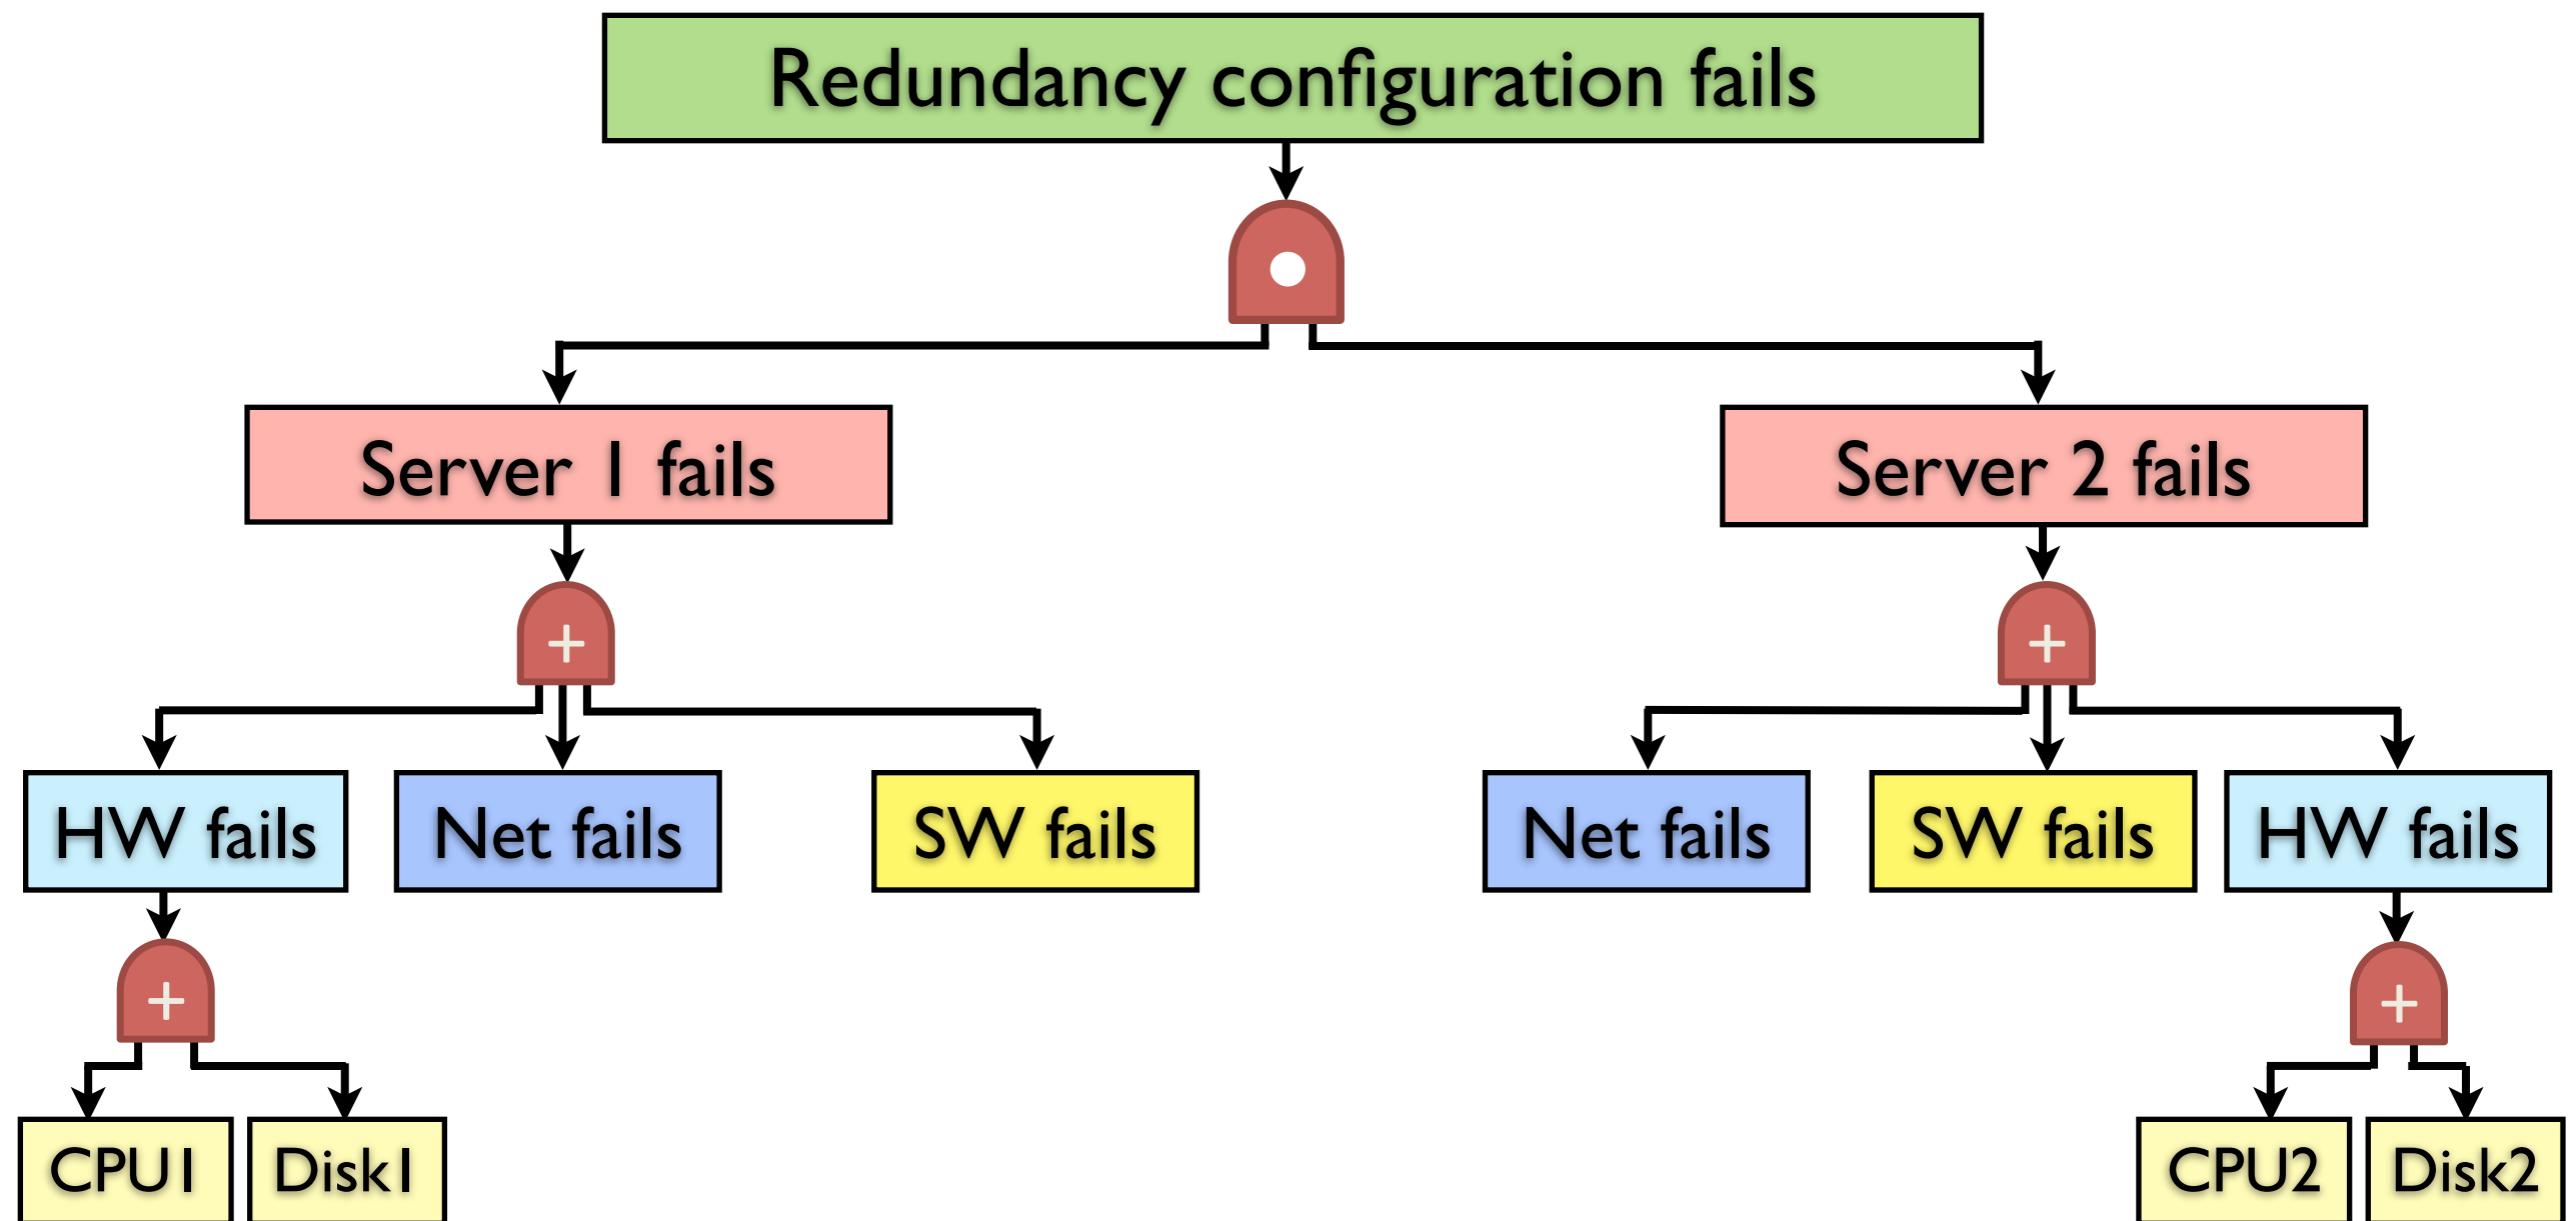
# Step3: Dependency Nodes



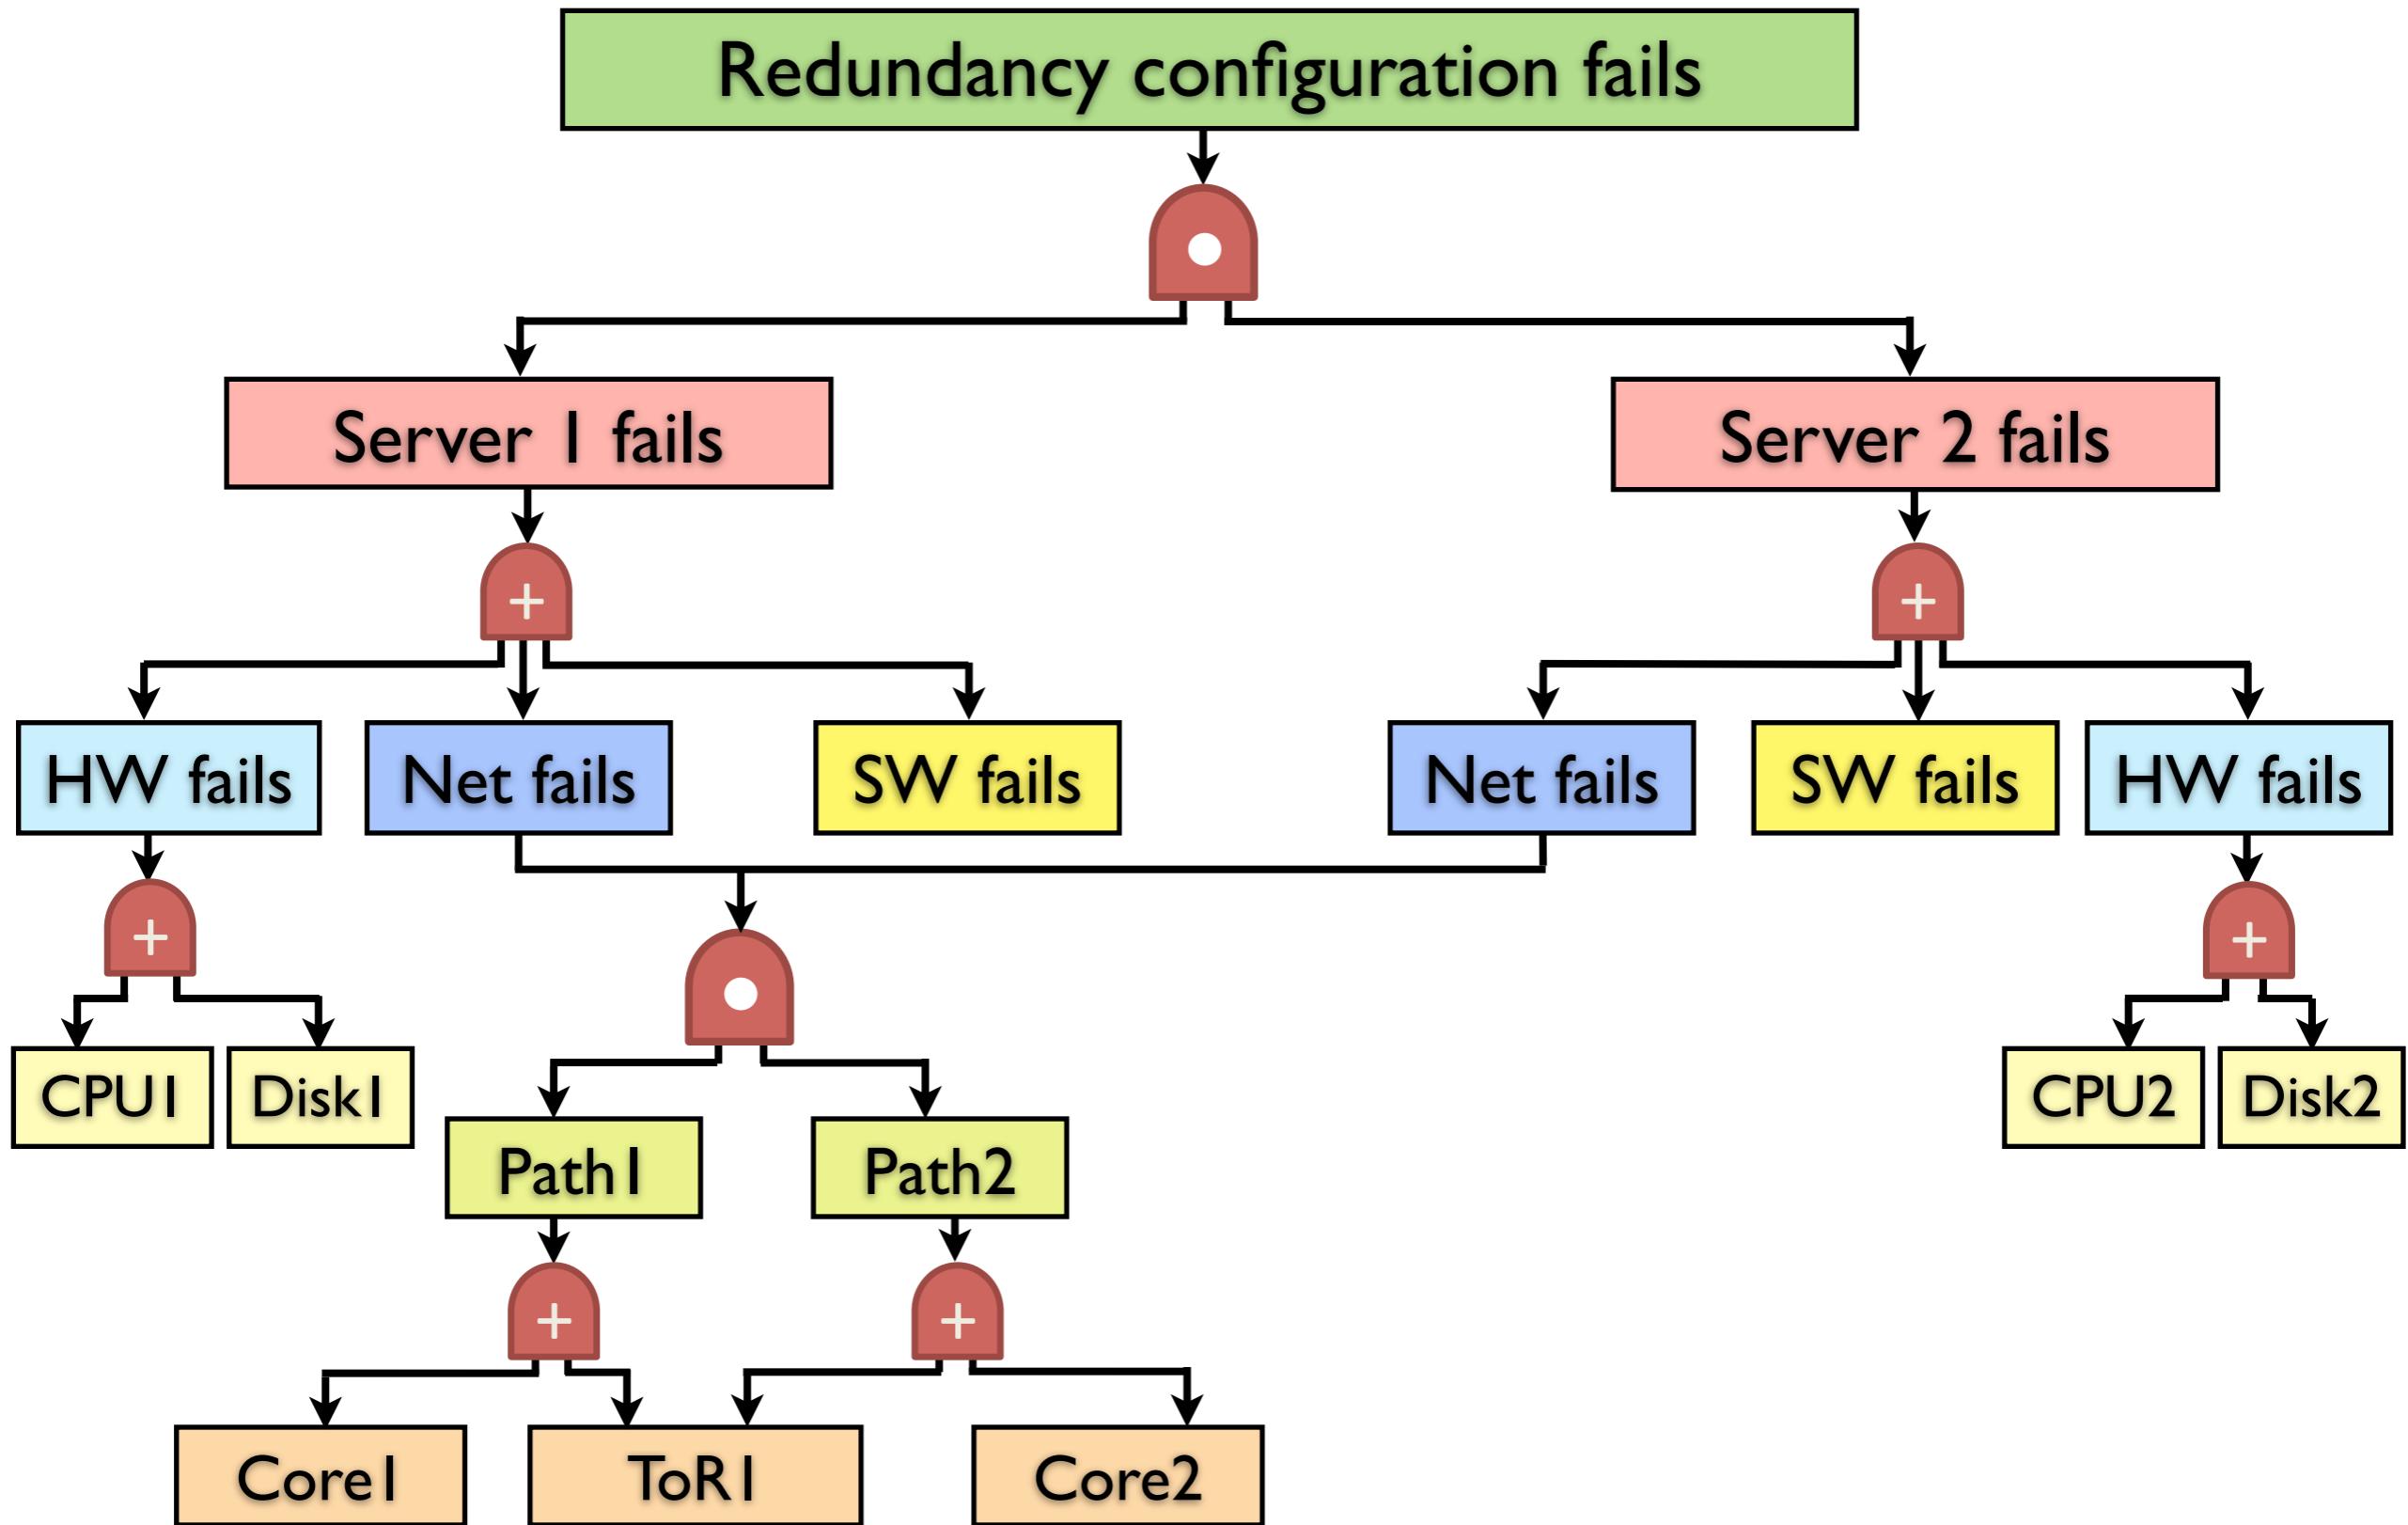
# Step3: Dependency Nodes



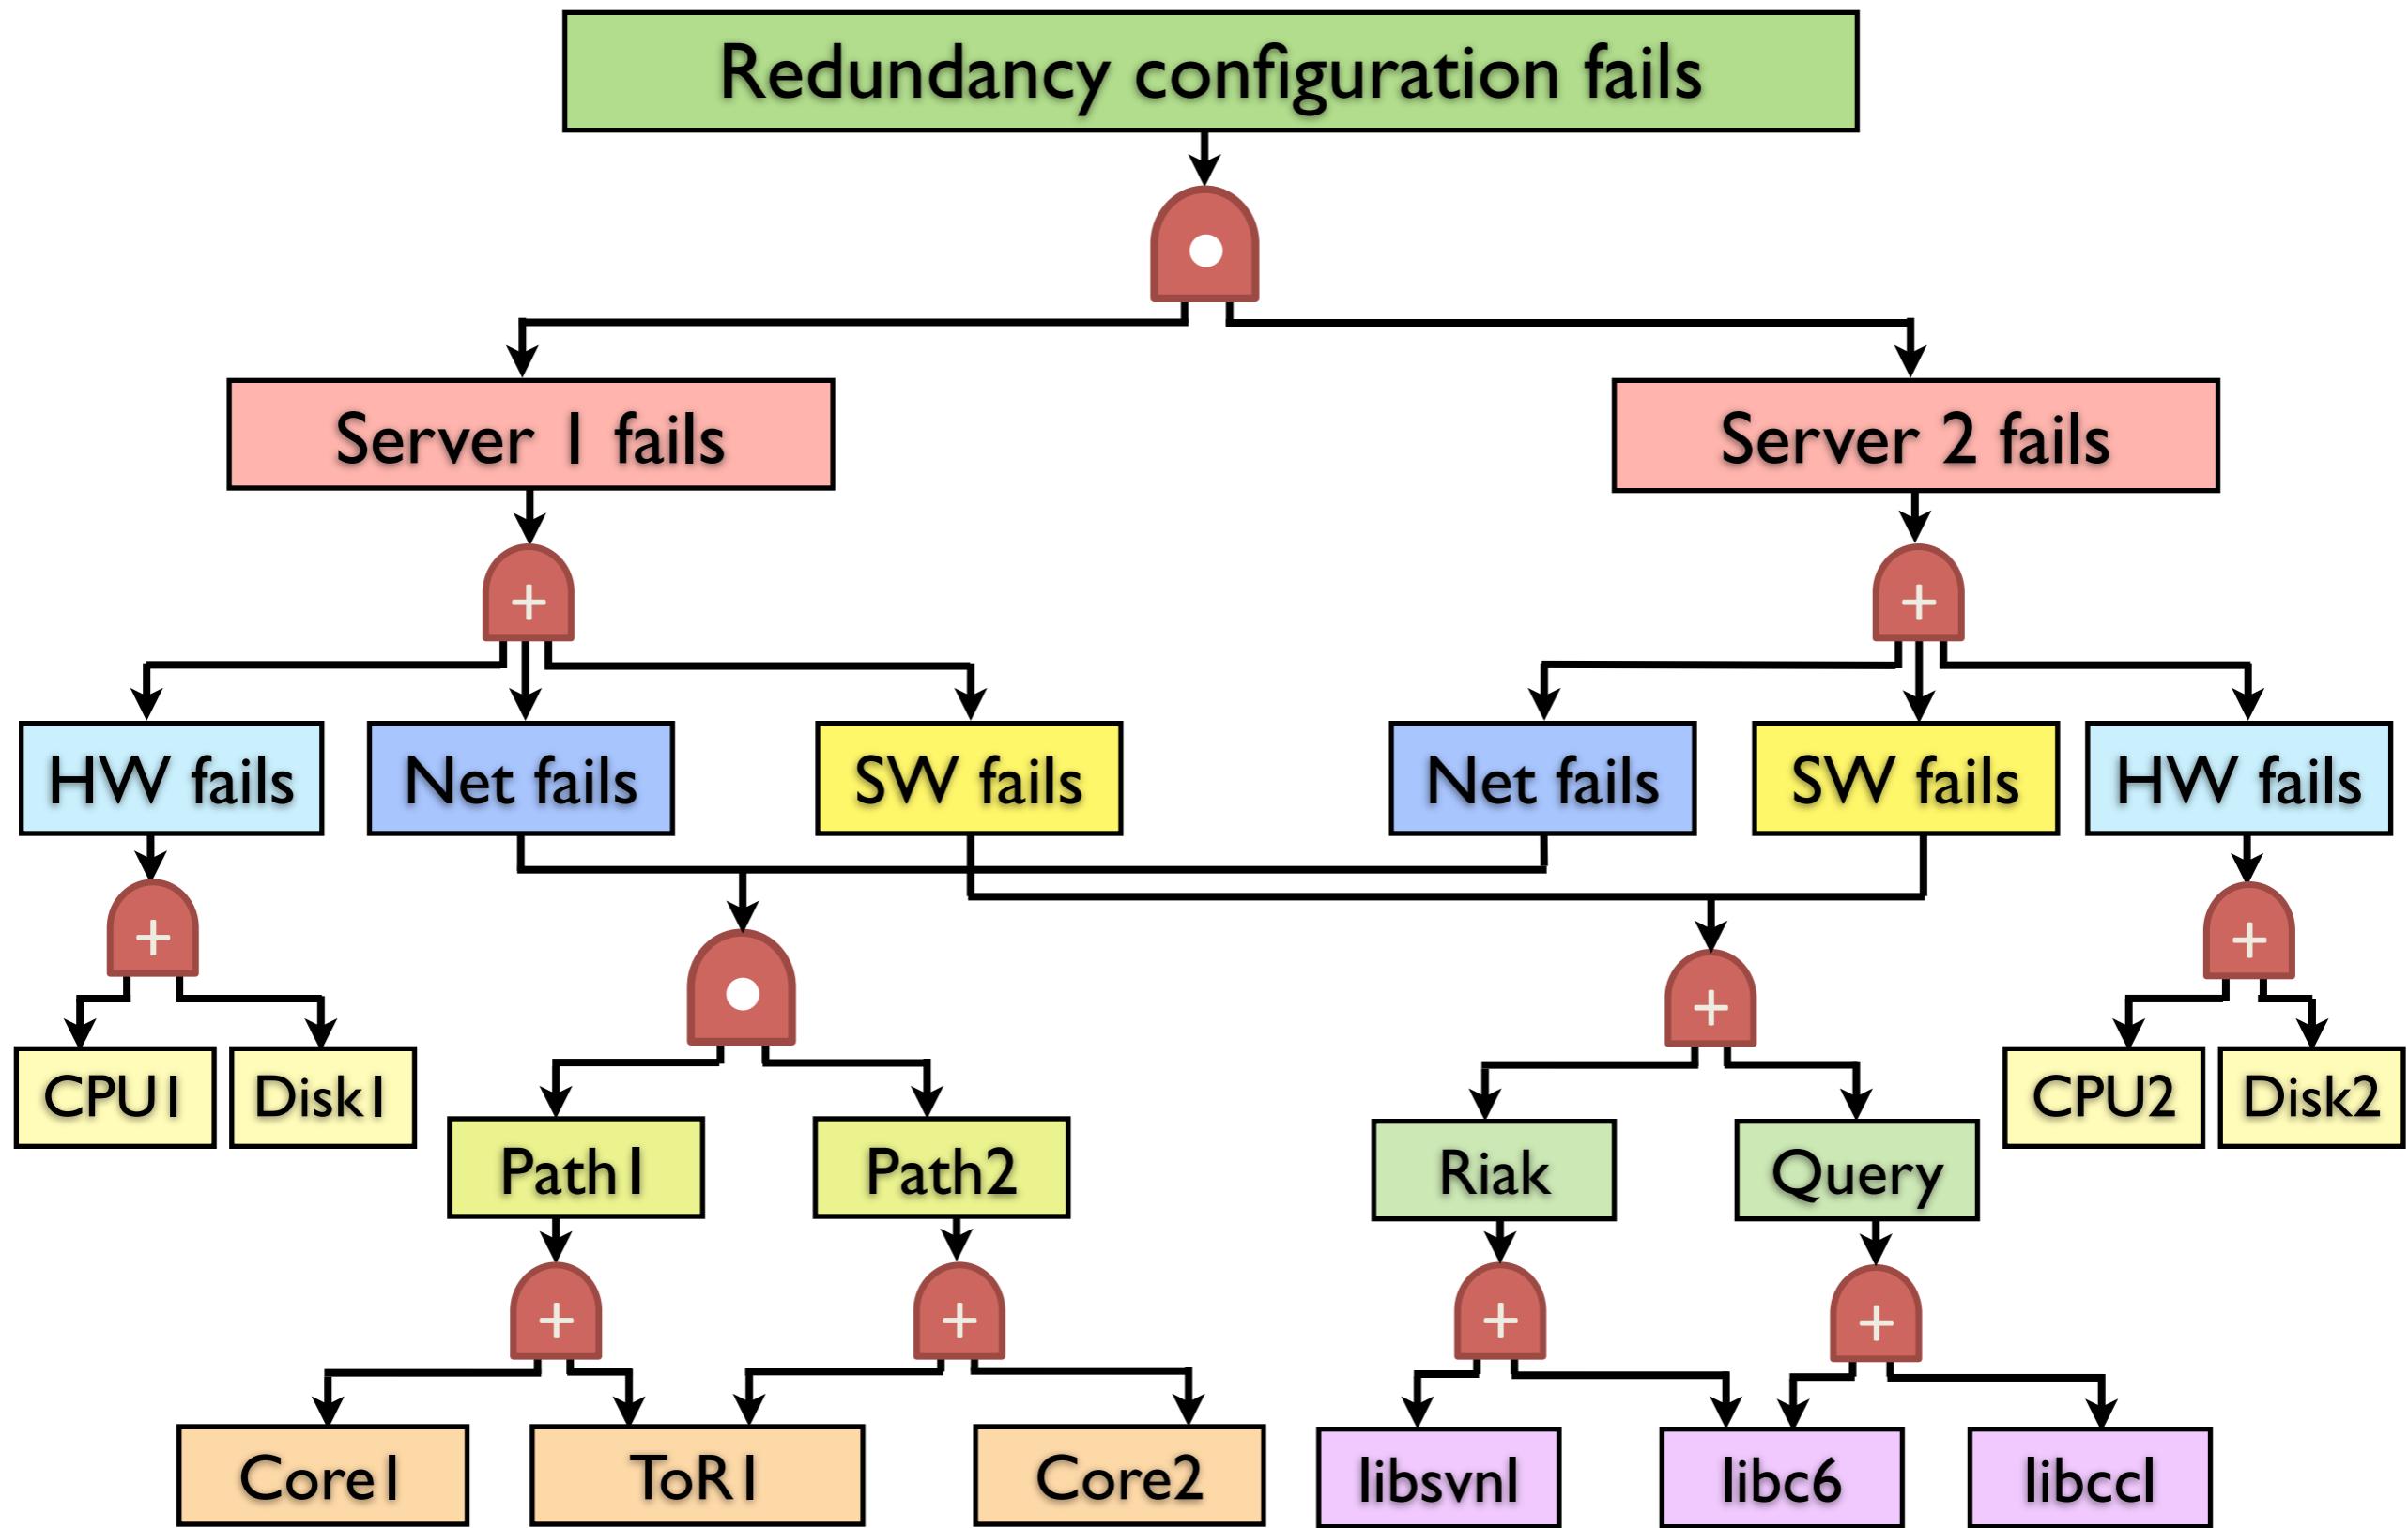
# Step4: Hardware Dependency



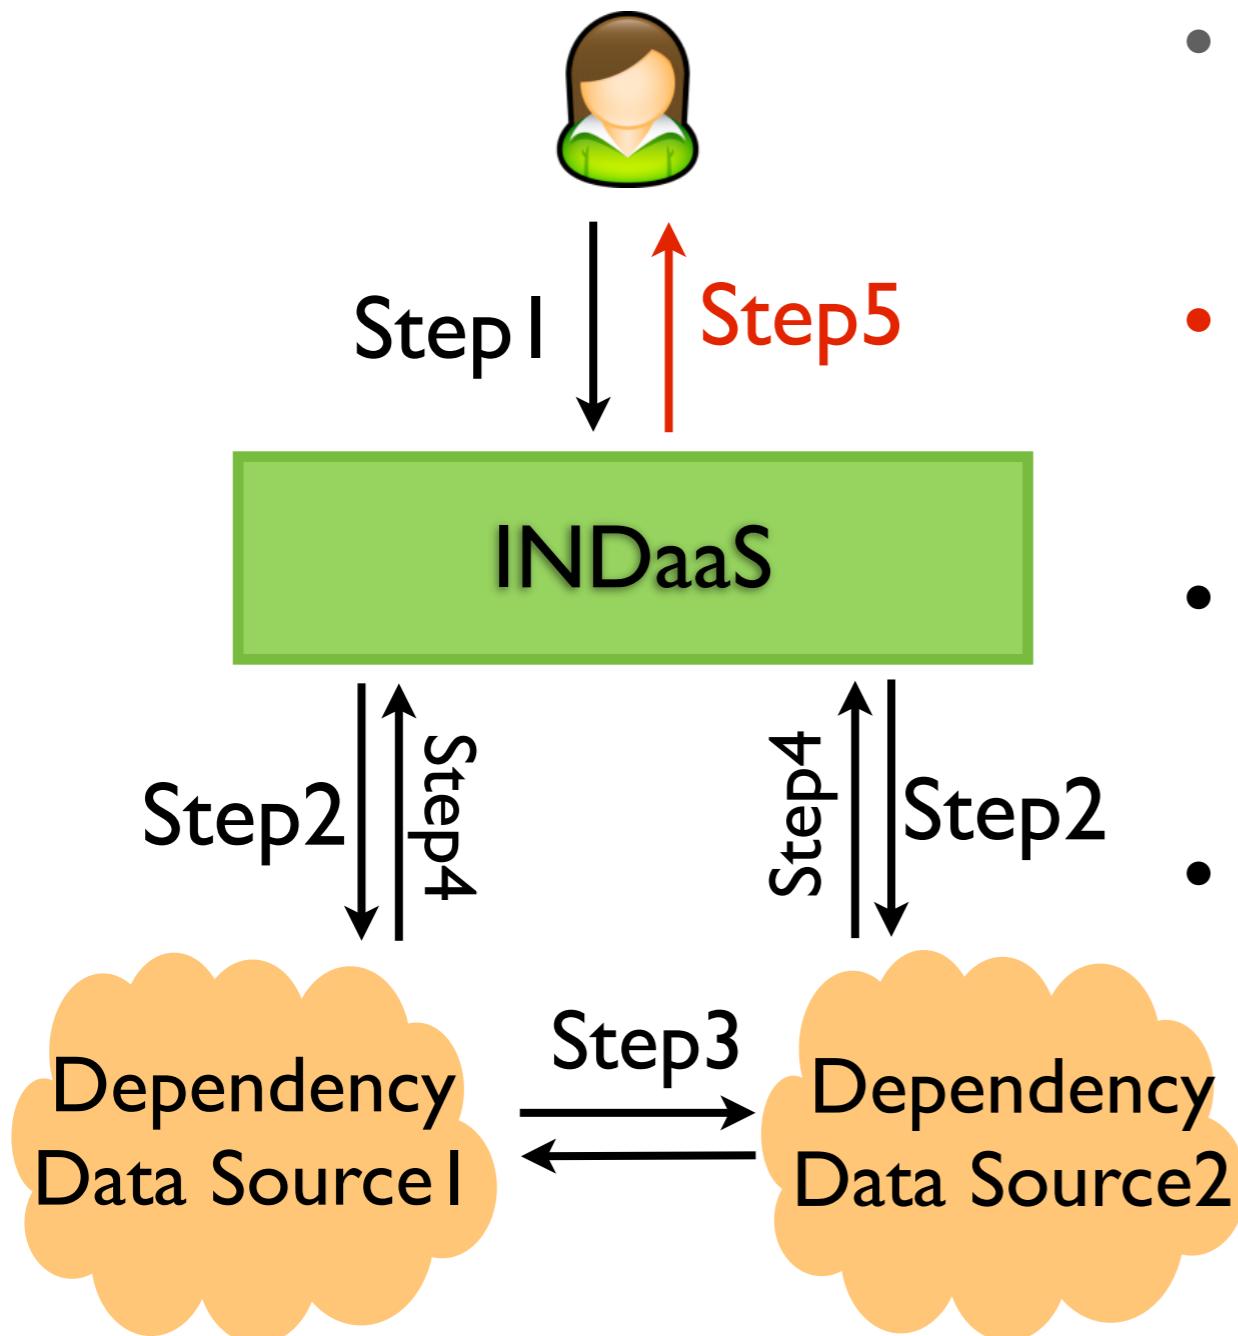
# Step5: Network Dependency



# Step6: Software Dependency

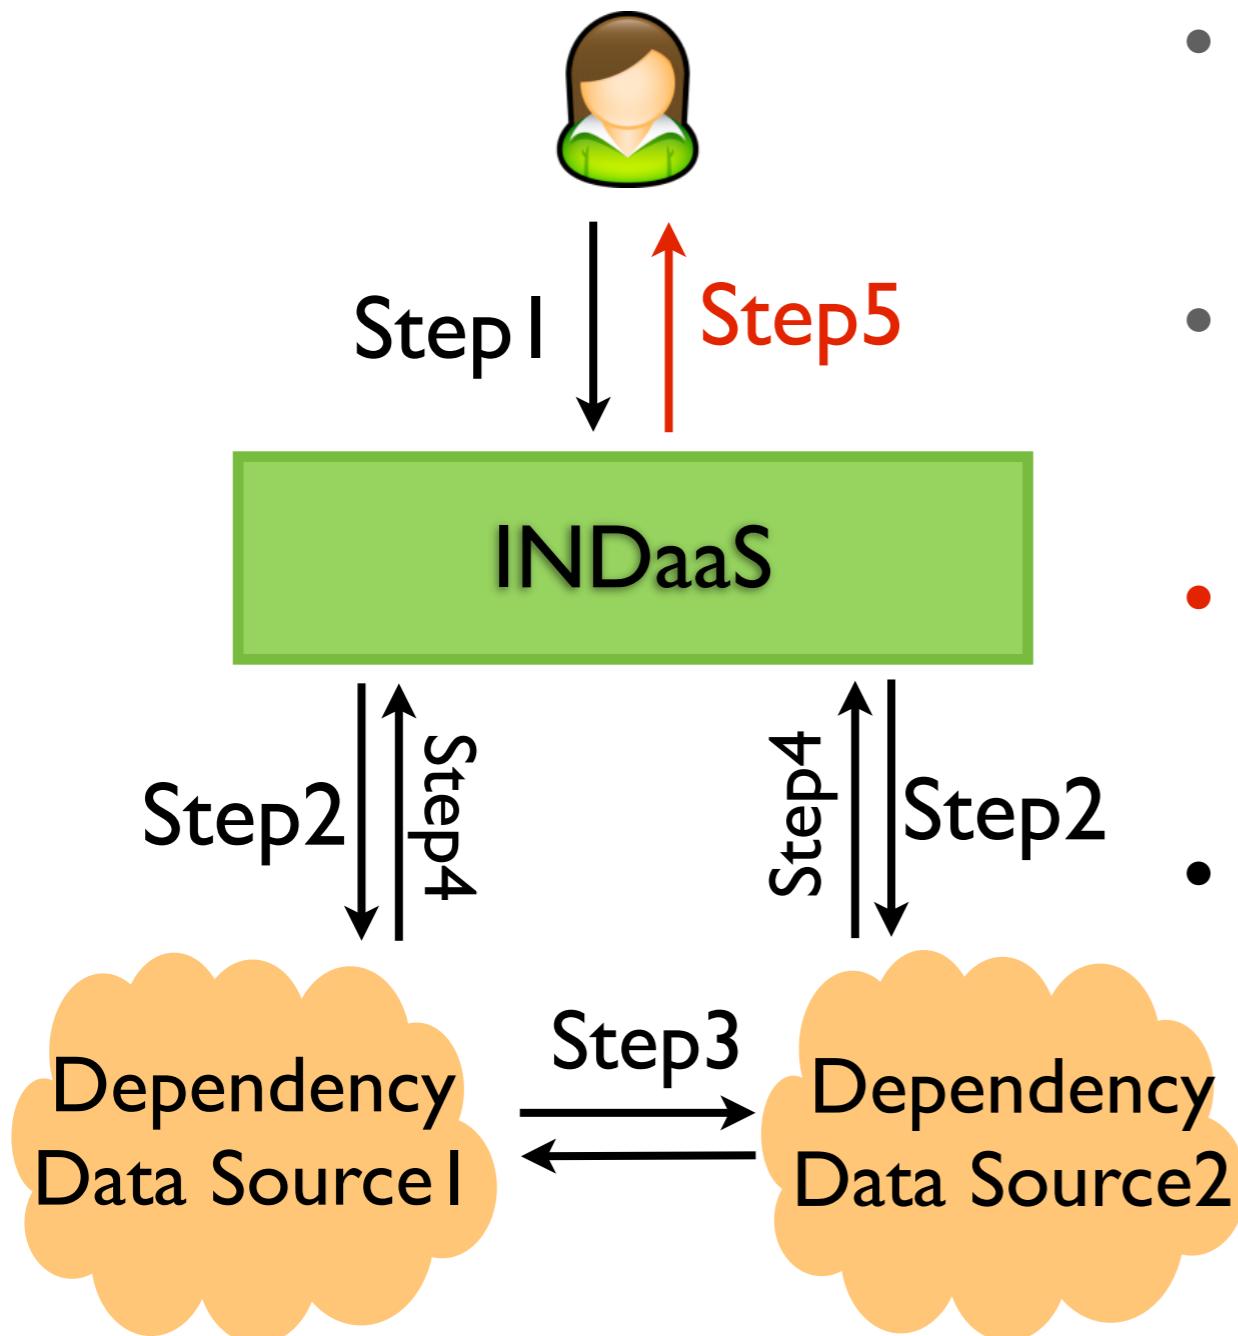


# RoadMap



- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

# RoadMap



- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

# Efficient Auditing

- Two algorithms balancing cost and accuracy:
  - Minimal fault set algorithm
  - Failure sampling algorithm

# Efficient Auditing

- Two algorithms balancing cost and accuracy:
  - Minimal fault set algorithm
  - Failure sampling algorithm

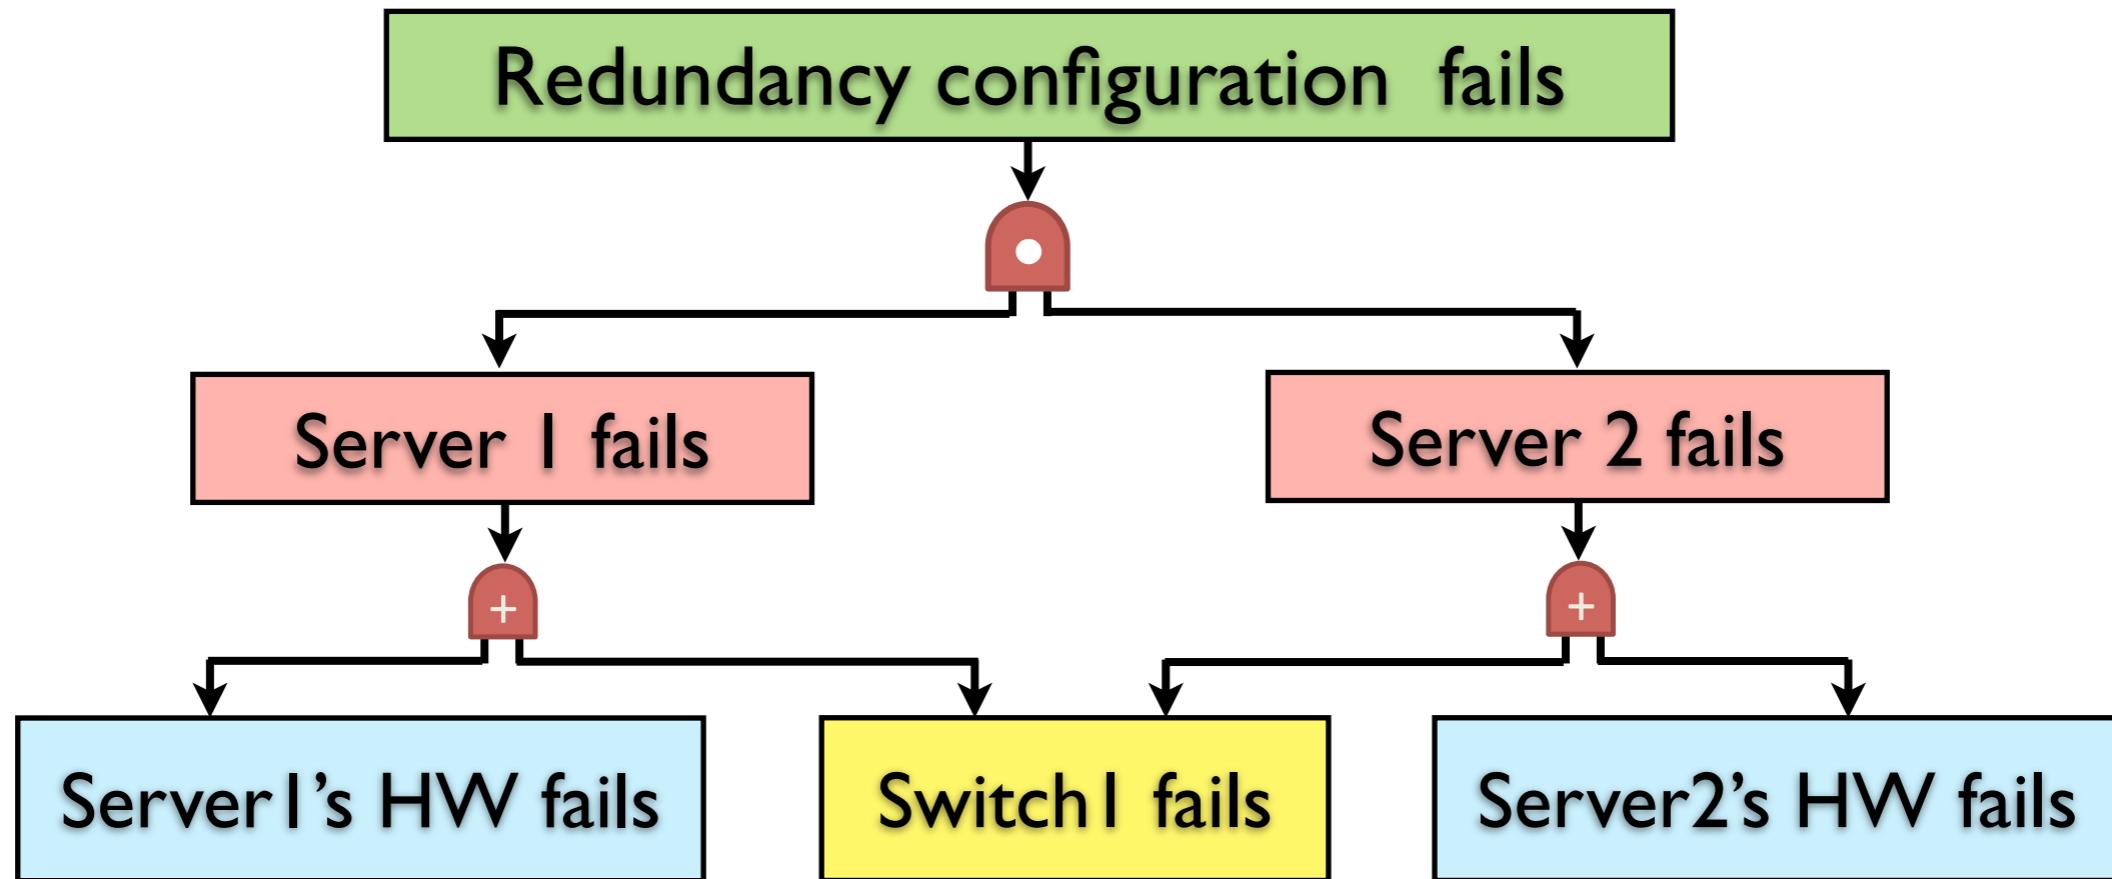
# Minimal Fault Set Algorithm

- Traditional algorithm in safety engineering
  - Exponential complexity (NP-hard)
- We are the first to apply it in Cloud area:
  - Analyzing a fat tree with 30,528 with ~40 hours

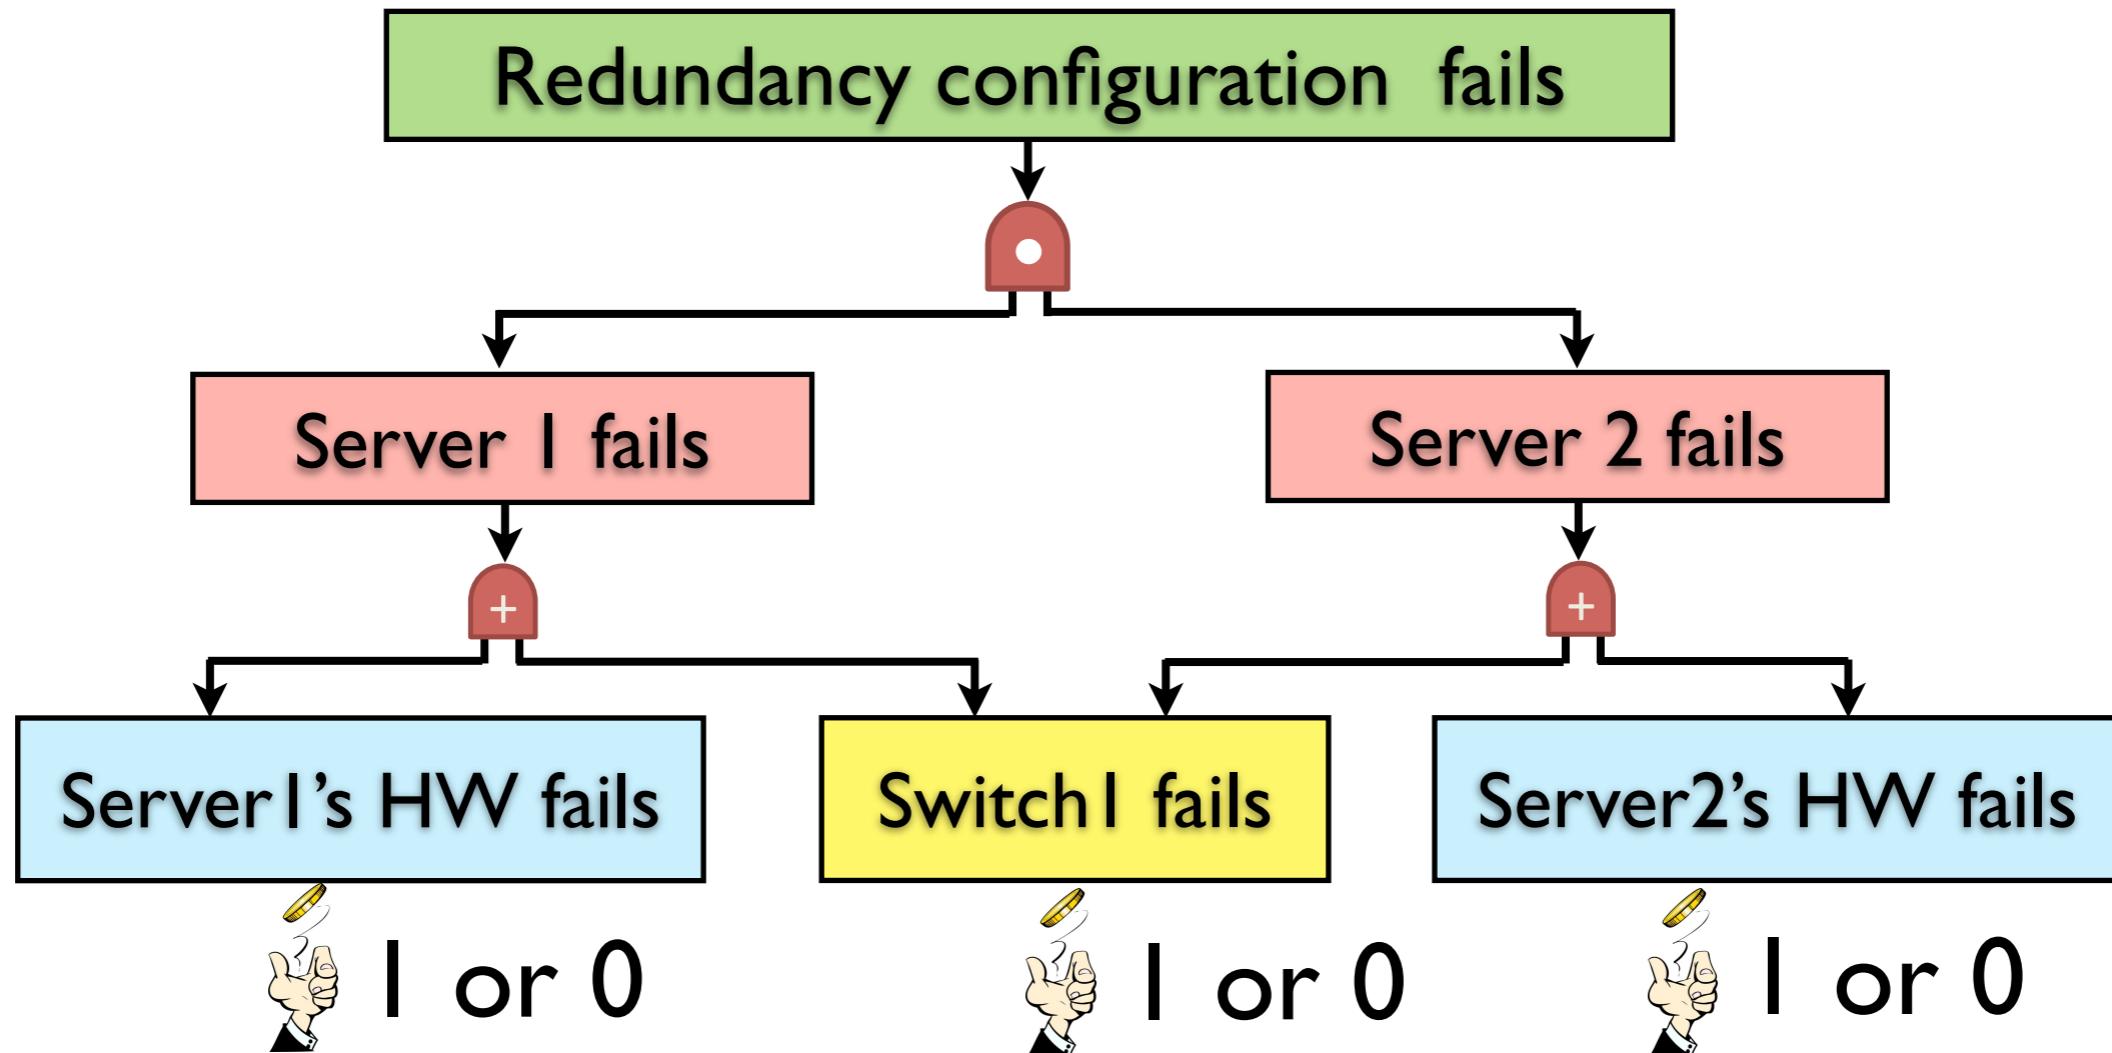
# Minimal Fault Set Algorithm

- Traditional algorithm in safety engineer
  - Exponential complexity (NP-hard)
- We are the first to apply it in Cloud area:
  - Analyzing a fat tree with 30,528 with ~40 hours
- We propose efficient failure sampling algorithm.

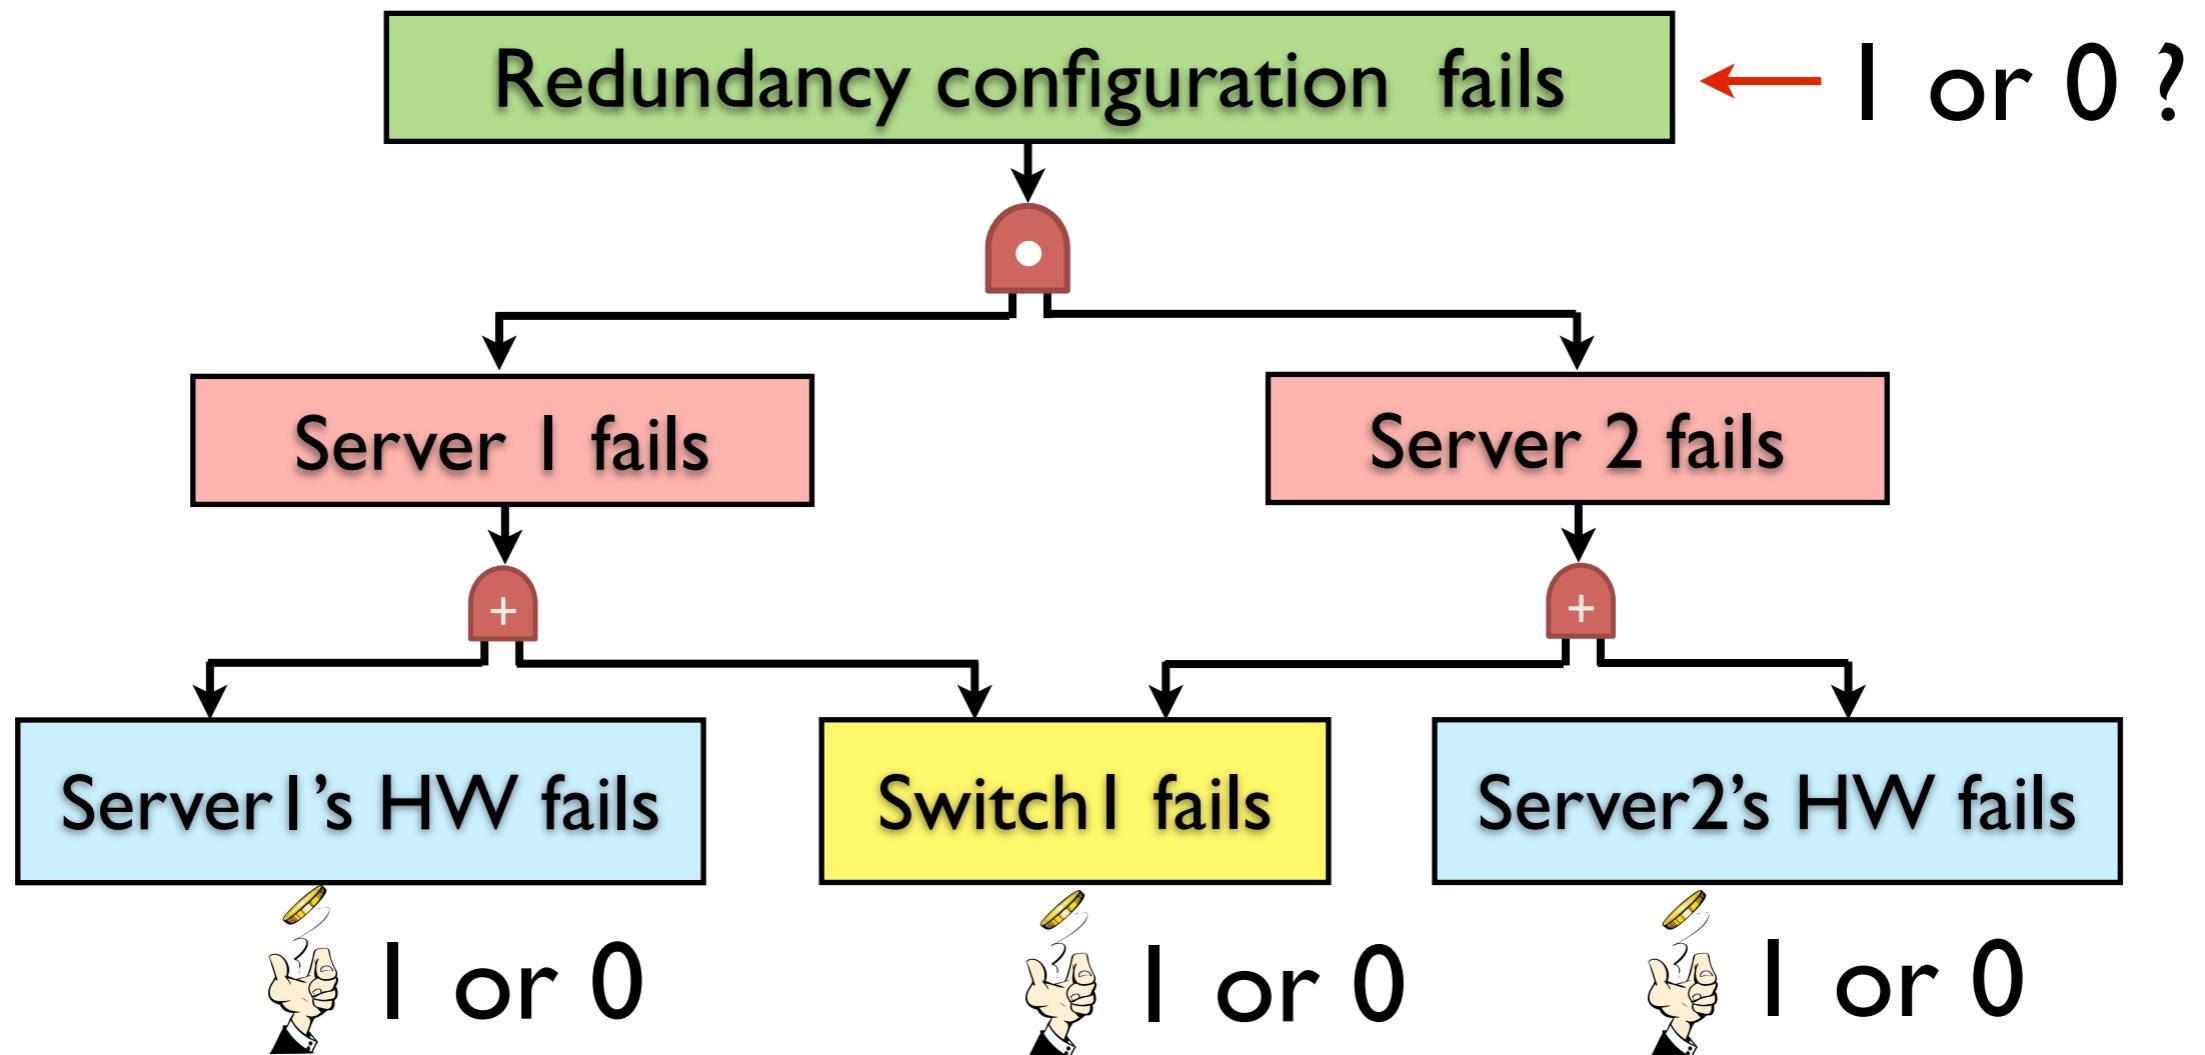
# Failure Sampling Algorithm



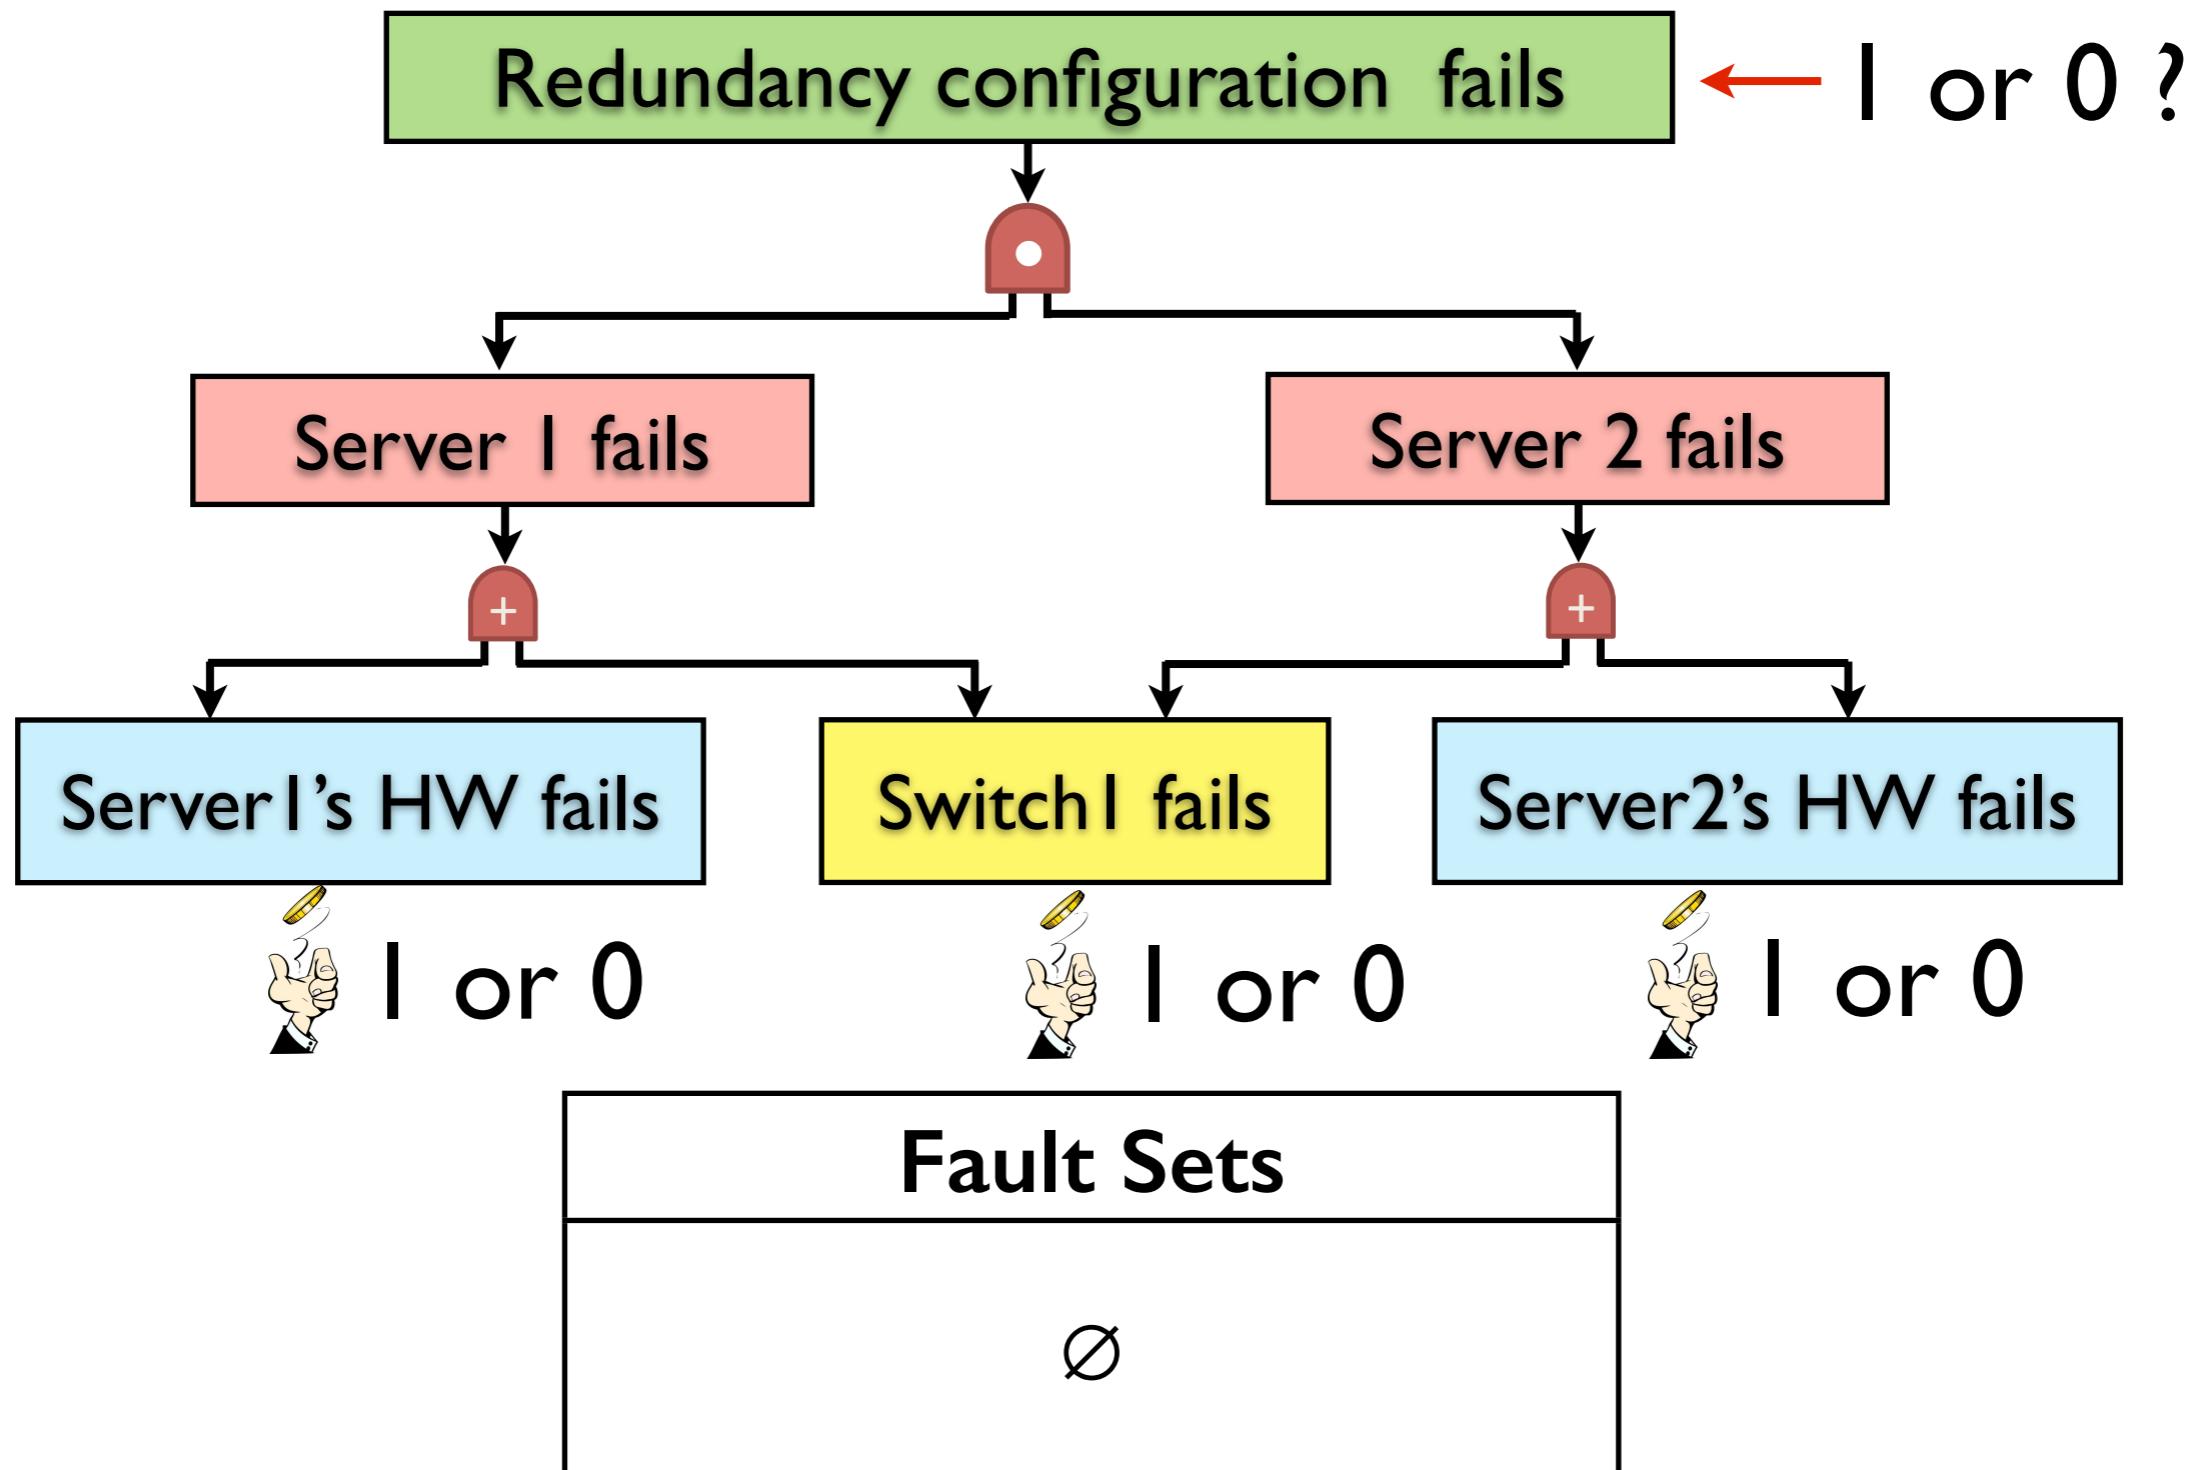
# Failure Sampling Algorithm



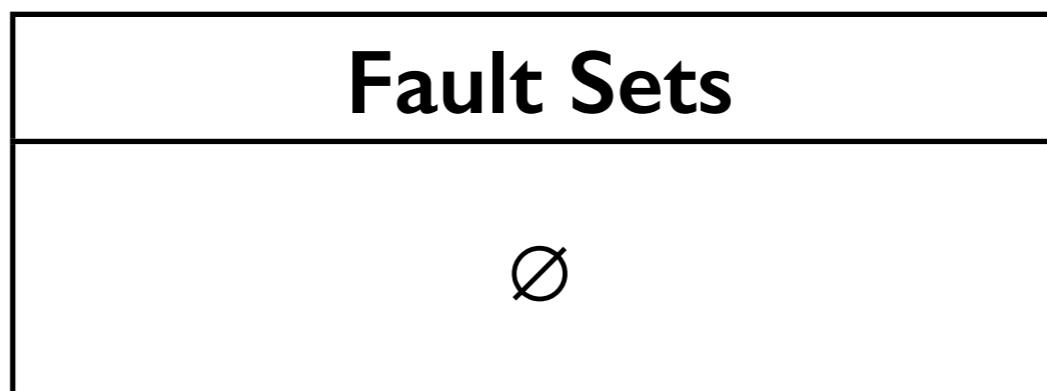
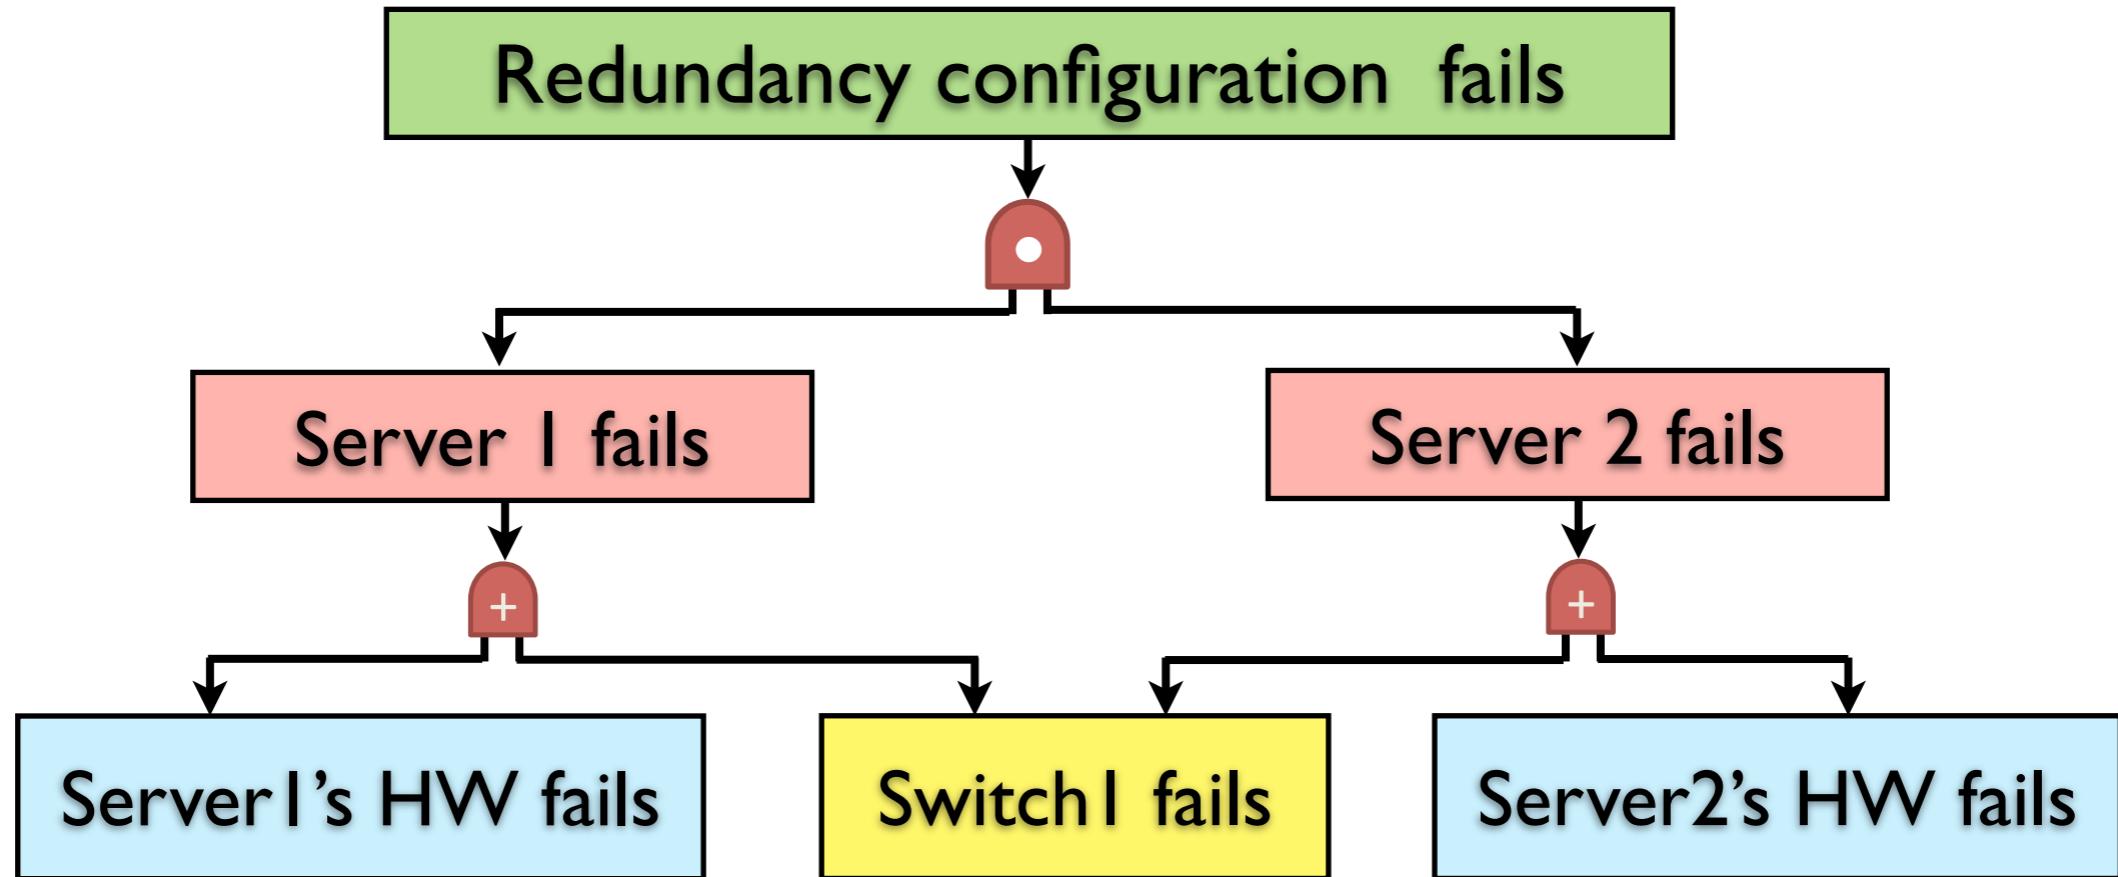
# Failure Sampling Algorithm



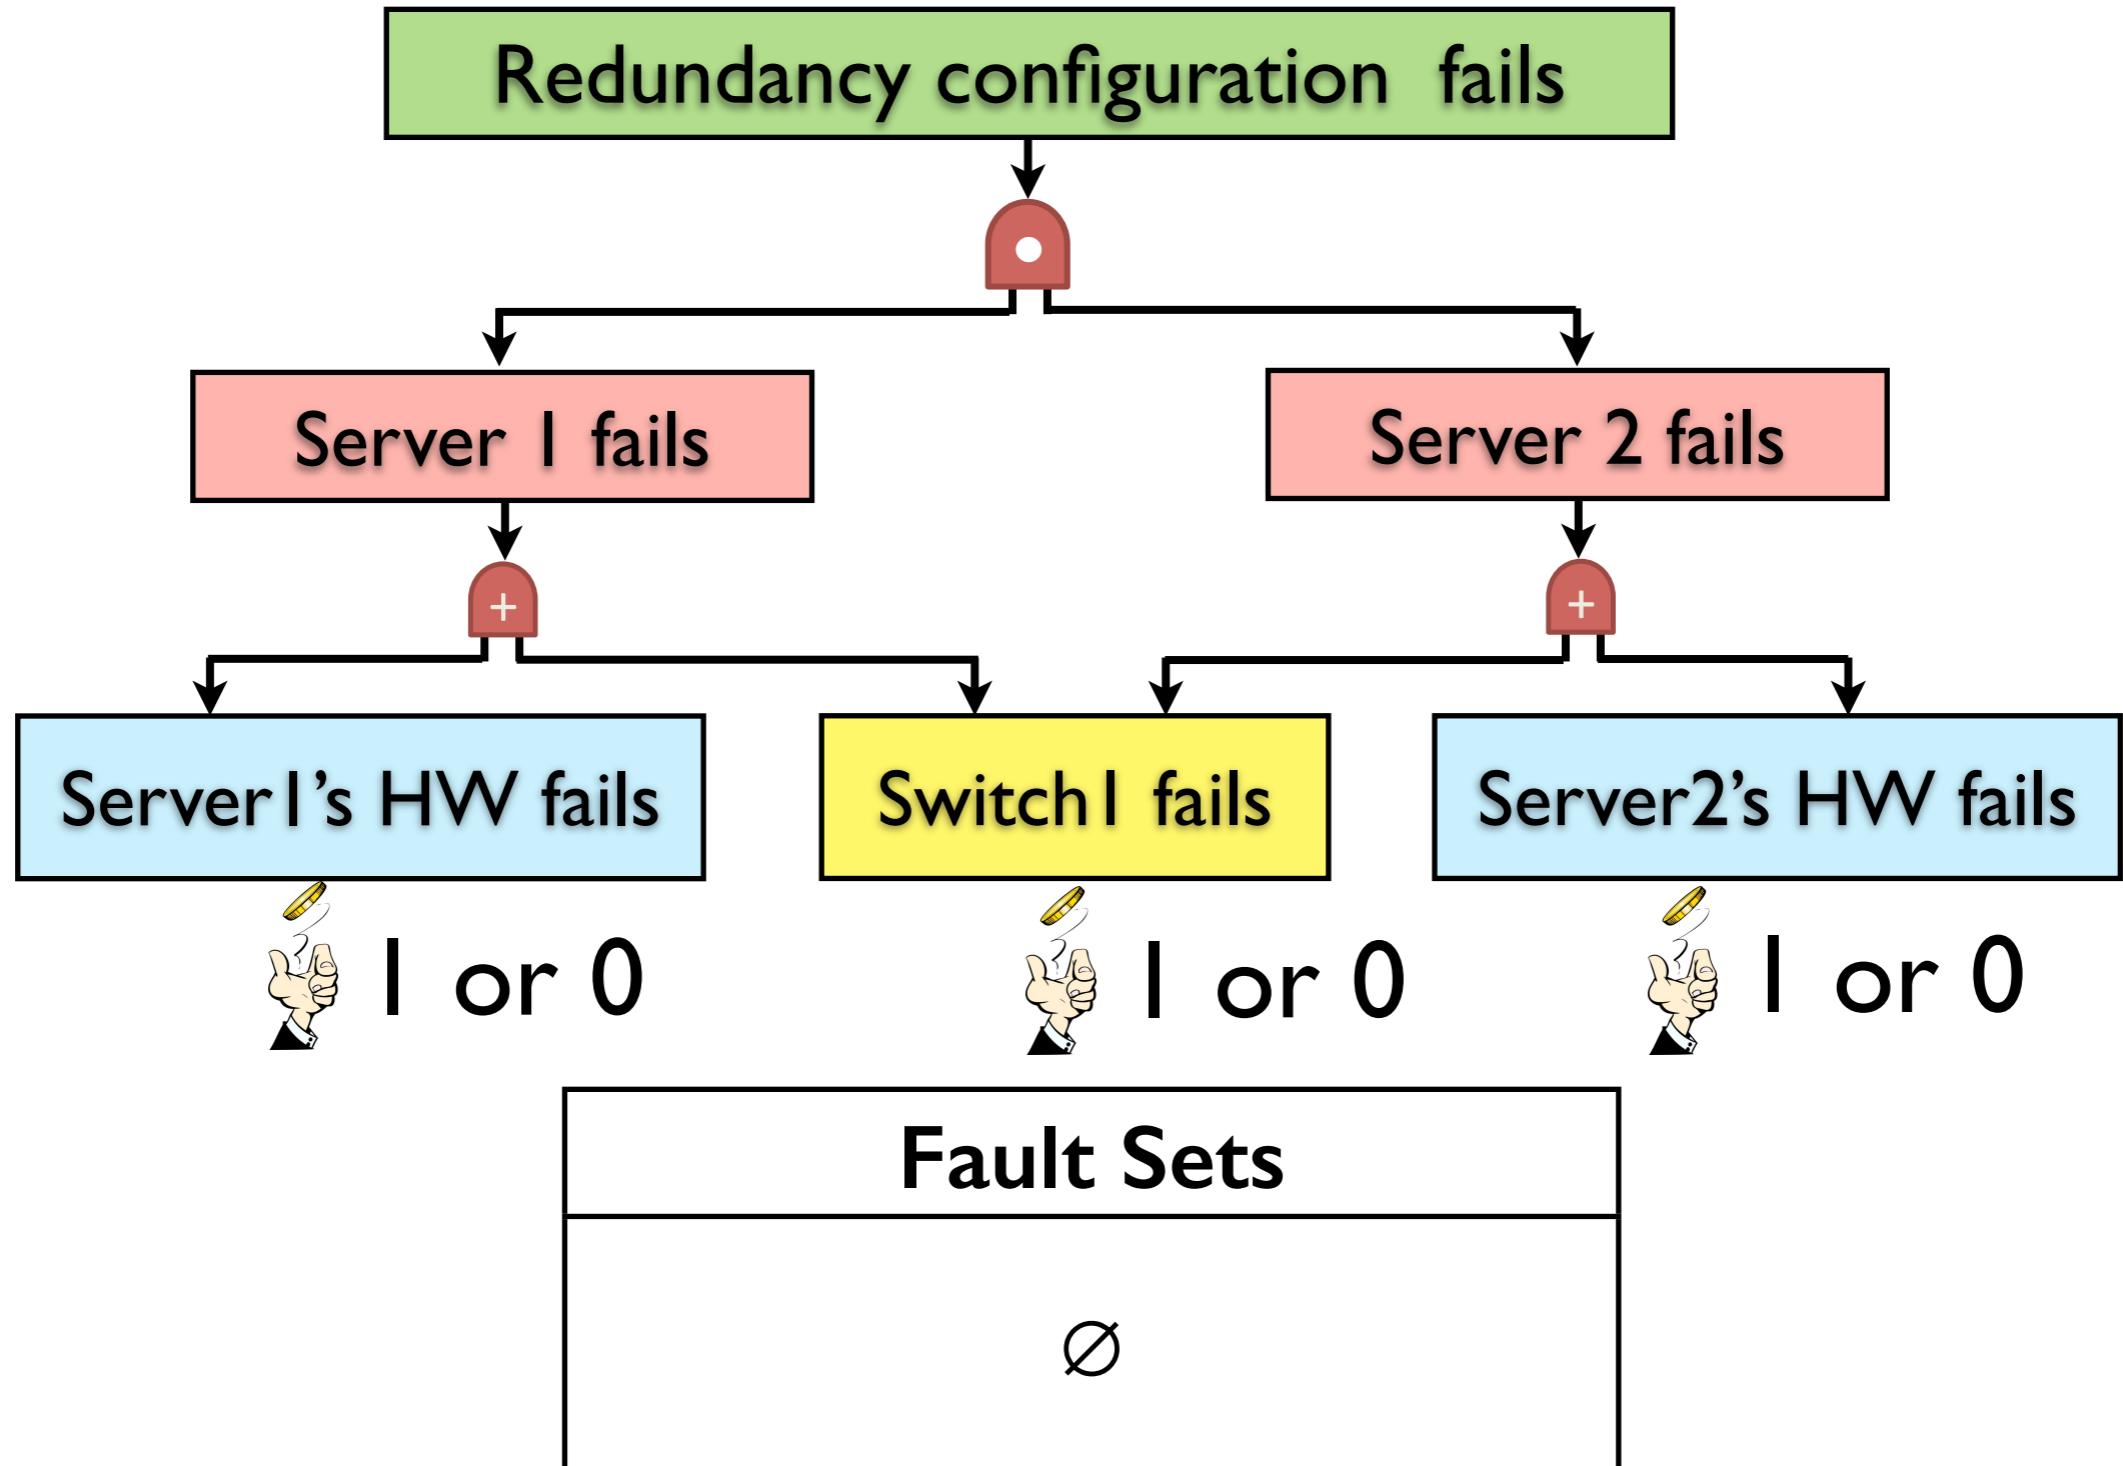
# Failure Sampling Algorithm



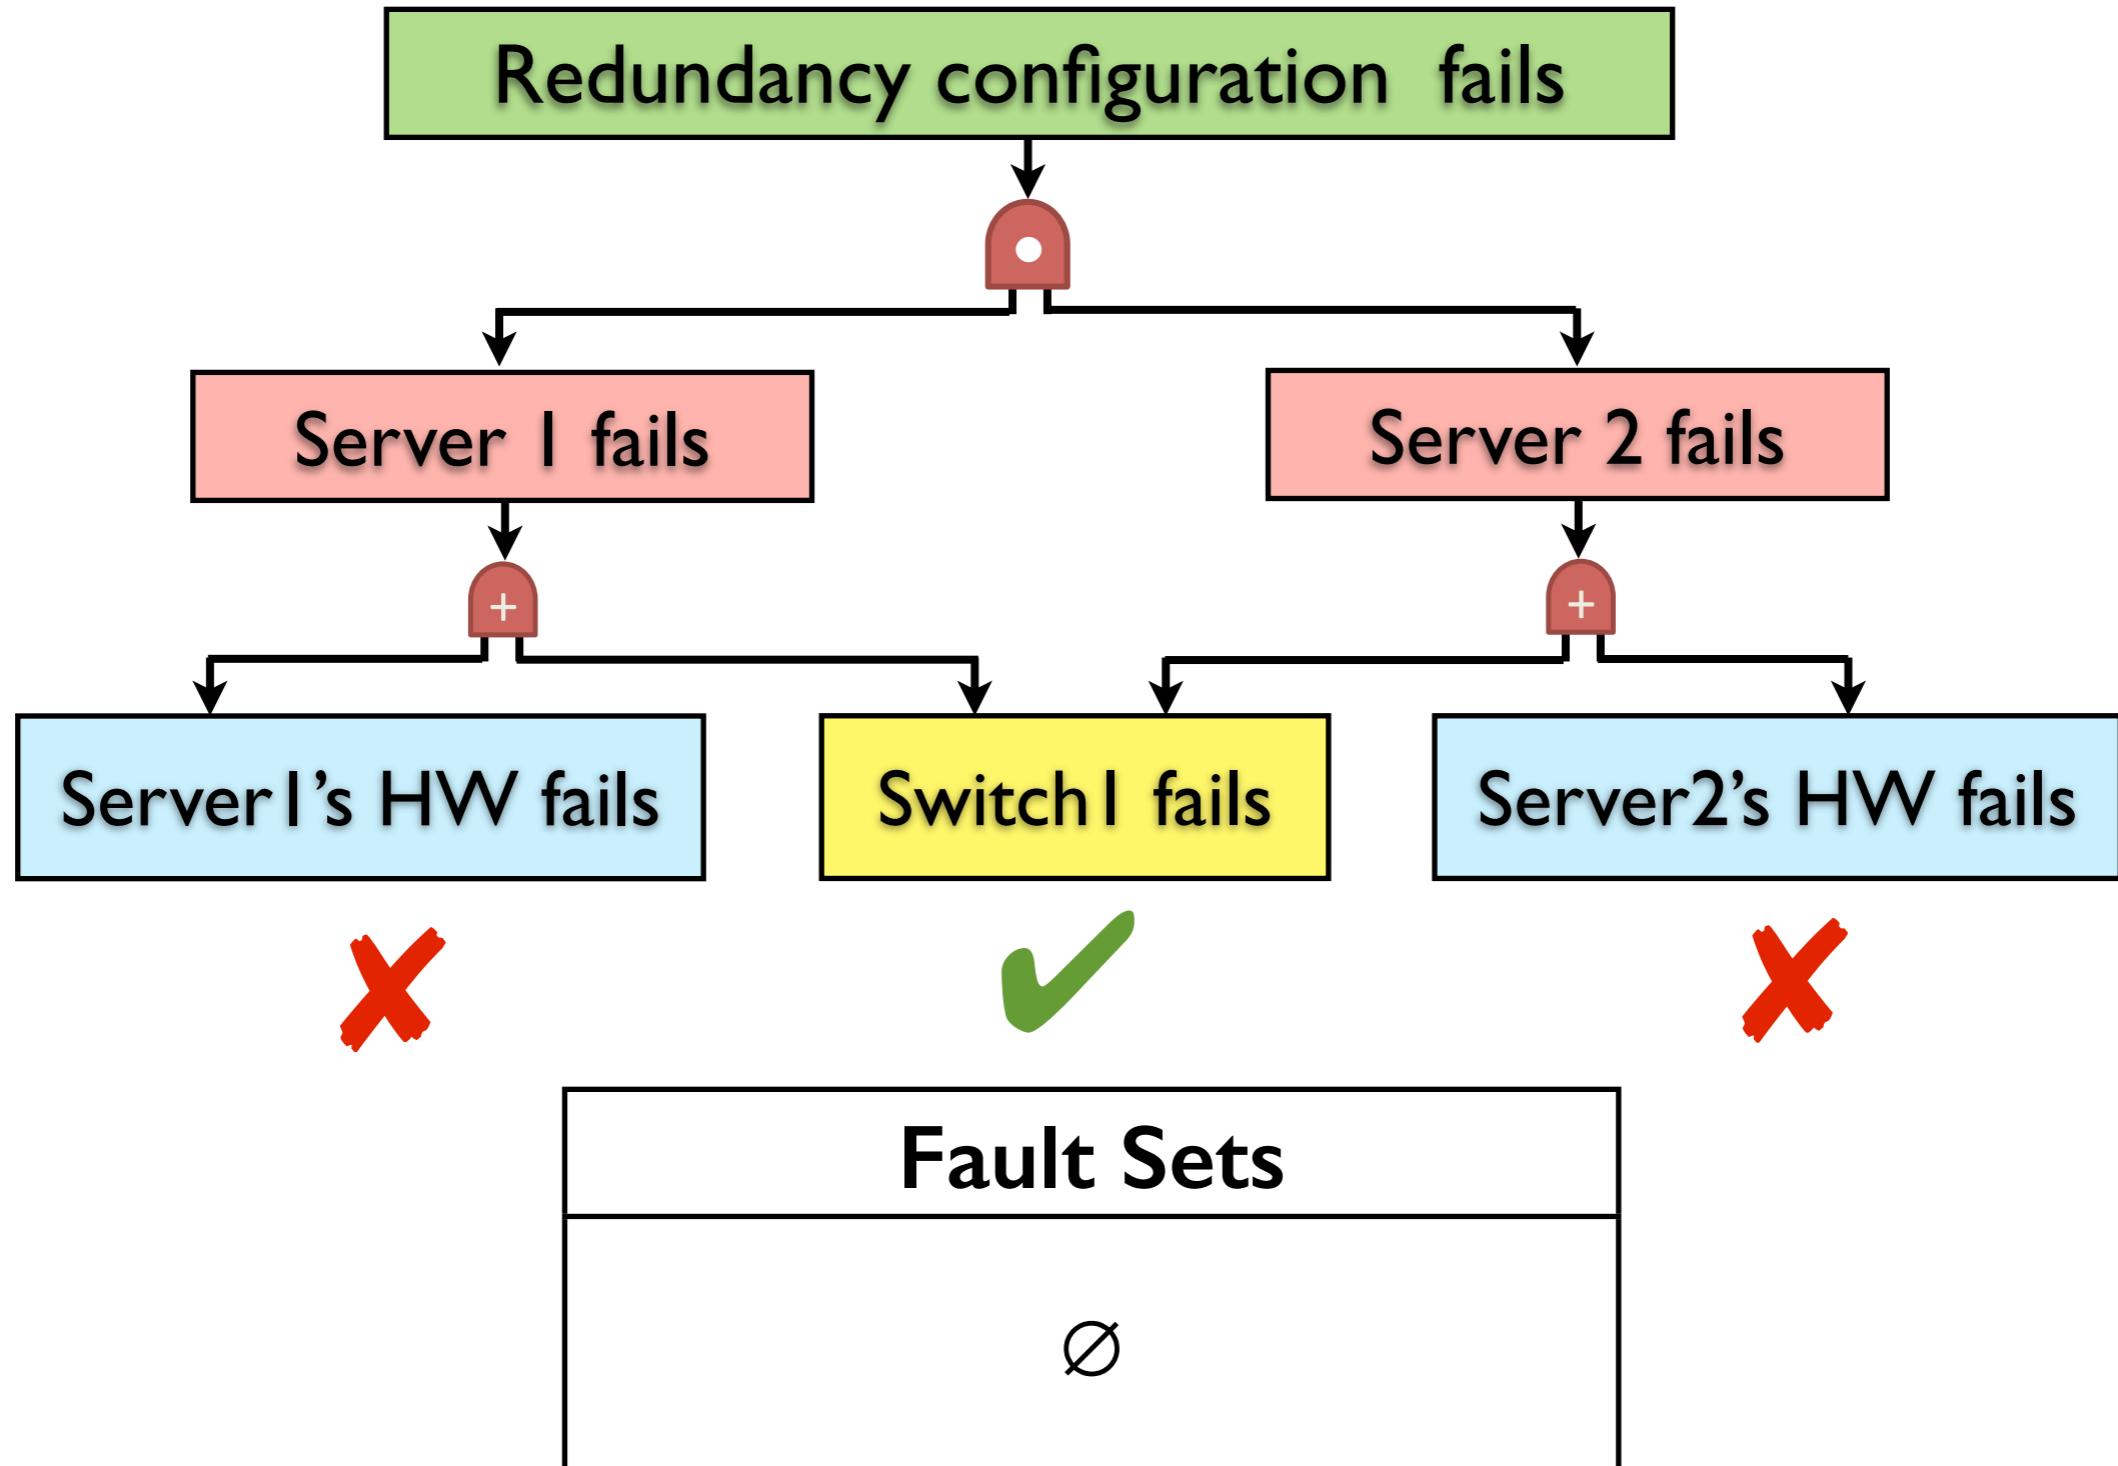
# The 1st Sampling Round



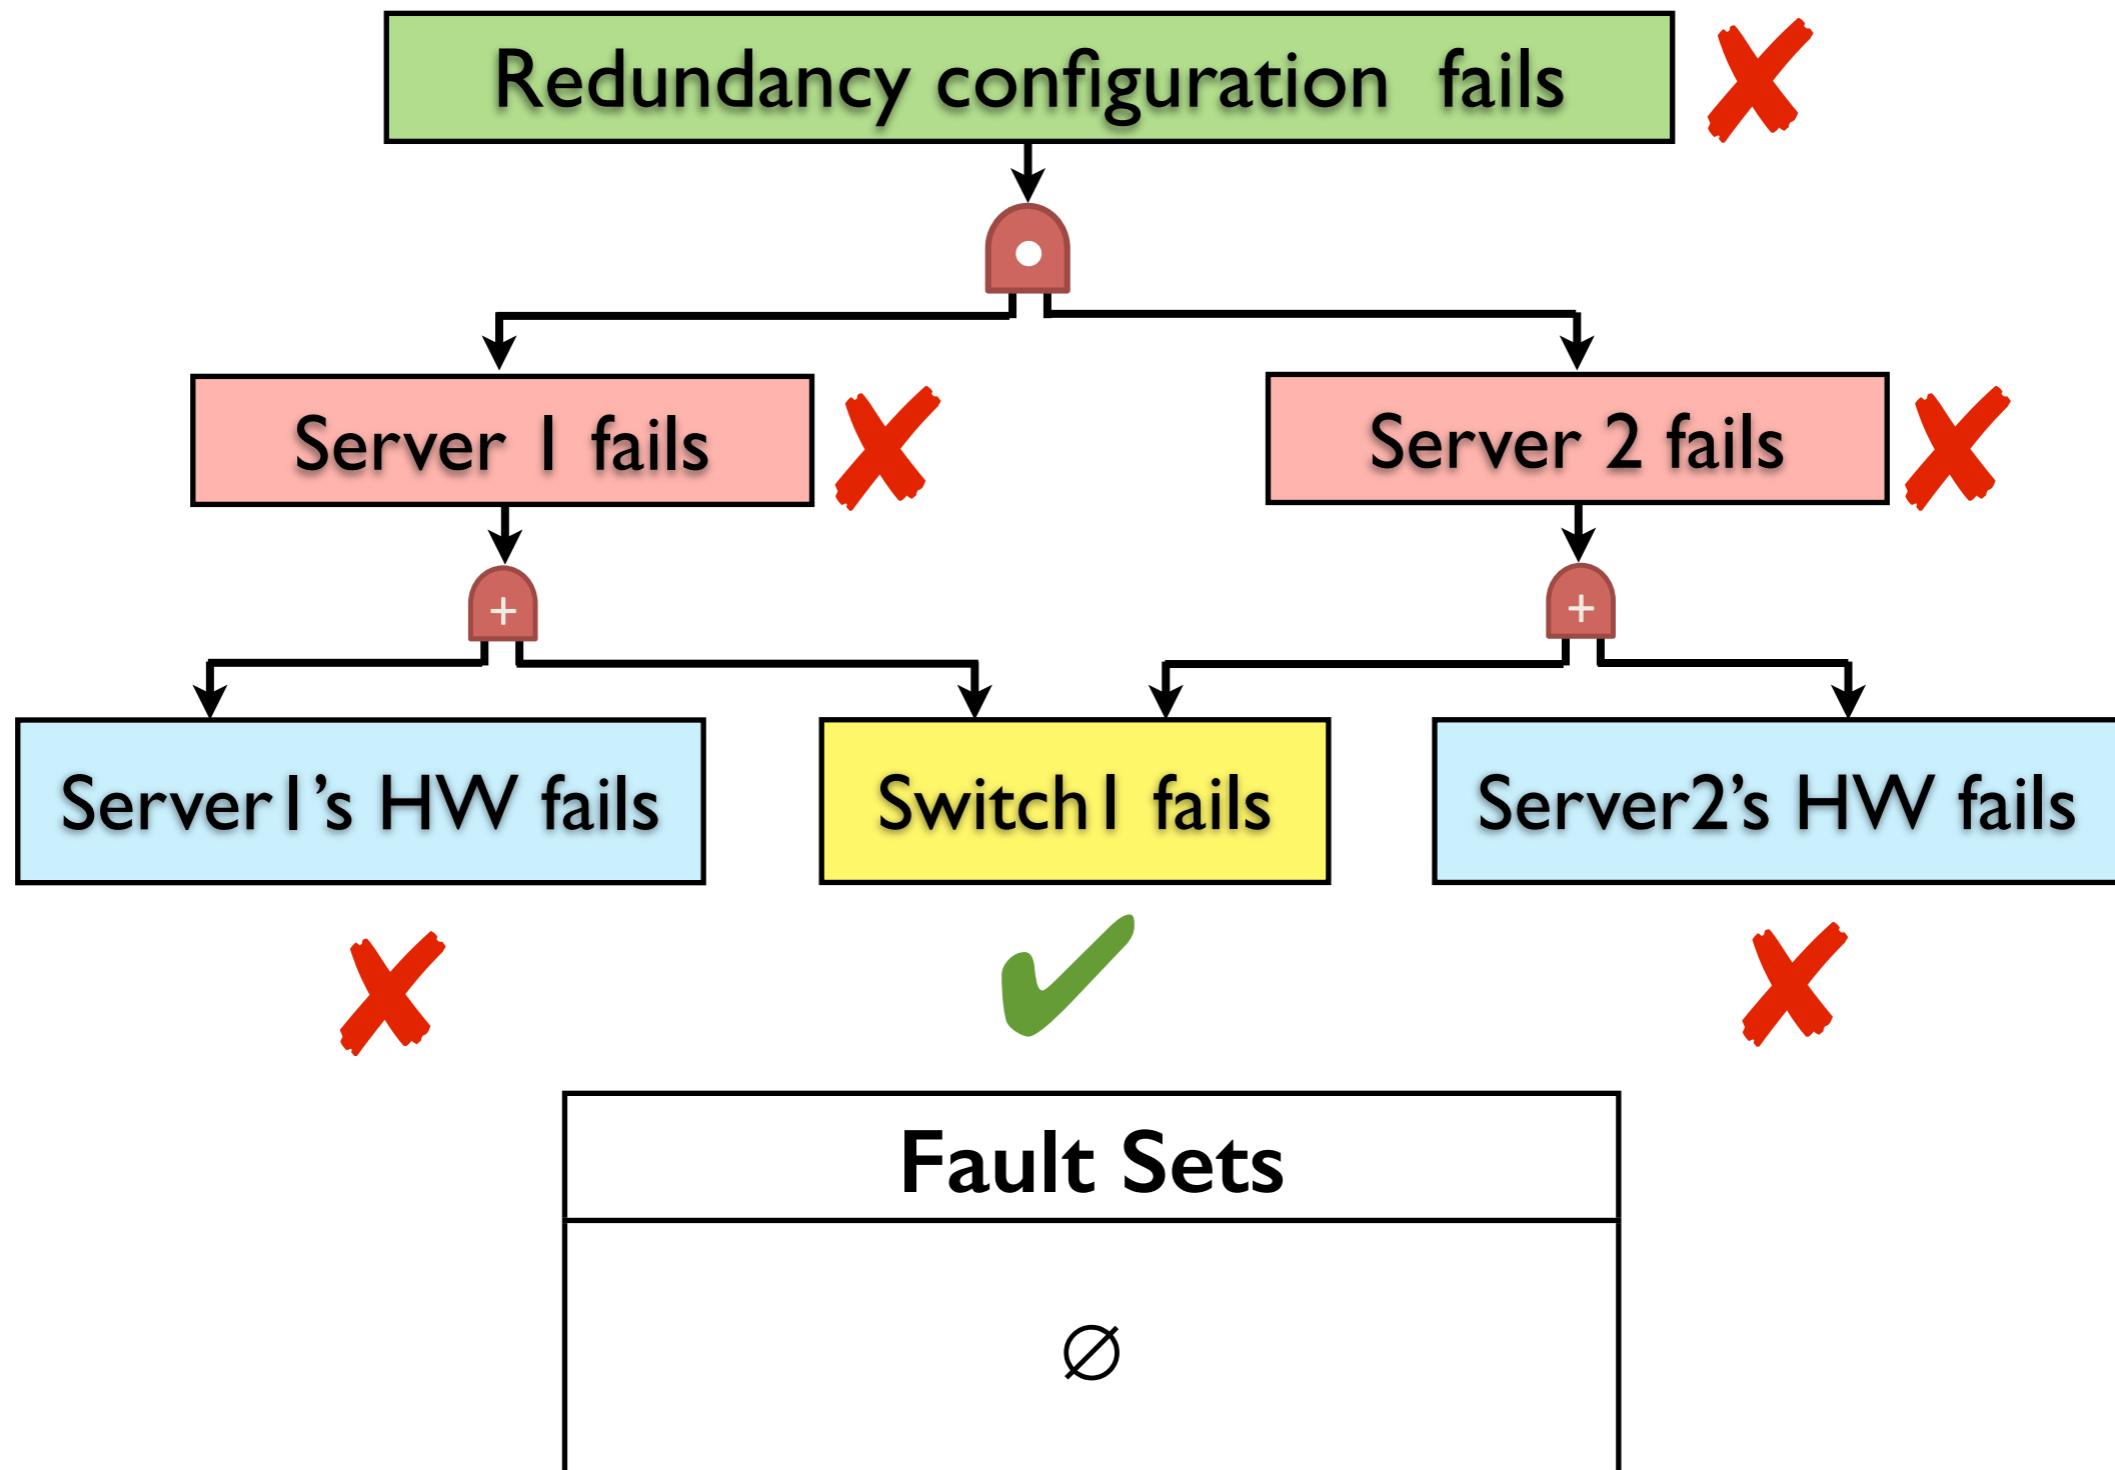
# The 1st Sampling Round



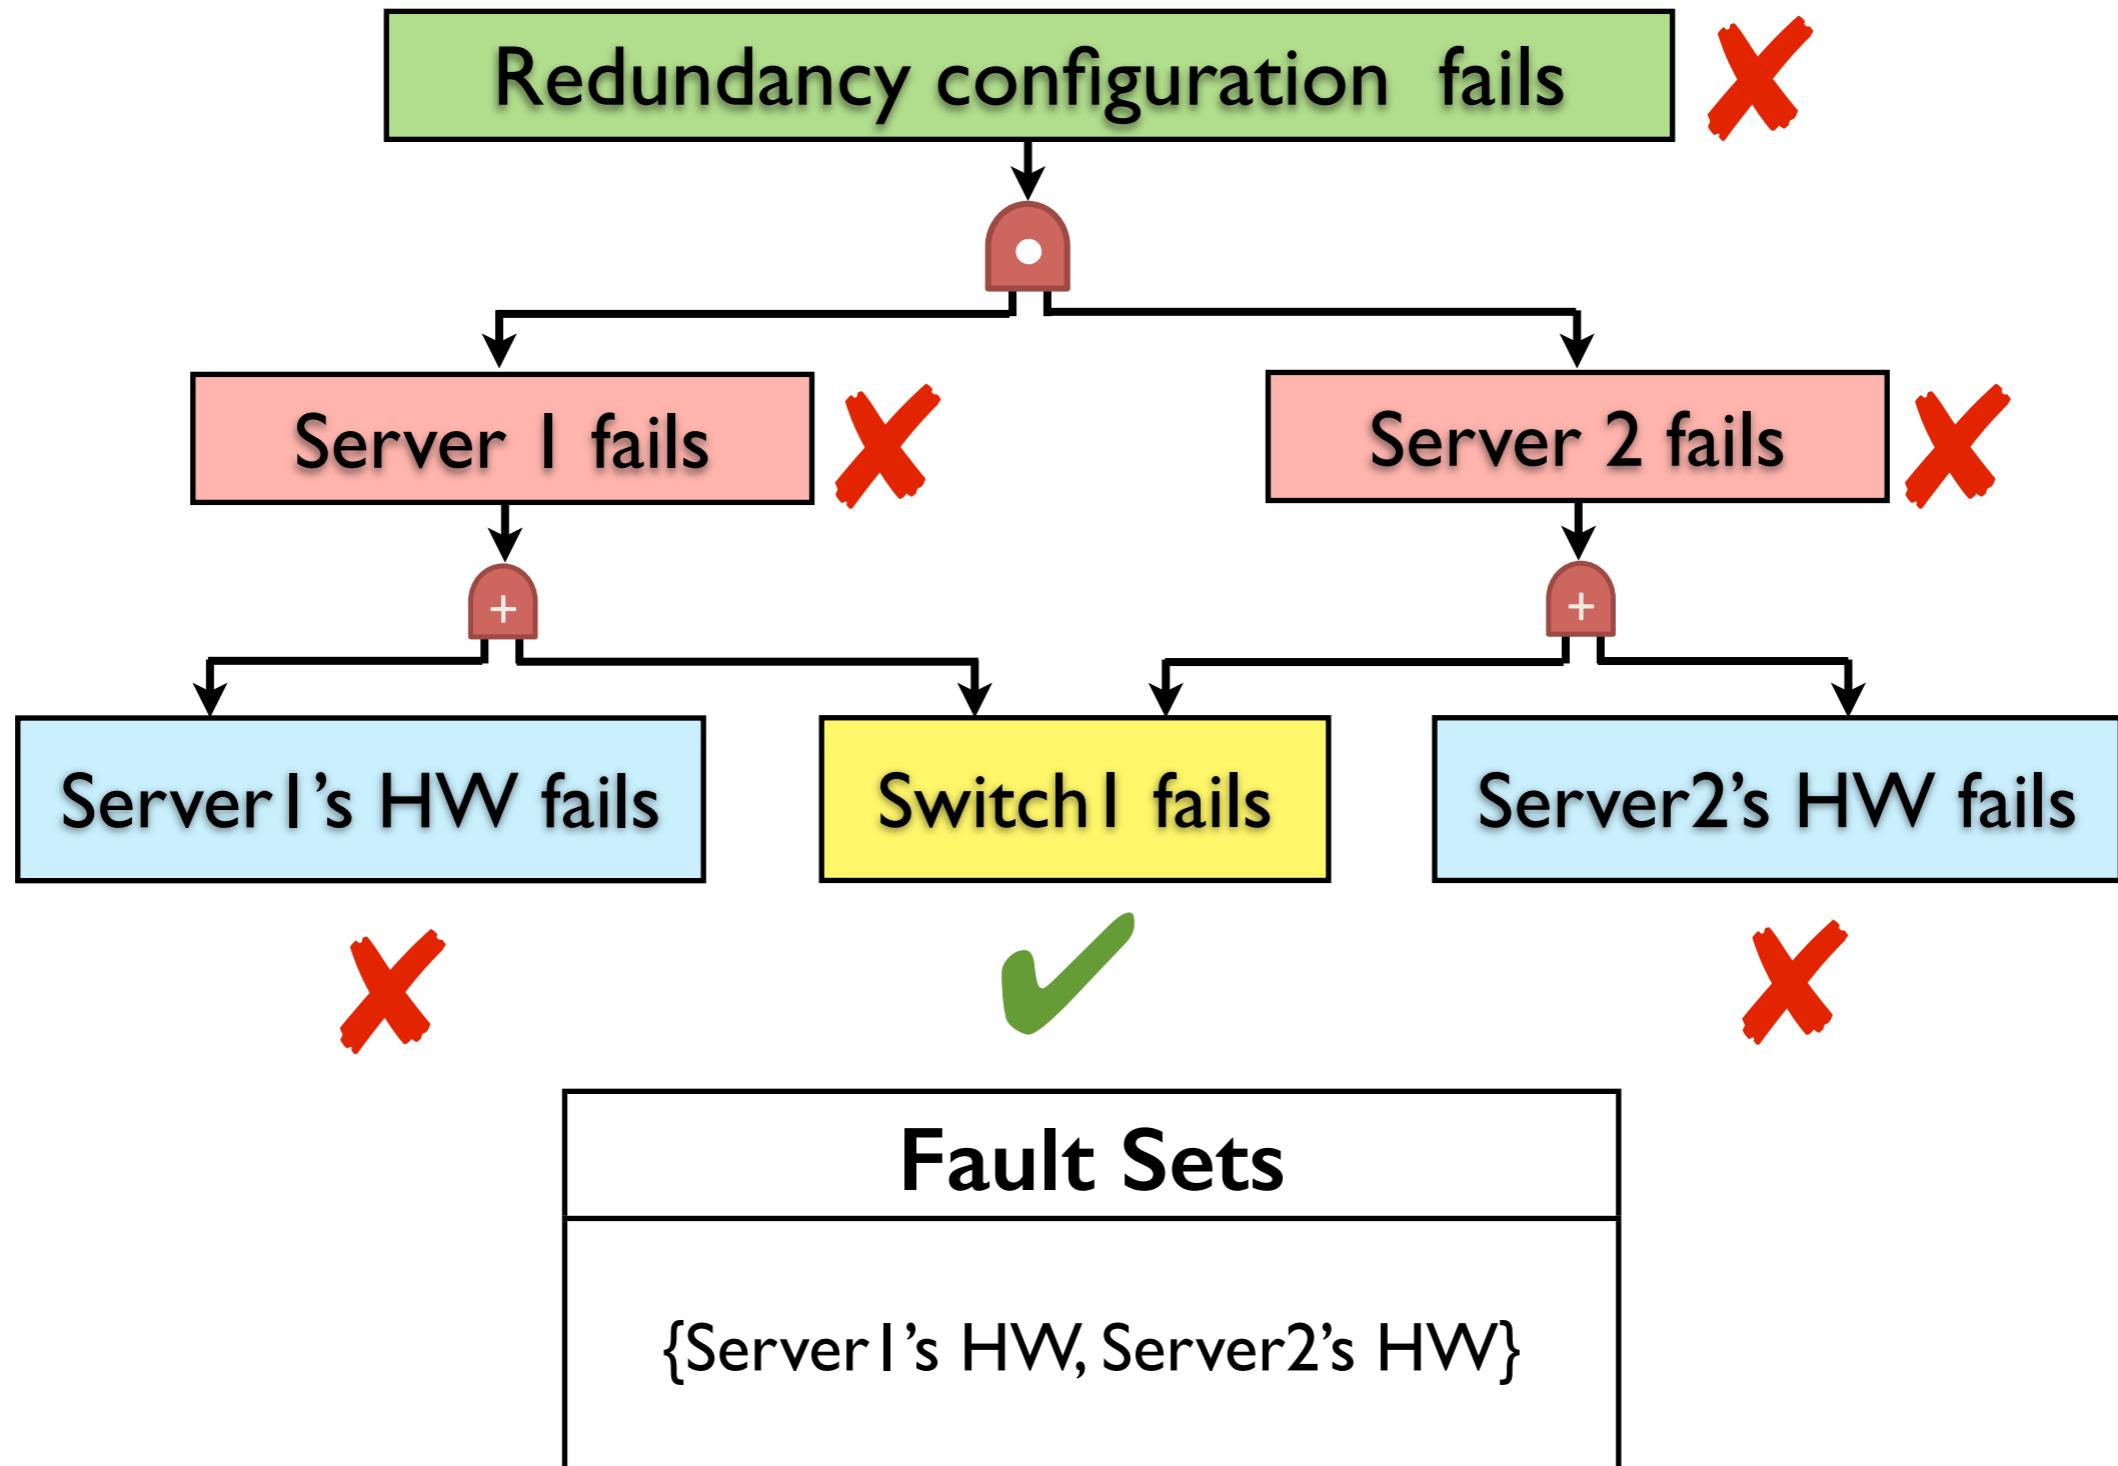
# The 1st Sampling Round



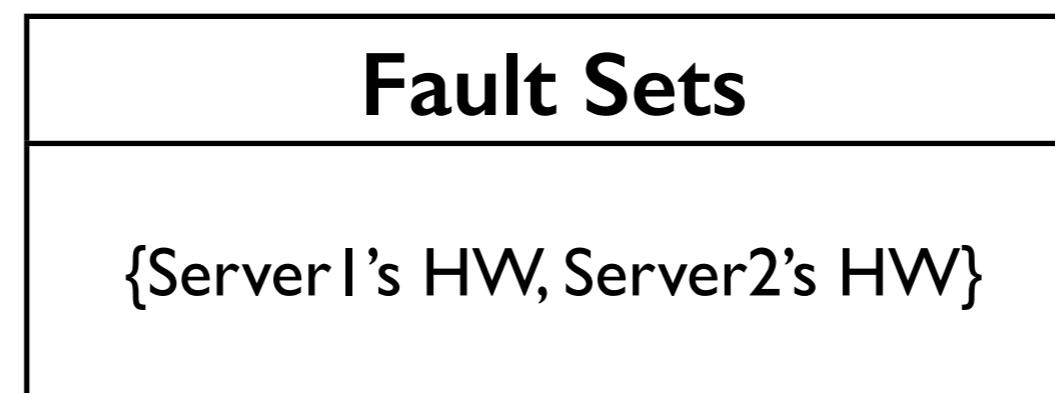
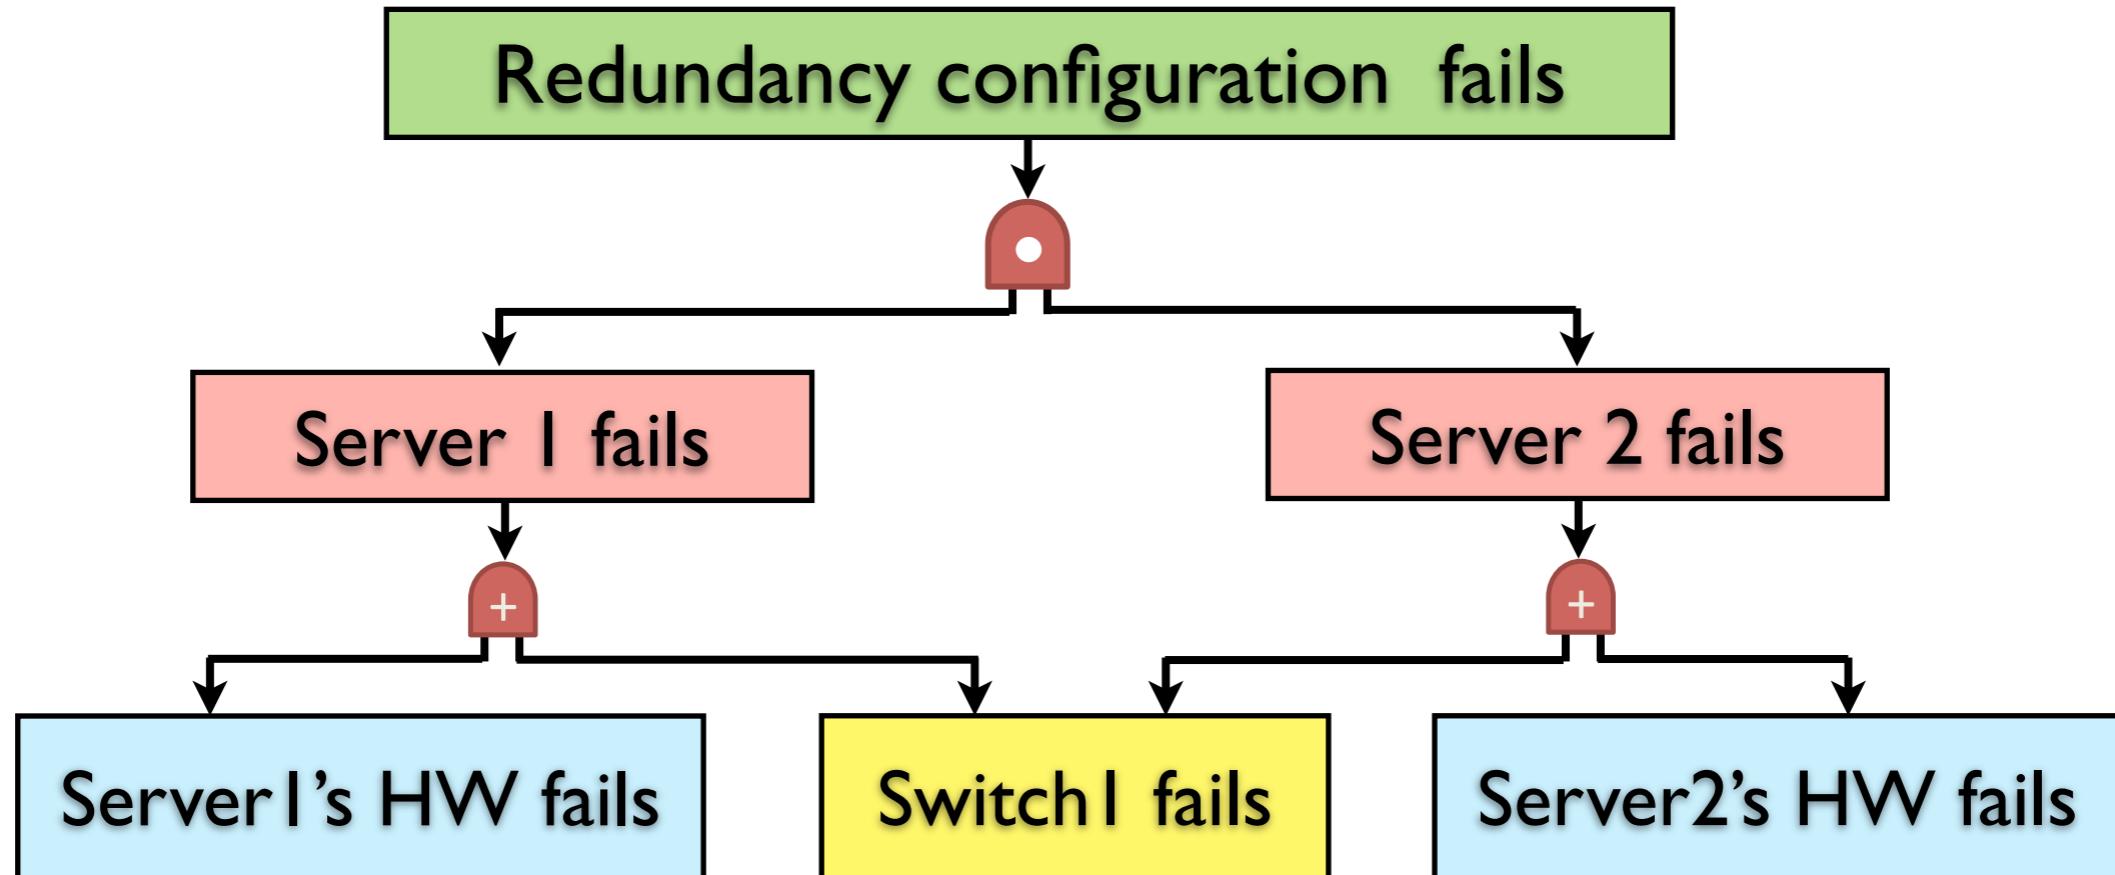
# The 1st Sampling Round



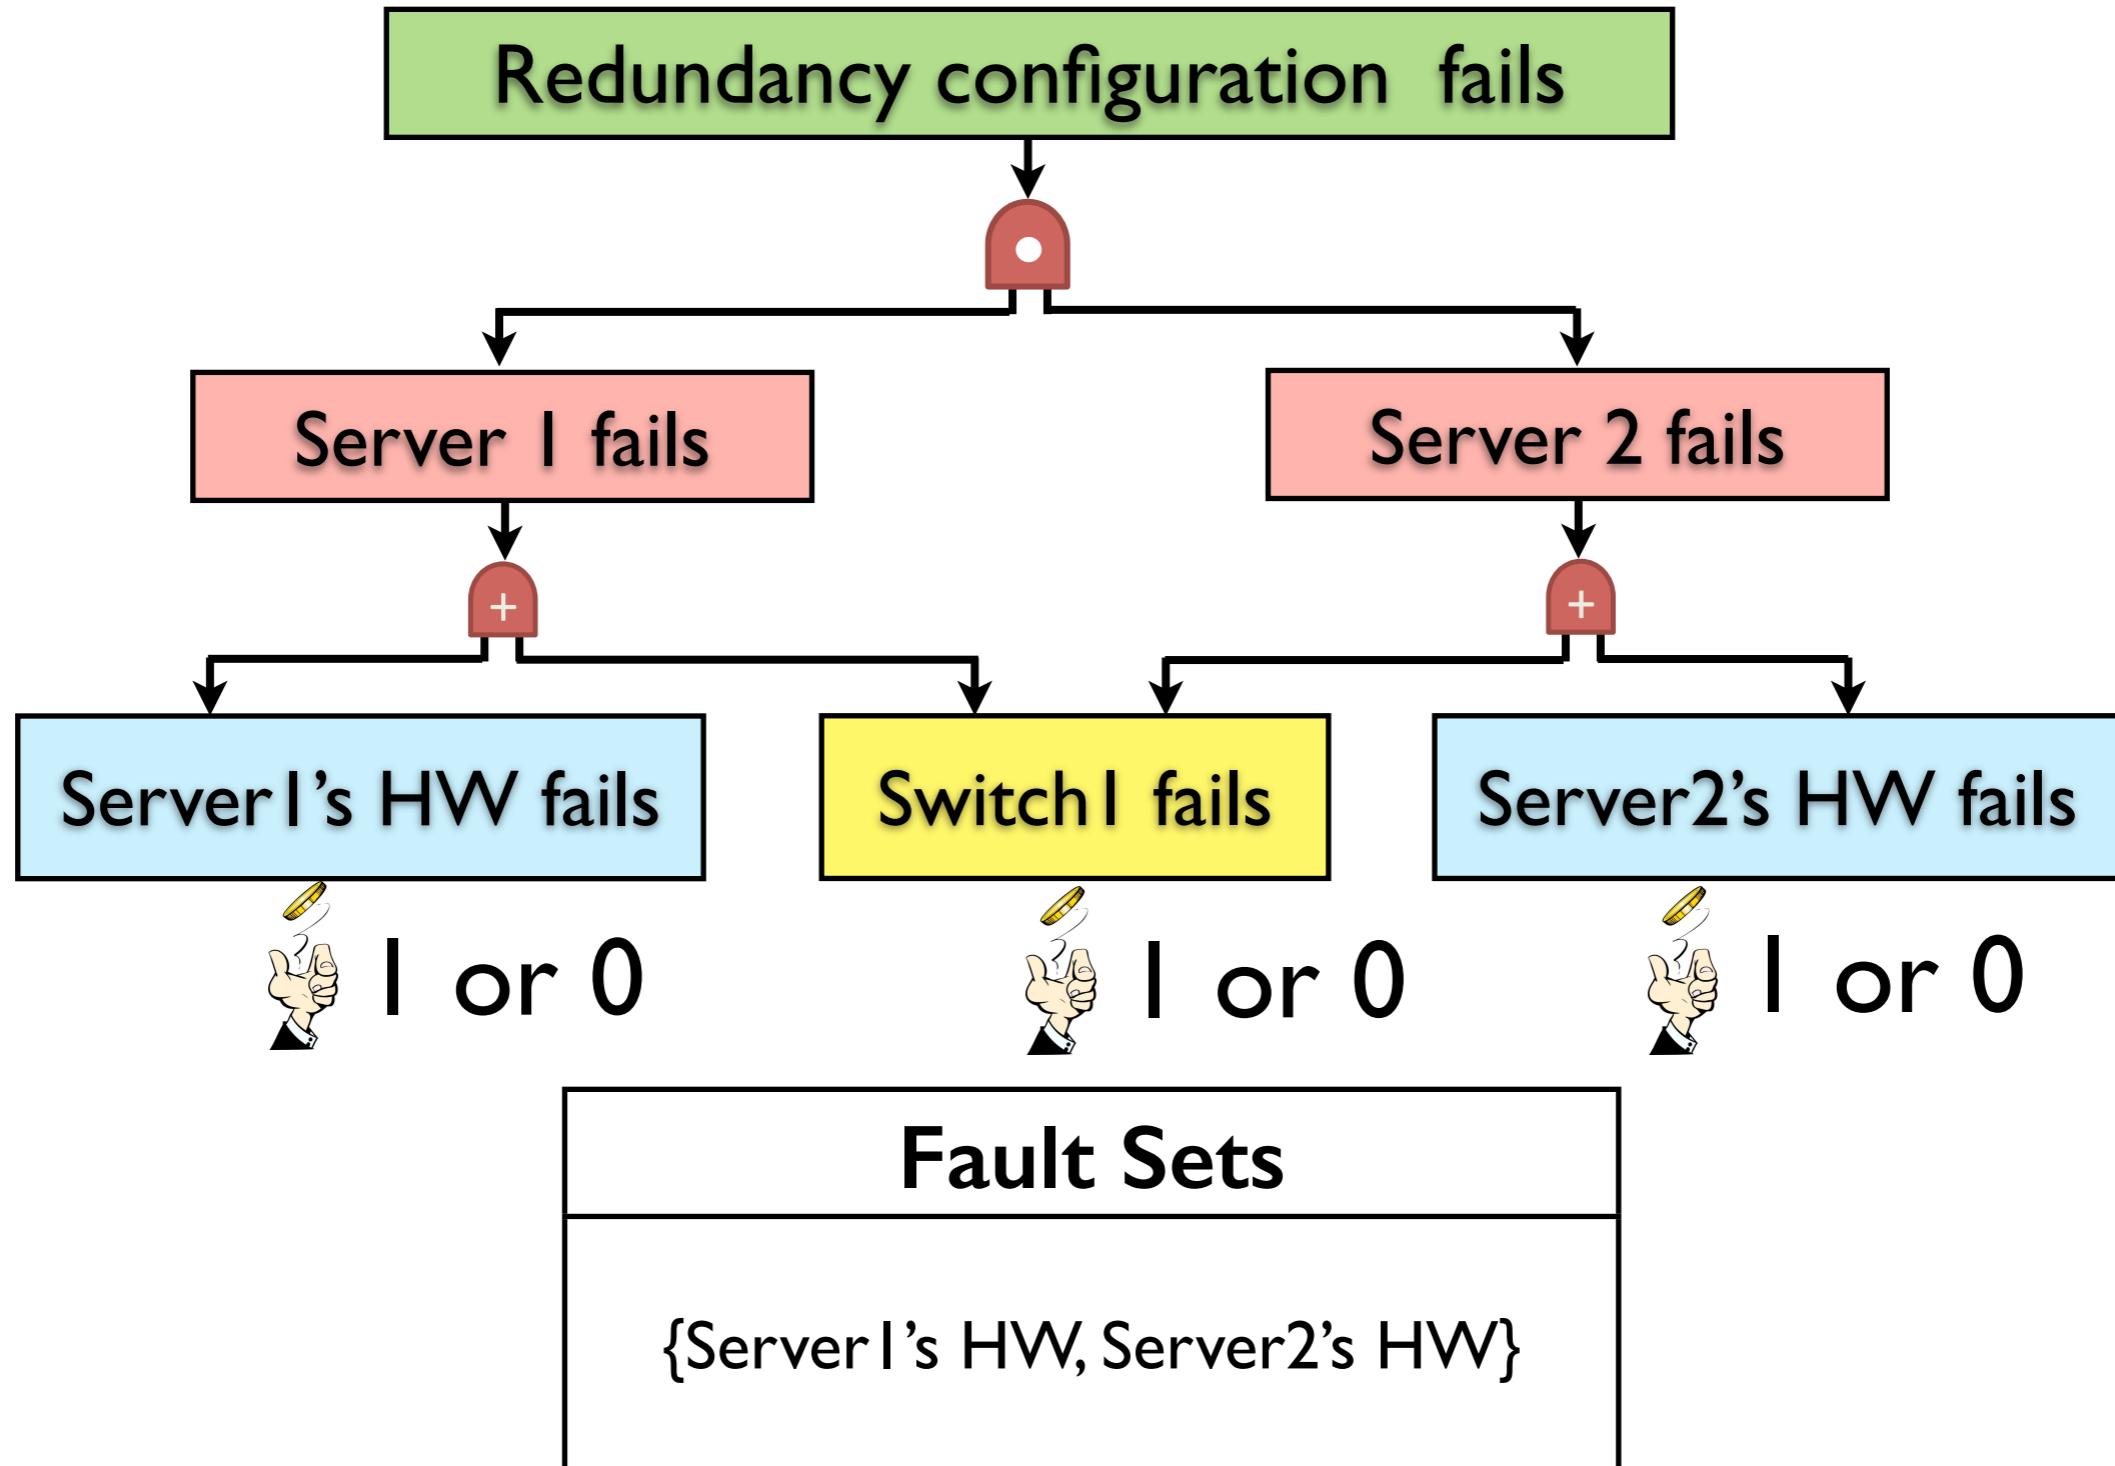
# The 1st Sampling Round



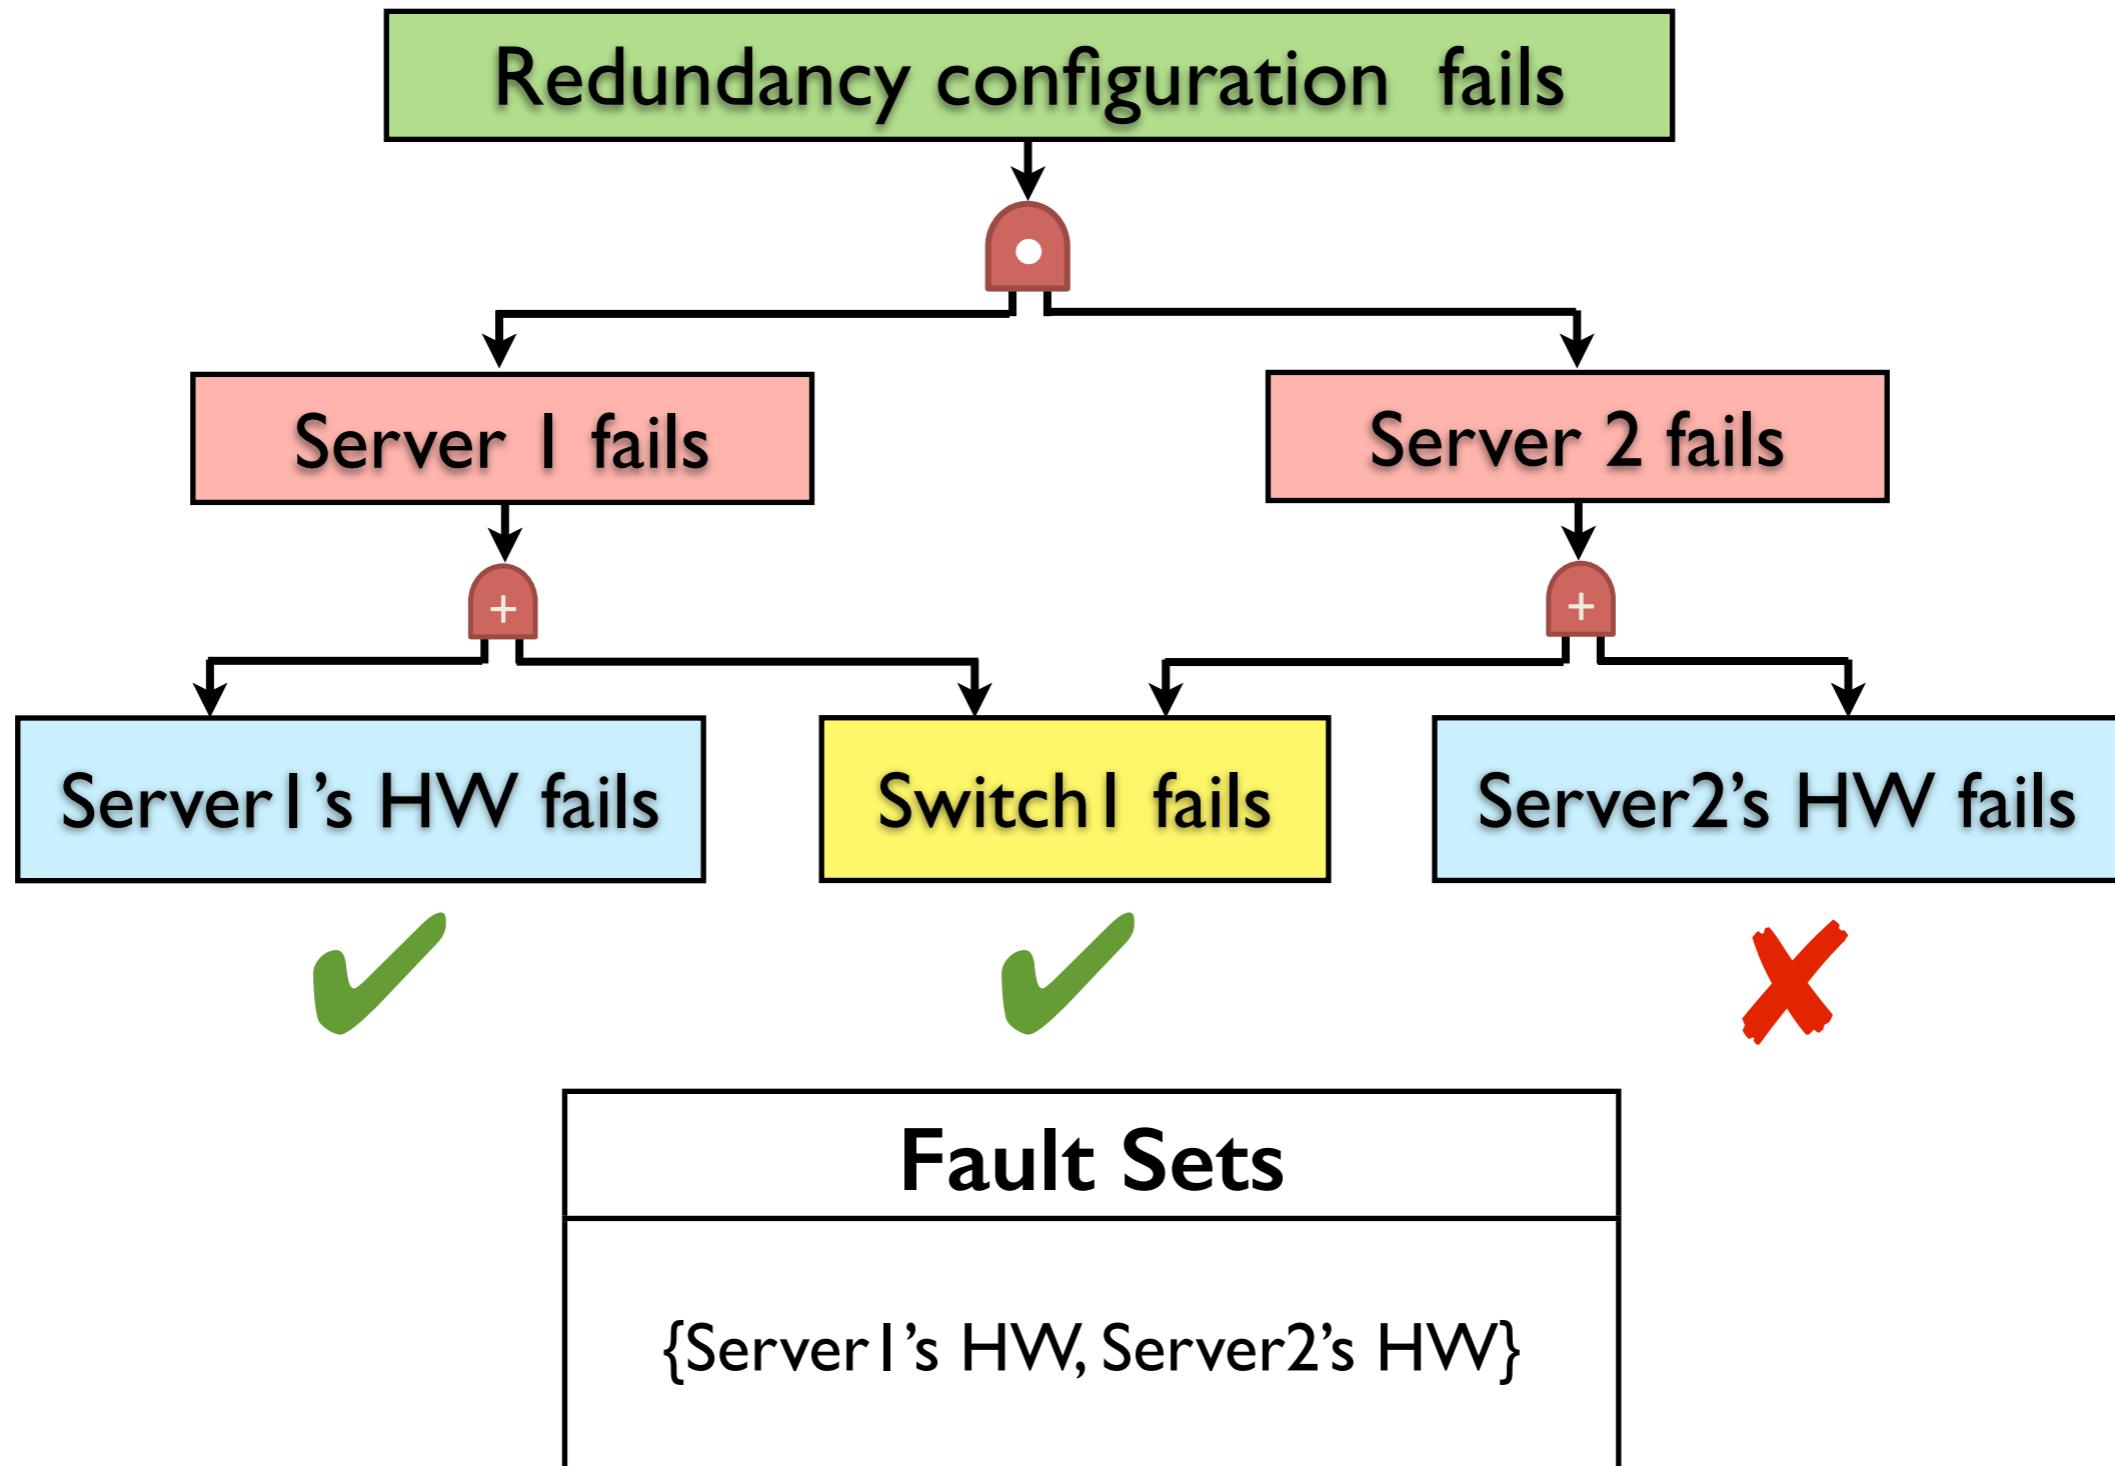
# The 2nd Sampling Round



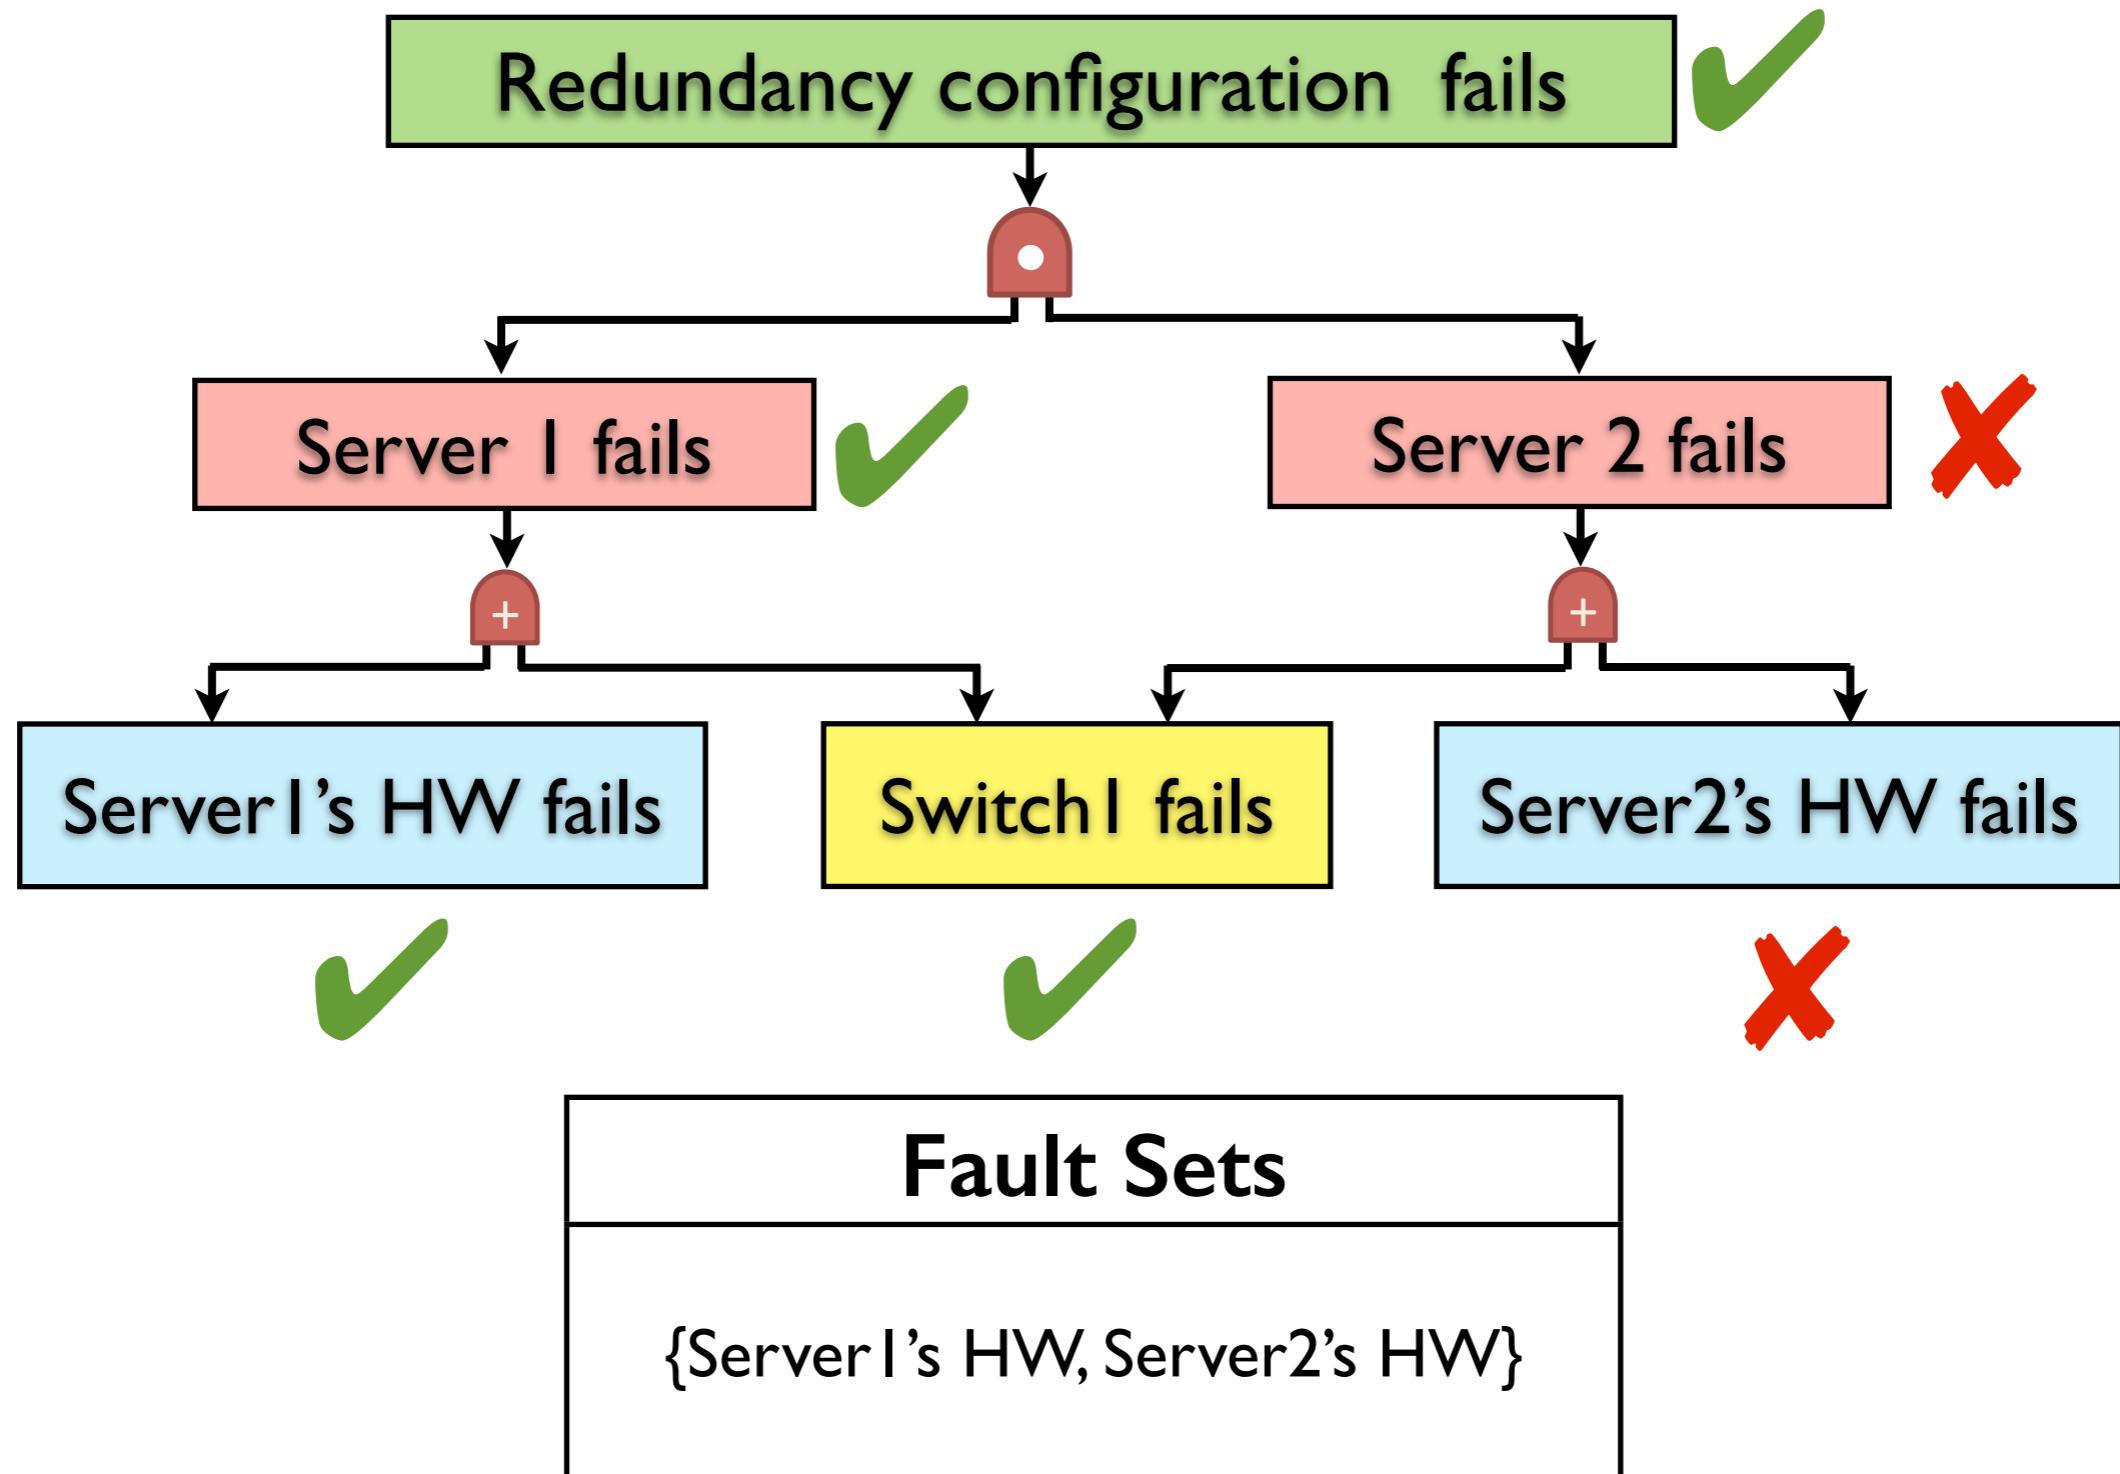
# The 2nd Sampling Round



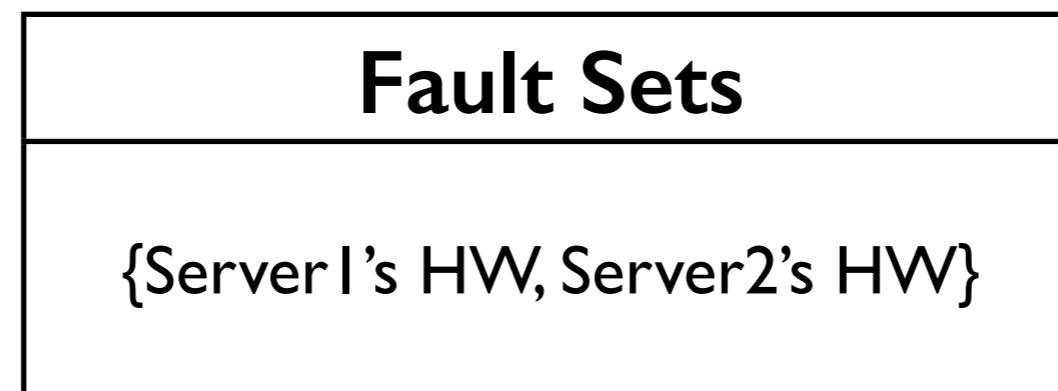
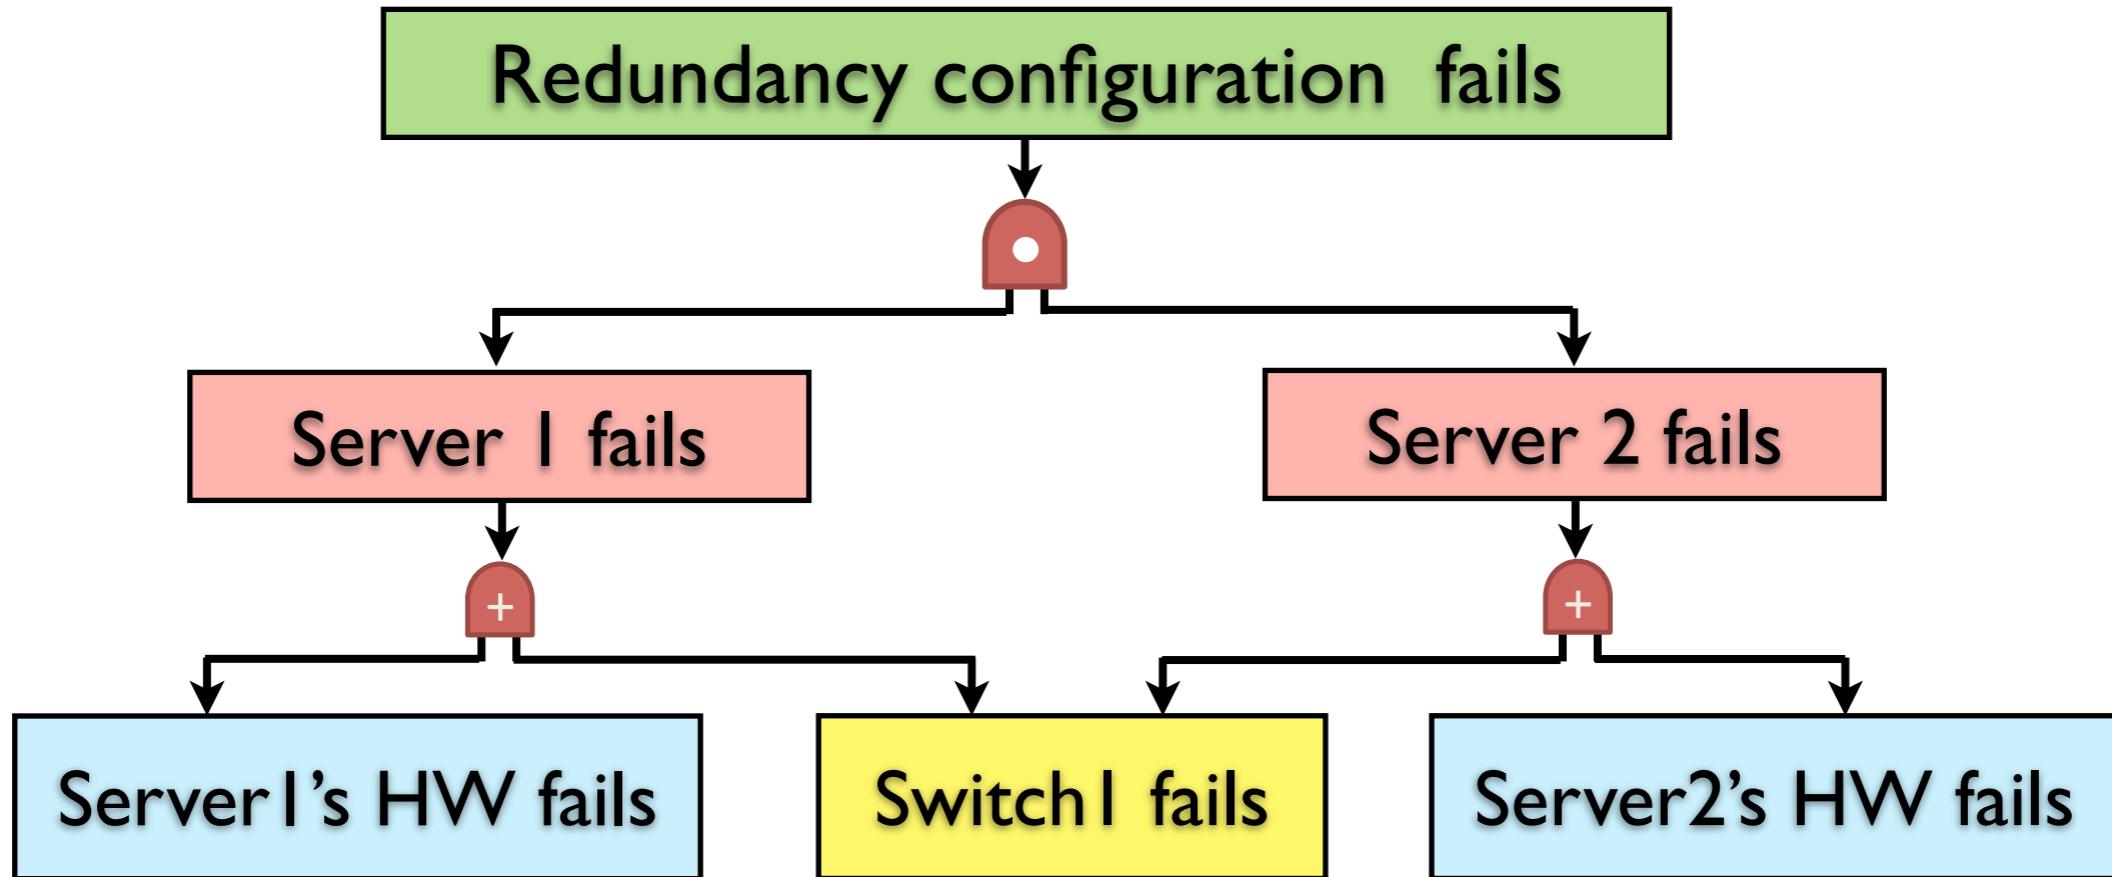
# The 2nd Sampling Round



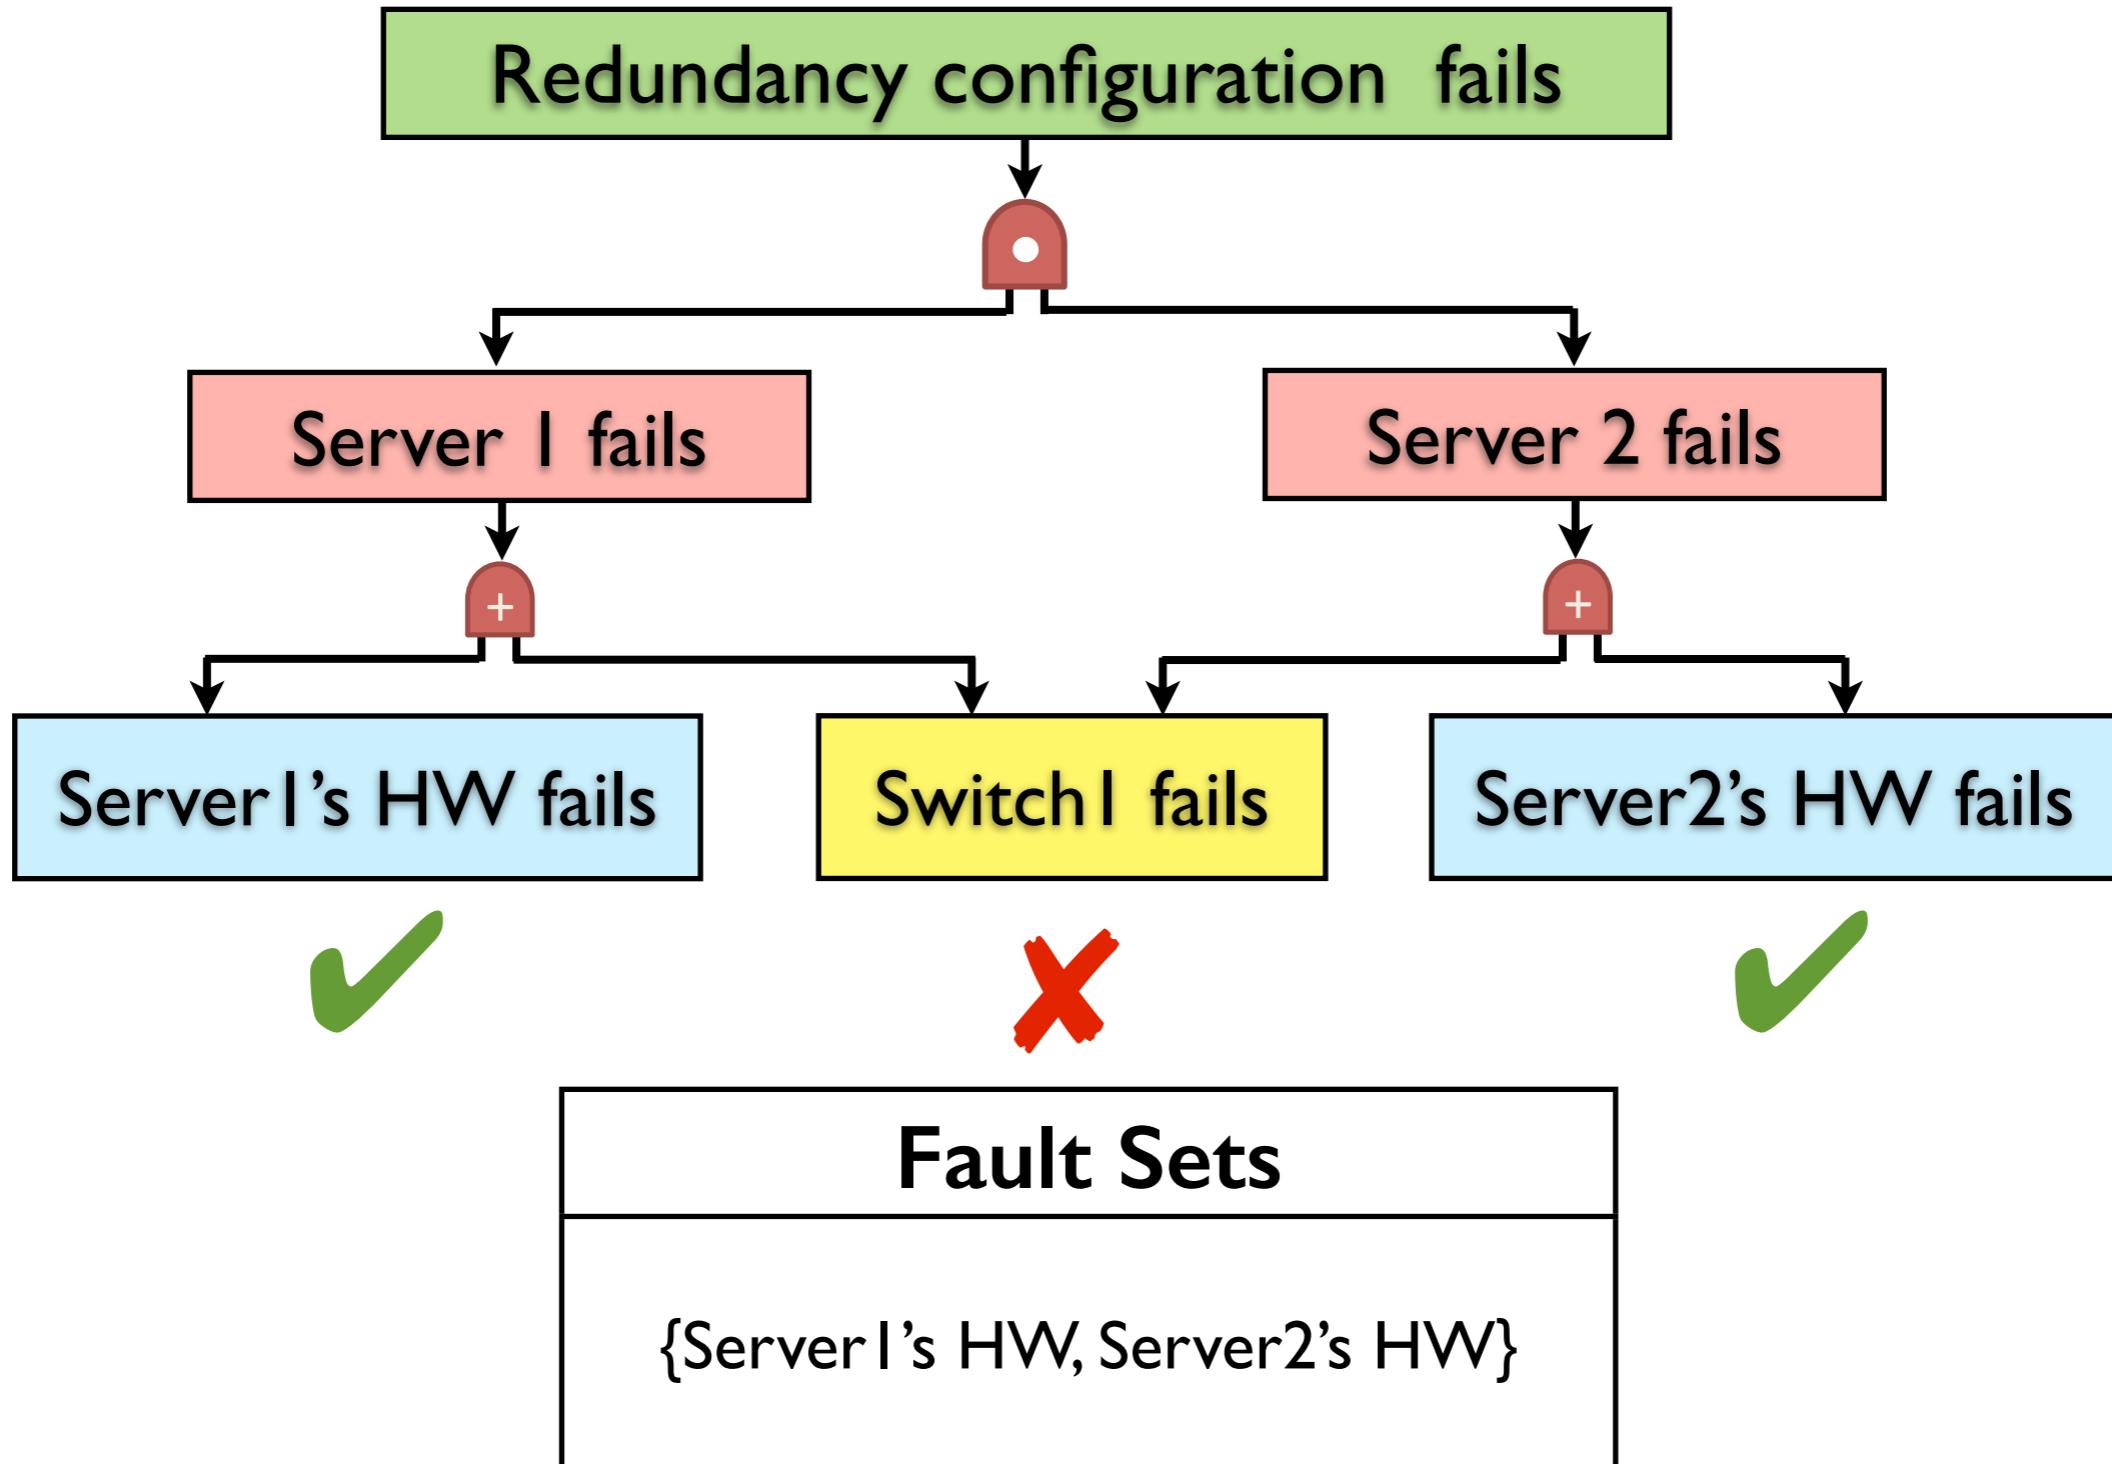
# The 2nd Sampling Round



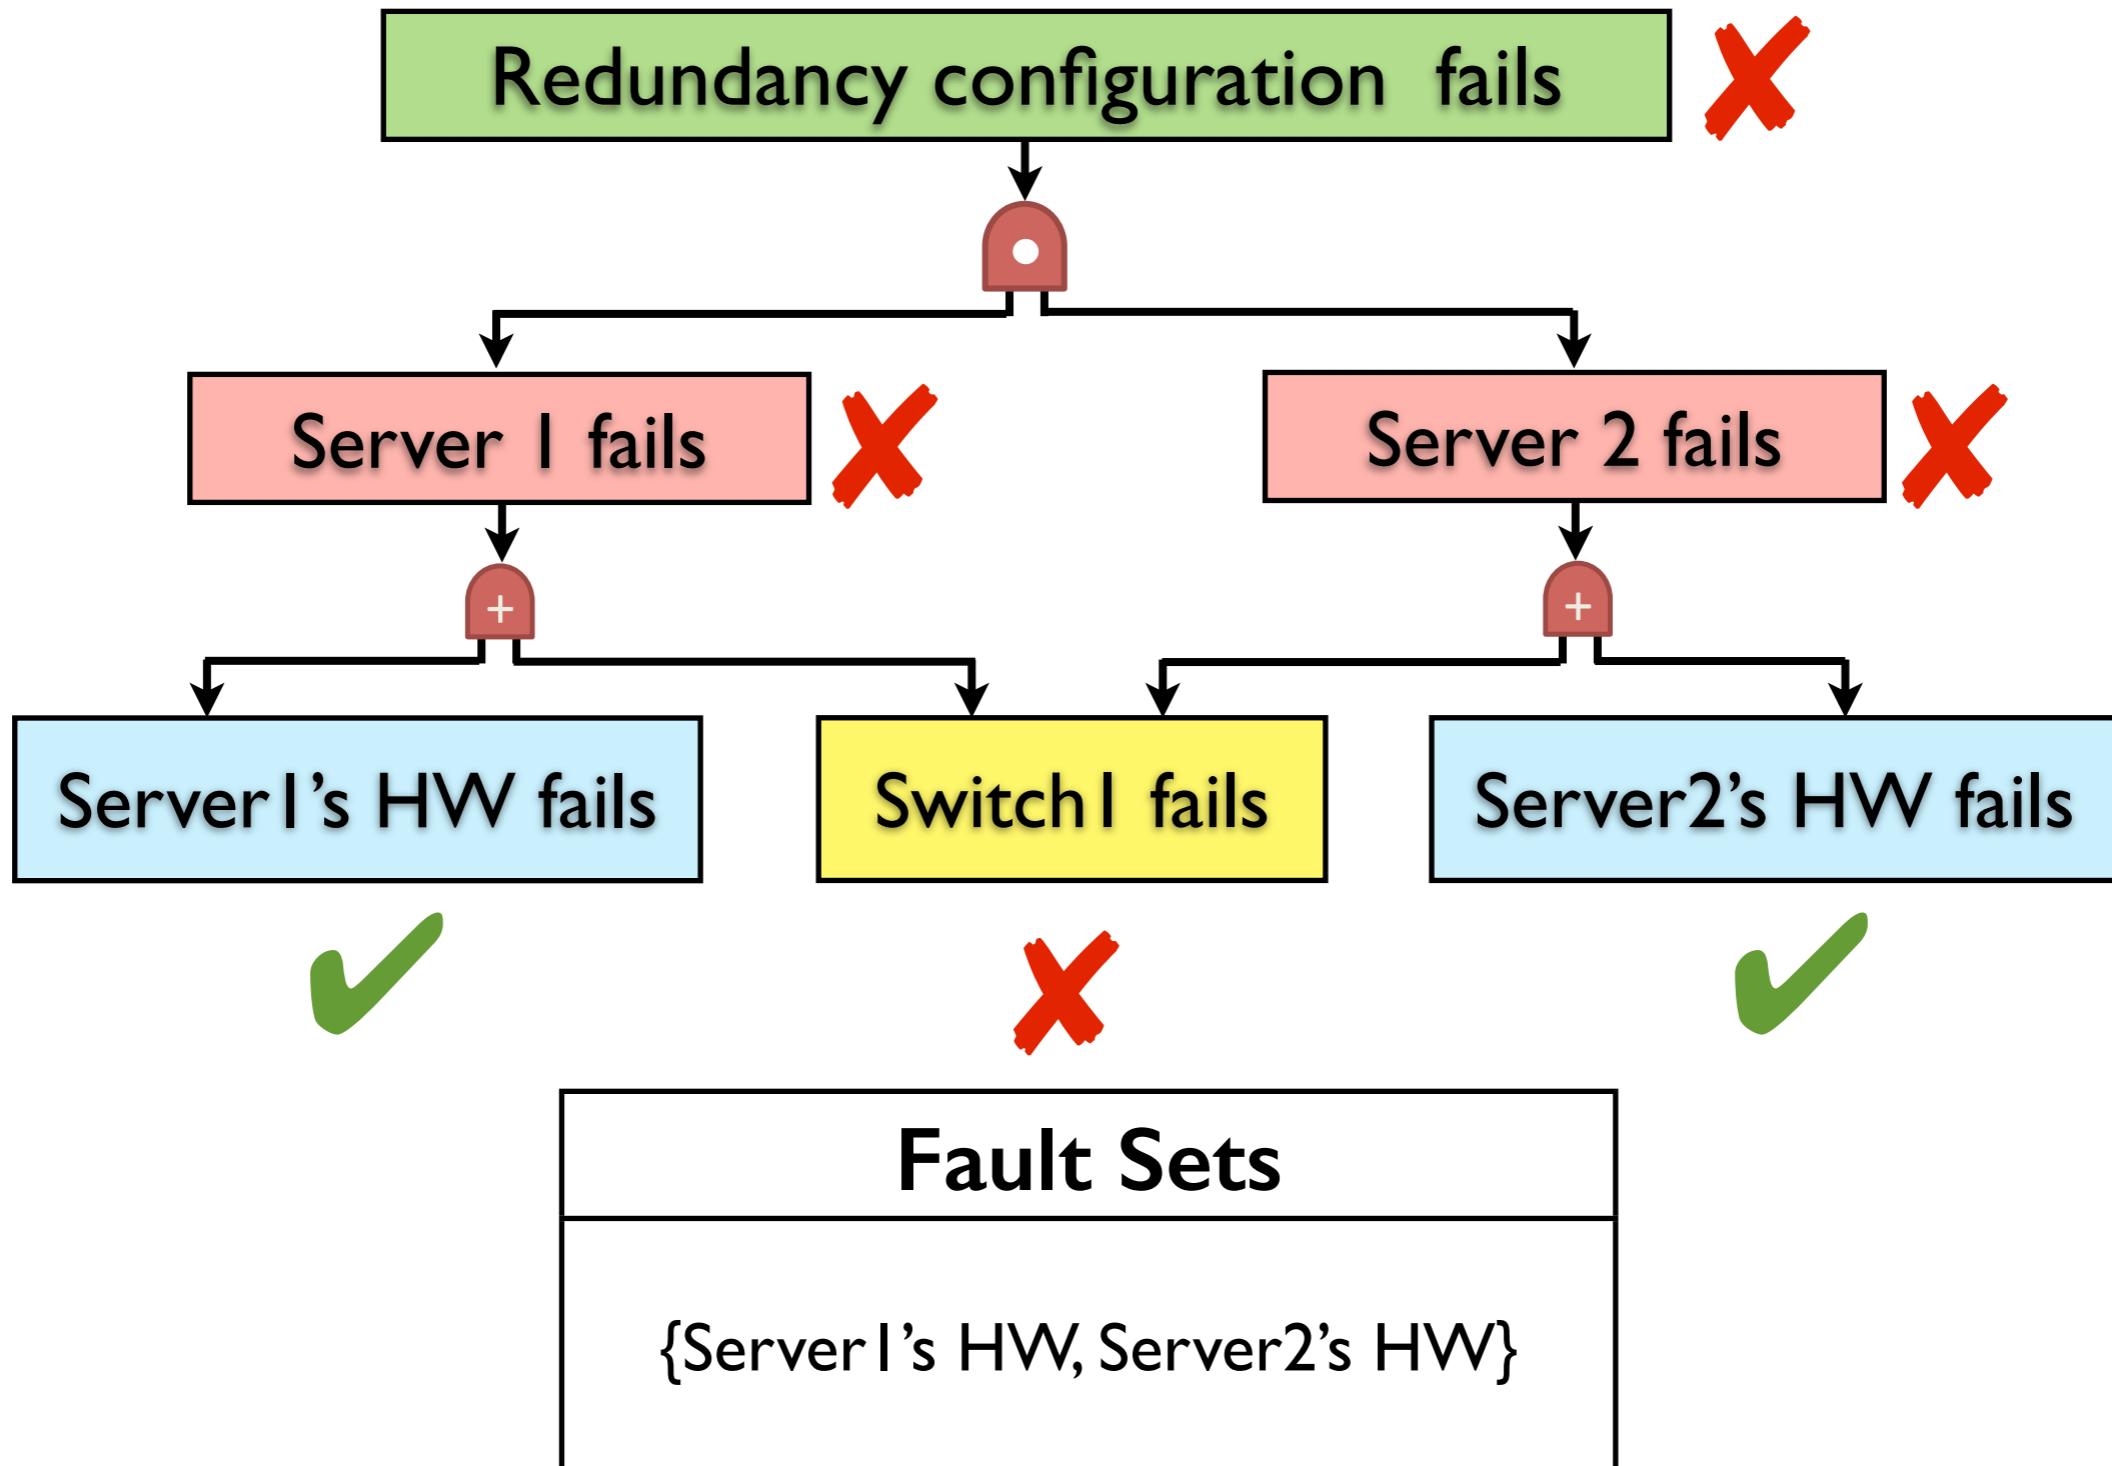
# The 3rd Sampling Round



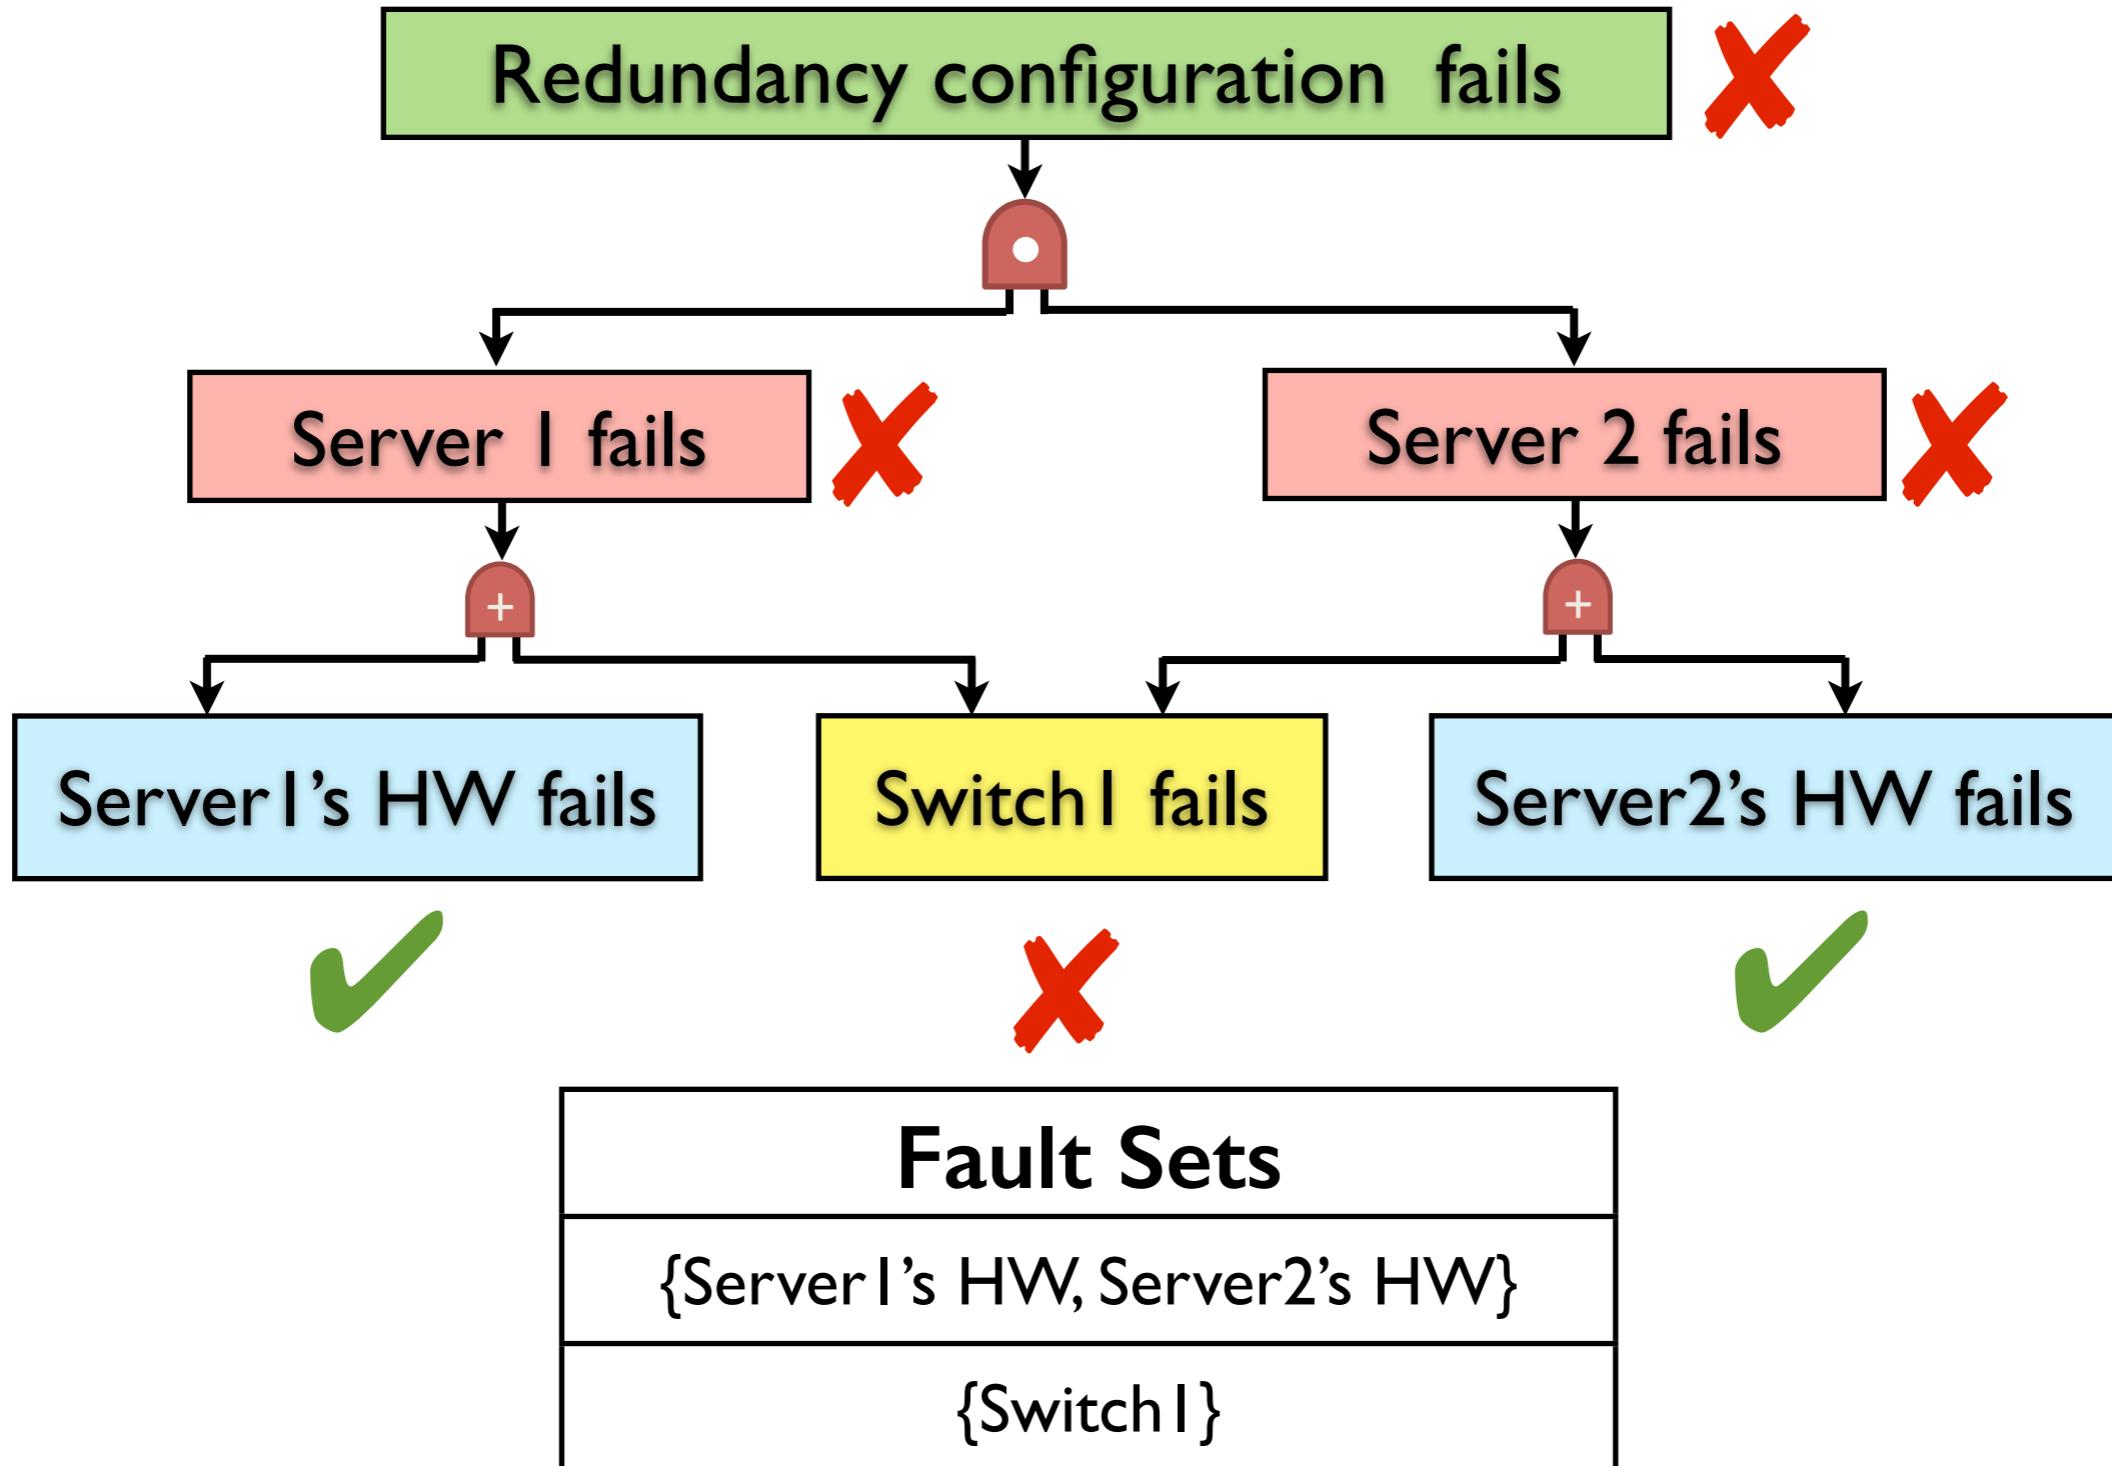
# The 3rd Sampling Round



# The 3rd Sampling Round



# The 3rd Sampling Round



# After Many (e.g., $10^7$ ) Rounds

| Fault Sets                   |
|------------------------------|
| {Server1's HW, Server2's HW} |
| {Switch1}                    |
| {Switch1, Server2's HW}      |
| {Switch1}                    |
| {Switch1, Server2's HW}      |
| ....                         |

# Size-Based Ranking

| Fault Sets                   |
|------------------------------|
| {Server1's HW, Server2's HW} |
| {Switch1}                    |
| {Switch1, Server2's HW}      |
| {Switch1}                    |
| {Switch1, Server2's HW}      |
| ....                         |

# Size-Based Ranking

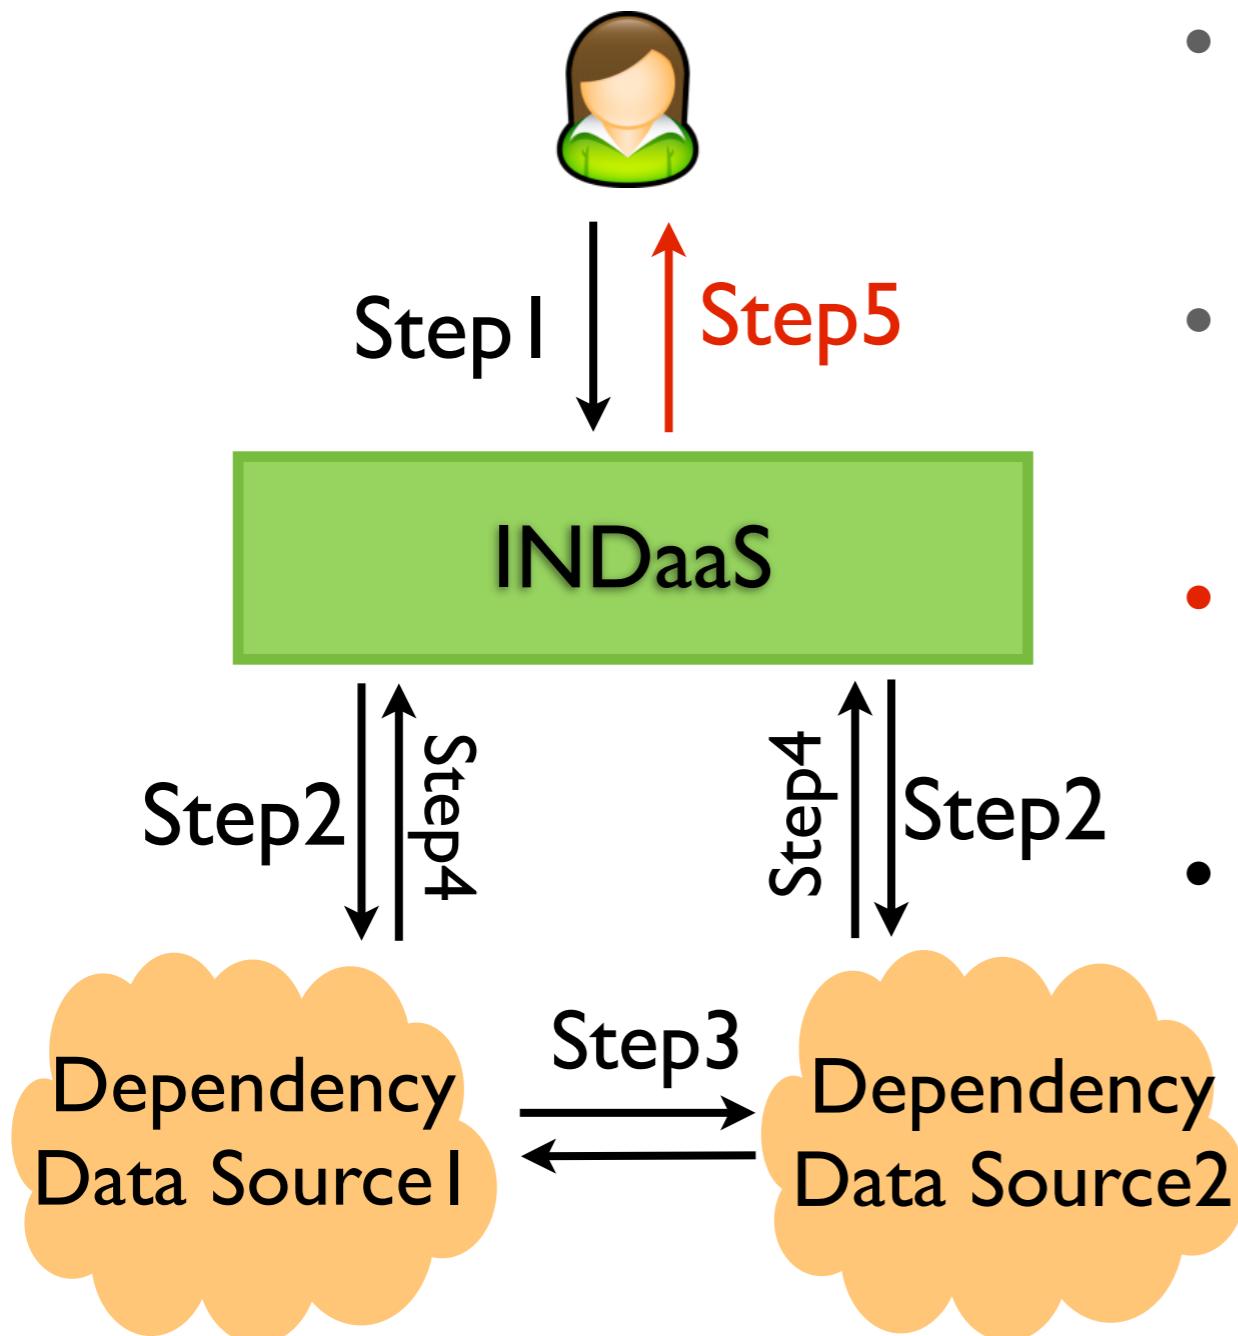
| Fault Sets                   |
|------------------------------|
| {Switch 1}                   |
| {Switch 1}                   |
| {Switch 1, Server2's HW}     |
| {Switch 1, Server2's HW}     |
| {Server1's HW, Server2's HW} |
| ....                         |

# Independence Evaluation

- Multiple equations for option:
  - summation of sizes
  - weighted average of sizes

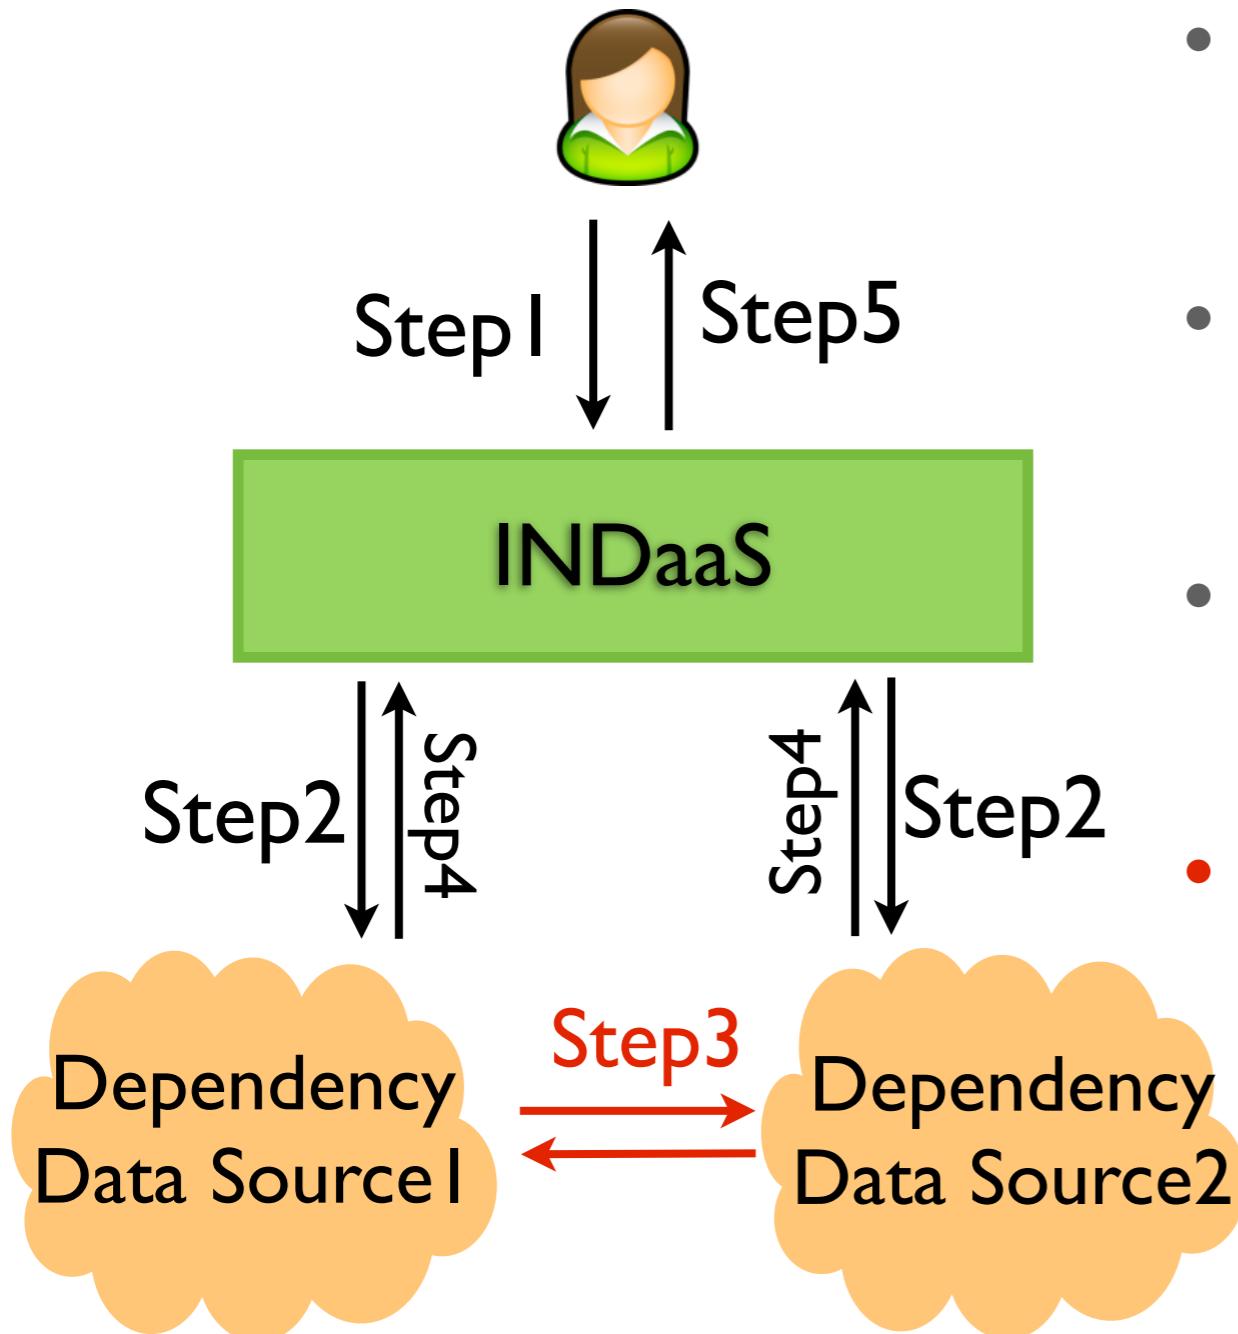
| Fault Sets                   |
|------------------------------|
| {Switch1}                    |
| {Switch1}                    |
| {Switch1, Server2's HW}      |
| {Switch1, Server2's HW}      |
| {Server1's HW, Server2's HW} |
| ... ...                      |

# RoadMap



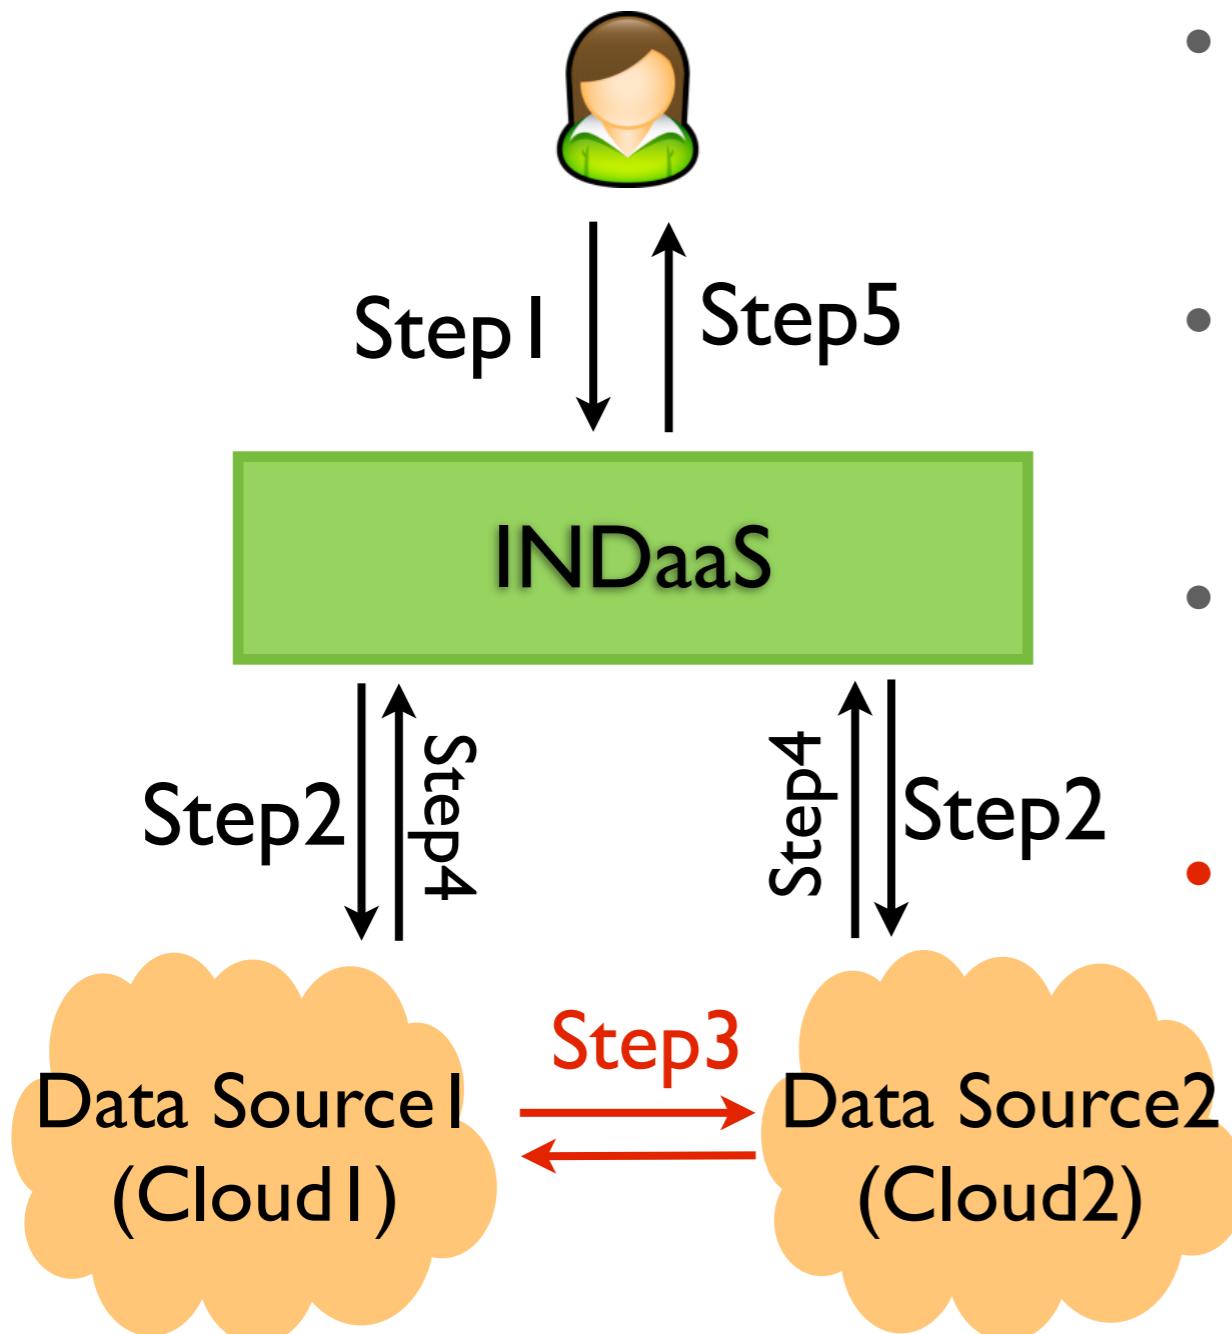
- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

# RoadMap



- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

# RoadMap



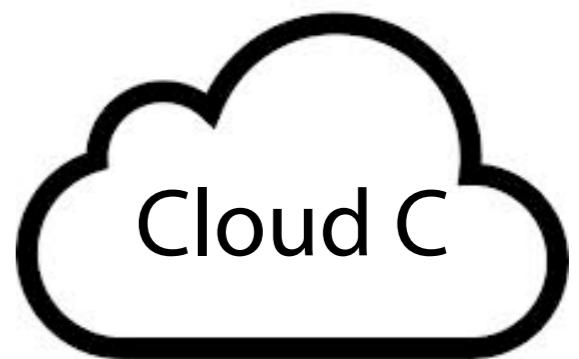
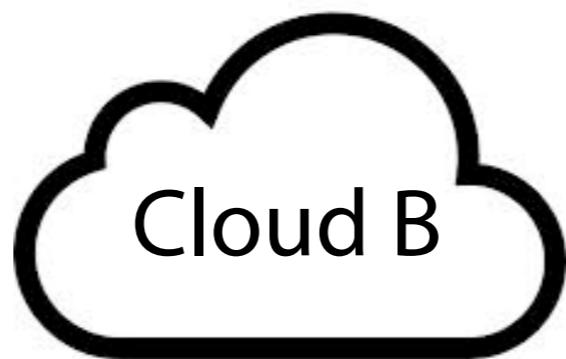
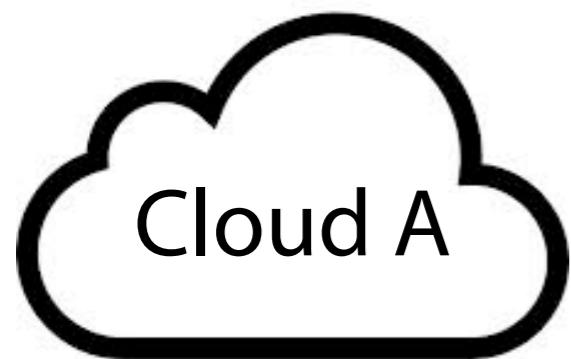
- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity



Service Provider



INDaaS Agent

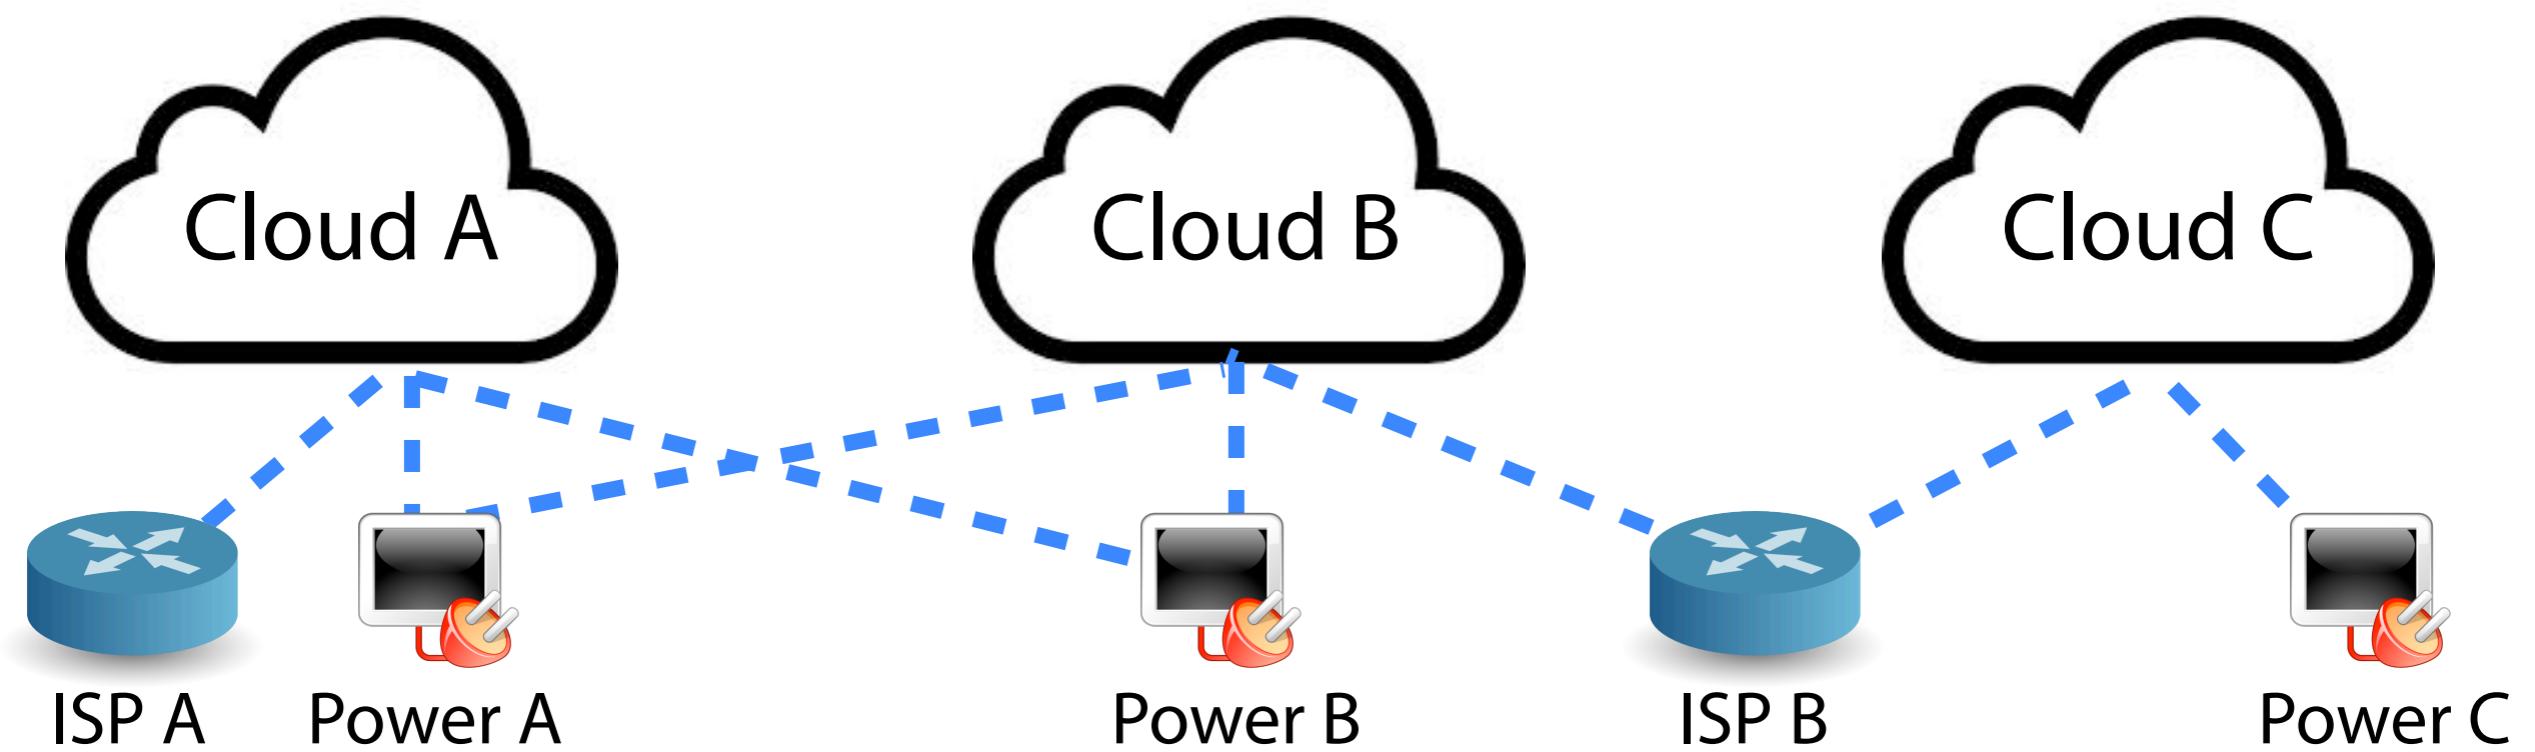


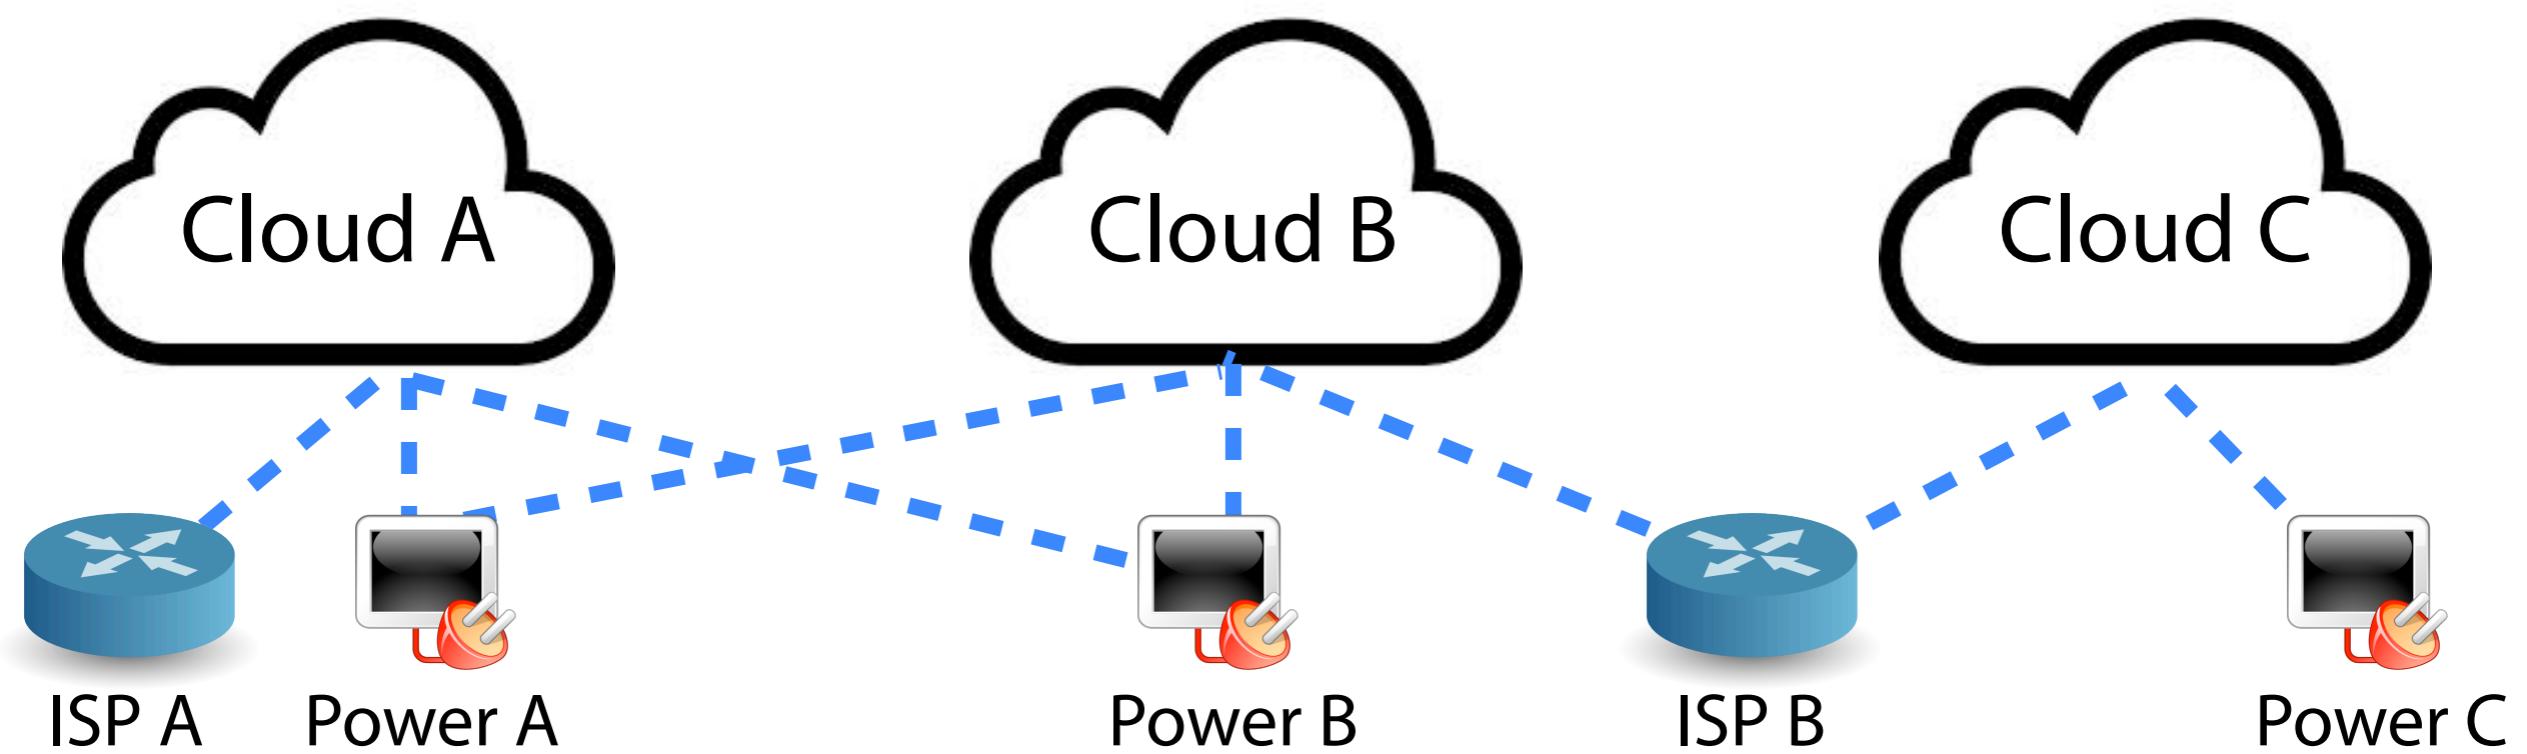
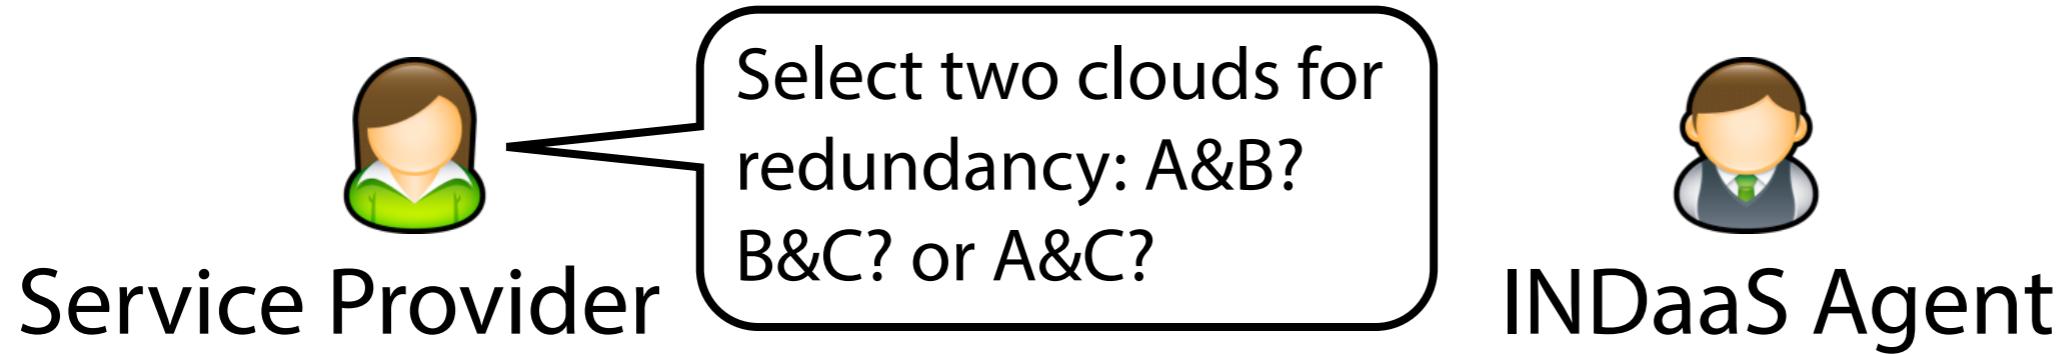


Service Provider

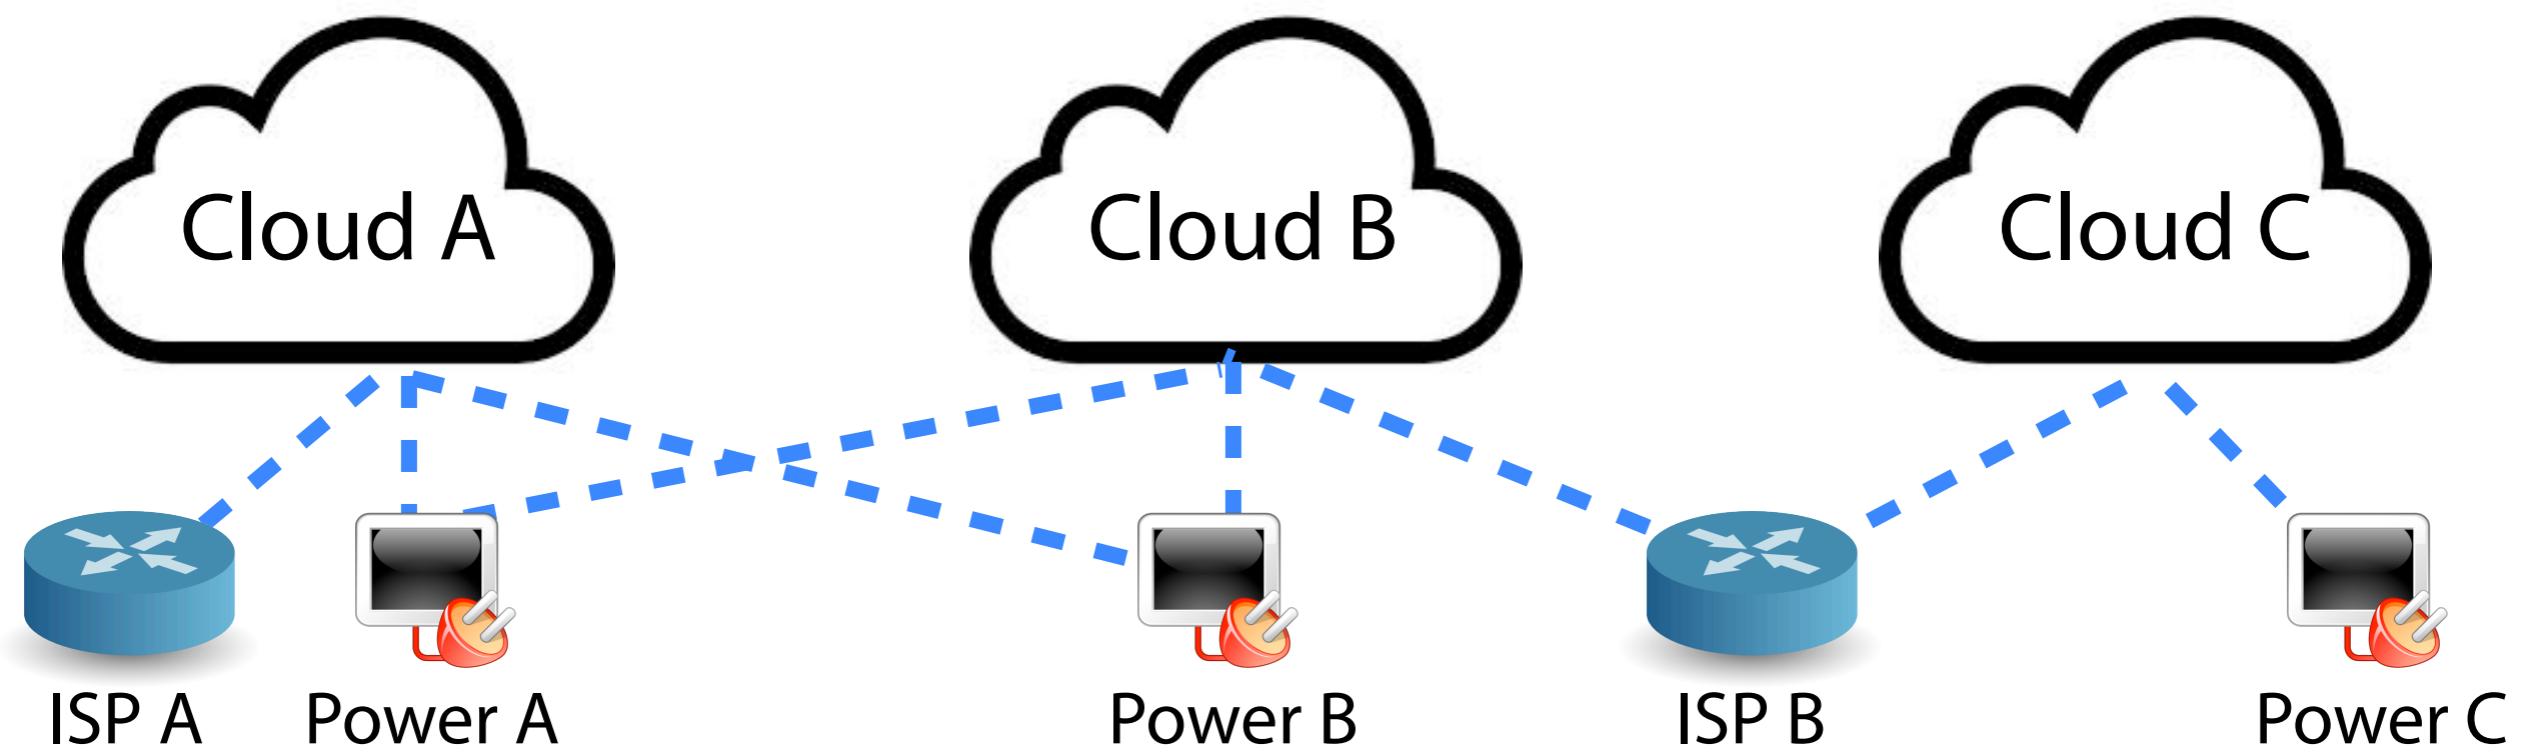
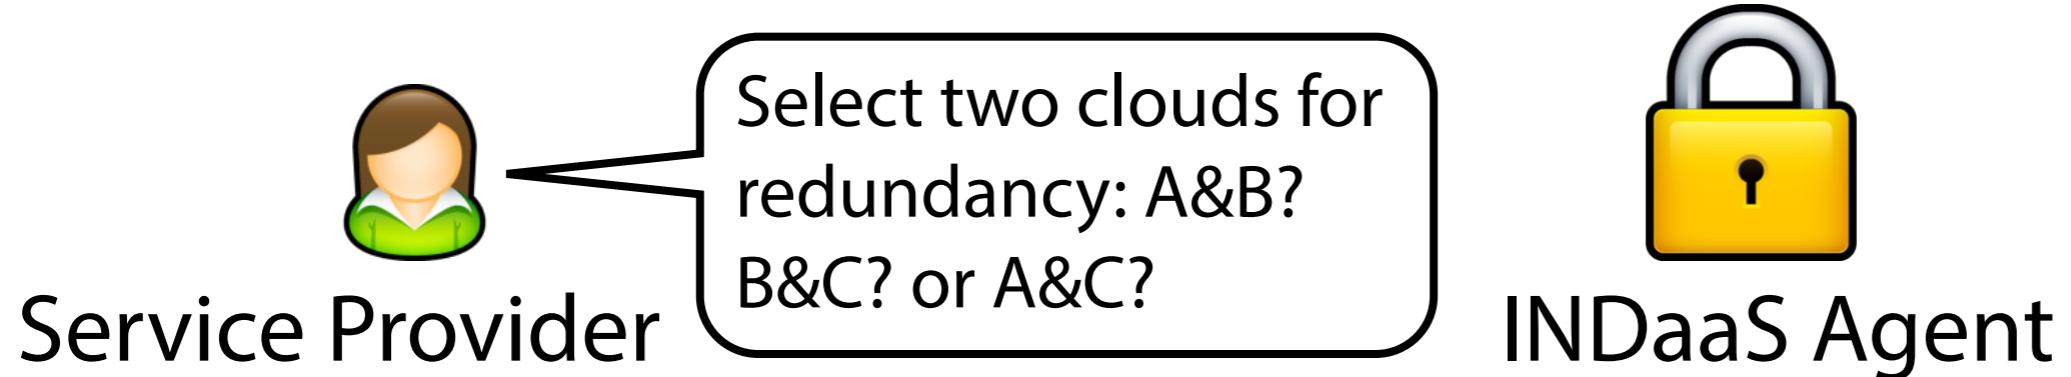


INDaaS Agent





# Trusted Third-Party



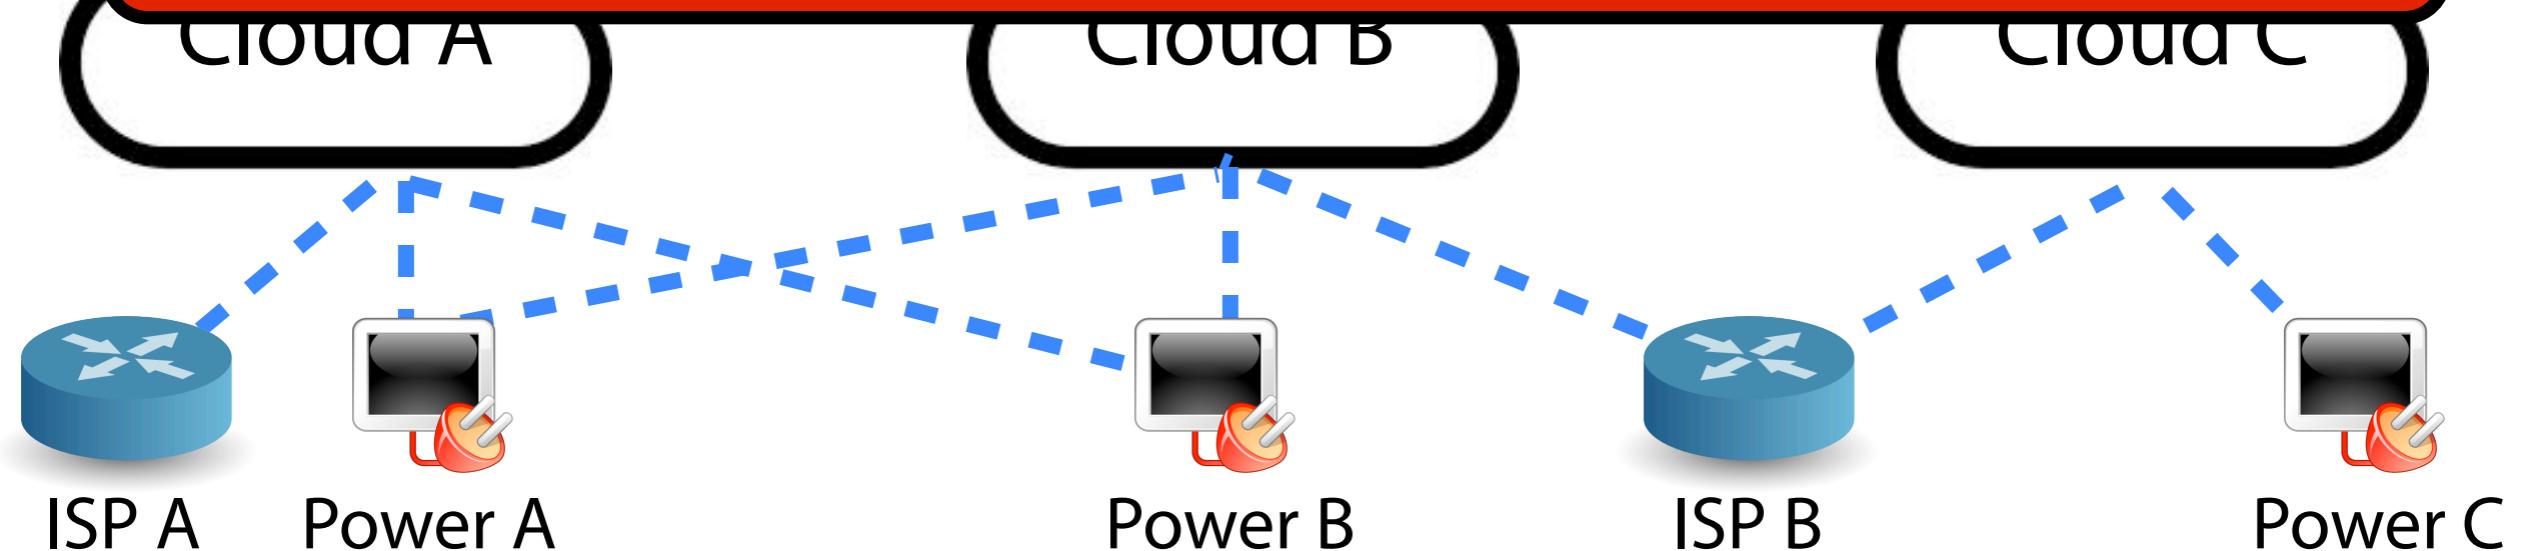
# Trusted Third-Party



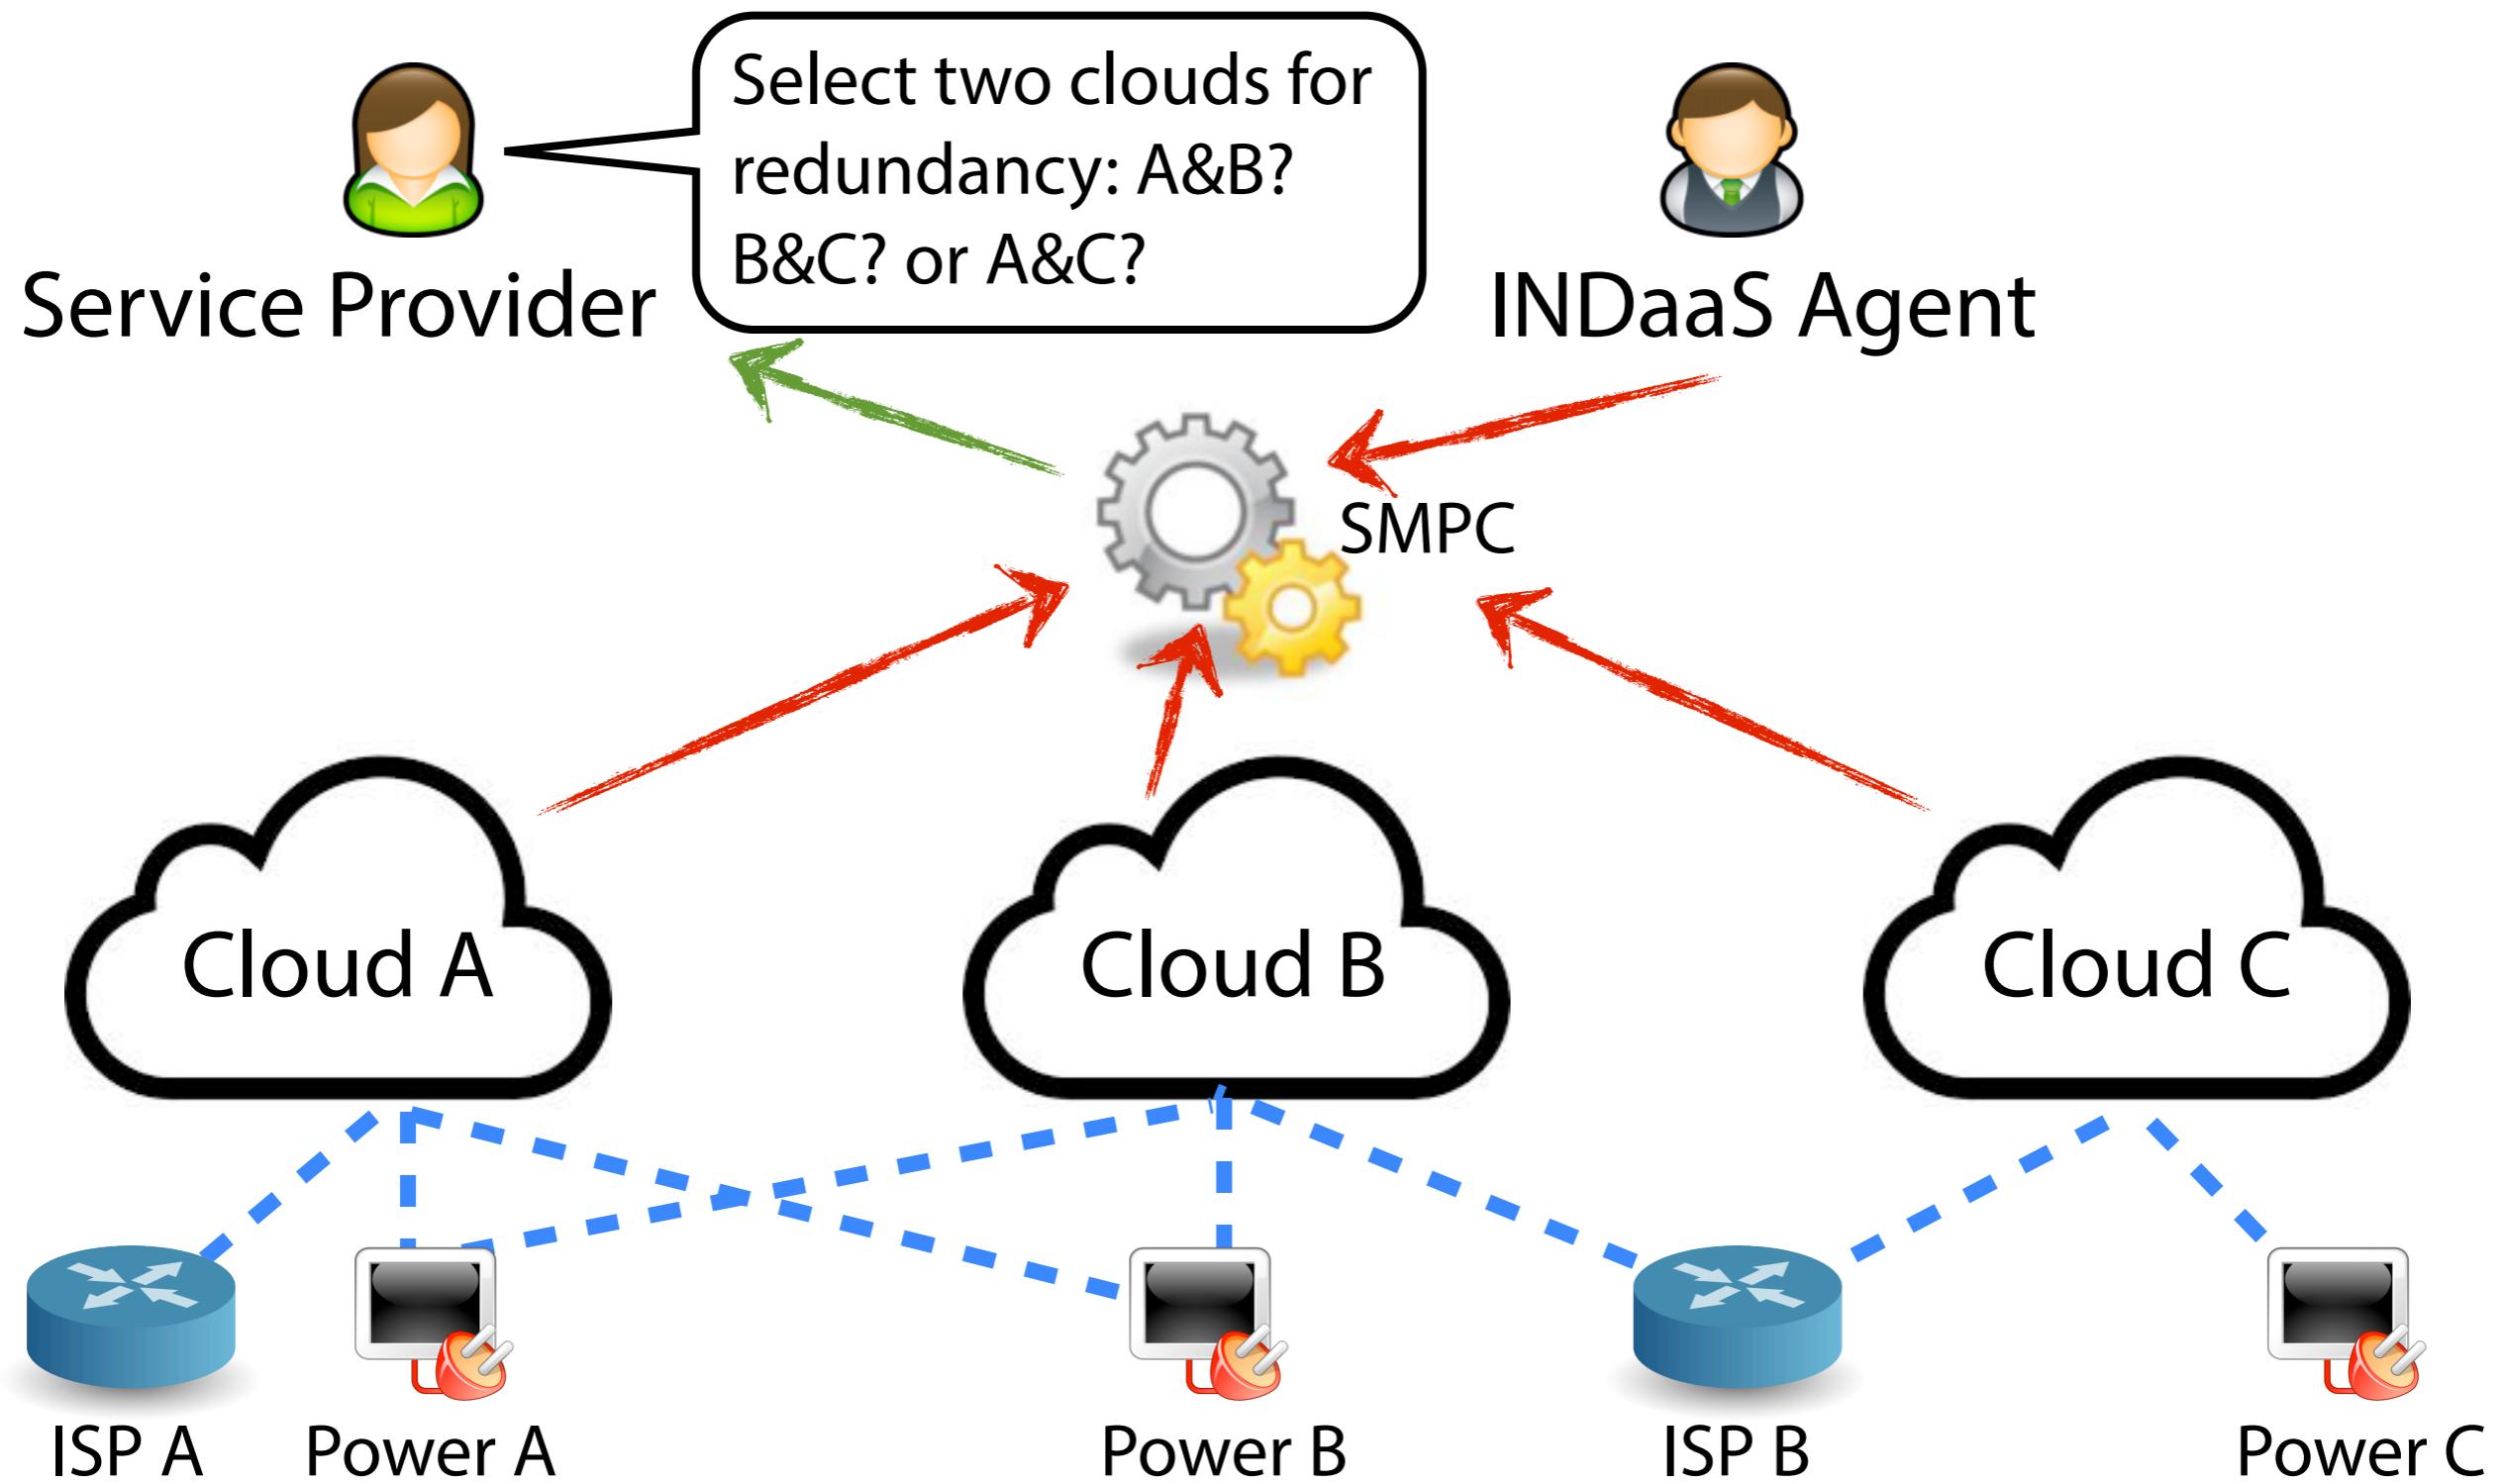
Select two clouds for redundancy: A&B?  
B&C? or A&C?



Cloud providers are reluctant  
to share this information!



# Secure Multiparty Computation (SMPC)



# Secure Multiparty Computation (SMPC)



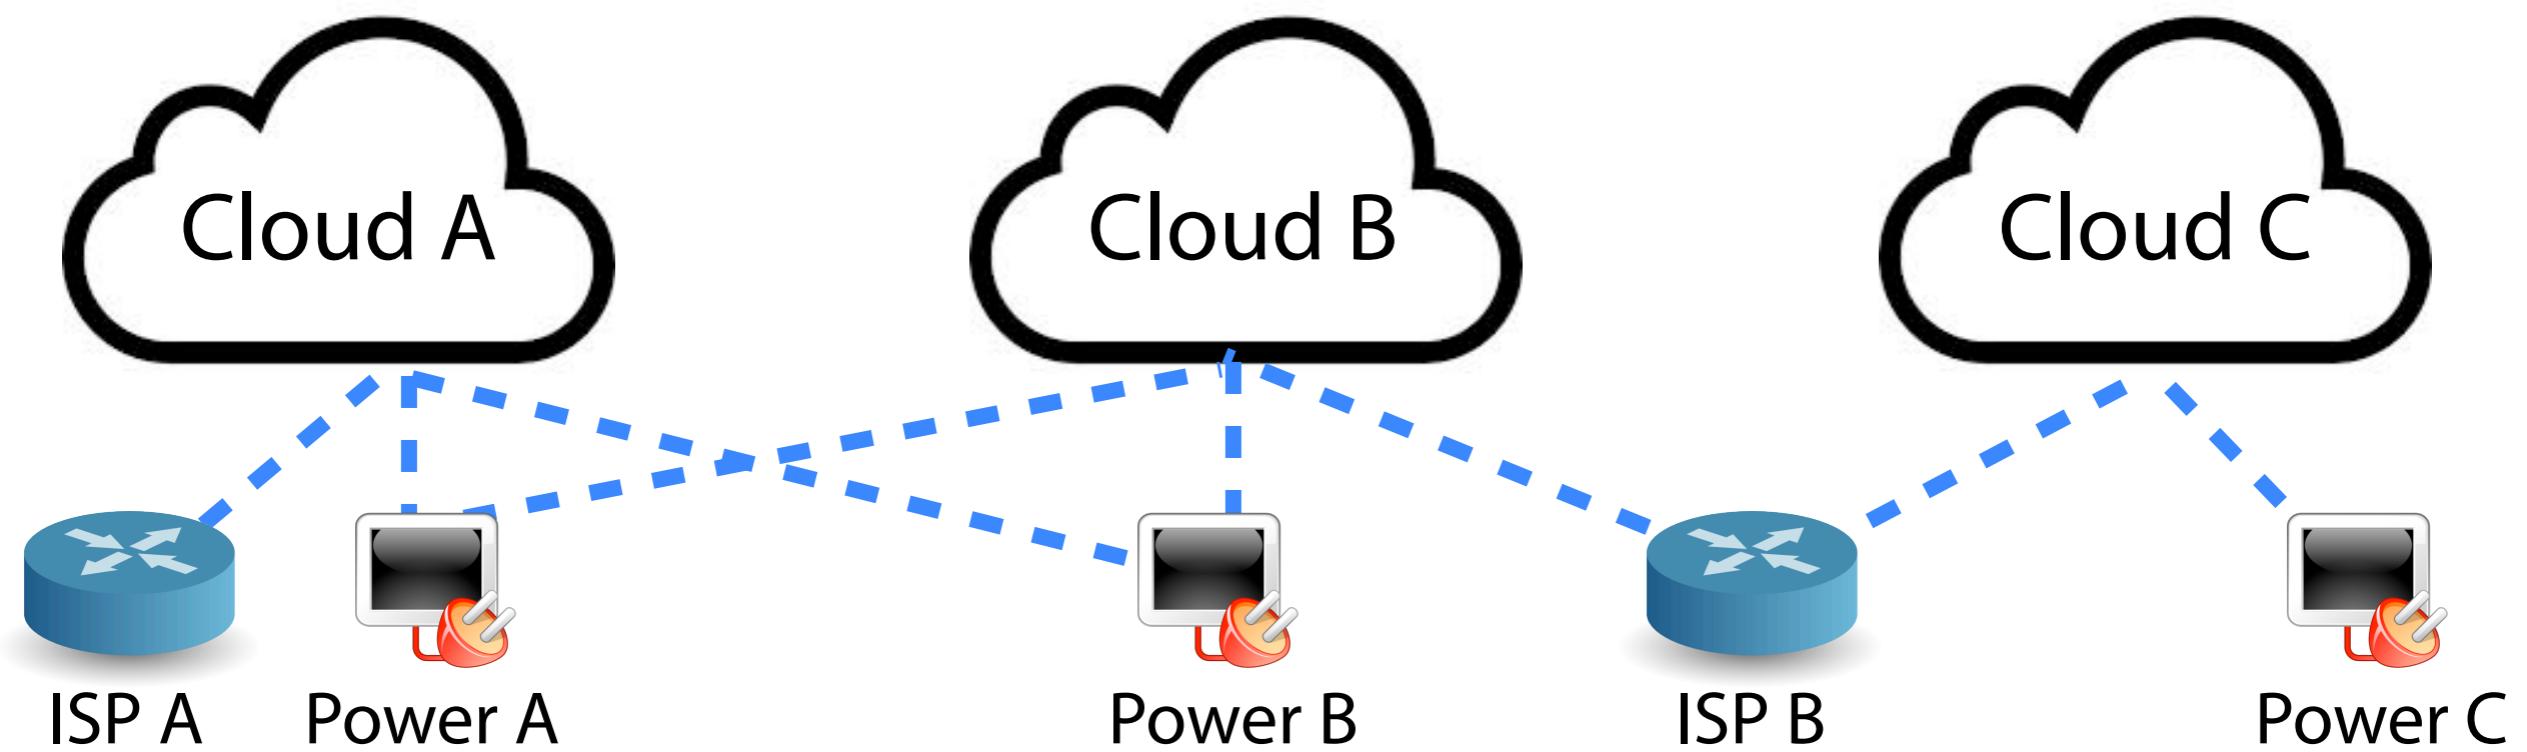
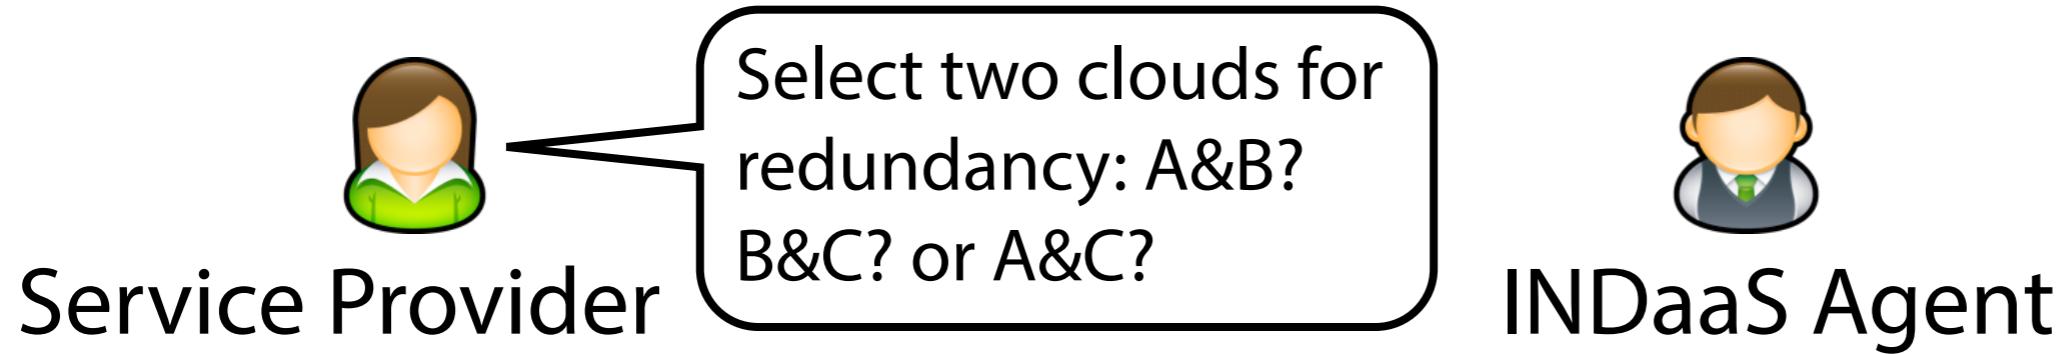
Select two clouds for redundancy: A&B?  
B&C? or A&C?

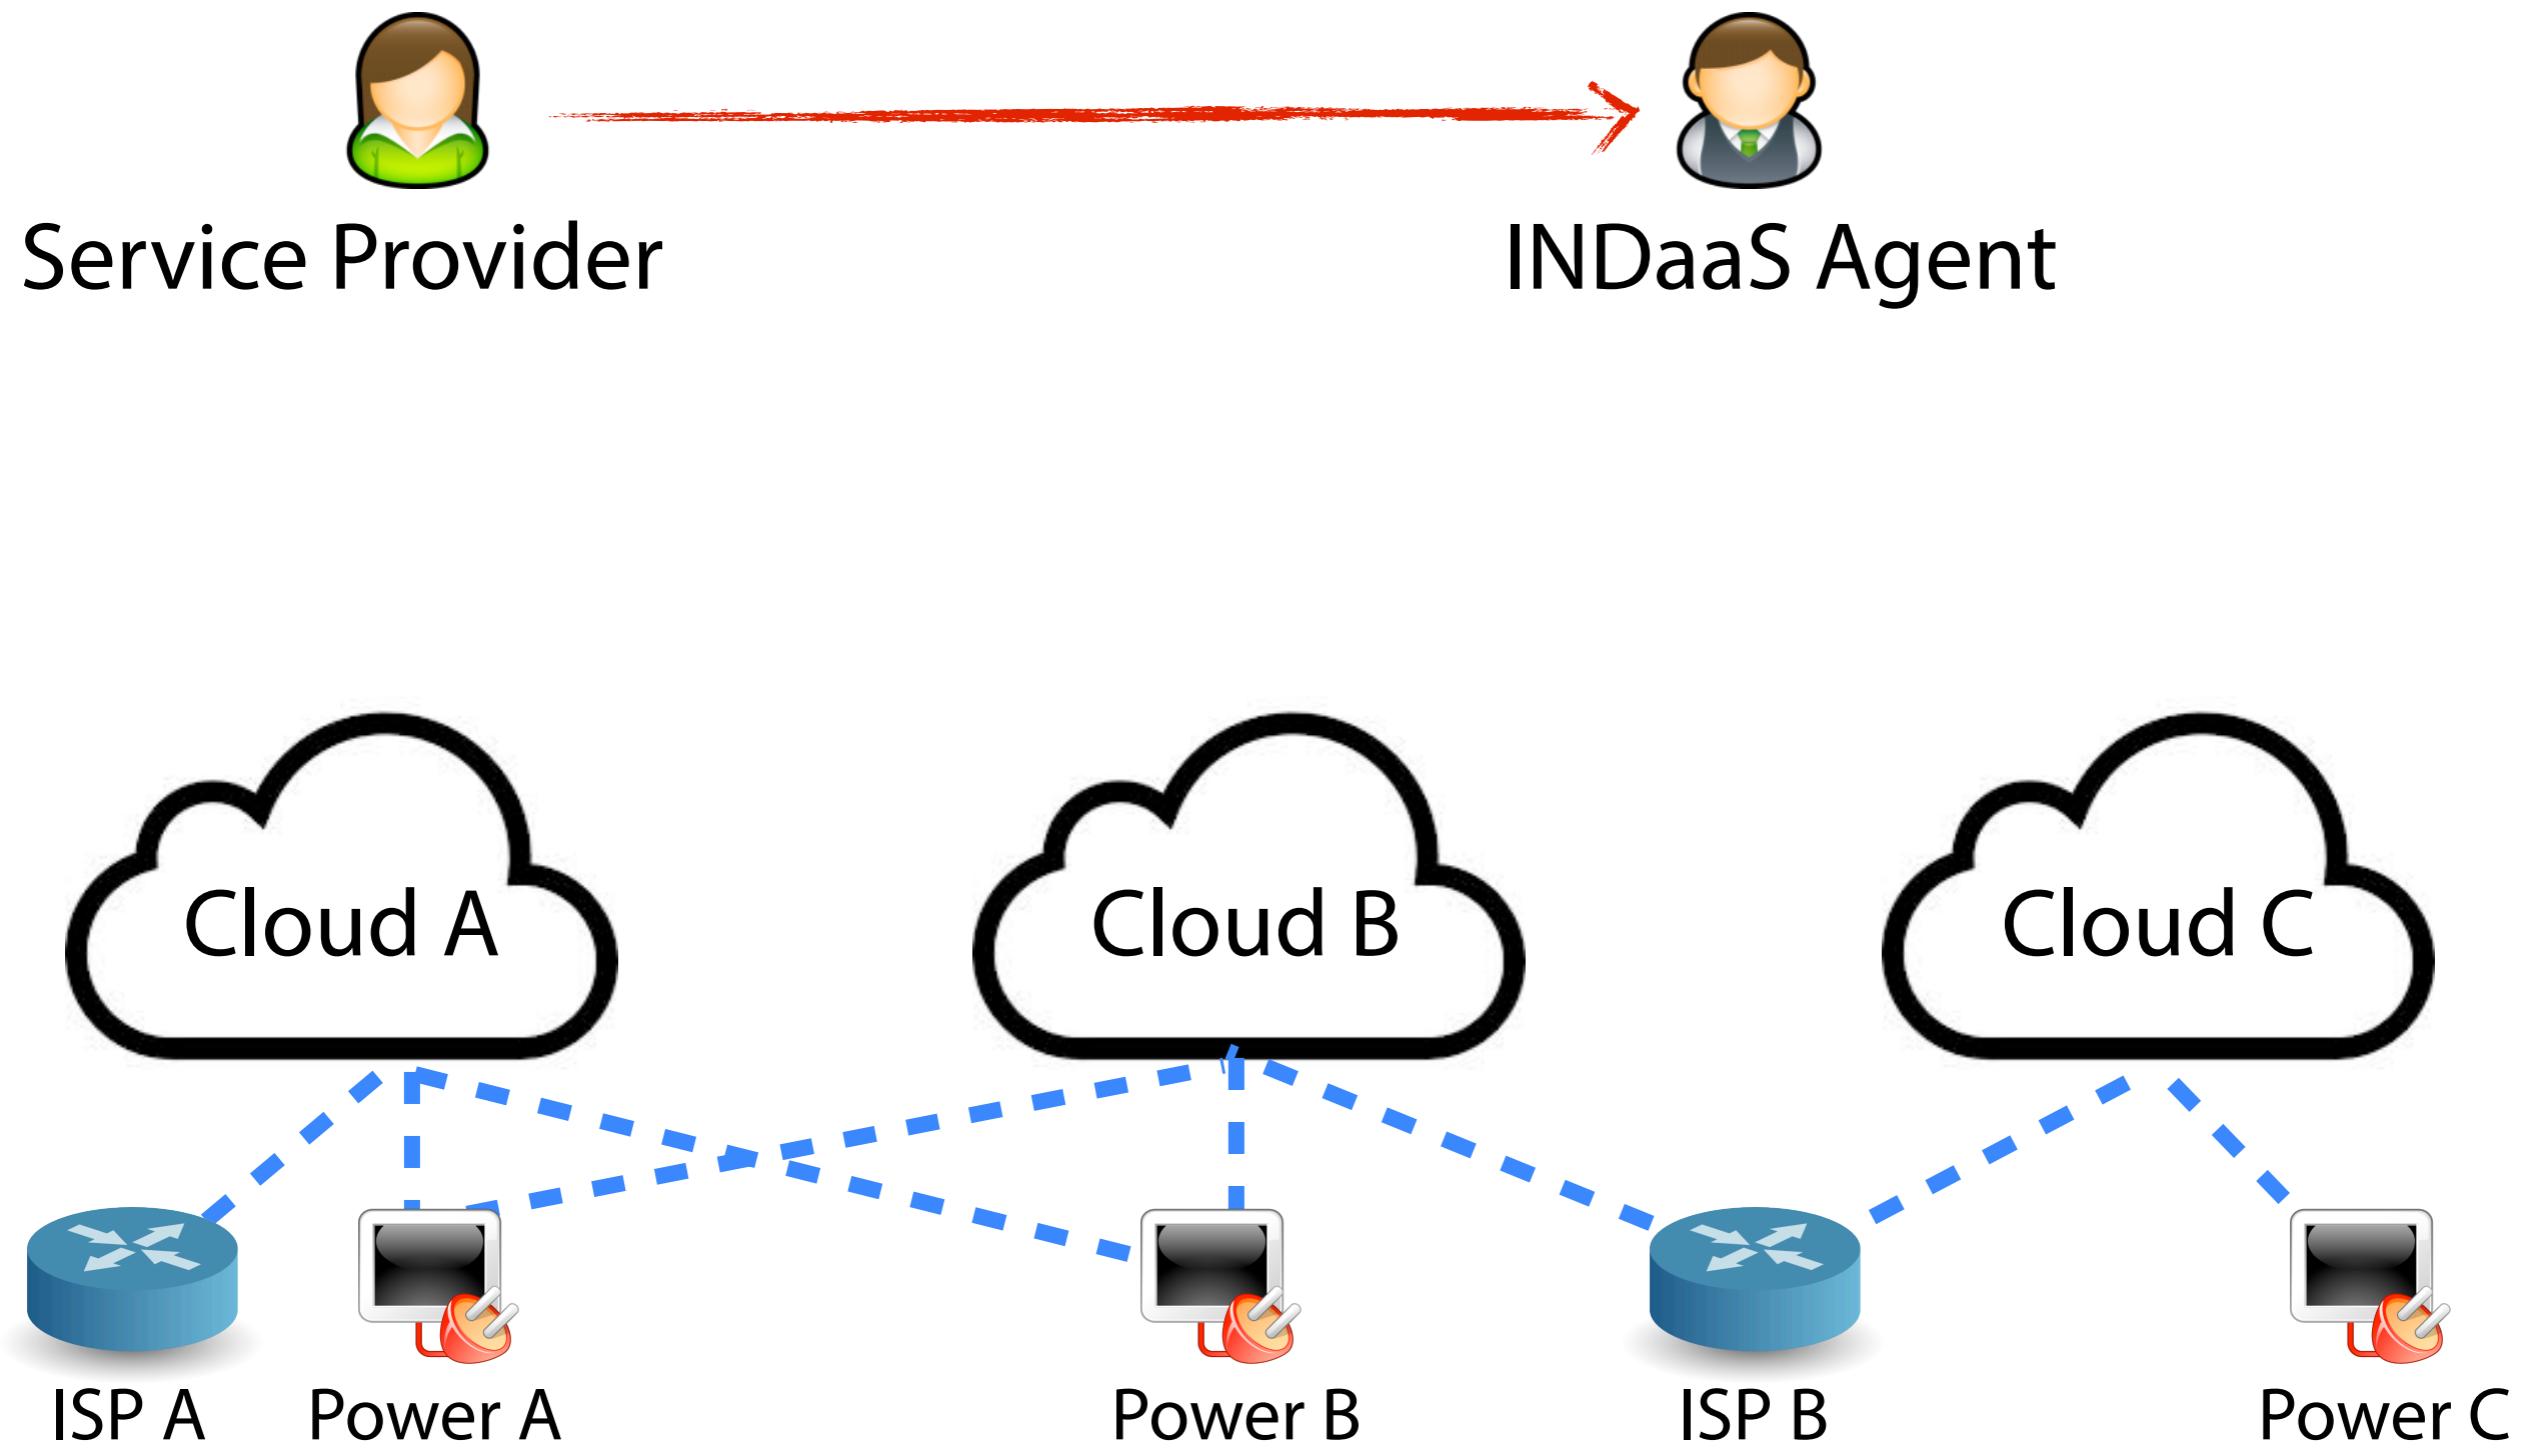


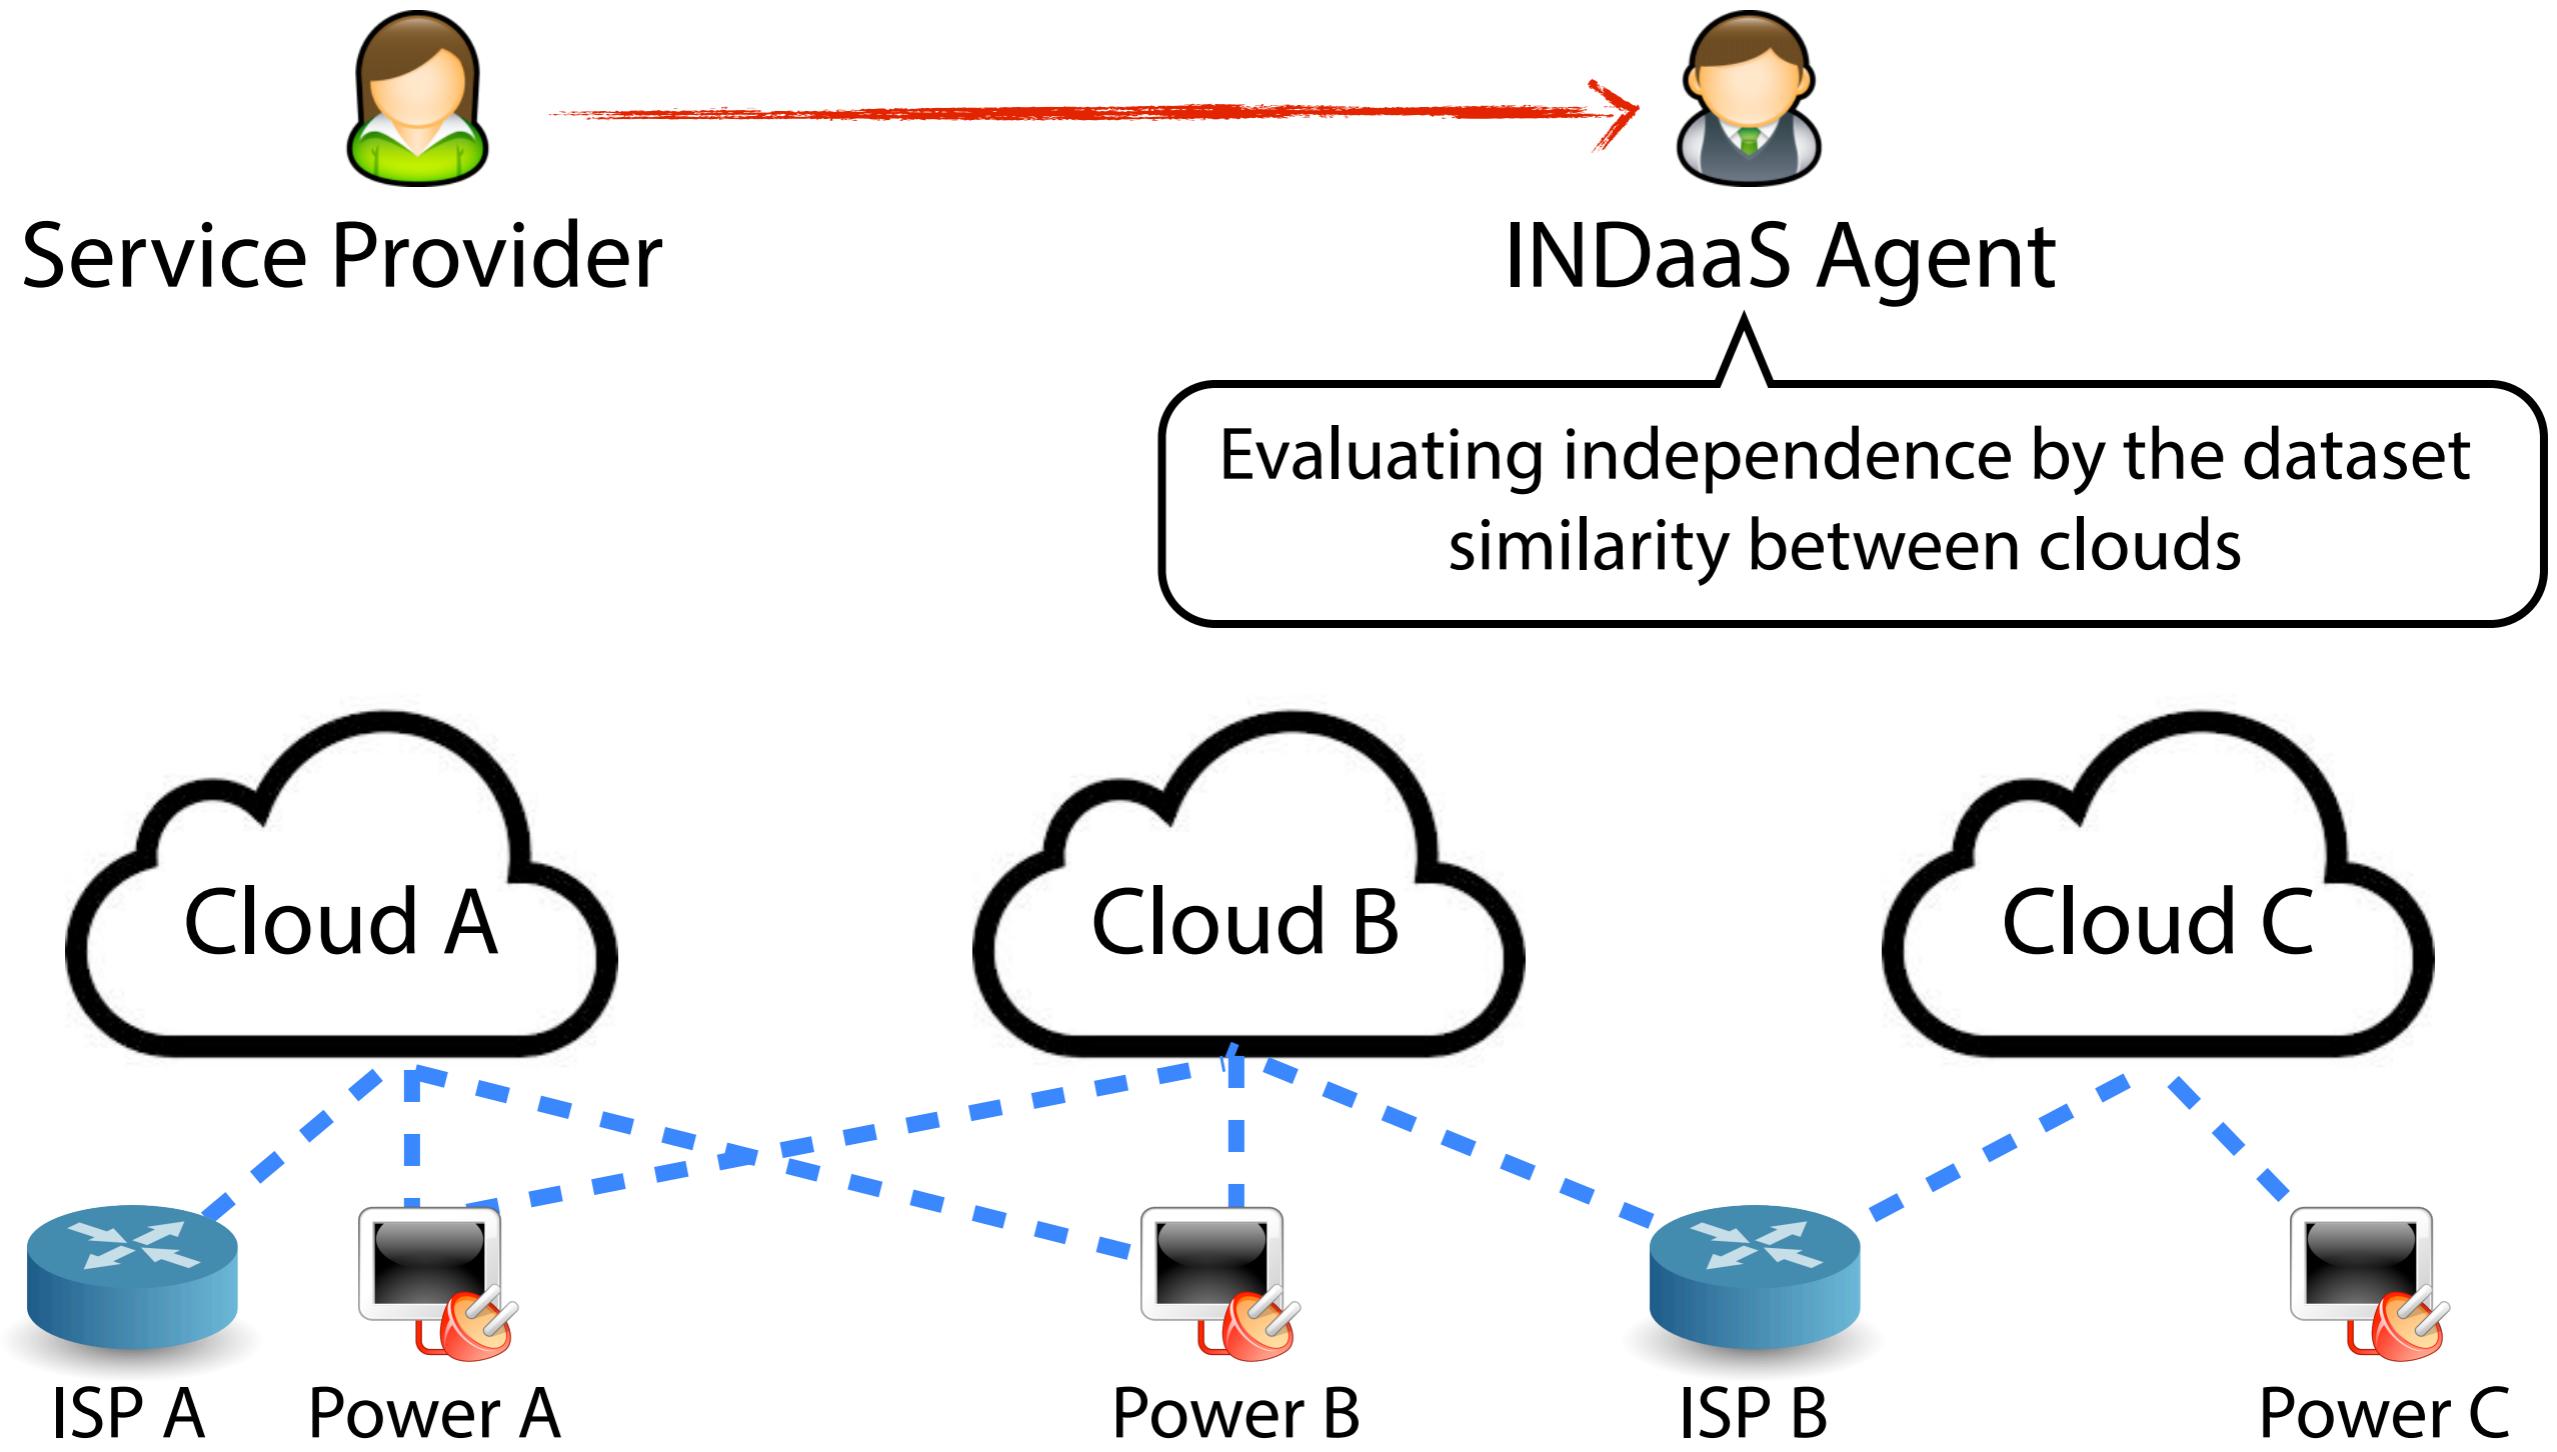
IND CLOUDS

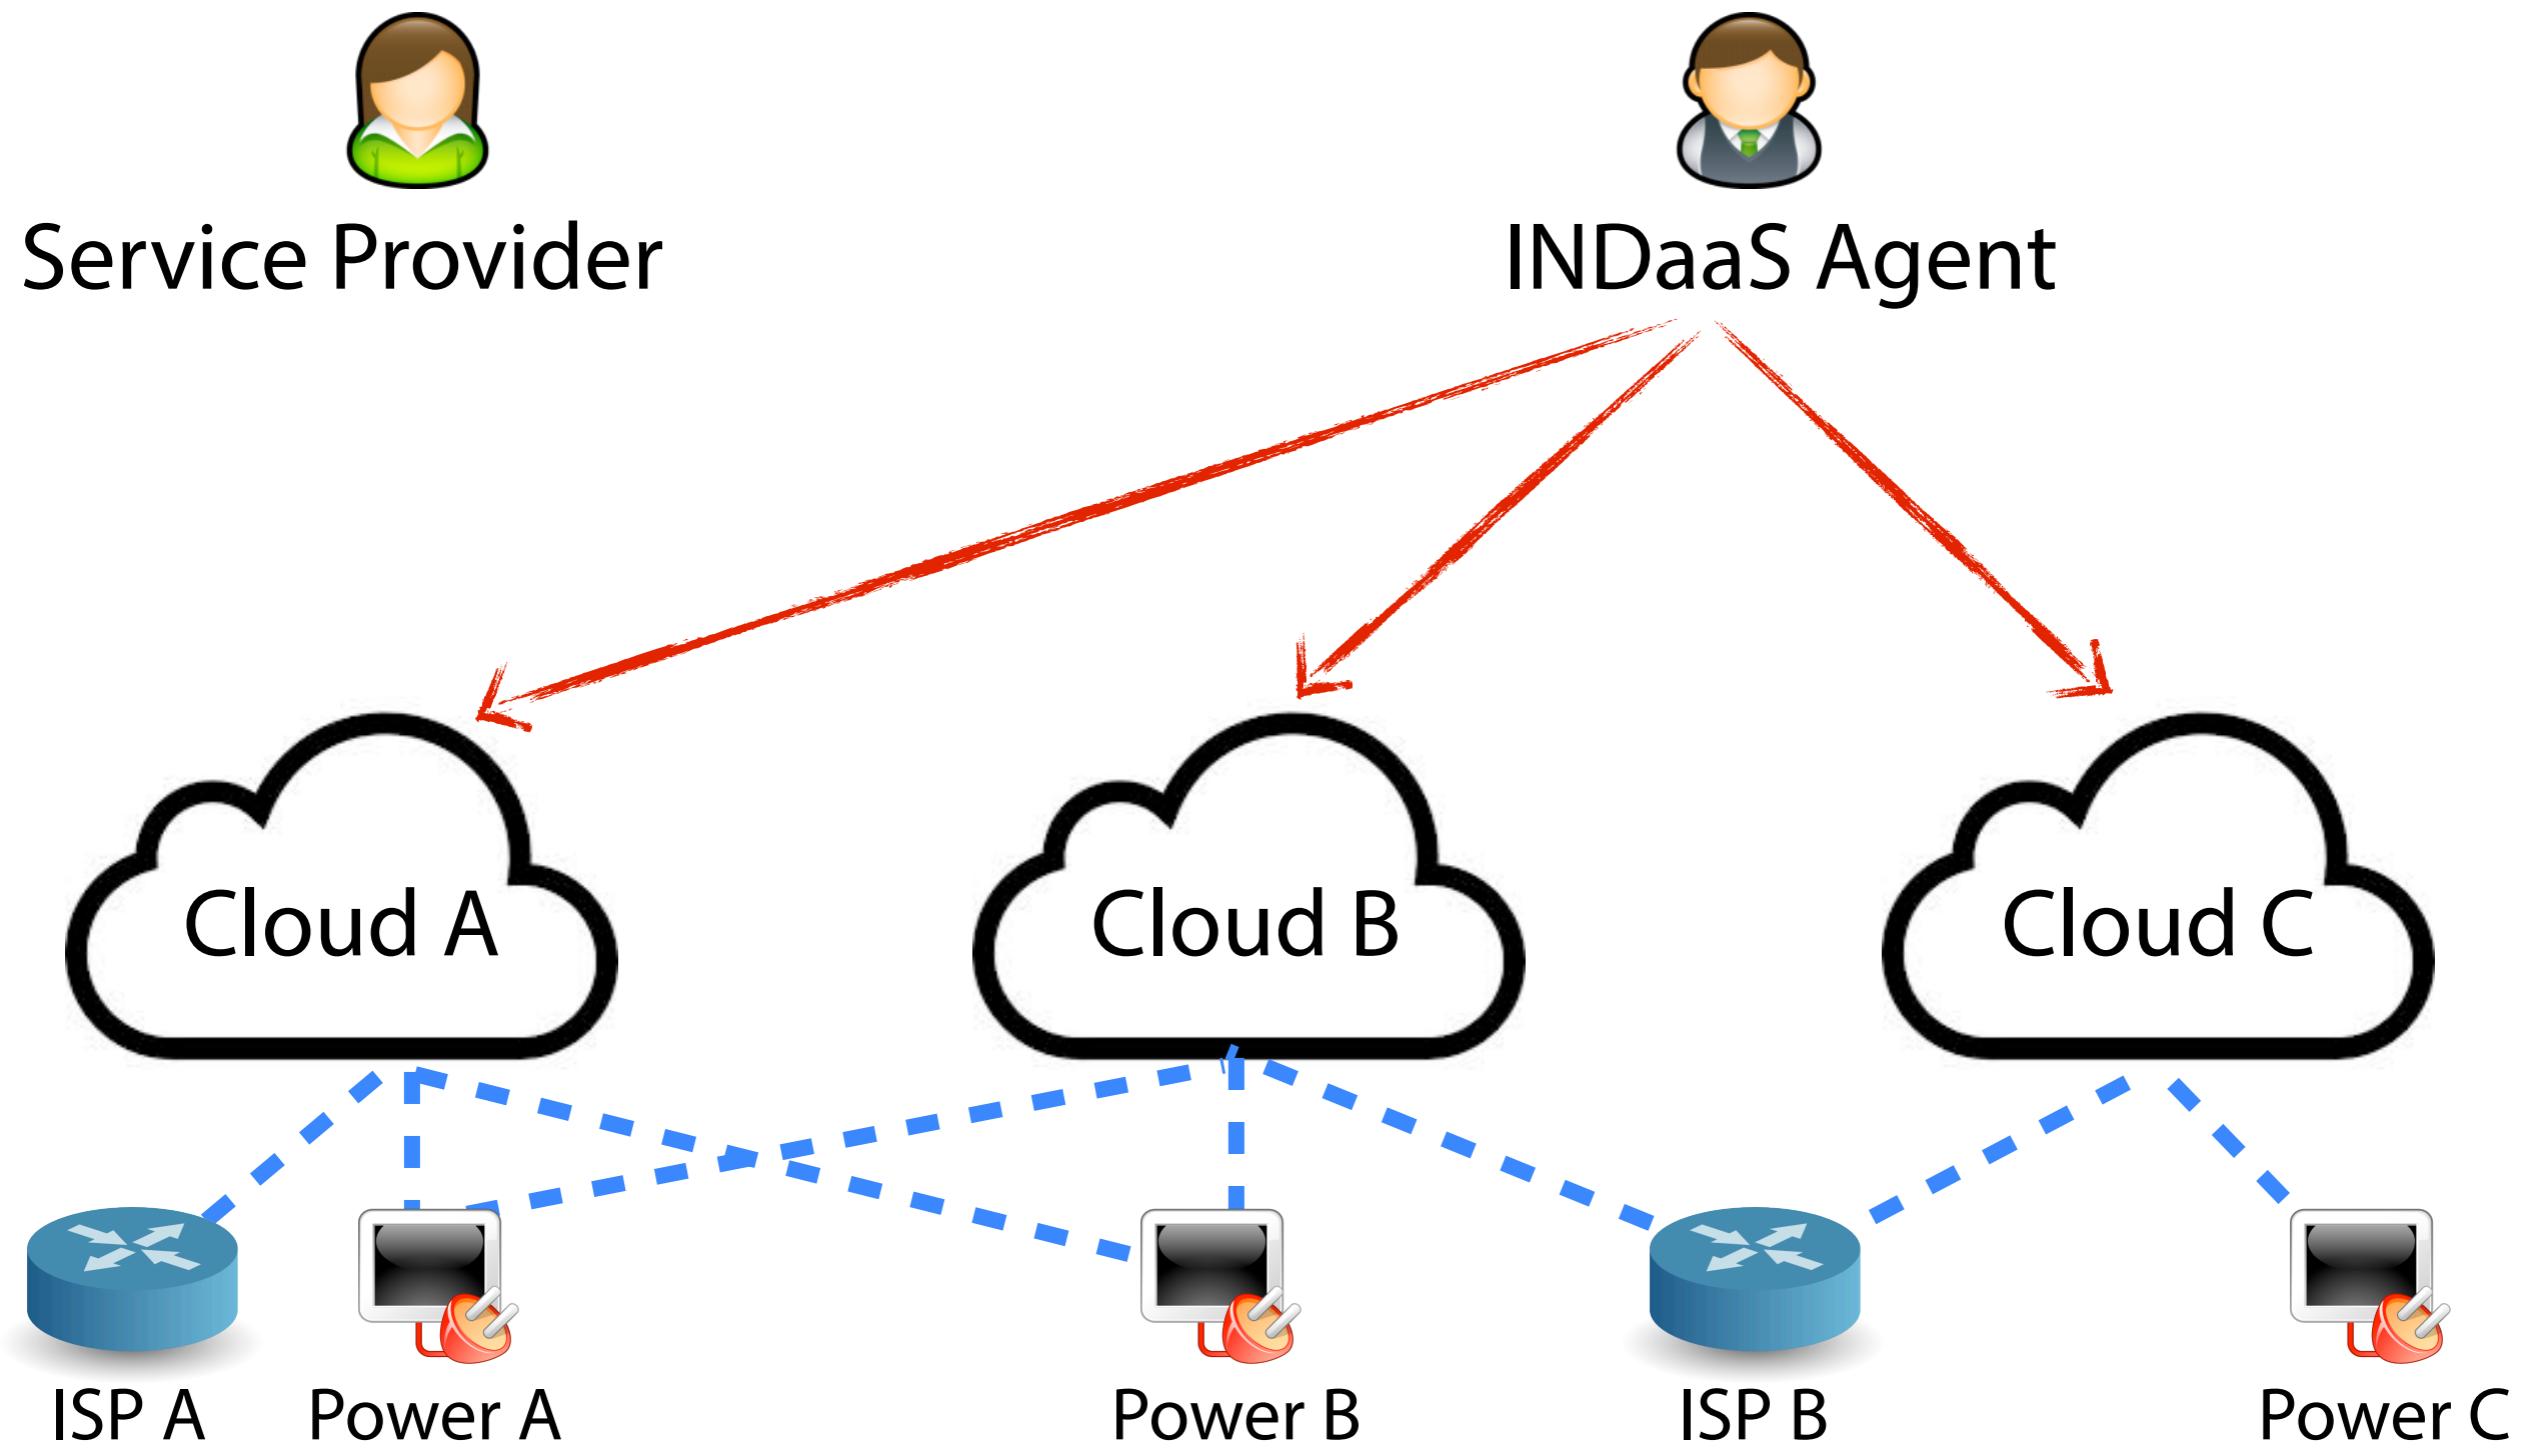
SMPC is hard to scale!  
[Xiao et al. CCSW'13]









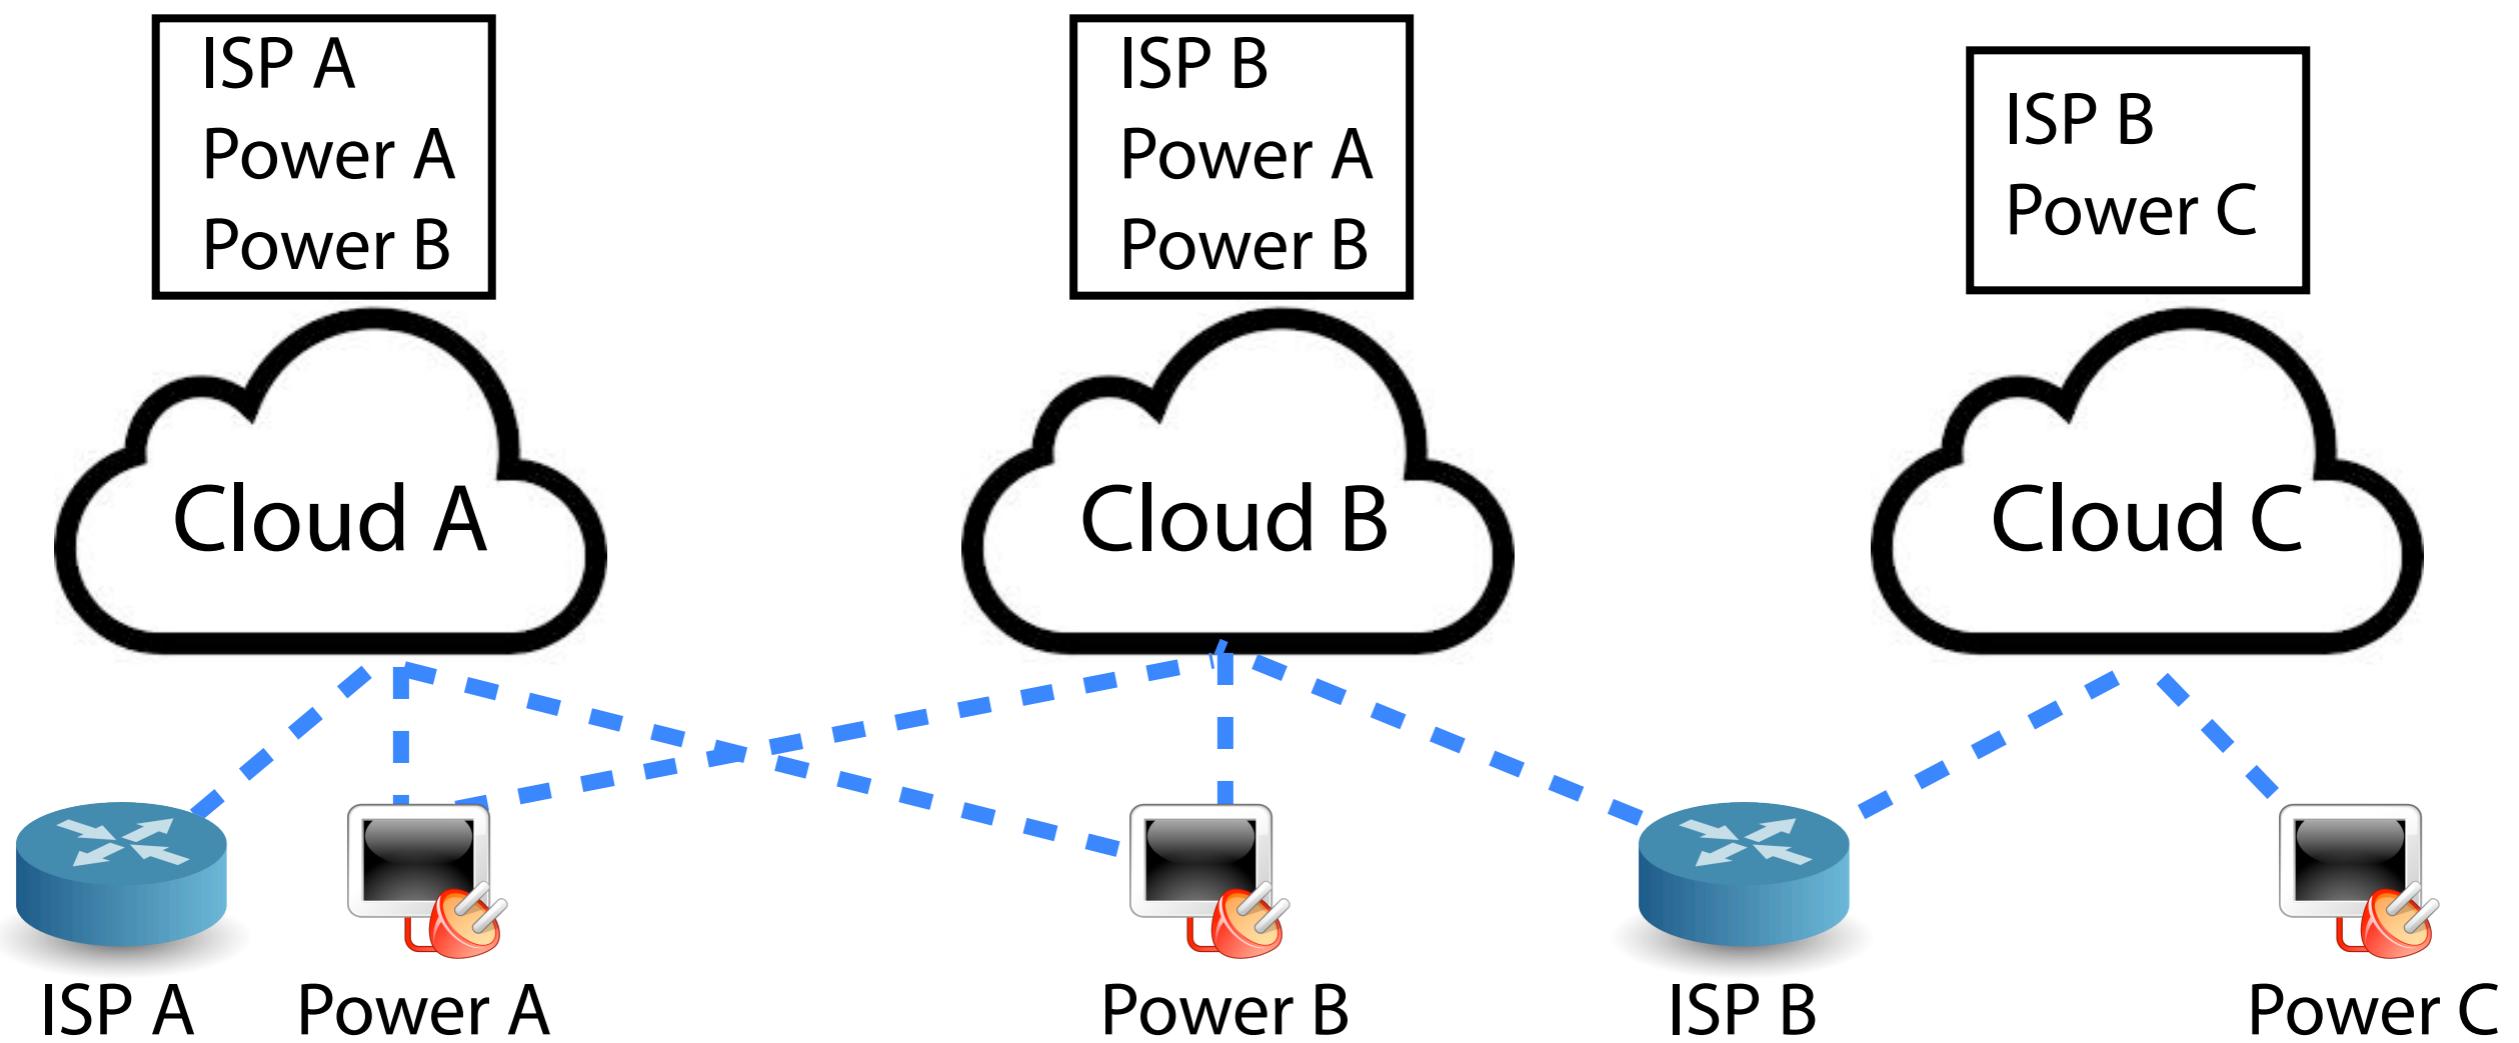




Service Provider



INDaaS Agent





Using Jaccard similarity to evaluate the independence of each redundancy configuration.

Cloud A

Cloud B

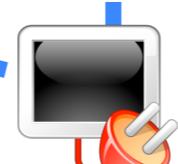
Cloud C



ISP A



Power A



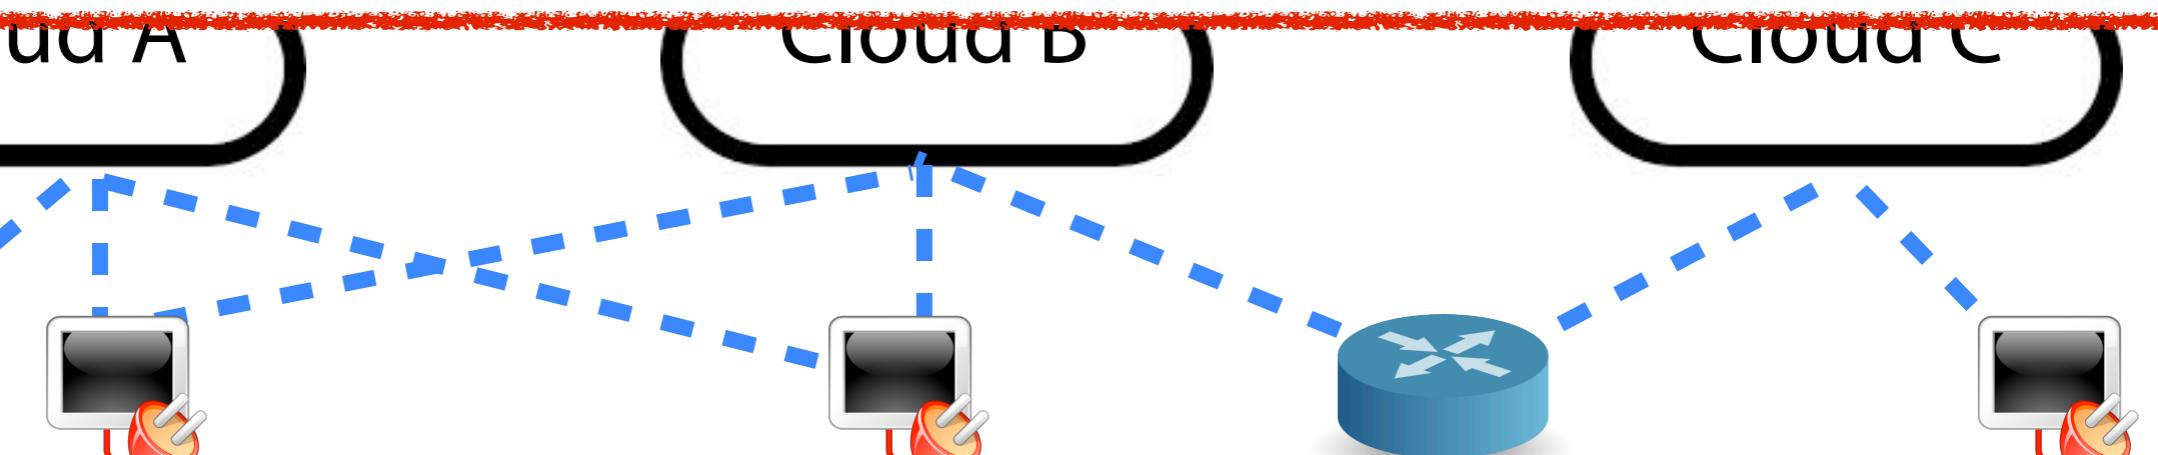
Power B



ISP B



Power C





$$J(S_1, S_2, \dots, S_n) = \frac{|S_1 \cap S_2 \cap \dots \cap S_n|}{|S_1 \cup S_2 \cup \dots \cup S_n|}$$

Cloud A

Cloud B

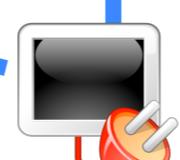
Cloud C



ISP A



Power A



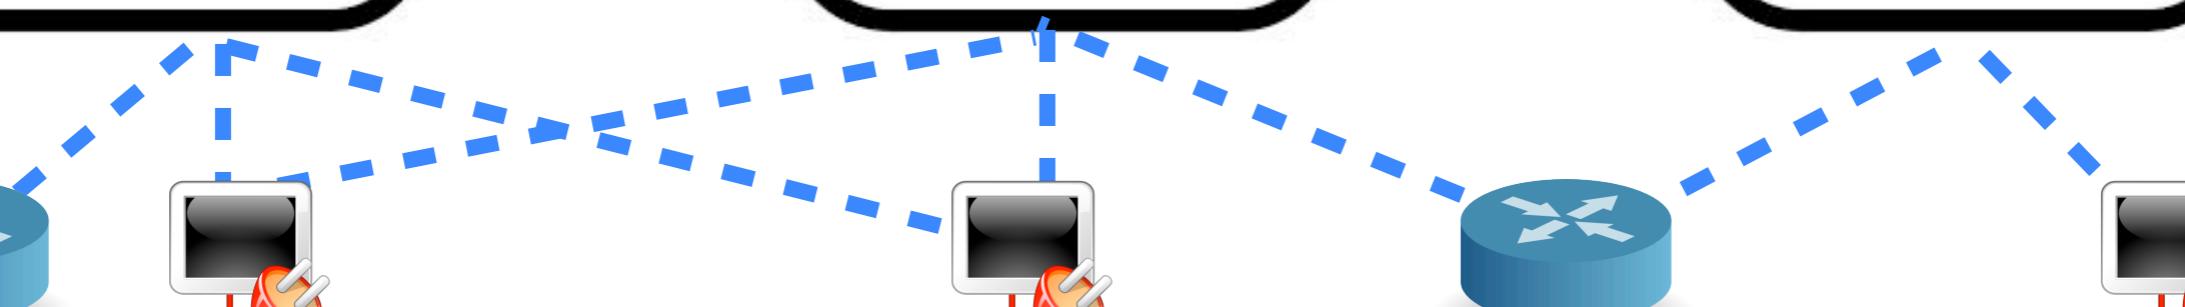
Power B



ISP B



Power C

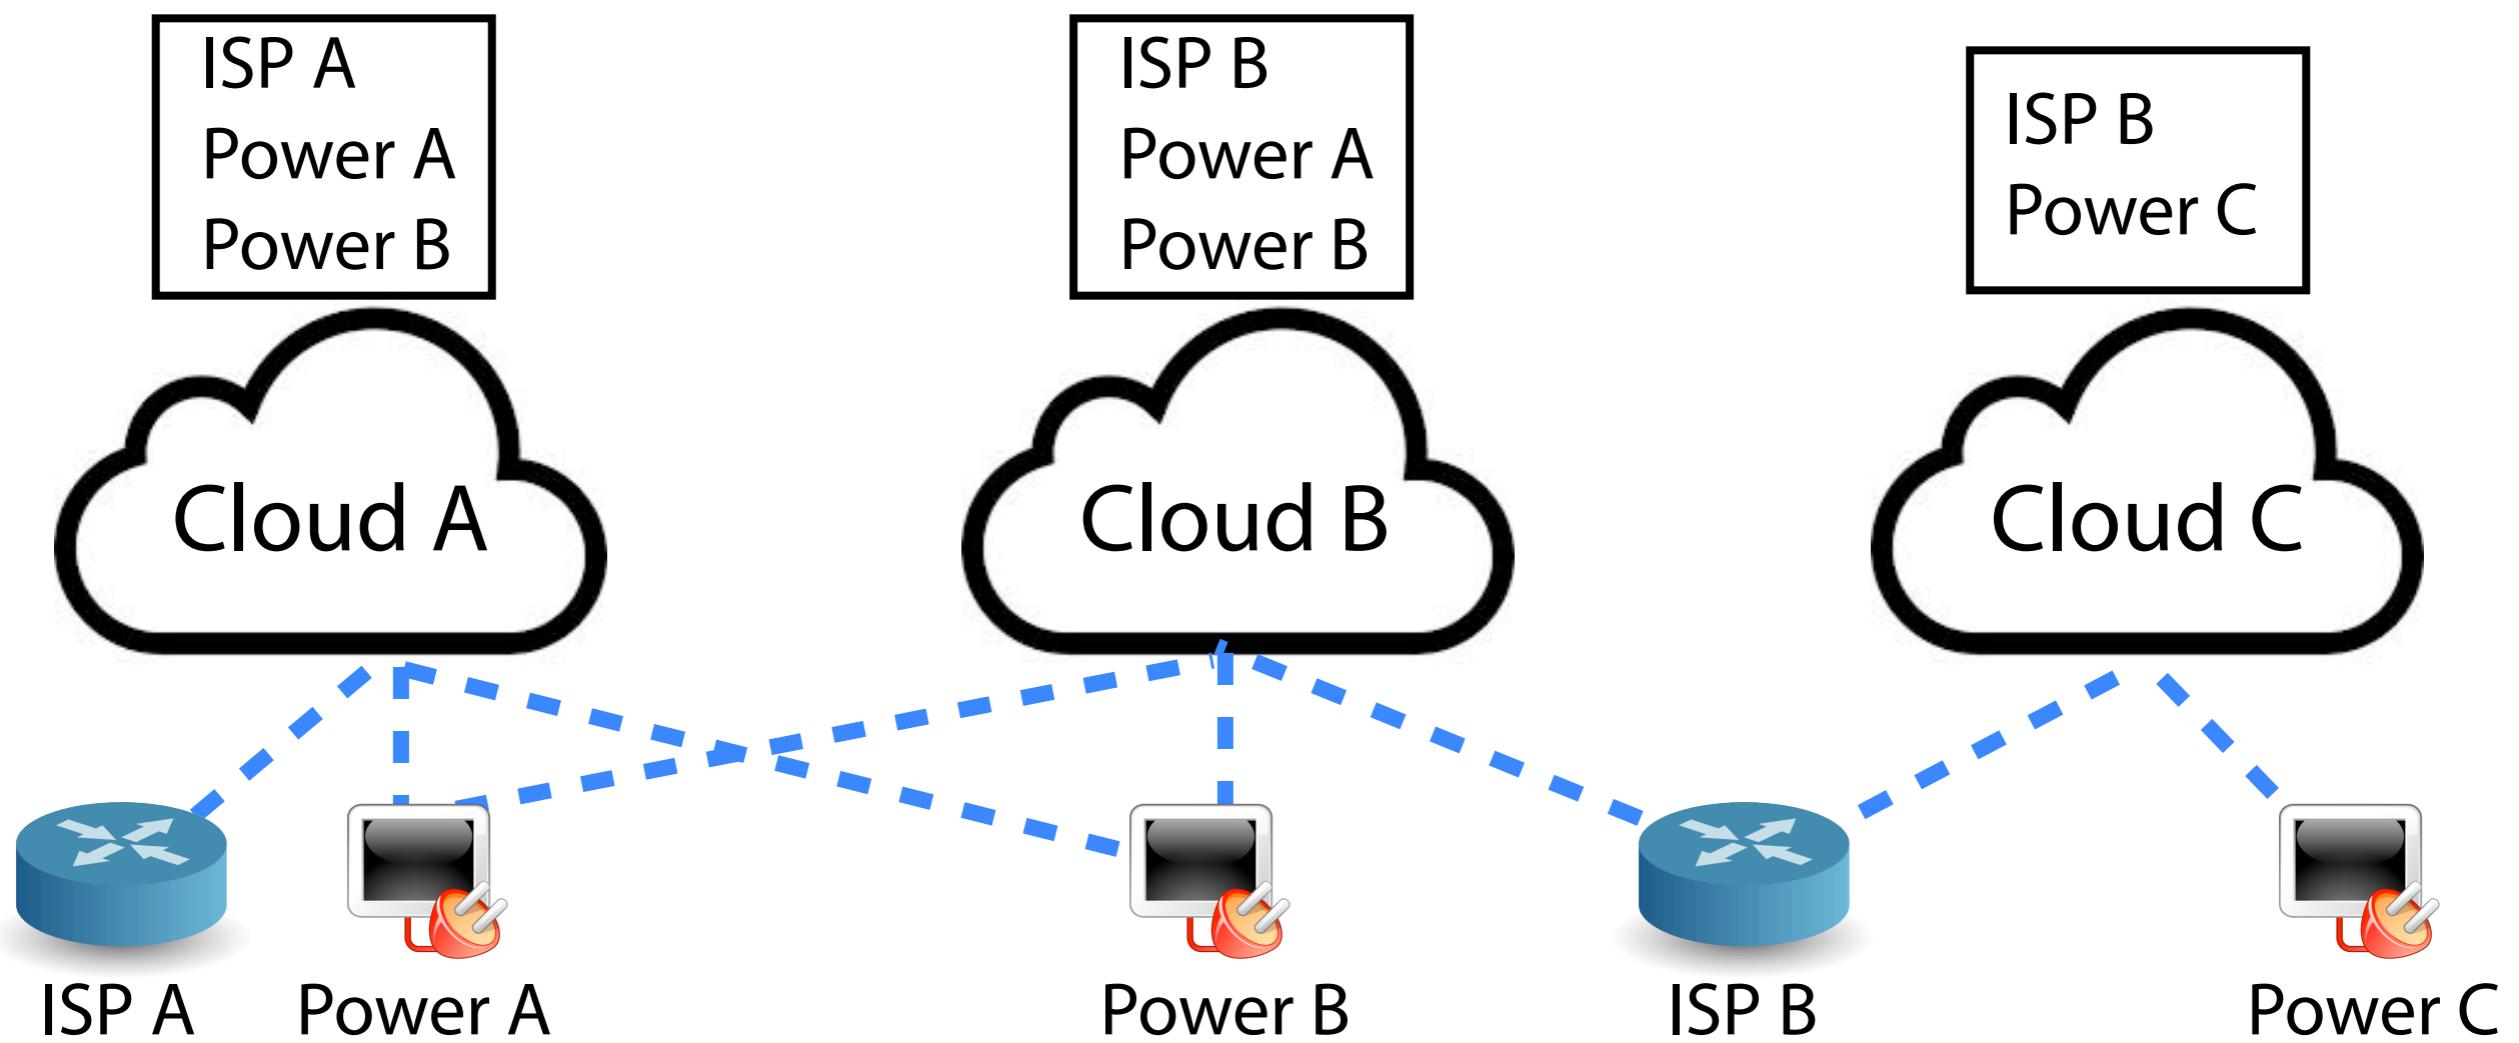




Service Provider



INDaaS Agent



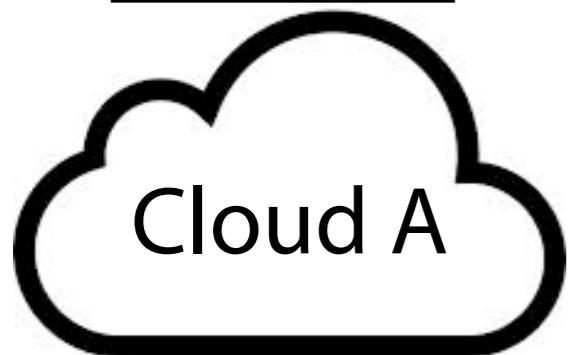


Service Provider

ISP A  
Power A  
Power B

$$\begin{matrix} 0 \\ \cup \\ =2 \\ =4 \end{matrix}$$

ISP B  
Power A  
Power B



ISP A

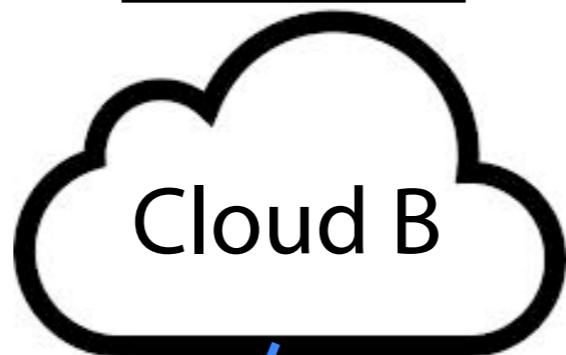


Power A

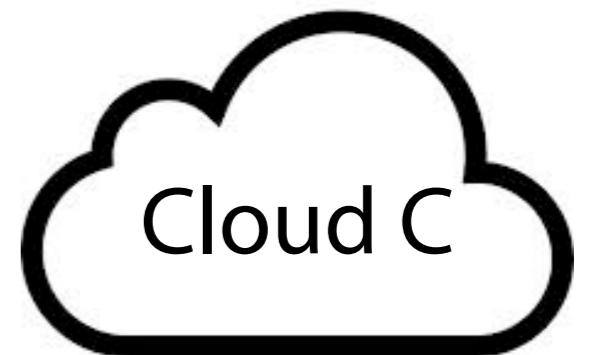


INDaaS Agent

ISP B  
Power C



Power B



ISP B



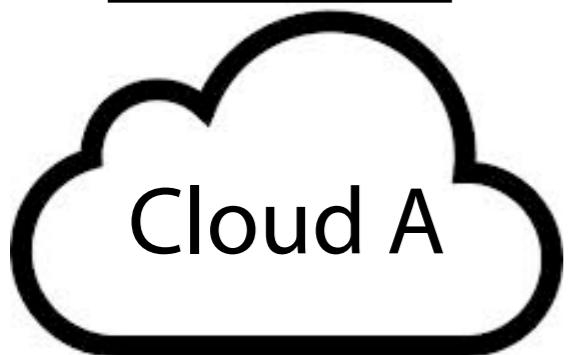
Power C



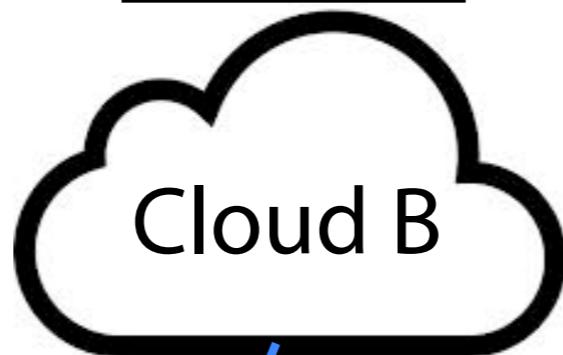
Service Provider

ISP A  
Power A  
Power B

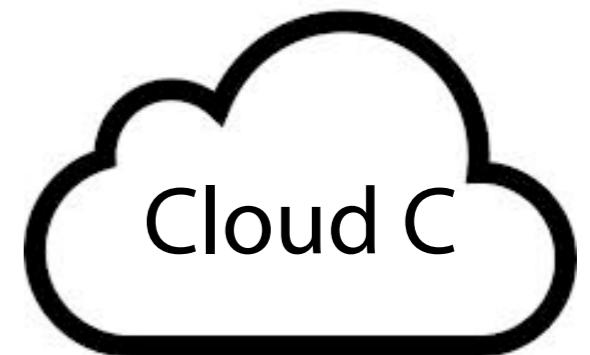
$$\begin{matrix} 0 \\ \cup \\ =2 \\ =4 \\ J = 2/4 \end{matrix}$$



ISP B  
Power A  
Power B



ISP B  
Power C



ISP A



Power A



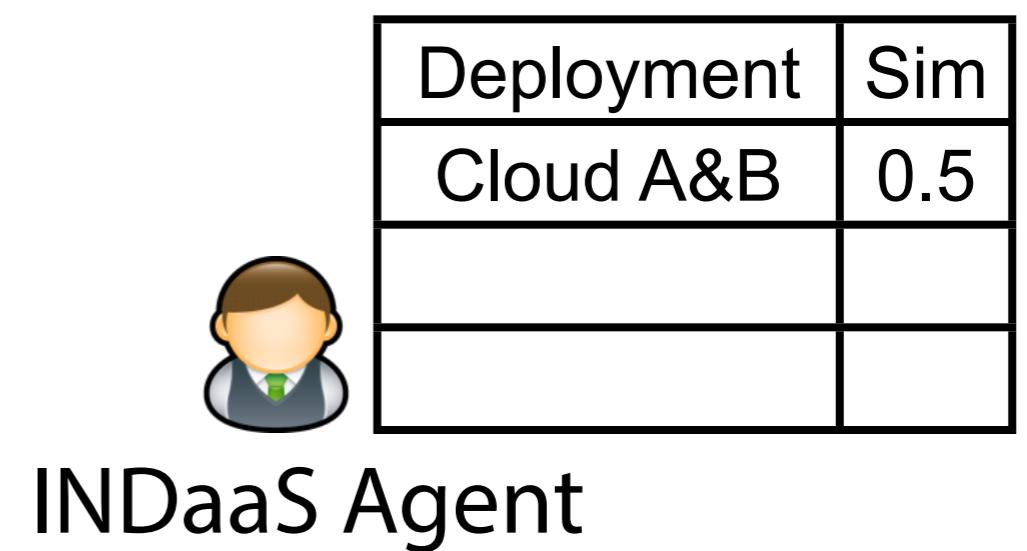
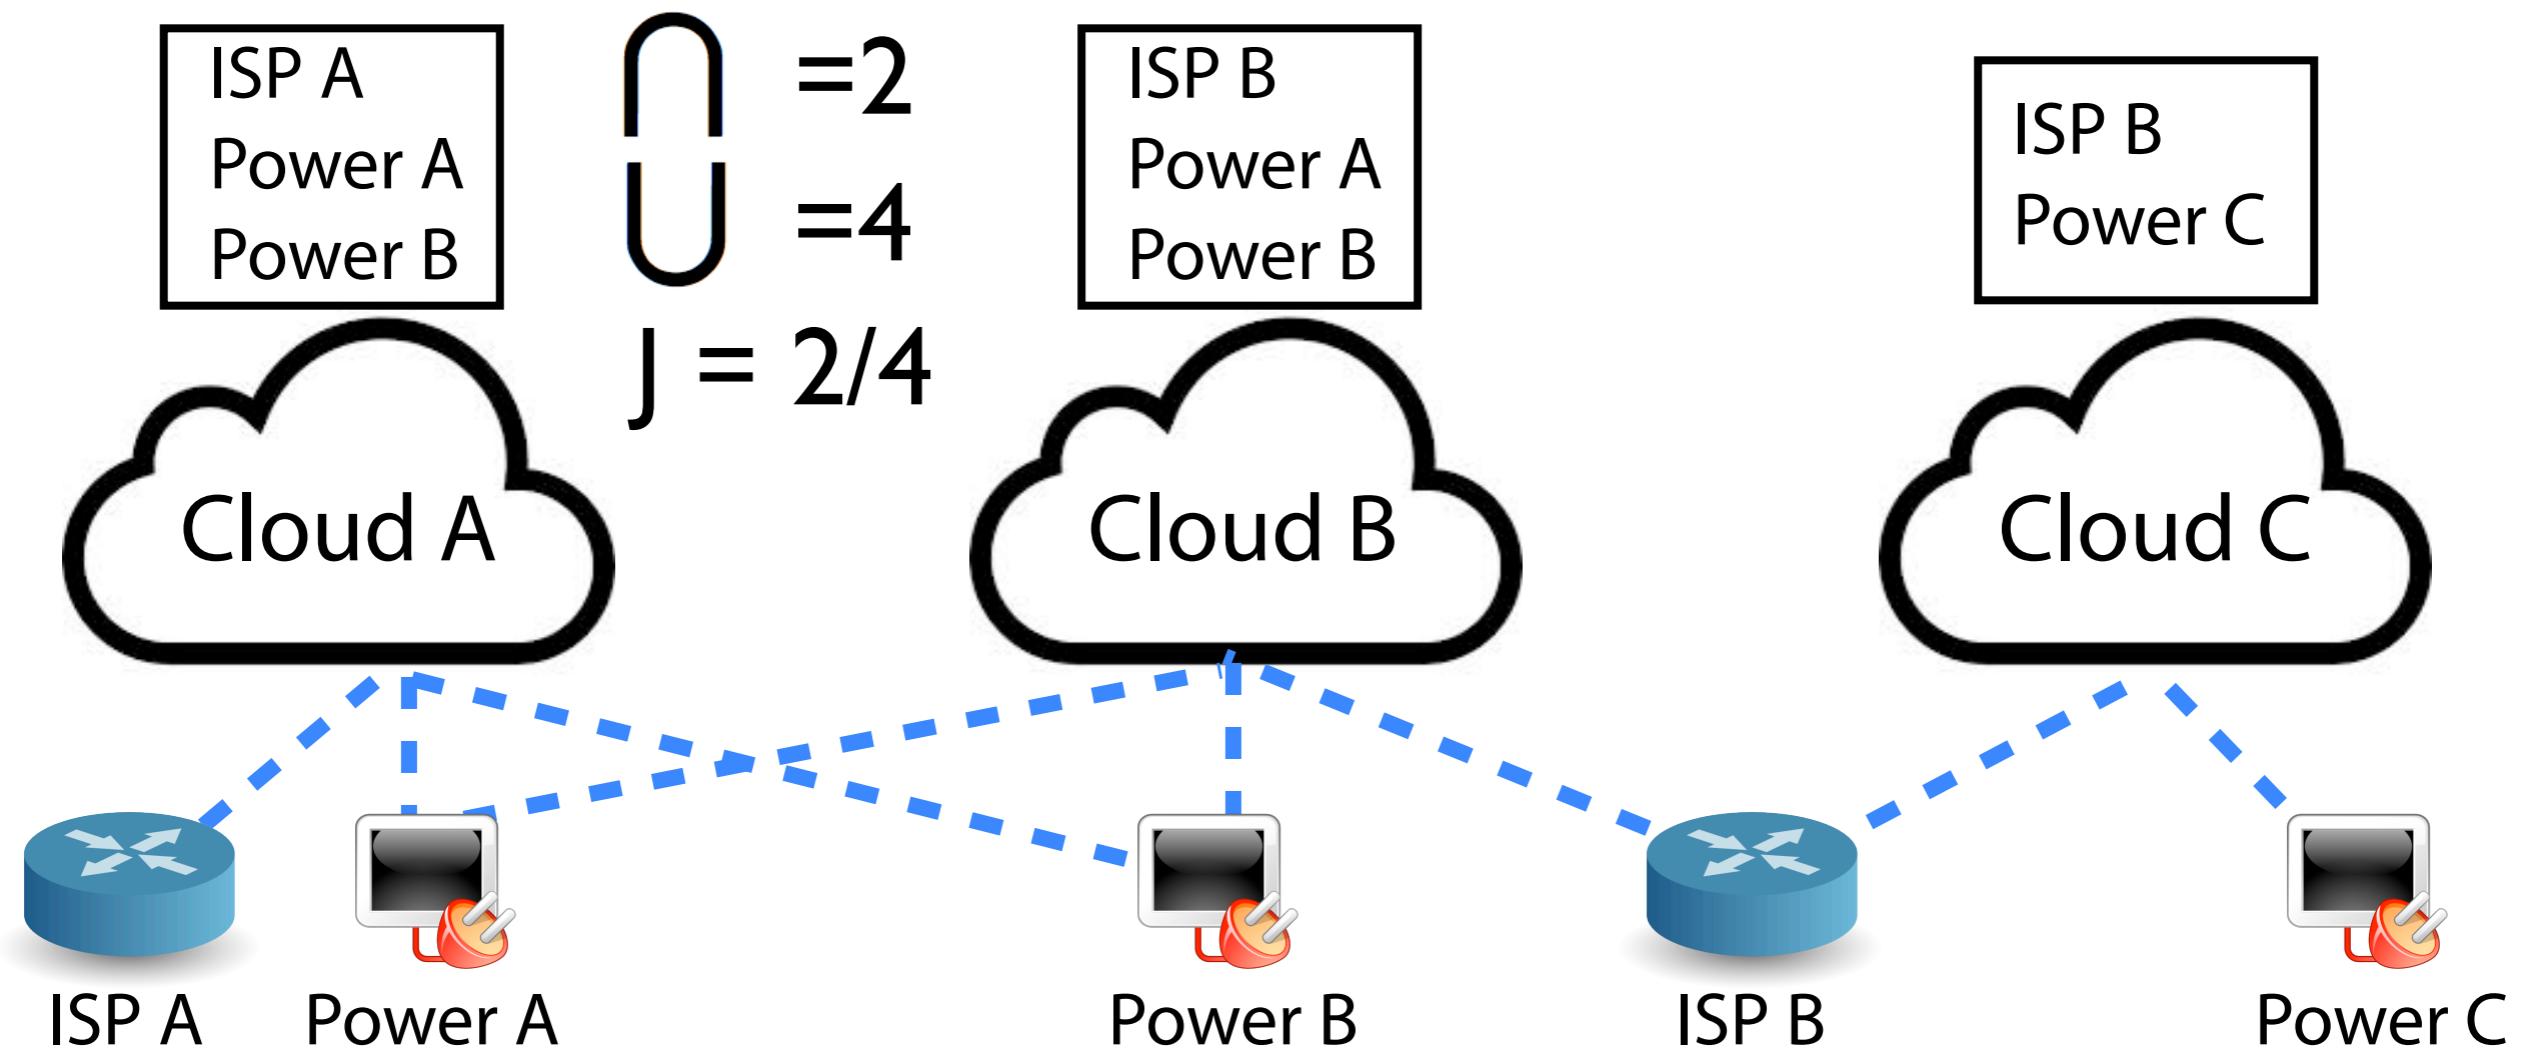
Power B



ISP B

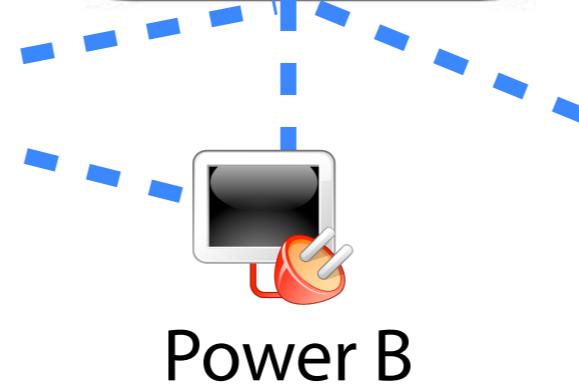
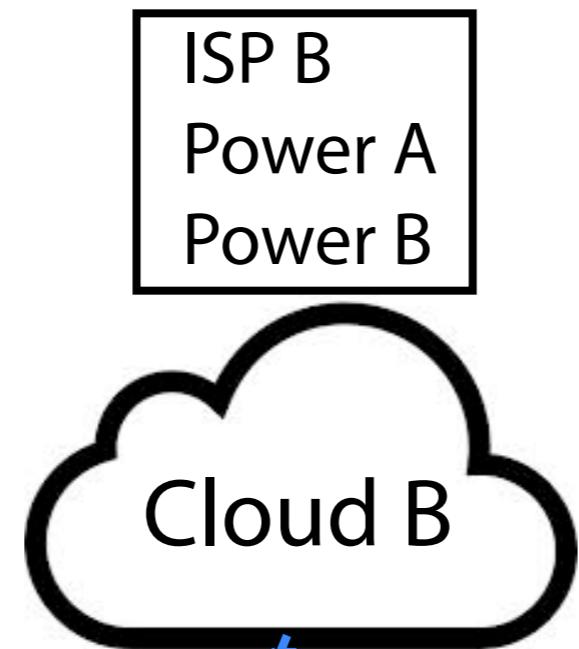
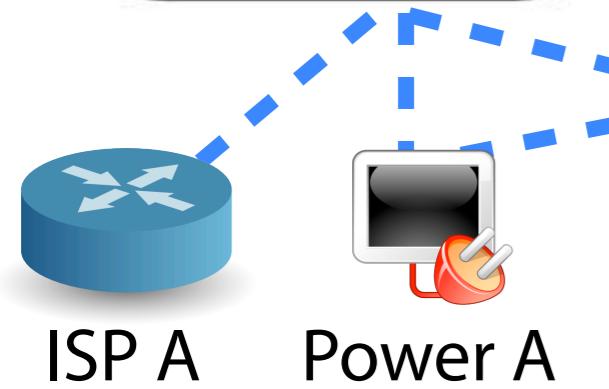
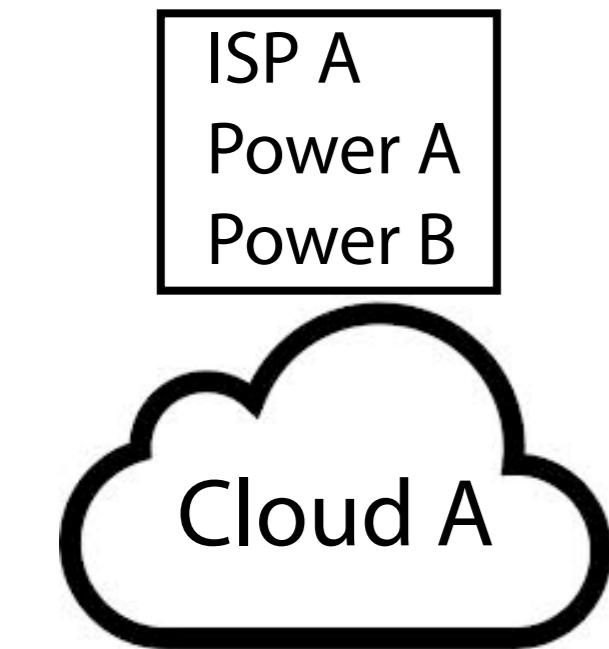


Power C





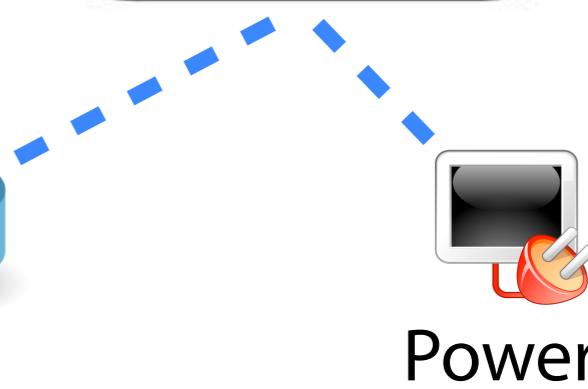
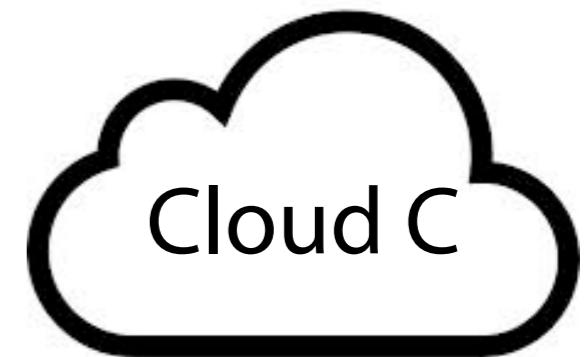
## Service Provider



| Deployment | Sim  |
|------------|------|
| Cloud A&B  | 0.5  |
| Cloud B&C  | 0.25 |
|            |      |

## INDaaS Agent

$$\begin{aligned} 0 &= I \\ &= 4 \\ J &= I/4 \end{aligned}$$



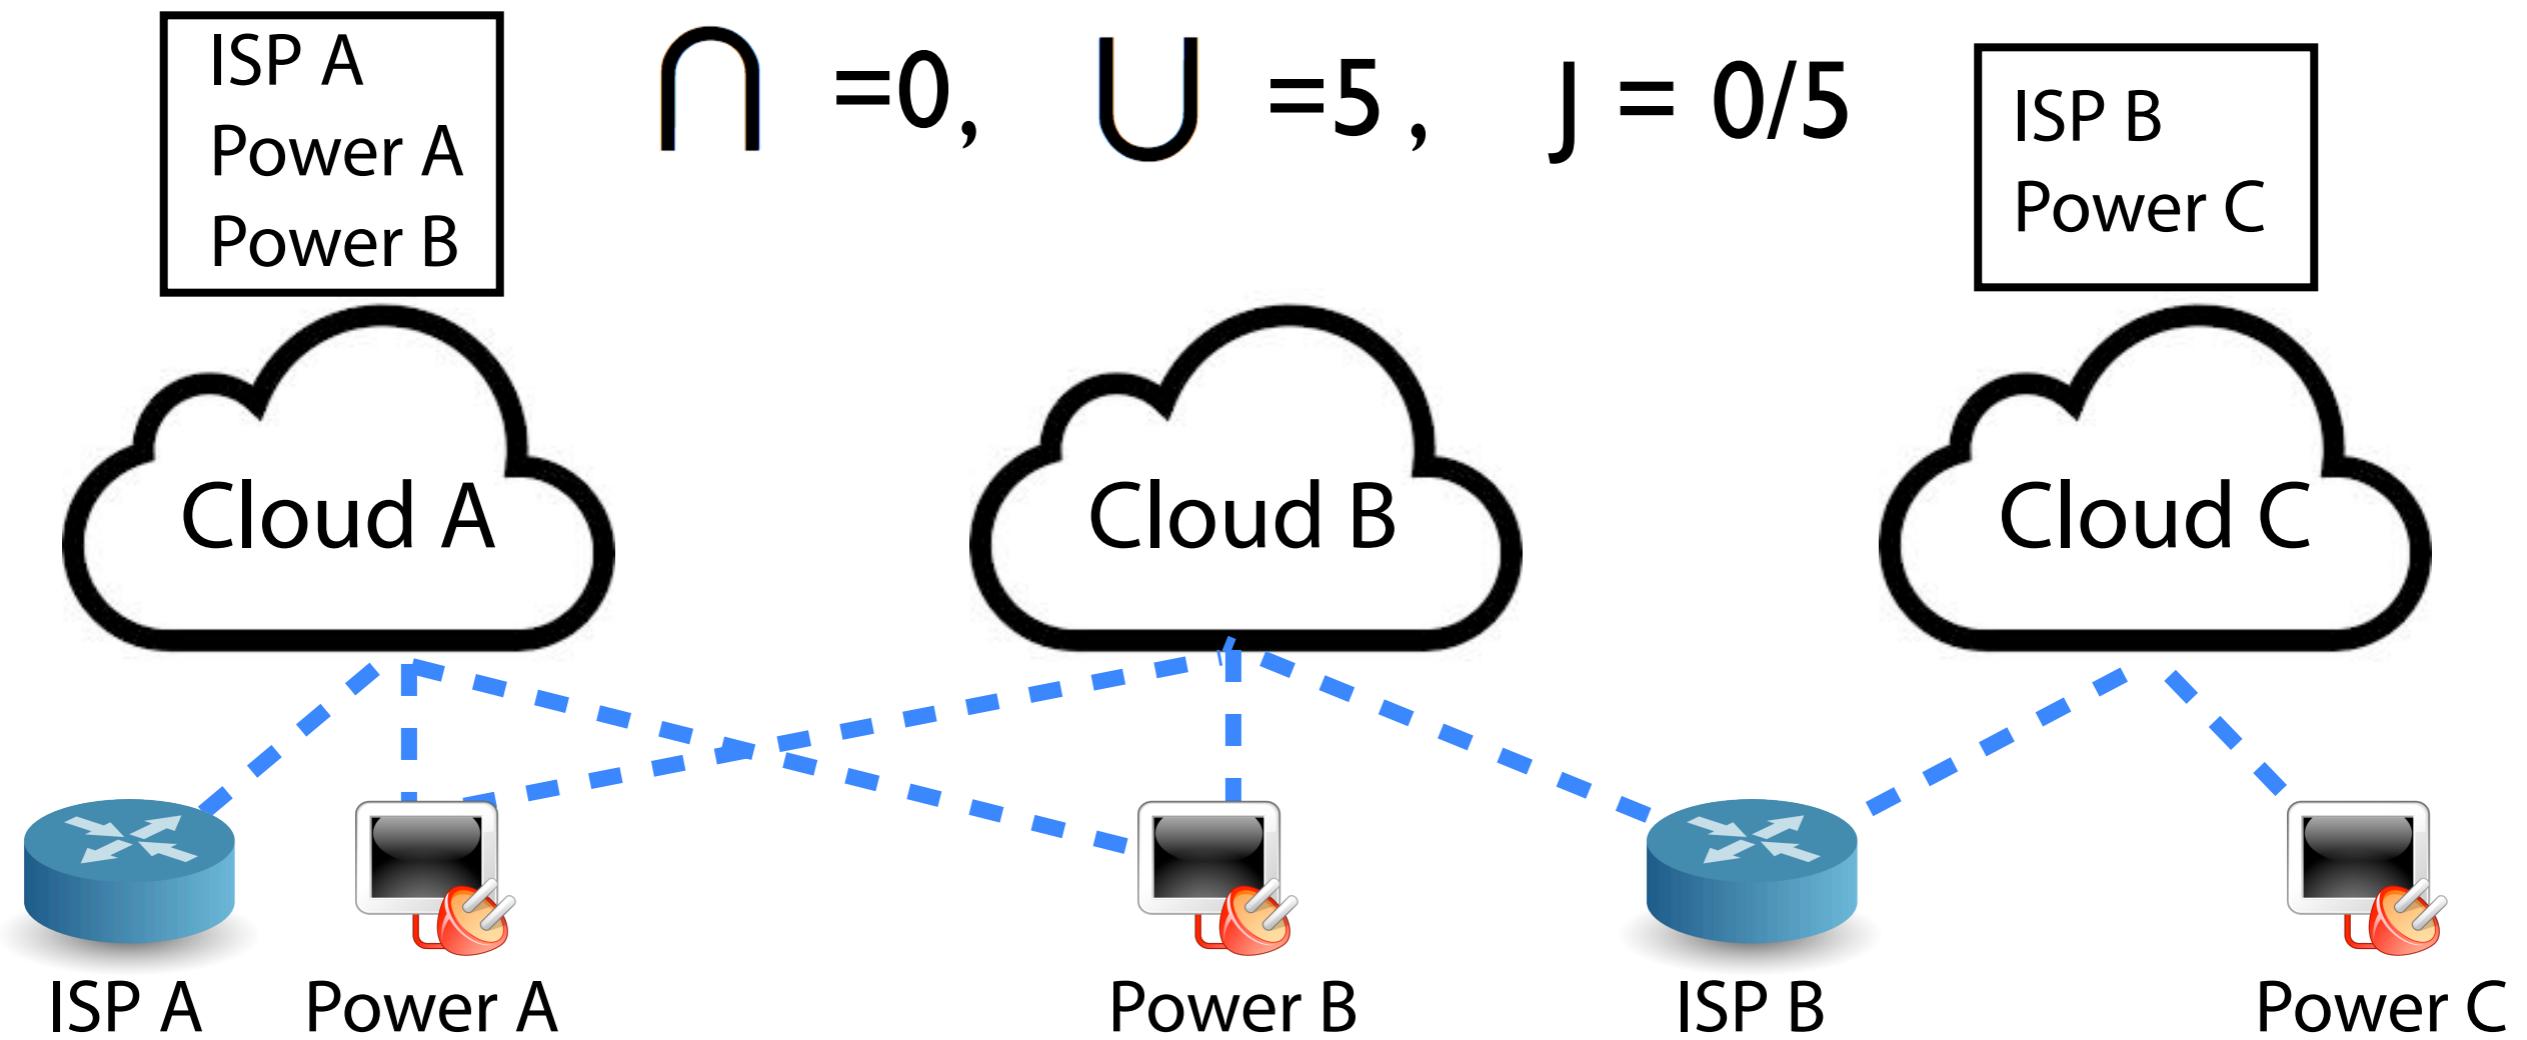


Service Provider



INDaaS Agent

| Deployment | Sim  |
|------------|------|
| Cloud A&B  | 0.5  |
| Cloud B&C  | 0.25 |
| Cloud A&C  | 0    |



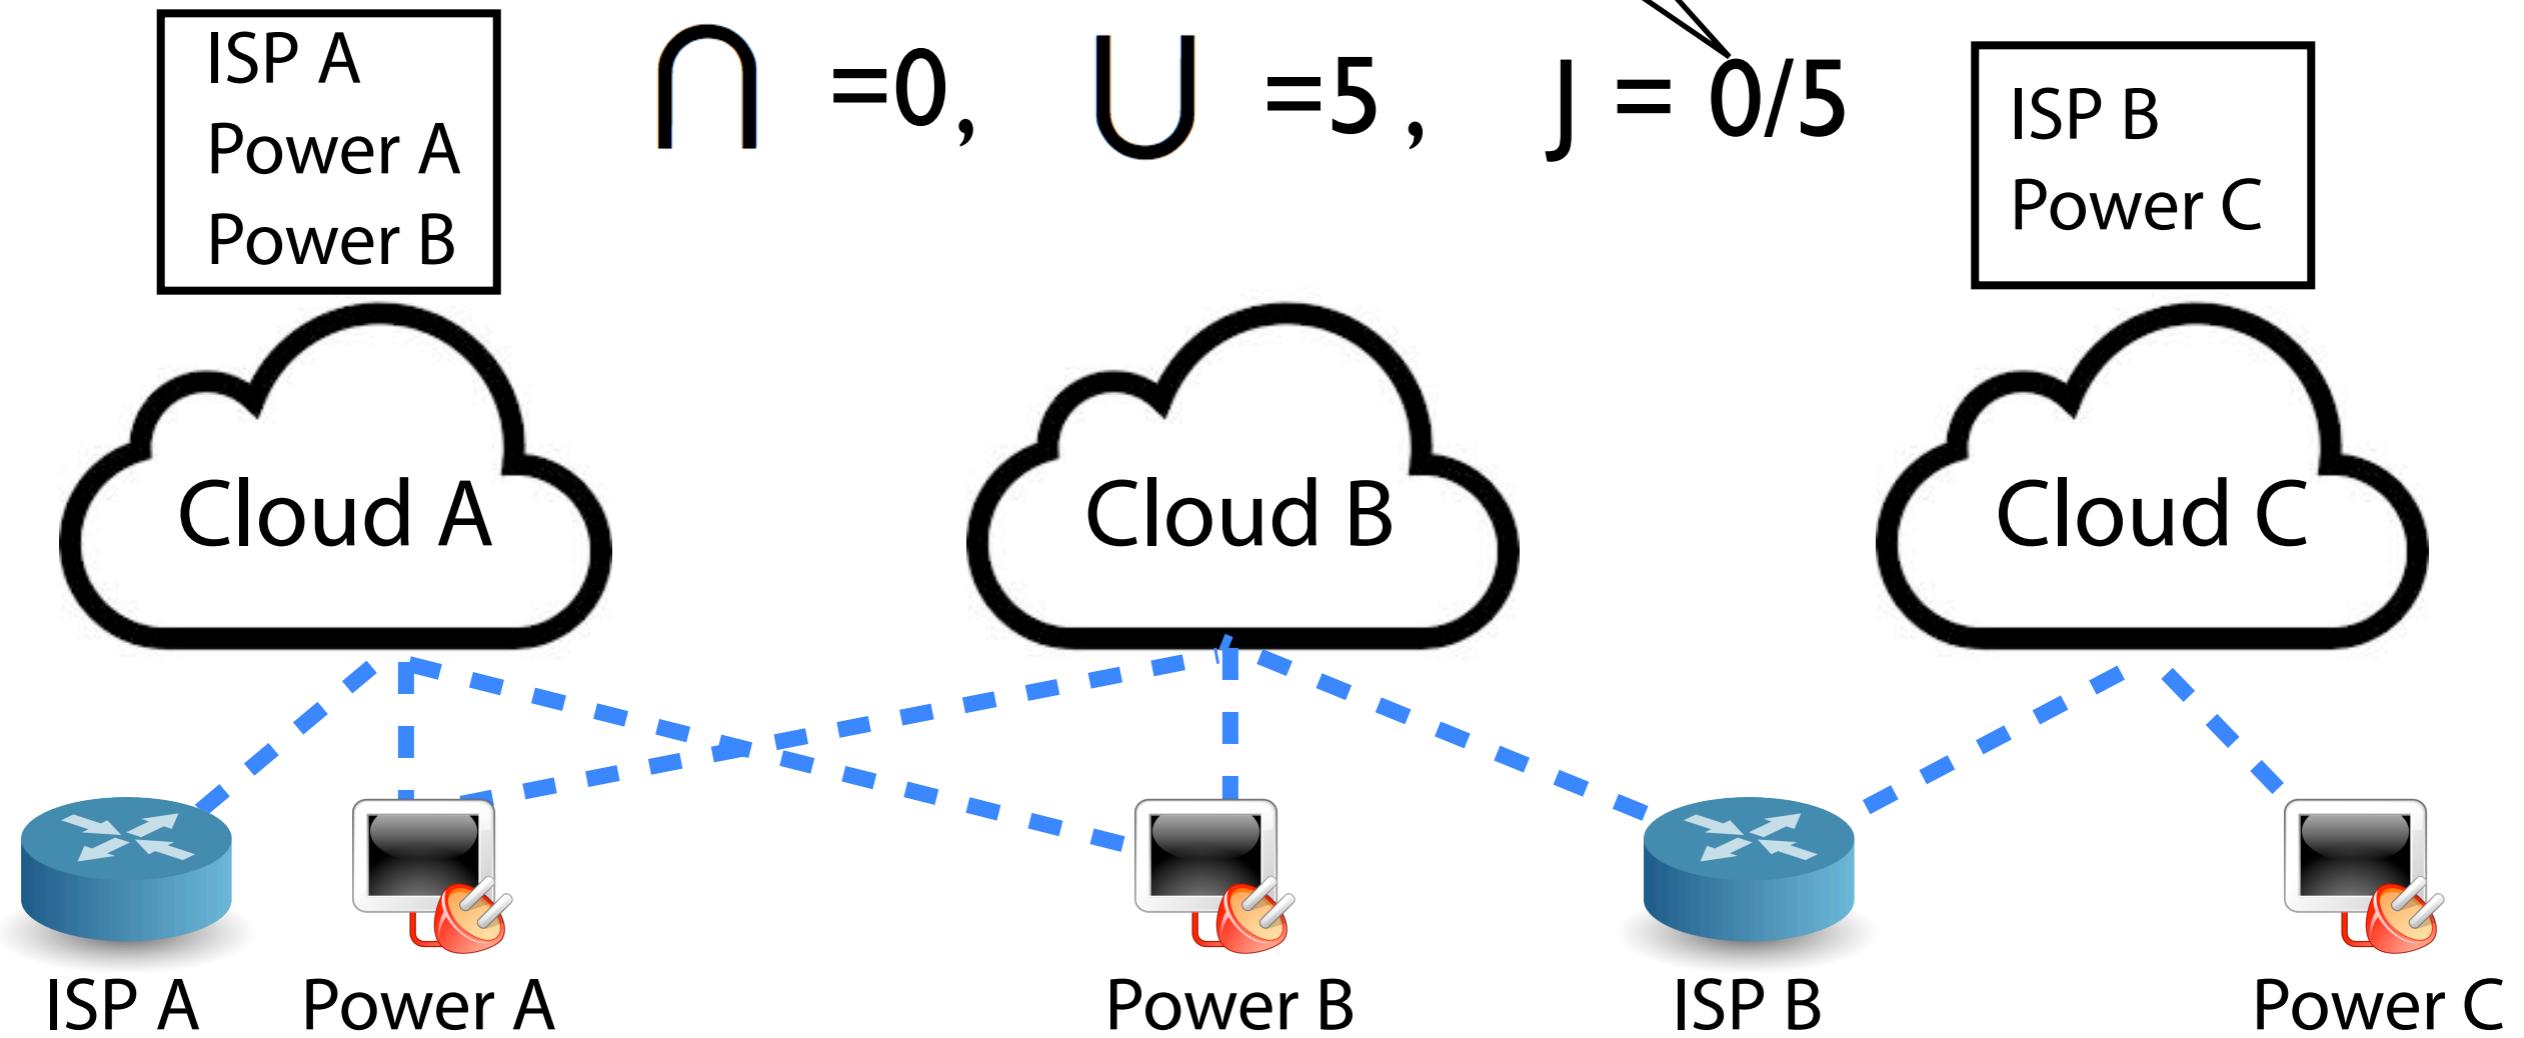


0 means fully independent



DaaS Agent

| Deployment | Sim  |
|------------|------|
| Cloud A&B  | 0.5  |
| Cloud B&C  | 0.25 |
| Cloud A&C  | 0    |





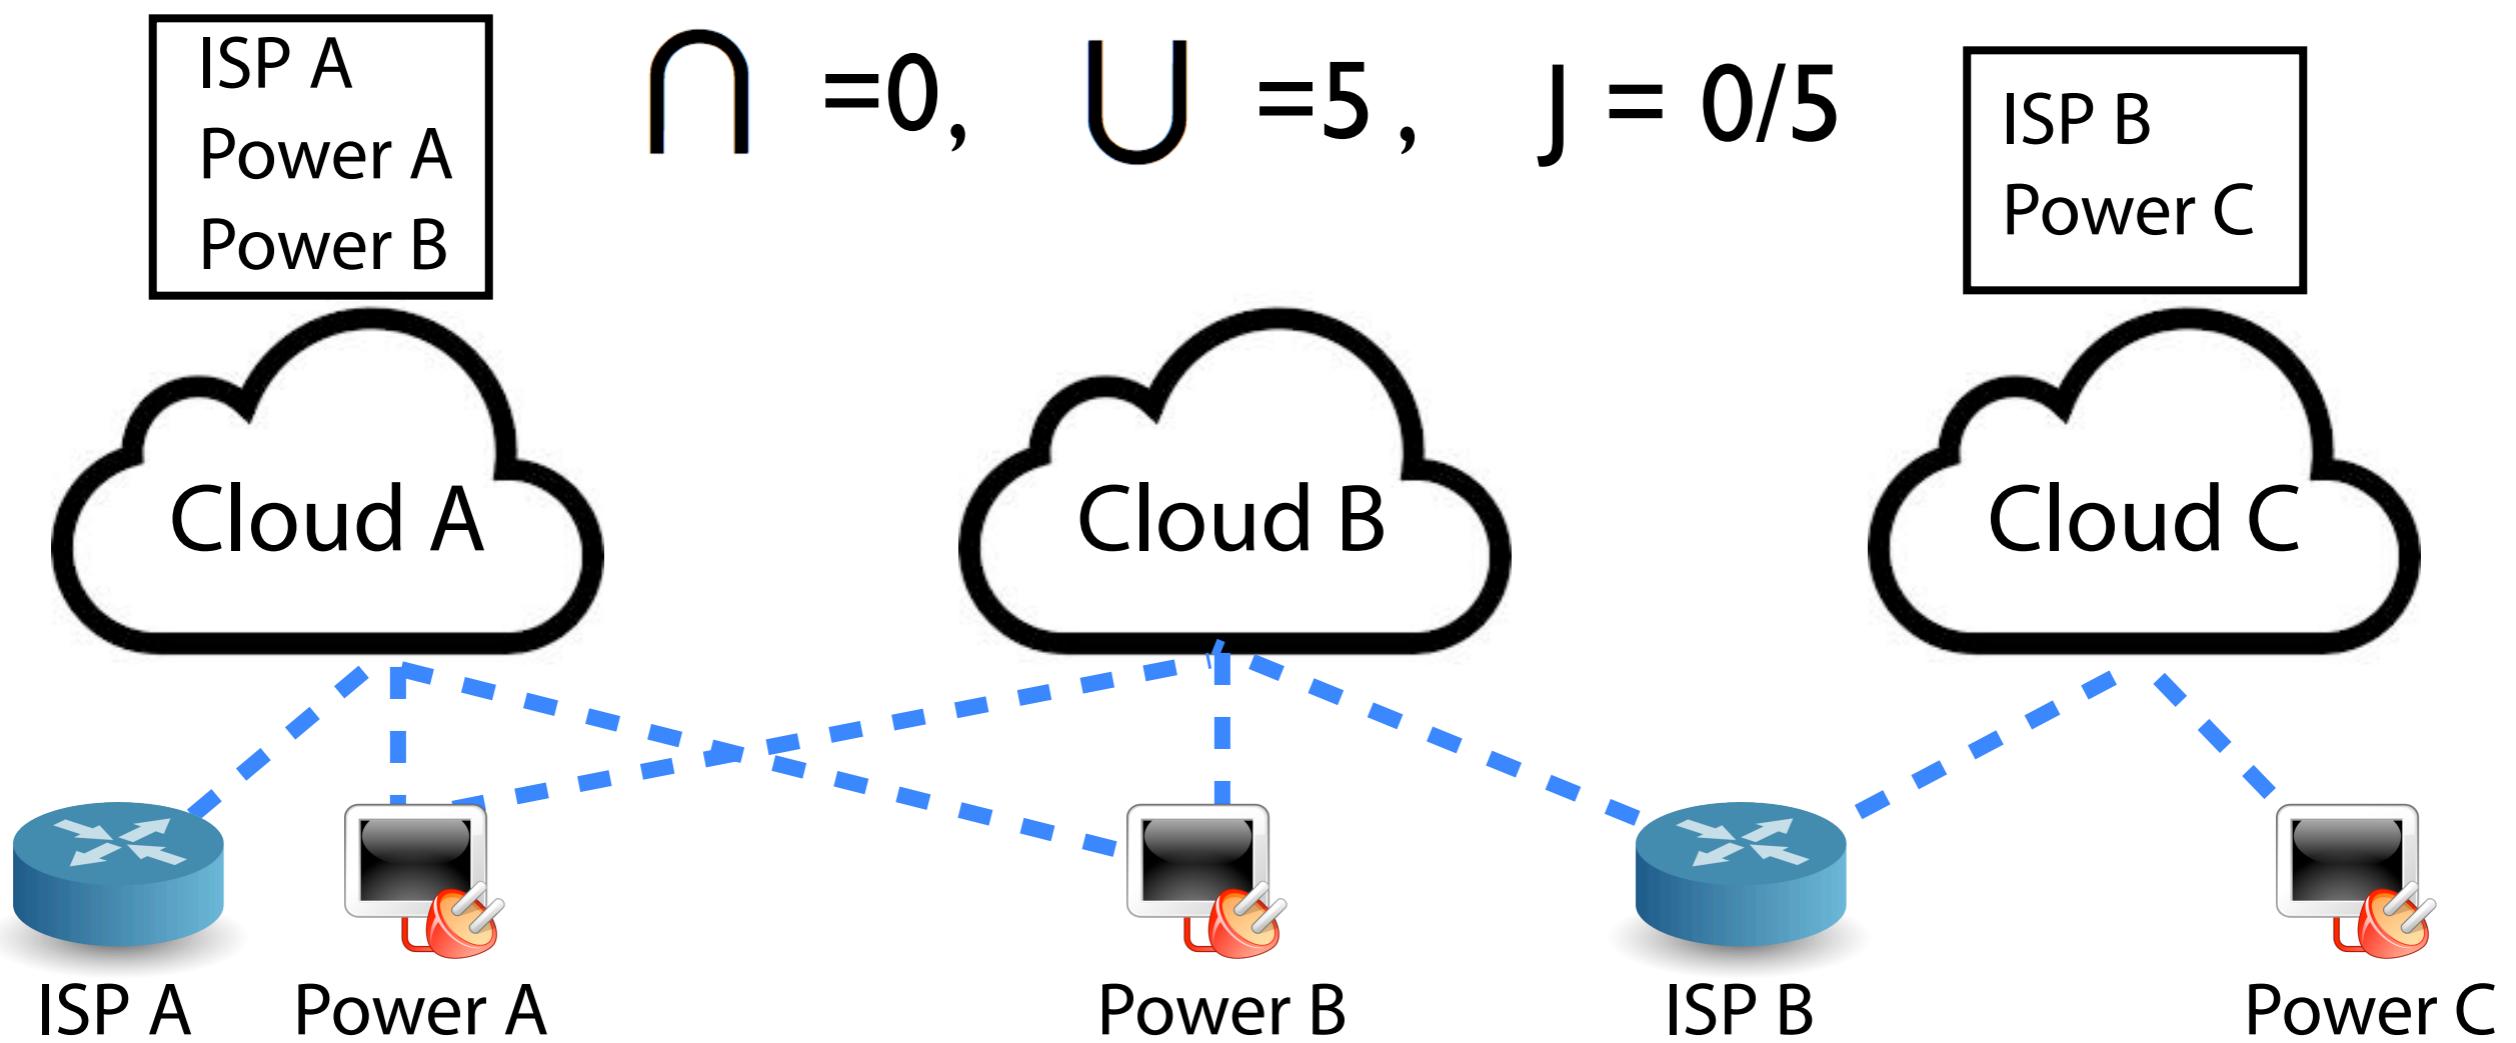
Service Provider



INDaaS Agent

| Deployment | Sim  |
|------------|------|
| Cloud A&B  | 0.5  |
| Cloud B&C  | 0.25 |
| Cloud A&C  | 0    |

$$\cap = 0, \cup = 5, J = 0/5$$



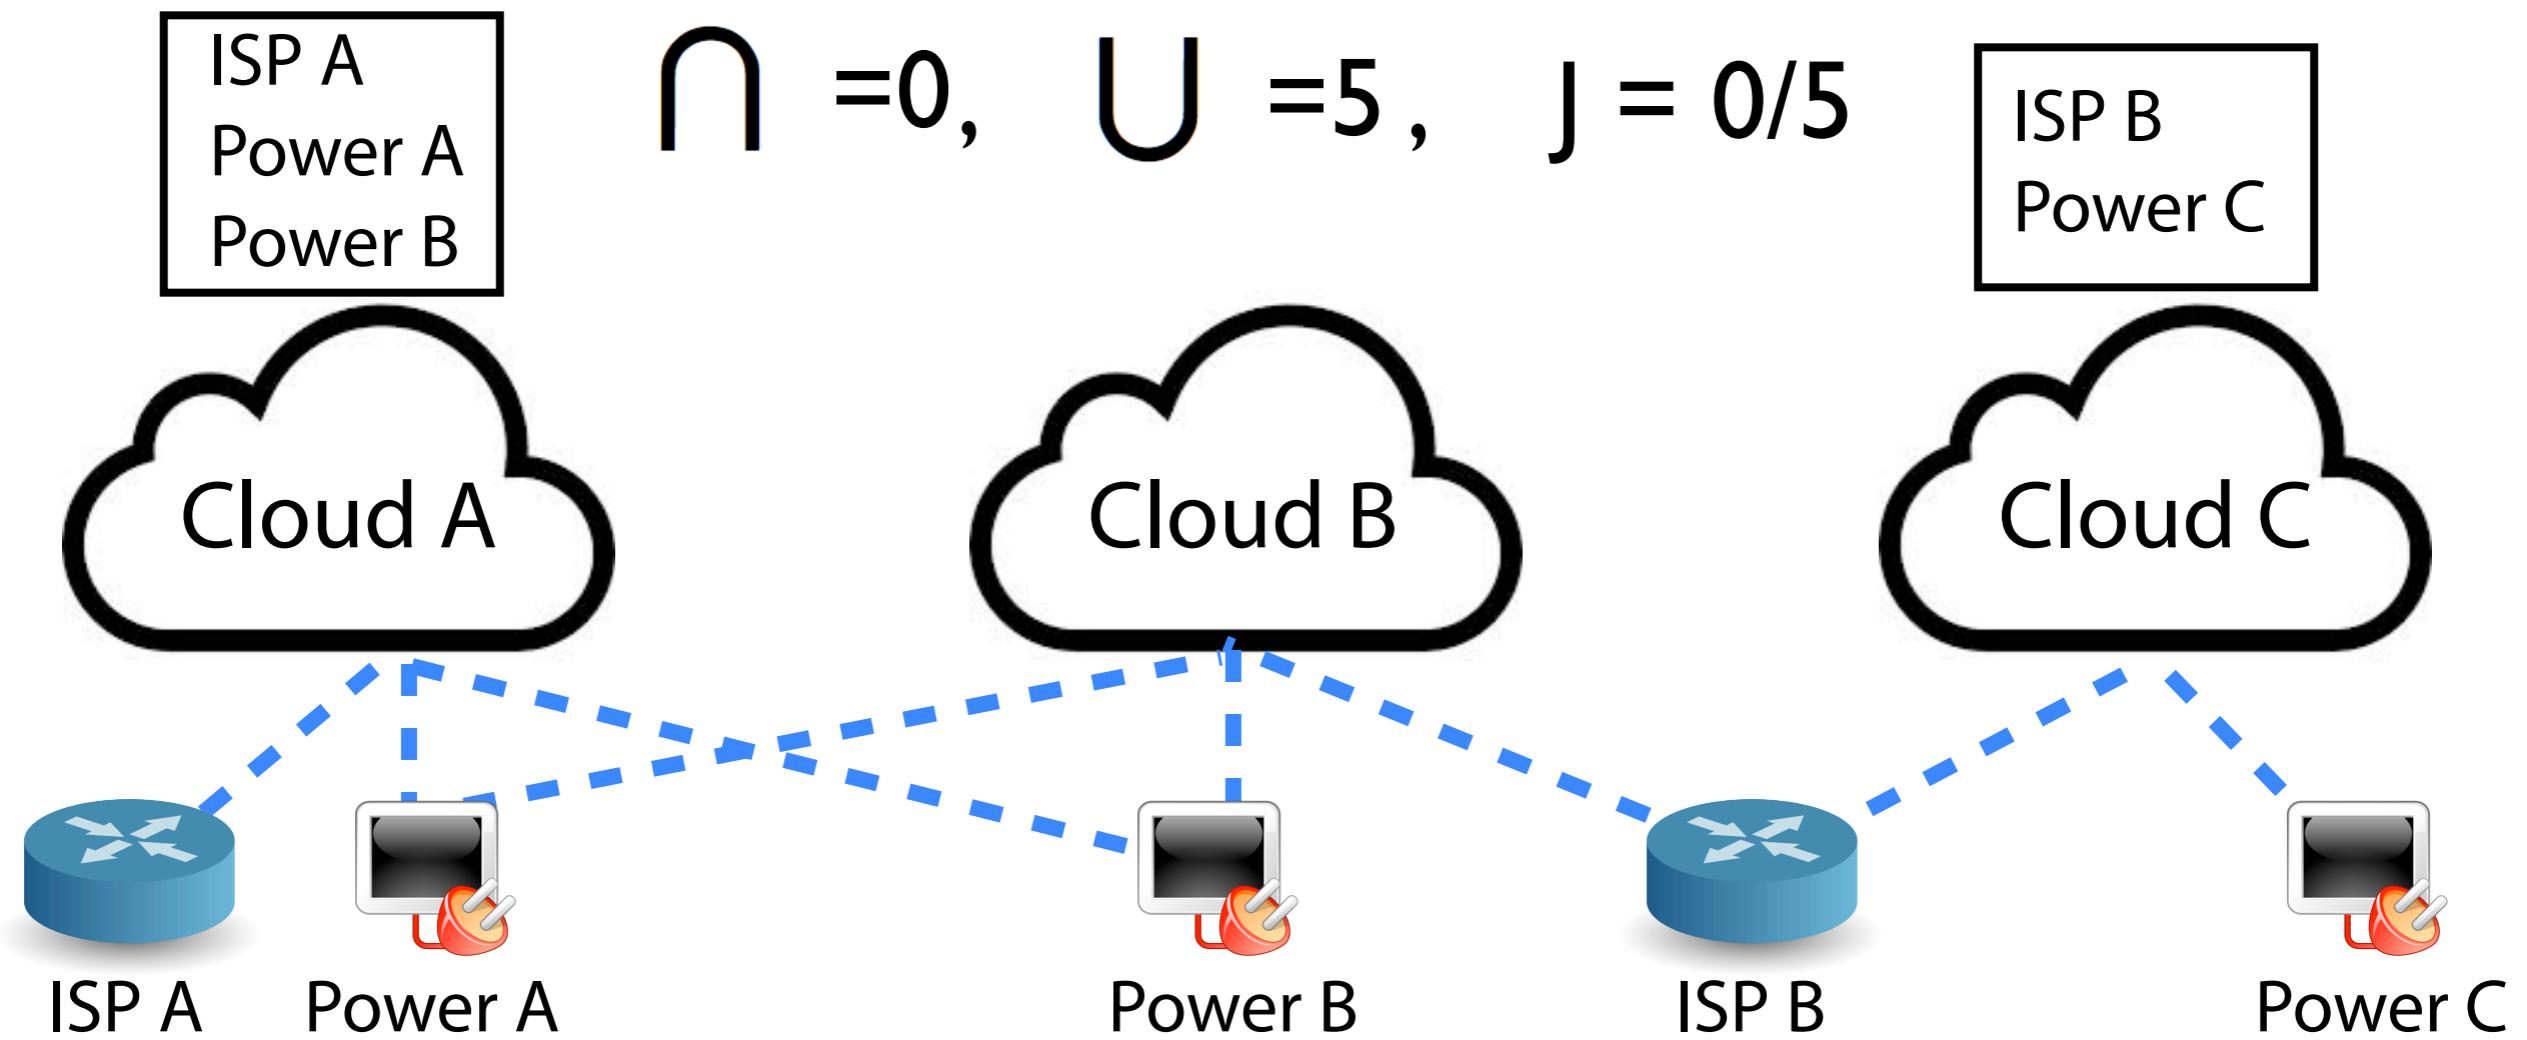


Service Provider



INDaaS Agent

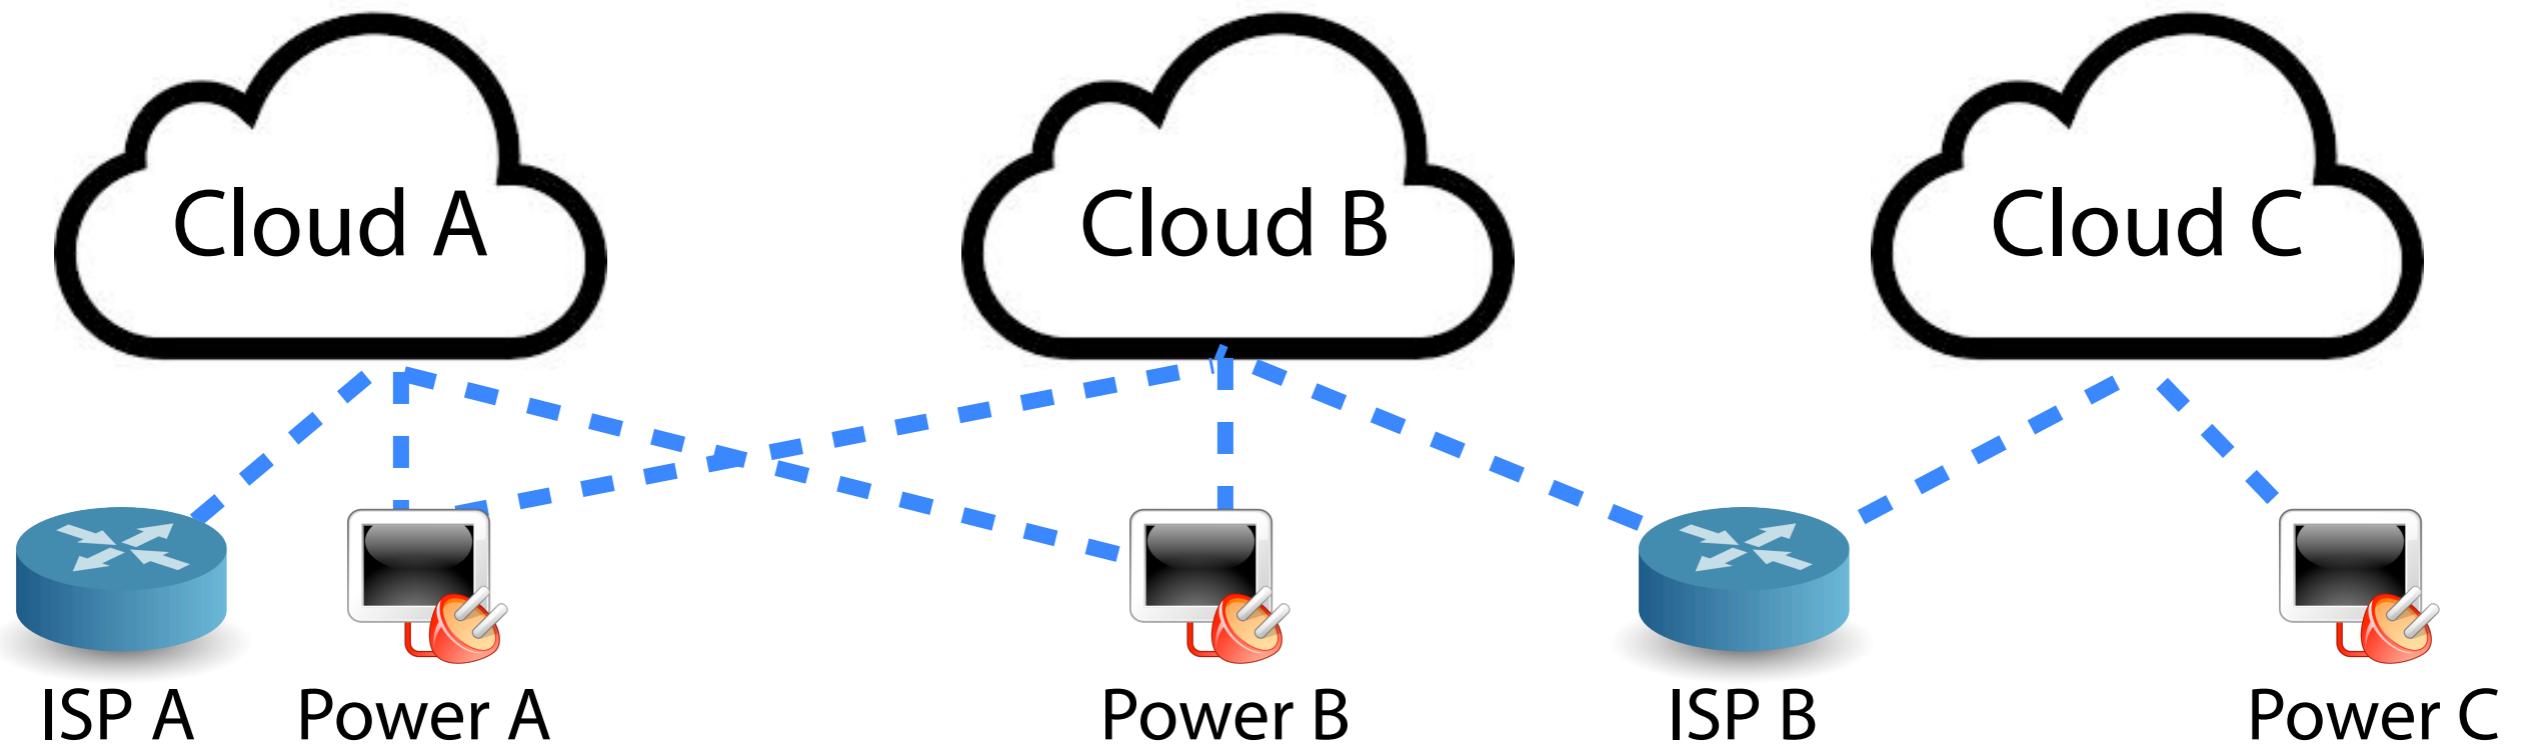
| Deployment | Sim  |
|------------|------|
| Cloud A&C  | 0    |
| Cloud B&C  | 0.25 |
| Cloud A&B  | 0.5  |



| Deployment | Sim  |
|------------|------|
| Cloud A&C  | 0    |
| Cloud B&C  | 0.25 |
| Cloud A&B  | 0.5  |

 Service Provider

 INDaaS Agent



# P-SOP [Vaidya et al. JCS05]

- We apply Private Set Operation Protocol (P-SOP):
  - Private set intersection cardinality.
  - Private set union cardinality.

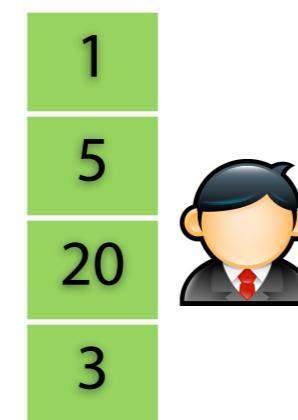
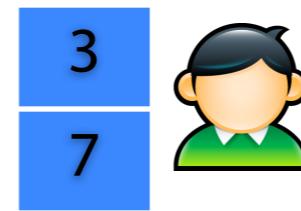
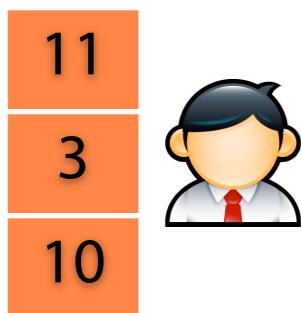
$$J(S_1, S_2, \dots, S_n) = \frac{|S_1 \cap S_2 \cap \dots \cap S_n|}{|S_1 \cup S_2 \cup \dots \cup S_n|}$$

# P-SOP [Vaidya et al. JCS05]

- Allow  $k$  parties to compute both intersection and union cardinalities without learning other information.

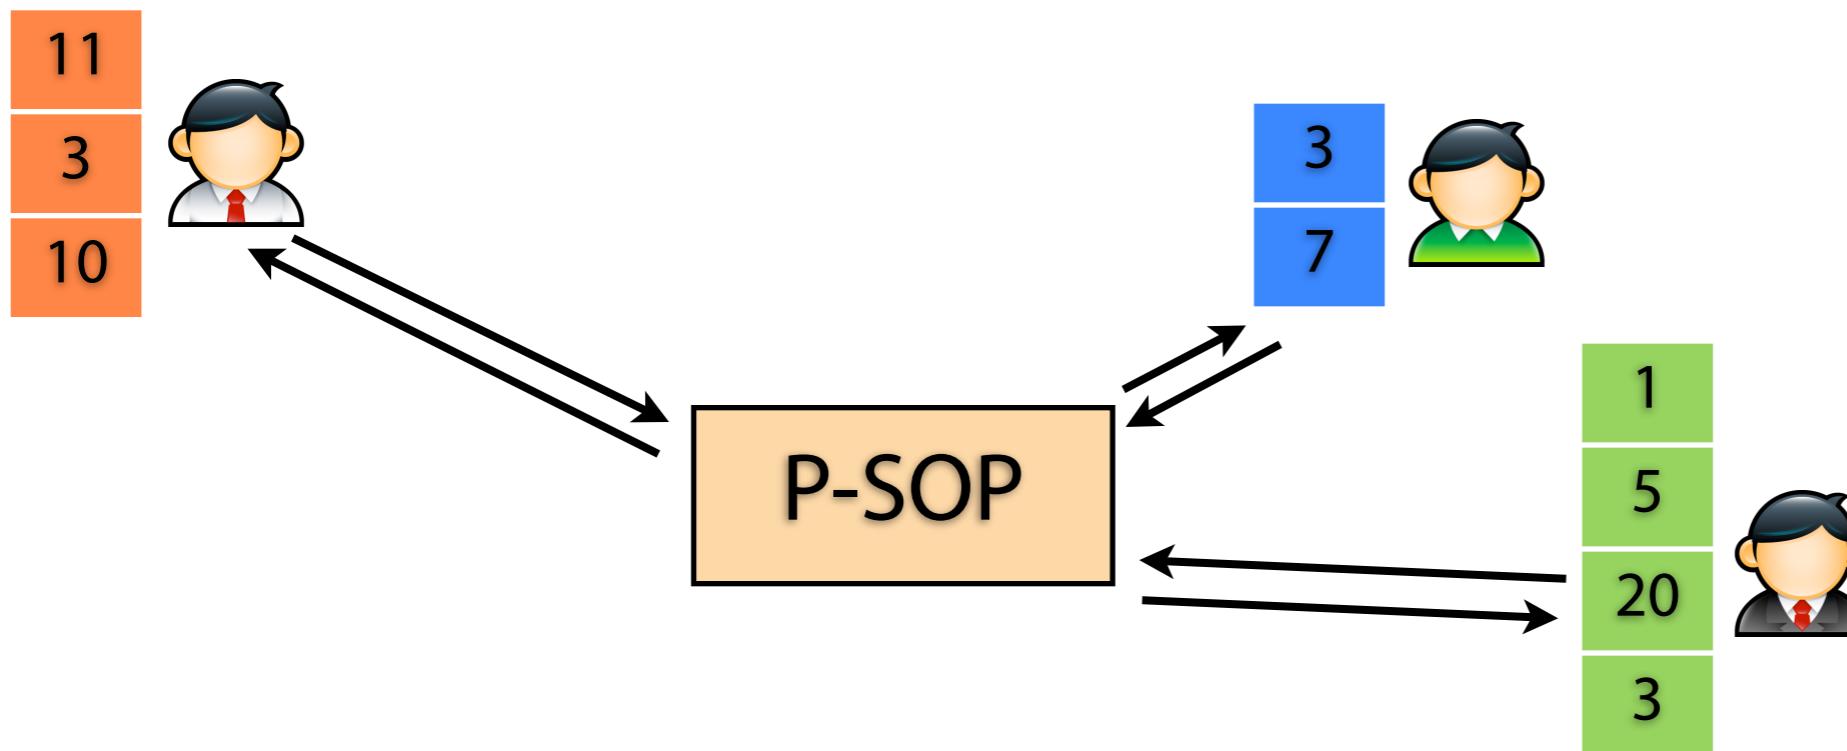
# P-SOP [Vaidya et al. JCS05]

- Allow  $k$  parties to compute both intersection and union cardinalities without learning other information.



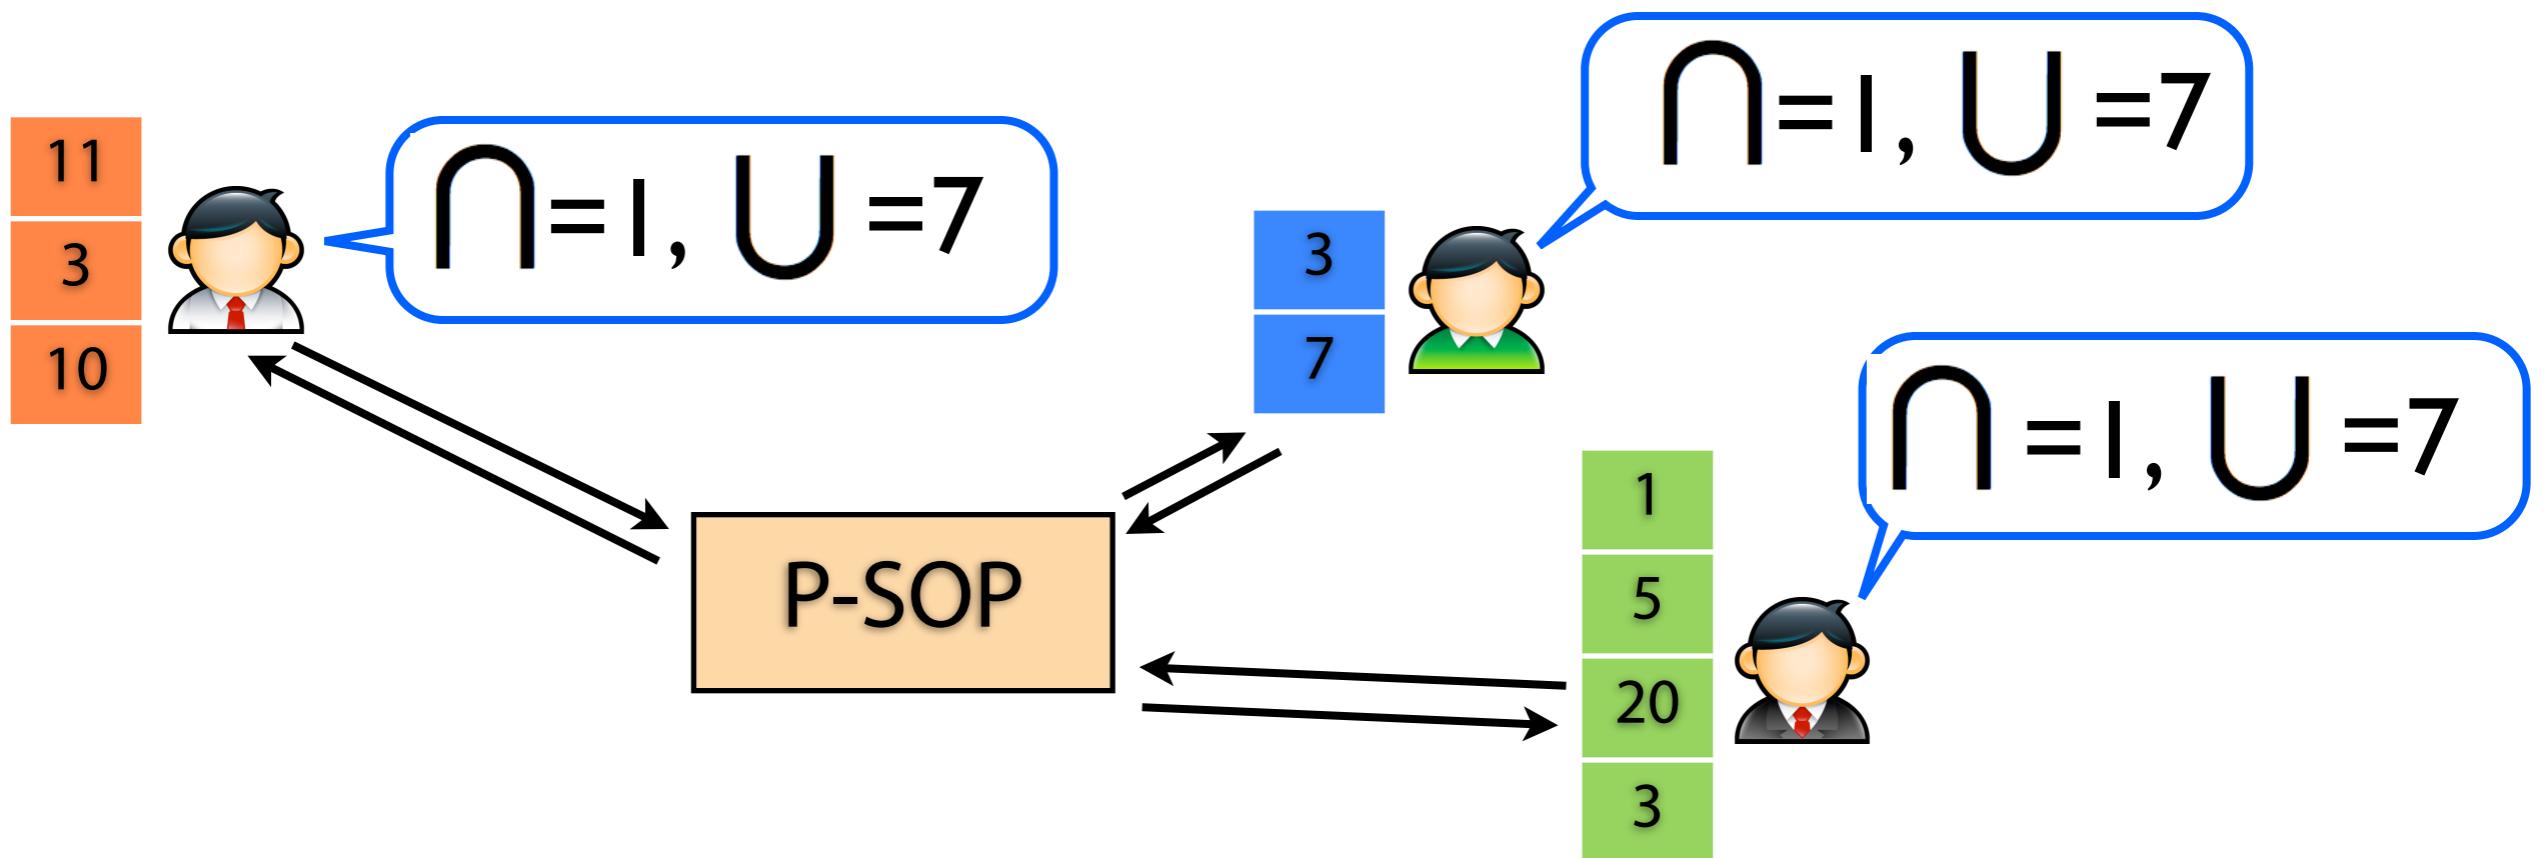
# P-SOP [Vaidya et al. JCS05]

- Allow  $k$  parties to compute both intersection and union cardinalities without learning other information.



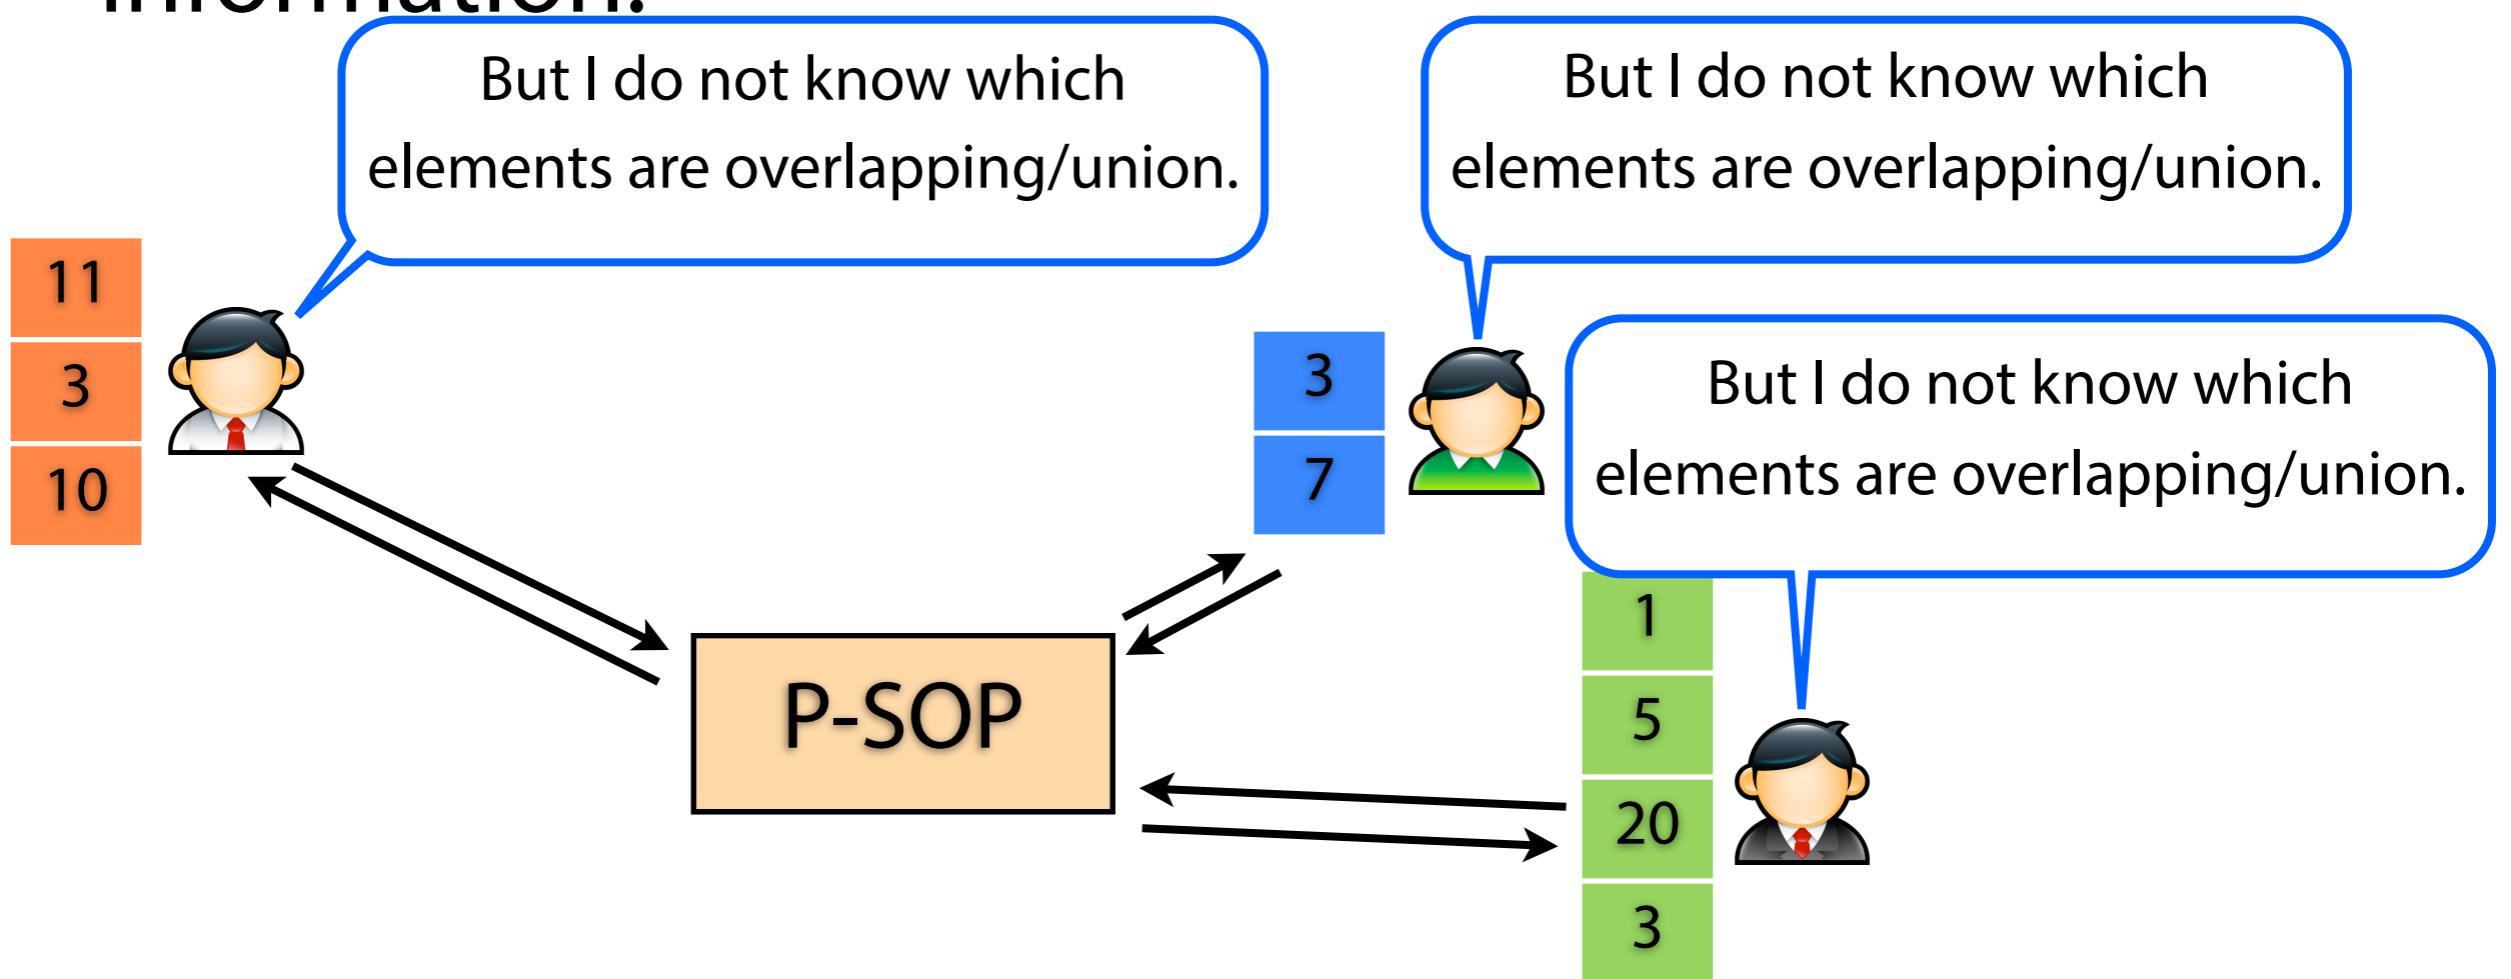
# P-SOP [Vaidya et al. JCS05]

- Allow  $k$  parties to compute both intersection and union cardinalities without learning other information.

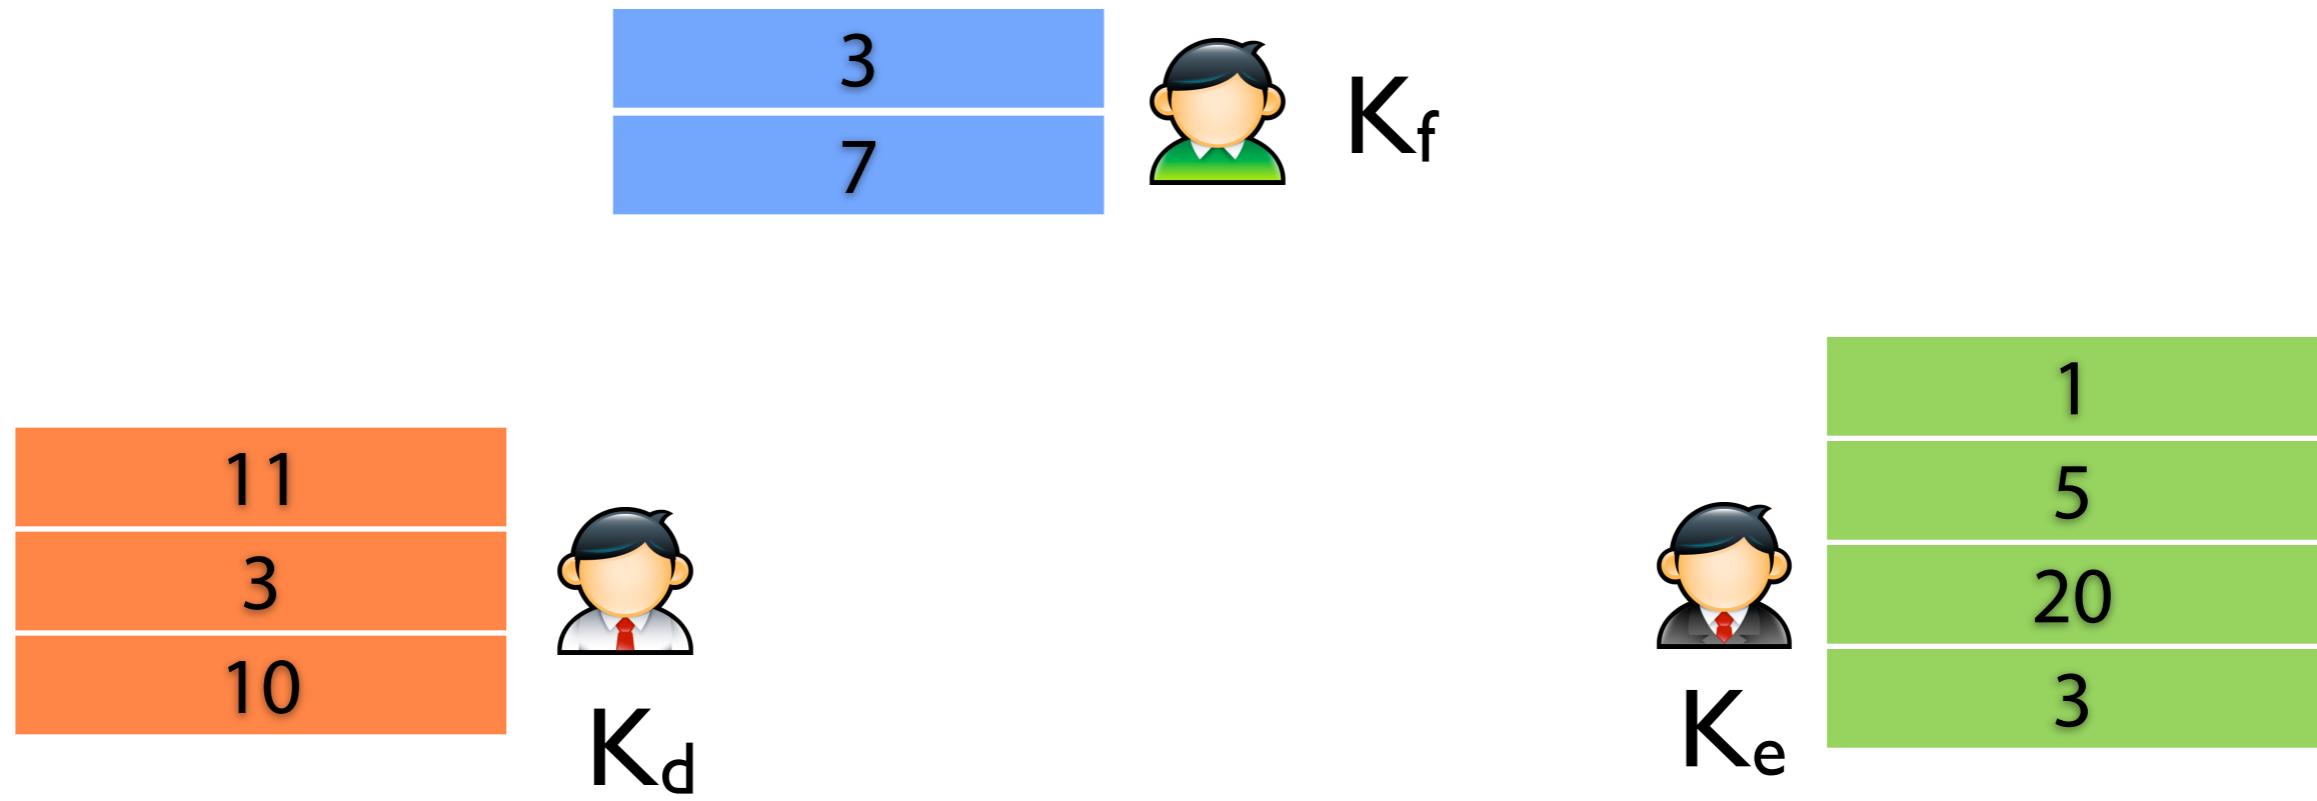


# P-SOP [Vaidya et al. JCS05]

- Allow  $k$  parties to compute both intersection and union cardinalities without learning other information.



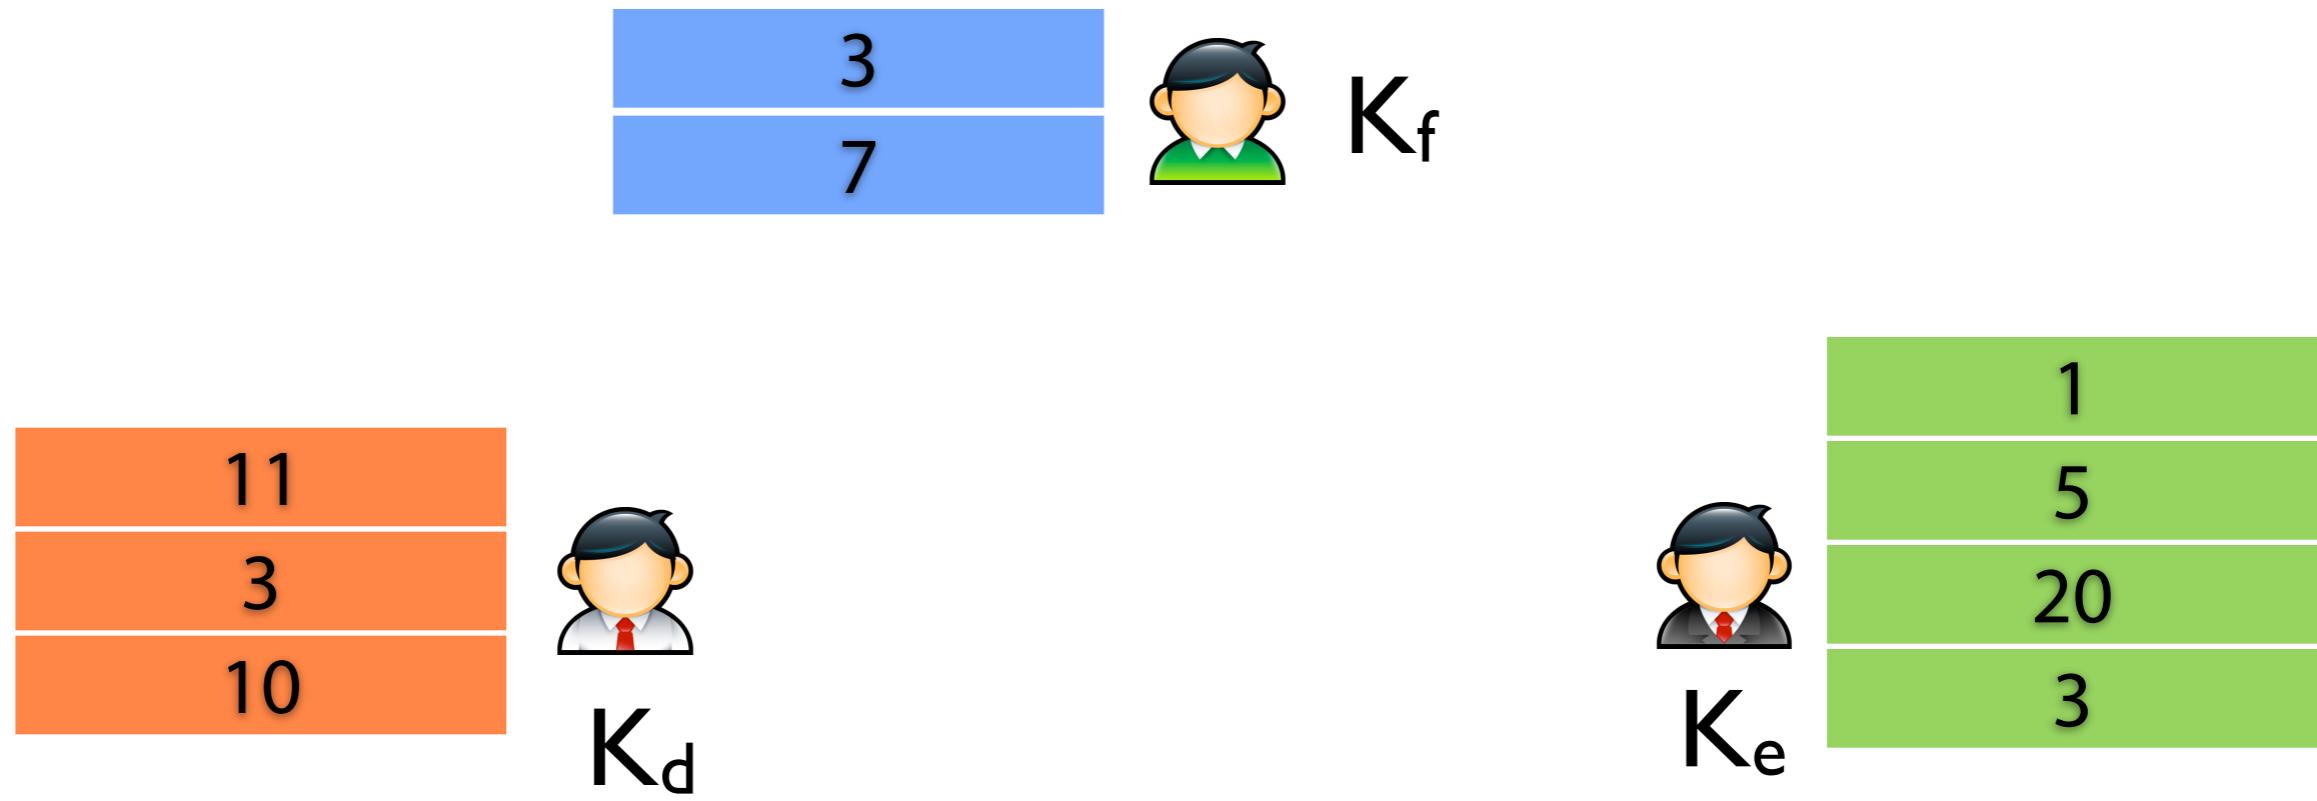
# P-SOP [Vaidya et al. JCS05]



Each party maintains a commutative encryption key

Commutative encryption holds:  $E_x(E_y(m)) = E_y(E_x(m))$

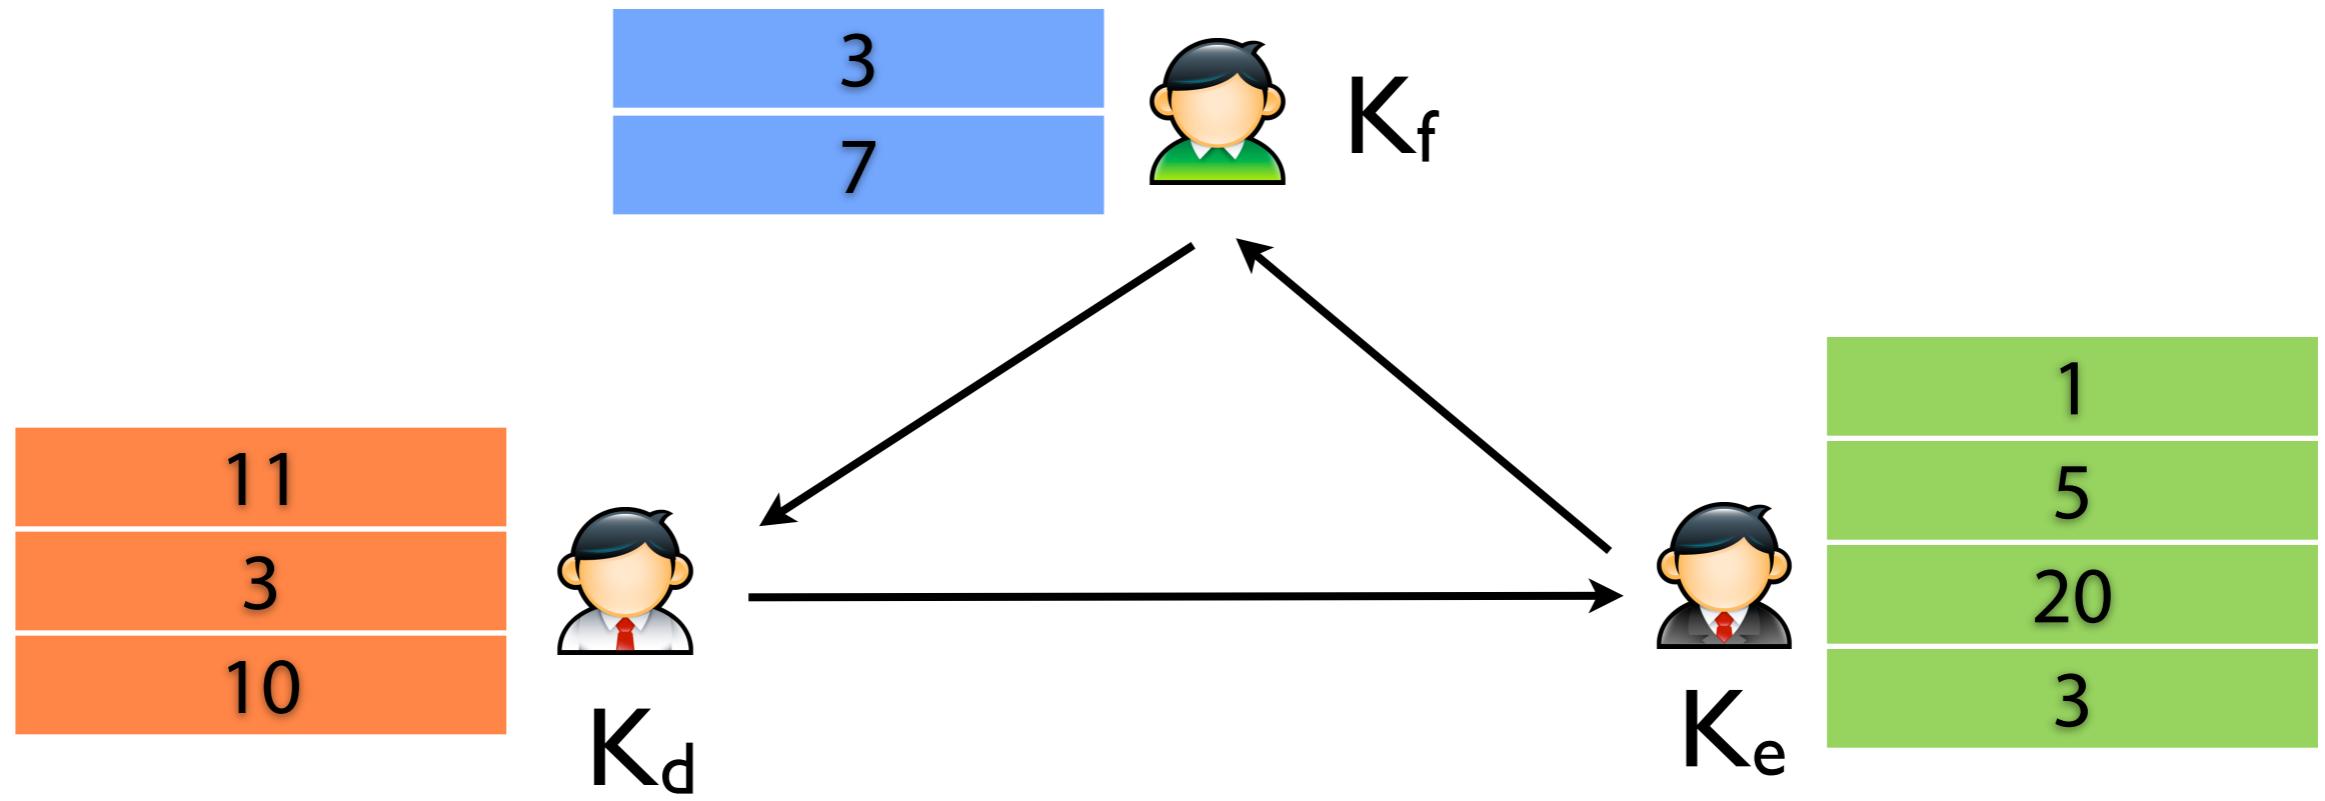
# P-SOP [Vaidya et al. JCS05]



Each party maintains a commutative encryption key

Commutative encryption holds:  $E_x(E_y(m)) = E_y(E_x(m))$

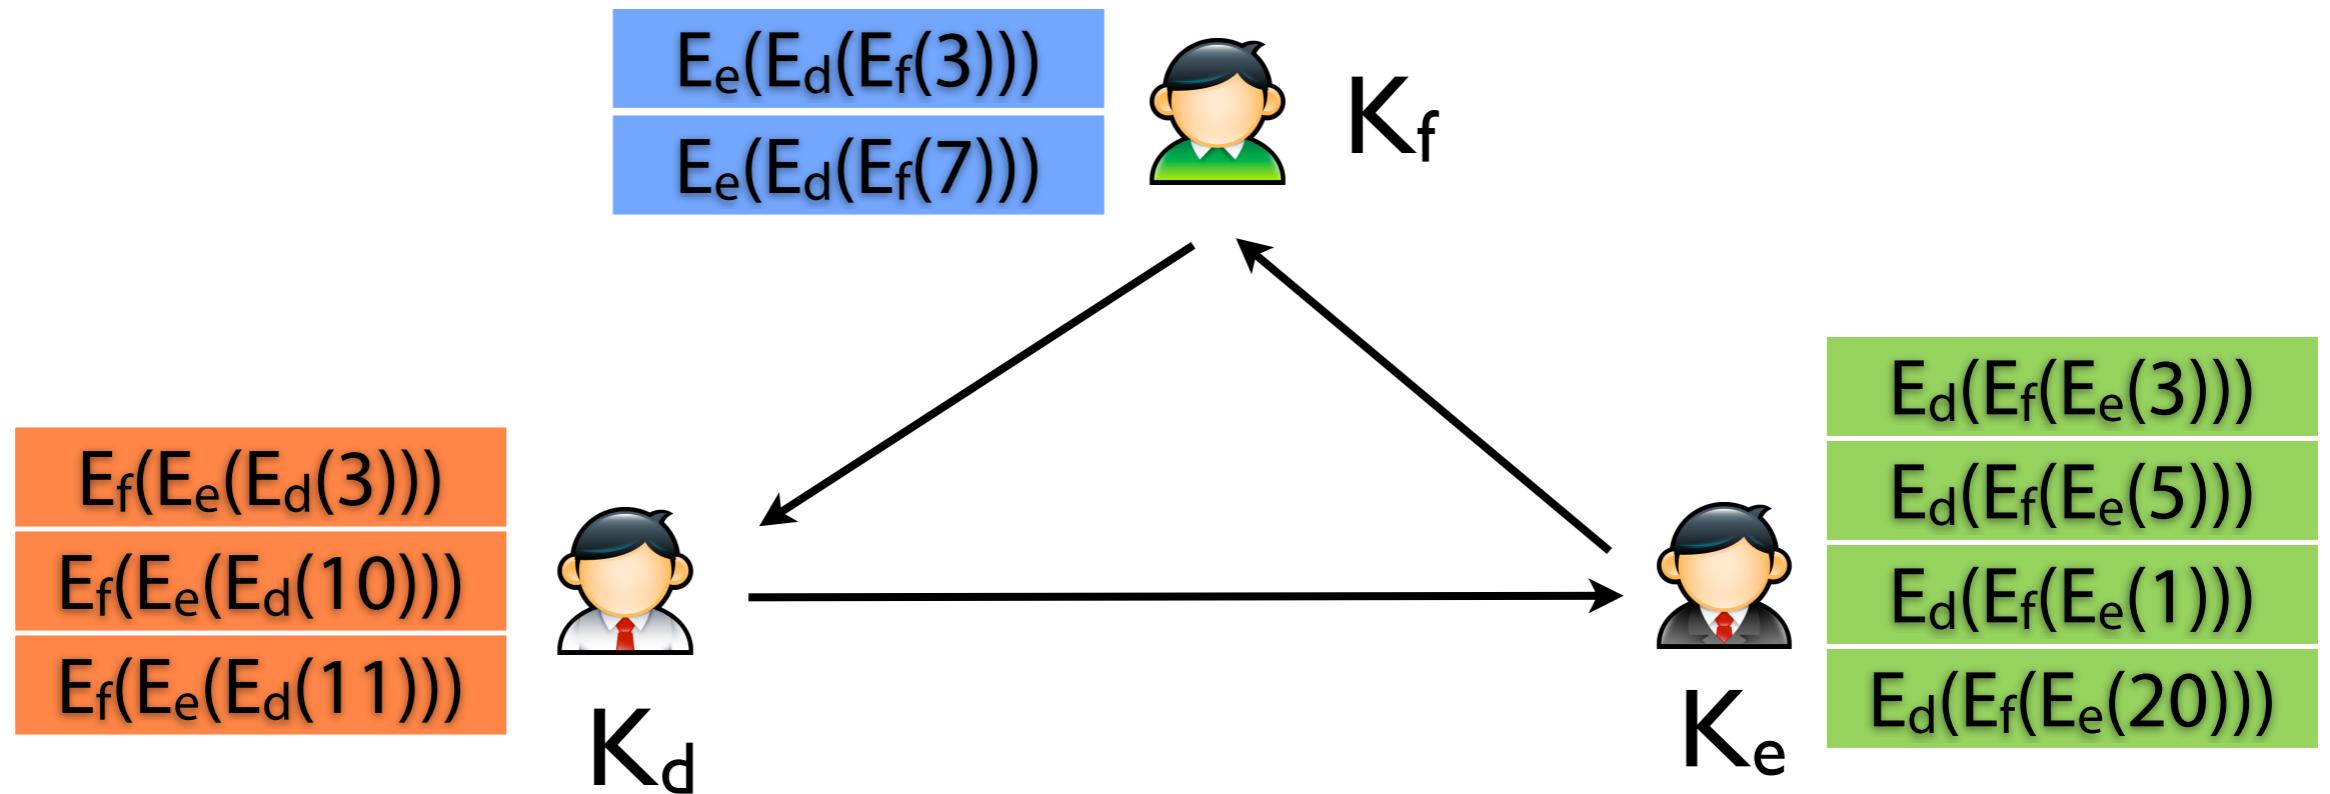
# P-SOP [Vaidya et al. JCS05]



Each party maintains a commutative encryption key

Commutative encryption holds:  $E_x(E_y(m)) = E_y(E_x(m))$

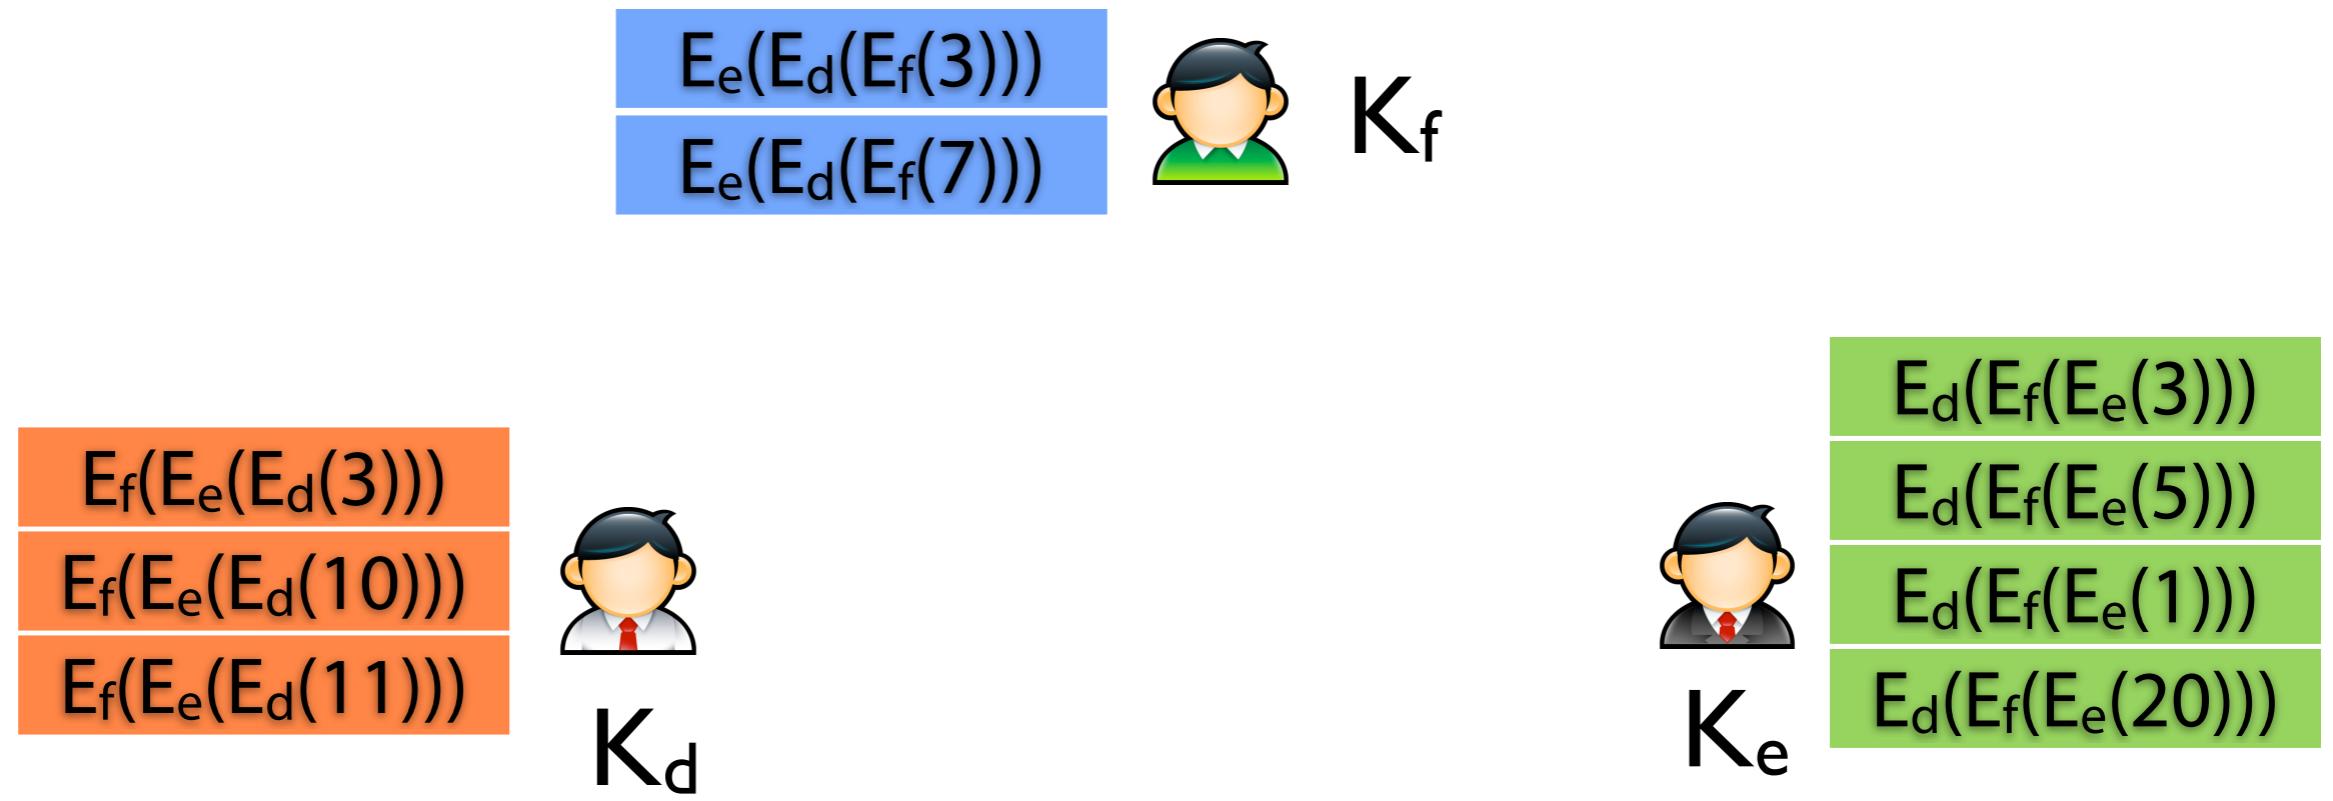
# P-SOP [Vaidya et al. JCS05]



Each party maintains a commutative encryption key

Commutative encryption holds:  $E_x(E_y(m)) = E_y(E_x(m))$

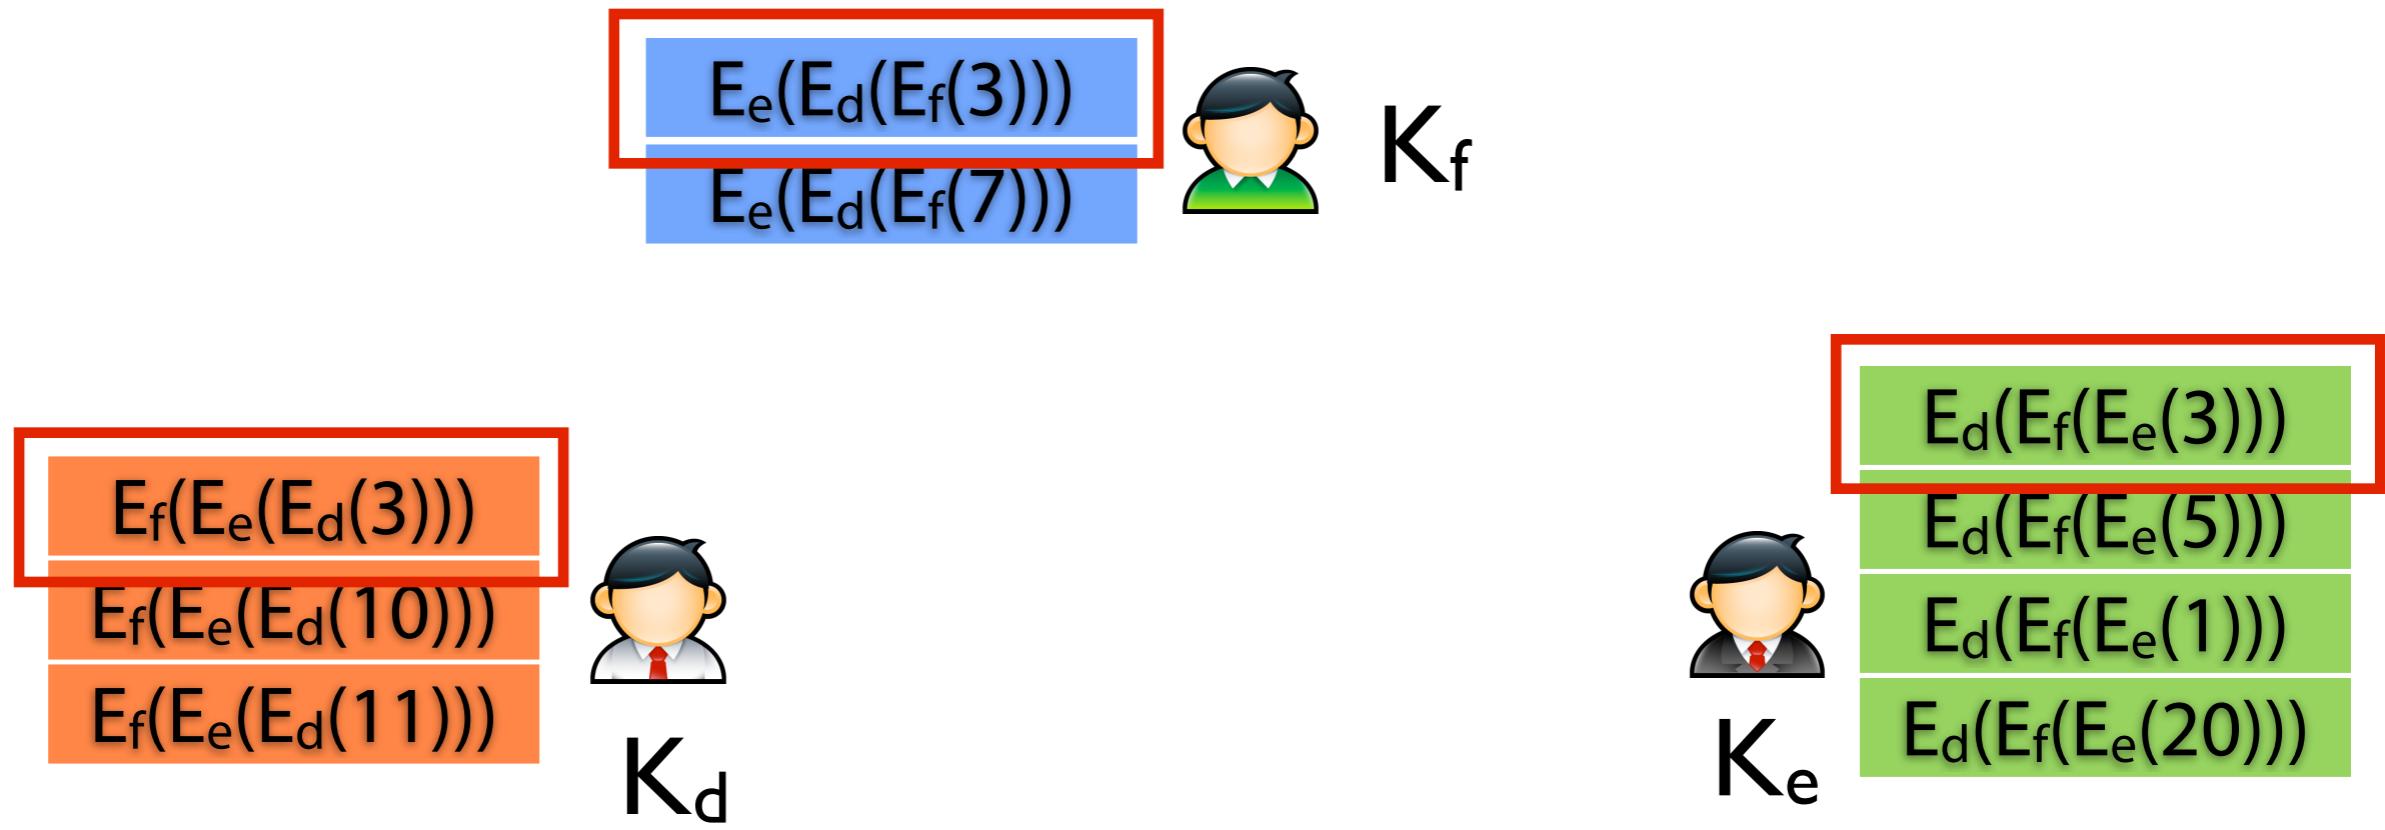
# P-SOP [Vaidya et al. JCS05]



Each party maintains a **commutative encryption key**

Commutative encryption holds:  $E_x(E_y(m)) = E_y(E_x(m))$

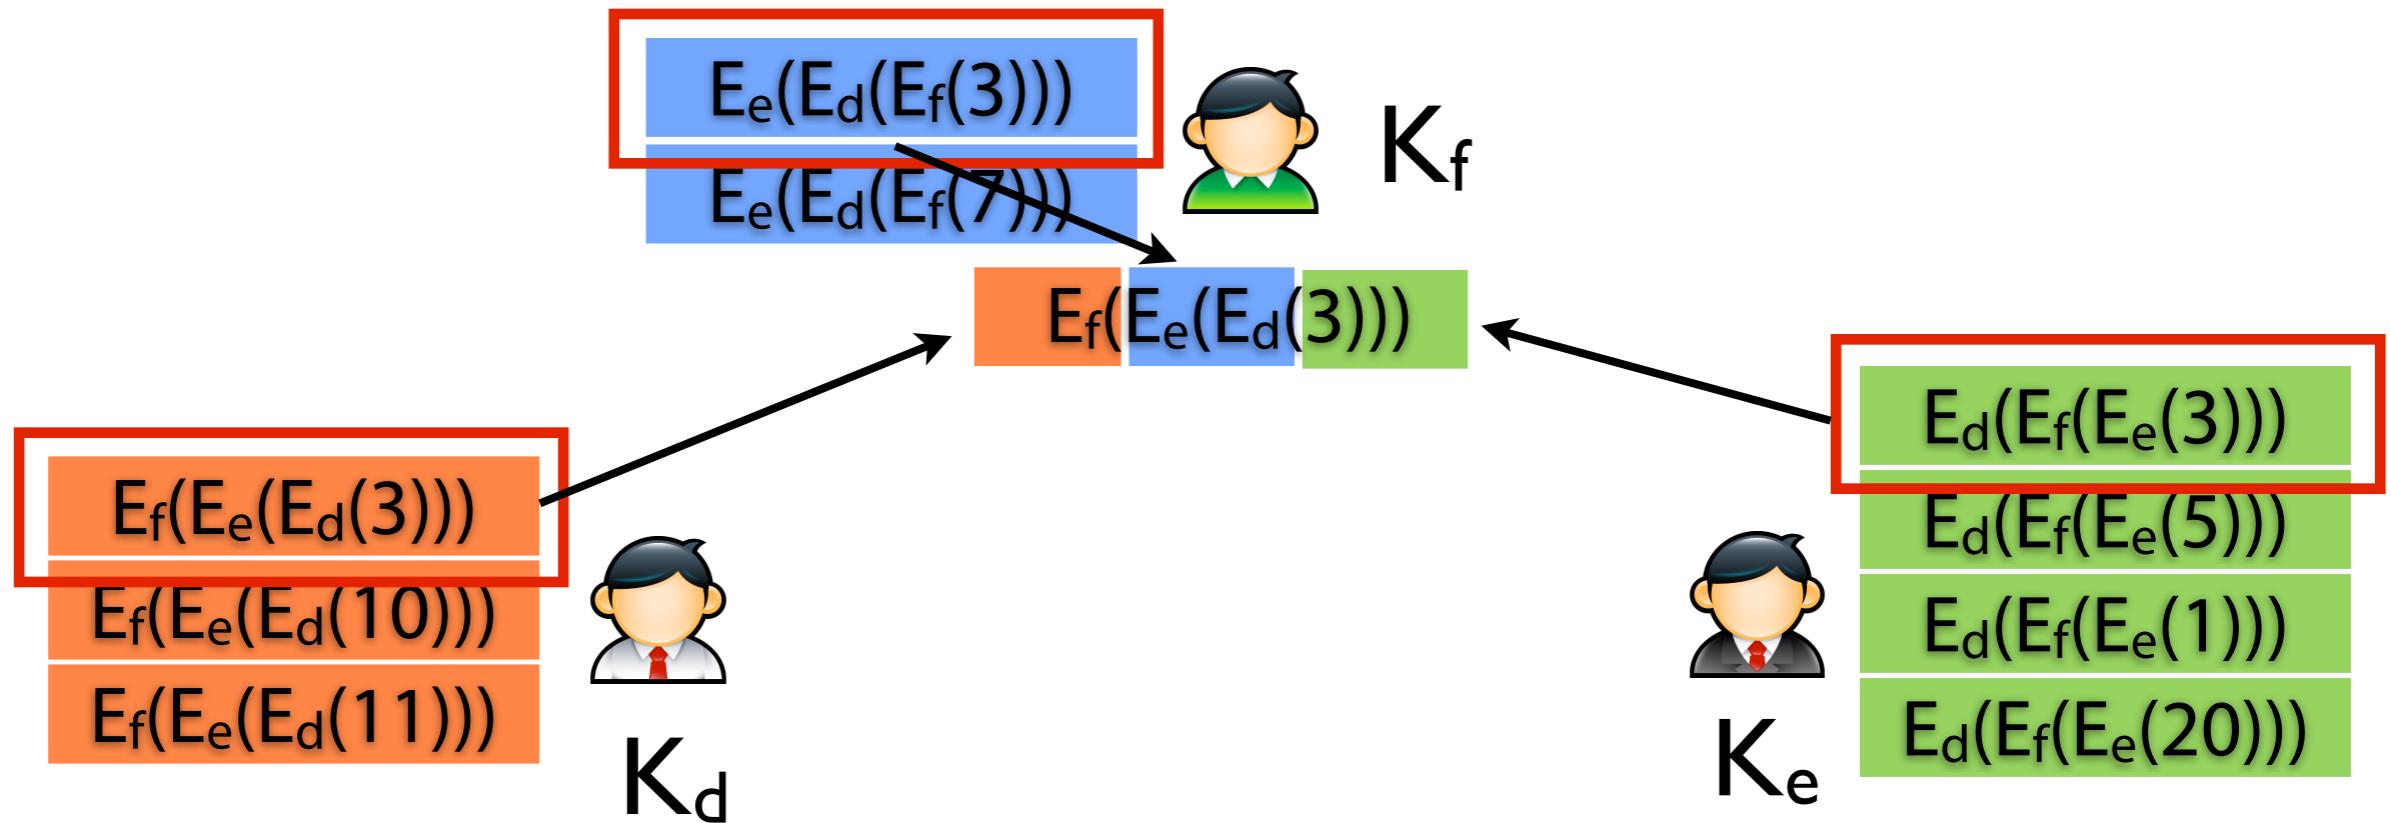
# P-SOP [Vaidya et al. JCS05]



Each party maintains a **commutative encryption key**

Commutative encryption holds:  $E_x(E_y(m)) = E_y(E_x(m))$

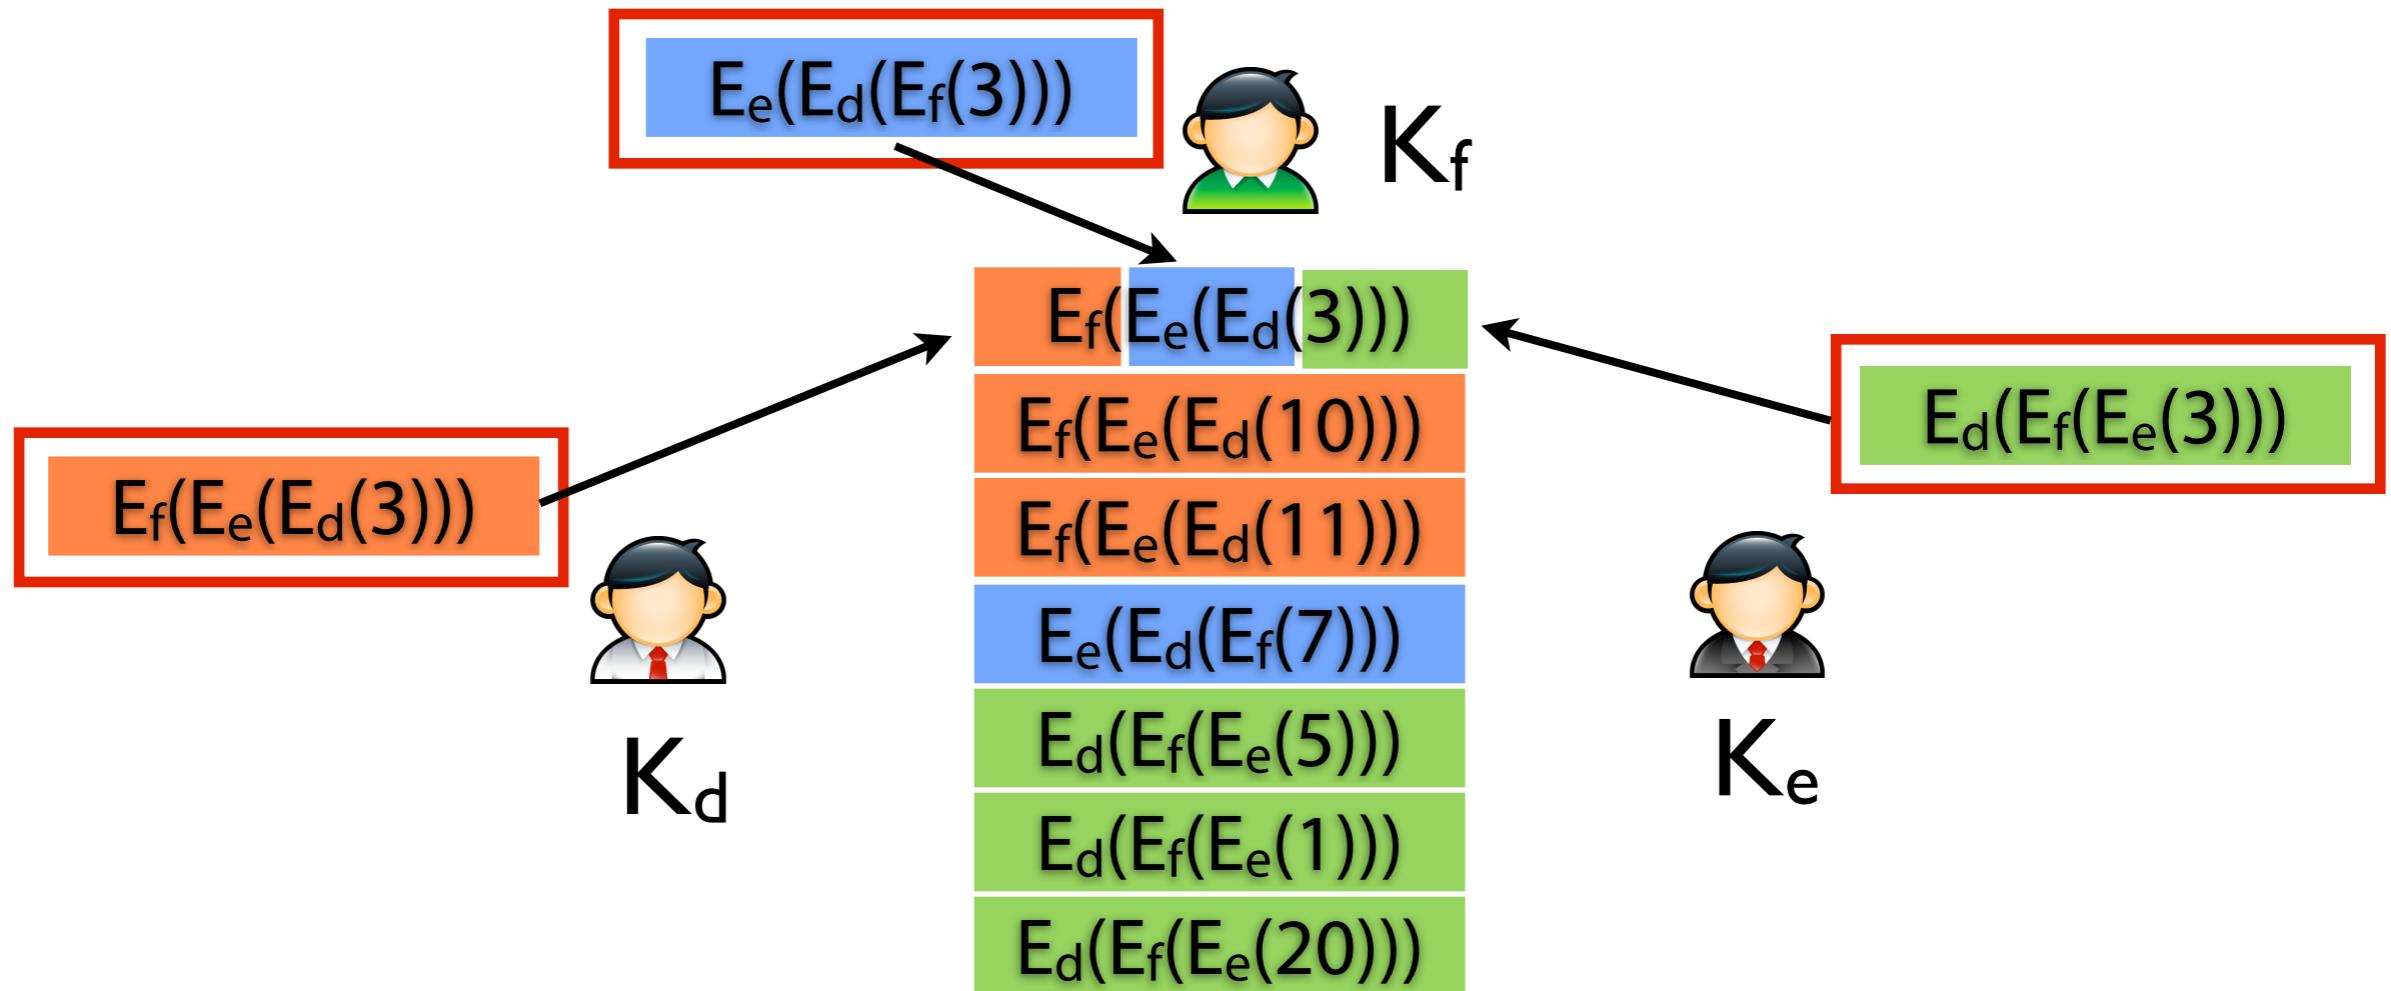
# P-SOP [Vaidya et al. JCS05]



Each party maintains a commutative encryption key

Commutative encryption holds:  $E_x(E_y(m)) = E_y(E_x(m))$

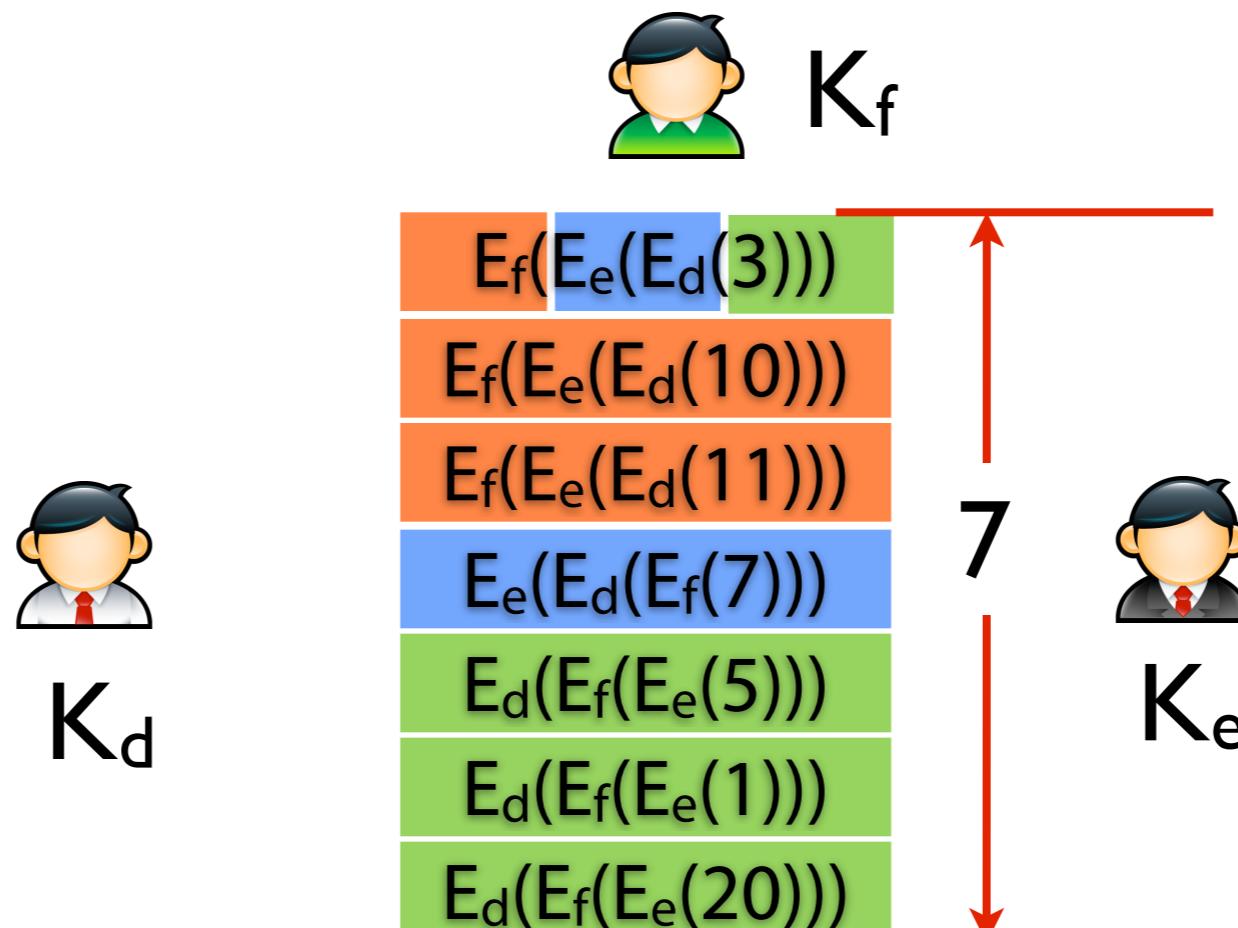
# P-SOP [Vaidya et al. JCS05]



Each party maintains a **commutative encryption key**

Commutative encryption holds:  $E_x(E_y(m)) = E_y(E_x(m))$

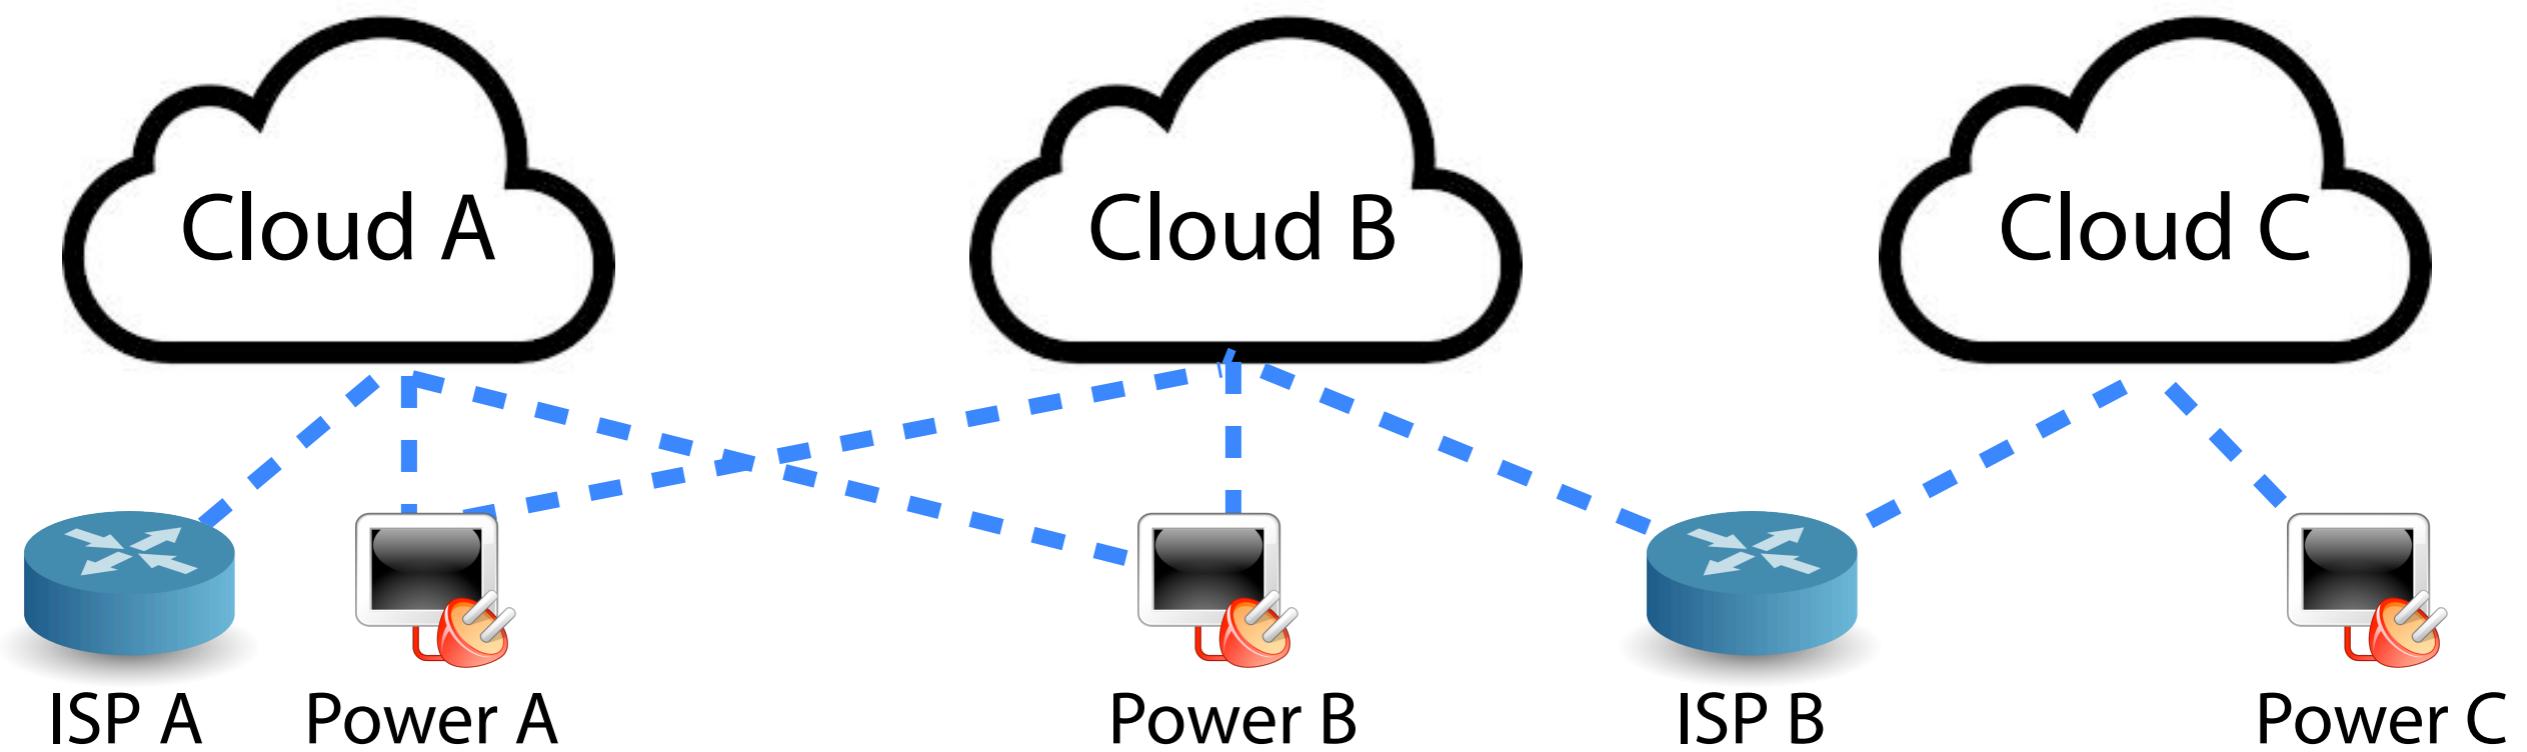
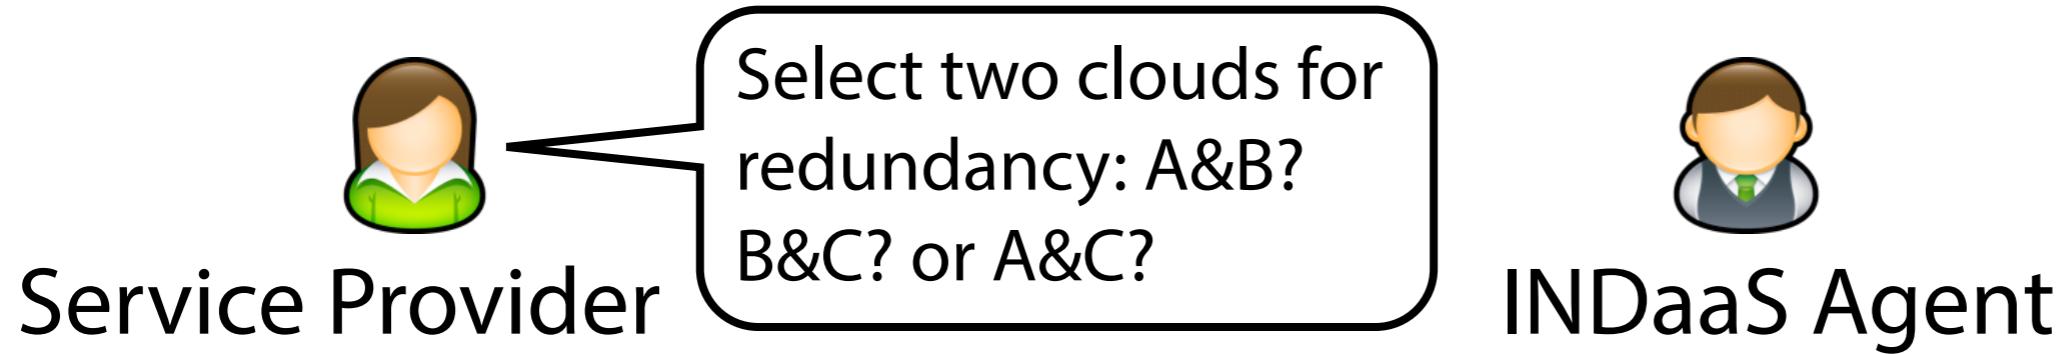
# P-SOP [Vaidya et al. JCS05]

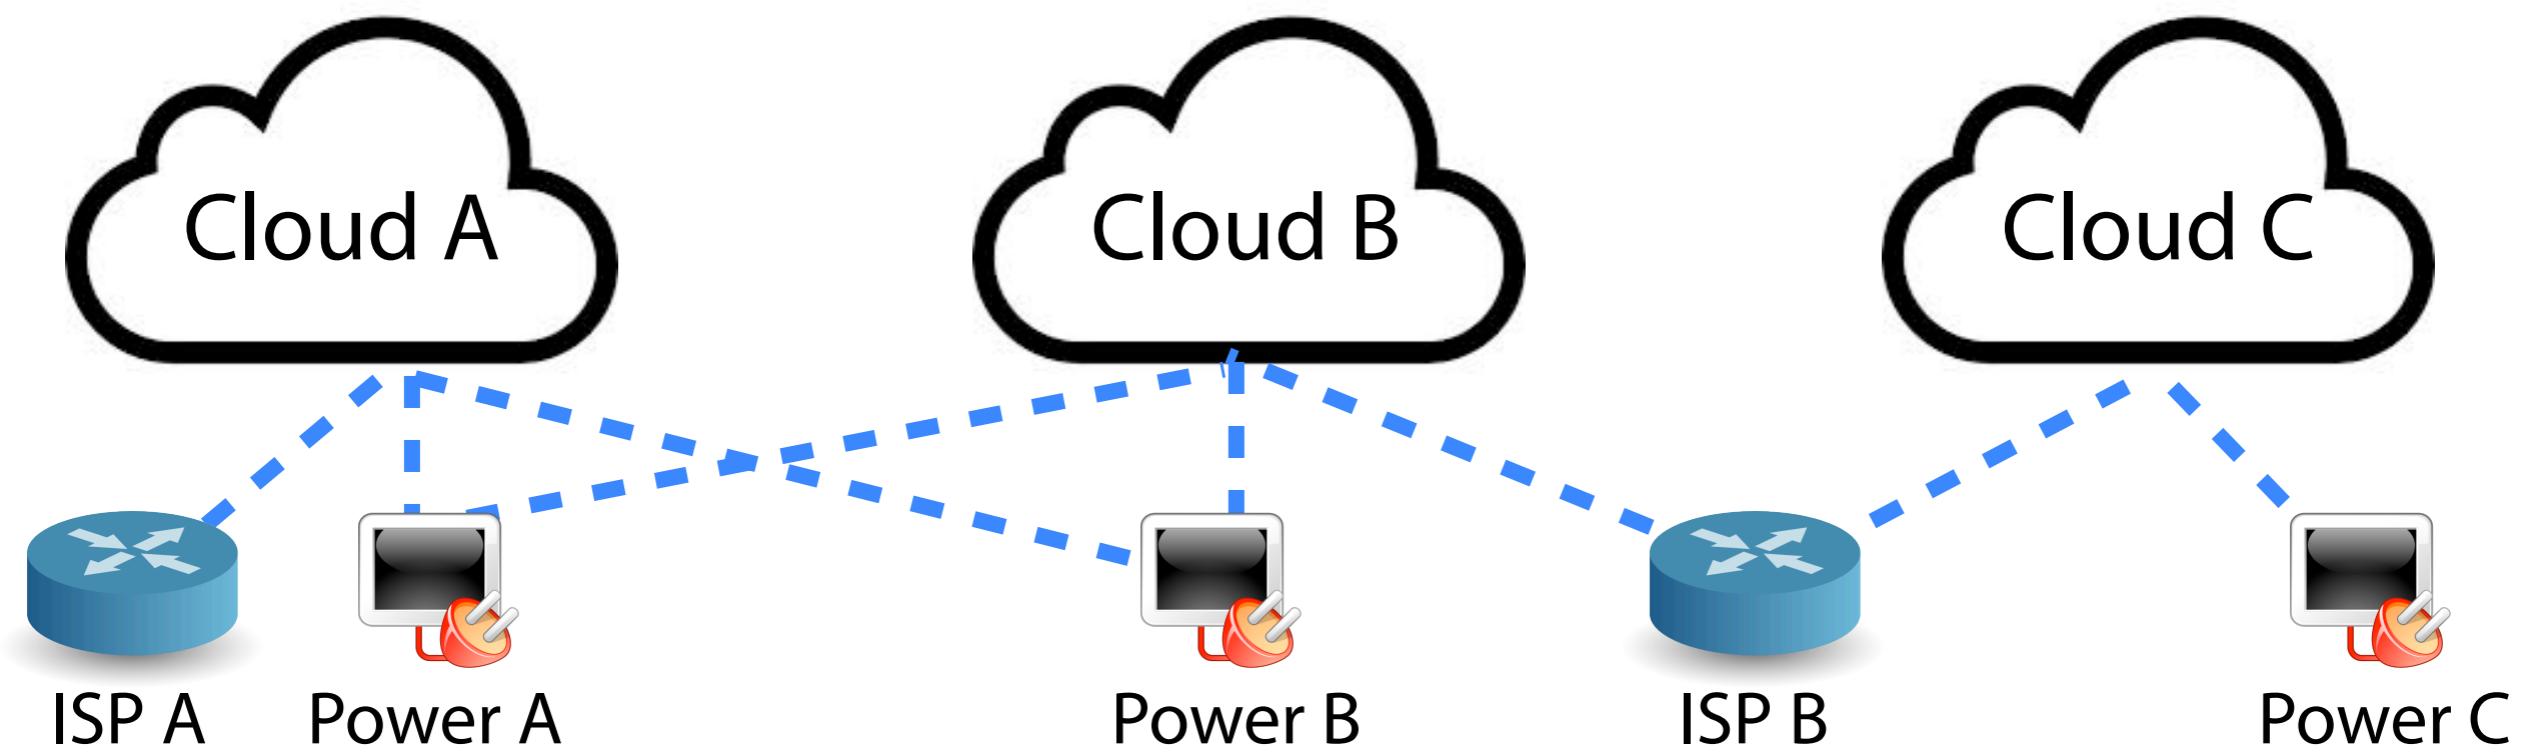
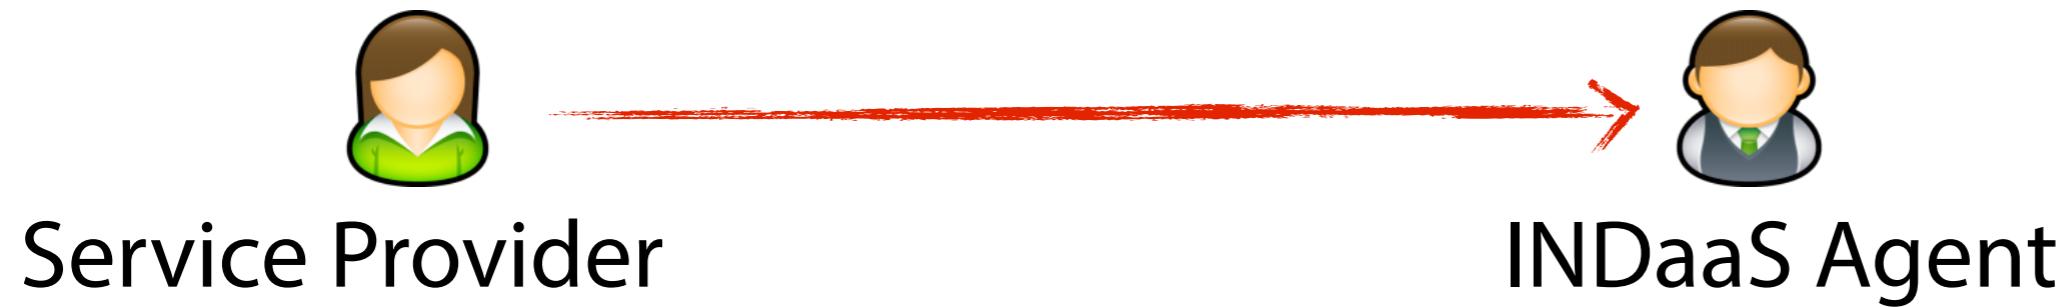


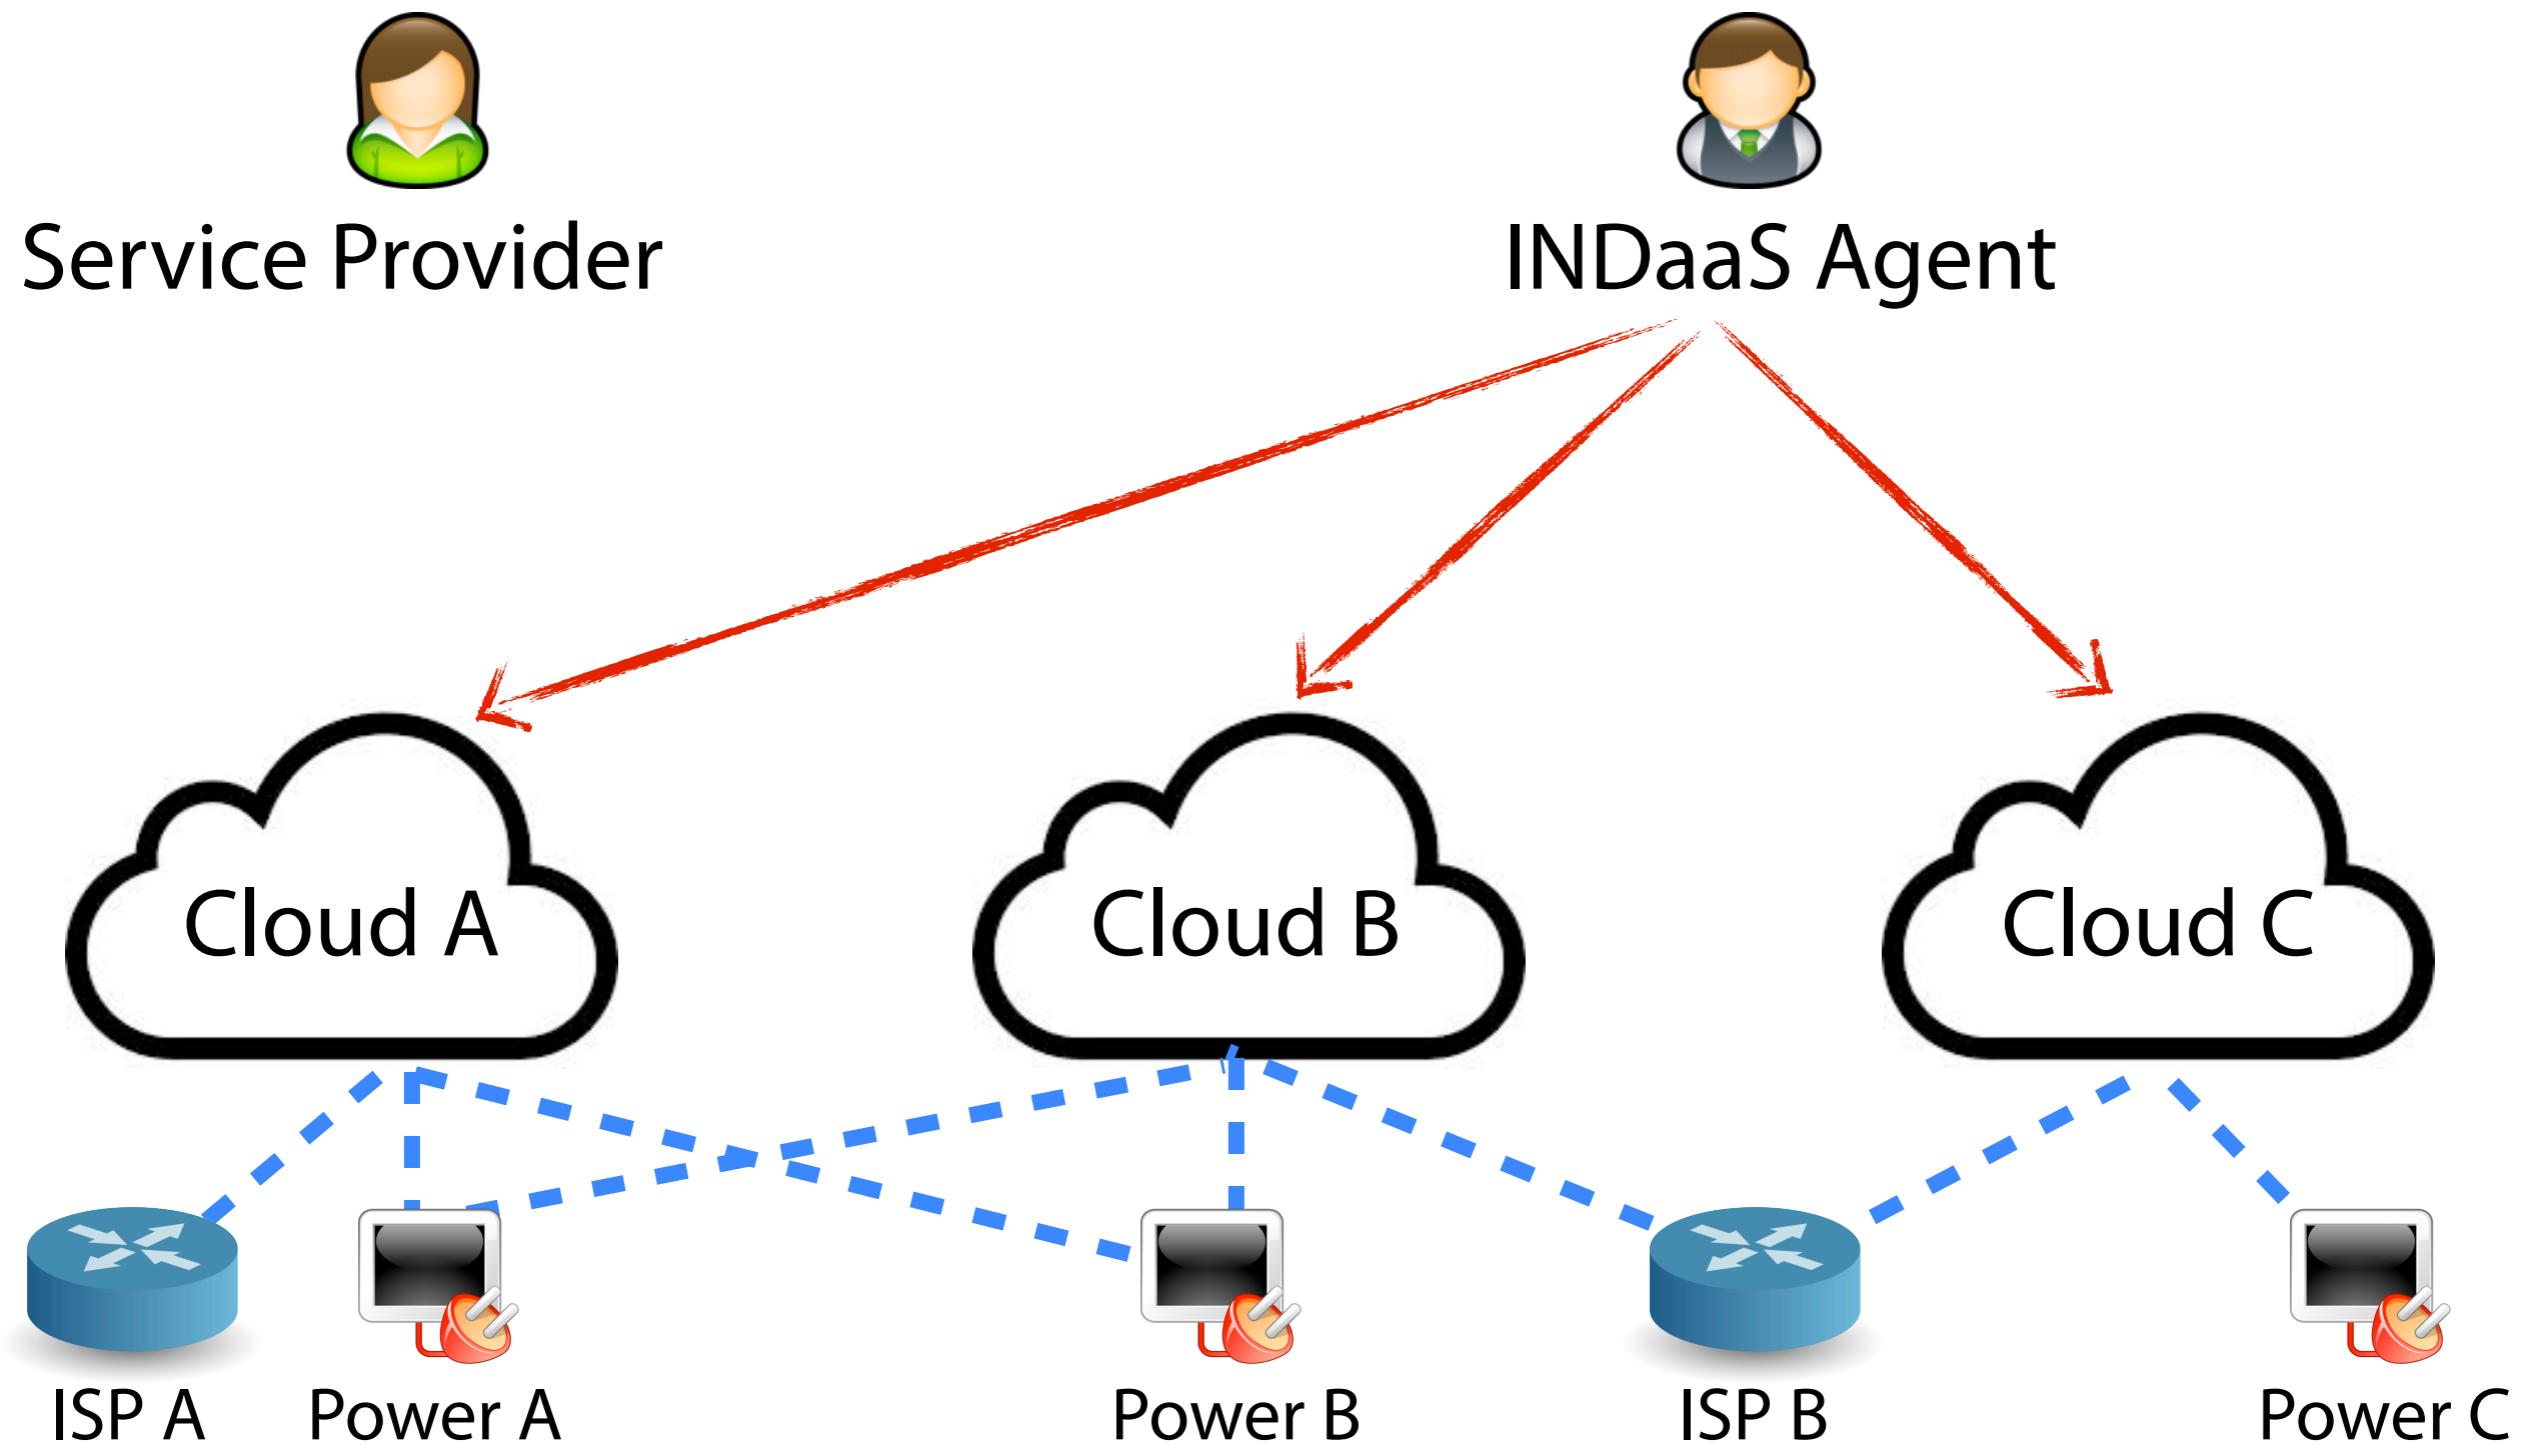
Each party maintains a commutative encryption key

Commutative encryption holds:  $E_x(E_y(m)) = E_y(E_x(m))$

# Private Independence Evaluation





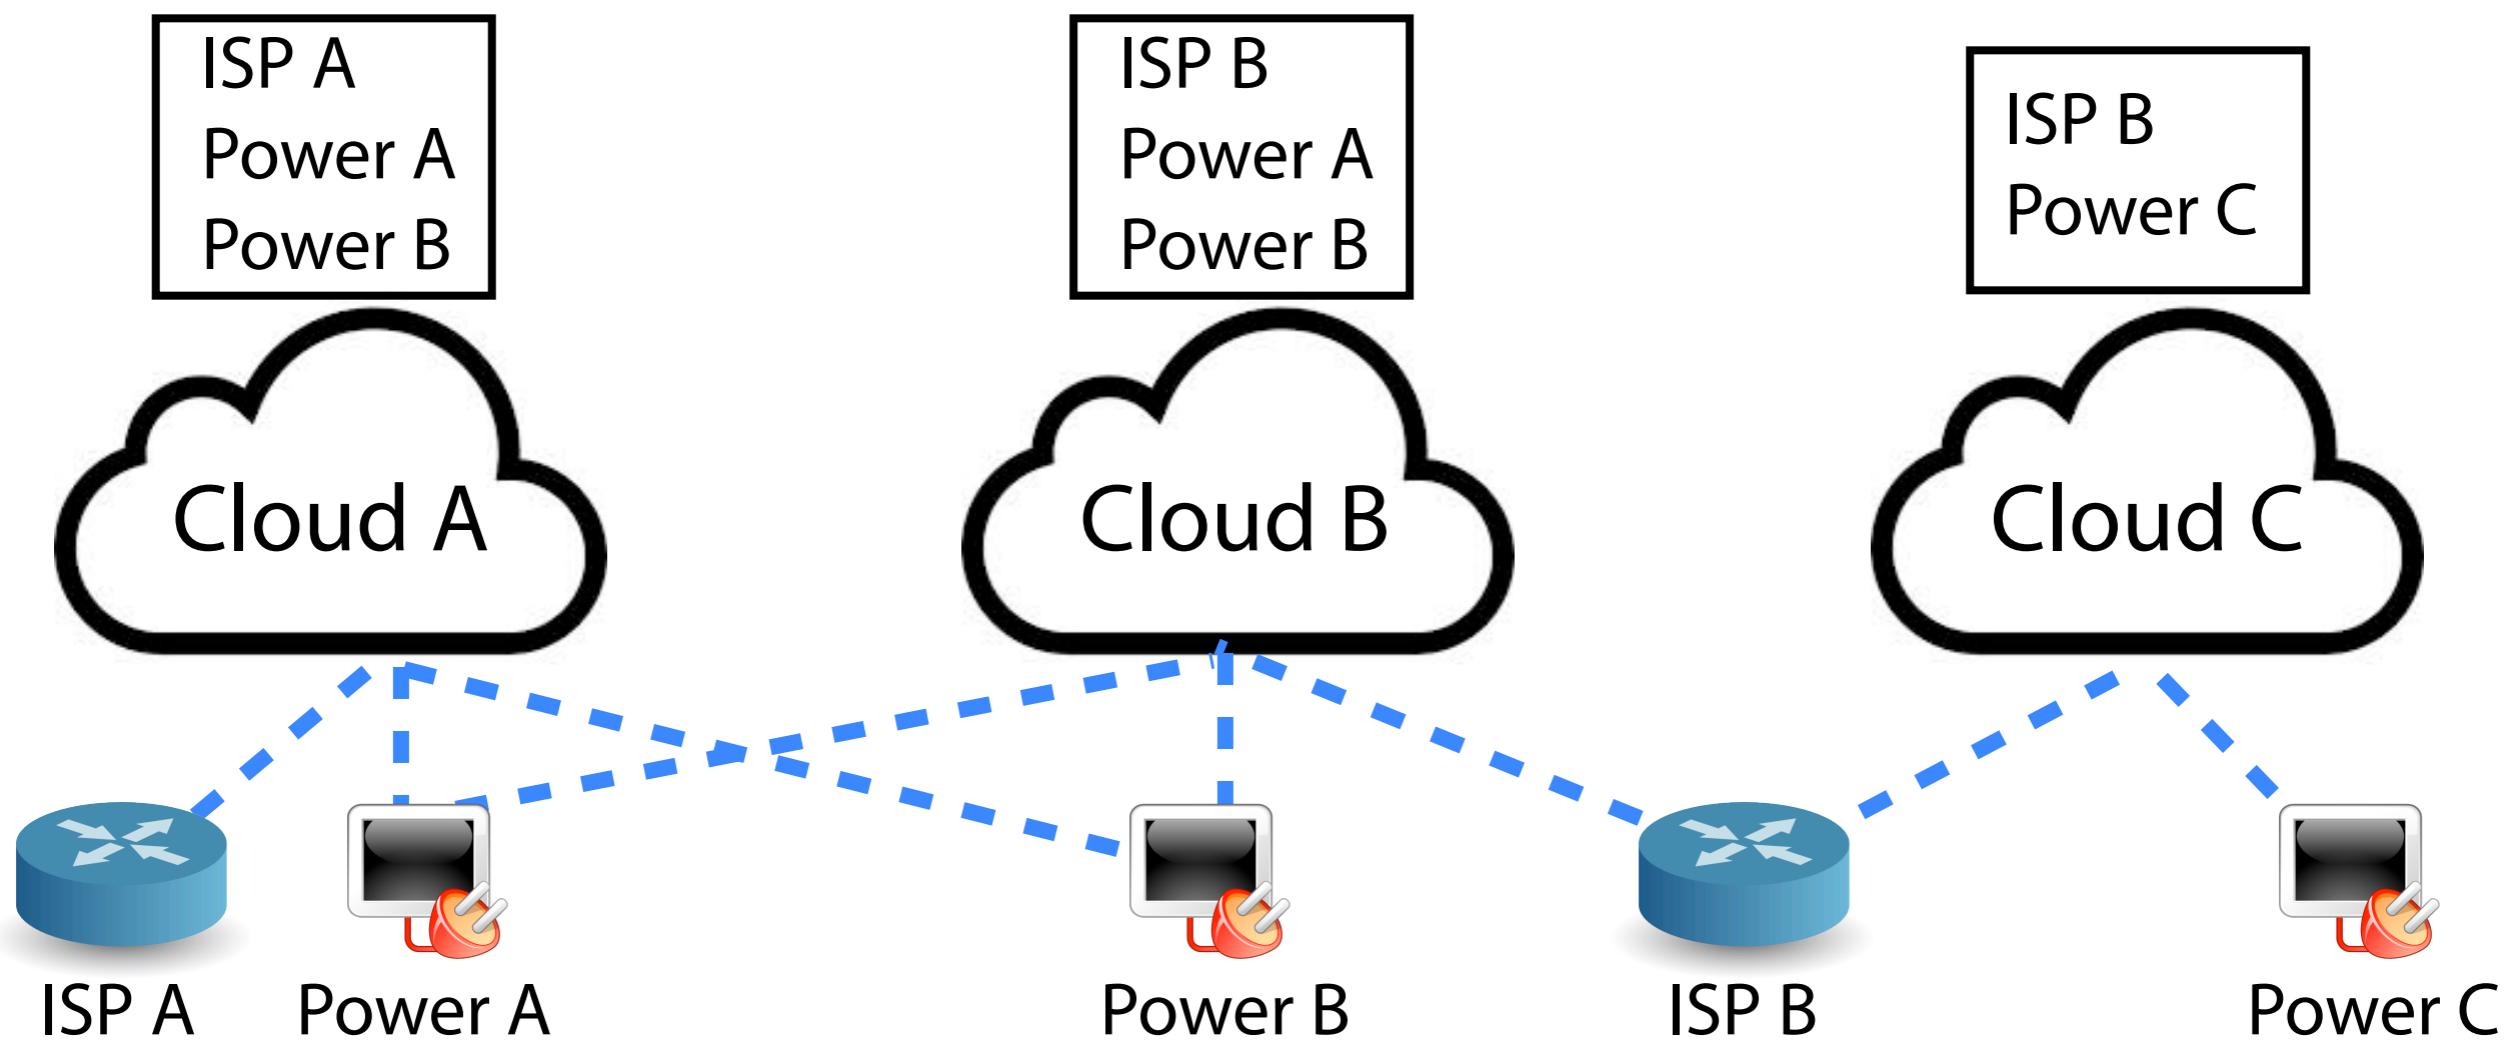


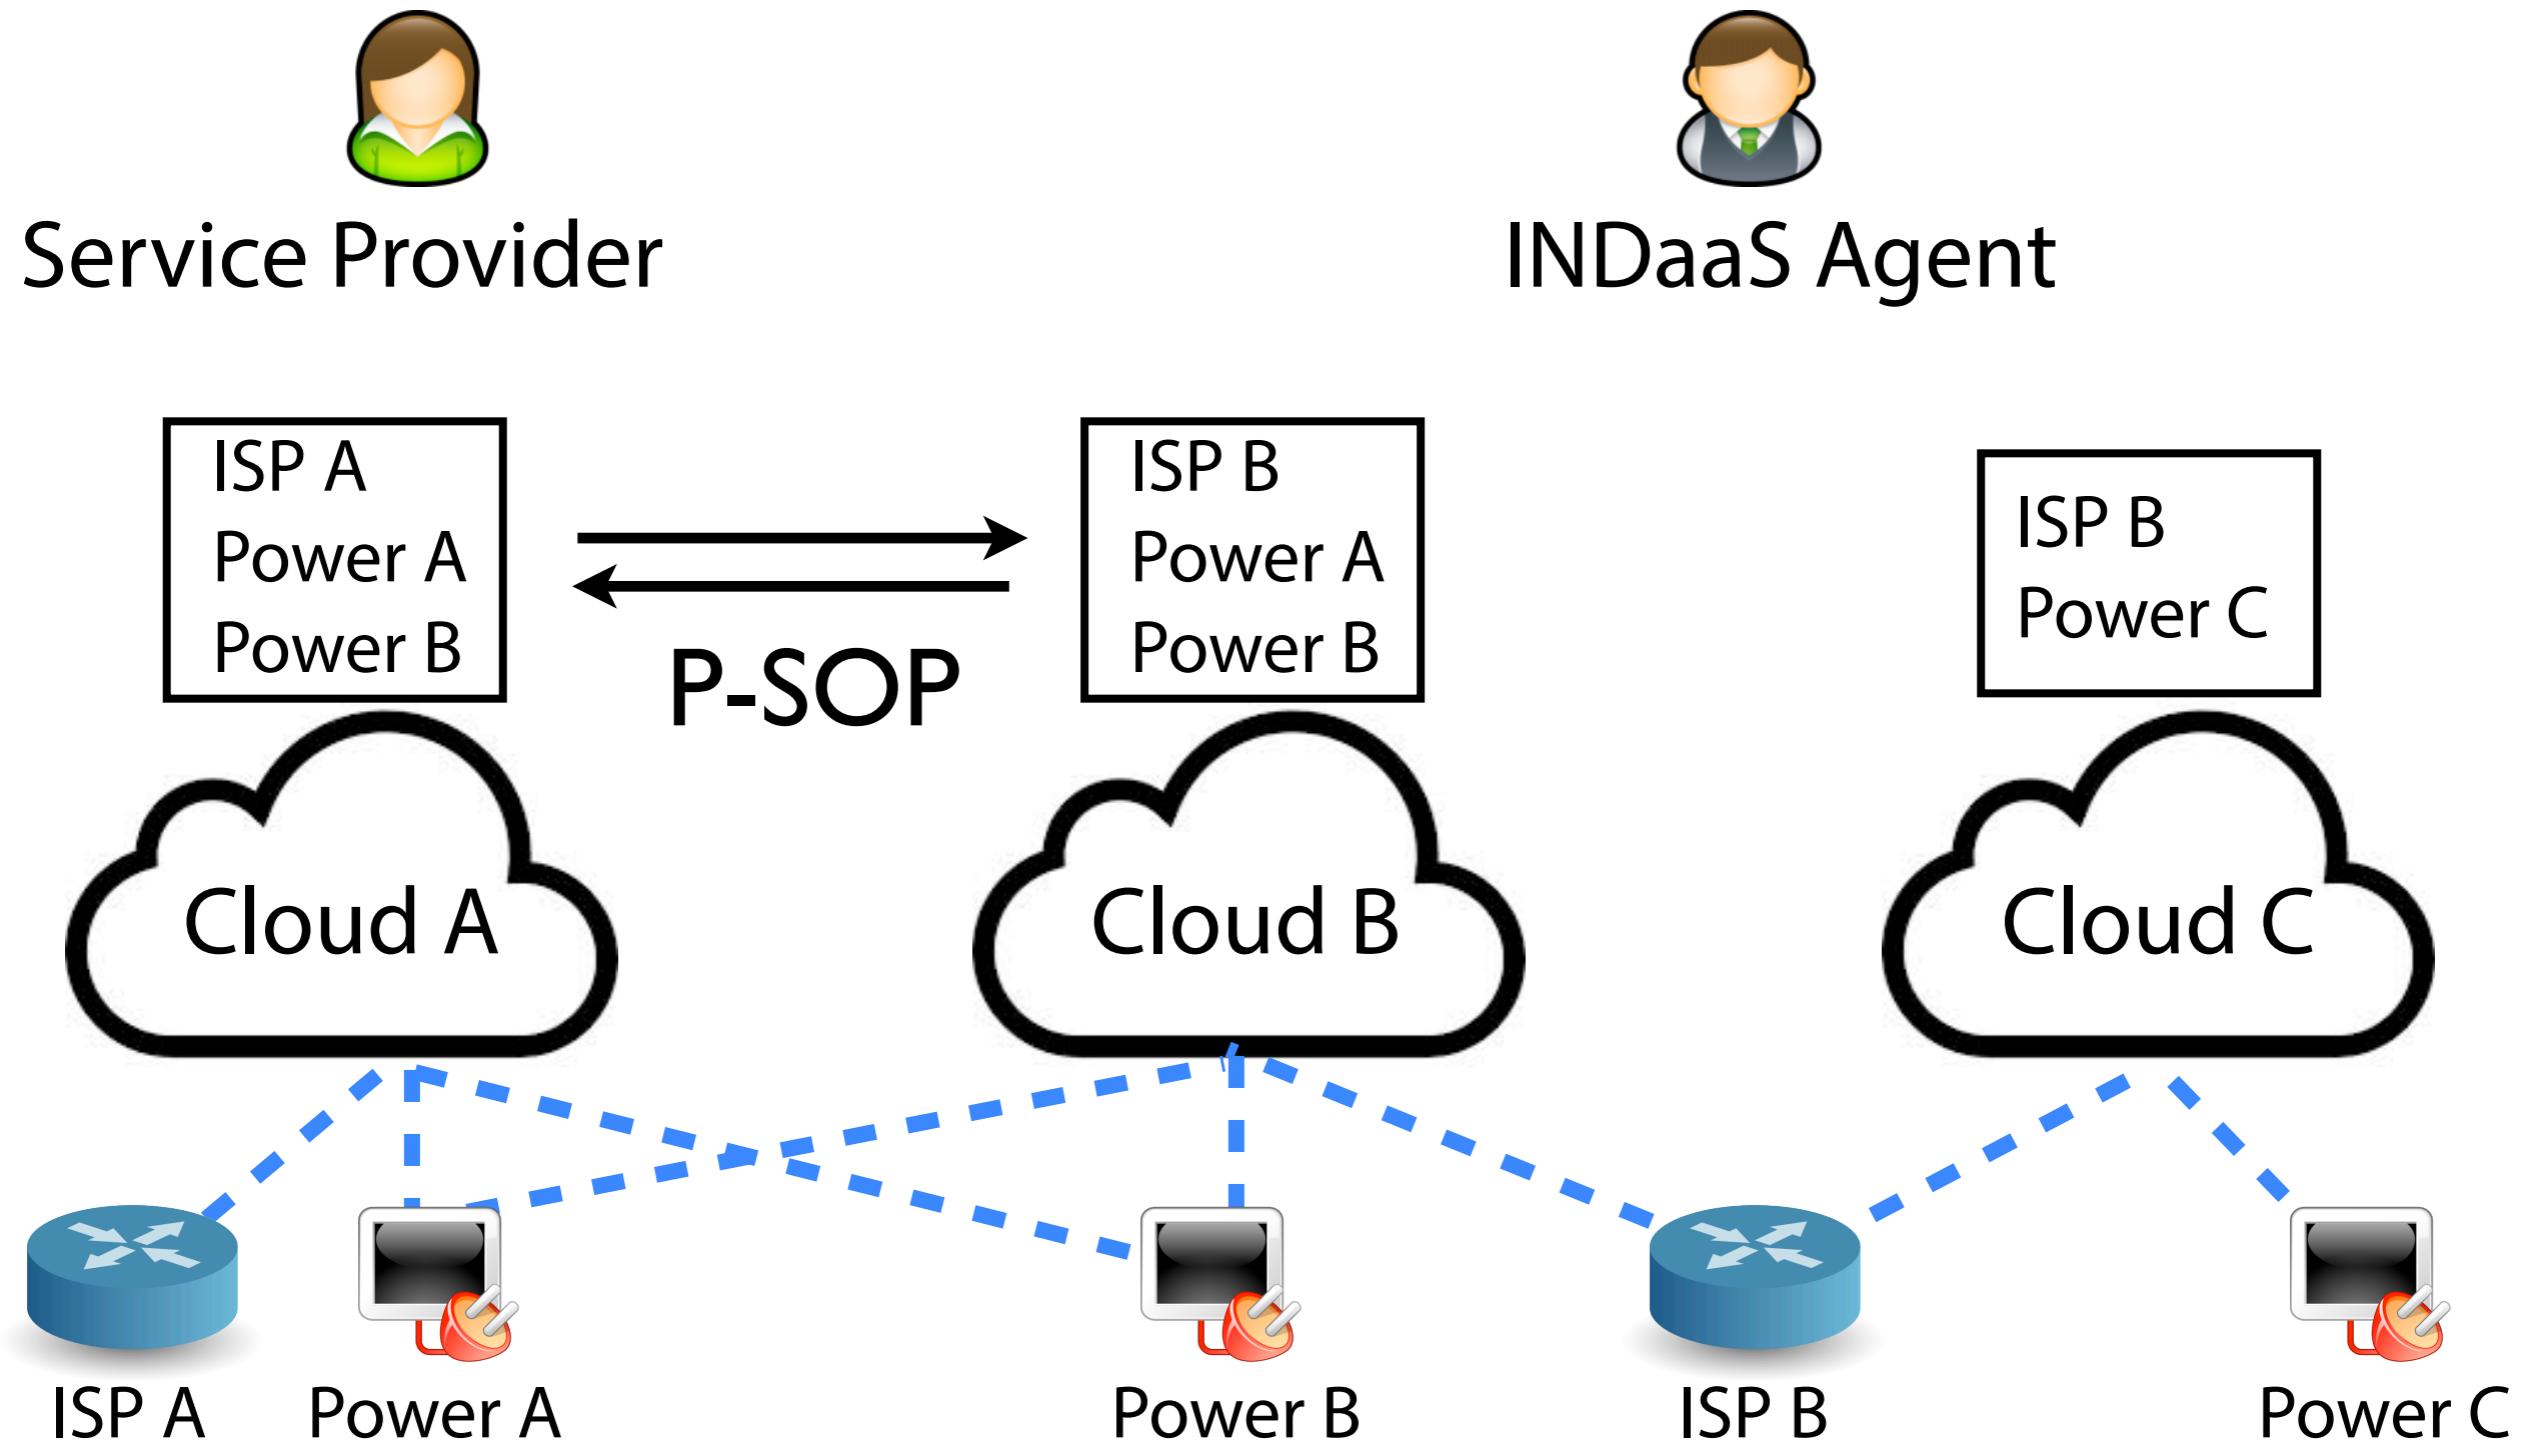


Service Provider



INDaaS Agent



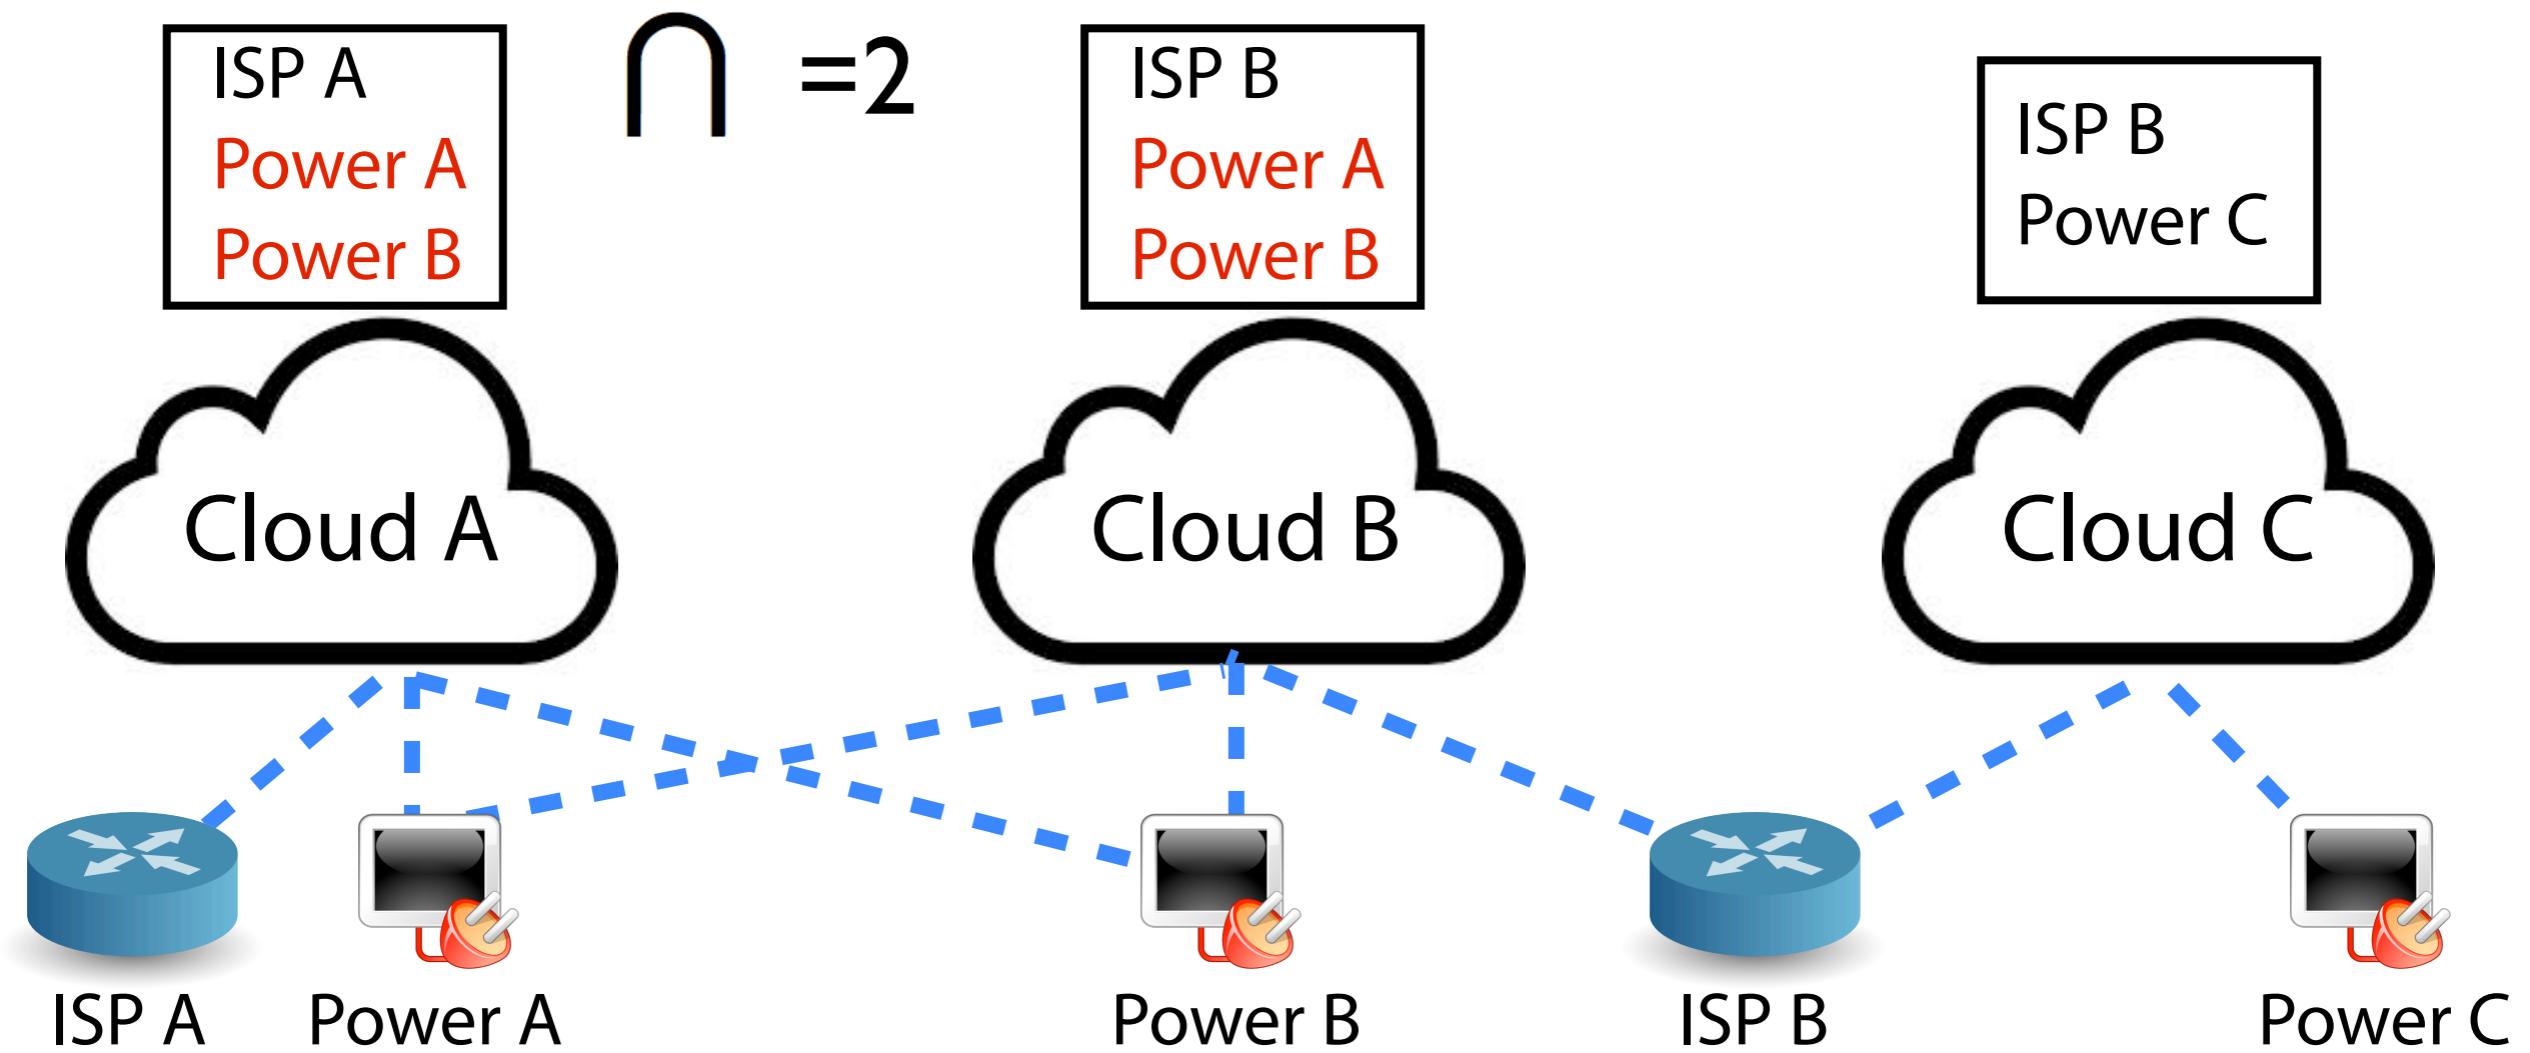




Service Provider



INDaaS Agent

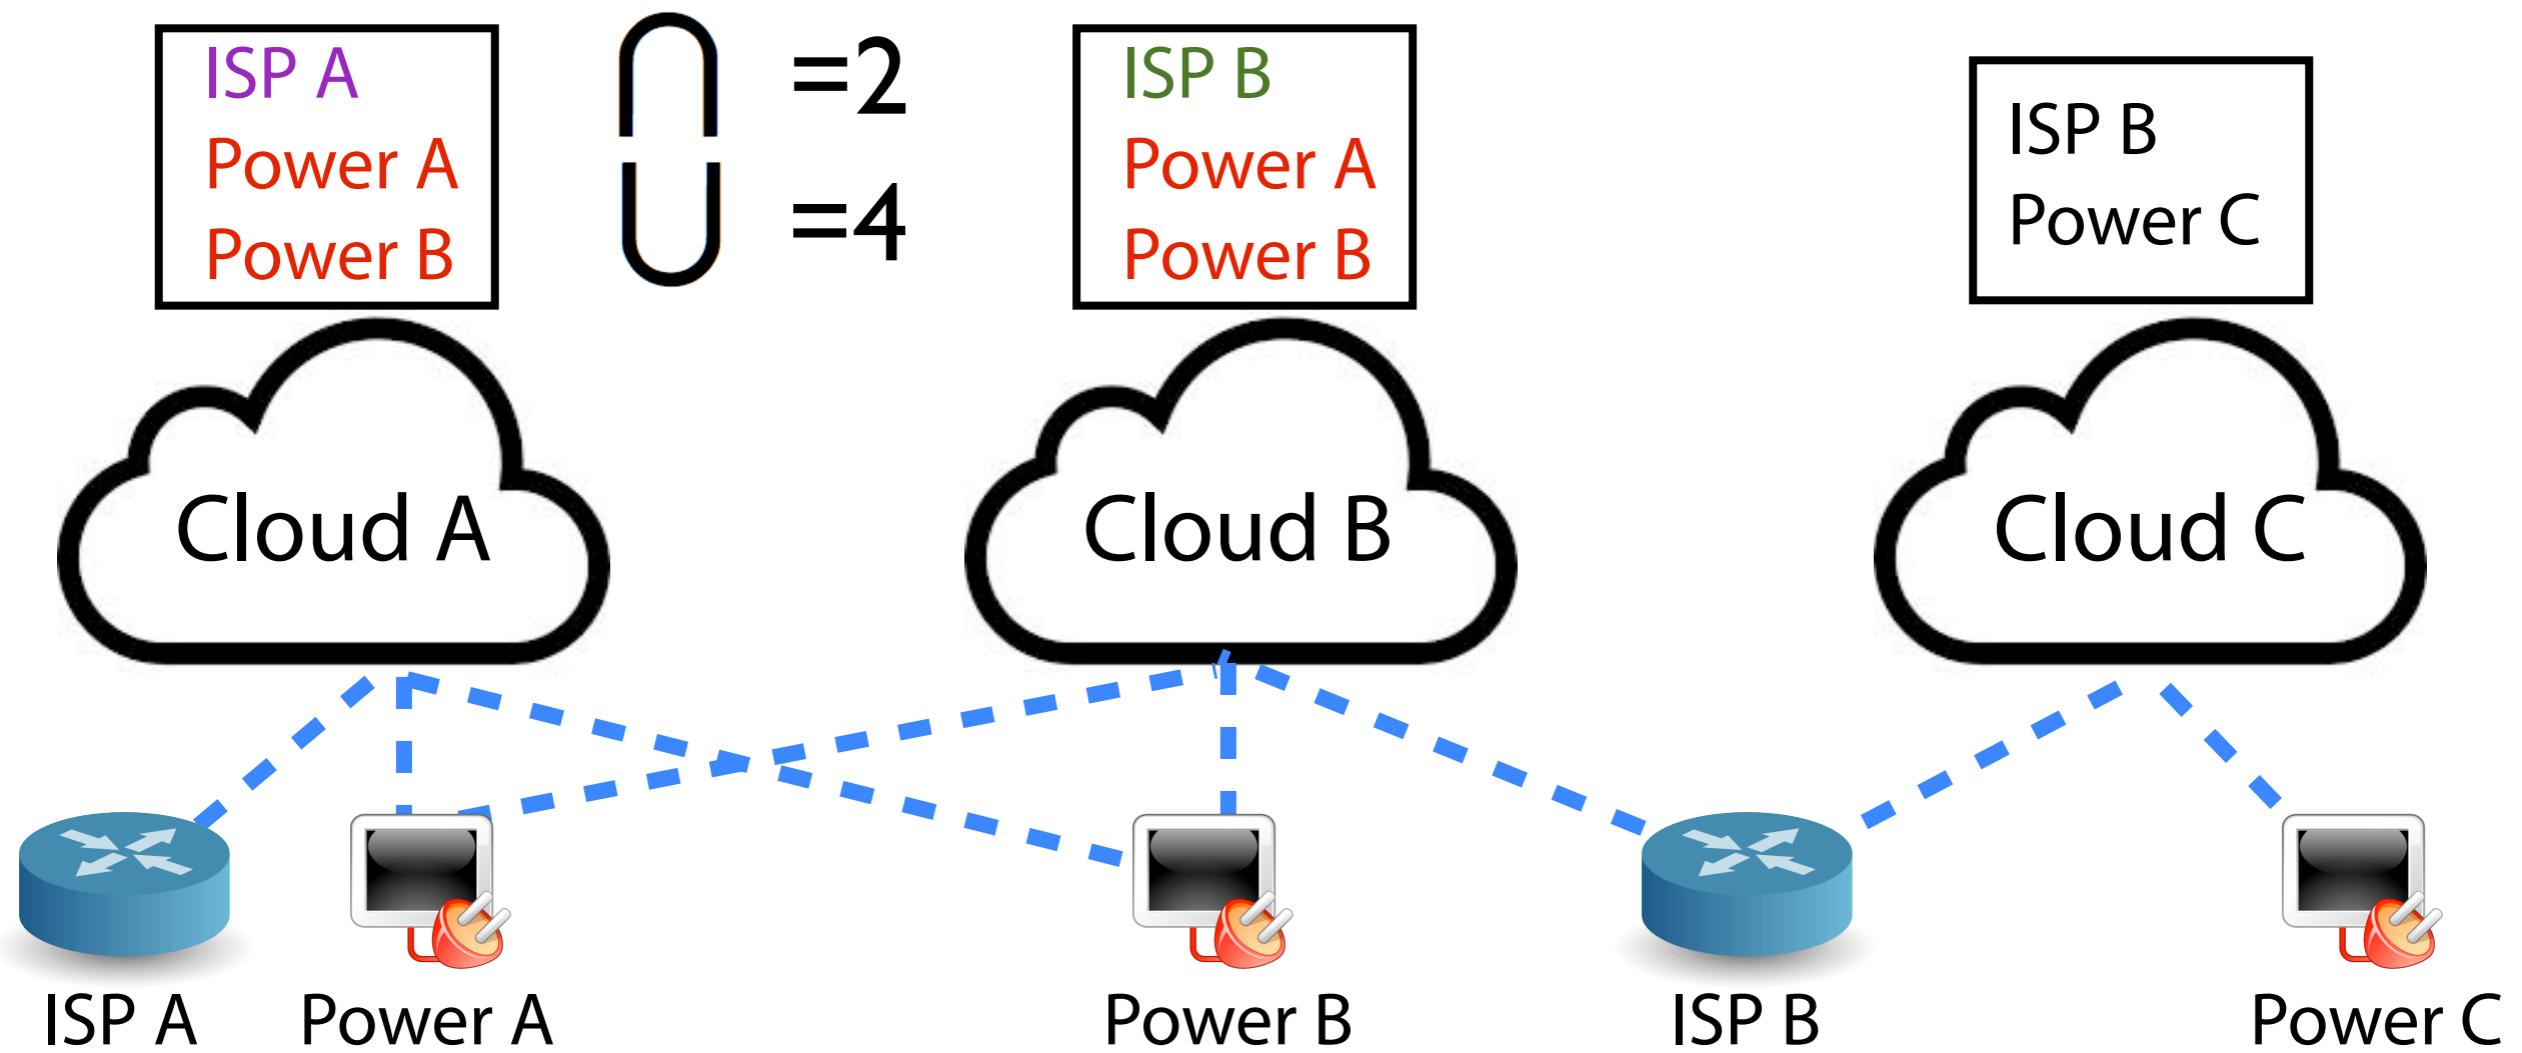




Service Provider

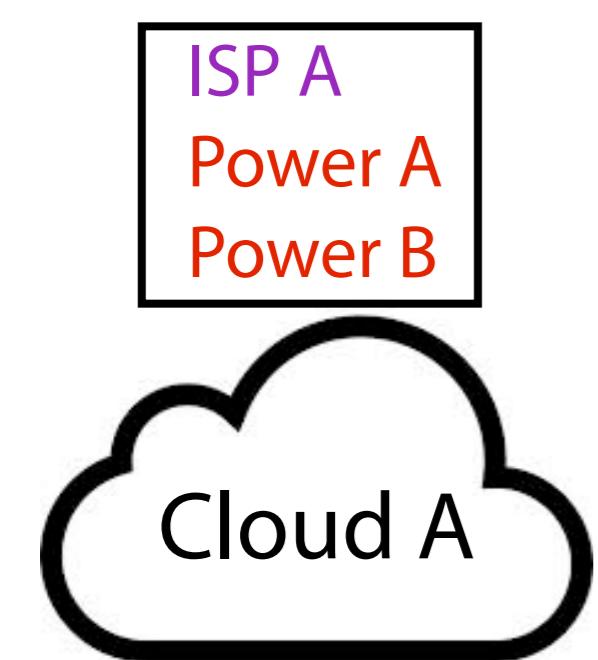


INDaaS Agent

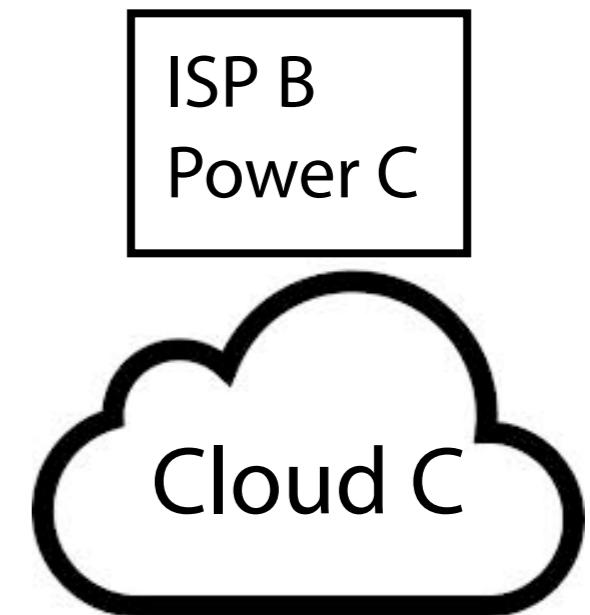
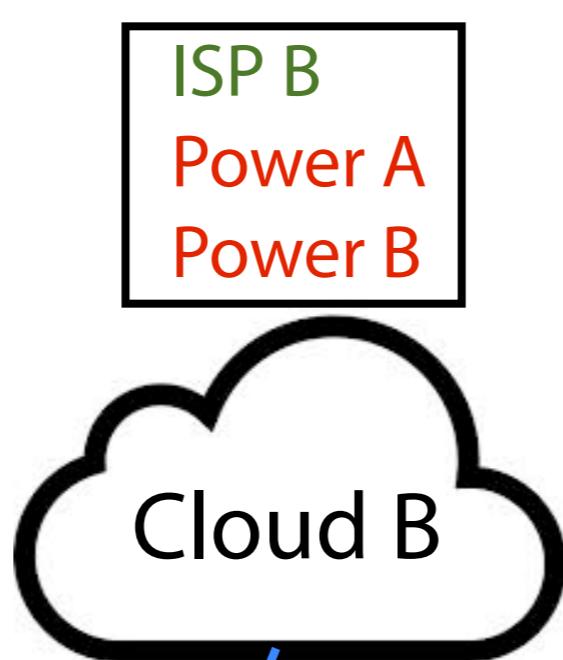




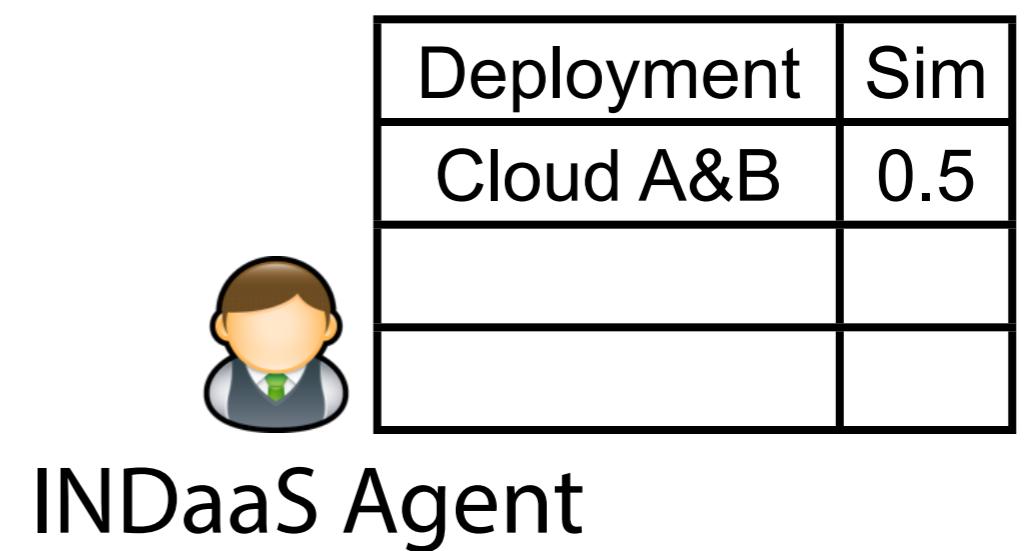
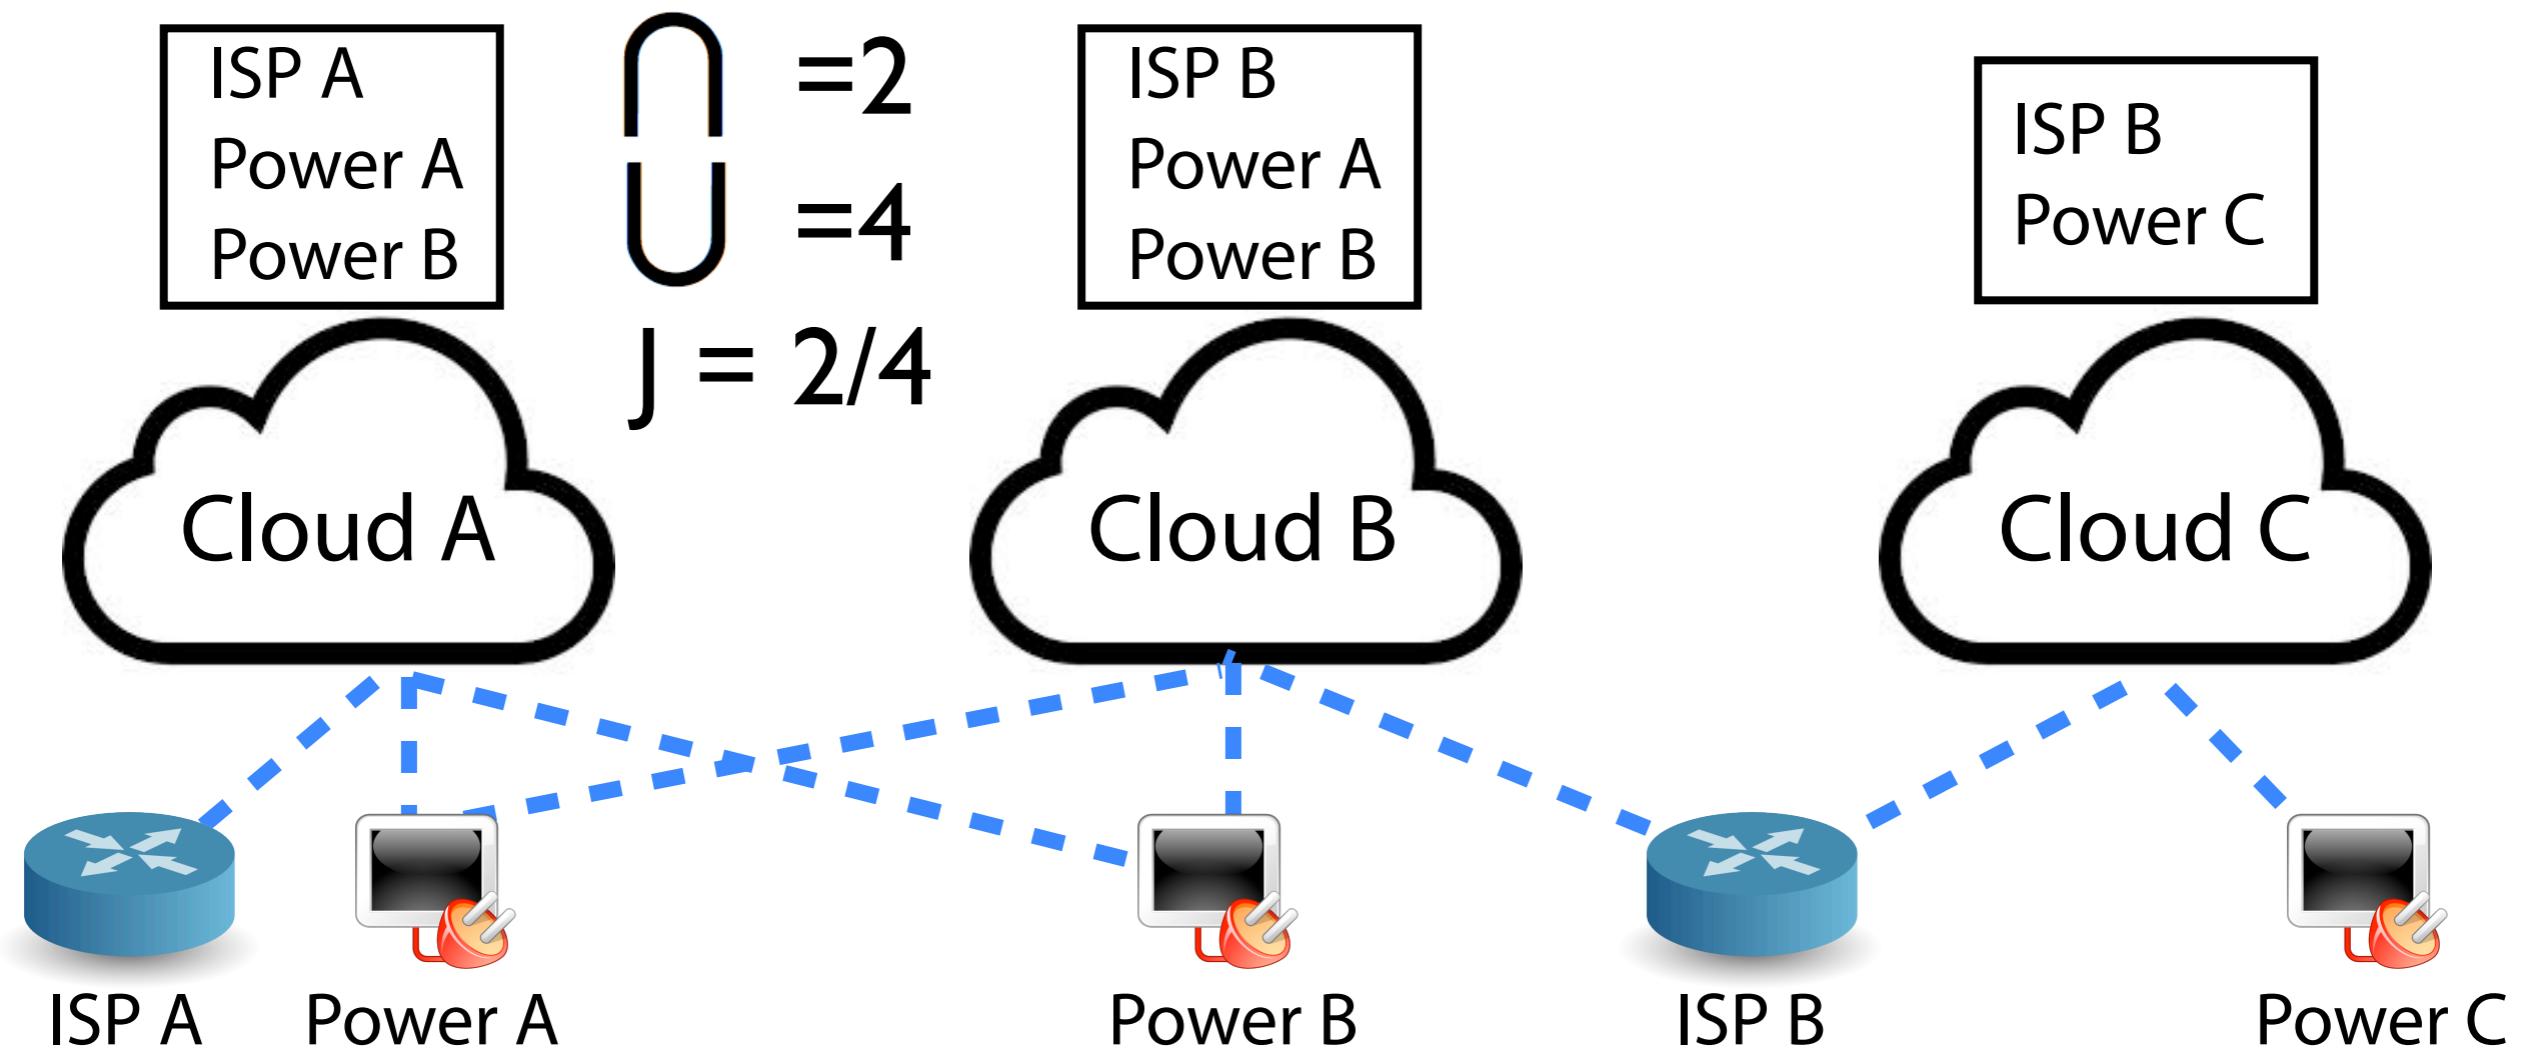
Service Provider



$$\begin{matrix} 0 \\ \cup \\ =2 \\ =4 \\ J = 2/4 \end{matrix}$$

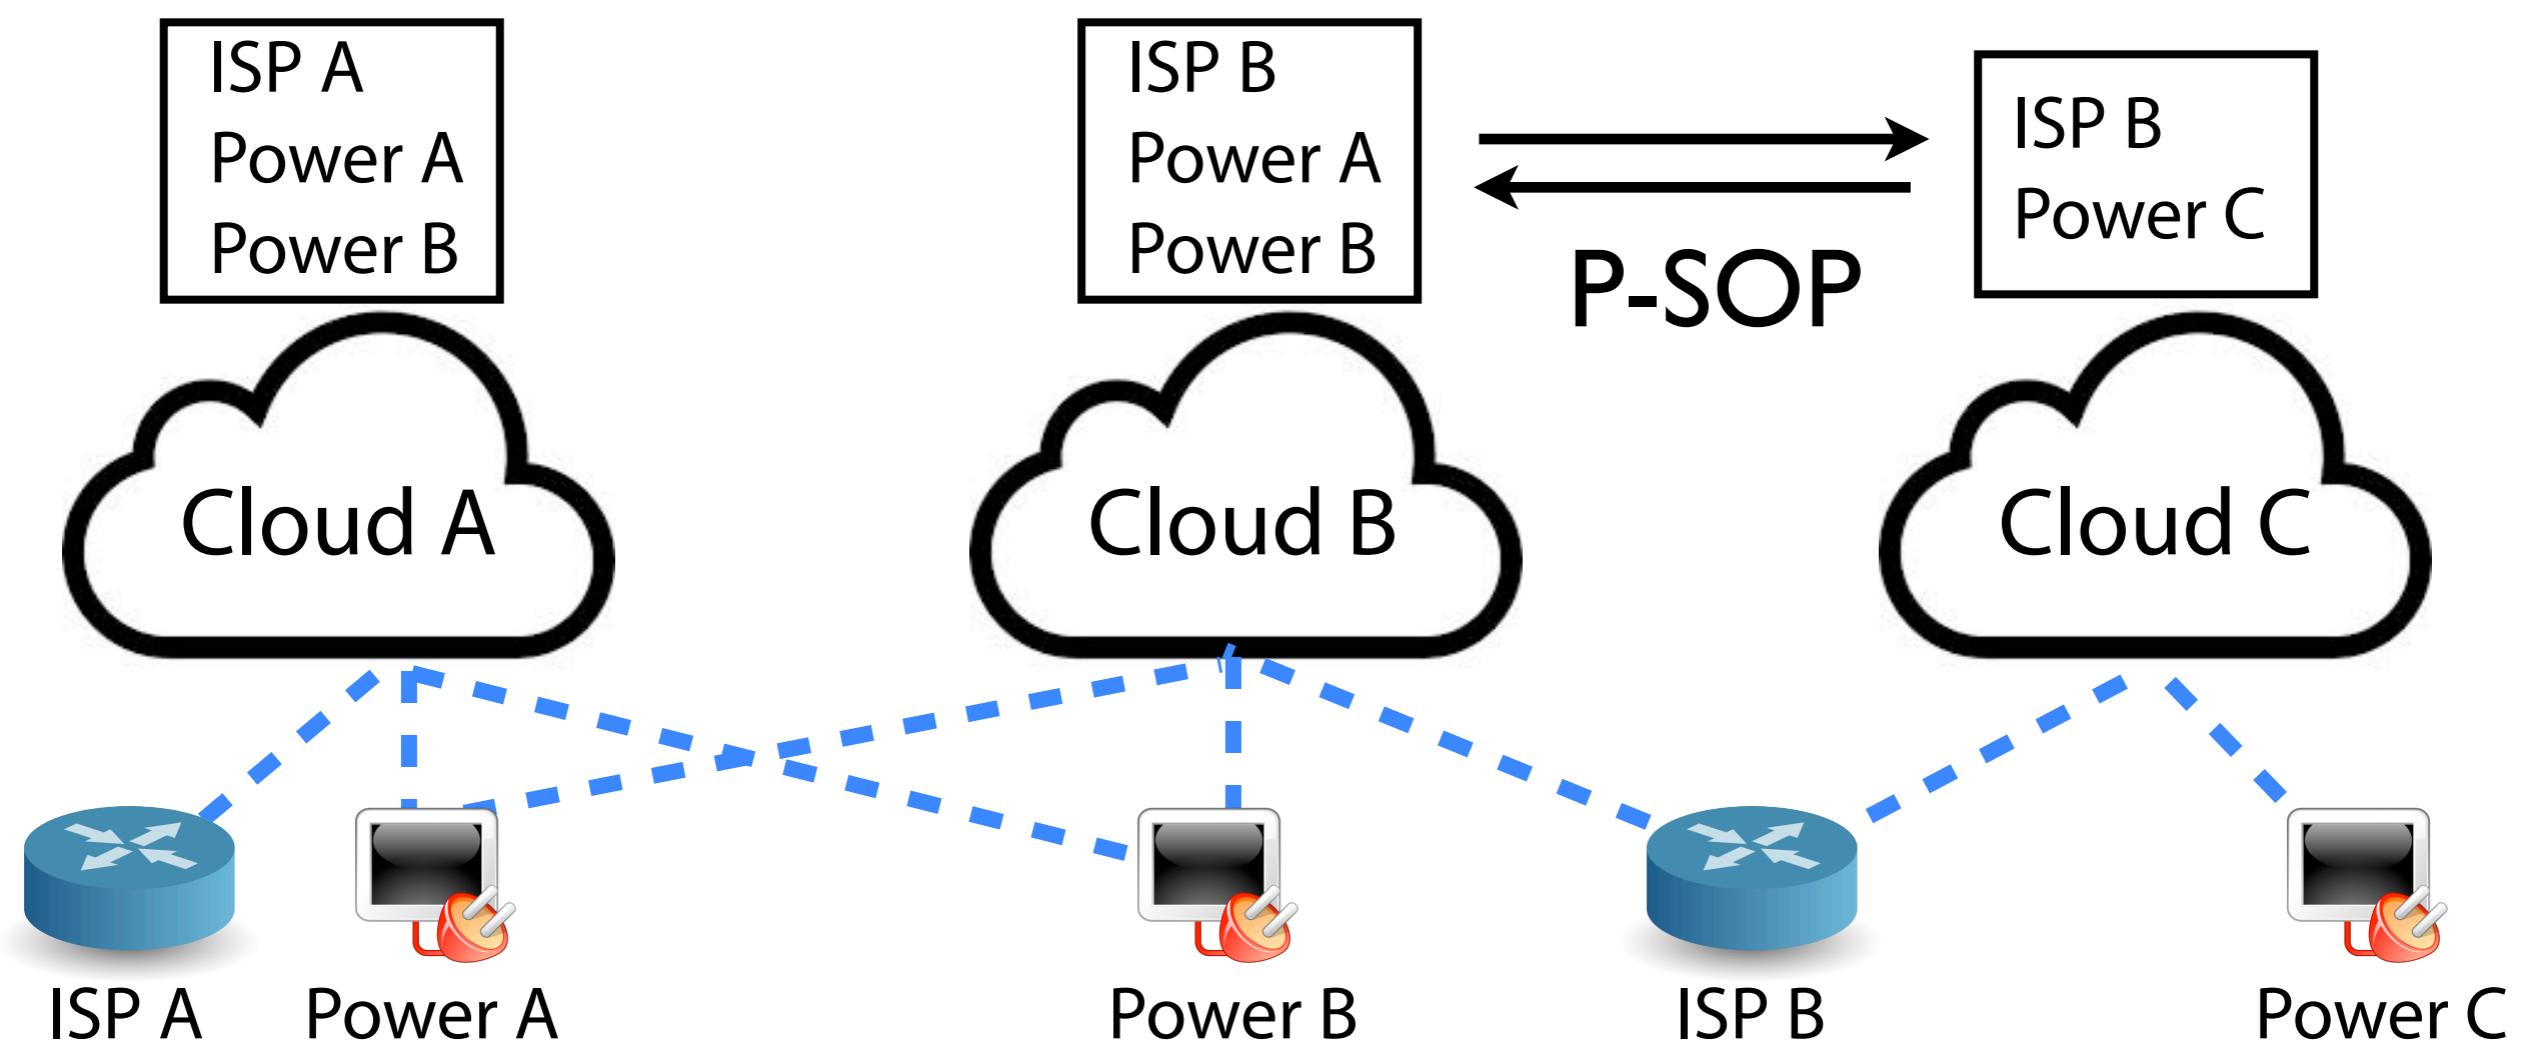
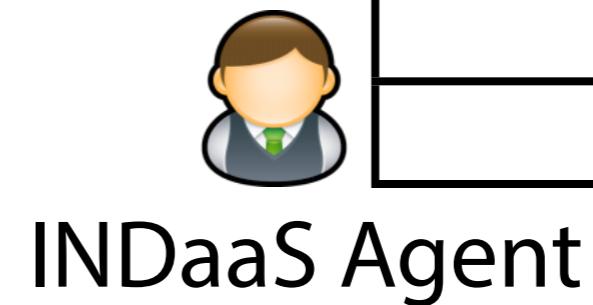


INDaaS Agent



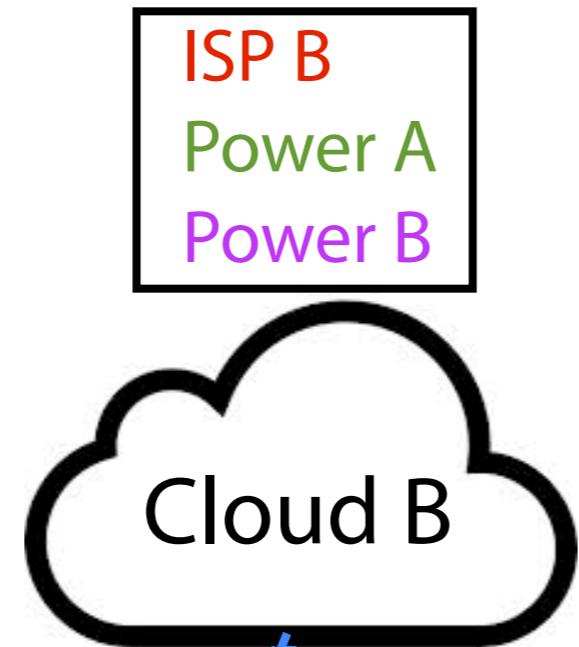
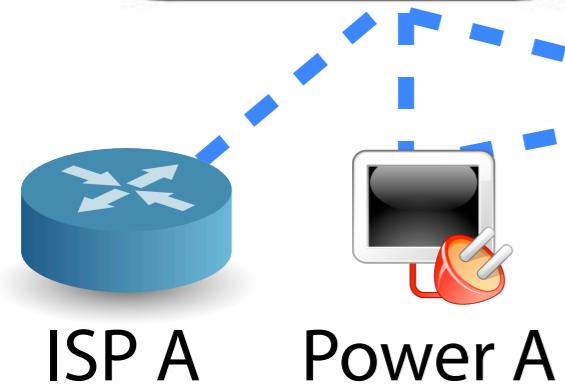
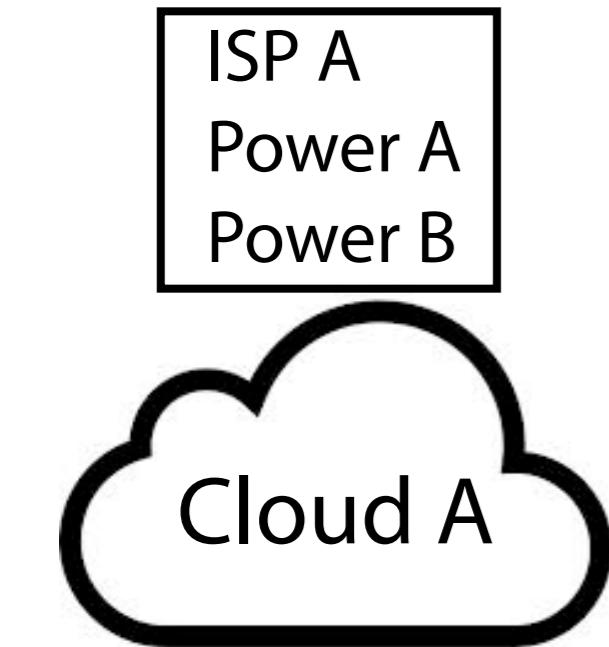


| Deployment | Sim |
|------------|-----|
| Cloud A&B  | 0.5 |
|            |     |
|            |     |



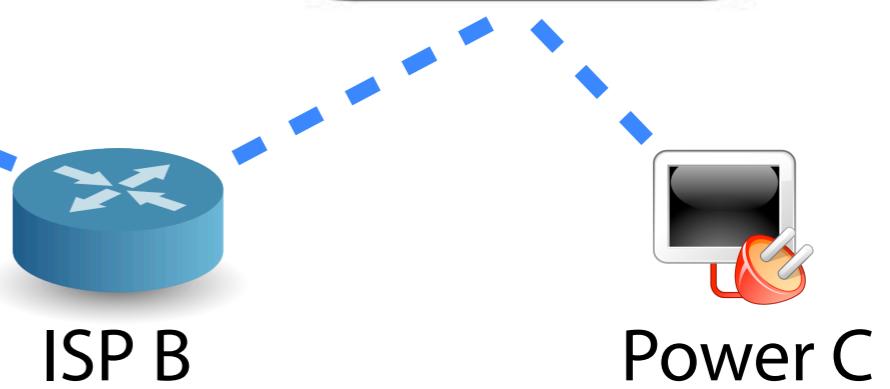
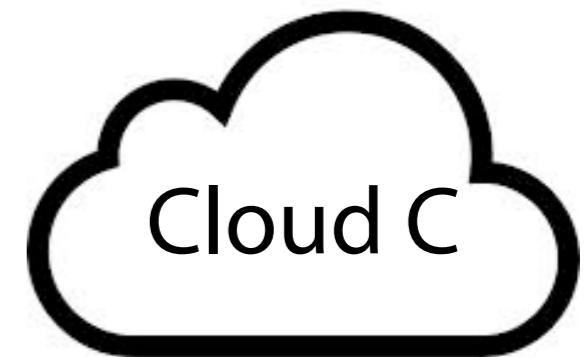


## Service Provider



## INDaaS Agent

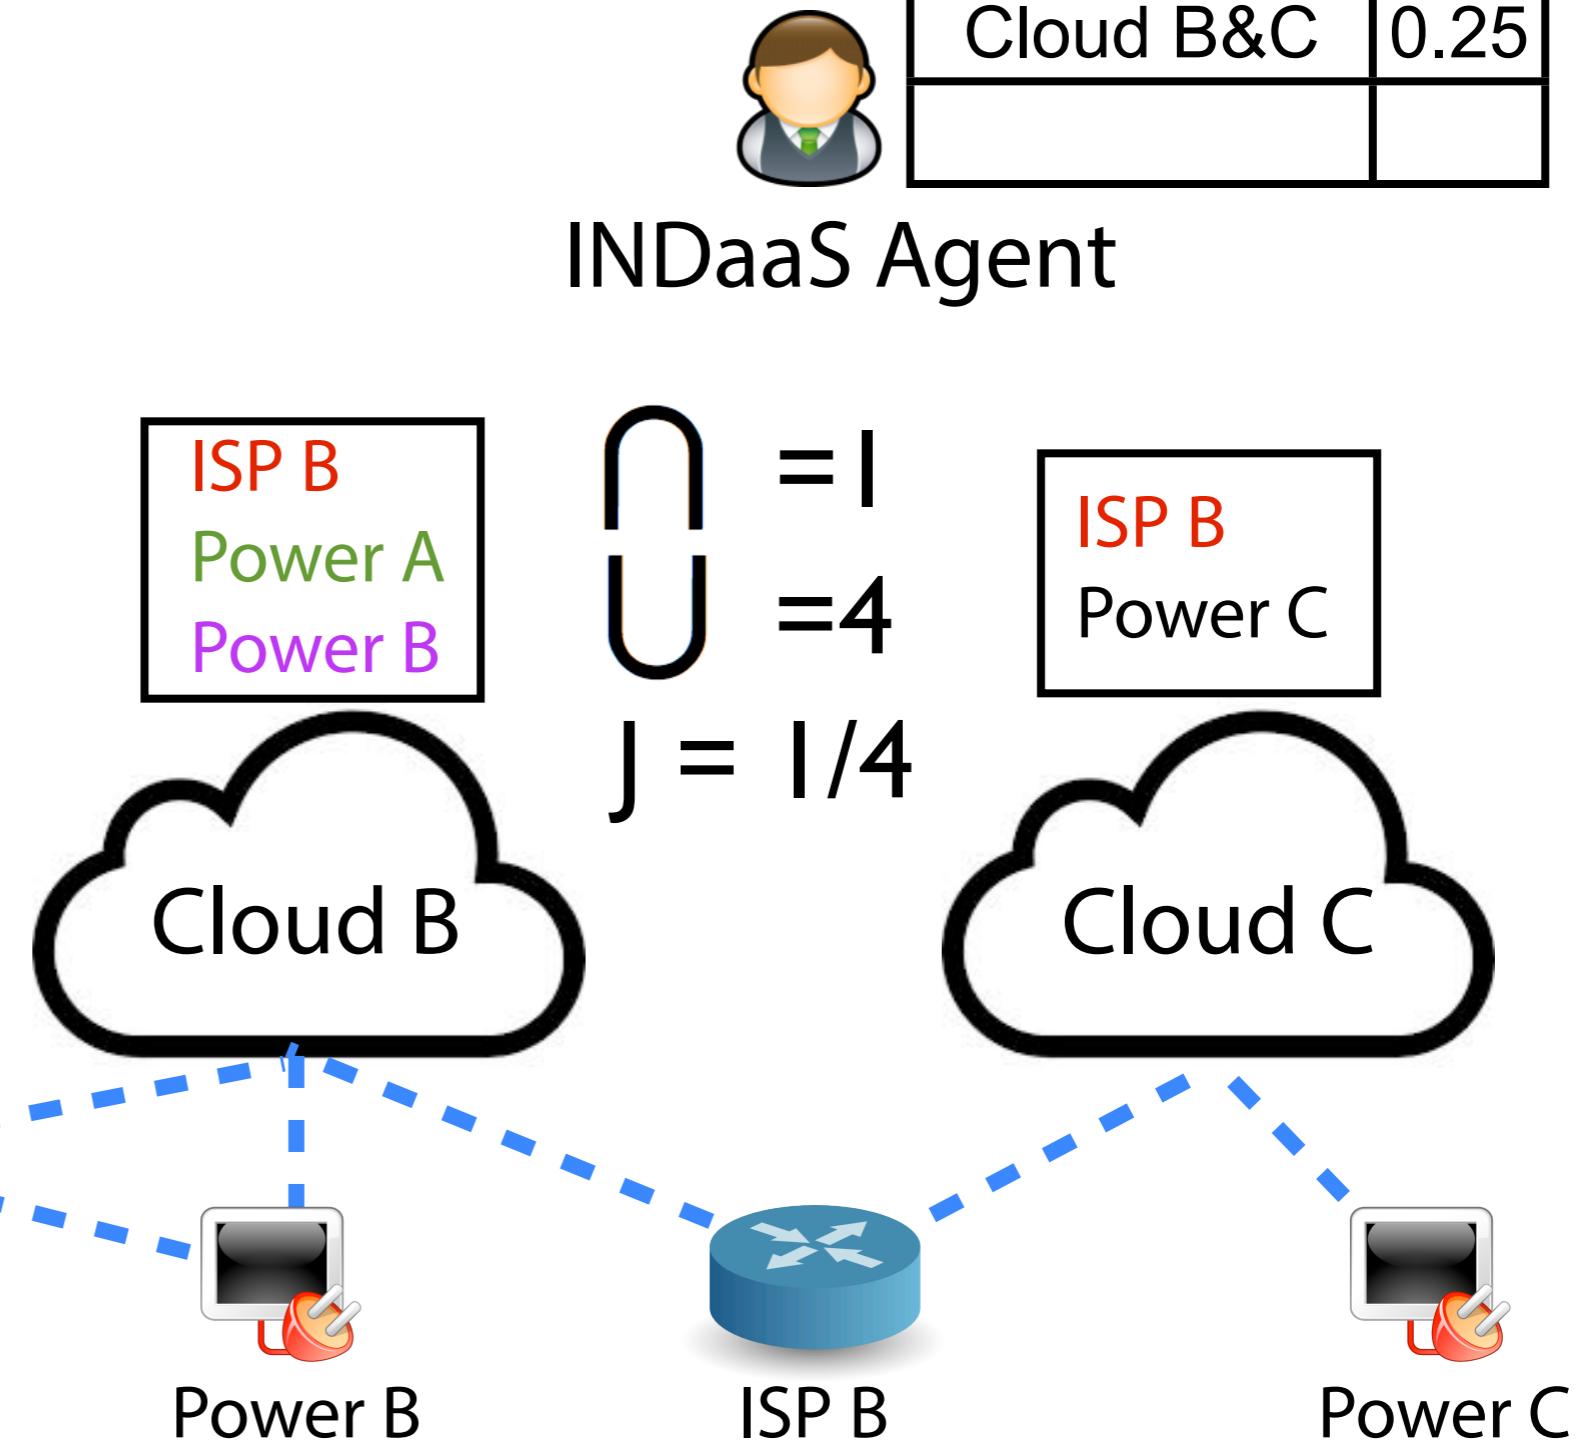
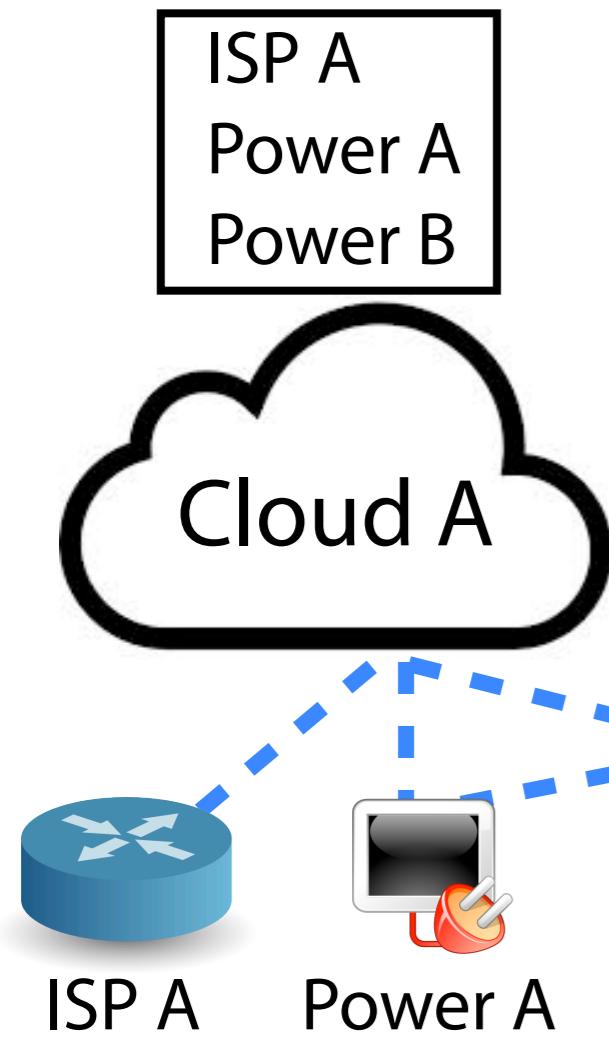
$$\begin{aligned} 0 &= 1 \\ &= 4 \\ J &= 1/4 \end{aligned}$$



| Deployment | Sim |
|------------|-----|
| Cloud A&B  | 0.5 |
|            |     |
|            |     |

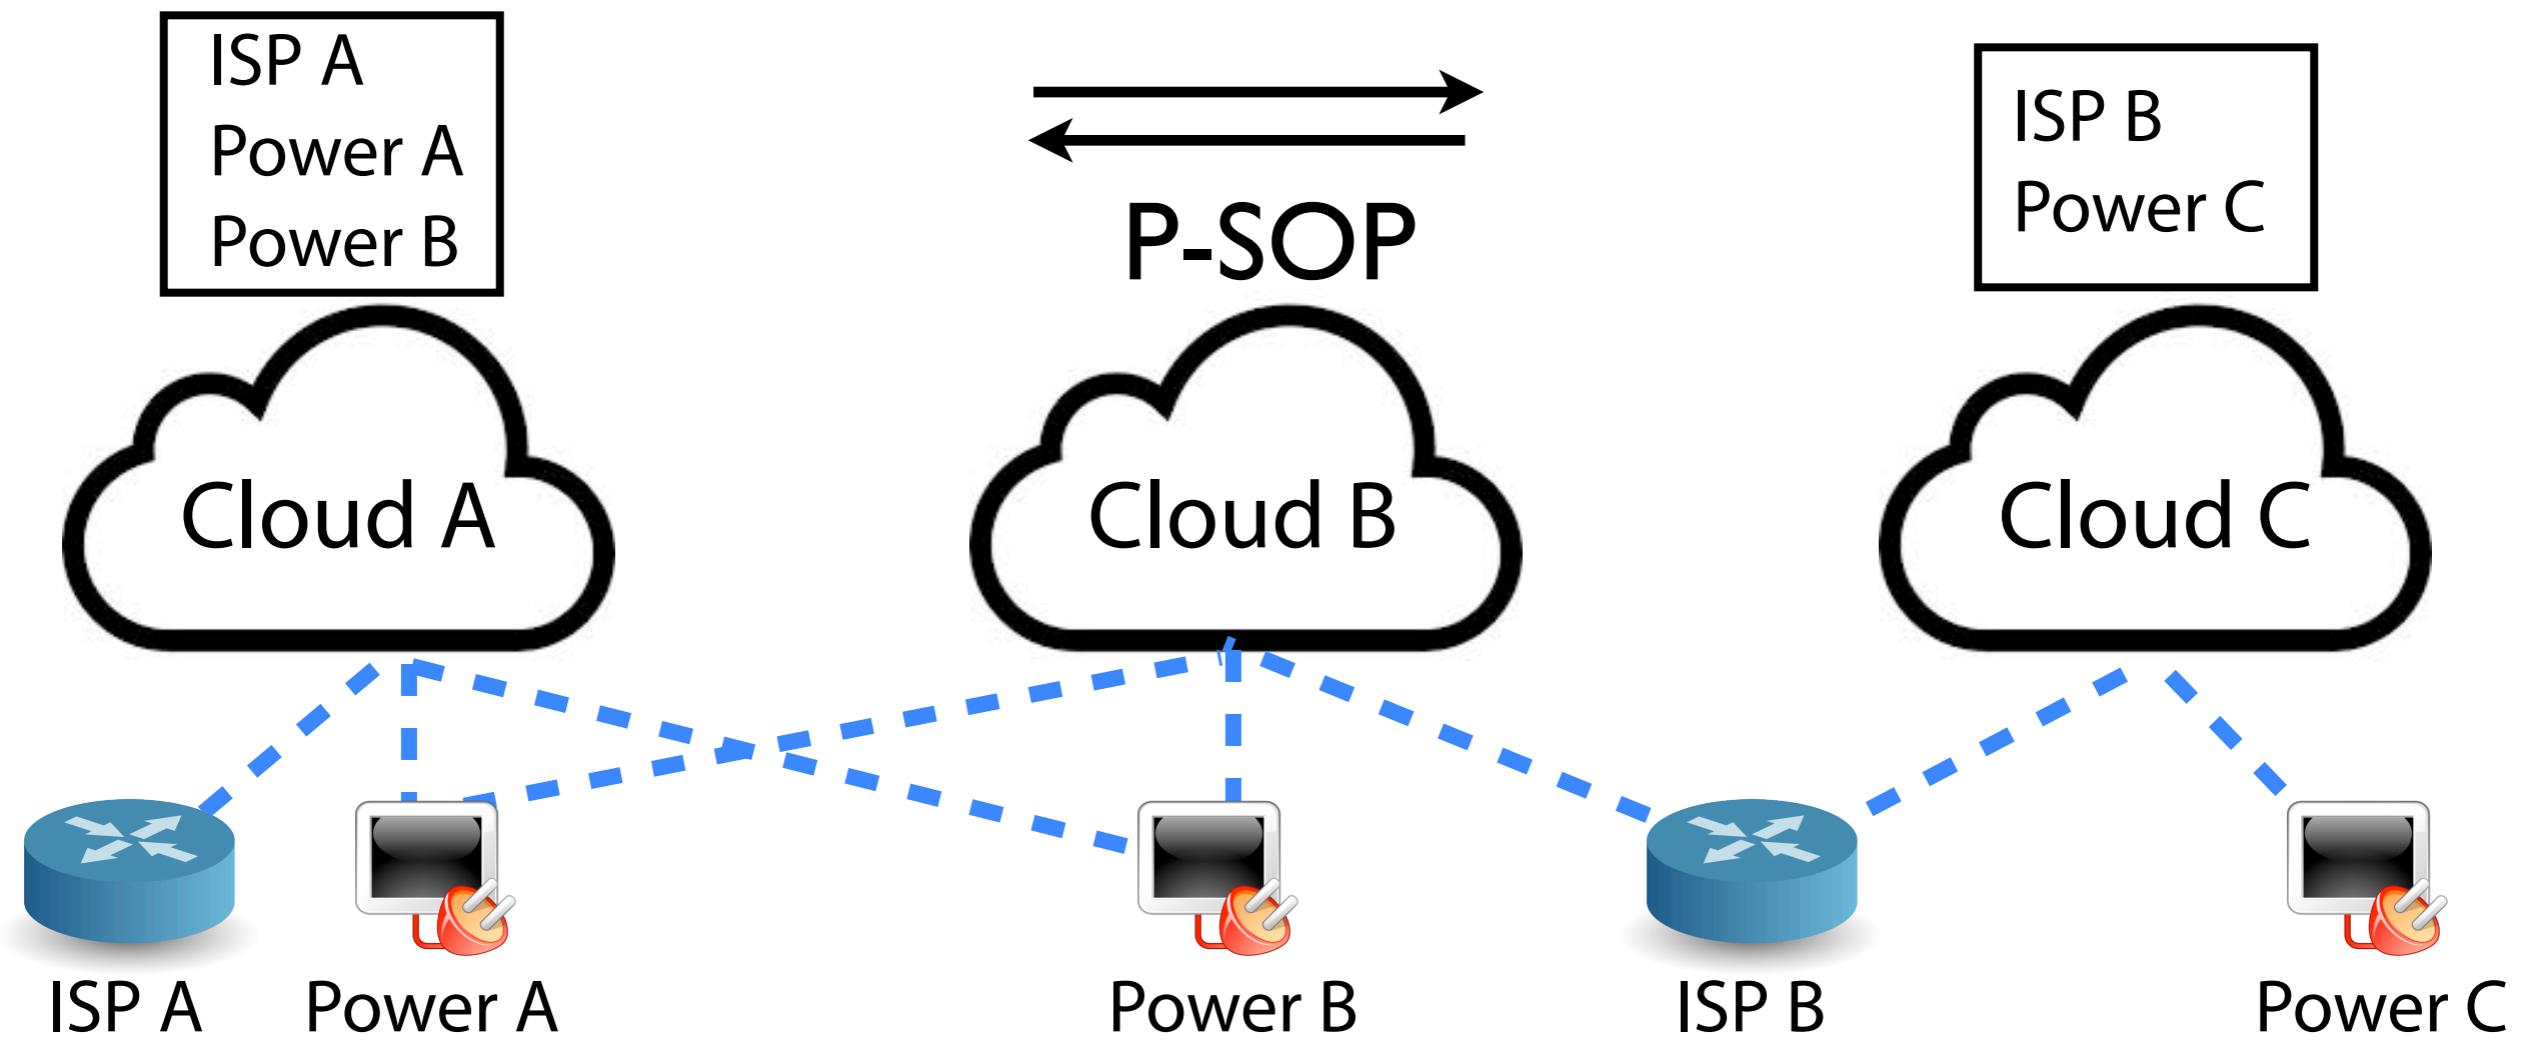


## Service Provider



 Service Provider

 INDaaS Agent



| Deployment | Sim  |
|------------|------|
| Cloud A&B  | 0.5  |
| Cloud B&C  | 0.25 |
|            |      |

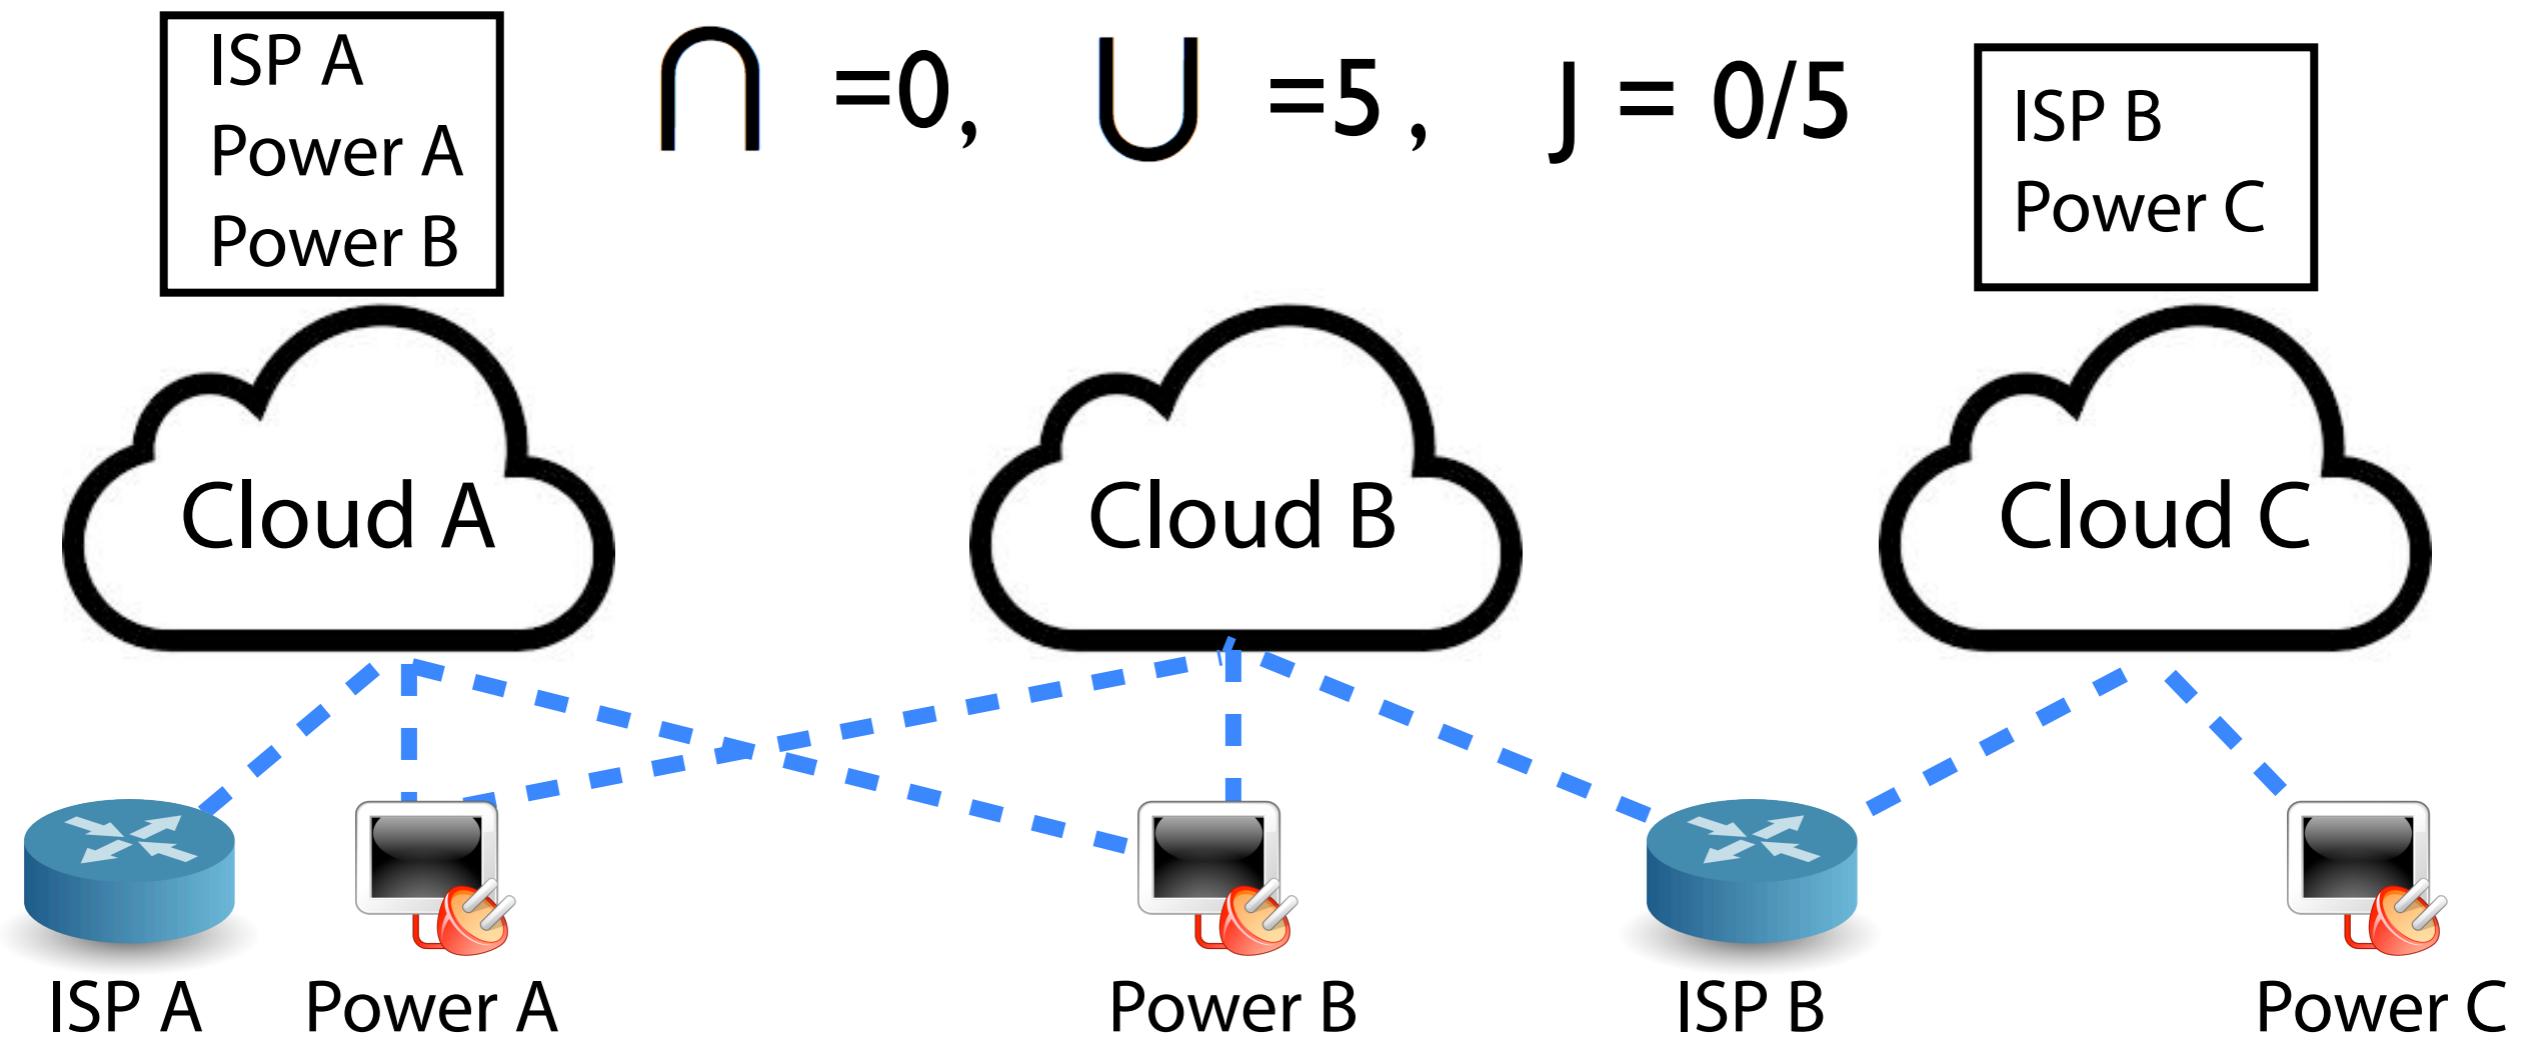


Service Provider



INDaaS Agent

| Deployment | Sim  |
|------------|------|
| Cloud A&B  | 0.5  |
| Cloud B&C  | 0.25 |
|            |      |



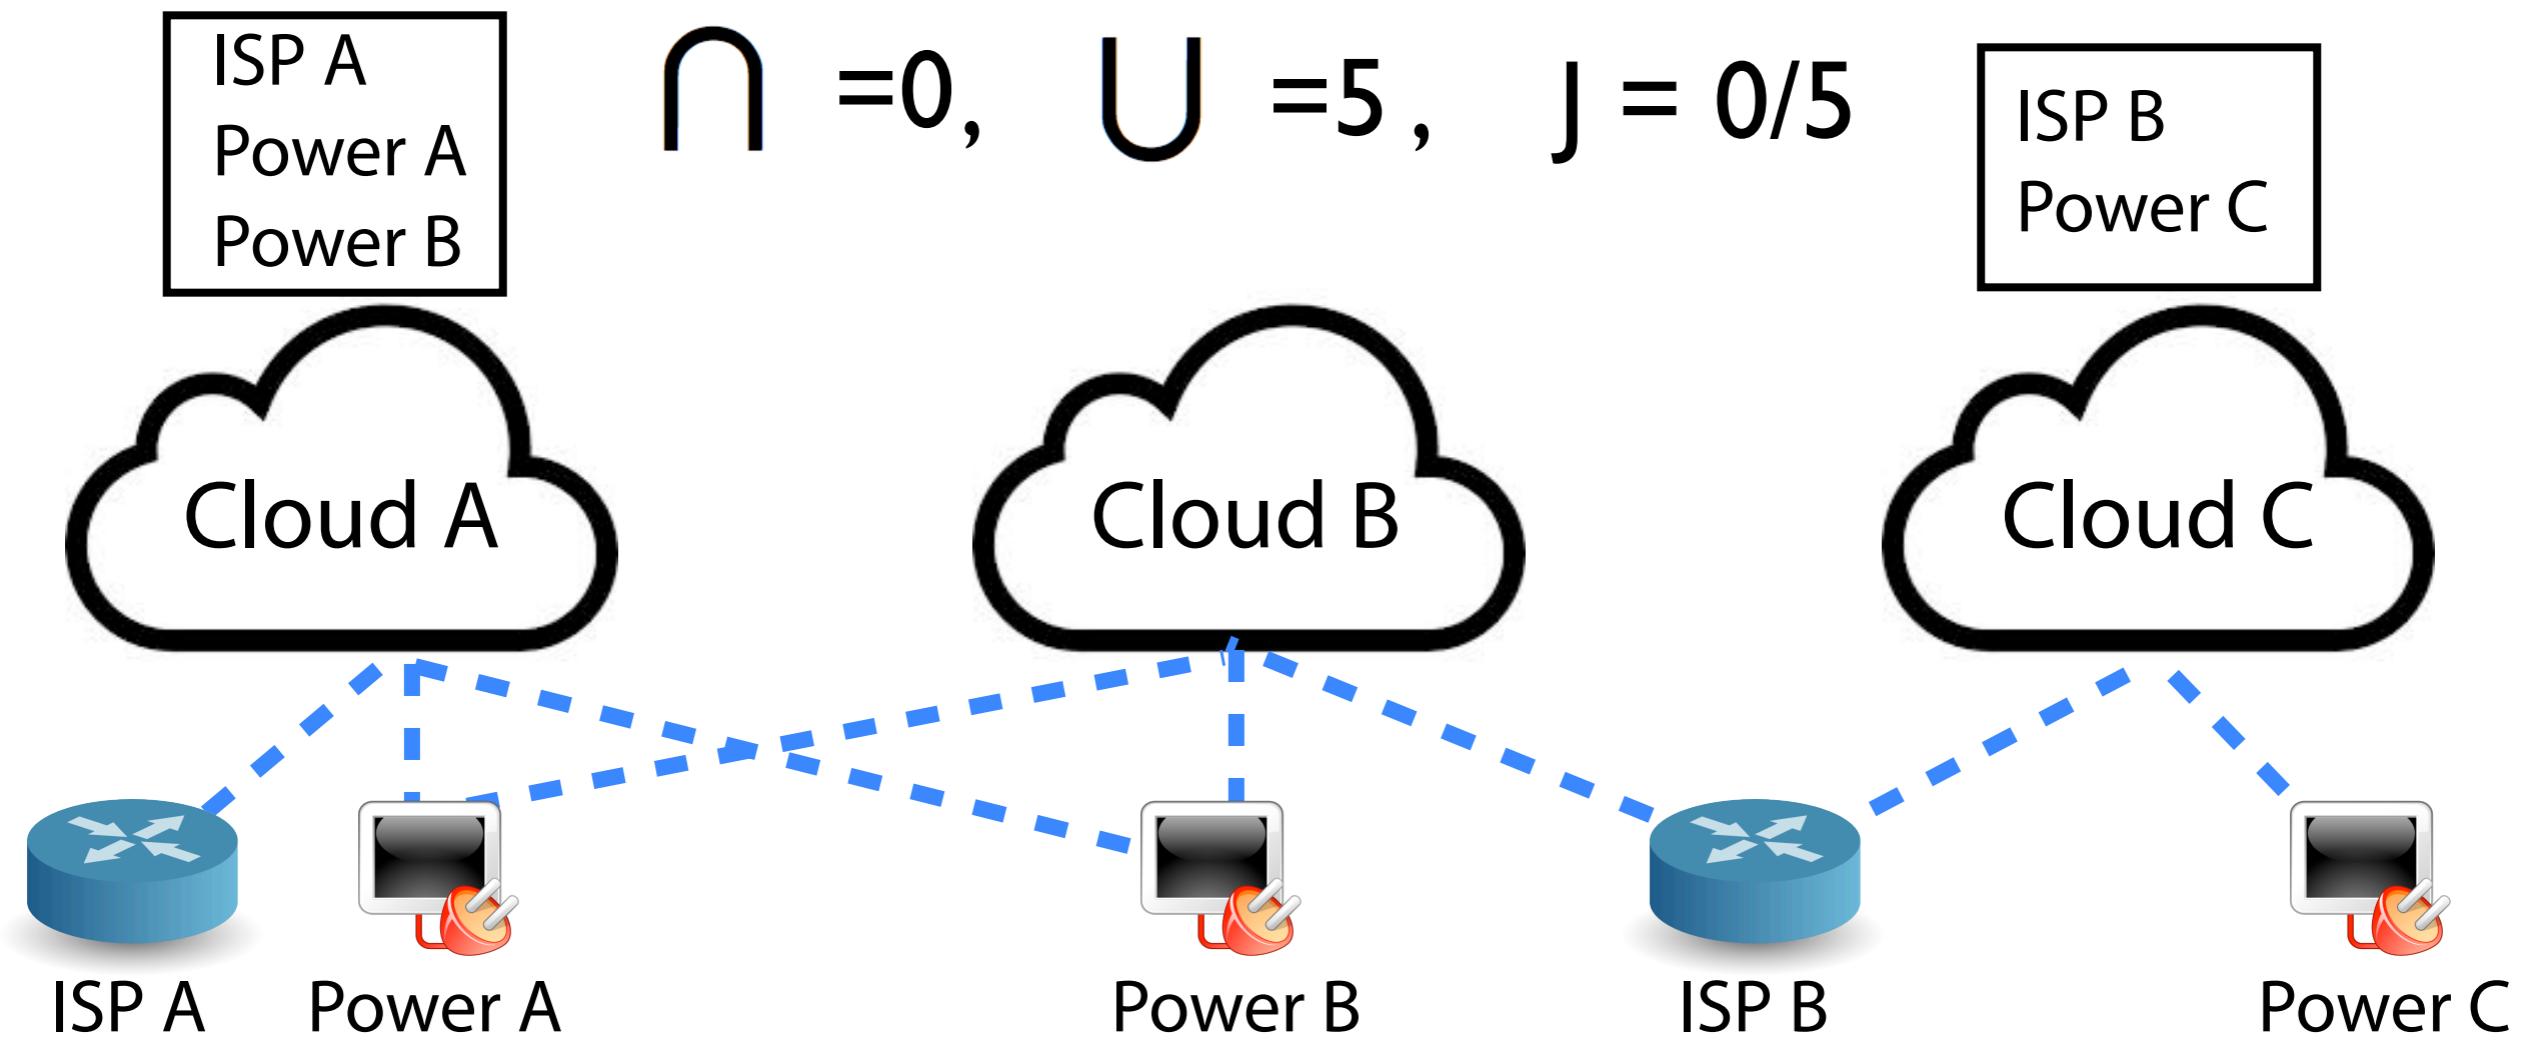


Service Provider



INDaaS Agent

| Deployment | Sim  |
|------------|------|
| Cloud A&B  | 0.5  |
| Cloud B&C  | 0.25 |
| Cloud A&C  | 0    |





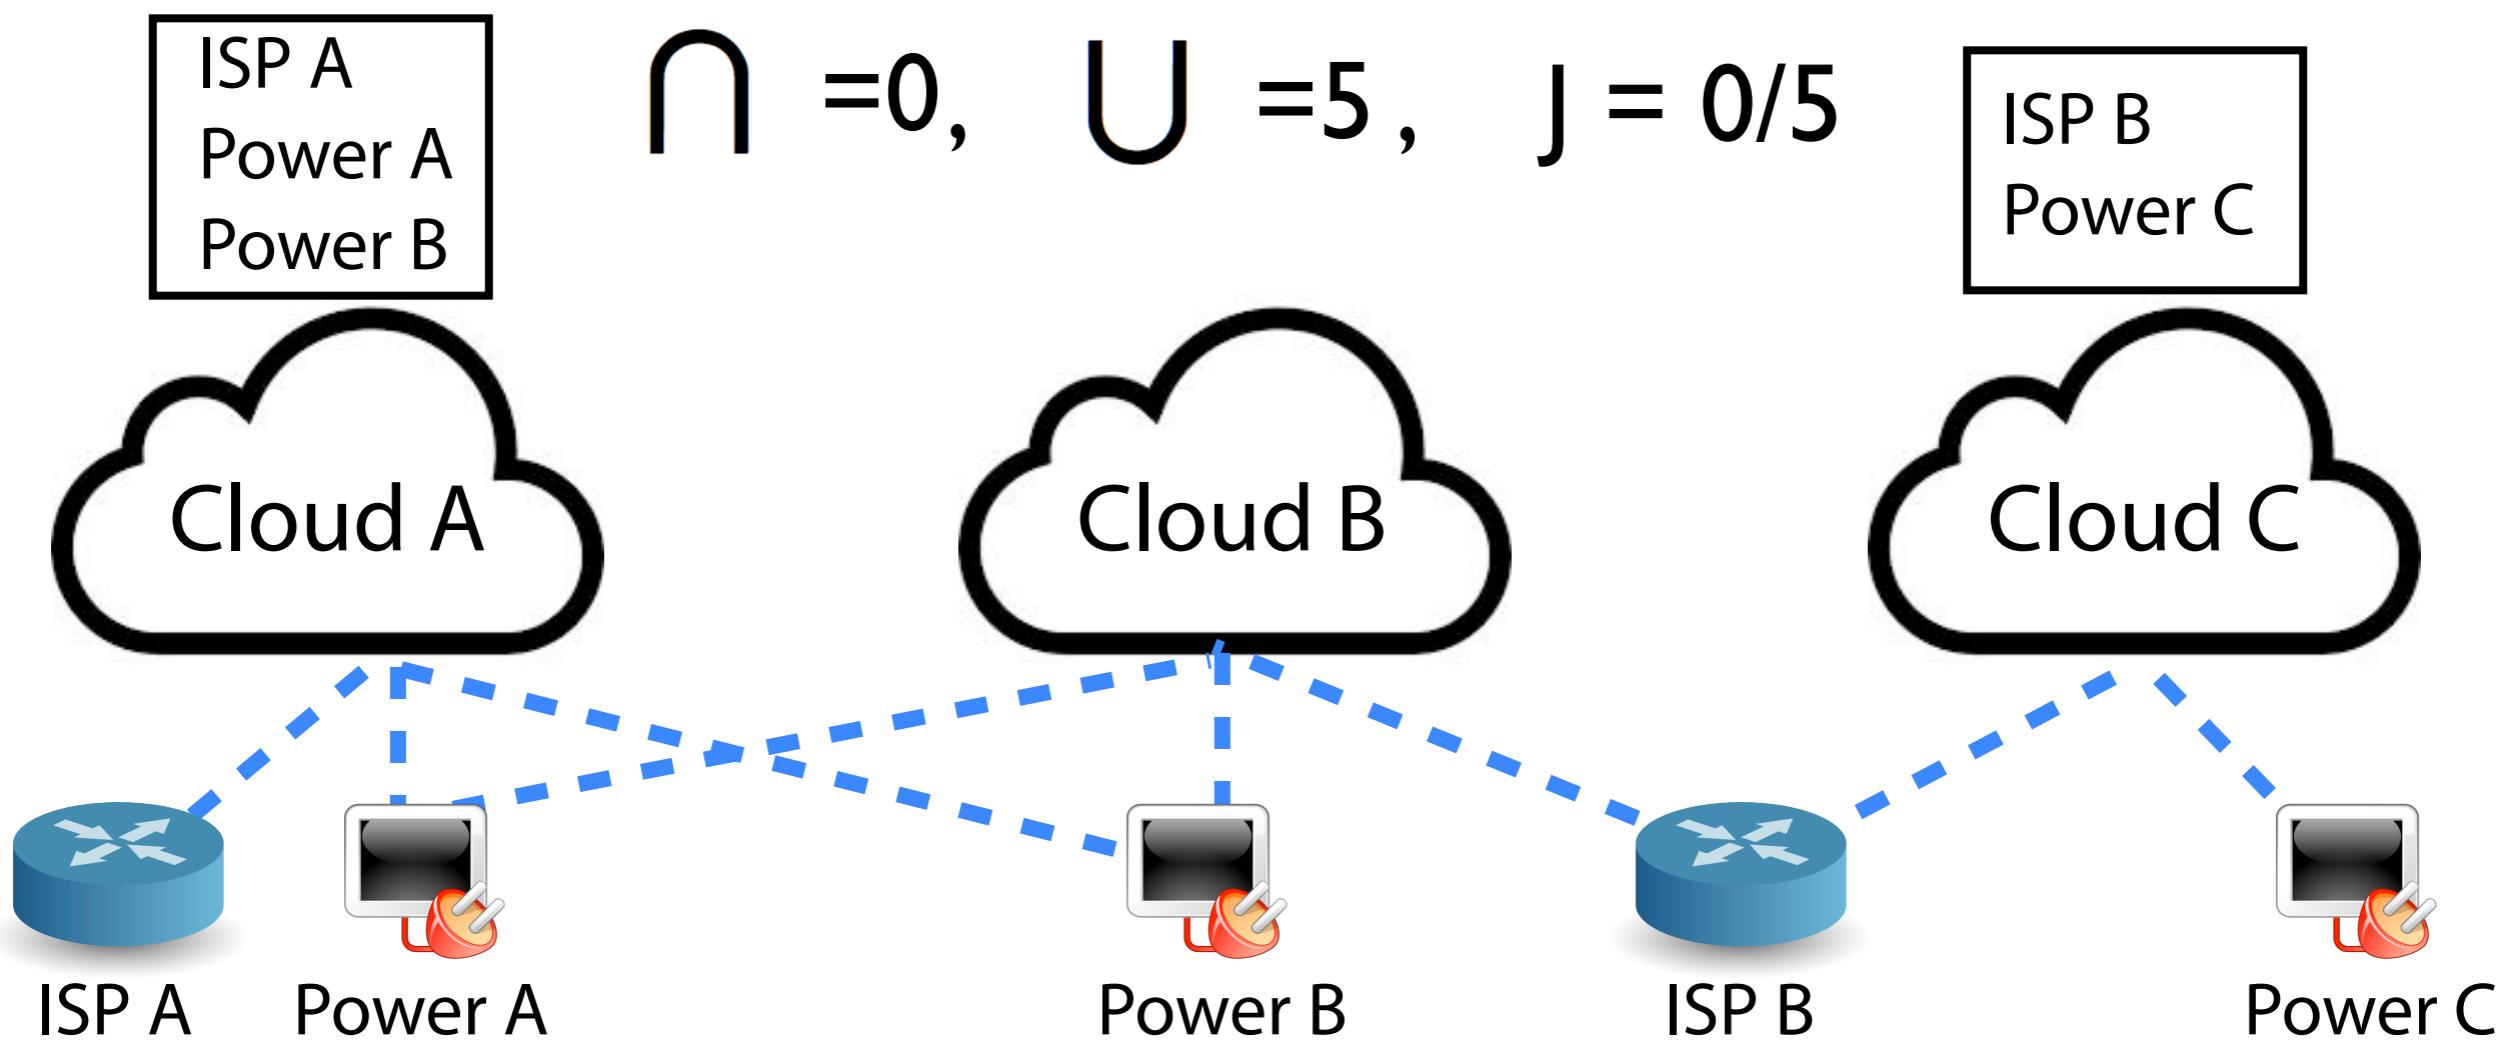
Service Provider



INDaaS Agent

| Deployment | Sim  |
|------------|------|
| Cloud A&B  | 0.5  |
| Cloud B&C  | 0.25 |
| Cloud A&C  | 0    |

$$\cap = 0, \quad U = 5, \quad J = 0/5$$



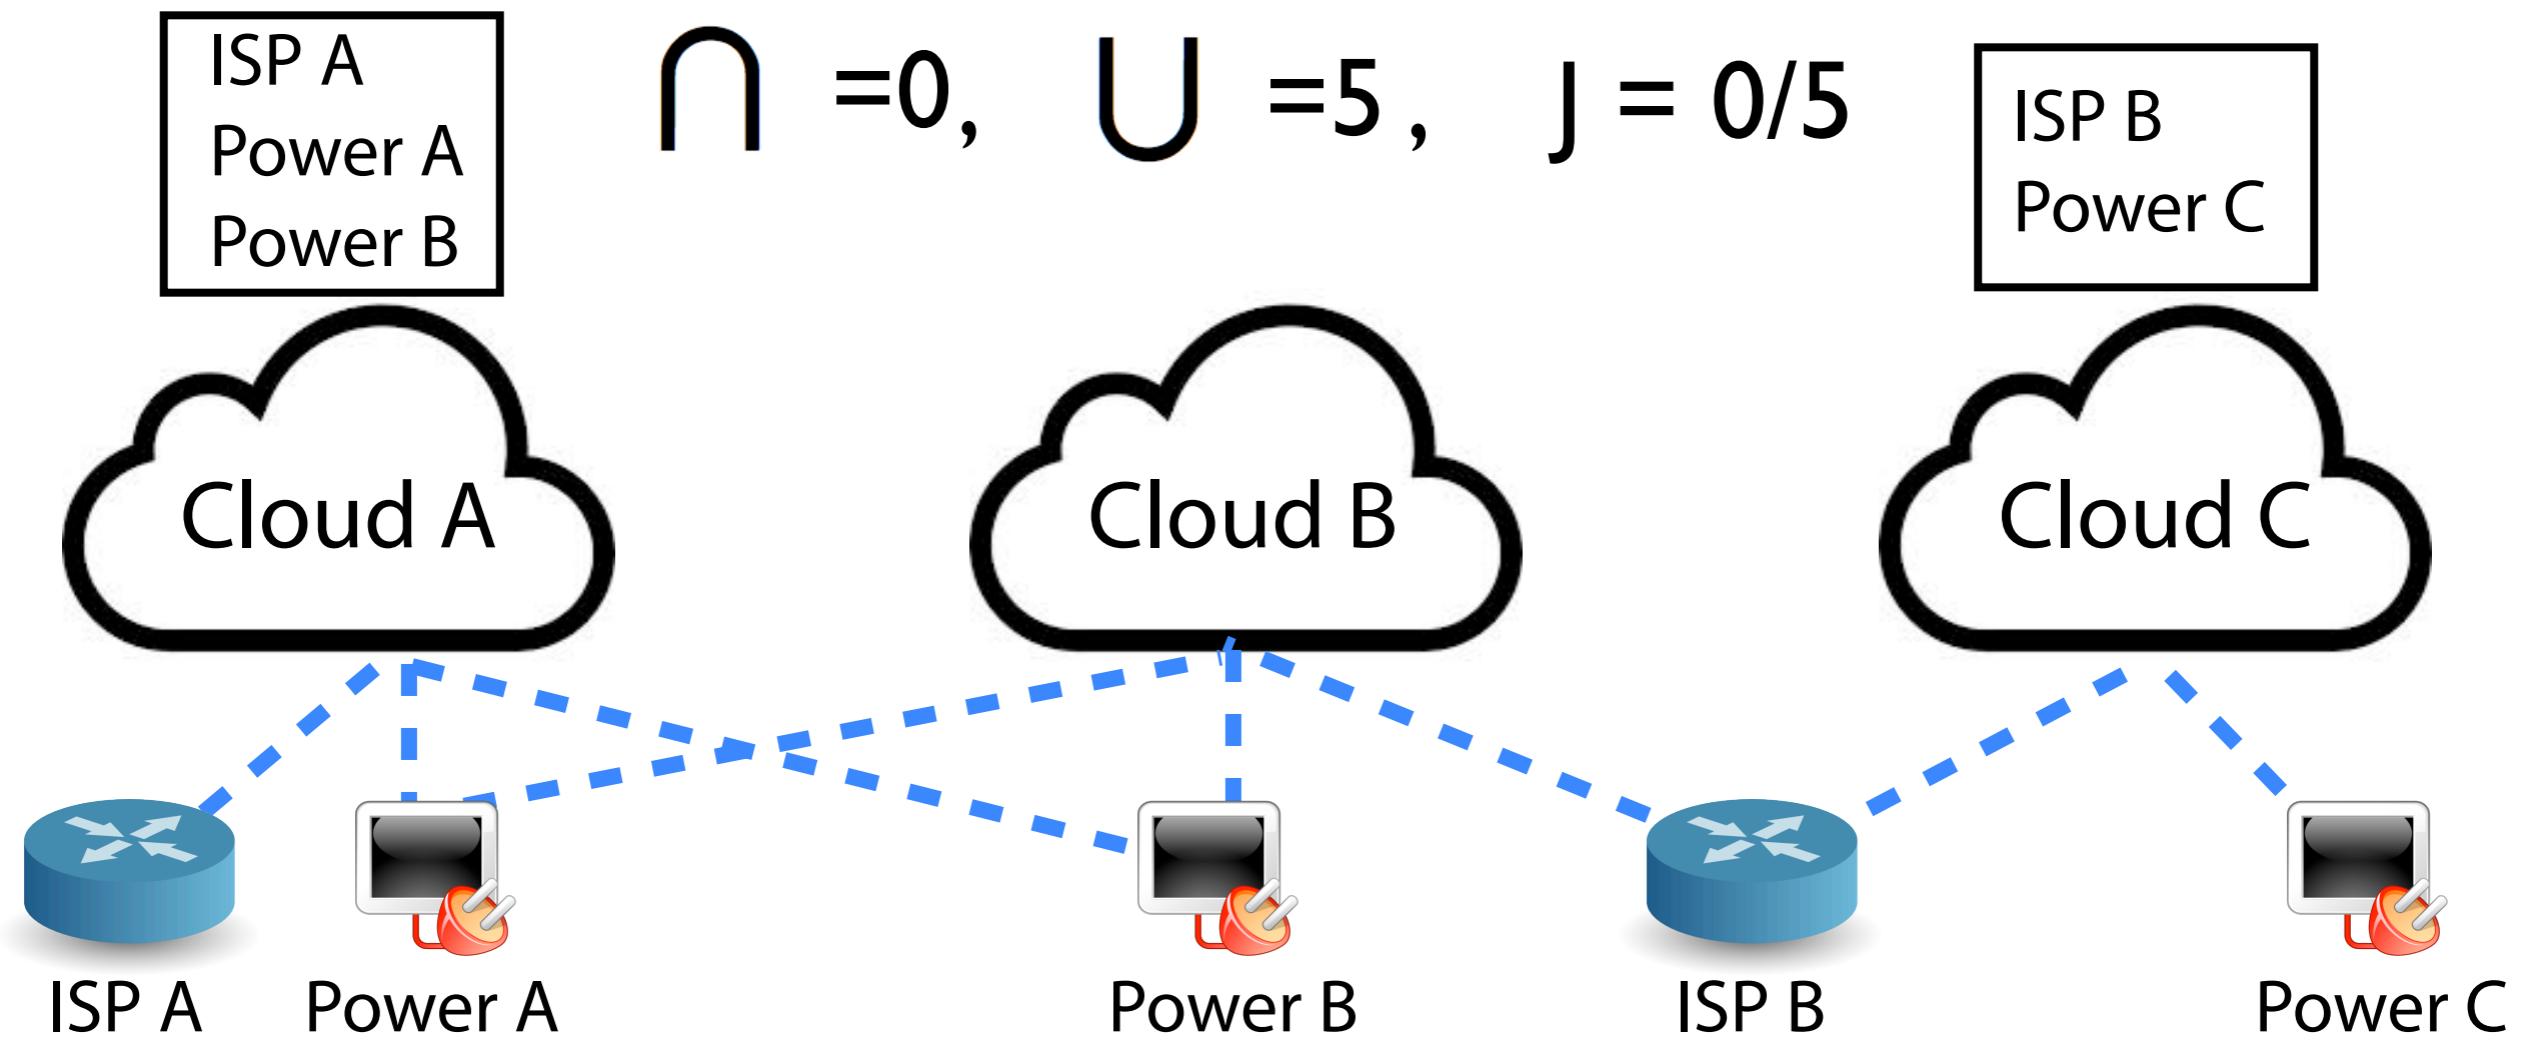


Service Provider



INDaaS Agent

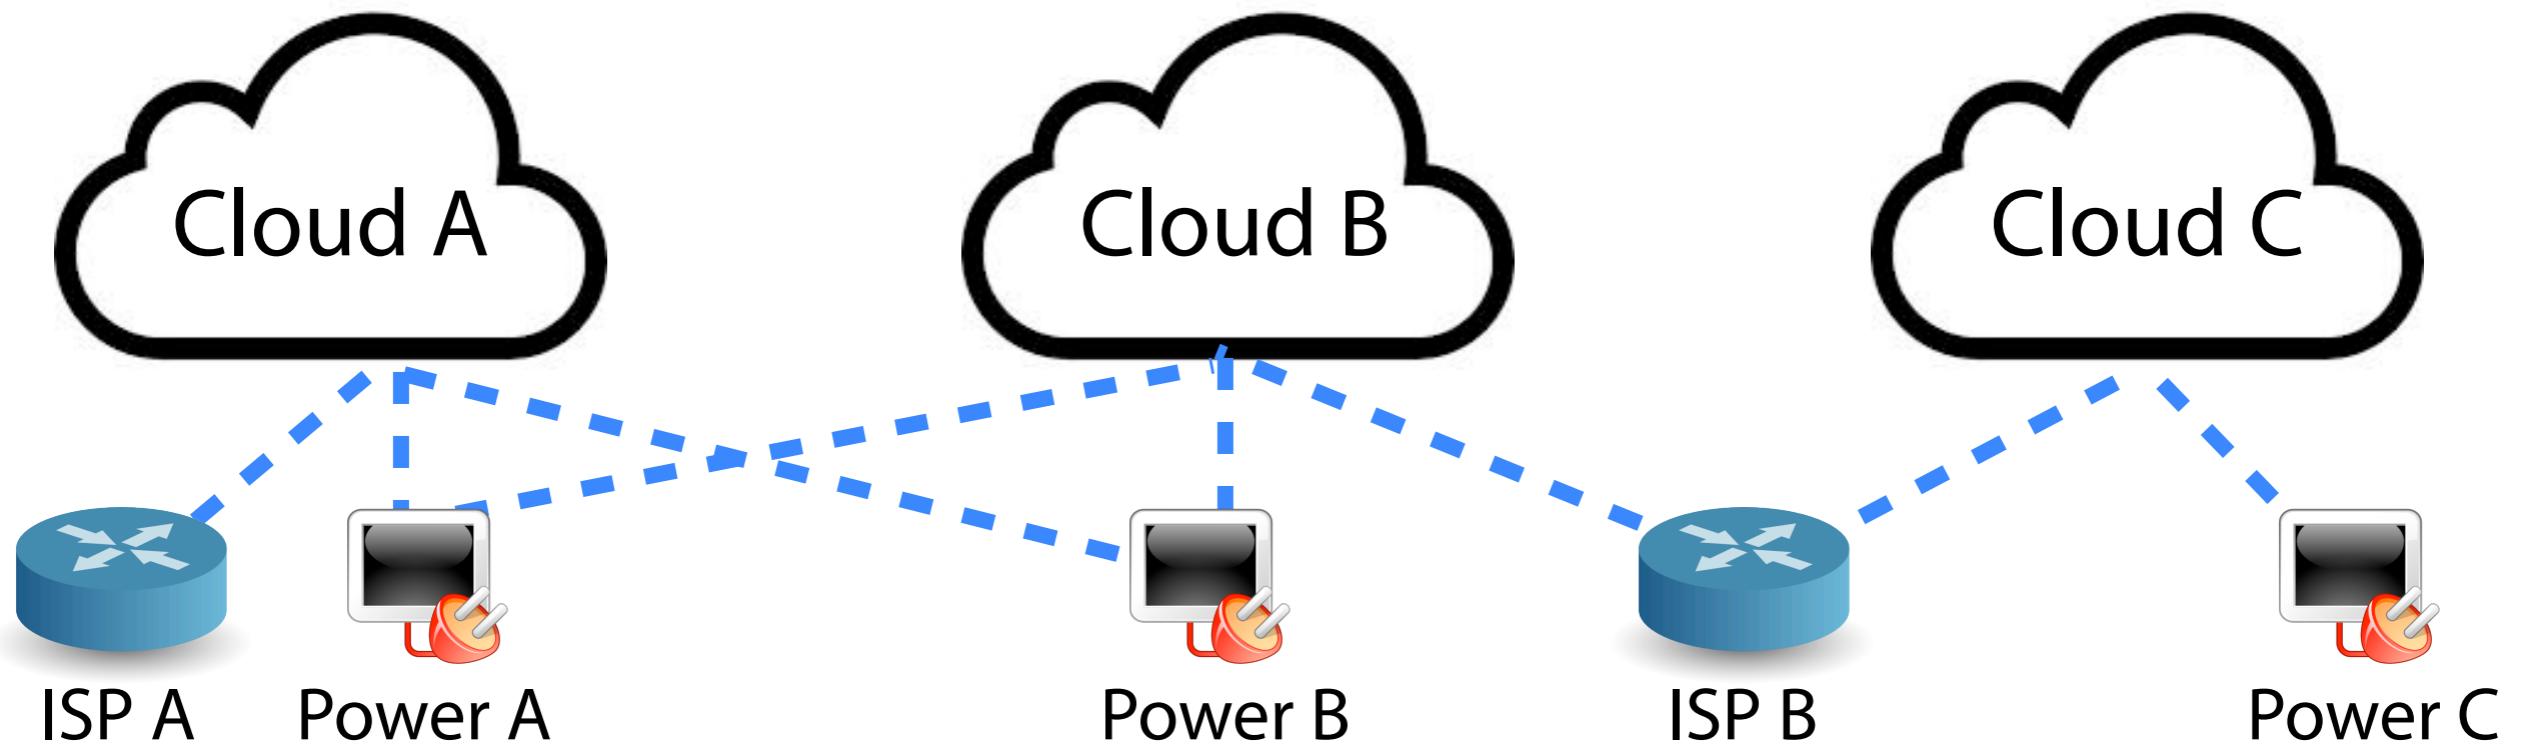
| Deployment | Sim  |
|------------|------|
| Cloud A&C  | 0    |
| Cloud B&C  | 0.25 |
| Cloud A&B  | 0.5  |

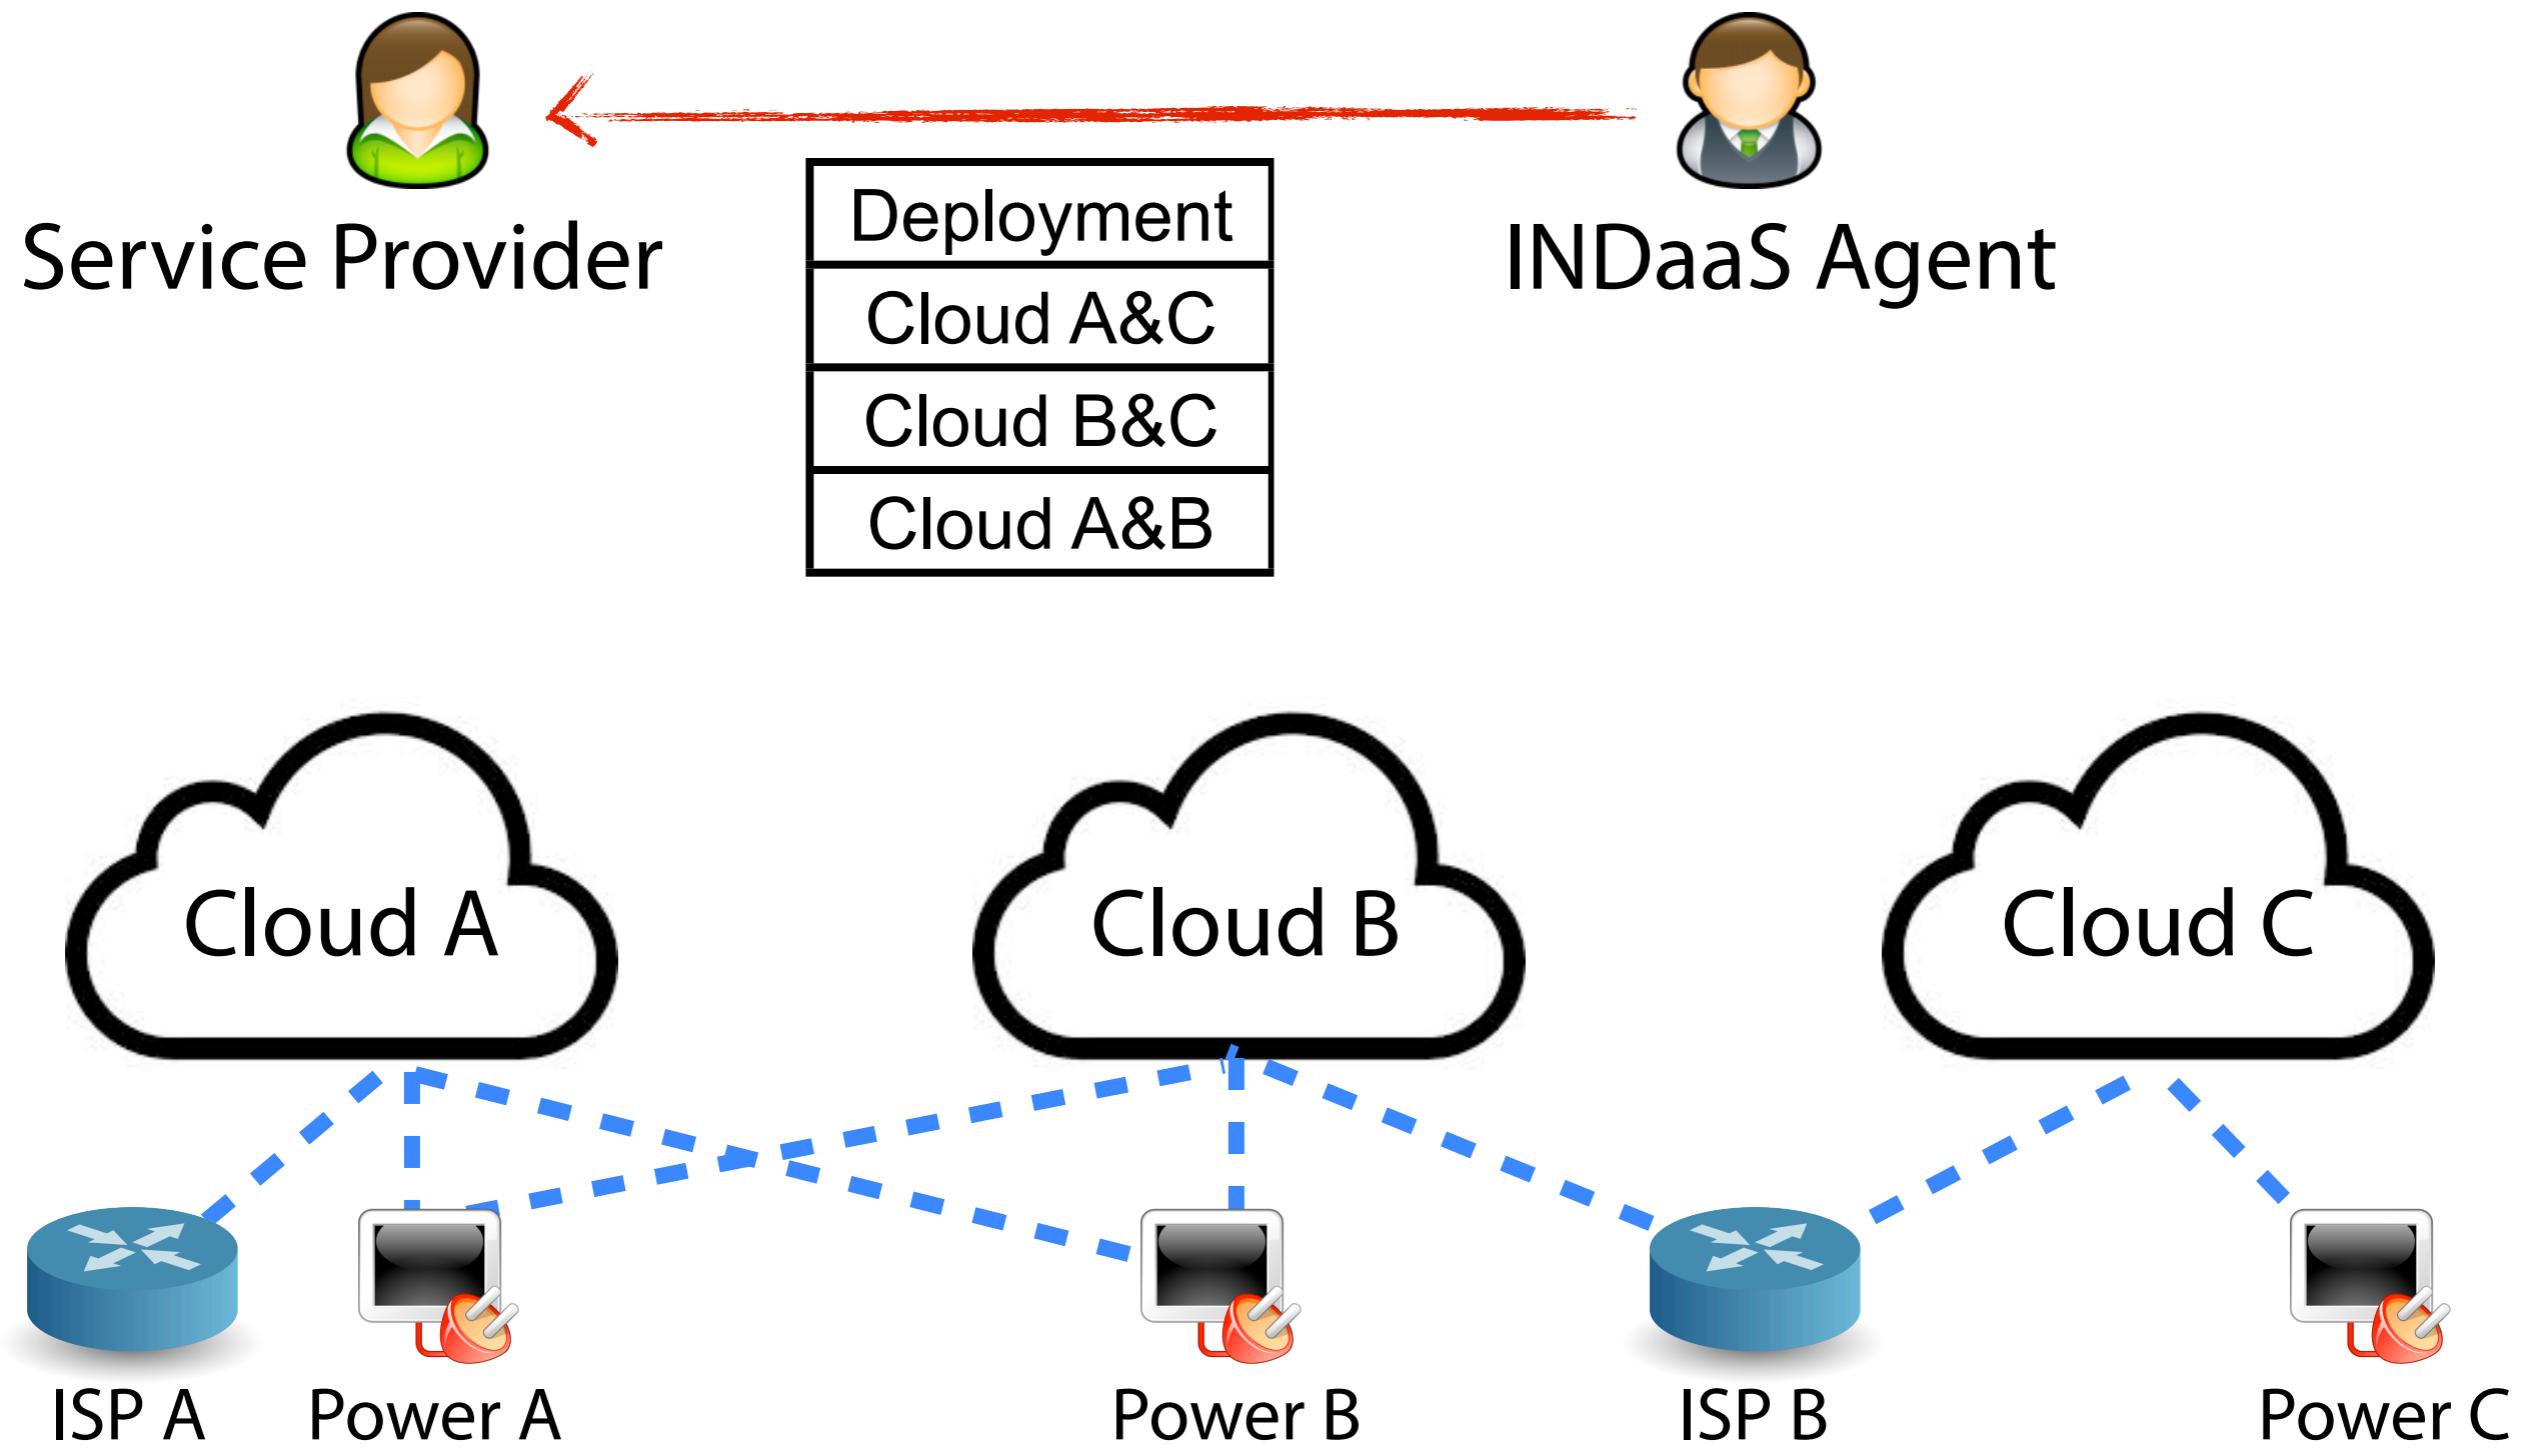


| Deployment | Sim  |
|------------|------|
| Cloud A&C  | 0    |
| Cloud B&C  | 0.25 |
| Cloud A&B  | 0.5  |

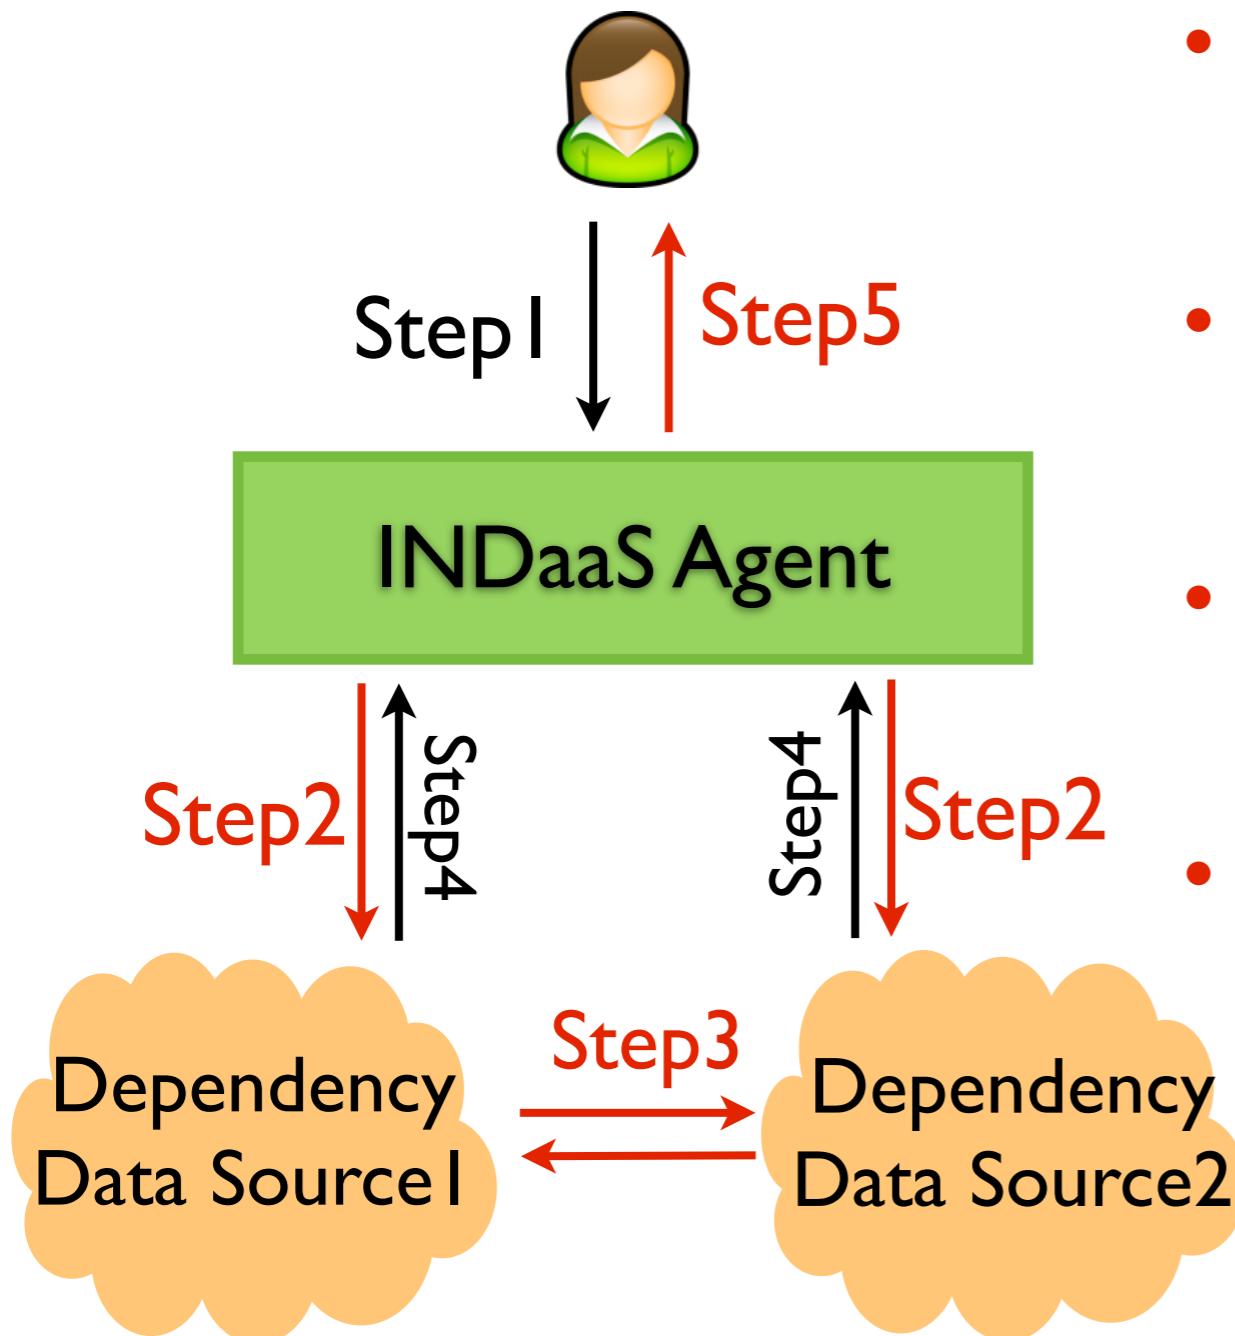
 Service Provider

 INDaaS Agent





# RoadMap



- #1: Dependency collections
  - Solution: Reusing existing tools
- #2: Dependency representation
  - Solution: Fault graphs
- #3: Efficient auditing
  - Solution: Failure sampling algorithm
- #4: Private independence audit
  - Solution: Private Jaccard similarity

# Evaluation

# Evaluation

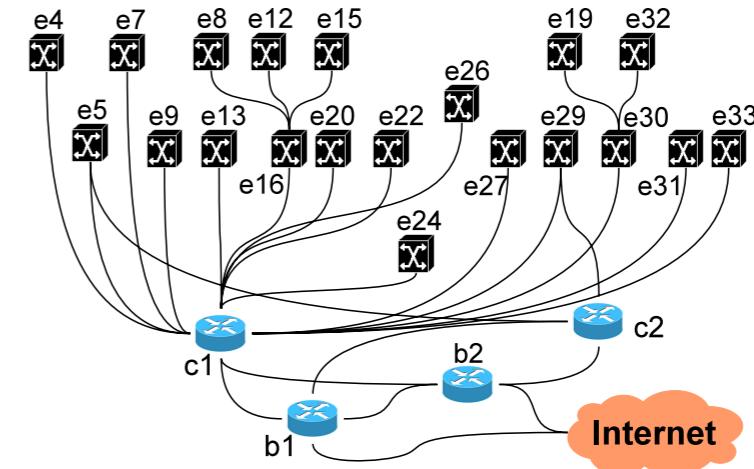
- Three realistic case studies.
- Tradeoff between auditing algorithms
- Overhead of P-SOP protocol

# Evaluation

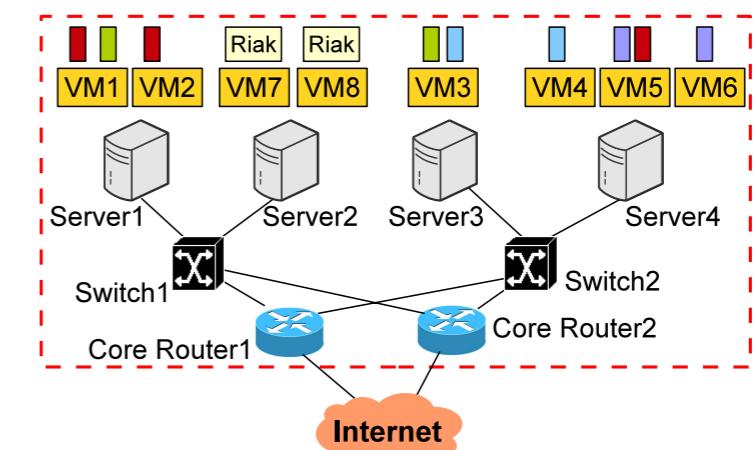
- Three realistic case studies.
- Tradeoff between auditing algorithms
- Overhead of P-SOP protocol

# Three Case Studies

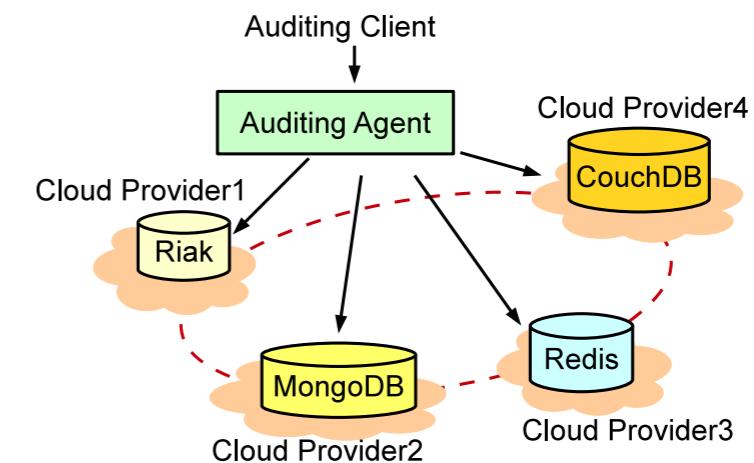
- Common network dependency



- Common hardware dependency



- Common software dependency



Please see our paper for more details

# Evaluation

- Three realistic case studies.
- Tradeoff between auditing algorithms
- Overhead of P-SOP protocol

# Tradeoff Evaluation

- We evaluate efficiency/accuracy tradeoff.
- We generate topology based on fat tree model.

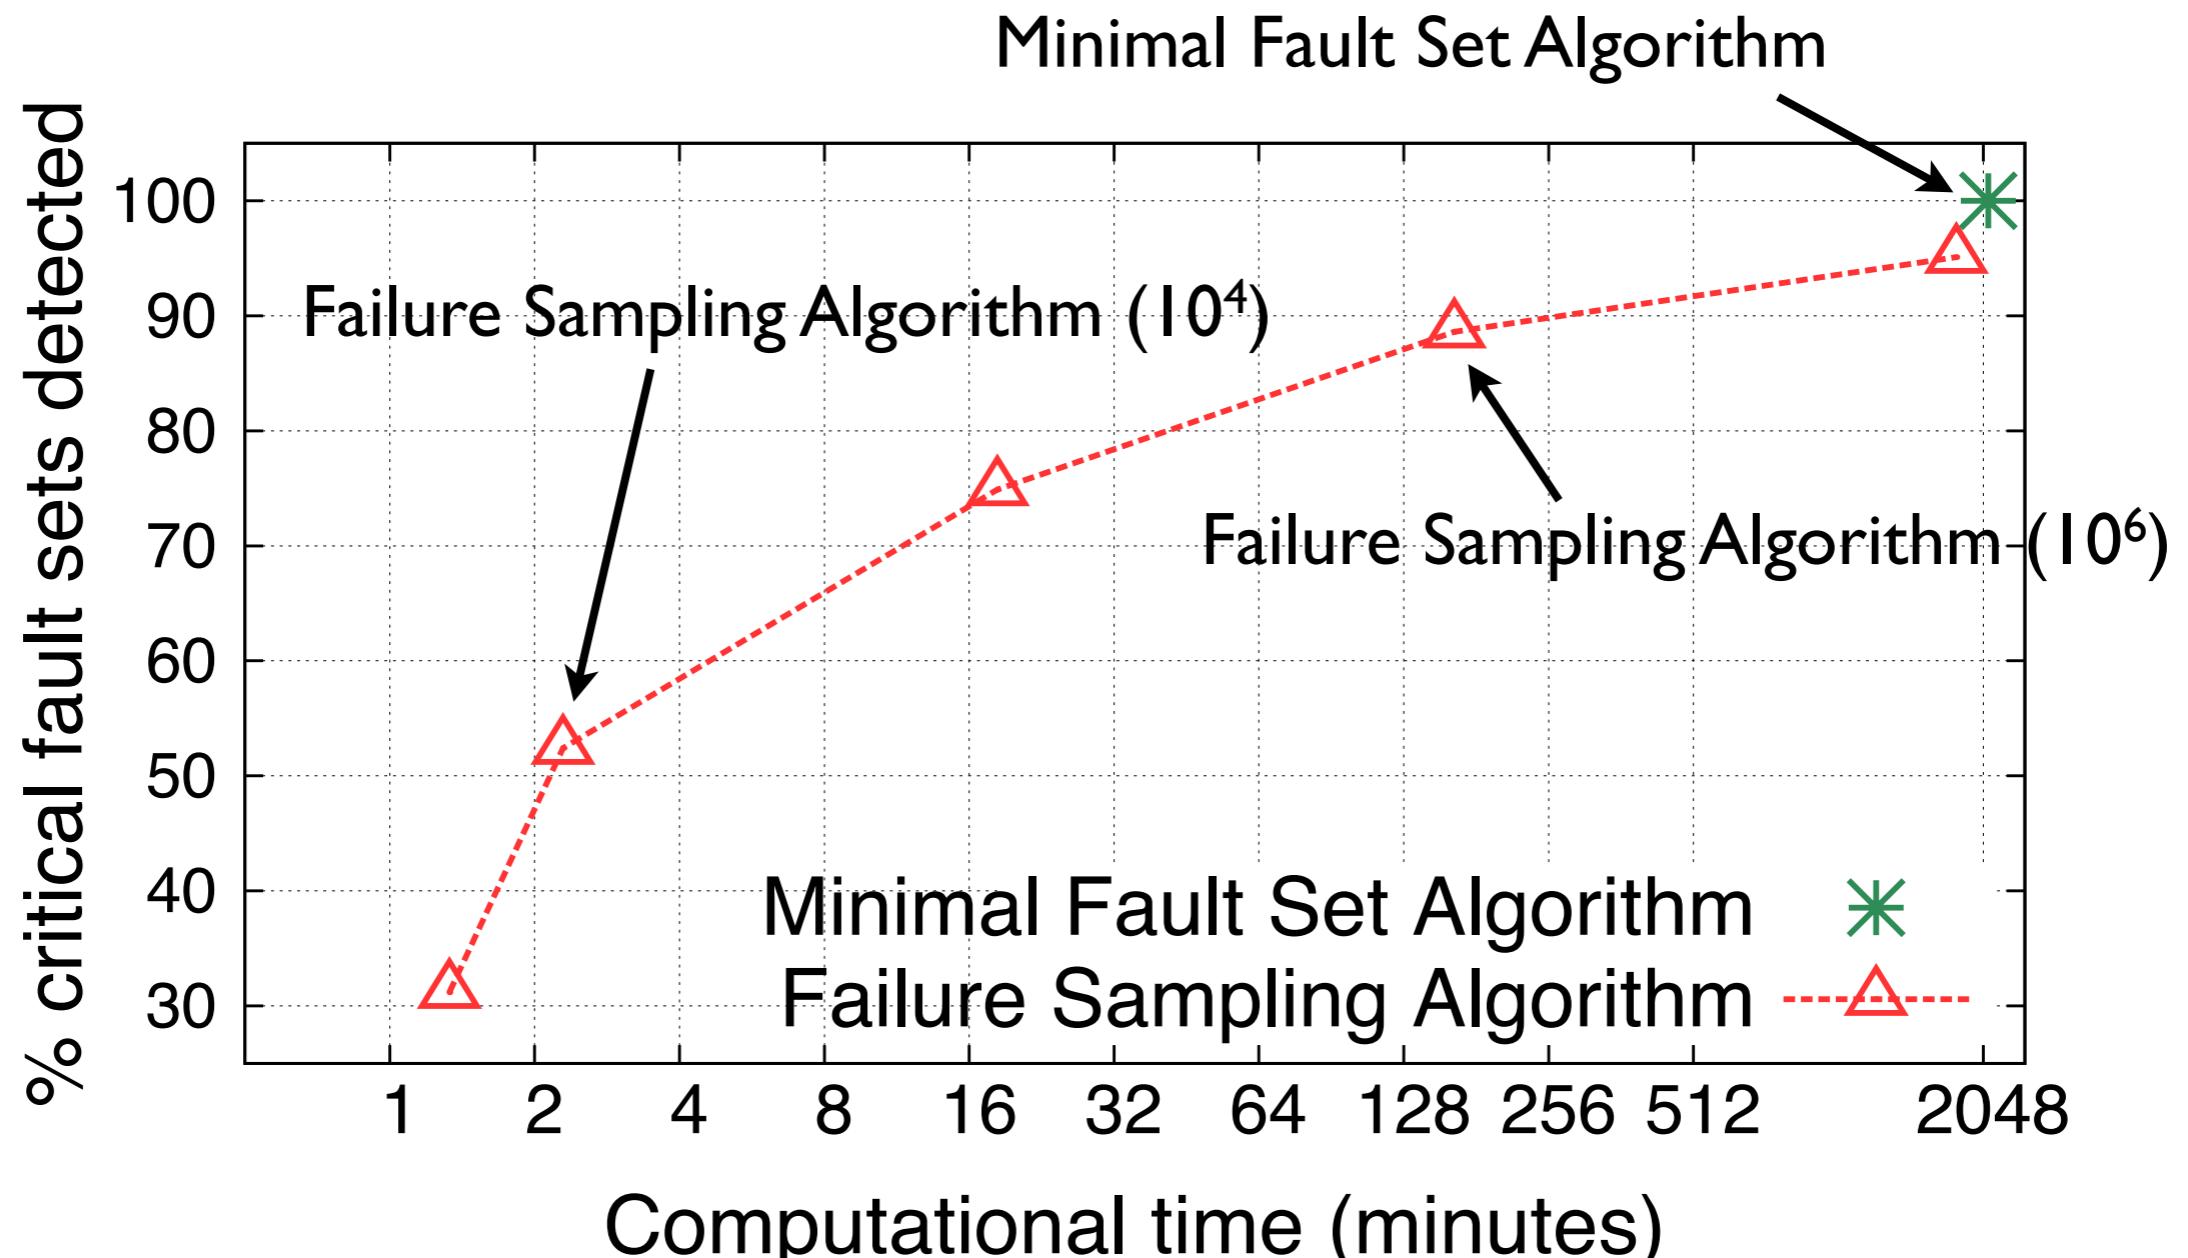
# Tradeoff Evaluation

|                    | Topology A | Topology B | Topology C |
|--------------------|------------|------------|------------|
| # of Core Routers  | 64         | 144        | 576        |
| # of Agg Switches  | 128        | 288        | 1,152      |
| # of ToR Switches  | 128        | 288        | 1,152      |
| # of Servers       | 1,024      | 3,456      | 27,648     |
| Total # of devices | 1,344      | 4,176      | 30,528     |

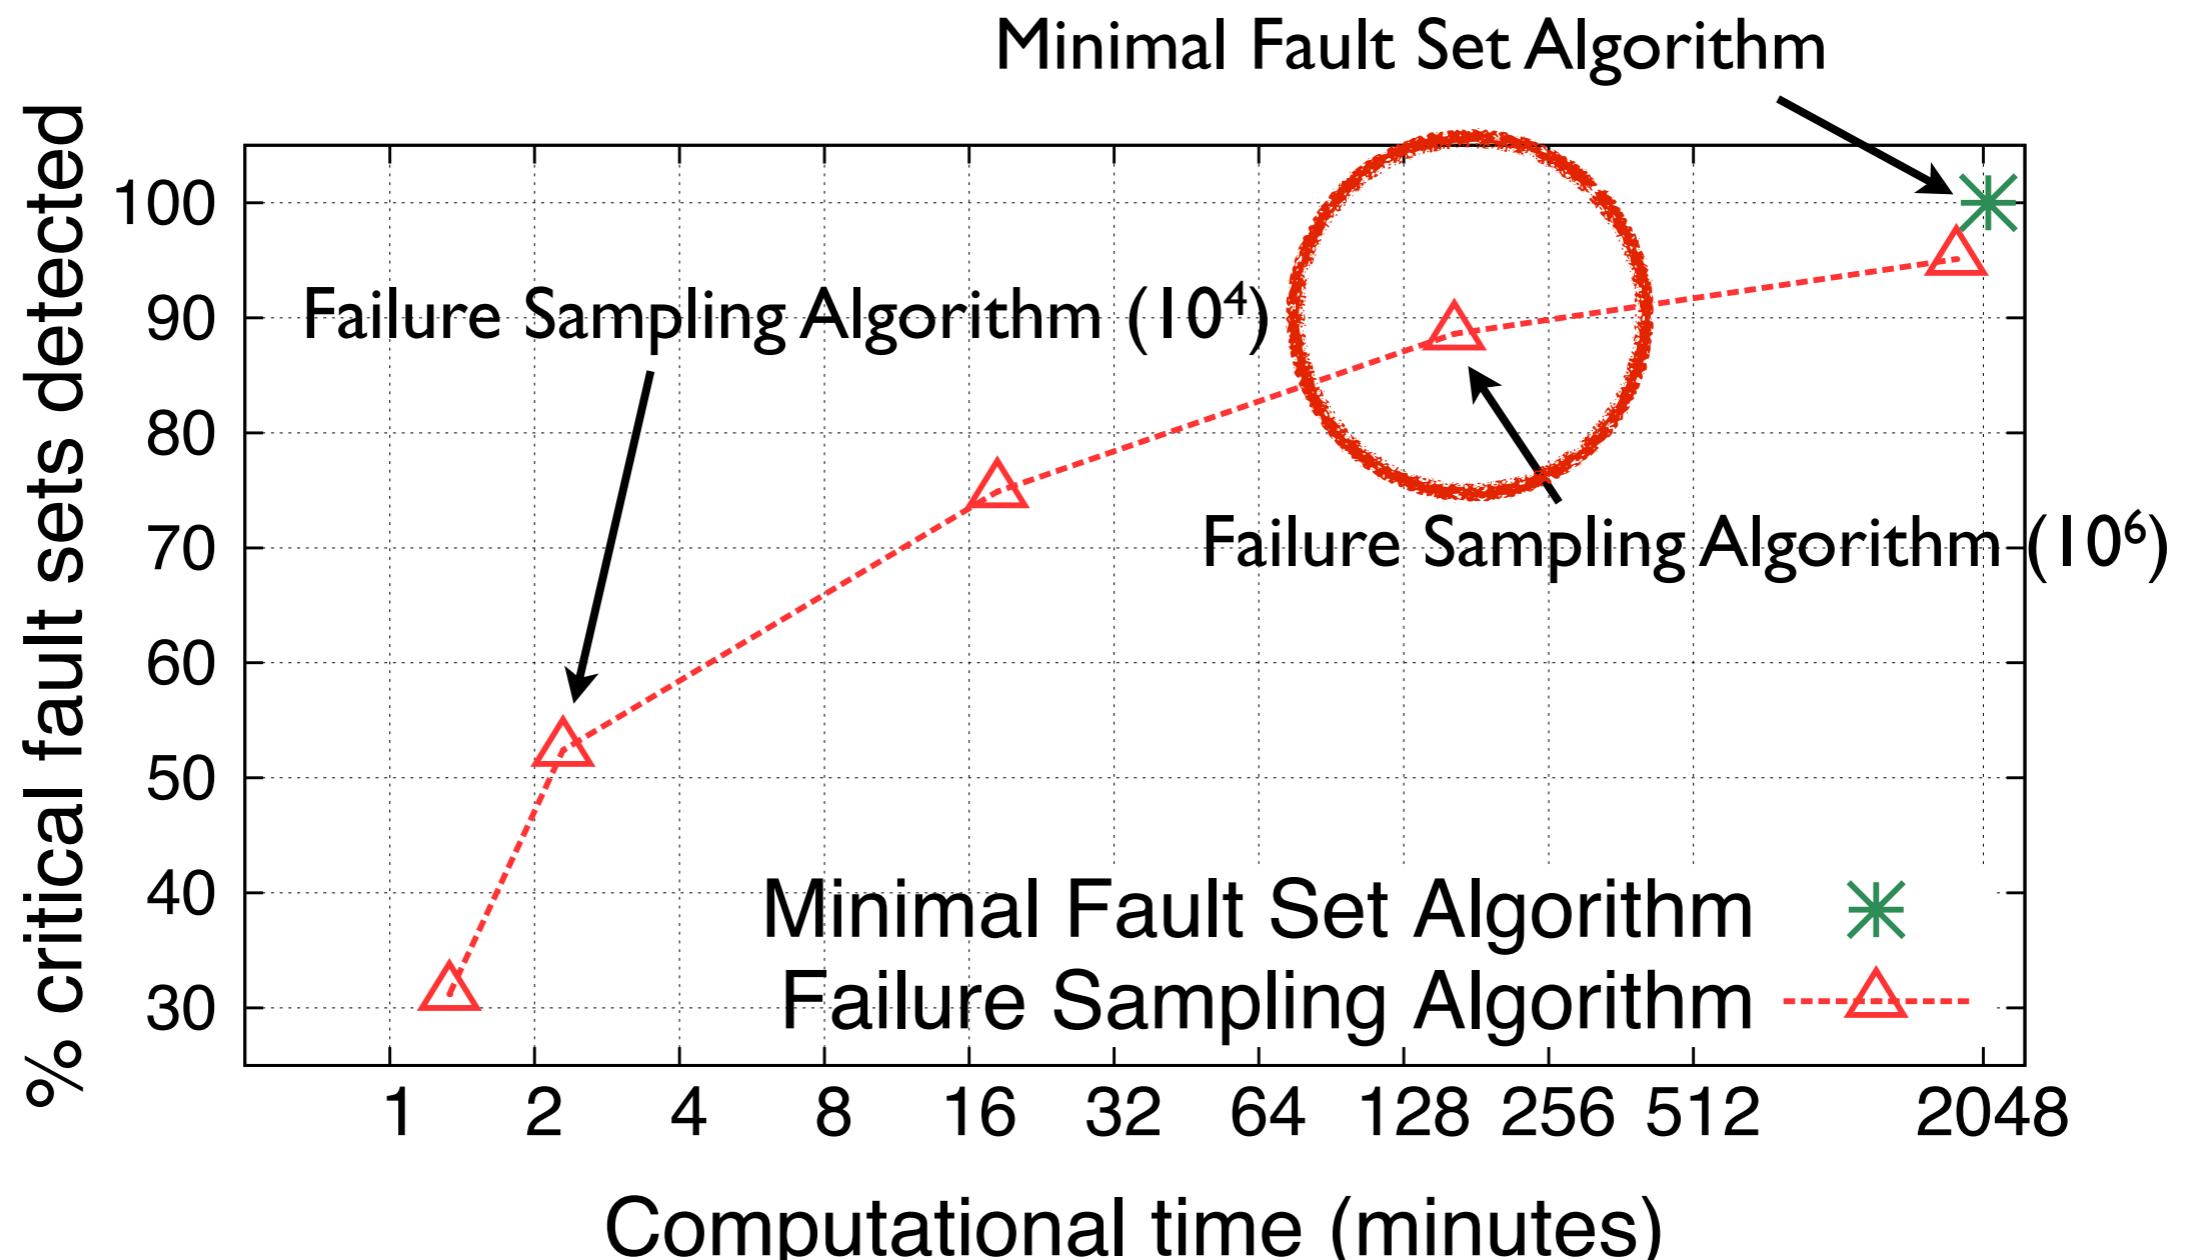
# Tradeoff Evaluation

|                    | Topology A | Topology B | Topology C |
|--------------------|------------|------------|------------|
| # of Core Routers  | 64         | 144        | 576        |
| # of Agg Switches  | 128        | 288        | 1,152      |
| # of ToR Switches  | 128        | 288        | 1,152      |
| # of Servers       | 1,024      | 3,456      | 27,648     |
| Total # of devices | 1,344      | 4,176      | 30,528     |

# Topology C: 30,528 Devices



# Topology C: 30,528 Devices



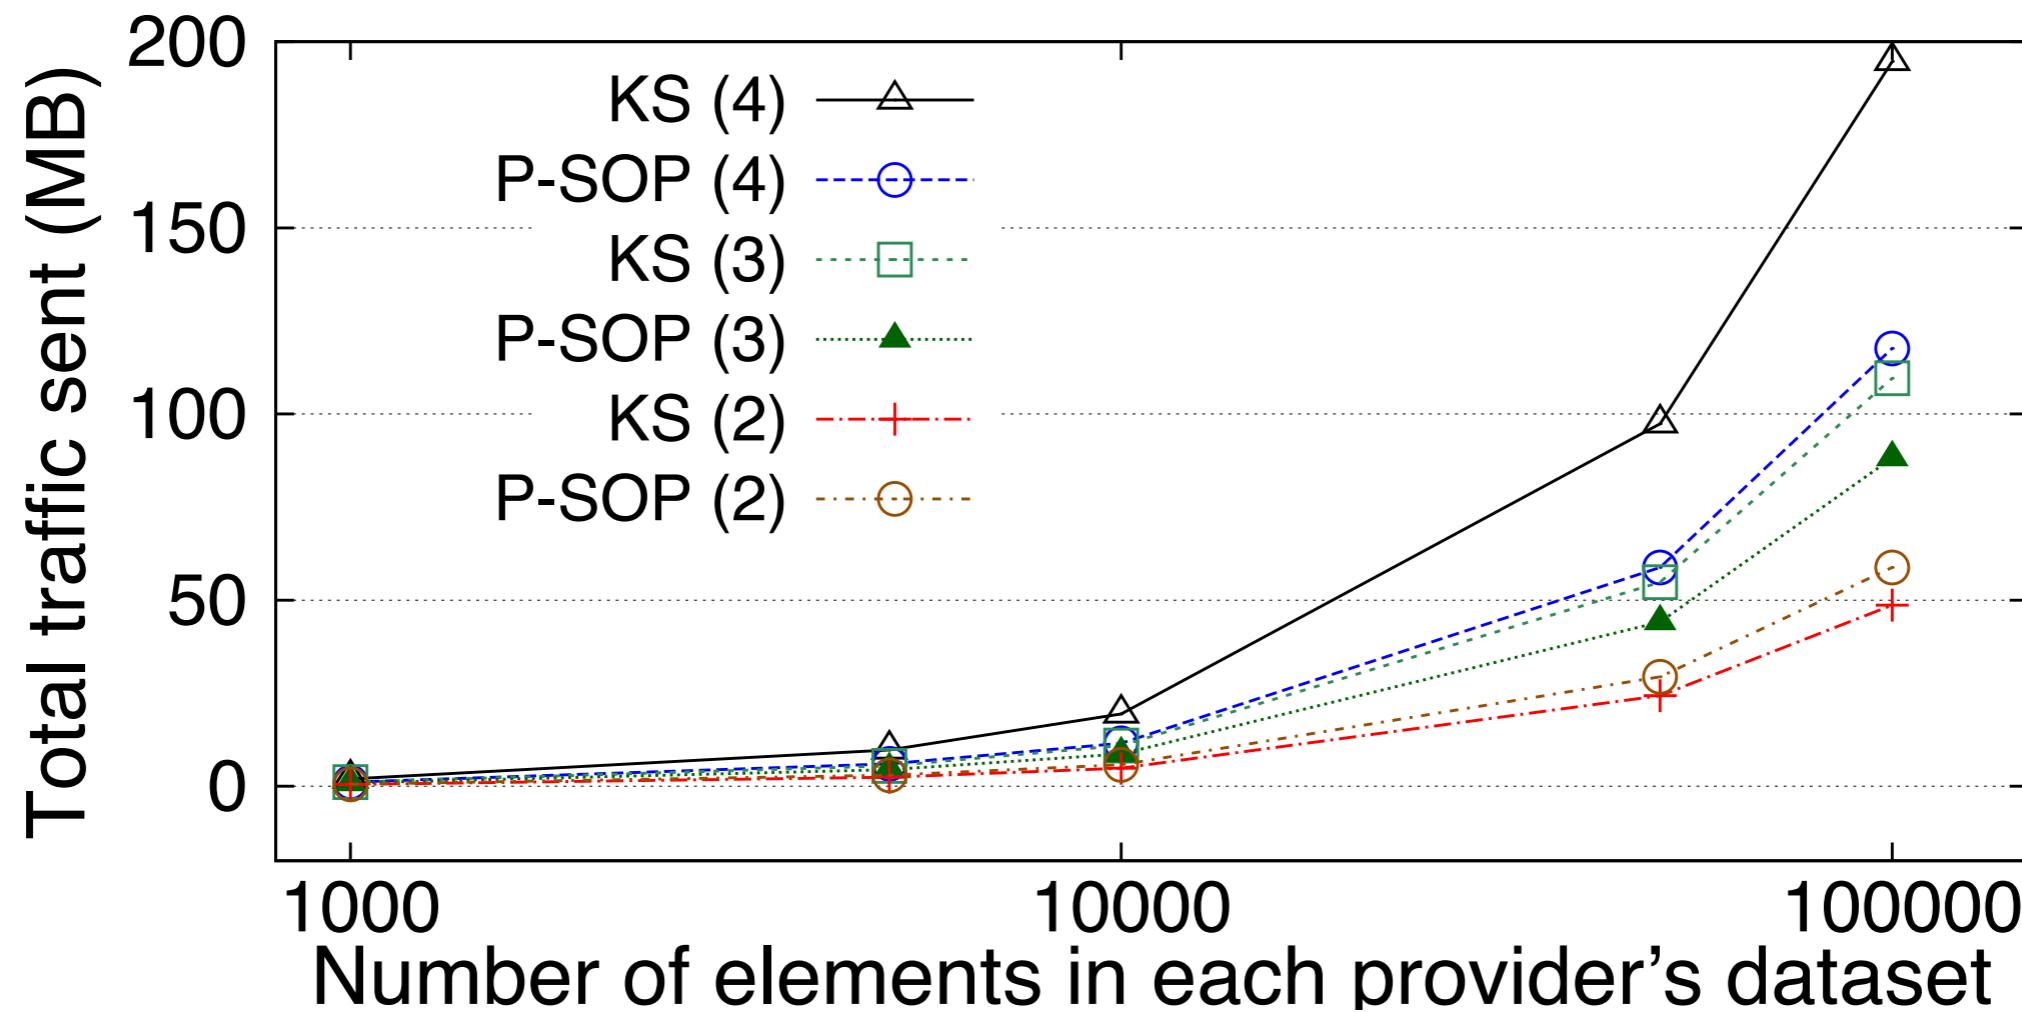
# Evaluation

- Three realistic case studies.
- Tradeoff between auditing algorithms
- Overhead of P-SOP protocol

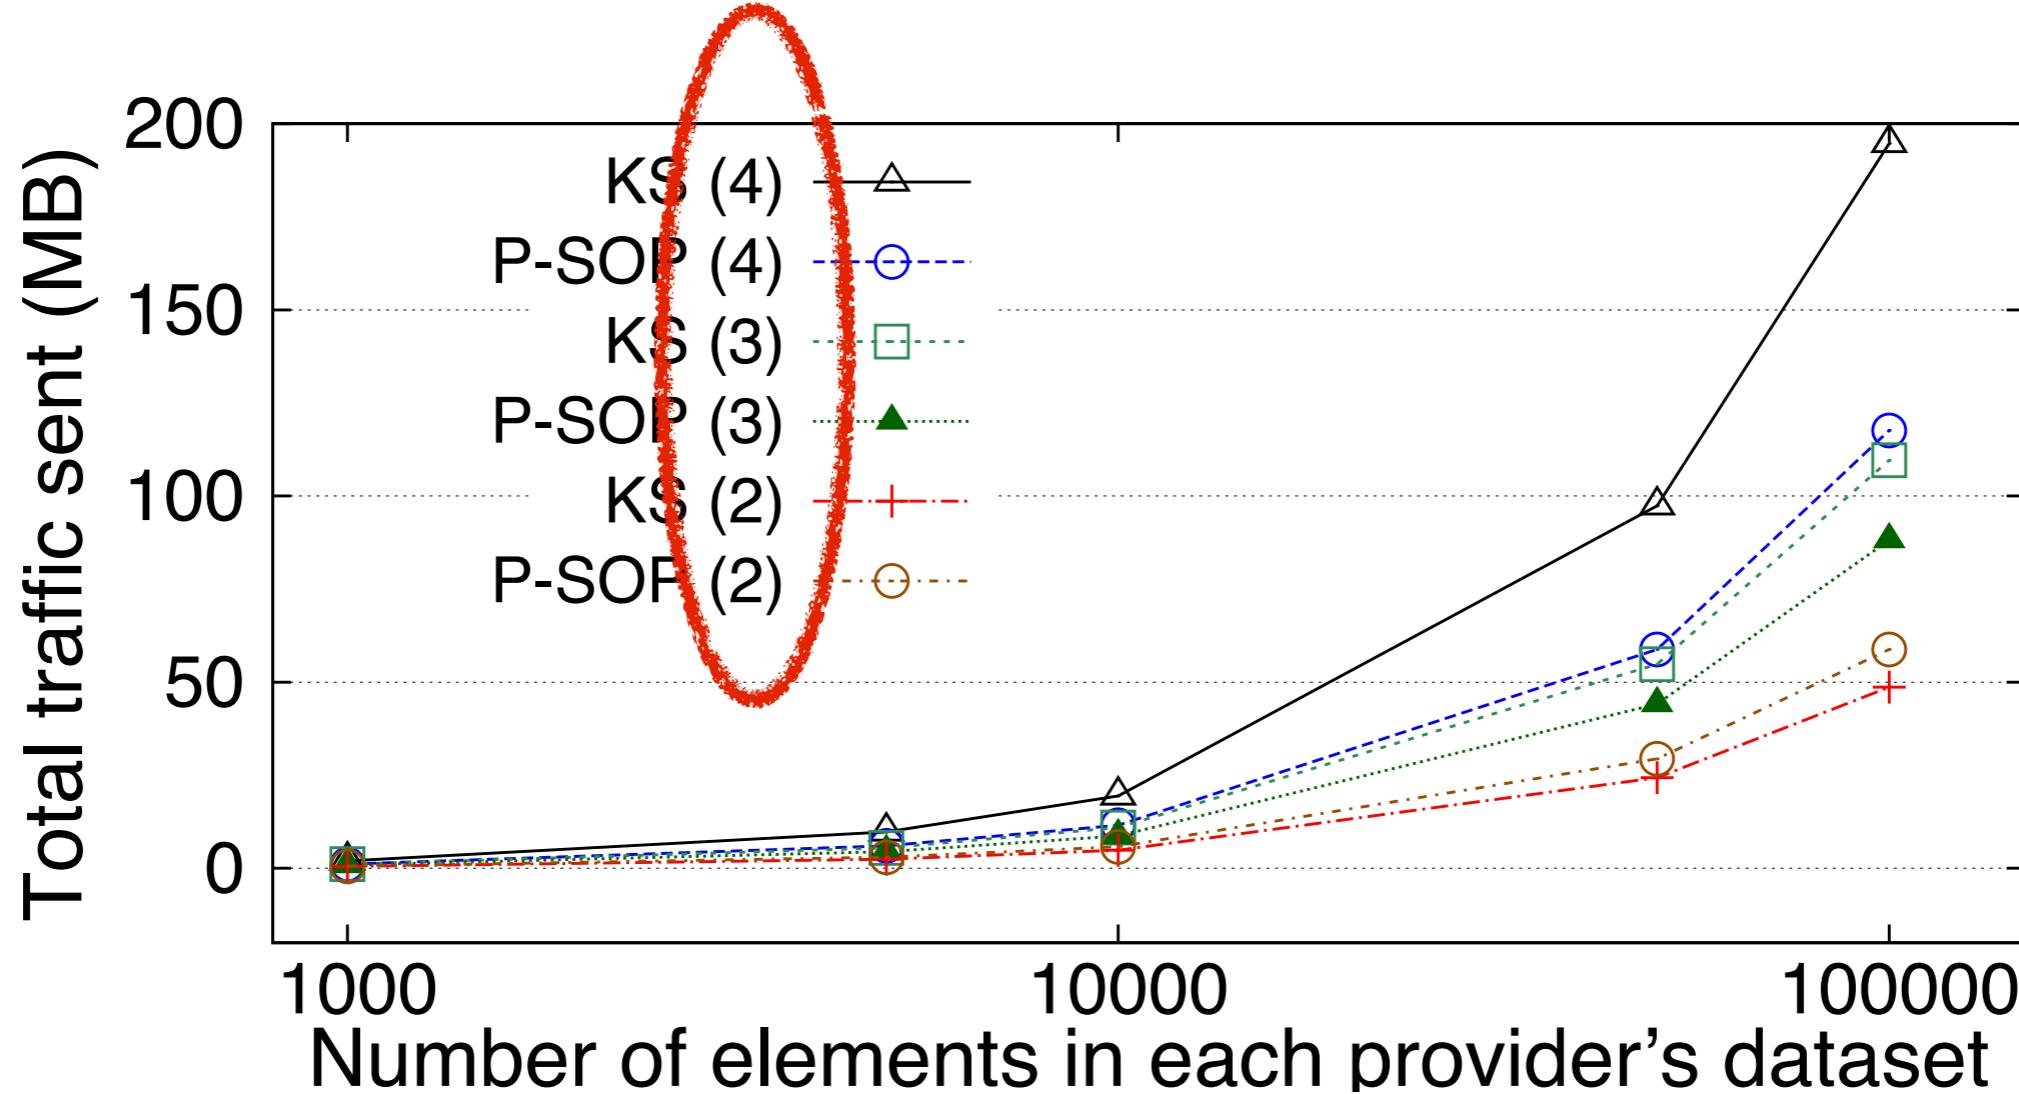
# What we evaluate?

- Kissner and Song (KS) protocol for comparison.
- Bandwidth overhead of P-SOP and KS.
- Computational overhead of P-SOP and KS.

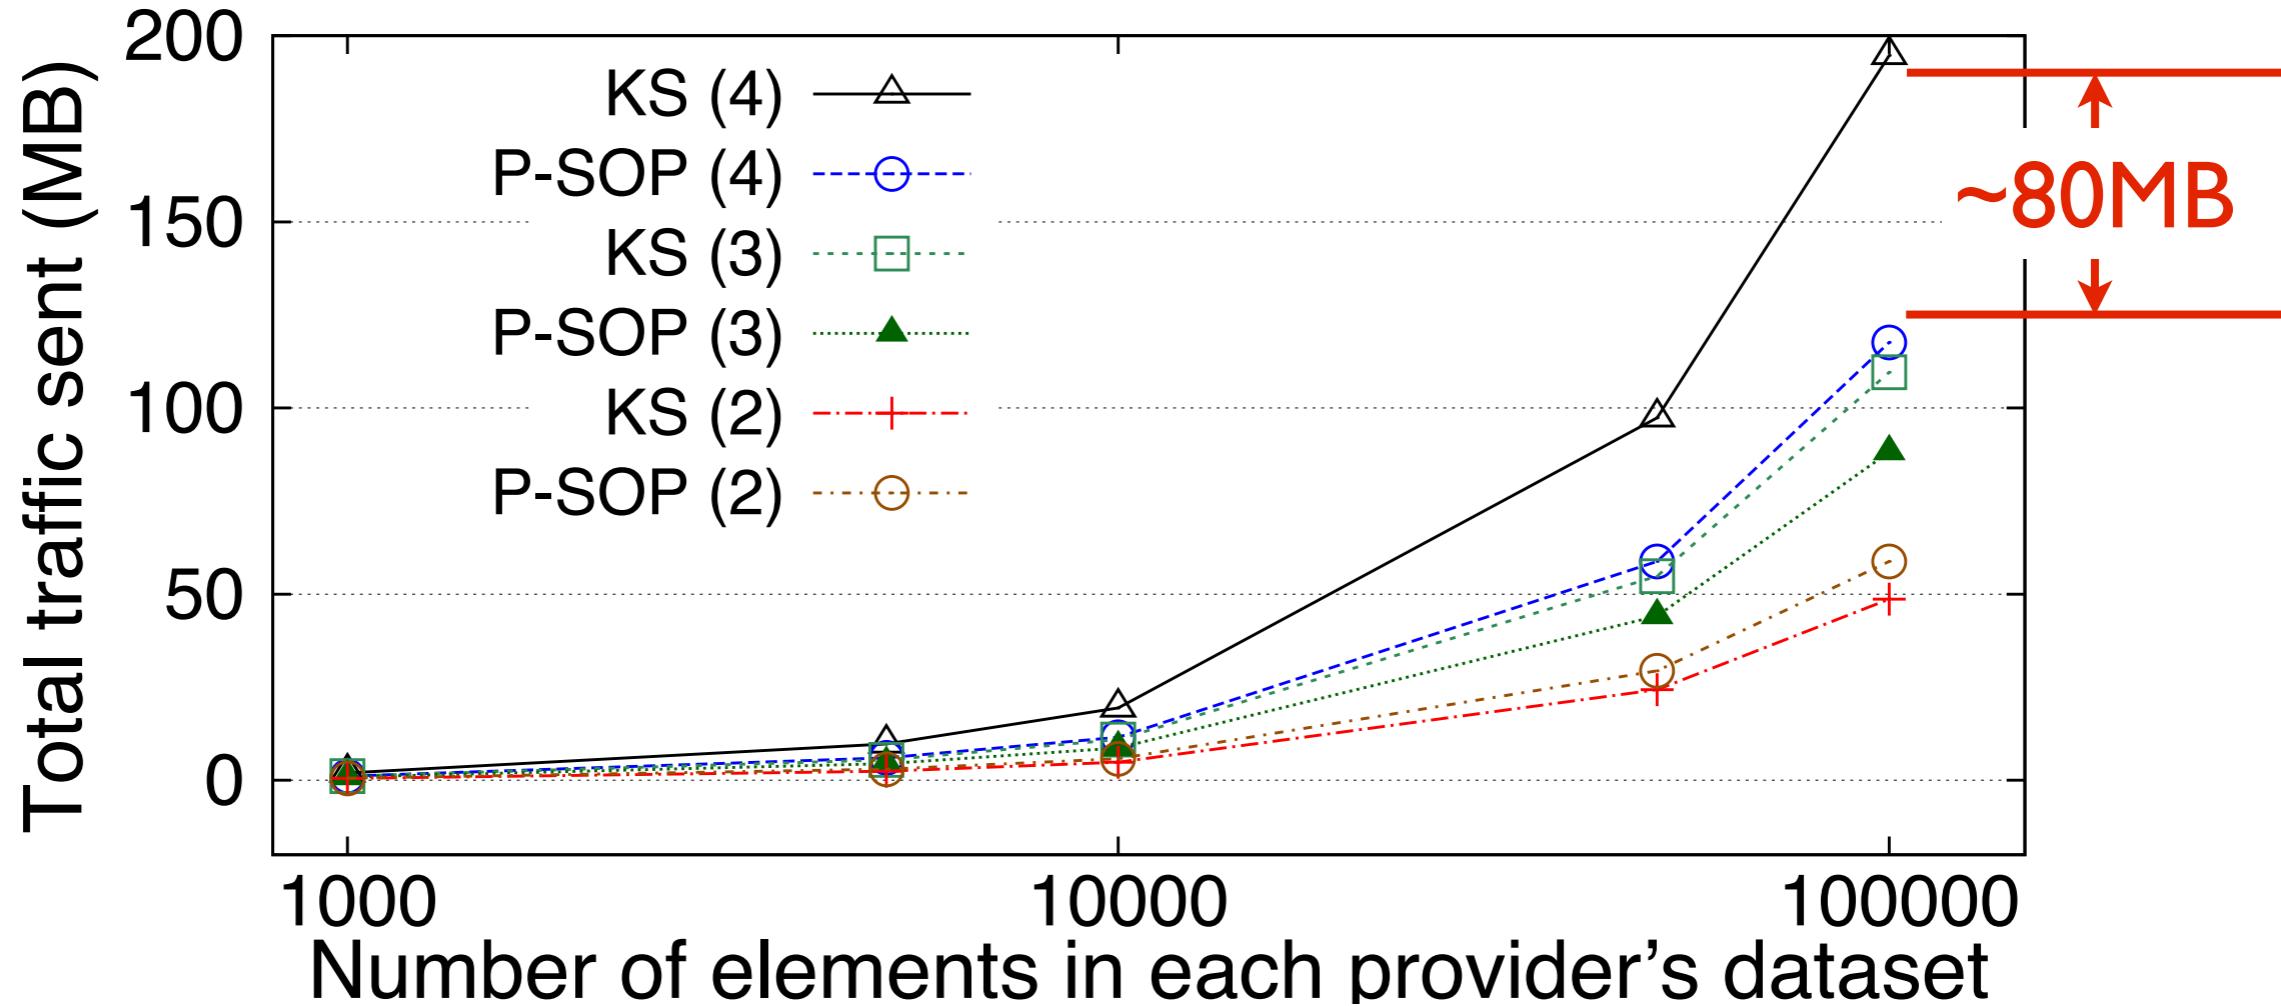
# Bandwidth Overhead



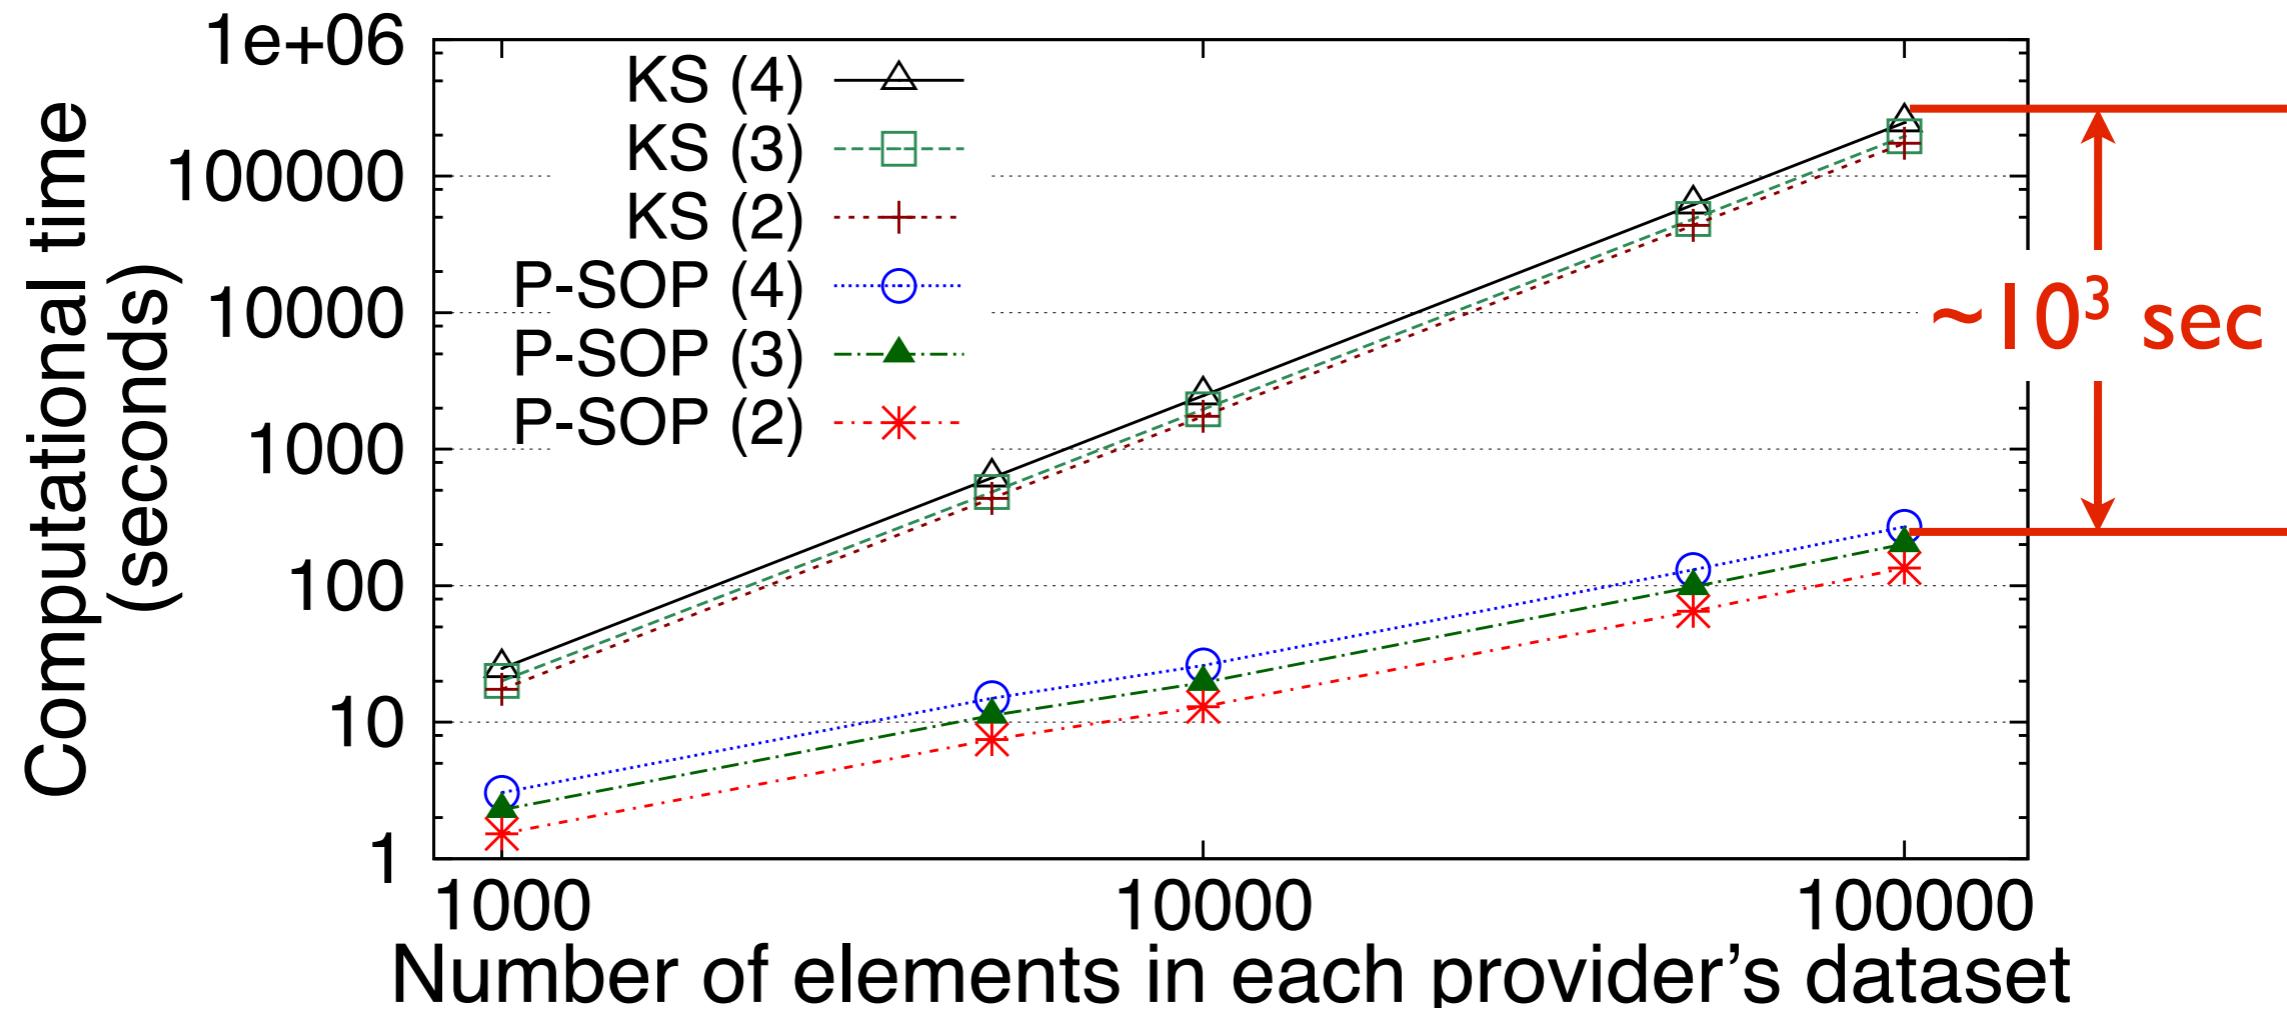
# Bandwidth Overhead



# Bandwidth Overhead



# Computational Overhead



# Conclusions

- INDaaS is a first step towards reliable clouds:
  - Dependency collections
  - Dependency representations
  - Efficient auditing
  - Private independence auditing
- We evaluated INDaaS with three realistic case studies and large-scale simulations.

# Thanks, questions?

- INDaaS: Heading off correlated failures in clouds
- Find out more at:
  - <http://dedis.cs.yale.edu/cloud/>
- We will be at the poster session tonight.