

Automated and Scalable QoS Control

- For Network Convergence

Wonho Kim (Princeton Univ.) Puneet Sharma, Jeongkeun Lee,
Sujata Banerjee, Jean Tourrilhes, Sung-Ju Lee,
and Praveen Yalagandula (HP Labs)

INM/WREN'10 Proceedings of the 2010 internet network management conference on
Research on enterprise networking

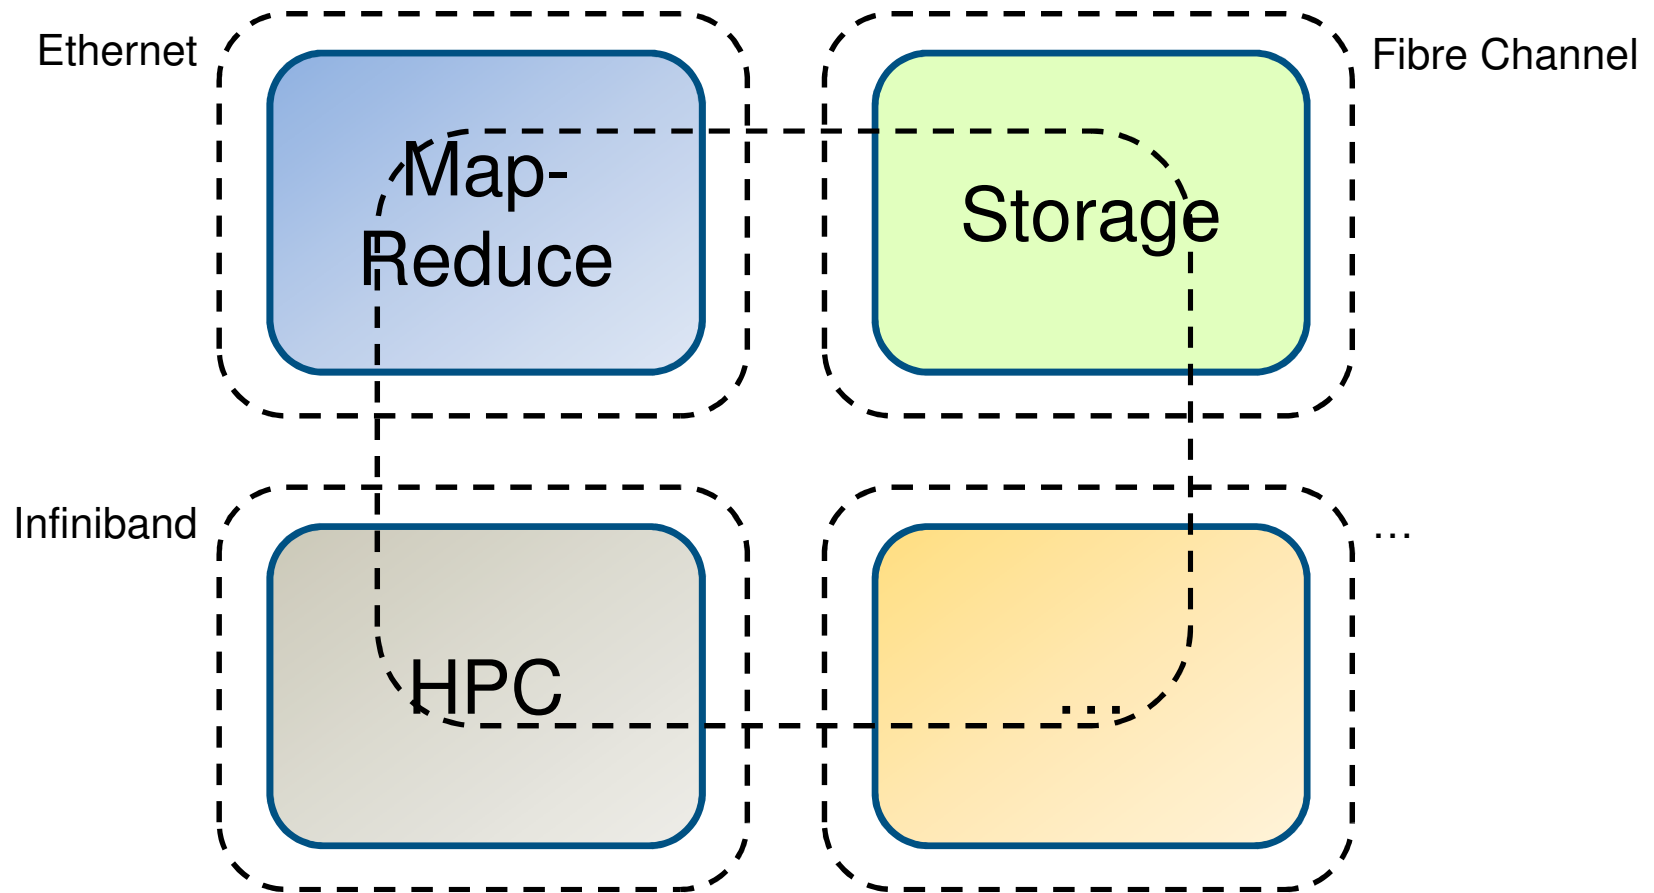




Motivation

- Why do we care about QoS control?
 - Network convergence
 - Multi-tenancy networks
- **Automated** QoS control is needed


Network convergence



Different protocols, adapters, switches, and configuration



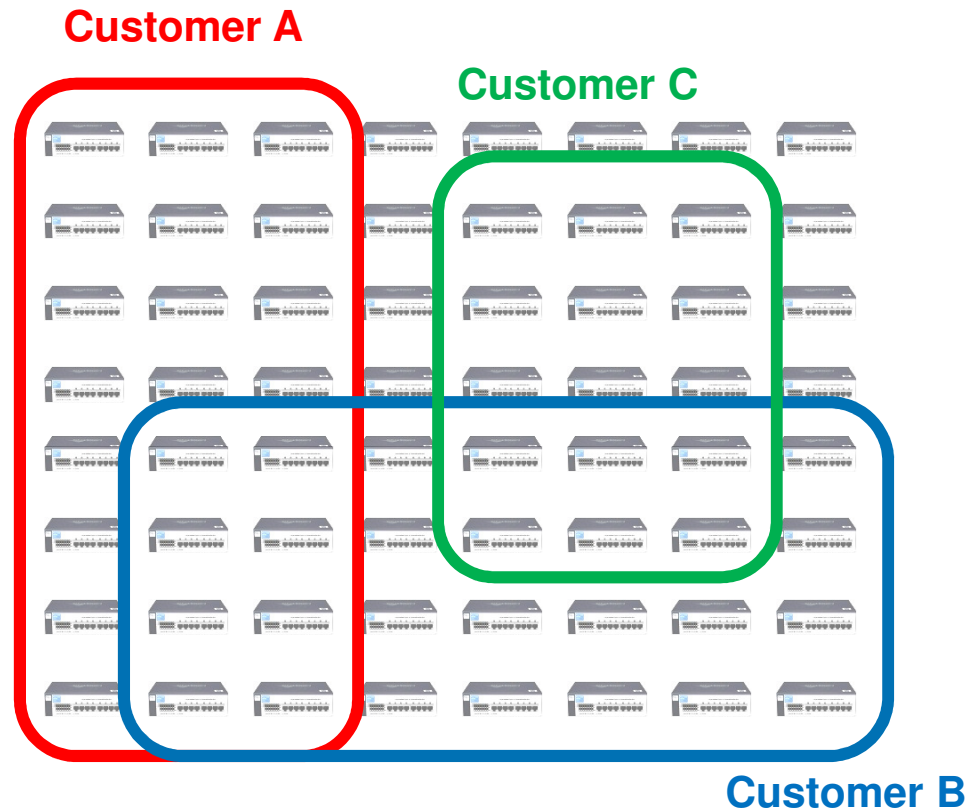
Network convergence



Fewer switches, ports, adapters, cables
Reduced power, equipment, cooling cost
Simpler topology
I/O consolidation
Unified resource management



Converged Enhanced Ethernet (CEE)
Data Center Ethernet (DCE)
Data Center Bridging (DCB)
Fibre Channel over Ethernet (FCoE)
Fibre Channel over CEE (FCoCEE)

Multi-tenancy networks





- Serve multiple customers with a single fabric
- Better utilization of network infrastructure

Performance isolation







Performance isolation

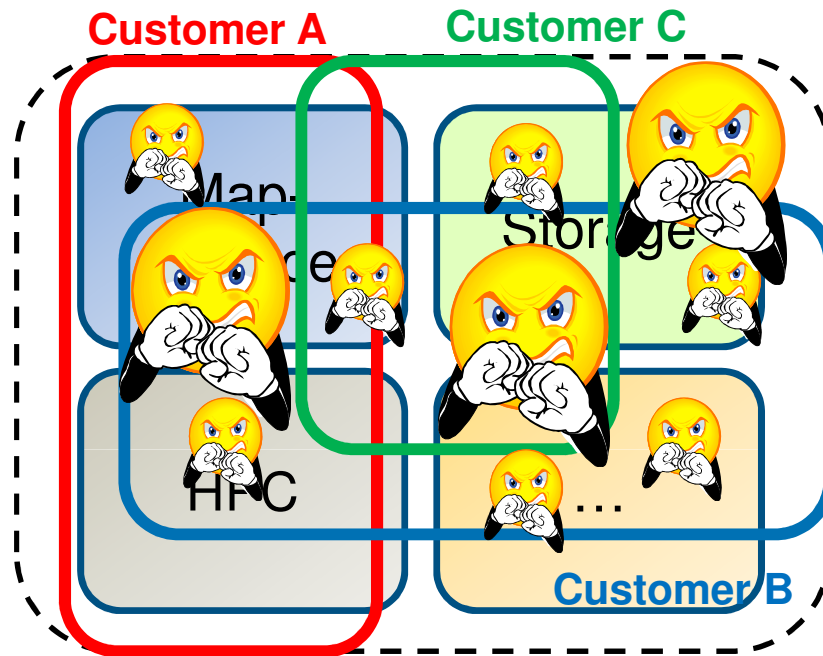




Performance isolation



Performance isolation



Virtualized Servers


Variable Workloads

Bugs, malicious attack

- Need virtual network slices
- Need fine-grained performance isolation

Goal

- Enables performance isolation with QoS control






Good news

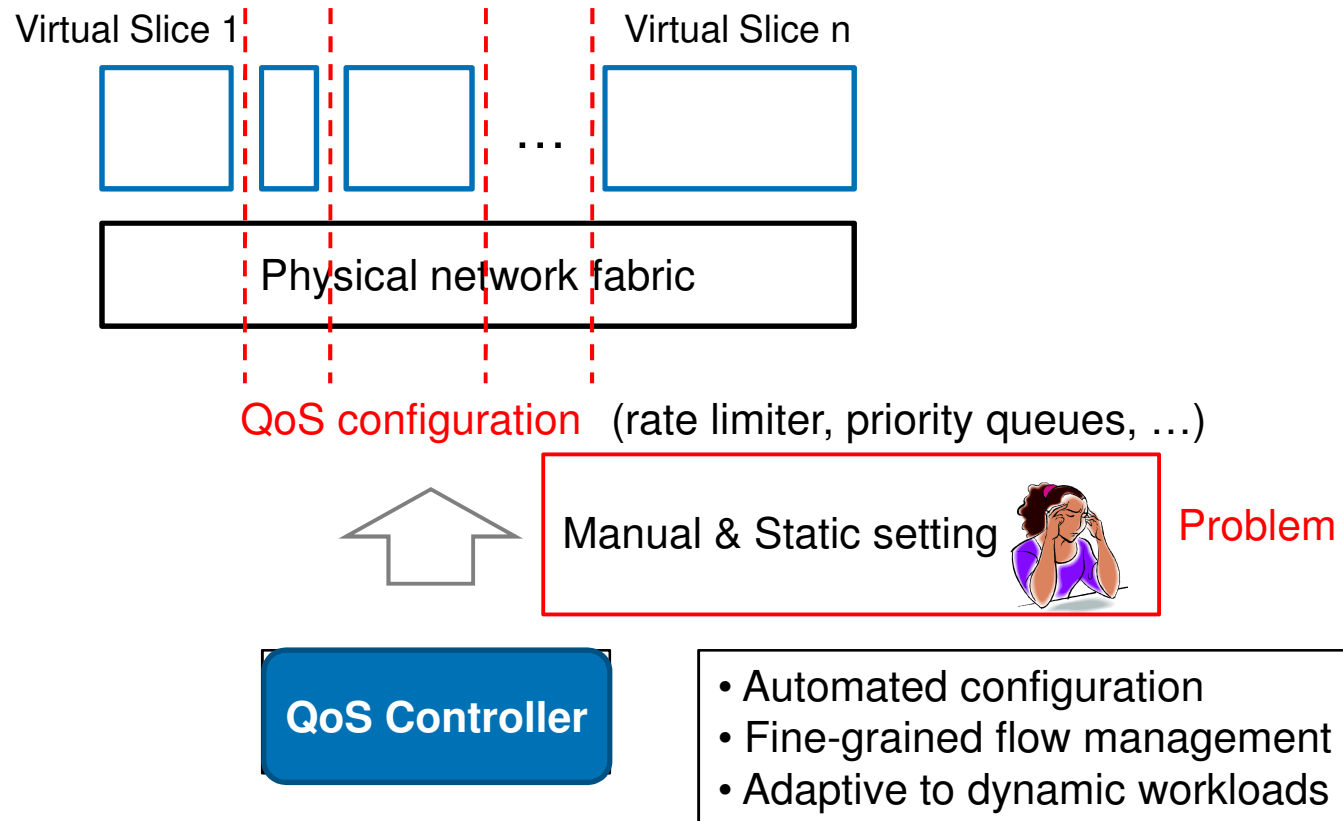
- Most commodity switches have QoS knobs
 - rate limiter
 - priority queues
 - schedulers
- Single network domain
 - datacenters, enterprise networks, ...
 - free from Layer-8 issues (billing, collaborations, ...)
 - fine-grained control becomes feasible

Challenges

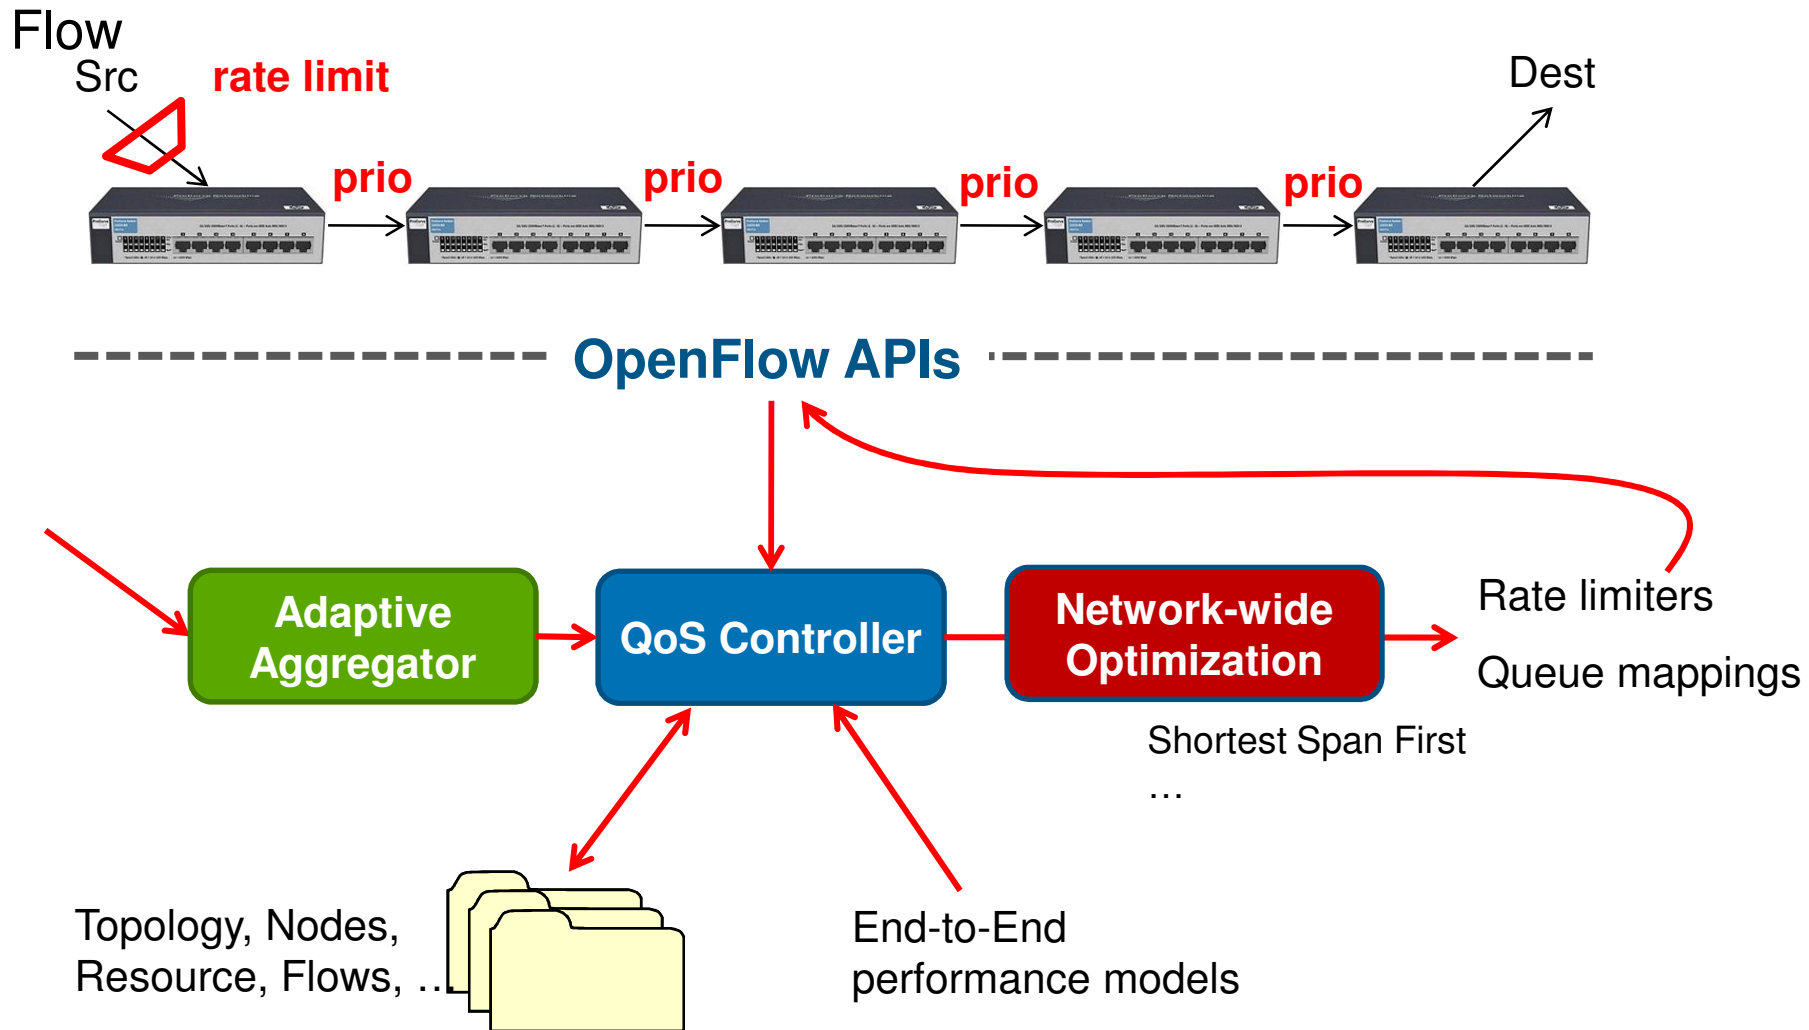
- Coarse-grained QoS knobs
 - designed for distributed management
 - class-based
 - no e2e performance
- Manual configuration
 - no standards for classifiers
 - error-prone
 - static (not adaptive)



Our Solution: OpenFlow QoS Controller



Overview of OpenFlow QoS controller

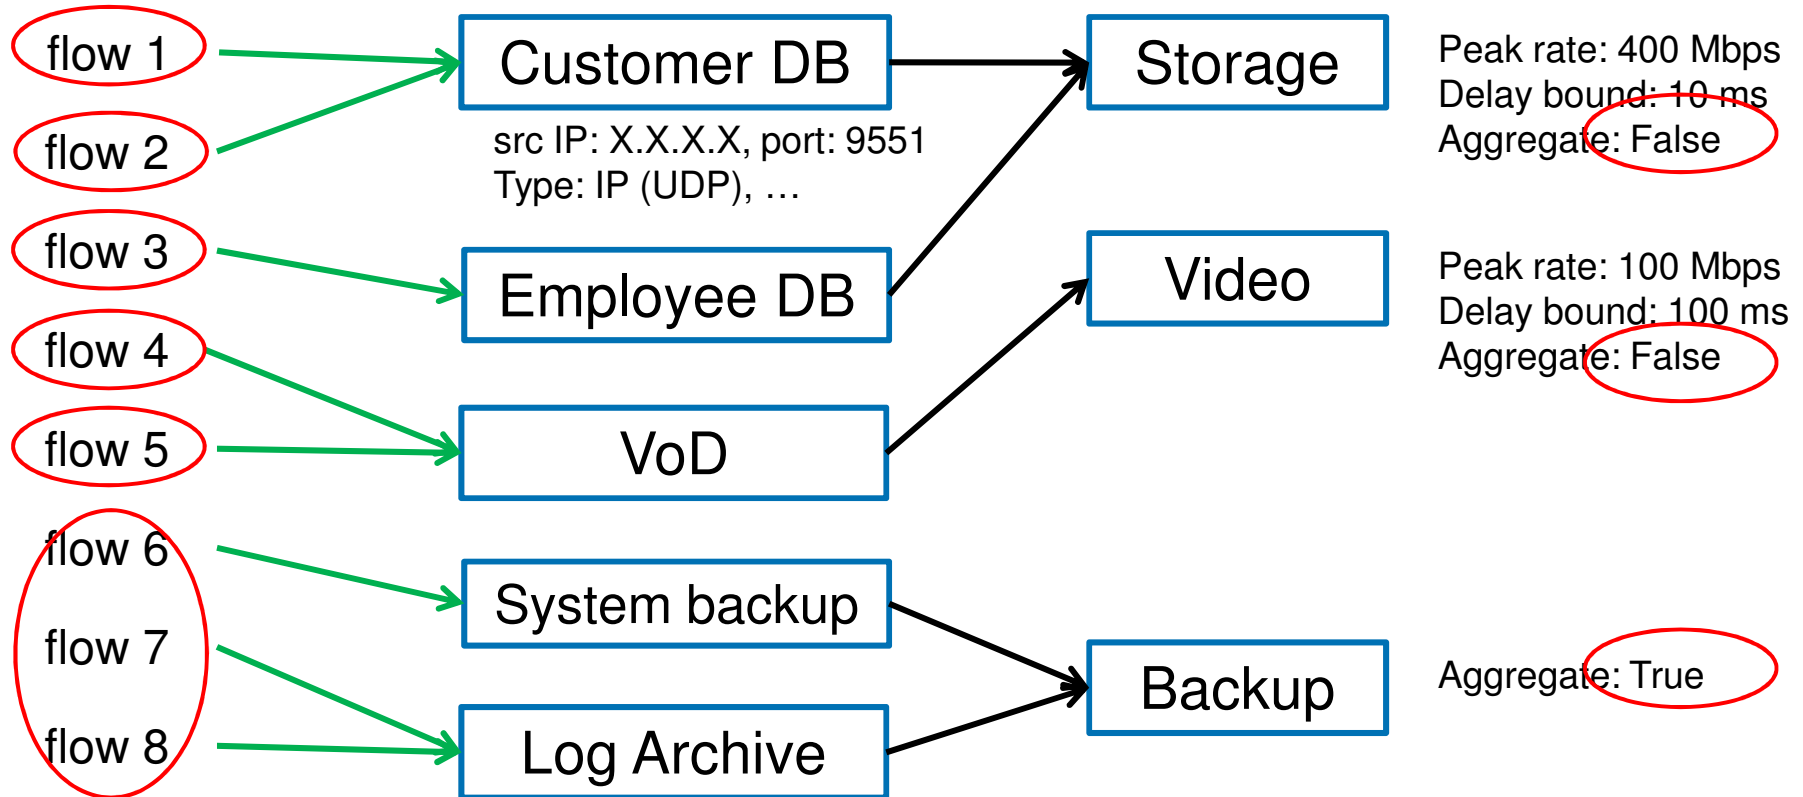




Adaptive aggregation


Flow specs

Slice specs






Available QoS Knobs (Priority queue)





Available QoS Knobs (Rate limiter)



OpenFlow QoS APIs

Rate limiter

**Priority
queue mapping**



- Extension of OpenFlow specification
- Expose QoS capability in switches



OpenFlow QoS APIs

- With OpenFlow flow control
 - fine-grained control of flows
 - automated flow management
- With OpenFlow QoS APIs
 - uniform control of QoS knobs
 - configure QoS for individual (or aggregate) flows



Admission Control


- Input
 - new flow arrival event
 - performance requirements (peak rate, e2e delay)
 - database for the current network state
 - end-to-end performance model
- Output
 - admission control result (accept/reject)
 - priority queue assignment, rate limiter settings
 - path selection



Admission Control

- Two conditions should be satisfied
 - satisfy f 's performance requirement
 - not violate existing flows in the networks

Difficulties in queue assignment



We should consider interactions between


- flows in a switch
- flows in multiple switches



Admission control heuristic

- Goal
 - increase the ratio of admitted flows
 - lower the complexities in queue allocation
- Shortest Span First (SSF)
- Basic ideas
 - estimate affordable options for a flow
 - try first switches more likely to reject flow


Highest level & Lowest level



- Highest level: not violate existing flows
- Lowest level: not violate the new flow
- Span: available options for f




Shortest Span First (SSF)



- **Step 1**: compute highest & lowest levels independently




Shortest Span First (SSF)



- **Step 2:** sort switches in order of the span




Shortest Span First (SSF)



- **Step 2:** sort switches in order of the span



Shortest Span First (SSF)



- **Step 3:** try highest level at each hop
 - try first a switch more likely to reject flow

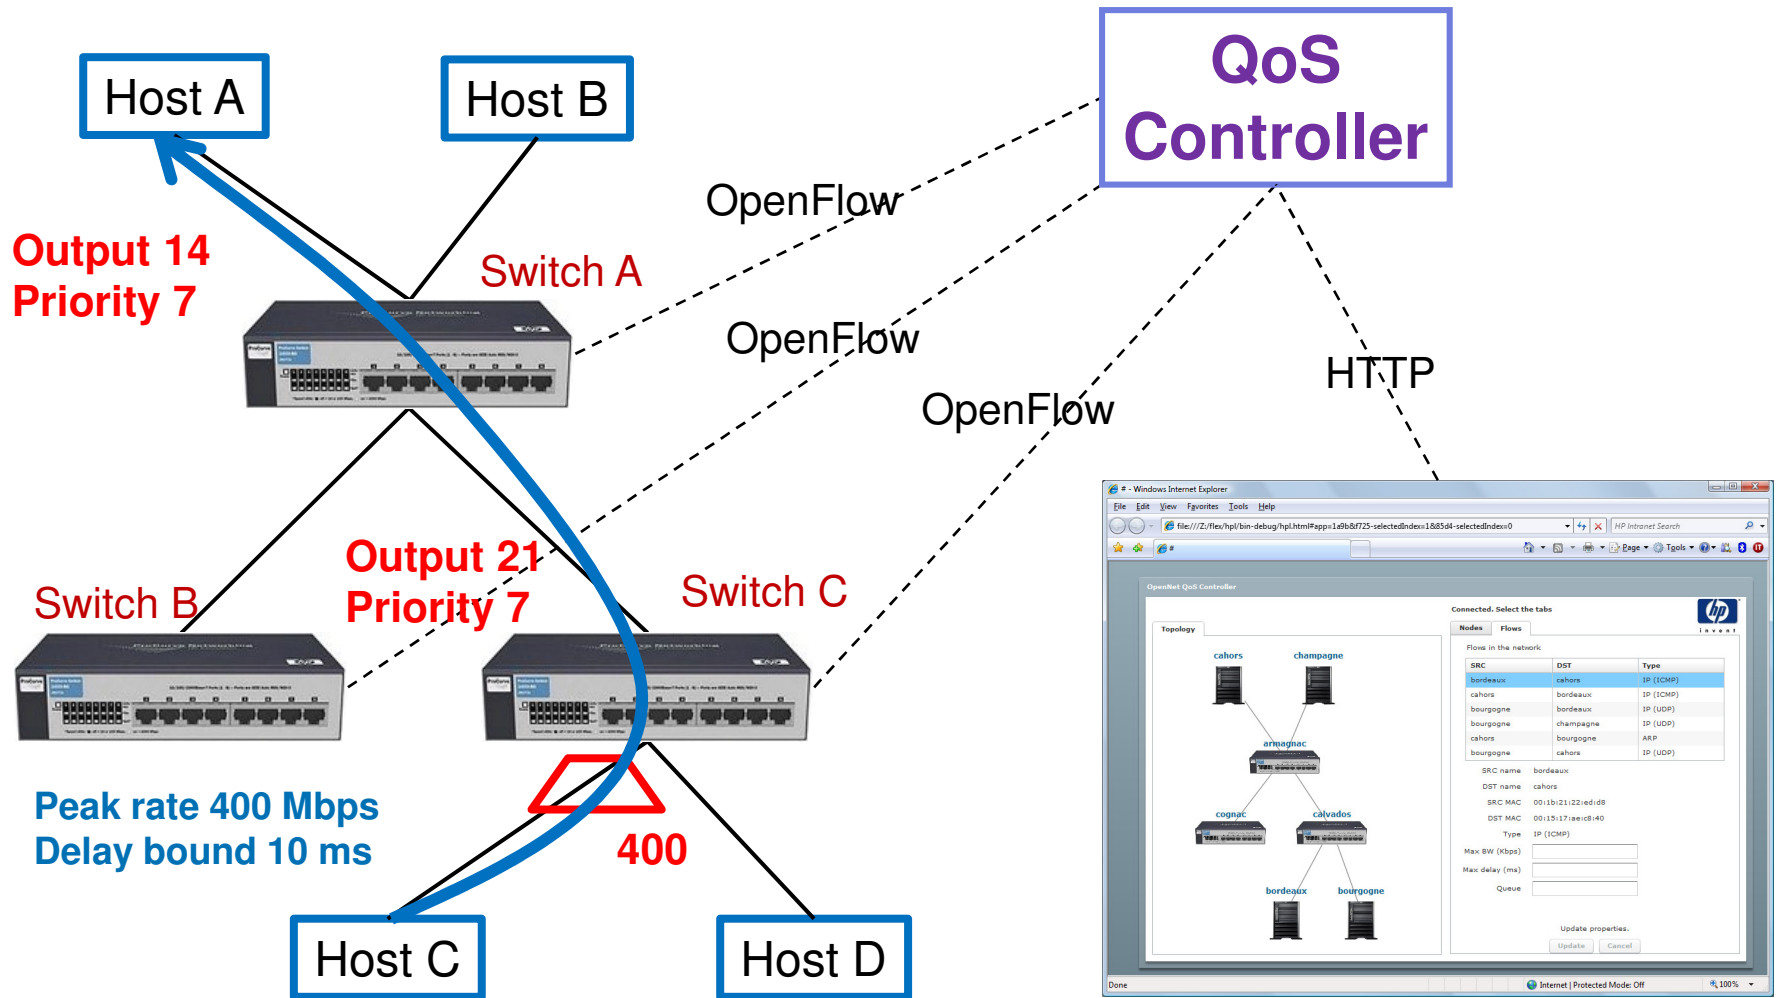


Implementataion

- QoS APIs implemented on
 - hardware switch (HP ProCurve 5406zl)
 - software switch (Open vSwitch)
- QoS Controller implemented on top of NOX
 - open-source OpenFlow controller
 - <http://noxrepo.org>
- QoS Controller web interface



Prototype






Evaluation

- Traffic generation
 - generate 3 guaranteed flows from emulated services (UDP)
 - generate cross traffic (UDP, TCP)
- Disable/Enable QoS controller
- Measured throughput and packet loss in testbeds

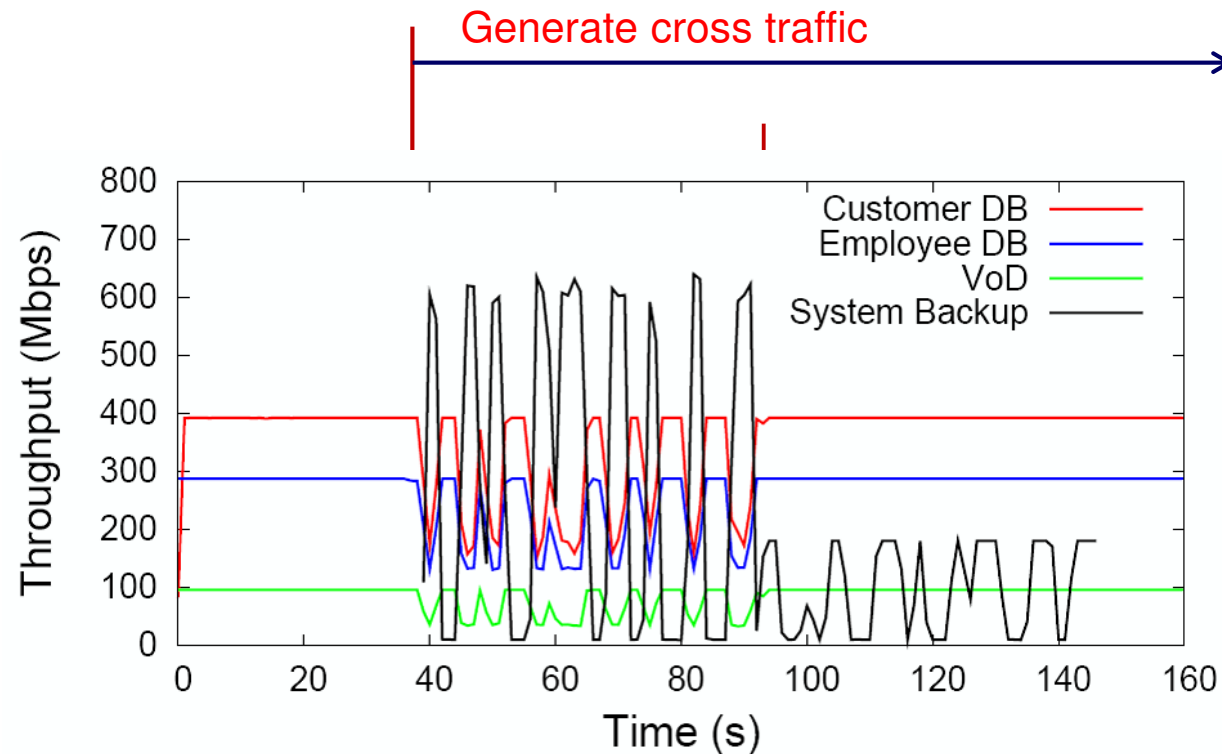
Throughput with UDP cross traffic



Flow name	Route (queue assignment)
Customer DB	H3 – S3(8) – S1(8) – H1
Employee DB	H4 – S3(8) – S1(8) – H2
VoD	H3 – S3(7) – S1(7) – H1
System Backup	H4 – S3(1) – S1(1) – H2

QoS controller protects guaranteed flows in congestion


Throughput with UDP cross traffic



QoS controller protects guaranteed flows in congestion



Packet loss with TCP cross traffic



QoS control is needed even when most traffic in network is TCP



Future works

- Evaluations
 - effectiveness of admission control heuristics (ratio of admitted flows)
 - compare with offline optimal assignment
 - simulations on a variety of datacenter networks (e.g., Hierarchical, FatTree, ...)
- Deployment
 - extend deployment to large networks
 - test with mixture of services



Conclusion

- Single integrated network fabric is desirable
- We need fine-grained automated QoS control
- Contributions
 - Design & Implement OpenFlow QoS APIs
 - QoS controller: automated QoS control for network slicing



Thank you