

ANCS 2008


Compact Architecture for High-Throughput Regular Expression Matching on FPGA

Yi-Hua E. Yang, Weirong Jiang and Viktor K. Prasanna
EE, USC



Kyle Wang, NSLAB, Tsinghua University

[Outline



Introduction

Background

Basic Architecture on FPGA

Architectural Optimizations

Performance Evaluation

Conclusion and Future Work

Introduction

■ Basic regular expression example (**What**)

- $b^*c \cdot (a|b)^*[ac]\#$

■ RE-NFA & RE-DFA (**How**)

- regular expression matching (**RE**M) architecture based on (**N**on-)Deterministic Finite Automaton

■ Bottleneck for DPI in IDS (**Why**)

- large number of patterns to scan for and increasing bandwidth of network traffic

■ An improved RE-NFA approach on FPGA

- n -state m -character input REME (REM Engine) with $O(n \times \log_2 m)$ constructing time and $O(n \times m)$ logic units


[Background]

- Challenges to perform REM on hardware
 - processing large numbers of patterns in parallel
 - obtaining high *concurrent throughput*
 - throughput of the input stream processed concurrently by N REMEs
 - RE size/complexity, amount of resources, achieved clock frequency
- Prior work upon implementation of RE-NFA on hardware
 - “The Compilation of Regular expressions into Integrated Circuits” (*Floyd et al, J.ACM 1982, Cited by 85*)
 - “Fast Regular Expression Matching Using FPGAs” (*Sidhu et al, FCCM 2001, Cited by 272*)
 - optimizations (e.g. utilizing BRAM).....

Basic Architecture on FPGA

From Regular Expression to NFA(1/3)

- Parse the regular expression into a tree structure




The parse tree for regular expression $b^*c(a|b)^*[ac]\#$


Basic Architecture on FPGA

From Regular Expression to NFA(2/3)


■ Modified McNaughton-Yamada construction




(a) [c]




(b) $r1 \cdot r2$




(c) $r1 | r2$



(d) $r3^*$



(c) $r1 | r2$




(d) $r3^*$

dashed ellipses indicate other units
dashed lines indicate ϵ -transitions

Basic Architecture on FPGA

From Regular Expression to NFA(3/3)

- Constructing the “modular” NFA using the modified McNaughton-Yamada rules




The RE-NFA for regular expression $b^*c(a|b)^*[ac]\#$

Basic Architecture on FPGA

From RE-NFA to HDL(1/2)

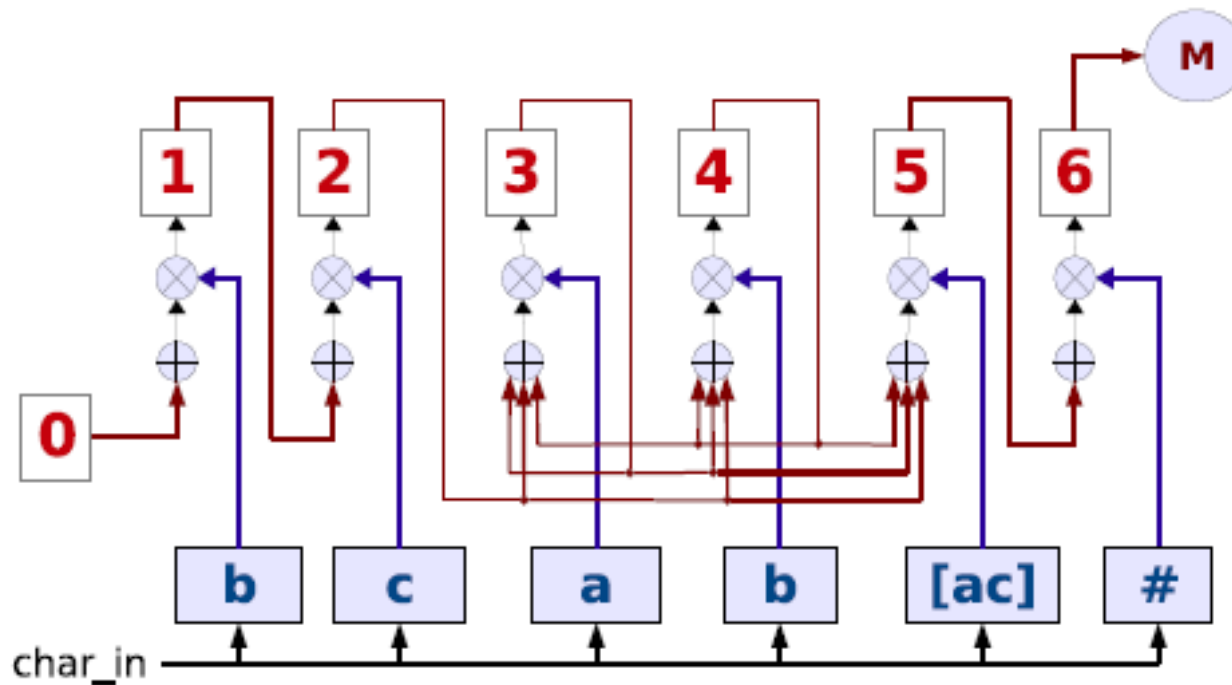
- Mapping state directly into logic module



Basic Architecture on FPGA

From RE-NFA to HDL(2/2)

- Mapping NFA directly to HDL



Matching circuit for regular expression $b^*c(a|b)^*[ac]\#$

Basic Architecture on FPGA

BRAM-based Character Classification(1/2)

■ Character classification unit

○ 256 bits for specifying all 8-bit characters

- e.g. character **a** can be classified by

$\overbrace{000 \dots 0100 \dots 0000}^{91\#}$
256#

- max $256 \times n$ bits BRAM usage for n -state RE-NFA
- Less if two states use the same character


○ cons.

- redundancy in BRAM
- max 2^{256} possible **character classes** (e.g. simple [ac])
- optimization is necessary

Basic Architecture on FPGA

BRAM-based Character Classification(2/2)

- Dual-4bit selector
 - 2*16 bits
 - only 'and gate' delay
 - why not use?
 - logic slices vs. BRAM
 - clock frequency influence
- For character 'a'
 - 0110 0001
 - if ' ' ?




[Architectural Optimizations]

- Three available optimizations to improve the basic design above
 - Multi-Character Input Matching
 - Centralized Character Classification
 - Staging and Pipelining
- These techniques are unique to our design and take advantage of the modularity of the proposed architecture.
 - Scalable pattern matching for high speed networks(multi-character)
 - Regular Expression Matching for Reconfigurable Packet Inspection(centralized)
 - Assisting Network Intrusion Detection with Reconfigurable Hardware(pipeline)

Architectural Optimizations

Multi-Character Input Matching(1/2)

- 2-character input (2-Stride Accelerating)

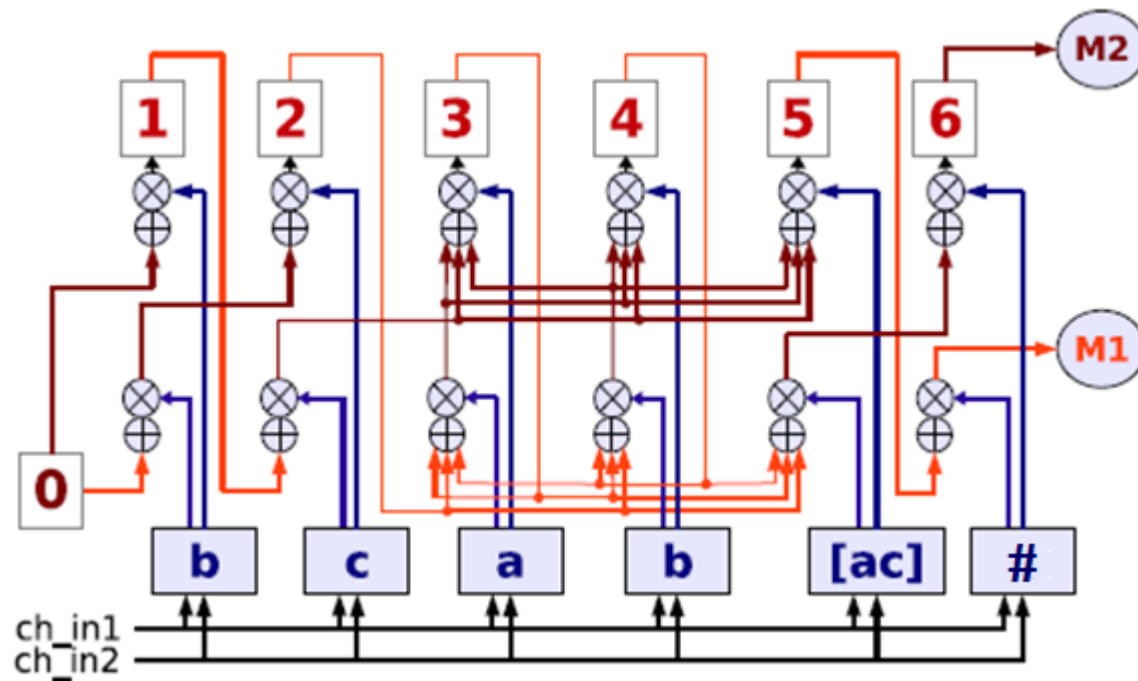


Multiple outputs may be at 'and gate' or 'state register'

Architectural Optimizations

Multi-Character Input Matching(2/2)

- $O(n \times \log_2 m)$ time Alg. for m -character input



2-input matching circuit for regular expression $b^*c(a|b)^*[ac]\#$

Architectural Optimizations

Centralized Character Classification(1/2)

- Call function to examine and compare each state's character class to the character class entries (CCE) in BRAM so far
 - **if new:** a new 256-bit entry is added into BRAM
 - **if old:** proper connection is made from the BRAM output of the previous CCE to the input of the current state
 - $O(n \times \omega)$ time complexity
 - ω is the number of distinct character classes among the n state, in the worst case, ω could be linear in n

Centralized Character Classification(2/2)

-
- Character Classes for all REMEs



Architectural Optimizations

Staging and Pipelining

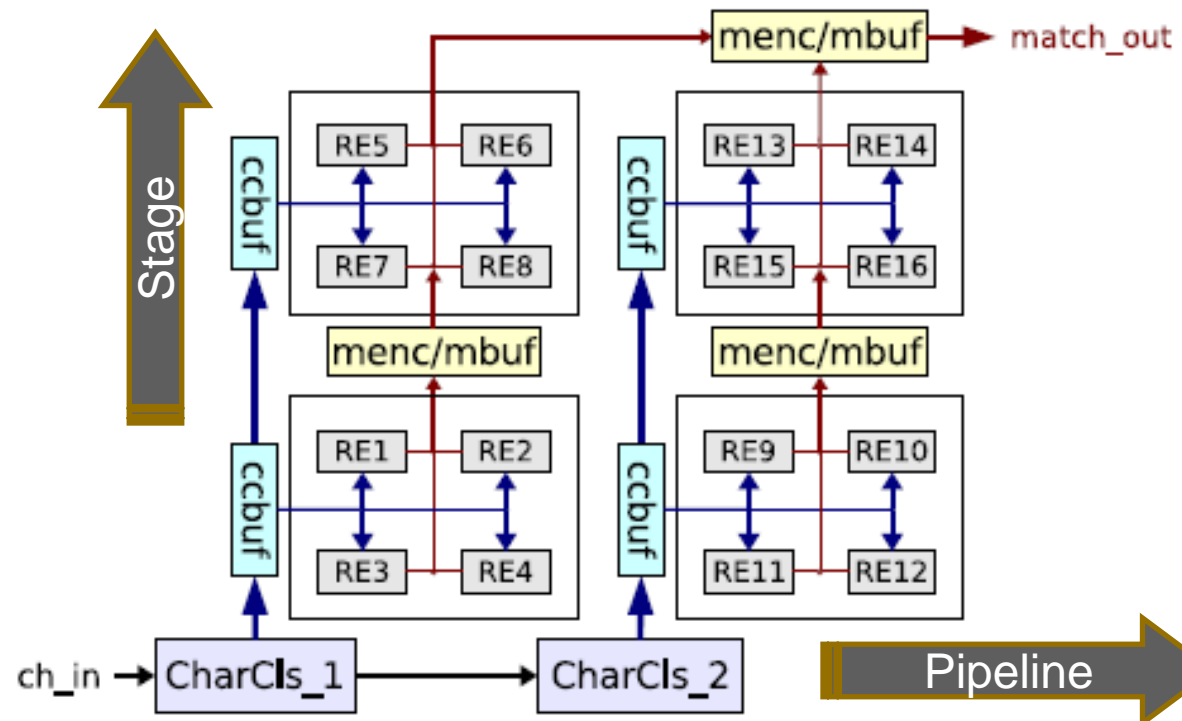
- Against the decline in achievable clock frequency with larger numbers of REMEs

fixed delay:

p (pipeline)

$+ s$ (stage)

clock cycles



Performance Evaluation

Hardware Complexity of Regular Expression

- Metrics to quantify the complexity of RE
 - **area requirement**
 - **State count:** Total number of states needed by REME
 - **logic complexity**
 - **State fan-in:** maximum number of states that can immediately transition to any state in the REM
 - **State fan-out:** maximum number of states to which any state in the REME can immediately transition
 - **routing complexity**
 - **Loop size:** total number of transitions within a loop of state transitions
 - **Branch-size delta:** difference in number of transitions between two state transition paths with the same first and final states

[Performance Evaluation]

Implemented Regular Expressions(1/3)

- Patterns selected for RE-NFA evaluation
 - avoid patterns too short or too simple
 - can be dealt with pattern-level optimizations
 - count identical patterns in different rules as one
 - should not be used to inflate the REME numbers
 - avoid patterns with long quantified repetitions
 - can be easily dealt with shift-register
 - avoid patterns requiring backreference
 - requires capability beyond regular languages
 - is beyond the scope of this paper

Performance Evaluation

Implemented Regular Expressions(2/3)

- Snort rule categories (partially implemented)

Snort rule cat. implemented	# of pat.	# of states	# of states/pat.
backdoor	154	3648	23.7
chat/ftp	18	163	9.1
deleted	23	607	26.4
smtp/pop2/pop3	22	306	13.9
spyware-put	446	11720	26.3
web-misc/web-php	51	1082	21.2
(others)	46	1052	22.9
760-REME	760	18578	24.4
523-REME	523	13379	25.6
267-REME	267	6551	24.5



Performance Evaluation

Implemented Regular Expressions(3/3)

- State fan-in and fan-out values


value v	1	2	3	4	5	6	7	8	9
# of state fan-in = v	52	656	23	5	14	3	3	4	0
# of state fan-out = v	45	663	31	7	7	2	0	2	3

- Although **a few** states reach 8 or 9, the **slowest** REME determines the overall clock frequency when implementing in the same clock domain



Performance Evaluation

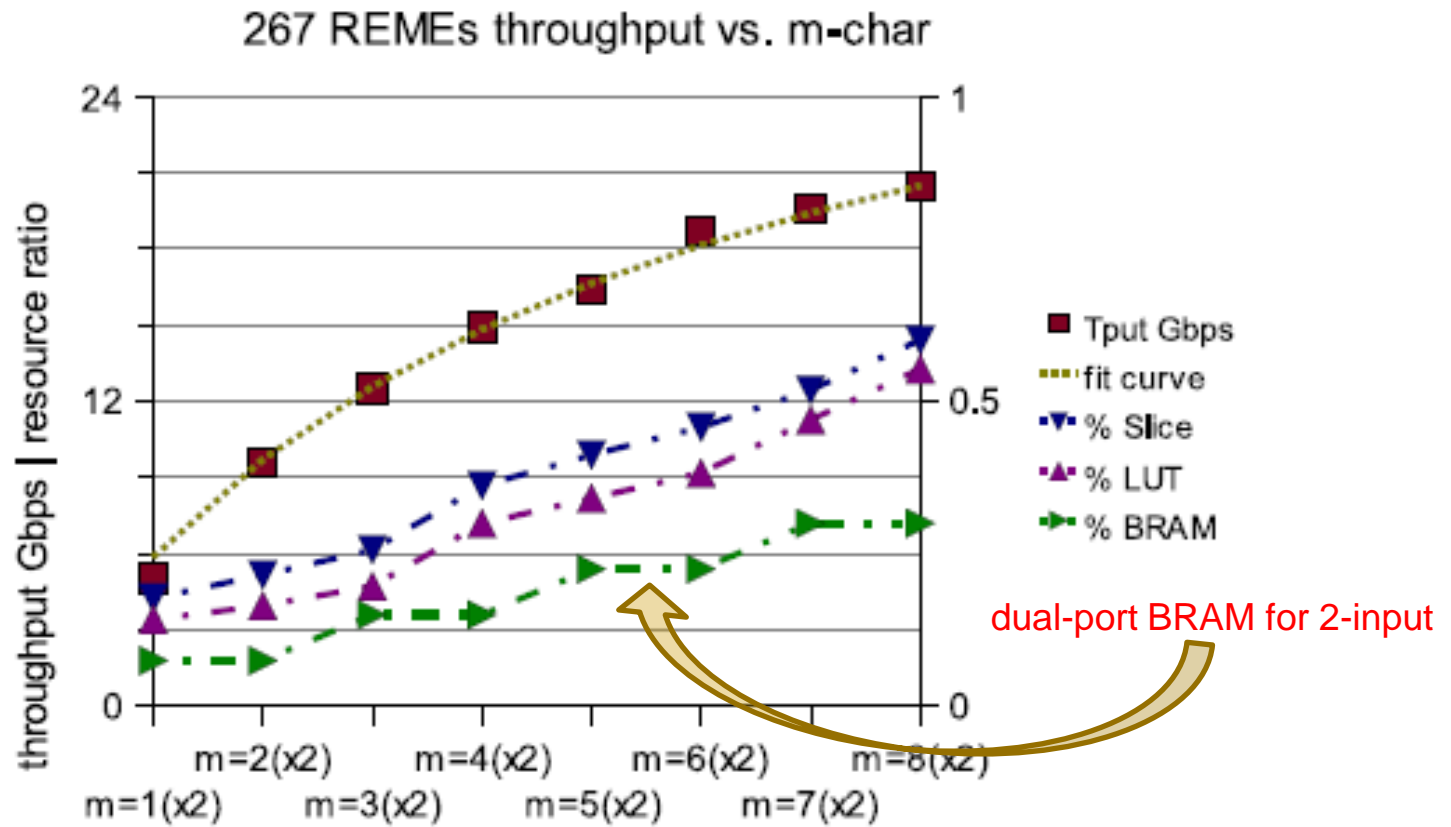
Implementations Results(1/3)



Throughput scaling of 760 REMEs on Xilinx Virtex 4 LX-100-12

Performance Evaluation


Implementations Results(2/3)



Throughput scaling of 267 REMEs on Xilinx Virtex 4 LX-100-12

Performance Evaluation

Implementations Results(3/3)



Clock rate and LUT usage of 267 REMEs on **Xilinx Virtex 4 LX-100-12**

Performance Evaluation

Performance Comparison

■ Throughput efficiency

- concurrent throughput of the circuit divided by the number of LUTs the circuit uses per state
- a metric to evaluate performance in various conditions

	# non-meta char.	Multi-char. m	Tput. (Gbps)	# LUT per <i>state</i>	Tput. efficiency (Gbps* <i>state</i> /LUT)
760-REME-2	~ 20k	2	4.8	1.27	3.8
523-REME-2	~ 15k	2	4.88	1.24	3.9
523-REME-4	~ 15k	4	7.46	2.2	3.4
267-REME-4	~ 8k	4	7.5	2.25	3.3
Bispo <i>et al.</i> [5]	19580	1	2.9	1.28	2.3
Clark <i>et al.</i> -1 [6]	17537	1	2.0	1.7	1.9
Clark <i>et al.</i> -4 [6]	17537	4	7.0	3.1	2.3
Mitra <i>et al.</i> [13]	N.A.	1	12.8/16	~ 2.3	~ 0.35

NOTE: The FPGA platform is not the same, but the throughput efficiency can still illustrate something

[Conclusion and Future Work]

■ Conclusion

- compact and easy-mapping architecture
- stack character spatially to match multiple input
- utilize block memory (BRAM) available
- 2D staging and pipelining organization

■ Future Work

- group REMEs into stages intelligently (Fang Yu)
 - exploit pattern-level properties such as the common prefix extraction
- separate simple and complex REs into different clock domains
 - supposedly a pipeline of simple REMEs can be clocked higher than complex REMEs

[Thanks for your attention!]

Let's Talk.